

Advanced Microservices Caching Patterns

Natan Silnitsky

Backend Infra TL,
Wix.com

natansil.com [twitter@NSilnitsky](https://twitter.com/NSilnitsky) [linkedin/natansilnitsky](https://linkedin.com/natansilnitsky) github.com/natansil


Wix has a huge scale of traffic. More than 500 billion HTTP requests and more than 1.5 billion Kafka business events per day. This talk goes through 4 Caching Patterns that are used by Wix's 1500 microservices in order to provide the best experience for Wix users along with saving costs and increasing availability. A cache will reduce latency, by avoiding the need of a costly query to a DB, a HTTP request to a Wix servicer, or a 3rd-party service. It will reduce the needed scale to service these costly requests. It will also improve reliability, by making sure some data can be returned even if aforementioned DB or 3rd-party service are currently unavailable. The patterns include:

- Configuration Data Cache - persisted locally or to S3
- HTTP Reverse Proxy Caching - using Varnish Cache
- Kafka topic based 0-latency Cache - utilizing compact logs
- (Dynamo)DB+CDC based Cache and more - for unlimited capacity with continuously updating LRU cache on top.

Each pattern is optimal for other use cases, but all allow to reduce costs and gain performance and resilience.



>200 Million registered users from 190 countries



Unreliable Requests

Stores
Checkout
Service




WIX APP
MARKET

1. Get app identification
2. Start




Request Overload

Stores
Catalog
Service



Store
Catalog



- No Cache** - Increase in
- loading and response
 - Network costs
 - DB failure

- With Cache** -
- reduce latency
 - reduce the needed scale
 - improve reliability

When & how to Cache

Examples from 2000 microservices of various use cases

#1
high risk/cost
of network
failure

Like caching
external critical
configuration
data.

#2
To improve
average latency
for data layer
access

DynamoDB +
CDC +
LRU cache.

#3
Some very high
external traffic
cases

Caches are very
important -
cache before
the service.

#4
When NOT to -
young products


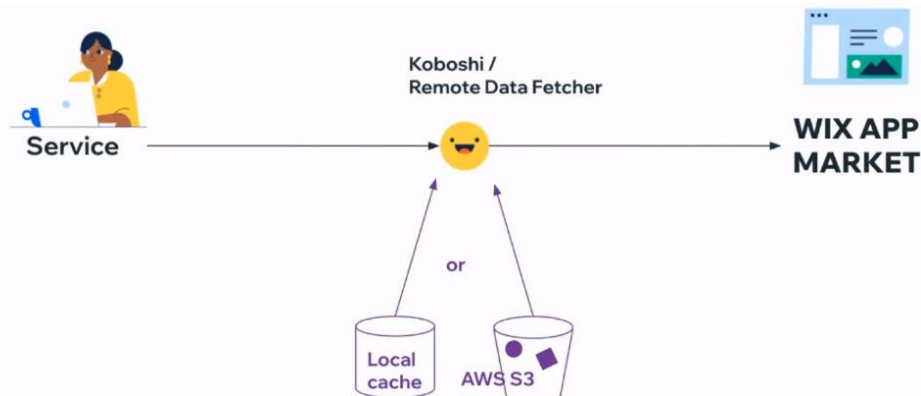

Don't cache
prematurely

Agenda

3 Caching Patterns:

1. S3 backed static cache
2. DynamoDB+CDC based cache
3. HTTP Reverse Proxy cache

for Unreliable Requests



Cache Type

Read-through cache
Before the 3rd party service

Cache Size

Data has to
fit into memory

Best For

External source
of static configuration

Not For

Dynamically
updating data



Next

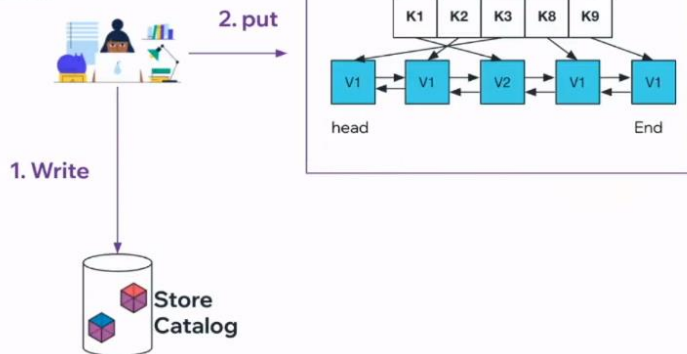
3 Caching Patterns:

1. S3 backed static cache
2. DynamoDB+CDC based cache
3. HTTP Reverse Proxy cache

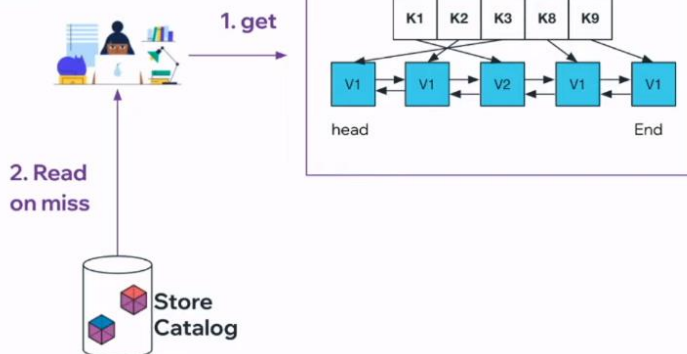
for Request Overload



Stores Catalog Service



Stores Catalog Service



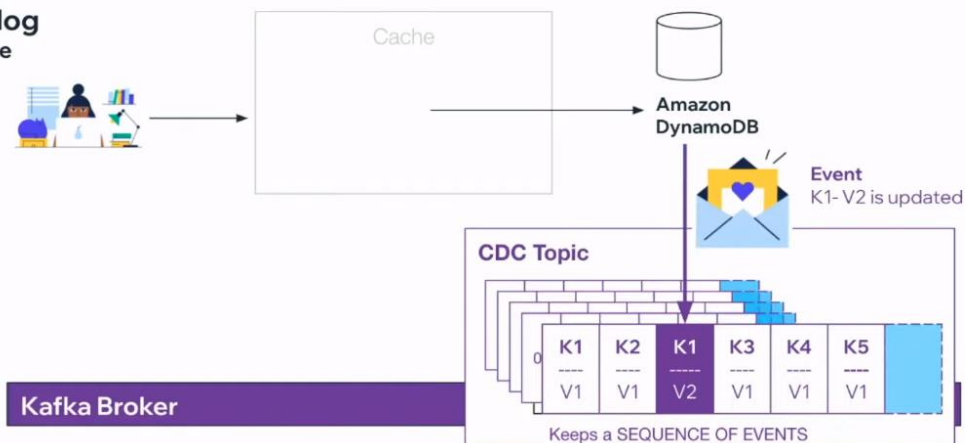
Stores Catalog Service



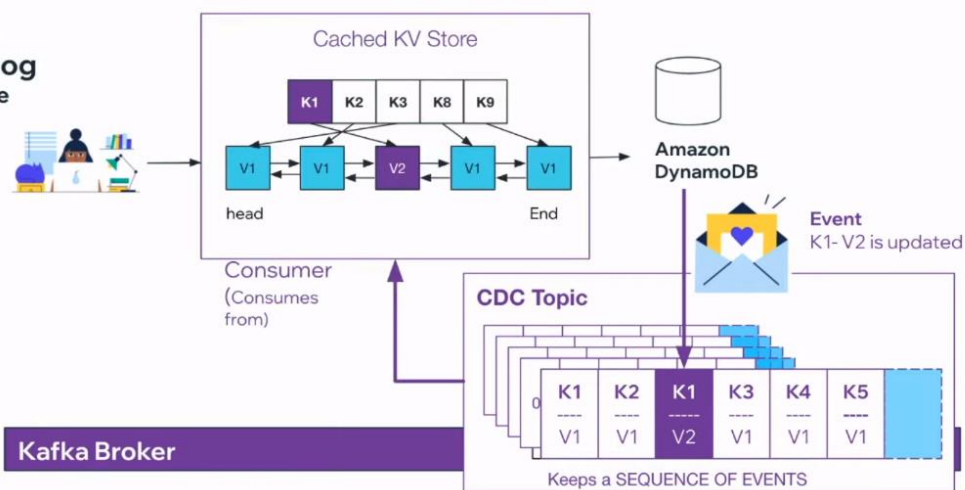
Stores Catalog Service



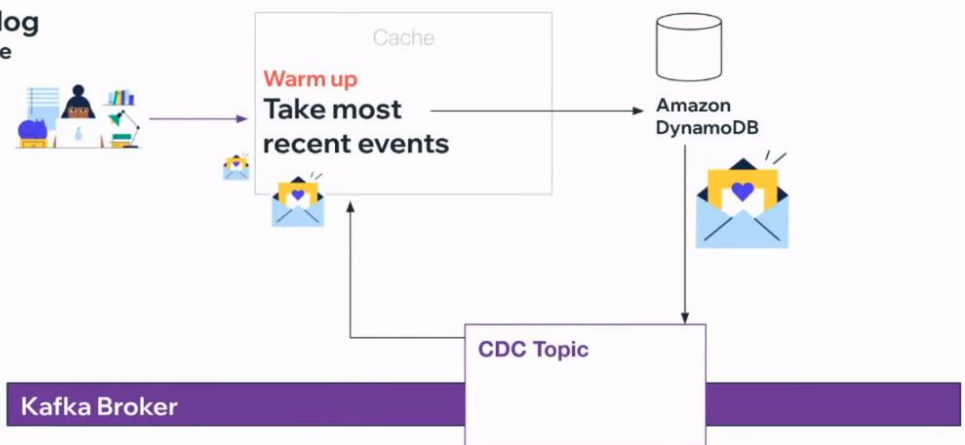
Stores Catalog Service



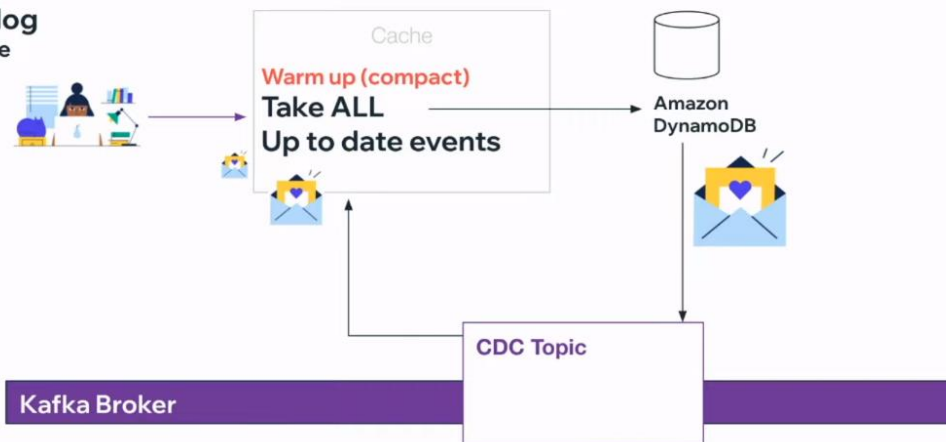
Stores Catalog Service



Stores Catalog Service



Stores
Catalog
Service



Cache Type

Cache Size

Best For

Not For

Cache-Aside
Before the DB

No limitation on DB data size
LRU cache size is configurable

Make DB read access faster
for "popular" values

Highly critical application
startup information



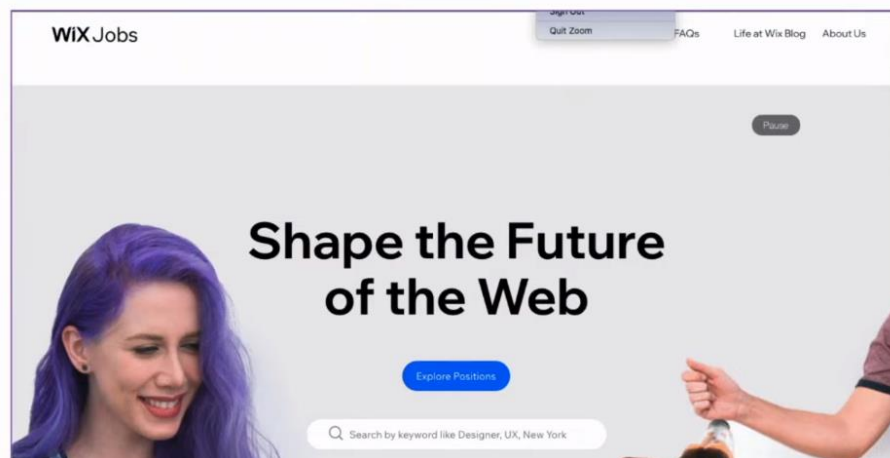
Next


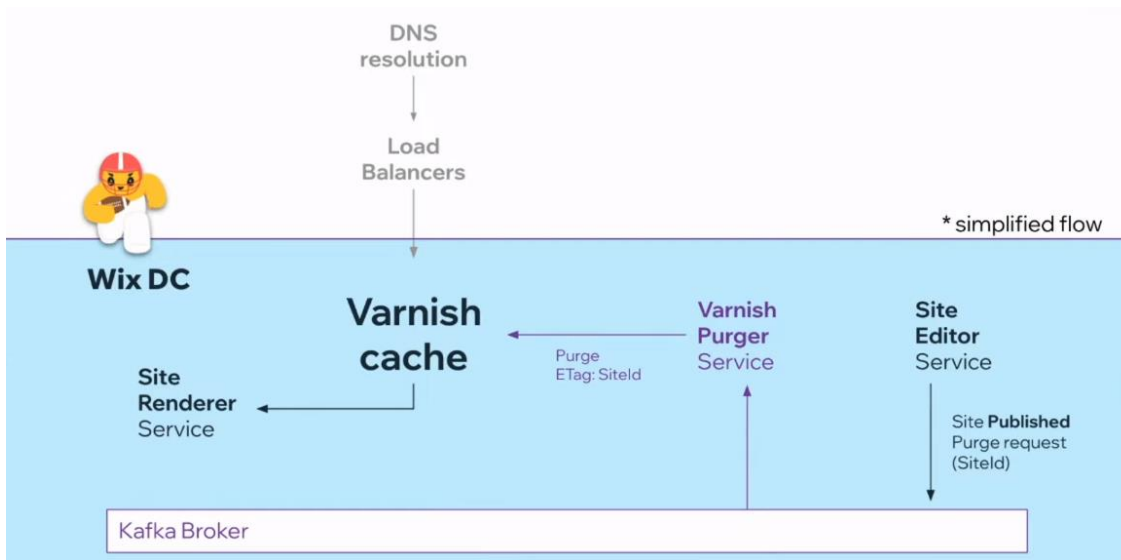
3 Caching Patterns:

S3 backed static cache

DynamoDB+CDC based cache

HTTP Reverse Proxy cache





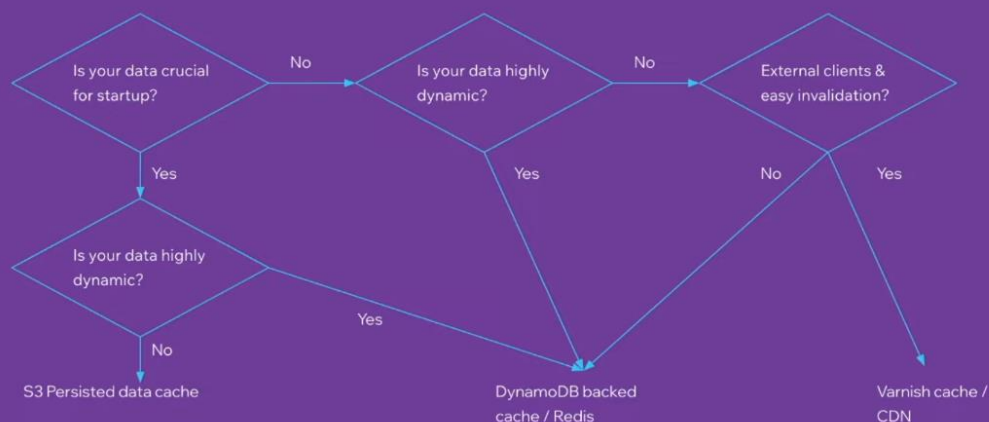
Size limit depends on the configured
Storage backends:

1. **malloc** is memory based — so the size is limited by (more) expensive memory cost
2. **file** is memory-backed-by-disk based — which is limited by less expensive disk cost, but can cause performance penalty due to fragmentation...
3. Massive Storage Engine (**MSE**) — also backed by file, but has a "fragmentation proof algorithm," part of Varnish Cache Plus (paying customers only).

Cache Type	Cache Size	Best For	Not For
Read-through before the app	Depends on storage backend	Improving latency for "stable" HTTP responses	Hard to invalidate aggregated data



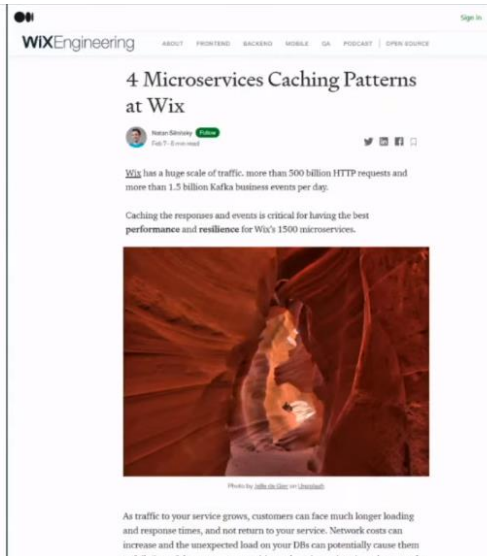

Choose Cache Pattern



The blog post

<https://medium.com/wix-engineering/4-microservices-caching-patterns-at-wix-b4dfee1ae22f>

@NSilnitsky



github.com/wix/greyhound

A Java/Scala high-level SDK for Apache Kafka.

0.2 is out!