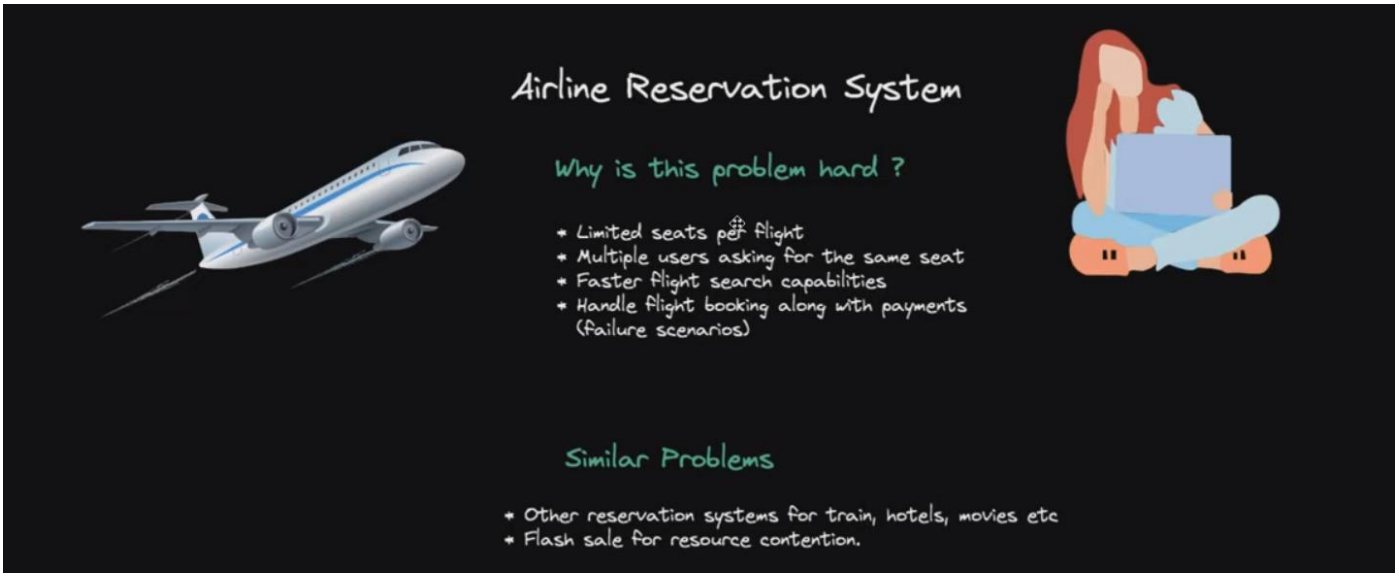

Suggestions - for flight information table, we have primary key as flight id but most relational databases create unique index on primary key across partitions so we shall better than flight id and departure date as composite primary key. Though I like the idea of departure date as partition key. Here, we are assuming that if same plane does multiple round trips between two destination, each will have a separate flight id on given day. Also, instead of booking open, we can call that column as booking status - instead of just open and close, there could be more status like flight cancelled.

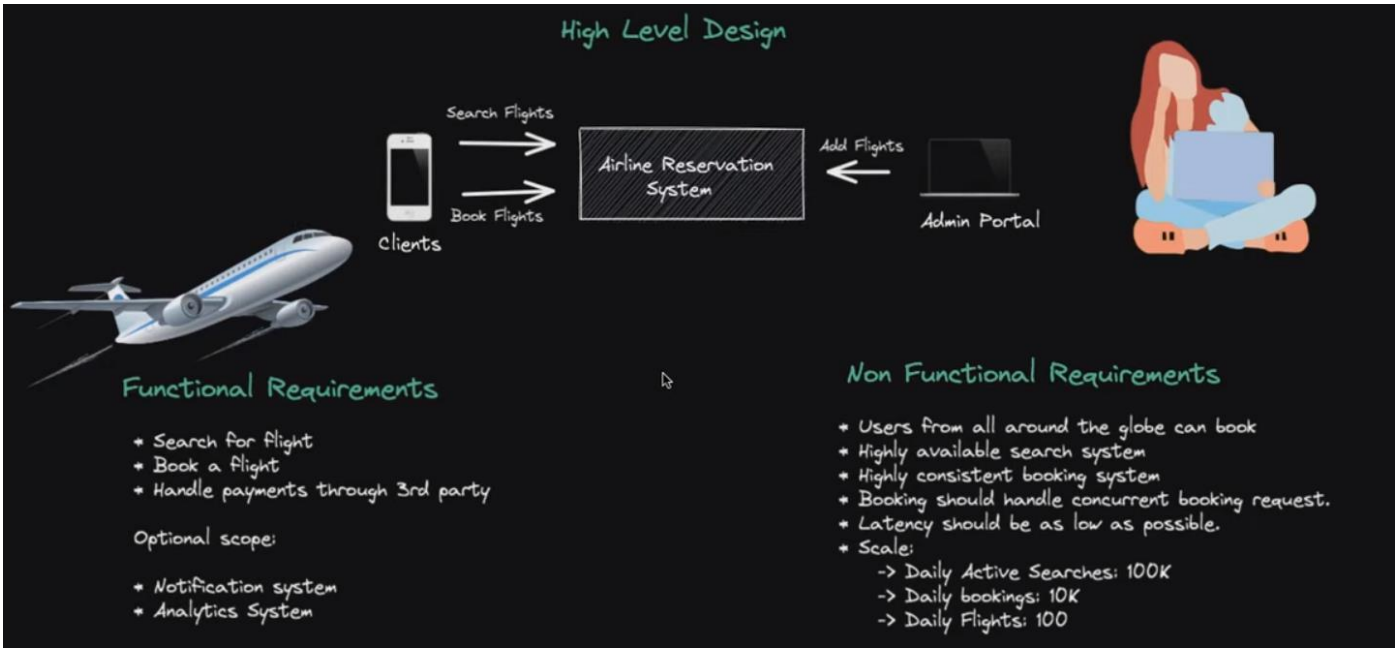

APIs & Capacity Estimates

* Scale:
  -> Daily Active Searches: 100K
  -> Daily bookings: 10K
  -> Daily Flights: 100

APIs

Search: search(from, to, Date)

Book: book(flight-id, customer-id, seat-id)

TPS Estimation

Search: 100k/86400 = ~ 1 TPS
Book: 10K/86400 = ~ 0.1 TPS

Read Heavy system on search

DB schema & Capacity Estimates

* Scale:
  -> Daily Active Searches: 100K
  -> Daily bookings: 10K
  -> Daily Flights: 100

Flight Information

flight id, (PK)
src-destination, (Index)
from (Src),
to (destination),
departure date, (Partition Key)
departure time,
duration,
airlines,
booking open (Filter)

Flight Seat Information

flight-id (PK)
seat-id (PK)
booking-status (Filter)
  * available
  * booked

Booking Information

booking-id (PK)
customer-id
flight-id
seat-id
payment-transaction-id

Storage Estimates

Flight Information:
 -> record size: 1KB
 -> (1KB x 100 x 365) = 36.5 MB/year
 -> ~ 5GB storage for 5 years

Flight Seat Information:
 -> record size: 100B
 -> (100B x 100 x 100 x 365) = 3.6 GB/
 -> ~ 15 GB storage for 5 years

This problem has now got tricky ???

Considering the scale for requests & storage is very less the interviewer can ask you do you need distributed systems and distributed storage ?

Distributed Storage

Why we usually need distributed systems:

* Scaling reads --> We need replicas/caching
* Scaling writes & reads --> We need sharding
* To prevent Single Points of Failure.
  --> serving requests
  --> storage

Distributed Systems

Why we need micro-service architecture:

* Systems should scale independently
* Prevent single points of failure

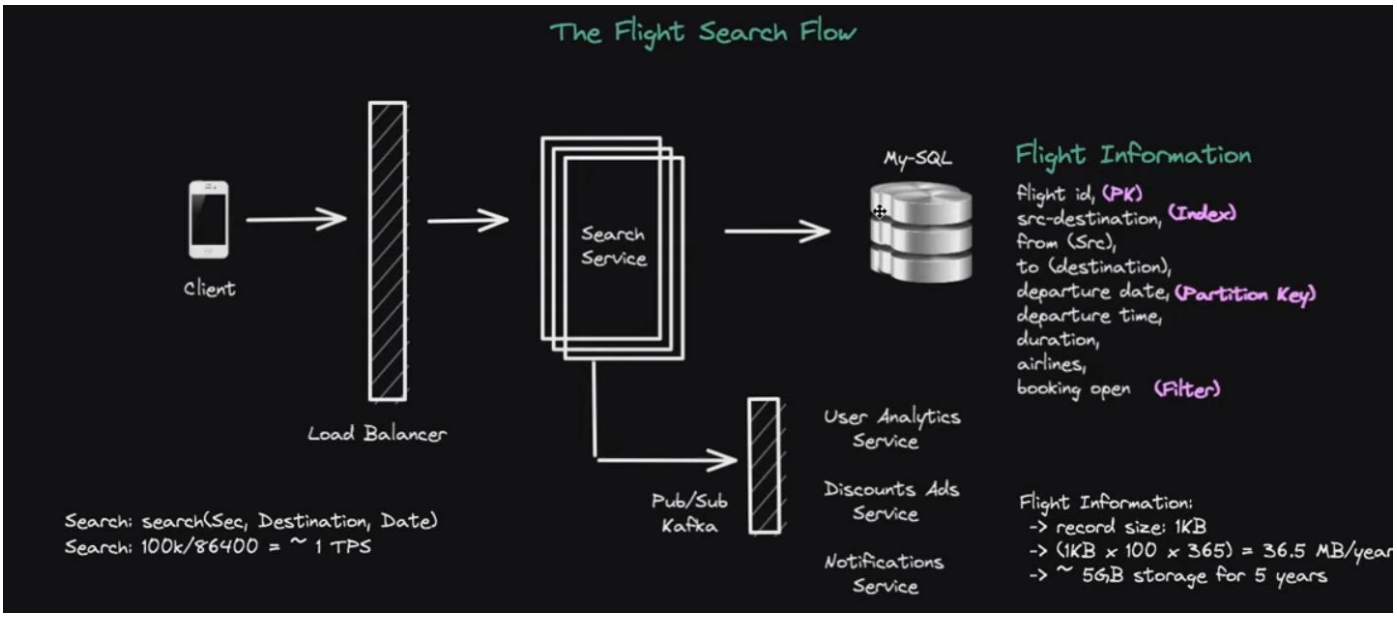
The Flight Search Flow
Client
Load Balancer
Search Service
My-SQL
Flight Information
flight id, (PK)
src-destination, (Index)
from (Src),
to (destination),
departure date, (Partition Key)
departure time,
duration,
airlines,
booking open (Filter)
Pub/Sub Kafka
User Analytics Service
Discounts Ads Service
Notifications Service
Search: search(Sec, Destination, Date)
Search: 100k/86400 = ~ 1 TPS
Flight Information:
-> record size: 1KB
-> (1KB x 100 x 365) = 36.5 MB/year
-> ~ 5GB storage for 5 years

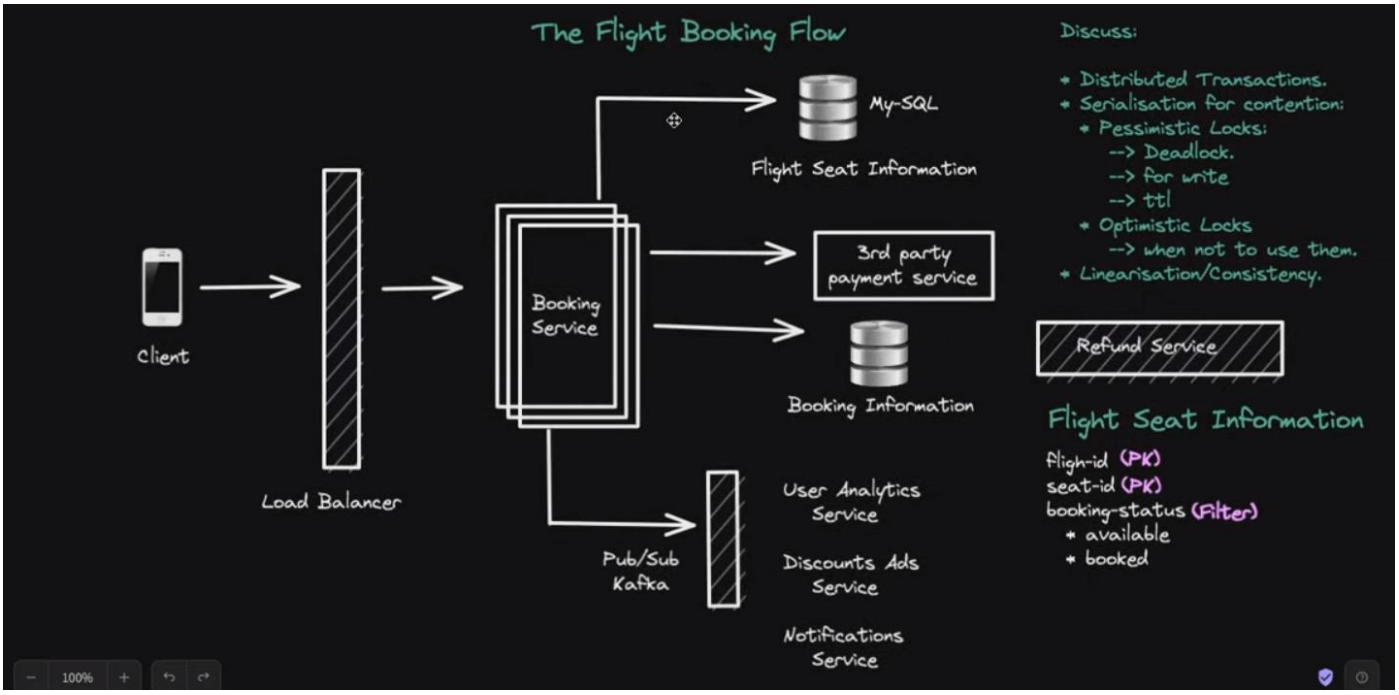
The Flight Booking Flow
Client
Load Balancer
Booking Service
My-SQL
Flight Seat Information
3rd party payment service
Booking Information
Refund Service
Pub/Sub Kafka
User Analytics Service
Discounts Ads Service
Notifications Service
Discuss:
* Distributed Transactions.
* Serialisation for contention:
* Pessimistic Locks:
--> Deadlock.
--> for write
--> ttl
* Optimistic Locks
--> when not to use them.
* Linearisation/Consistency.
Flight Seat Information
fligh-id (PK)
seat-id (PK)
booking-status (Filter)
* available
* booked
100%