

# Beyond Microservices: Streams, State and Scalability

Gwen Shapira, Engineering Manager  
@gwenshap


Microservices have been a popular architecture choice for at least 5 years by now. Over these years we've adopted microservices architectures to ever growing set of use-cases and different development and deployment strategies. Lessons were learned and our ability to design, develop, deploy and operate microservices has improved.

This presentation will give an opinionated view of how microservices evolved in the last few years, based on experience gained while working with companies using Apache Kafka to update their application architecture. We'll discuss the rise of API gateways, service mesh, state management and serverless architectures - what works well, and in which cases. We'll show real-world examples of how applications become more resilient and scalable when new patterns are introduced, and make sure to include caveats - because patterns...

## In the beginning...

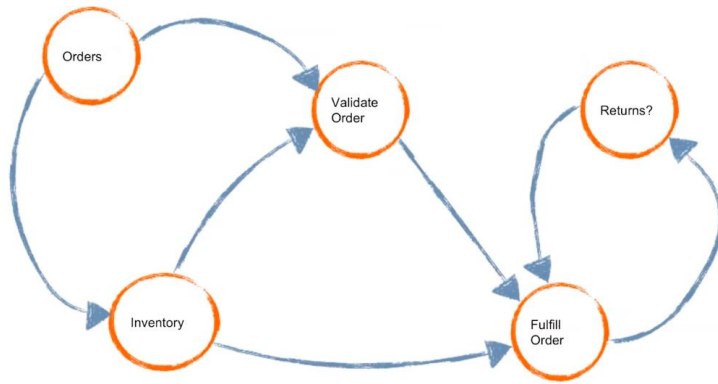


We Have Microservices



They need to communicate


I know! I'll use REST APIs




You now have distributed monoliths

Synchronous request-response communication  
Leads to  
Tight point-to-point coupling


Clients know too much



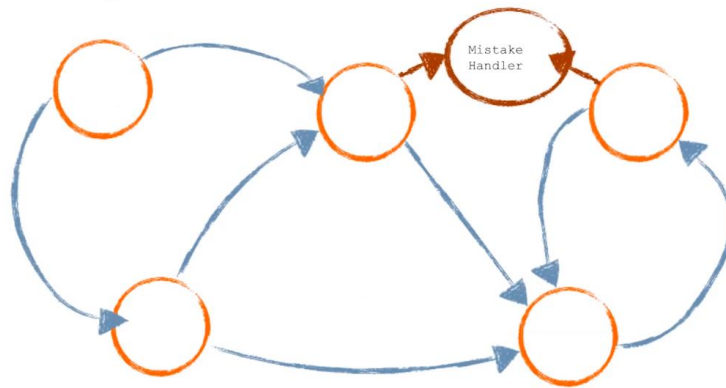
shifts in responsibility, redundancy



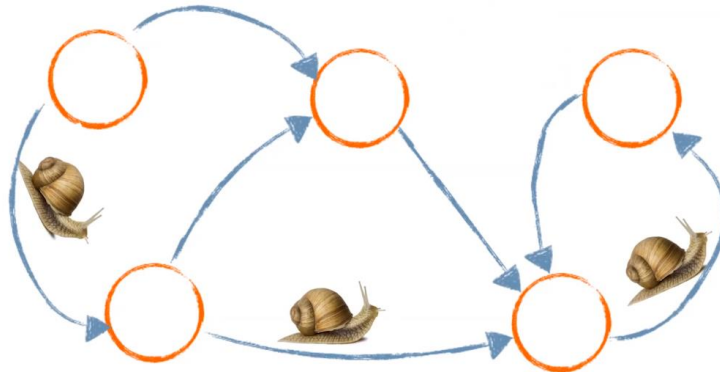
## Making Changes is Risky



## Adding Services Requires Explicit Calls




REST = HTTP + JSON = SLOW



**We can do better.**


## Nice to meet you!

- Moving data around for 20 years
- Engineering Manager at Confluent
- Apache Kafka Committer
- Wrote a book or two
- Tweets a lot - @gwenshap

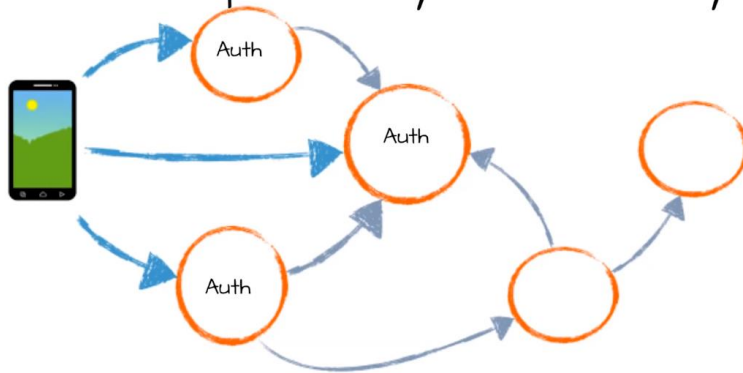


## API Gateway

Clients know too much




Shift in responsibility, redundancy



We don't need all microservices having their own authentication logic

# API Gateway


x



## API Gateway Responsibility

- Authentication
- Routing
- Rate Limiting
- Logging and analytics

## Anti-pattern




The risk is that every microservice wants to use the API Gateway and this becomes a bottleneck and dependency

## Service Mesh


Service Mesh is your distributed internal API Gateway

## North-South Traffic



We want to take this north-south traffic and turn it into an east-west traffic using the Service Mesh

## East-West Traffic





## Side-car




The side-cars is ran next to every microservice in a separate container to implement the distributed API Gateway

## Proxy as sidecar:




The sidecar automatically knows about the new IP via the Service Mesh's discovery/control plane/layer called Istio, the microservice does not care at all.




The sidecar has the logic to limit retries automatically using the configured setting from the control plane Istio

## Event Driven

### Making Changes is Risky




### Adding Services Requires Explicit Calls



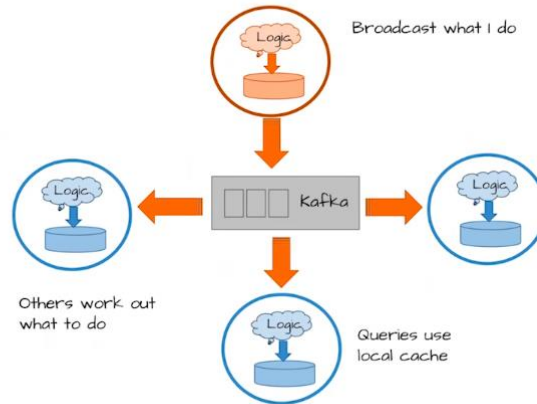
We can solve the 2 problems above by using the event-driven pattern to just tell what it known (events) below



## Request Driven



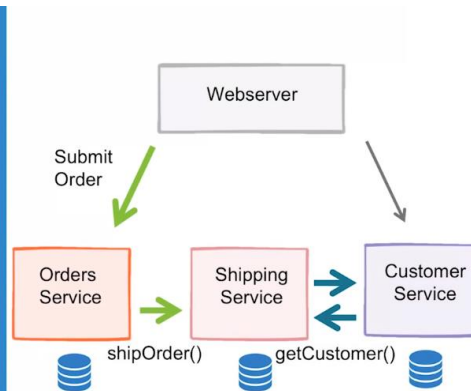
## Event Driven



## Events are both facts and triggers

### Buying an iPad (with REST)


- Orders Service calls Shipping Service to tell it to ship item.
- Shipping service looks up address to ship to (from Customer Service)



What if the Shipping Service is down? Does the Orders Service have to keep retrying?

### Using events for Notification

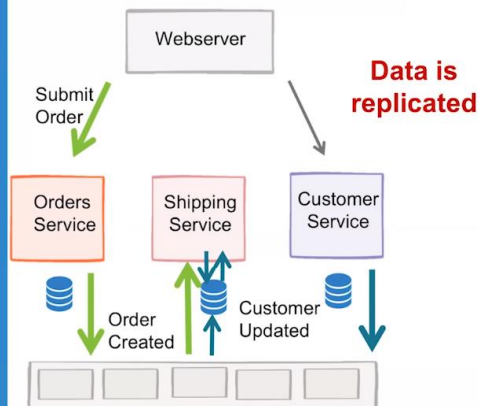
- Orders Service no longer knows about the Shipping service (or any other service). Events are fire and forget.






## Using events to share facts

- Call to Customer service is gone.
- Instead data is replicated, as events, into the shipping service, where it is queried locally.




## DB for Each Microservice?




- **It is safe:** They are all derived from same stream of events
- **Custom projection** just the data each service needs.
- Reduced dependencies
- Low latency

## Event Driven Microservices are Stateful

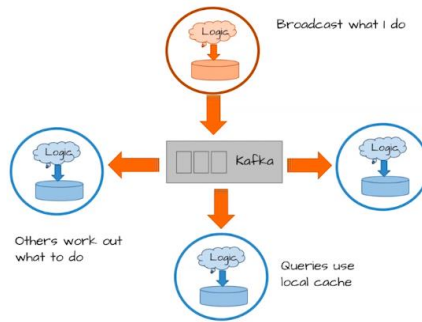


## Schema

## Request Driven



## Event Driven



Event driven world runs on events with messages


**The medium is not the message.**

### This is a message


```
{
  sessionId: 676fc8983gu563,
  timestamp: 1413215458,
  viewType: "propertyView",
  propertyId: 7879,
  loyaltyId: 6764532
  origin: "promotion",
  ..... lots of metadata....
}
```



REST = HTTP + JSON = SLOW





Making Changes is Risky

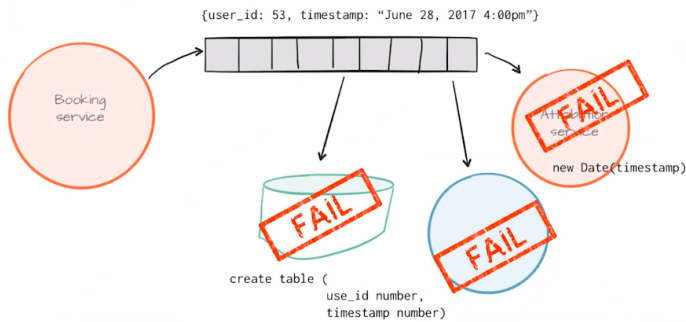


JSON is slow and has very little validation and events can be written with new, unvalidated data fields that can break our microservices. We need to have another data format for our messages or add JSON validations using schemas.


There are lots of dependencies



Moving fast and breaking things




**APIs between services are Contracts**  
**In Event Driven World – Event Schemas ARE the API**



The Confluent Schema Registry now supports writing schemas in Avro, JSON, and ProtoBuf.


So the flow is...




Every schema change can now be validated for compatibility before adding it to the schema registry

## Serverless

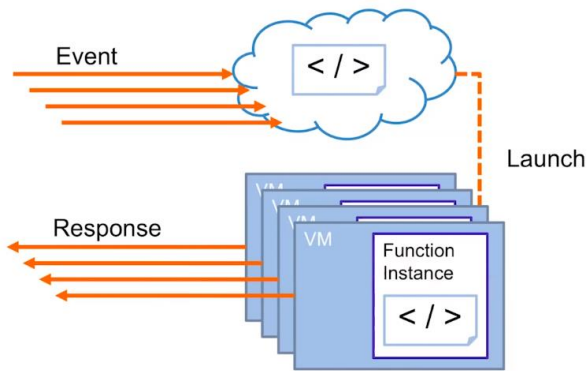
### Function as a Service



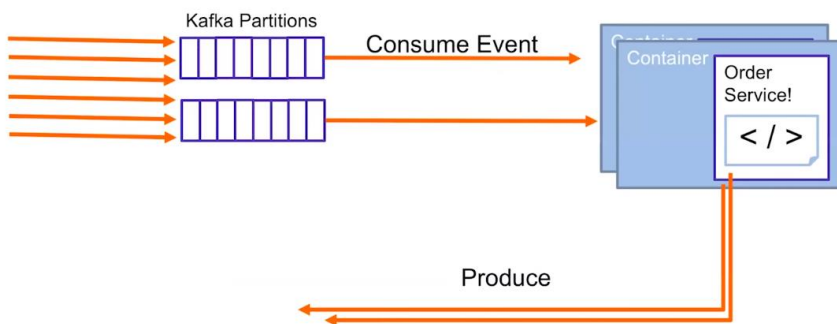
### When nothing happens



## At scale



## Wait, this is super familiar



Event-driven is all about events, the only missing thing is our state

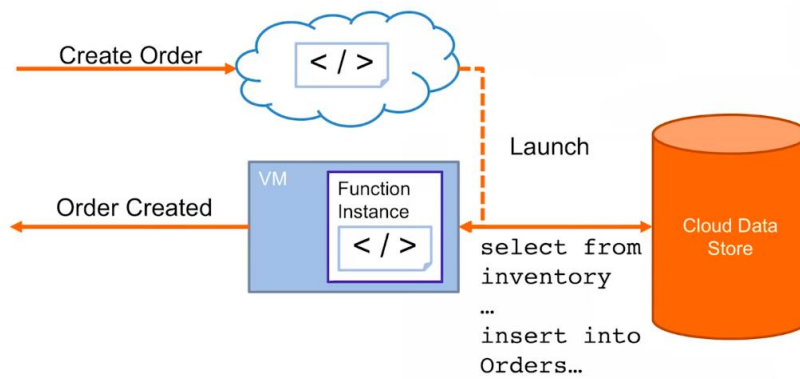
## Up Next: Stateful Serverless

### State is required

- Dynamic Rules
- Event enrichment
- Joining multiple events
- Aggregation

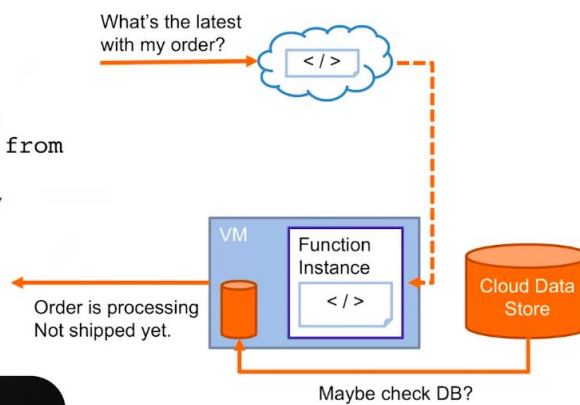


# How You Probably Do State




## We can do a bit better

```
Select order_id,  
customer_name, product,  
quantity, price, state from  
orders  
where state != "CLOSED"
```




goto;

## But I really want this back:





# Stateful Serverless



## What's Still missing?

- Durable functions everywhere
- Triggers and data from data stores to functions
- Unified view of current state