


FROM **TIC TAC TOE**



TO

ALPHA ZERO



Google DeepMind's AlphaGo was an extraordinary breakthrough for Artificial Intelligence. The game of Go has 1.74×10^{172} unique positions and is about a 'googol' times harder to calculate than chess. Experts thought it would take at least another decade before A.I. would be able to beat the best human players. So how did DeepMind tackle this problem? What algorithms did they [...]

TOPICS


- SIMPLE TREE GAME (TREE SEARCH, MINI-MAX)
- NOUGHTS AND CROSSES (PERFECT INFORMATION, GAME THEORY)
- CHESS (FORWARD/BACKWARD AND ALPHA/BETA PRUNING)
- GO (MONTE CARLO TREE SEARCH, NEURAL NETWORKS)

I WANNA PLAY A GAME...


- TREE-STRUCTURE
- **YOU** ALWAYS START
- HIGHEST SCORE **WINS**



DEPTH: N=1



DEPTH: N=3




 @ROYVANRI.IN

This is no longer trivial and we can use the minimax algorithm for this

MINIMAX

- MINIMISE THE MAXIMUM SCORE (WHEN IT IS THE OPPONENTS TURN)
- MAXIMISE THE MINIMUM SCORE (WHEN IT IS OUR TURN)
- THIS SIMULATES 'PERFECT PLAY'

DEPTH: N=3




@ROYVANRIJN

We generally start at the bottom and evaluate all the nodes to pick the node with the highest value

```
int minimax(Node node, boolean maximizingScore) {  
    if(node.isEndNode()) {  
        return node.evaluate();  
    }  
  
    int bestScore = maximizingScore ? Integer.MIN_VALUE : Integer.MAX_VALUE;  
    for(Node child: node.getChildren()) {  
        int score = minimax(child, !maximizingScore);  
        if(maximizingScore) {  
            bestScore = Math.max(score, bestScore);  
        } else {  
            bestScore = Math.min(score, bestScore);  
        }  
    }  
    return bestScore;  
}
```

SIMPLE TREE GAME

- BRANCHING FACTOR: 2
- GAME DEPTH: 3
- PERFECT INFORMATION



@ROYVANRIJN



PLAYING A GAME

using the computer



- A WAY TO **GENERATE** ALL (VALID) **MOVES** (CREATE THE TREE)
- A WAY TO **EVALUATE** NODES
- A WAY TO PICK A **PATH** IN THIS **TREE**

NOUGHTS AND CROSSES

butter, cheese and eggs?



@ROYVANRIJN

NOUGHTS AND CROSSES

butter, cheese and eggs?



@ROYVANRIJN

You can get very large possibilities for the choice here, but you can eliminate some duplicates if you take the 1st choice and rotate it to get the 3rd, 7th and 9th choices. You can also mirror the 2nd choice to get the 8th choice.

NOUGHTS AND CROSSES

butter, cheese and eggs?



- **BOTTOM NODES HAVE VALUES:**

WIN: **1**
TIE: **0**
LOSE: **-1**

We can evaluate the nodes using the above approach

NOUGHTS AND CROSSES

butter, cheese and eggs?




- **BRANCHING FACTOR:** **5** = $(9+8+7+6+5+4+3+2+1)/9$
- **DEPTH:** MAX **9** MOVES
- **REMOVING SYMMETRIES THERE ARE 138** TERMINAL POSITIONS
× WINS 91 TIMES, ○ WINS 44 TIMES, 3 DRAWS

This is a perfect information game where we can calculate the entire game easily, we always end up in one of the terminal positions in at most 9 moves.



THE GAME OF CHESS



Chess is no longer trivial and takes much longer to play

'PERFT'

Depth	Nodes	Captures	E.p.	Castles	Promotions	Checks	Checkmates
0	1	0	0	0	0	0	0
1	20	0	0	0	0	0	0
2	400	0	0	0	0	0	0
3	8,902	34	0	0	0	12	0
4	197,281	1576	0	0	0	469	8
5	4,865,609	82719	258	0	0	27351	347
6	119,060,324	2812008	5248	0	0	809099	10828

This is the Perft table for the chess start board. A game engine's suitability is in how fast it can generate all the possible moves at each depth and choose the best next move.

CHESS

- **NO** PERFECT INFORMATION
- HOW DO WE **EVALUATE** A NODE?
- **COUNT** THE PIECES
- **EVALUATE** PIECE **POSITIONS**/LIBERTIES


This allows us to write an evaluation function that we can use to evaluate the board positions up to a certain depth, assign it a value,


HORIZON PROBLEM



PRUNING

- WE NEED TO **CUT** BACK THE TREE
- **FORWARD** PRUNING: RISKY
- **BACKWARD** PRUNING: SAFE

 @ROYVANRIJN




FORWARD PRUNING

- IF A **MOVE** IS TOO **BAD**: STOP EVALUATING
- IF A **MOVE** IS TOO **GOOD**: STOP EVALUATING




ALPHA/BETA PRUNING



@ROYVANRIJN

ALPHA/BETA PRUNING



@ROYVANRIJN

```
int minimax(Node node, boolean maximizingScore) {
    if(node.isEndNode()) {
        return node.evaluate();
    }

    int bestScore = maximizingScore ? Integer.MIN_VALUE : Integer.MAX_VALUE;
    for(Node child: node.getChildren()) {
        int score = minimax(child, !maximizingScore);
        if(maximizingScore) {
            bestScore = Math.max(score, bestScore);
        } else {
            bestScore = Math.min(score, bestScore);
        }
    }
    return bestScore;
}
```



```

int alphaBeta(Node node,                      boolean maximizingScore) {
    if(node.isEndNode()) {
        return node.evaluate();
    }

    int bestScore = maximizingScore ? Integer.MIN_VALUE : Integer.MAX_VALUE;
    for(Node child: node.getChildren()) {
        int score = alphaBeta(child,                      !maximizingScore);
        if(maximizingScore) {
            bestScore = Math.max(bestScore, score);
                                                
        } else {
            bestScore = Math.min(bestScore, score);
                                                
        }
                                            
    }
    return bestScore;
}

```

@ROYVANRIJN

```

int alphaBeta(Node node, int alpha, int beta, boolean maximizingScore) {
    if(node.isEndNode()) {
        return node.evaluate();
    }

    int bestScore = maximizingScore ? Integer.MIN_VALUE : Integer.MAX_VALUE;
    for(Node child: node.getChildren()) {
        int score = alphaBeta(child, alpha, beta, !maximizingScore);
        if(maximizingScore) {
            bestScore = Math.max(bestScore, score);
            alpha = Math.max(alpha, bestScore);
        } else {
            bestScore = Math.min(bestScore, score);
            beta = Math.min(beta, bestScore);
        }
        if(beta <= alpha)
            break; /* stop evaluating */
    }
    return bestScore;
}


```

@ROYVANRIJN

Using pruning and an evaluation function to pick a move, we can now have a much deeper search and a much better chess engine

PLAYING CHESS

using the computer



- A WAY TO **GENERATE** ALL VALID **MOVES** -> CHESS ENGINE
- A WAY TO **EVALUATE** NODES -> COUNT PIECES
- A WAY TO PICK A **PATH** IN THIS **TREE** -> ALPHA/BETA SEARCH


CHESS



- AVERAGE BRANCHING FACTOR: 35
- AVERAGE GAME DEPTH: 40-50 MOVES
- EVALUATION FUNCTION: RELATIVELY EASY
- ADVANCED CHESS A.I. CAN LOOK 20+ MOVES AHEAD



THE GAME OF GO

 @ROYVANRIJN

ABOUT THE GAME

- BOARD: 19x19
- BLACK AND WHITE STONES
- SURROUND AND CAPTURE AREAS



COMPLEXITY OF GO



- FIRST PROBLEM: BRANCHING FACTOR: +/- 250
- SECOND PROBLEM: GAME DEPTH: 300+ MOVES
- THIRD PROBLEM: EVALUATION FUNCTION:


COMPLEXITY OF GO



1.74×10^{172}

(LARGER THAN THE AMOUNT OF ATOMS IN THE ENTIRE UNIVERSE)

MONTE CARLO TREE SEARCH



@ROYVANRIJN

How can we also use random search to make a better Go playing engine?

MONTE CARLO TREE SEARCH

- PICK A **NODE**
- PLAY (SEMI-) **RANDOM** MOVES TO THE END
(AS OFTEN AS POSSIBLE)
- THIS GIVES A STRONG **INDICATION** OF THE STRENGTH



EXPERTS IN 2015:

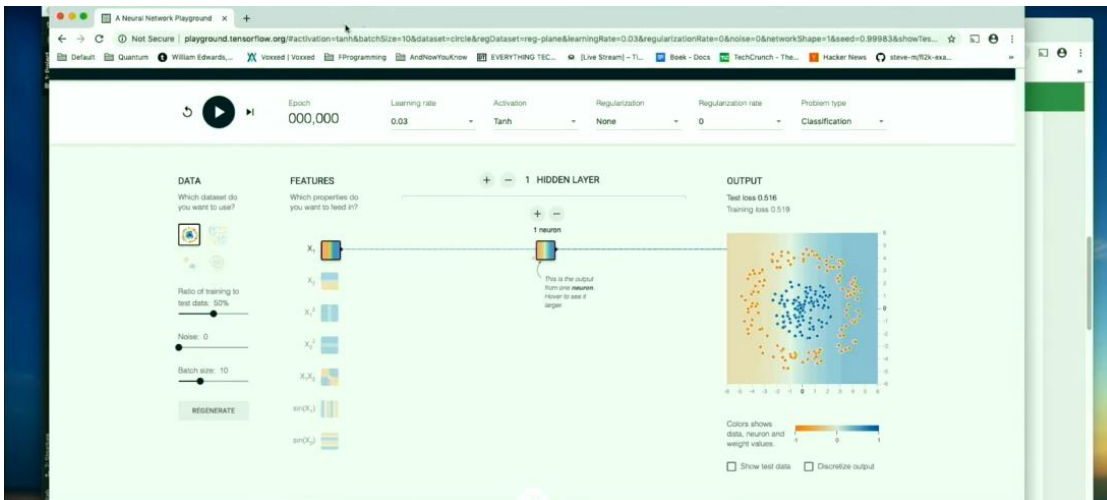
"IT WILL PROBABLY TAKE **10** TO **15** YEARS BEFORE A
COMPUTER CAN BEAT A PROFESSIONAL GO PLAYER"



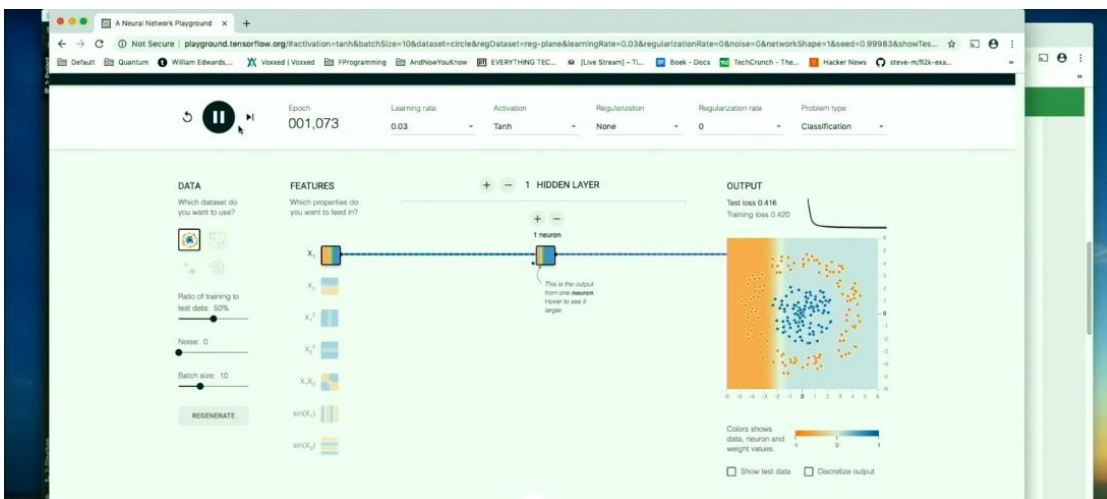
NEURAL NETWORK

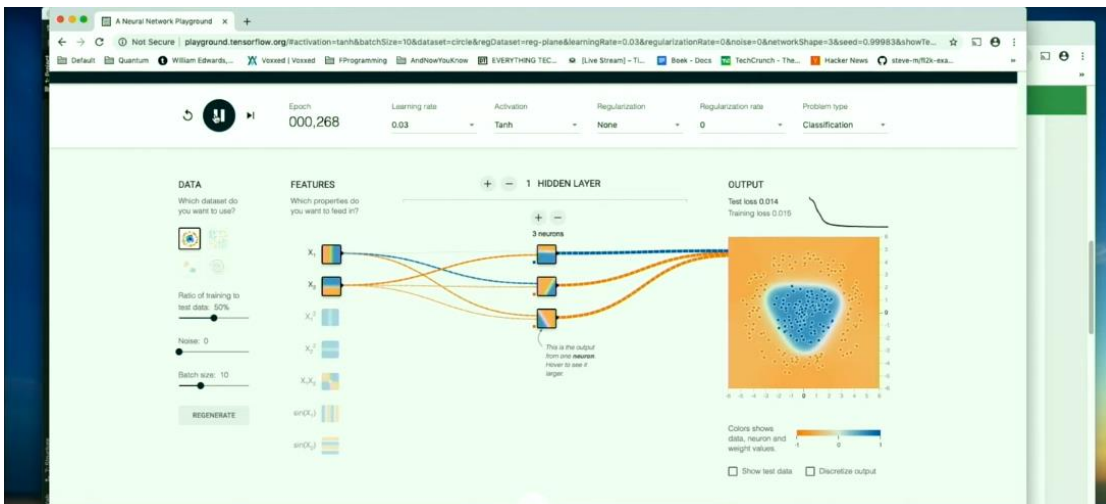
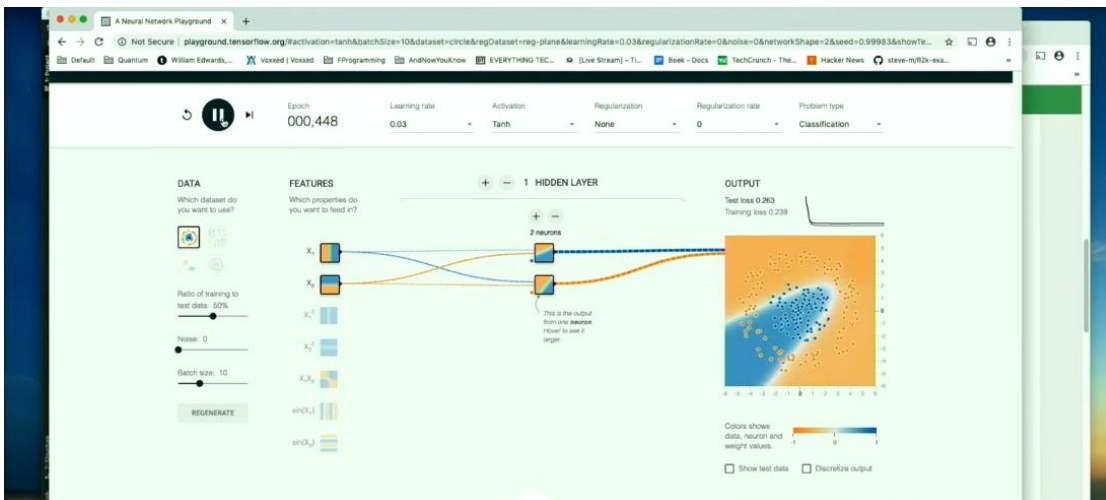
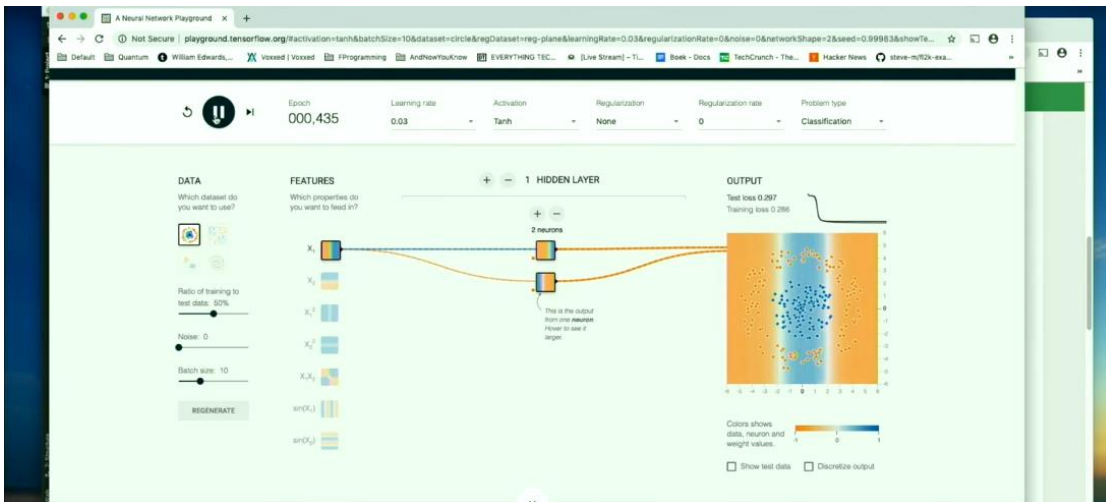
A NEURAL NETWORK IS A **COMPUTER MODEL** DESIGNED TO SIMULATE
THE BEHAVIOUR OF **BIOLOGICAL NEURAL NETWORKS**

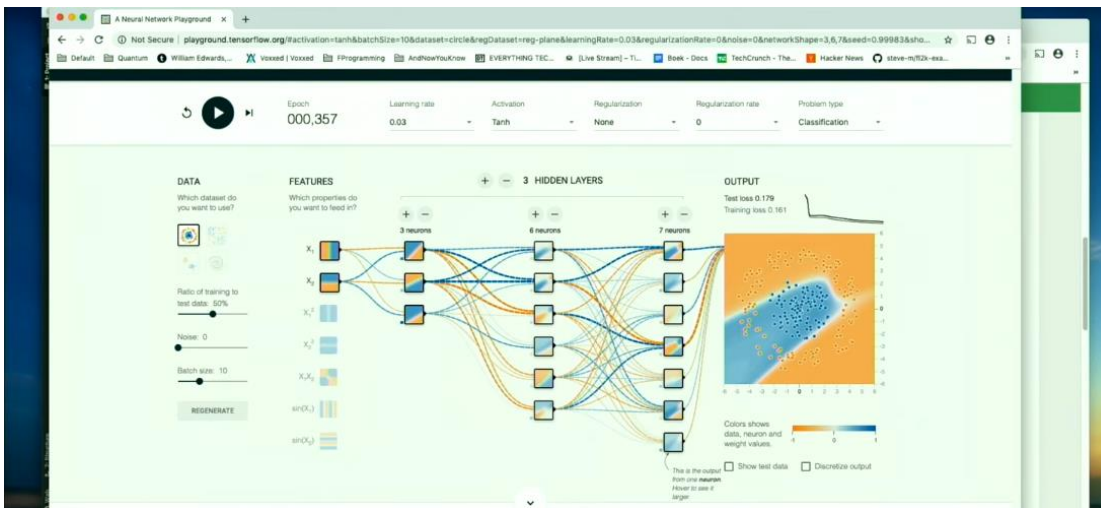




You can build a NN graphically and solve a sorting problem







MORE INFORMATION

- **TENSORFLOW** (<https://www.tensorflow.org>)
- **DEEPLARNING4J** (<https://deeplearning4j.org/>)
- **CAFFE, TORCH, THEANO, ETC**



@ROYVANRIJN

TELL ME HOW!



AlphaGo

@ROYVANRIJN

NEURAL NETWORKS IN ALPHAGO

- **CONVOLUTIONAL NEURAL NETWORKS**
- **LEARNING IS SUPERVISED**
- **HAS HIDDEN 13-LAYERS**



#1 SUPERVISED LEARNING POLICY NETWORK

- 30 MILLION AMATEUR MATCHES
- GOAL: PREDICT THE NEXT MOVE
- RESULT: 57% CORRECT



#2 REINFORCED LEARNING POLICY NETWORK

- COPY OF SUPERVISED NETWORK
- NEW GOAL: PREDICT THE *BEST* MOVE
- NETWORK PLAYED ITSELF 1.2 MILLION TIMES (TOOK ONE DAY)
- PLAYS PACHI AND WINS: 85% OF THE TIME (WITHOUT SEARCH!)



#3 FAST ROLLOUT POLICY NETWORK

- THE REINFORCED NETWORK IS SLOW: 3_{MS}
- TOO SLOW FOR MONTE CARLO SEARCH
- THIS IS SMALLER, BUT FASTER: 2_{μs} 1500_x



#4 VALUE NETWORK



- TRAINED USING THE SAME 30 MILLION GAMES
- PREDICTS THE WINNER BASED ON CURRENT BOARD
- INITIALLY HAD ERROR OF 0.37 (0.5 IS RANDOM)
- AFTER SELF-PLAY ERROR CAME DOWN TO ~0.23

#4 VALUE NETWORK



- TESTING THE VALUE NETWORK
- FOR A GIVEN BOARD, GENERATE ALL MOVES
- FOR ALL MOVES, EVALUATE AND PICK THE BEST NEXT MOVE
- BEATS THE STRONGEST KNOWN A.I. STILL WITHOUT TREE-SEARCH (!!!)

COMBINING ALL THE PIECES

- USE POLICY NETWORK TO LOOK AT THE CURRENT BEST MOVES
- FOR THOSE MOVES, USE THE VALUE NETWORK TO DOUBLE CHECK
- USE FAST ROLLOUT NETWORK FOR MONTE CARLO TREE SEARCH

THE CHALLENGE

- LEE SEDOL: THE **BEST** GO PLAYER OF THIS DECADE
- BEST OF **5** GAMES WINS
- WINNER GETS \$**1,000,000.-**



@ROYVANRIJN

THE CHALLENGER

- **DISTRIBUTED** ALPHAGO:
- 1202 CPU's



@ROYVANRIJN

GAME 1, MOVE 102



@ROYVANRIJN

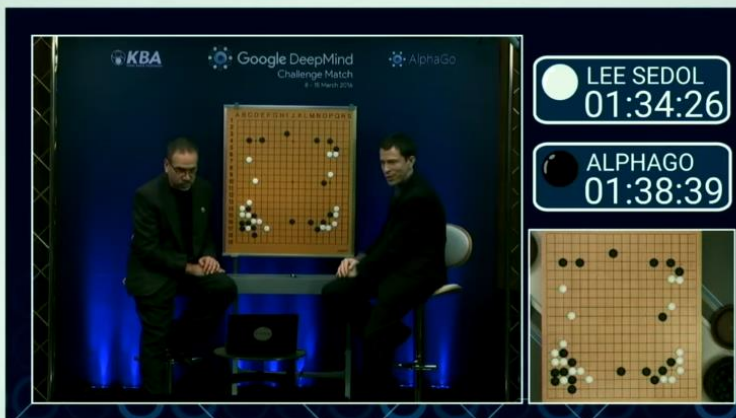
GAME 1, MOVE 102



GAME 2, MOVE 37



GAME 2, MOVE 37



EUROPEAN CHAMPION **FAN HUI**:

**“IT’S NOT A HUMAN MOVE.
I’VE NEVER SEEN A HUMAN PLAY THIS MOVE,
SO **BEAUTIFUL**.”**

GAME 4, MOVE 78


GU LI (LEE’S ARCHRIVAL):

“THIS MOVE WAS MADE WITH **THE HAND OF GOD.”**

RESULTS

ALPHA GO 4 - LEE SEDOL 1



 @ROYVANRIJN

NOBODY TAUGHT ALPHAGO WHAT A **GOOD OR **BAD** MOVE IS**

NOBODY PROGRAMMED AN **EVALUATION FUNCTION FOR ALPHAGO**

ALPHAGO ISN'T AN **EXPERT SYSTEM**

ALPHAGO **LEARNED** BY WATCHING OTHERS AND SELF-PLAY

USING **GENERAL** MACHINE LEARNING TECHNIQUES
TO FIGURE OUT FOR **ITSELF** HOW TO WIN AT GO...

AS A RESPONSE TO THE SUCCESS OF ALPHAGO,
SOUTH KOREA ANNOUNCED ON 17 MARCH 2016 THAT IT
WOULD INVEST **\$863 MILLION** IN ARTIFICIAL-
INTELLIGENCE RESEARCH OVER THE NEXT FIVE YEARS.



- ALPHAGO **ZERO** VERSUS ALPHAGO: **100 - 0**
- SUPERHUMAN ABILITIES FOR: **CHESSE, SHOGI**
- PROTEIN FOLDING: ALPHAFOLD
- STARCRAFT: ALPHASTAR
- ULTIMATE GOAL: DEEPMIND **HEALTH...**

