

# The Java Cloud-Native Stack for Microservices & Serverless Architecture

Java has been leading for more than 20 years. However, Java is known as complicated, heavy-weight, slow, and for its slow startup time and high RAM consumption. Now, the cloud age has begun and the cloud changes everything. Is Java still suited for the cloud? Are traditional databases still suited for the cloud or is NoSQL the way to go? Can in-memory computing speed up my system and how does Java support it? Is serverless really interesting for Java developers and how do Java and serverless fit together?

No worries, the holy grail was already found. A completely new Java stack was created to build real cloud-native apps with Java. Native images with Java, millisecond startup-time, in-memory data processing, up to 1000x faster database queries, highly cost-efficient serverless infrastructure, everything built with pure Java. That's not a vision, it's reality. In this session, you will learn how to build ultra-fast cloud-native apps and microservices with core Java and modern Java micro-frameworks only.



**Markus Kett**  
CEO at MicroStream,  
Contributor to Helidon (Oracle)  
Editor in Chief at JAVAPRO Magazine  
Organizer JCON Conference  
Conference Speaker

MicroStream  
RAPIDclipse  
JPA-SQL  
JAVAPRO

## RapidClique 11

**Low-Code Visual Java IDE  
based on Eclipse.**

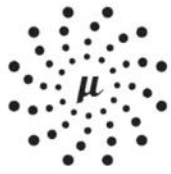
**Build high-performance in-memory database applications in record time.**  
FOR WINDOWS, LINUX, MAC – OPEN SOURCE FRAMEWORK – GET STARTED RIGHT AWAY.

OPEN SOURCE FRAMEWORK

## MicroStream is Part of Helidon



Helidon is a fast framework for developing modern cloud-native microservices with Java. Helidon is mainly developed by Oracle.



M I C R O N A U T™

## MicroStream is Integrated with Micronaut



Micronaut is a modern, JVM-based, full-stack framework for building modular, easily testable microservice and serverless applications.

### Why Performance?



New Innovation



Saving Resources



Saving Costs



Better CX



Higher Productivity

Amazon Study 2009:

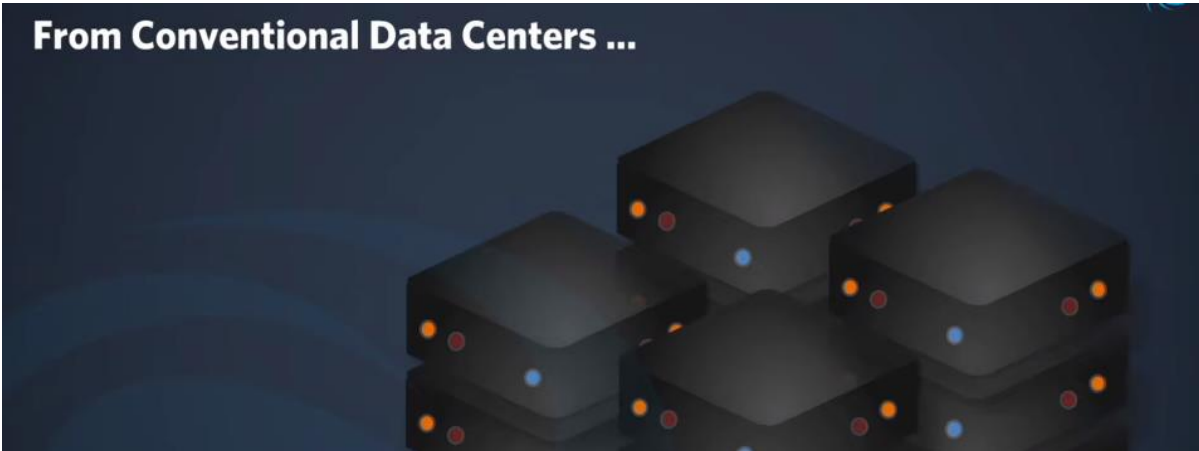
**100 ms of Latency  
Cost 1% in Sales**

[Source: gigaspaces.com/blog](http://gigaspaces.com/blog)

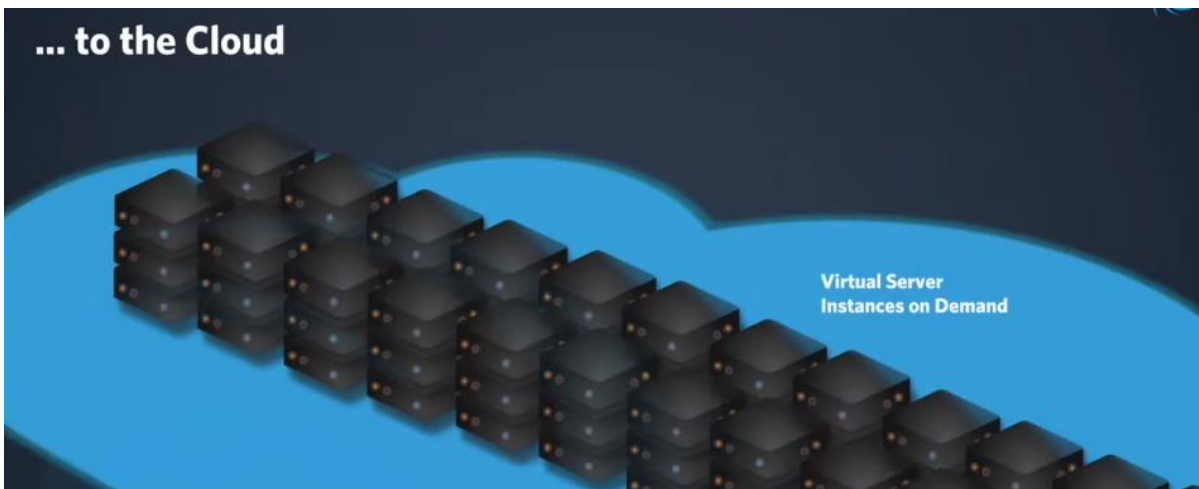


# Cloud has changed everything

From Conventional Data Centers ...

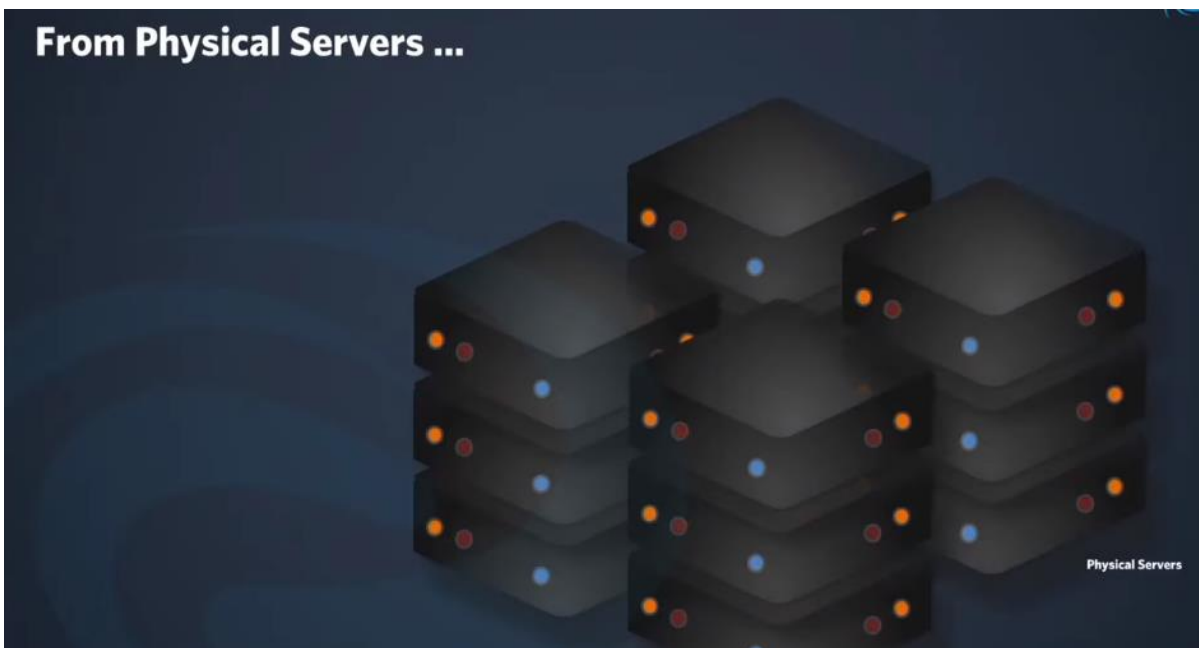


... to the Cloud

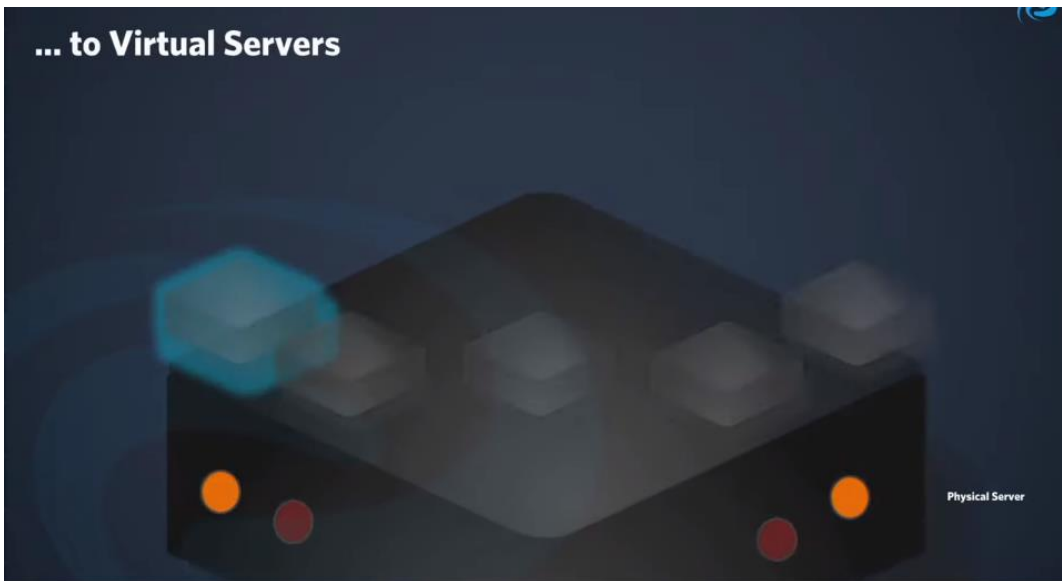


## Virtualization

From Physical Servers ...



## ... to Virtual Servers

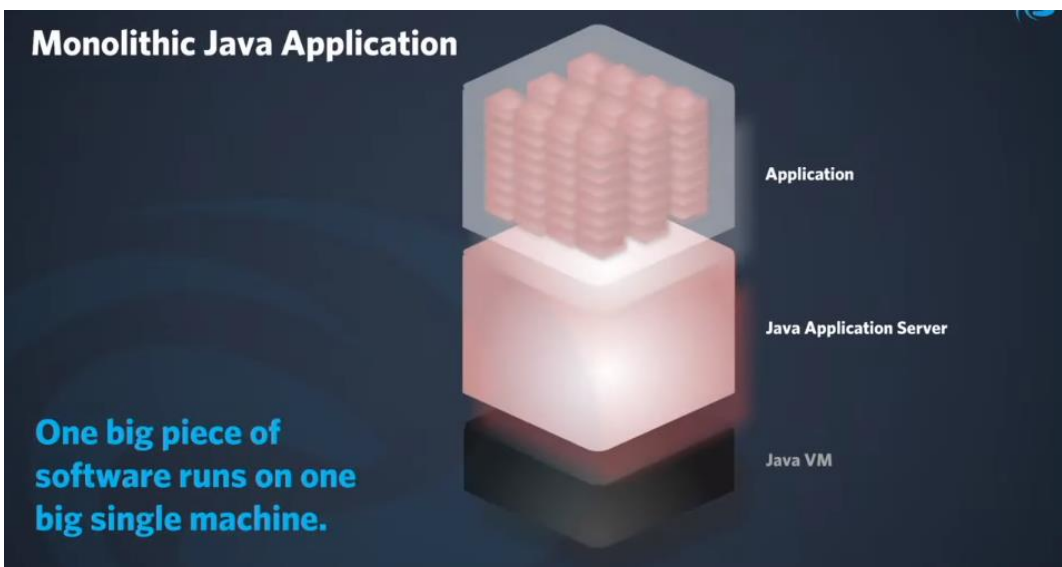


## ... to Container



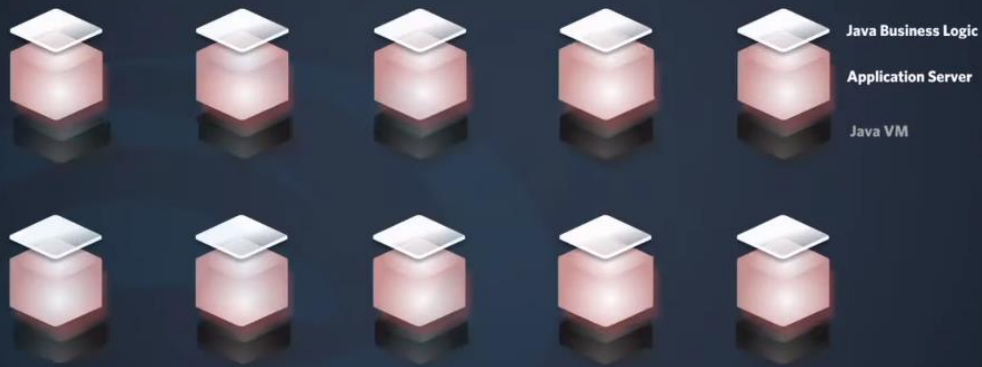
# MicroServices

## Monolithic Java Application





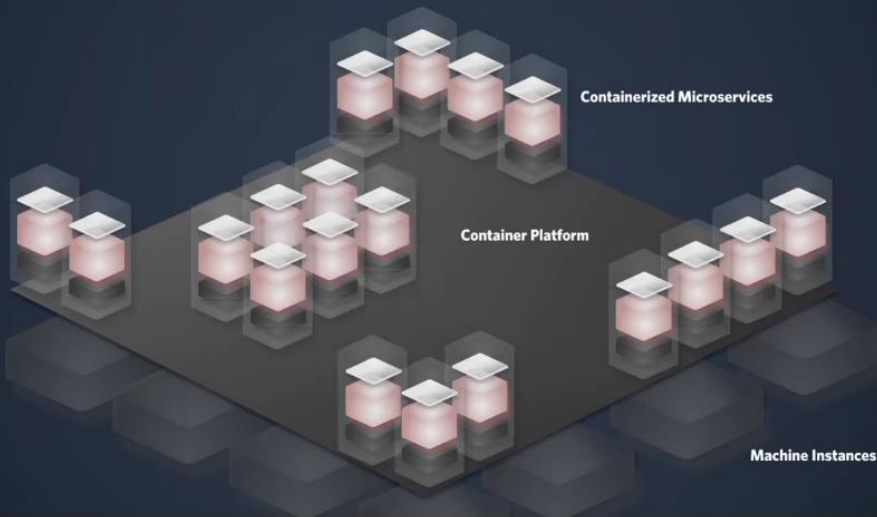
... Numerous of Single Microservices ...



... that Run Within Their Own Container ...



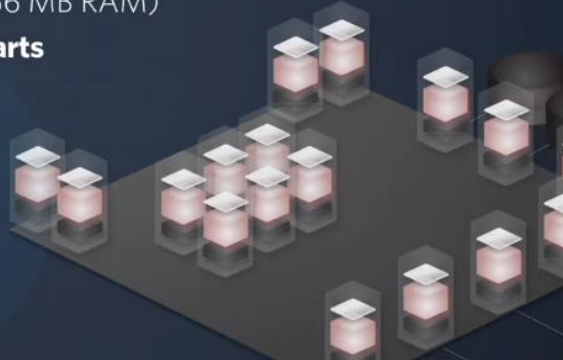
... On a Container Platform



**Challenges with Microservices Using the  
Traditional Java Stack**

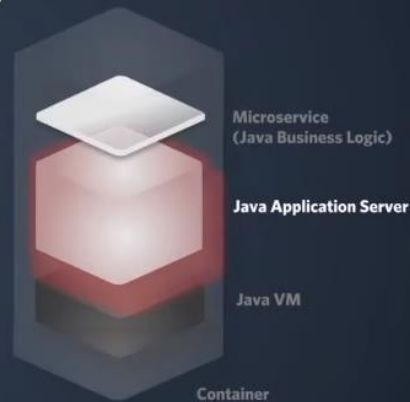
## Environment & Requirements to Microservices

- **Responsible for only 1 specific task** (separation of concerns)
- **Being as small as possible** (micro application)
- **Containerized - running in its own container**
- **Running on micro machines** (0.1 CPU, 256 MB RAM)
- **Must be built for regular crashes & restarts**
- **Fast startup time**
- **Low memory consumption**
- **Has its own database**



## Challenges with Traditional Java Application Server

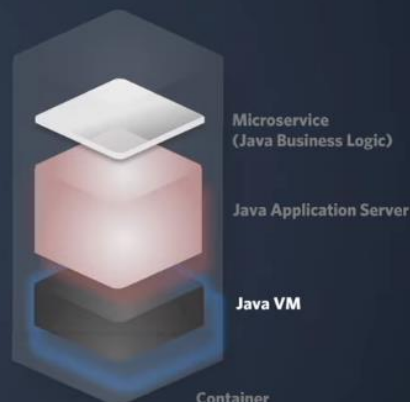
- **Built for running big monolithic Java applications**
- **Running multiple Java applications**
- **Sharing resources**
- **Slow startup time**
- **High memory consumption**
- **Heavyweight**



Should each microservice has its own application server?

## Challenges with Traditional Java VM

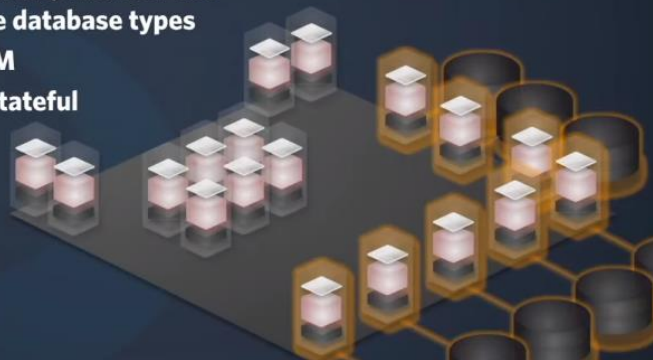
- **Built to run for long time periods**
- **Becomes extremely fast, but only after JIT optimization intervals** (Warm-up)
- **Slow startup time**
- **High memory consumption**



Startup time is too slow for microservices.

## Challenges with Traditional Java Persistence

- Each microservice that stores data has its own database
- The database should fit best to the specific domain model of the microservice. Thus, microservice infrastructure needs multiple database types
- Each microservice uses ORM
- Database connections are stateful
- High devops effort
- DBaaS is expensive

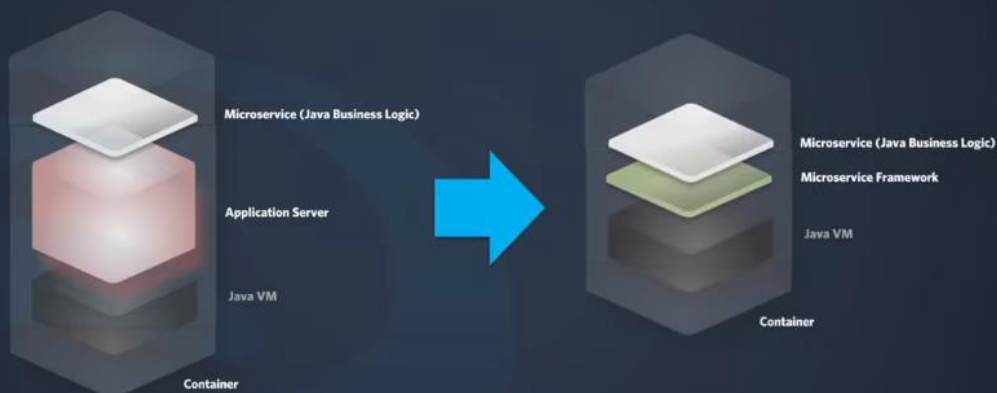


## From Java Application Servers to Microservice Frameworks / Micro Runtimes

### From Application Servers ..



### to Micro Runtimes (Microservice Framework)



# From Application Servers to Micro Runtimes

## Evolution of Java Runtimes – From Full Stack ...



## Evolution of Java Runtimes – From Full Stack to Microservice Frameworks



# From JVM to Native Images

## GraalVM™



Polyglot Programming



Native Images



Startup Time



Memory Footprint

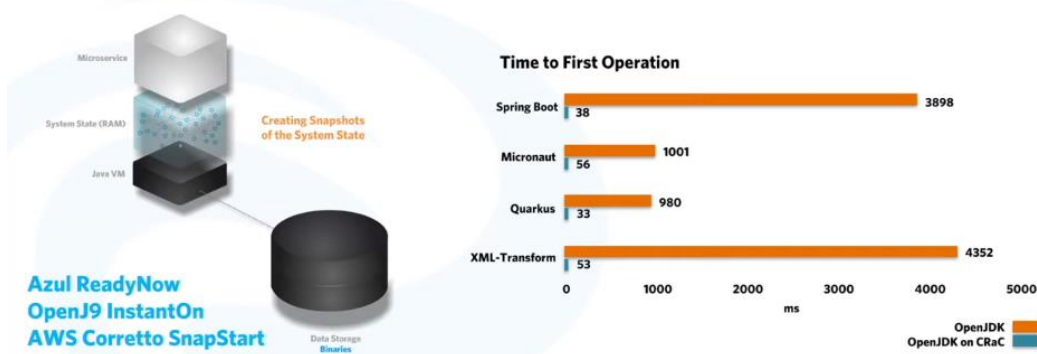


## Limitations to GraalVM Native Execution

- No JVMTI, Java Agents, JMX, JFR support
- Efficient only for smaller heap
- Reflection usage must be known at build time
- Ahead-of-time compiling is linked to longer build times
- Target platform must be known at build time
- Performance - JIT is faster than a native image
- No thread & heap dump support

## OpenJDK CRaC

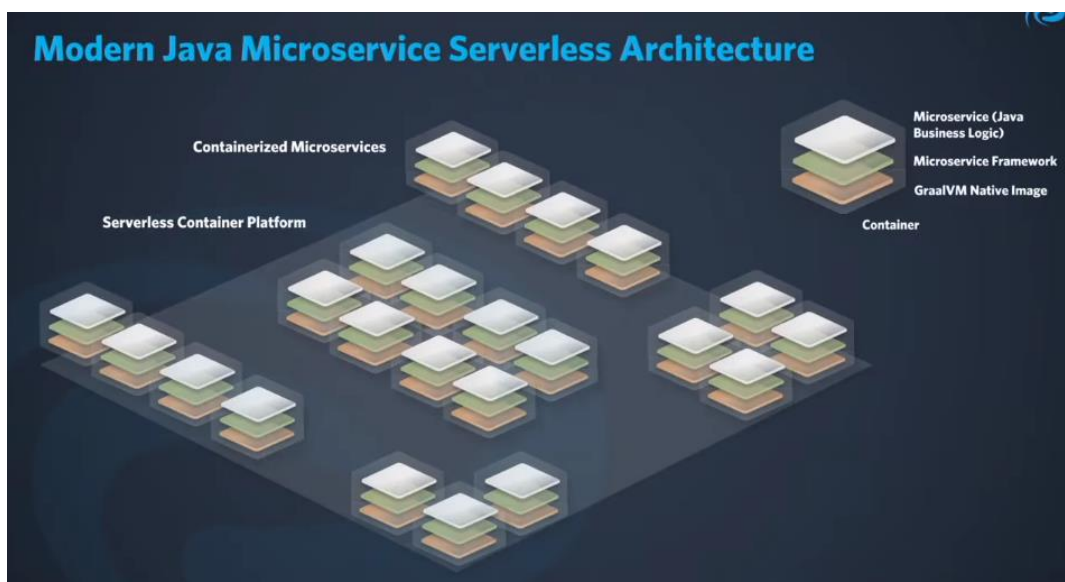
The CRaC (Coordinated Restore at Checkpoint) Project researches coordination of Java programs with mechanisms to checkpoint (make an image of, snapshot) a Java instance while it is executing. Restoring from the image could be a solution to some of the problems with the start-up and warm-up times. The primary aim of the Project is to develop a new standard mechanism-agnostic API to notify Java programs about the checkpoint and restore events. Other research activities will include, but will not be limited to, integration with existing checkpoint/restore mechanisms and development of new ones, changes to JVM and JDK to make images smaller and ensure they are correct.



This creates and stores a snapshot of your warmed up JVM (native or non-native) to allow faster startup times later

## Serverless Infrastructure ...

### Serverless Computing



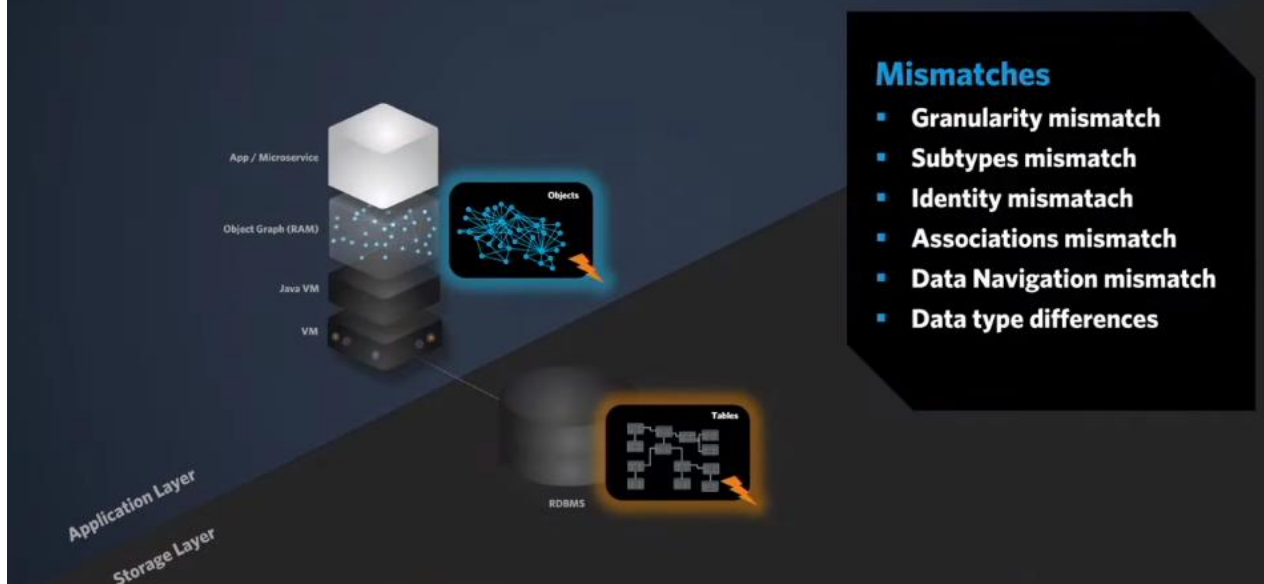
## Serverless Functions (e.g. AWS Lambda)

- Serverless computing platform
- Event-driven architecture
- Fully automated management of computer resources required by code
- No infrastructure configuration
- Supported programming languages: NodeJS, Python, Go, Ruby, C#, Java
- No DevOps effort
- Linux container, 128 MB - 10 GB RAM, 512 MB - 10 GB storage
- Pay only for what you use, price is based on RAM size (GB) and execution time in ms
- Idle resources are not charged
- 1M requests per month, 400,000 GB-seconds of compute time per month

## Java Persistence ...

### From Traditional Java Persistence to Micro Persistence

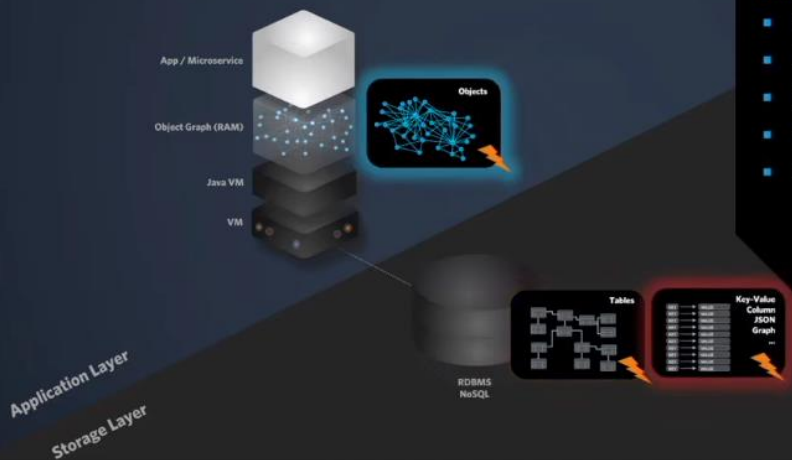
#### Java Persistence - Object-Relational Impedance Mismatch



## Java Persistence – NoSQL Database

### Mismatches

- Column Store
- JSON
- XML
- Graph (Proprietary)
- OO (Proprietary)

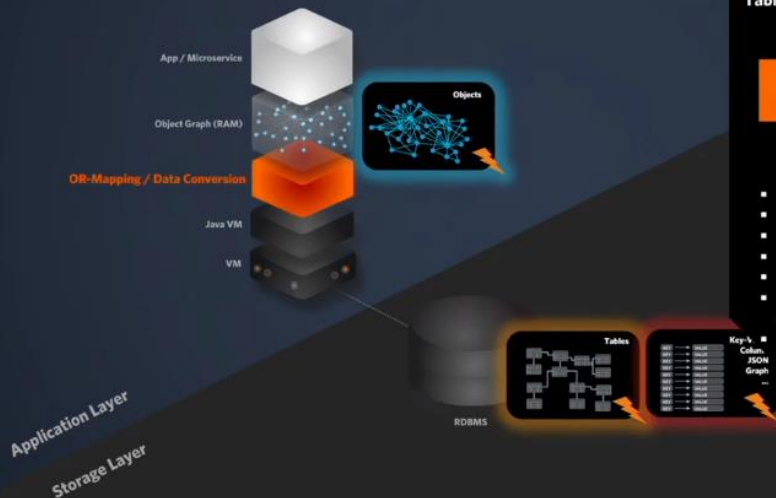


## Java Persistence – Object-Relational Mapping / Data Conversion

Challenge: Storing Objects into  
Tables / JSON / Key Value Stores / Graphs

**Data Conversion Through  
Every Single Read & Write !**

- Requires lots of CPU power
- Reduces your performance
- Expensive latencies
- Complex architecture
- Expensive development process
- Inefficient concept requires expensive cluster infrastructure
- Increase your costs of infrastructure



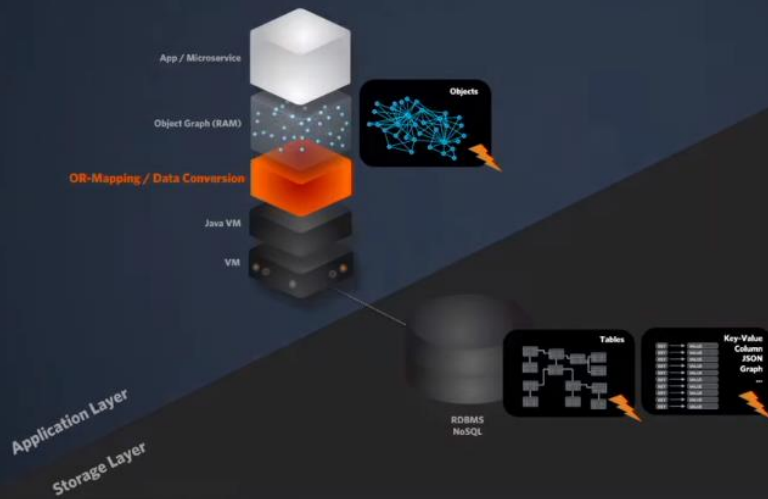
**Millisecond  
Query Time**

## The Problem of Incompatible Data Structures is Well Known as **Impedance Mismatch**

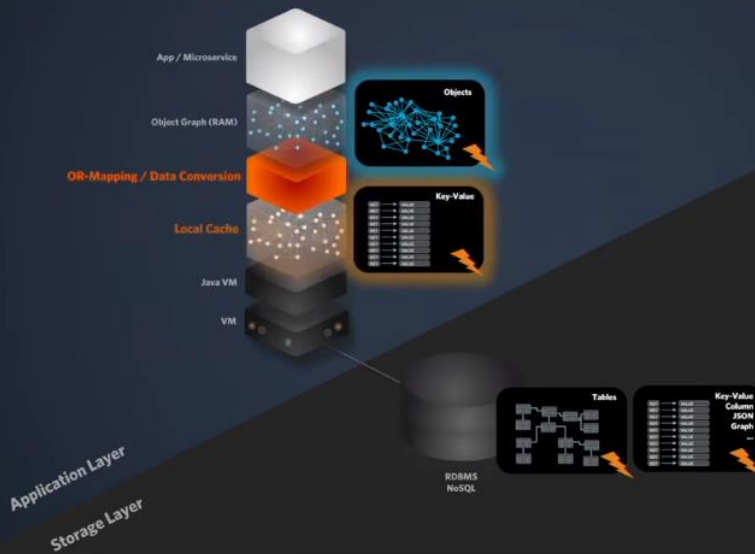


There are various **solutions**, but they are **only a more or less elegant way around the problem**. No matter which solution you choose - **as long as the systems are different**, every developer will **sooner or later** get to the point where his **solution no longer meets** one or more of the following points: **Maintainability, performance, intelligibility.**

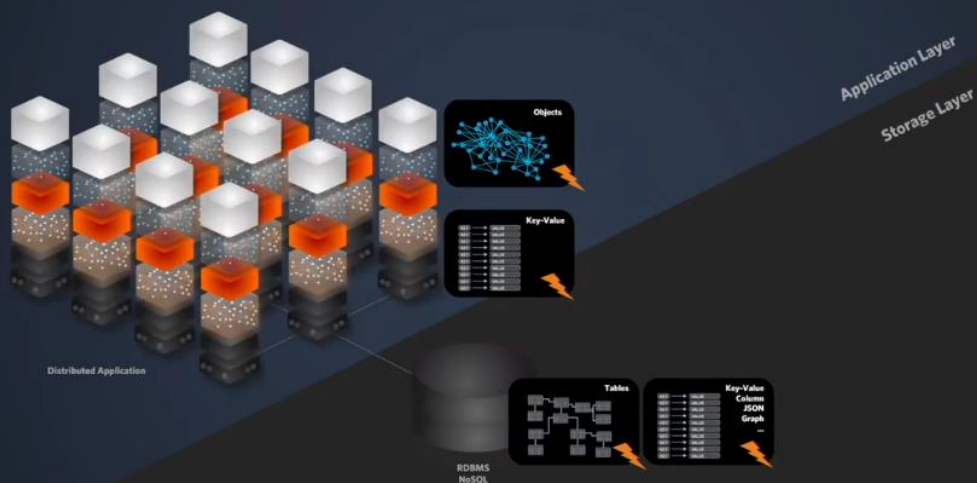
## Java Persistence



## Local Cache

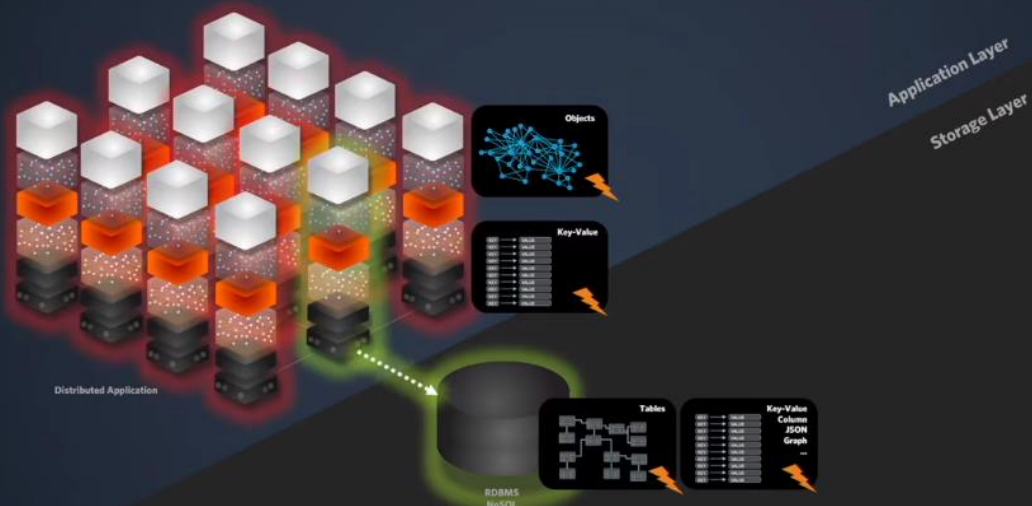


## Distributed Application with Local Cache

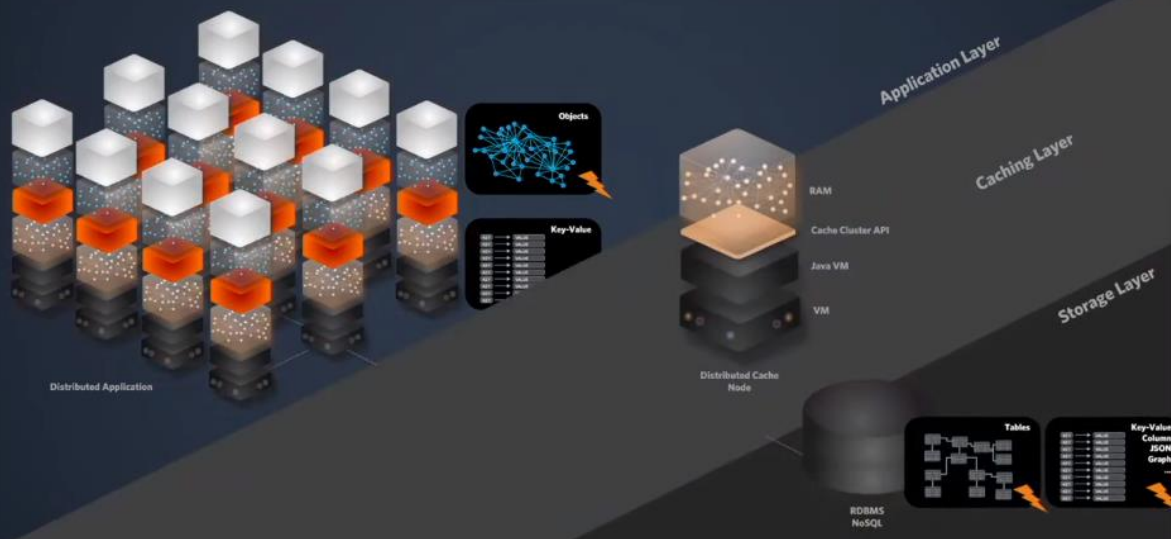




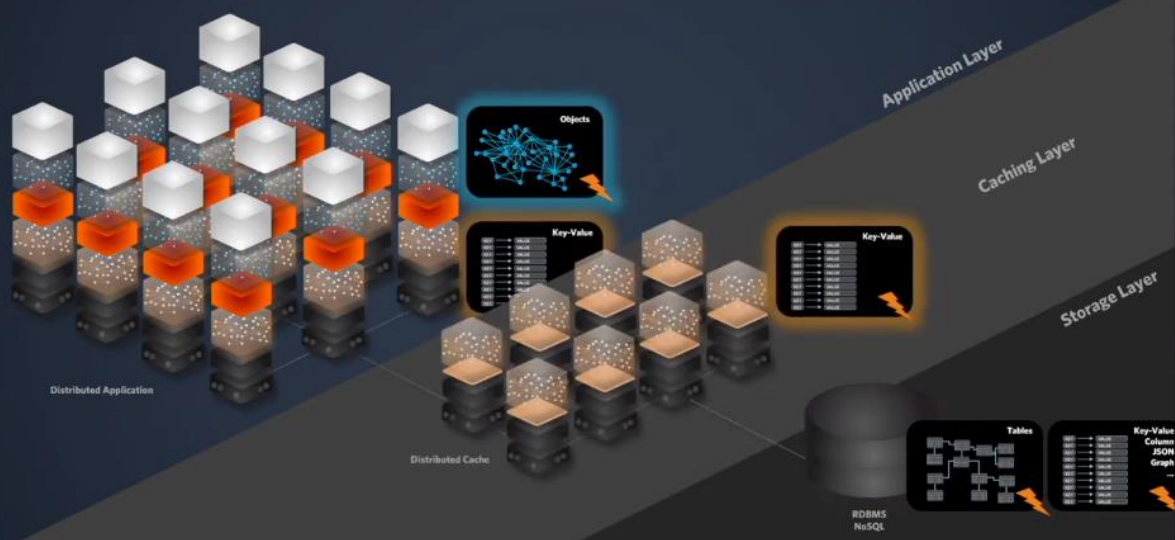
## Distributed Application with Local Cache

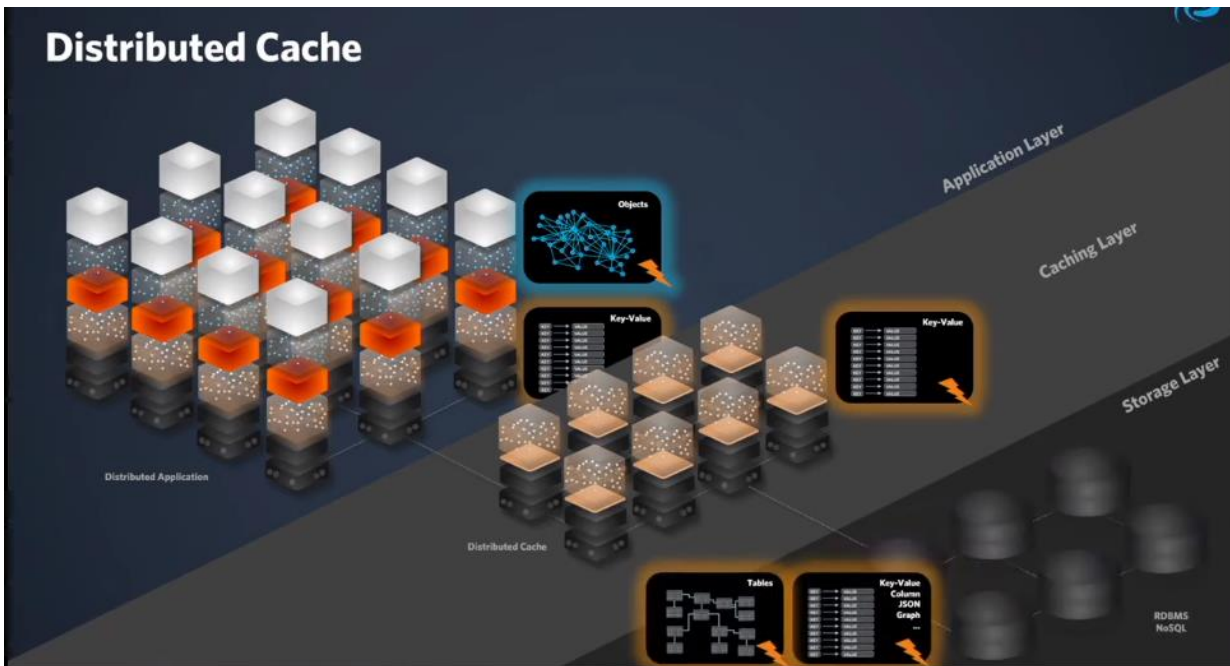


## Distributed Cache



## Distributed Cache





## Java In-Memory Data Processing

### Latency Numbers Every Programmer Should Know

Latency Comparison Numbers (~2012)

L1 cache reference	0.5 ns		
Branch mispredict	5 ns		
L2 cache reference	7 ns	14x L1 cache	
Mutex lock/unlock	25 ns		
Main memory reference	100 ns	20x L2 cache, 200x L1 cache	
Compress 1K bytes with Zippy	3,000 ns	3 us	
Send 1K bytes over 1 Gbps network	10,000 ns	10 us	
Read 4K randomly from SSD*	150,000 ns	150 us	~1GB/sec SSD
Read 1 MB sequentially from memory	250,000 ns	250 us	
Round trip within same datacenter	500,000 ns	500 us	
Read 1 MB sequentially from SSD*	1,000,000 ns	1,000 us	1 ms ~1GB/sec SSD, 4X memory
Disk seek	10,000,000 ns	10,000 us	10 ms 20x datacenter roundtrip
Read 1 MB sequentially from disk	20,000,000 ns	20,000 us	20 ms 80x memory, 20X SSD
Send packet CA->Netherlands->CA	150,000,000 ns	150,000 us	150 ms

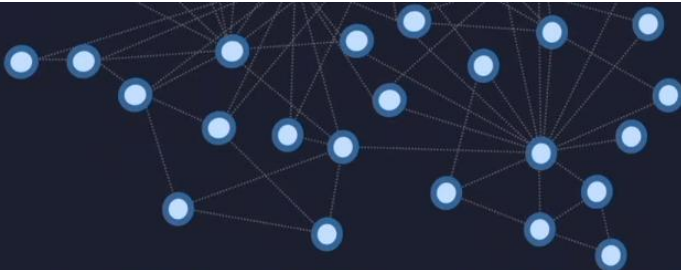
Source: <https://gist.github.com/jboner/2841832>





# GraalVM™

**GraalVM Enables to Compile  
Native Executables**



**Java's Object Graph Model is a  
Multi-Model Data Structure**

```
public static void booksByAuthor()
{
    final Map<Author, List<Book>> booksByAuthor =
        ReadMeCorp.data().books().stream()
            .collect(groupingBy(book -> book.author()));

    booksByAuthor.entrySet().forEach(e -> {
        System.out.println(e.getKey().name());
        e.getValue().forEach(book -> {
            System.out.print('\t');
            System.out.println(book.title());
        });
    });
}
```



**Searching & Filtering with  
Java Streams API**



**1000x**

**All Databases: Millisecond Query Time**  
**Java Streams API: Microsecond Query Time**



0 | 0 0  
0 0 0 0  
| 0 | |

Cloud BLOB  
Storages  
(AWS S3)

High-available, globally replication, auto backups, high security, fully managed, 99% cheaper than DBaaS.

**System Prevalence** is the Simplest and Fastest Way to Provide ACID Persistence For Java Objects



System prevalence is an architectural pattern that combines system images (snapshots) and transaction journaling to provide speed, performance scalability, transparent persistence and transparent live mirroring of computer system state.

In a prevalent system, state is kept in memory in native format, all transactions are journaled and system images are regularly saved to disk.

 MicroStream

MicroStream is a Java-native micro persistence layer that runs within your Java application or microservice. It enables you to store any object graph of any size and complexity or individual subgraphs persistently into any data storage. Vice versa MicroStream allows you to load any subgraph dynamically and restore it in RAM at any time. Loaded object-references are merged into the object graph in RAM automatically.

MicroStream is a tiny open-source Java library you can download via Maven. The code is available on Github.

[www.microstream.one](http://www.microstream.one)

It just serializes your Java data object graphs in-memory and into a blob store like S3 which you can search with the Java Streams API and also restore it easily when needed.



Product
Solutions
Open Source
Pricing

microstream-one / microstream
Public

master
22 branches
25 tags

Go to file
Code

about

High-Performance Java-Native-Persistence. Store and load any Java Object Graph or Subgraphs partially. Relieved of Heavy-weight JPA. Microsecond Response Time. Ultra-High Throughput. Minimum of Latencies. Create Ultra-Fast in-Memory Database Applications & Microservices.

microstream-one/

java
storage-engine
serializer
cache
persistence
in-memory-storage
in-memory-database
object-graph

Readme
EPL-2.0 license
Code of conduct
356 stars
9 watching
31 forks

Releases 30

07.00.00-MS-GA
on 28 Apr
9 releases

Contributors 11

zdenek-jonas	setup javac as default (#423)	7 days ago	2,578 commits
github	move CDI tests to integration tests + add github check (#404)	last month	
jmuh	Remove some warnings (#271)	11 months ago	
af5	closing file after deleting (#408)	27 days ago	
base	setup javac as default (#423)	7 days ago	
cache	Fix cache expiry handling with read through (#368)	4 months ago	
communication	use read/write methods with timeout	4 months ago	
configuration	Adapt to javac (#263)	4 months ago	
docs	Update housekeeping doc (#416)	14 days ago	
etc	Add readme	17 months ago	
examples	Fix usage of StorageManager MicroProfile Config converter and assign ...	last month	
integrations	setup javac as default (#423)	7 days ago	
is2	new dev version 6 (#350)	5 months ago	
persistence	GC Live Object Check (#402)	last month	
storage	GC Live Object Check (#402)	last month	
gitignore	License plugin update	14 months ago	
CODE_OF_CONDUCT.md	Create CODE_OF_CONDUCT.md	14 months ago	
CONTRIBUTING.md	add info about signed commits to CONTRIBUTING.md (#412)	21 days ago	
LICENSE	Update through build	14 months ago	
README.md	Java 11 (#358)	4 months ago	
pom.xml	setup javac as default (#423)	7 days ago	

## Maven Download

```

pom.xml
1 <repositories>
2   <repository>
3     <id>microstream-releases</id>
4     <url>https://repo.microstream.one/repository/maven-public</url>
5   </repository>
6 </repositories>
7 <dependencies>
8   <dependency>
9     <groupId>one.microstream</groupId>
10    <artifactId>storage.embedded</artifactId>
11    <version>04.01.00-MS-GA</version>
12  </dependency>
13 </dependencies>
14 <dependency>
15   <groupId>one.microstream</groupId>
16   <artifactId>storage.embedded.configuration</artifactId>
17   <version>04.01.00-MS-GA</version>
18 </dependency>
19 </dependencies>

```

## MicroStream is an Integral Part of Helidon & Micronaut

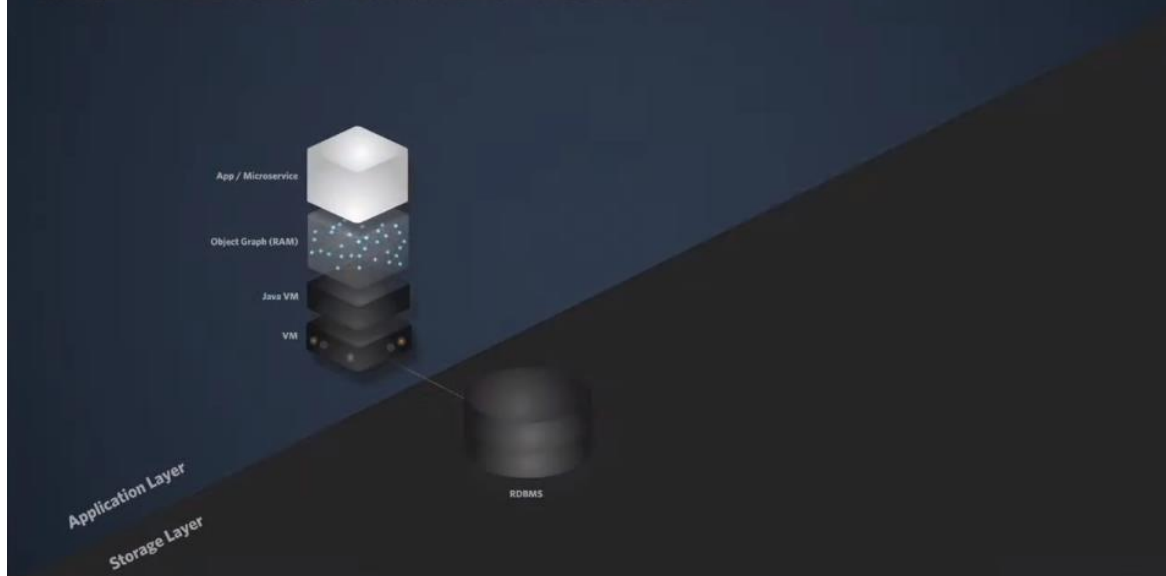
MicroStream is proud to be an integral part of and contribute to leading microservice frameworks. We are happy to deliver value to the Java Open Source community.

First-class support for MicroStream has been added to the Micronaut framework, and MicroStream is now a major sponsor of the Micronaut Open Source project.

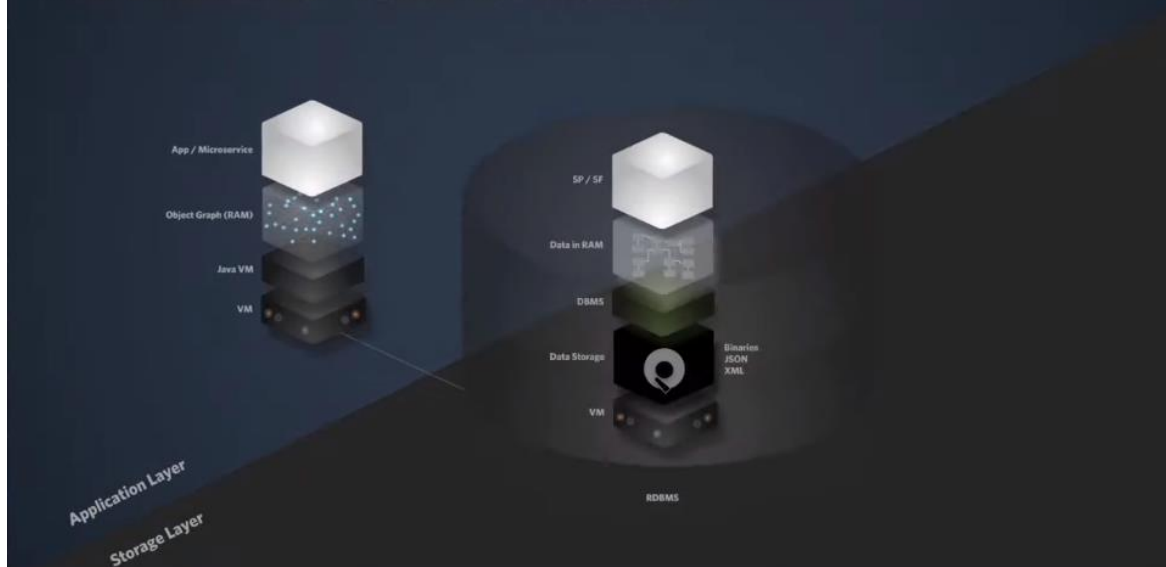
A seamless MicroStream integration has been added to both Helidon flavours, Helidon MP and SE. Helidon is a leading microservice framework for Java, powered by Oracle.

Our Partners

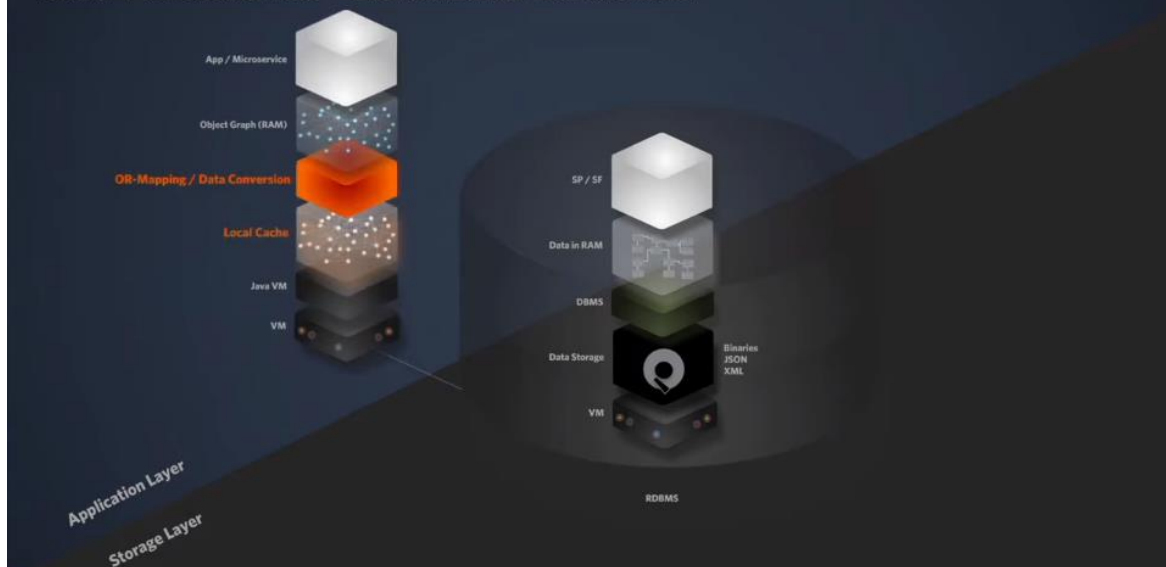
## Java Persistence - Relational Database



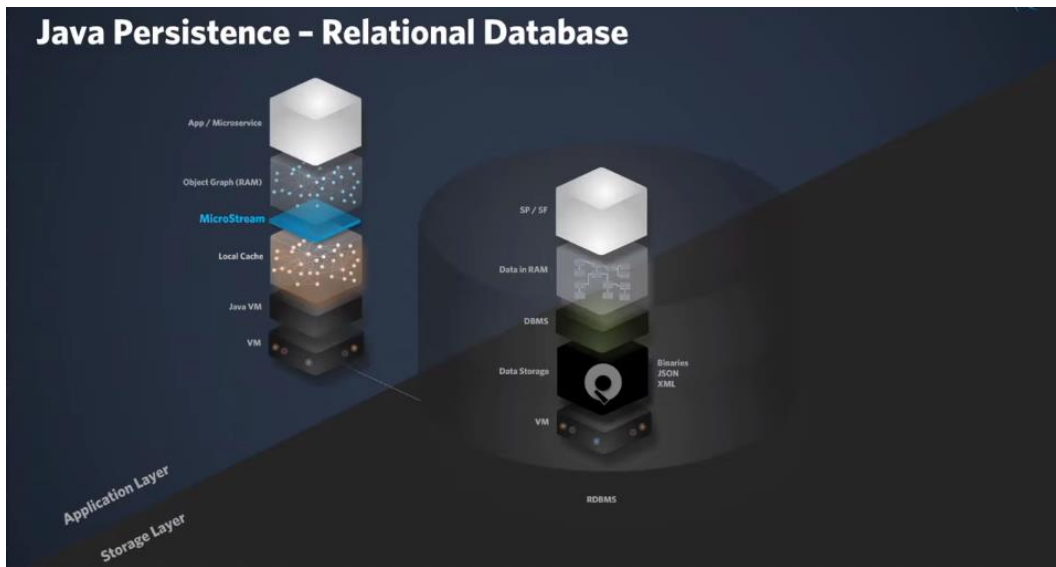
## Java Persistence - Relational Database



## Java Persistence - Relational Database

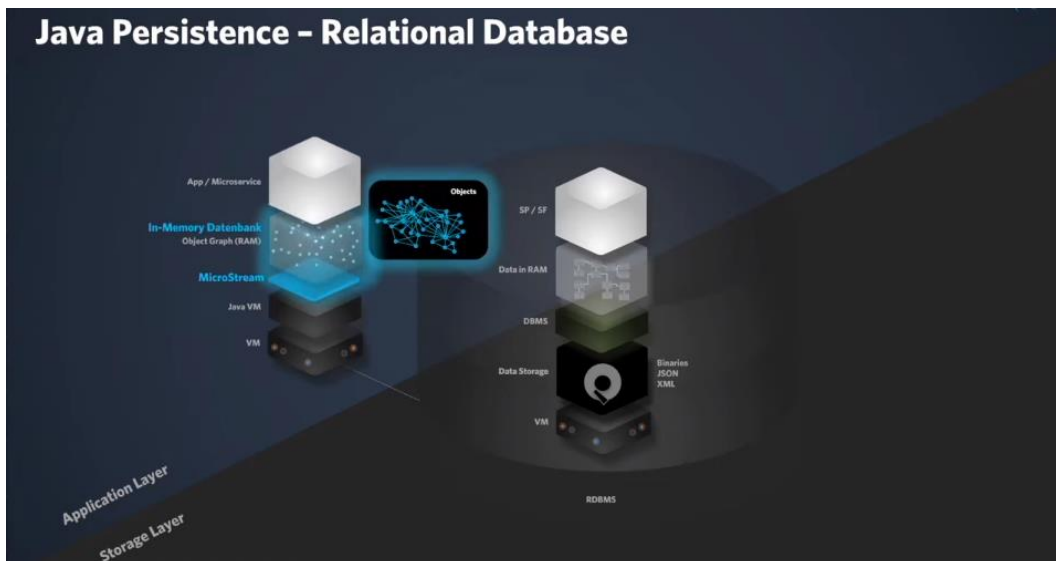


## Java Persistence - Relational Database



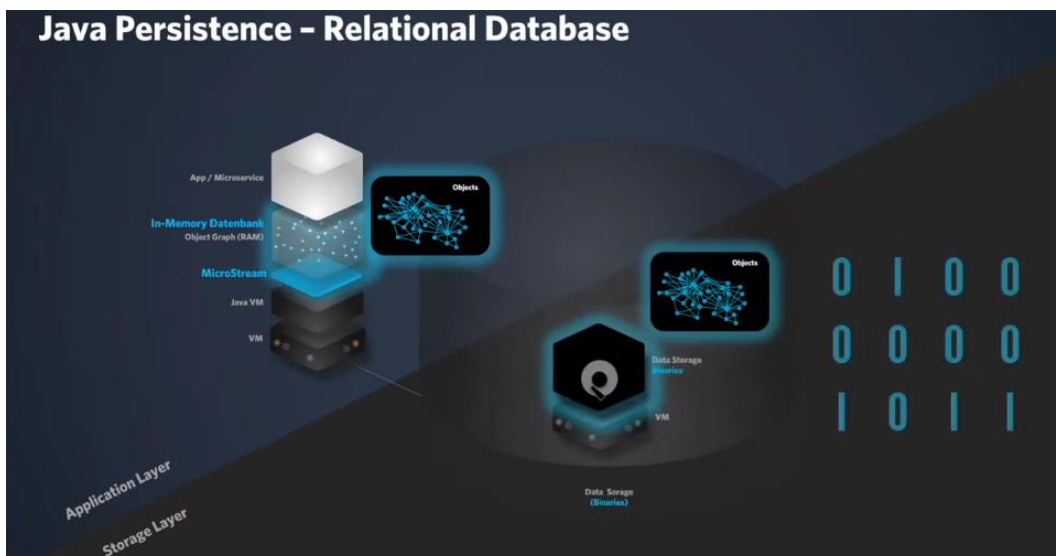
We have replaced object relational mapping with just serialization with **MicroStream**

## Java Persistence - Relational Database



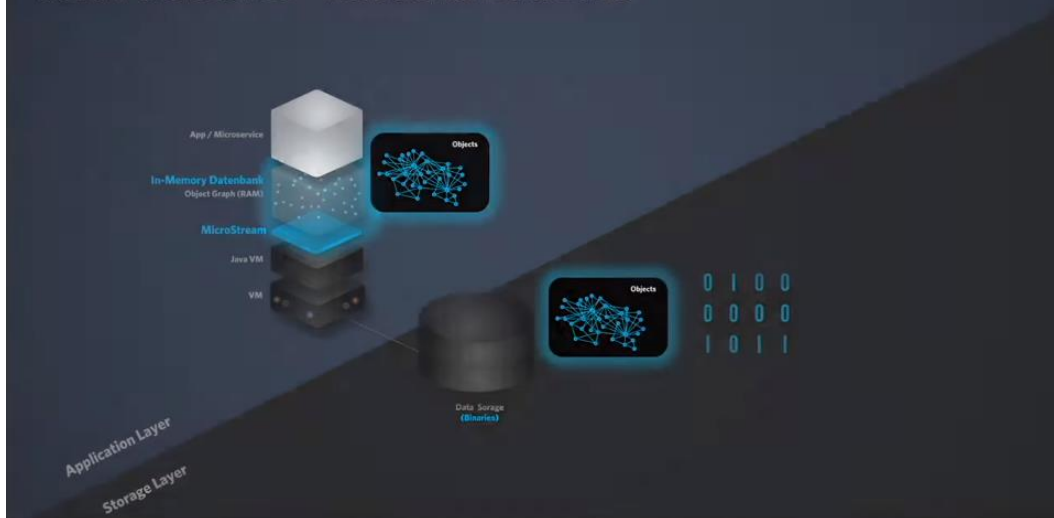
We then keep the data in-memory as a Java native in-memory database (this eliminates the need for a data cache!).

## Java Persistence - Relational Database



We now do not need data on the persistence side anymore but just replace it with the storage.

## Java Persistence - Relational Database



## Supported Storages



## Runs Wherever Java Runs

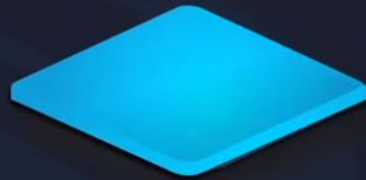


## MicroStream Components





# MicroStream



MicroStream Micro Persistence

## Accelerating Queries up to 1000x

Query: Revenue of the whole shop

JPA - Hibernate (Java Standard)	MicroStream	Factor
<b>439.05</b> Milliseconds 2.28 Queries / Second	<b>0.19</b> Milliseconds 190.11 Queries / Second	<b>2260x</b> 83x Queries / Second
Persistence: <b>Hibernate</b> Cache: <b>EHCache</b> Database: <b>Oracle DB</b>	Persistence: <b>MicroStream</b> Cache: <b>-</b> Database: <b>Oracle NoSQL</b>	

439.05 ms

0.19 ms

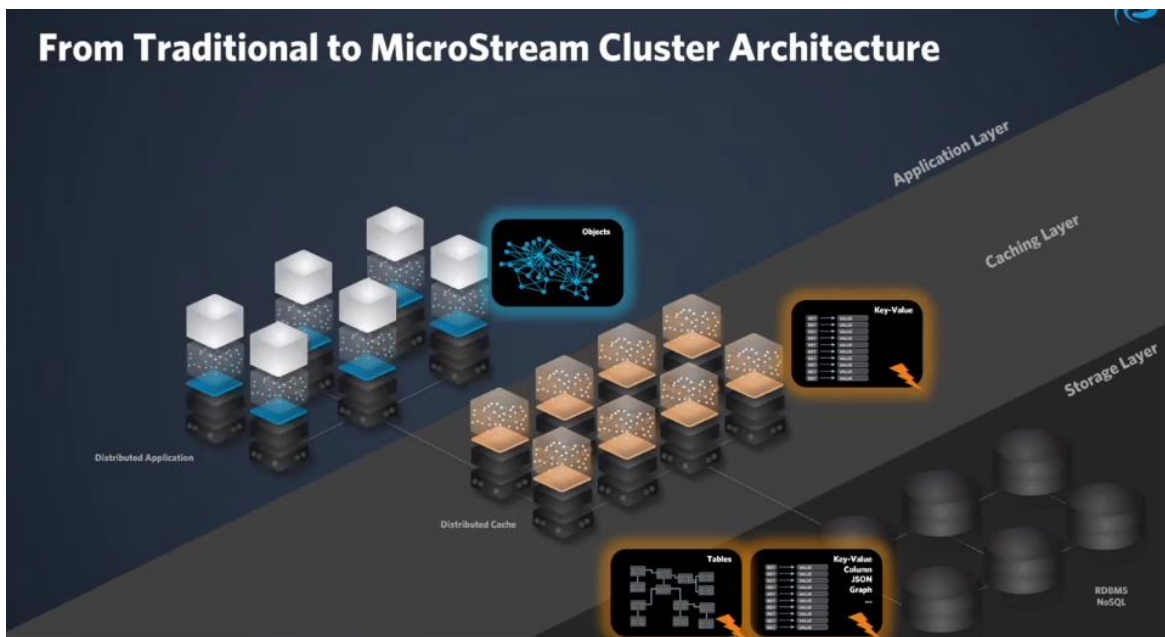
Live-Demo: [www.microstream.one](http://www.microstream.one)

## MicroStream Cluster

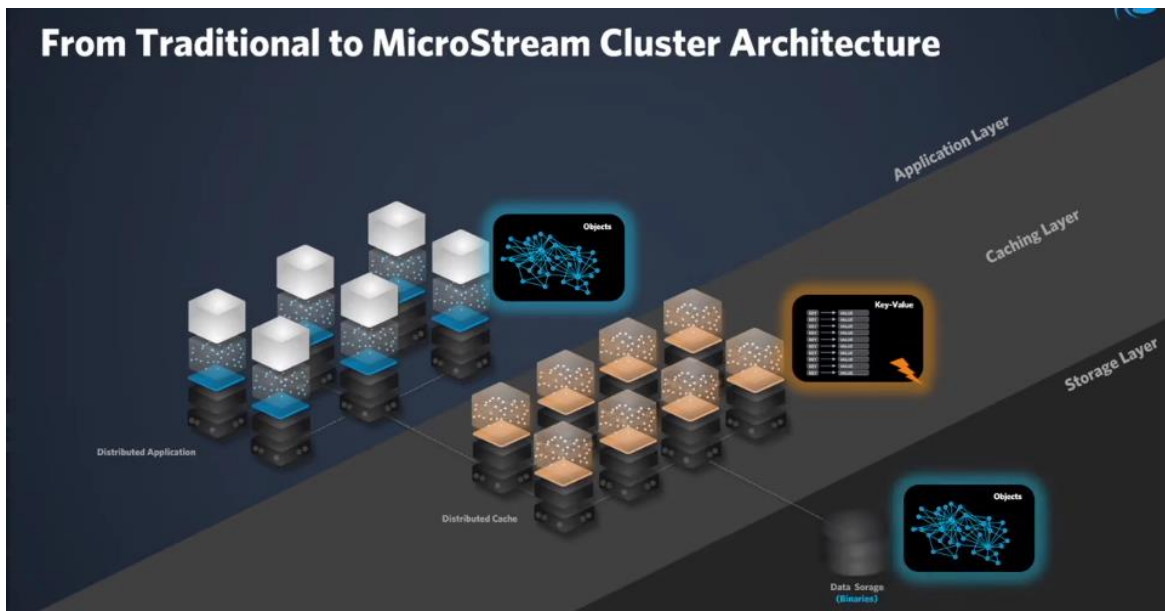
### From Traditional to MicroStream Cluster Architecture



## From Traditional to MicroStream Cluster Architecture



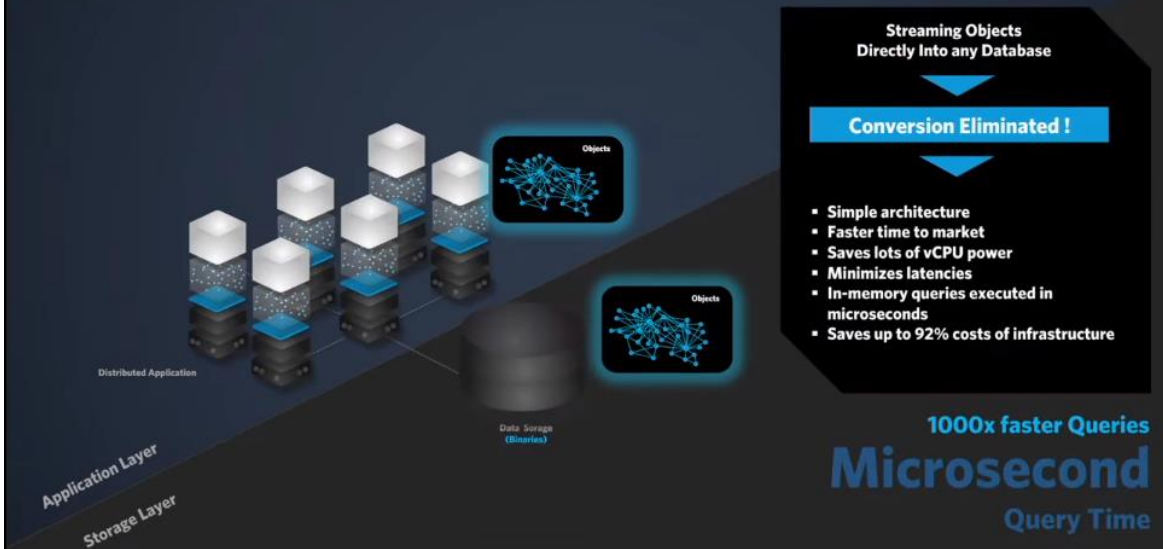
## From Traditional to MicroStream Cluster Architecture



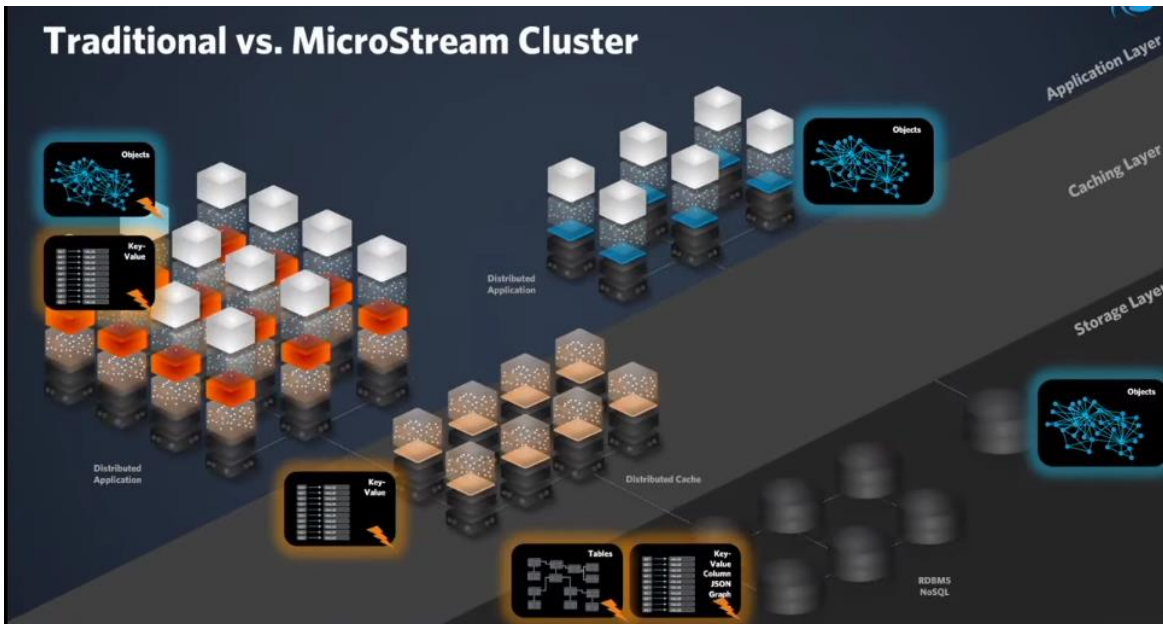
## From Traditional to MicroStream Cluster Architecture



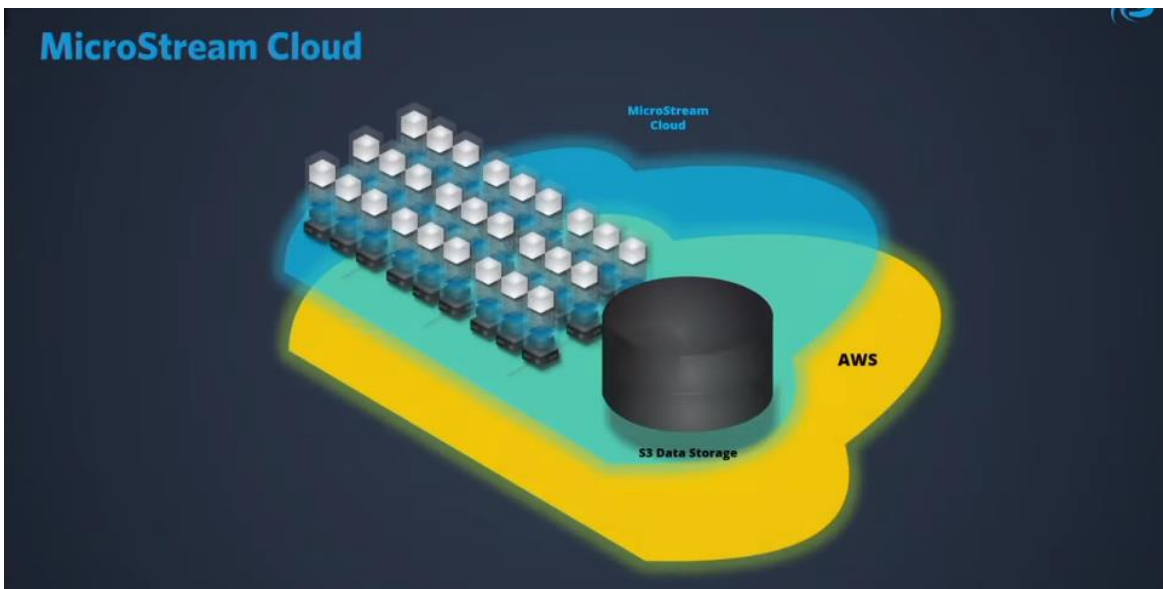
## MicroStream Cluster Architecture



## Traditional vs. MicroStream Cluster



## MicroStream Cloud





## Benefits - Development

- Simplest possible architecture & implementation
- No mappings or data conversion
- No specific requirements for your classes, just use POJOs
- No specific client API required
- No specific query language
- Core Java concepts only
- Get a clustered app / microservice out-of-the-box

## Benefits - Performance


- Ultra-fast in-memory data processing
- Microsecond query time
- Low-latency realtime responsiveness
- Gigantic workloads & throughput

## Benefits - Cost Savings

- 90% CPU consumption
- 90% energy
- 90% CO2 emission
- 99% cloud storage costs
- 90% cloud costs in total

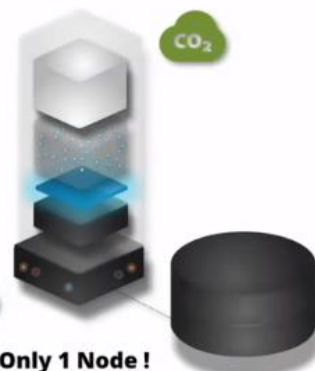
## Our Customers Save up to 89% CPU Power, Cloud Costs & CO2 Emissions

Example:

Allianz 

CPU / RAM / Storage  
**\$145k**  
Annually

Today with MicroStream:



Only 1 Node !

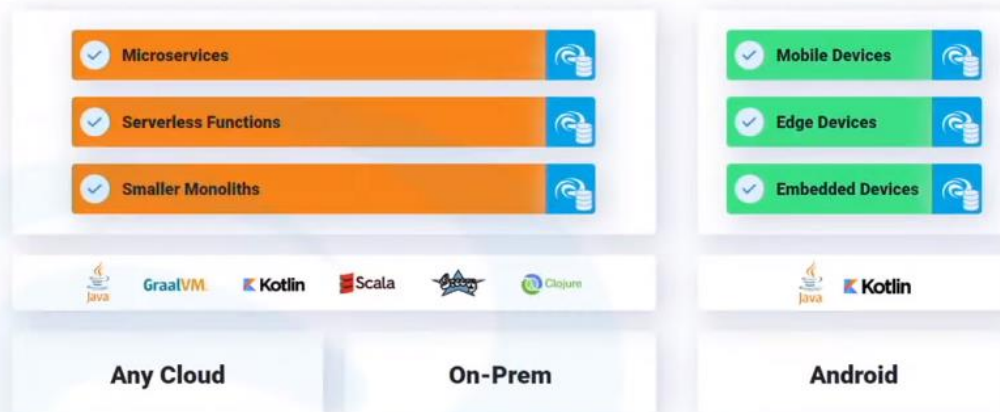
CPU / RAM / Storage  
**\$16k**  
Annually

**- 89%**  
Costs of Infrastructure  
annually

Conventional cluster for running a globally app



## Use-Cases



## Challenges with MicroStream

- Built for Java developers
- Paradigm shift in database programming
- No SQL support
- MicroStream is a storage engine, but not a DBMS
  - Your application must cover DBMS tasks
  - You must care vor validation
  - You have to care for concurrency
- Not suited for DBAs

## Challenges with MicroStream Cluster

- Creating cluster infrastructure is effortful
- DevOps resources required
- Konfiguration von Cluster-Nodes
- Kubernetes know-how required
- Kafka know-how required
- High effort for development & testing
- Monitoring
- No support

# Your Object Graph in RAM is Your Database

## Traditional Database Server Paradigm



## Java In-Memory Data Processing Paradigm



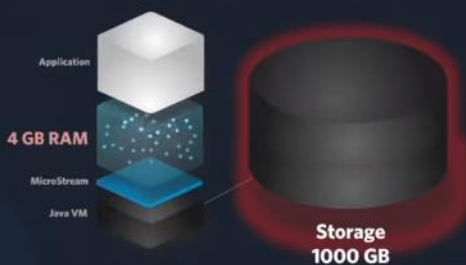
## All In-Memory or Lazy-Loading ?

### Enough RAM available:



- Load your entire DB into RAM
- Pure in-memory computing
- No latencies
- Super fast
- Lower startup time

### Data Storage is bigger than RAM:



- Preload most important data only (eager loading)
- Use lazy-loading to load data on demand only
- Clear lazy references which are not used anymore
- Faster startup time

## Get Started with MicroStream

Learn: [www.microstream.one](http://www.microstream.one)

GitHub: <https://github.com/microstream-one/microstream>

Doc: <https://manual.docs.microstream.one/data-store/getting-started>

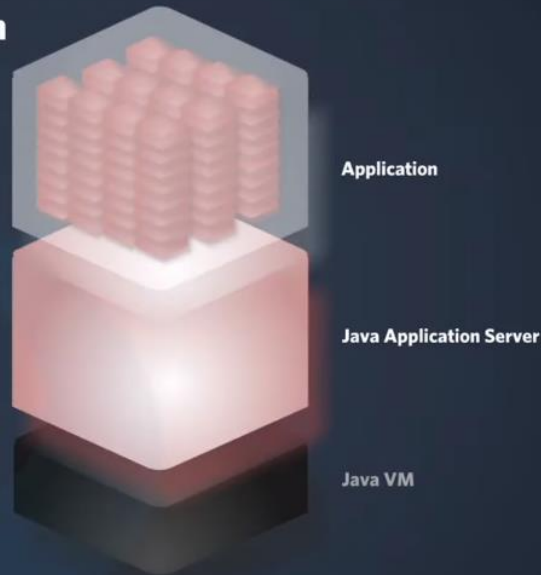
Videos on YouTube: <https://www.youtube.com/c/MicroStream/videos>

Free MicroStream training: <https://www.javapro.io/training>

## Conclusion

### Monolithic Java Application

One big piece of software runs on one big single machine.



### Modern Cloud-Native Microservice & Serverless Architecture

