

**AWS
re:Invent**

DEV301

Amazon CloudWatch Logs and AWS Lambda

A Match Made in Heaven

Bob Wilkinson, AWS

Robert Waugh, AWS

December 1, 2016

What to Expect from the Session



**Look at industry trends
impacting monitoring**



**Learn about CloudWatch
and CloudWatch Logs**



**Understand several key
monitoring use cases**



**See CloudWatch and
Lambda in action**

What to Expect: Scenario Preview

Centralize

Centralize logs from Elastic Load Balancing (ELB) using S3 bucket triggers



Customize

Customize alarms from CloudWatch to fit your specific needs



Analyze

Build an on-demand, scalable Elasticsearch cluster to solve a specific problem or to do analysis



Recognize This?



This dog has a monitoring problem

Day in life!

The story you about to hear is true (mostly)... Only the names have been changed to protect the innocent...

A customer writes a high severity ticket – your Application, ImportantApp, is down

John, the on-call developer, is paged through the ticketing system

None of your alarms fired

Blissful Ignorance

John engages and starts to scan service dashboards

He does see intermittent availability impact, but doesn't know how to assess impact to customers or where to begin troubleshooting

He decides to escalate to a manager on-call

Confusion

More customer tickets are pouring in

An escalation manager, Jane, joins the event and starts to assess the situation and impact

John and Jane's CTO happens to notice the problem.
Sends Jane an IM – "Jane, what's going on with ImportantApp?"

Stress

Jane and John recall a recent issue where certain customers started to issue "expensive" operations

John starts log diving on their production hosts

John identifies a suspect customer. Jane cuts a ticket and John prepares a configuration change to block the customer

False Hope

The other team engages and indicates they didn't change anything

Jane and John also confirm this when the availability impact persists after deploying the configuration change

Out of ideas, John suggests to fail over to the standby – “It can't hurt...”

Desperation

After the failover, ImportantApp recovers (Yay!)

Our root cause deep dive finds that a new JDBC version introduced a memory leak leading to Java heap exhaustion

We fix the leak, add new alarms on memory usage, and tune our service alarms

Enlightenment

Can we do better than that?

Day in the Life - Reflection

- We have missing alarms and some of the alarms we have are not actionable
- We do not always have the right logs and interacting with them can be tedious
- Our dashboards do not tell us enough about customer impact or behavior changes

Monitoring is really (really) hard!

Trends in Monitoring

Trend: Complexity Increasing

- Distributed **micro-services** based applications
- Applications are written in **different languages and frameworks**
- Workloads are increasingly running on transient resources such as **containers** and serverless compute
- Specialization in **persistence tier**

Trend: Applications are More Dynamic

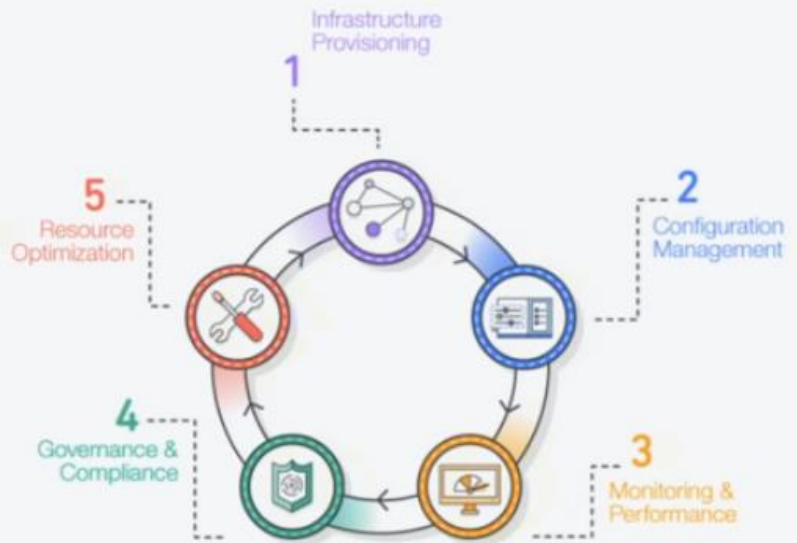
- Small changes are **continuously built**, tested, and **deployed**
- As the scale and design of applications are changing rapidly, so are the **infrastructure** needs
- Applications are **global** and customer behavior is **unpredictable**
- Increased role of **automation**

Trend: More Business Impact

- Increased role of applications in **business outcomes** (Revenue, Cost, SLA)
- Rapidly evolving applications are required to gain **competitive advantage**
- Increased **expectations** from customers

Trend: Integrated Management

- Monitoring is **no longer standalone**. It is one aspect of management and increasingly **integrated** into the lifecycle of application



See

React

Diagnose

Resolve

CloudWatch

CloudWatch (**CF**) is a portfolio of tools covering metrics, alarms, events, dashboards, etc.

See

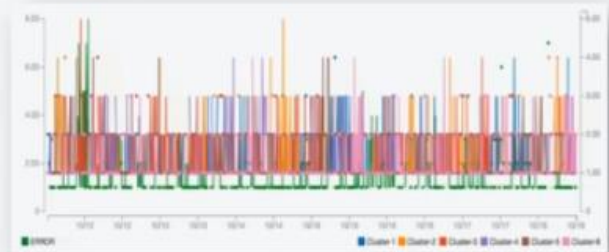
React

Diagnose

Resolve

Inspect, navigate, zoom, and correlate across time to investigate issues

Jump to your logs directly from your metrics to perform searches or generate additional metrics from log data



Time (UTC +00:00)	Message
2016-10-11	
19:48:30	2016-10-11T19:48:30.551Z b132298b-89b-11e6-aadd-3106b0d4821 ERROR Cluster0
19:48:30	2016-10-11T19:48:30.551Z b132298b-89b-11e6-aadd-3106b0d4821 undefined
19:48:30	END Request: b132298b-89b-11e6-aadd-3106b0d4821
19:48:30	REPORT Request: b132298b-89b-11e6-aadd-3106b0d4821 Duration: 41.50 ms Billed Duration: 1
19:48:30	Process exited before completing request
19:48:37	START Request: b132298b-89b-11e6-aadd-3106b0d4821 Version: SLATEST
19:48:37	2016-10-11T19:48:37.291Z b132298b-89b-11e6-aadd-3106b0d4821 DEBUG Cluster0
19:48:37	2016-10-11T19:48:37.348Z b132298b-89b-11e6-aadd-3106b0d4821 ERROR Cluster0
19:48:37	2016-10-11T19:48:37.348Z b132298b-89b-11e6-aadd-3106b0d4821 undefined

See

React

Diagnose

Resolve

Easily and automatically correct issues via common actions that you control

Define your own custom actions based on Lambda functions for more fine-grained control

Description: AutoScaling web servers on latency

Whenever: Latency

IN: TO:

For: consecutive periods

Actions

Define what actions are taken when your alarm changes state.

AutoScaling Action [Create](#)

Whenever this alarm: [Edit](#)

From resource type: [Edit](#)

From the: [Edit](#)

Take this action: [Edit](#)

Latency on 70

10.0
7.5
5.0
2.5
0.0

10/10 10/15 10/20 10/25

Homepage: [Edit](#)

Homepage: [Edit](#)

Metric Name: [Edit](#)

Period: [Edit](#)

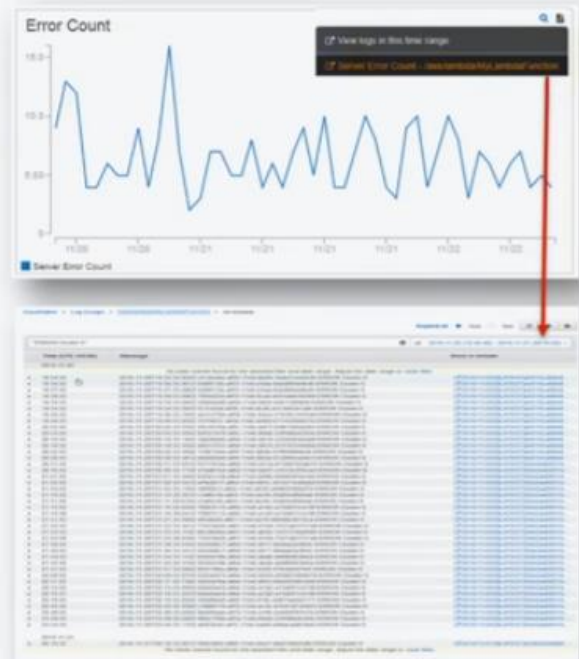
Statistics: [Edit](#)

```
function createEni(subnetId, securityGroups, callback){
  //Create a network interface, pass the interface ID to callback
  var eniCreationParams = {
    "SubnetId": subnetId,
    "Groups": securityGroups
  };
  console.log("DEBUG: CreateEni Params:\n", eniCreationParams);
  ec2.createNetworkInterface(enicreationParams, function createEniCallback(err, data) {
    if (err) {
      console.log("ERROR: ENI creation failed.\nData:\n", err);
      return callback(err, null);
    }
    console.log("INFO: ENI Created.\nData:\n", data);
    return callback(null, data.NetworkInterface.NetworkInterfaceId);
  });
}

function attachNetworkInterface (networkInterfaceId, instanceId, callback){
  //Attaches an ENI, passes the AttachmentId to callback.
  var nic_params = {
    "NetworkInterfaceId": networkInterfaceId,
    "InstanceId": instanceId,
    "DeviceIndex": 1 // Should be safe to assume index 1 is available
  };
}
```

Recent Improvements

- Metrics Price Drop
- More metrics, logs, events from AWS services: CloudTrail, Elastic Beanstalk, SES
- Simple navigation from Metrics to your Logs
- Upgraded metric retention from 2 weeks to up to 15 months



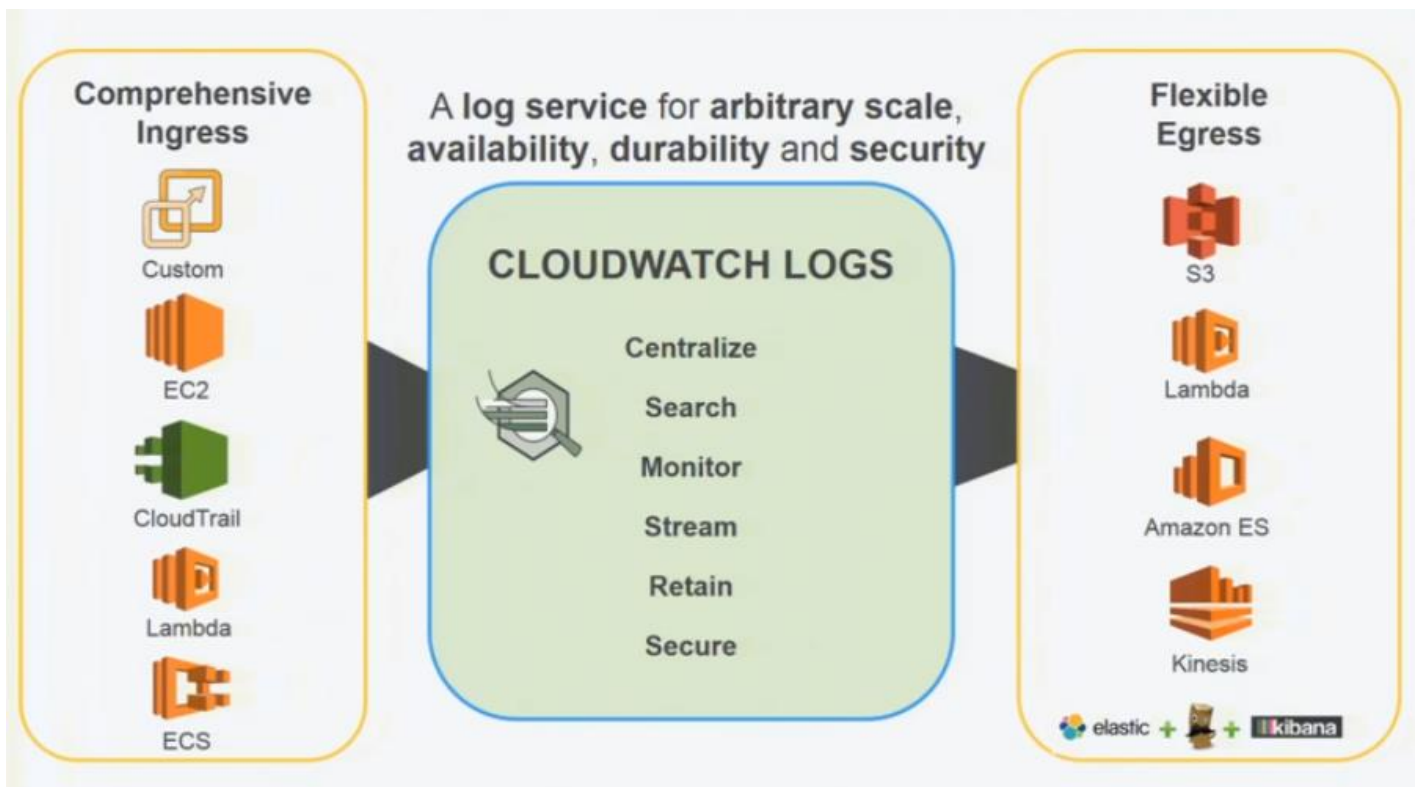
Recent Improvements (continued)

- Support for arbitrary metric percentiles
- collectd output plugin to simplify metric collection
- Improvements in Dashboards (new widgets, dark theme, Y axis limits)
- Improved Logs console experience



Not Just About What's Inside CloudWatch

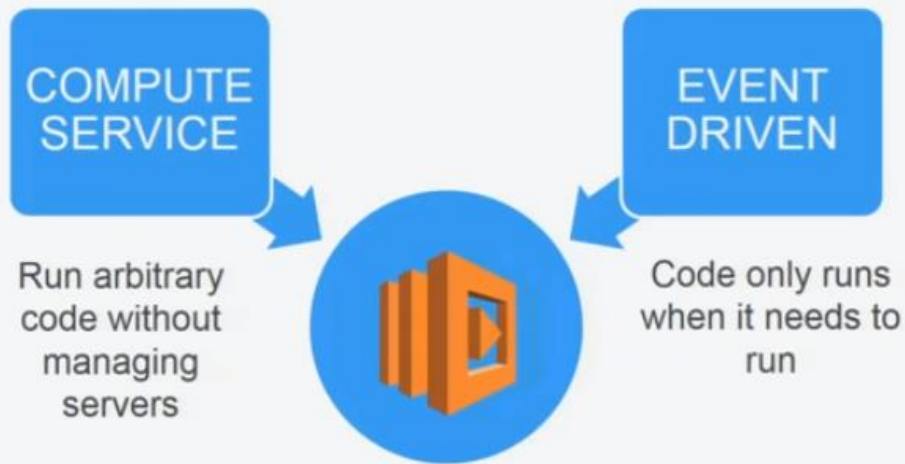
- Monitoring is hard (very hard)
- Every enterprise, team and situation has unique needs
- We have a rich partner ecosystem
- We give you the tools and flexibility to integrate with other AWS services



For the ingress option, we have a log agent that you can install on any host that allows you to point to a log file to ship logs from within your host and stream the logs out in near real-time. You can then run text search on your logs and also extract metrics from your logs. You can take a data element within your logs and have it emitted out as a metric that you can put on a graph, set an alarm on, etc.

On the egress side of logs, you can use your filtered logs to egress things to a Kinesis stream, Elasticsearch, pipe it through a lambda function, do batch export of your logs to an S3 bucket, etc

Serverless Compute: AWS Lambda



Having centralized logs from our instances and buckets in a single repository location with a consistent APIs, retention policy and access controls. You need to think about how to federate all that data into a single place.

Problem Statements

- Log data is scattered on instances and S3 buckets
- It would be better if it were centralized in CloudWatch Logs for searching and filtering
- Today CloudWatch provides an agent for instance logs, what about S3 delivered logs?

Flow of Events



You can configure ELB to send your access logs to an S3 bucket, we then set up an object creation notification event to lambda that will trigger a lambda function to read and publish the logs to the central location

Apache Servers

ELB & Lambda Configured

Lambda Triggered

ELB logs in CloudWatch

Provision three EC2 instances running Apache

EC2 Dashboard

Events

Tags

Reports

Limits

INSTANCES

Instances

Spot Requests

Reserved Instances

Scheduled Instances

Dedicated Hosts

IMAGES

AMIs

Bundle Tasks

ELASTIC BLOCK STORE

Launch Instance

Connect

Actions

search : apache

1 to 5 of 5

Name	elb-logs	Instance ID	Instance Type	Availa
apache-2a		i-0250179af2a646134	t2.micro	us-west
apache-2b		i-00255fd07ca832bd2	t2.micro	us-west
apache-2c		i-0fa0adccbf720442	t2.micro	us-west

Instance: i-00255fd07ca832bd2 (apache-2b)

Public DNS: ec2-54-190-48-183.us-west-2.compute.amazonaws.com

Description

Status Checks

Monitoring

Tags

Instance ID

i-

00255fd07ca832bd2

Public DNS

ec2-54-190-48-183.us-west-2.compute.amazonaws.com

Instance state

running

Public IP

54.190.48.183

Apache Servers

ELB & Lambda Configured

Lambda Triggered

ELB logs in CloudWatch

Create an ELB

EC2 Dashboard

Events

Tags

Reports

Limits

INSTANCES

Instances

Spot Requests

Reserved Instances

Scheduled Instances

Dedicated Hosts

IMAGES

AMIs

Bundle Tasks

ELASTIC BLOCK STORE

Create Load Balancer

Actions

Filter: apache

1 to 1 of 1

Name	DNS name	State	VPC ID	Availability Zones
ApacheLoadBalancer	ApacheLoadBalancer-12727...	active	vpc-b04225d4	us-west-2b, us-west-2a

Load balancer: ApacheLoadBalancer

Description

Listeners

Monitoring

Tags

Basic Configuration

Name: ApacheLoadBalancer

ARN: arn:aws:elasticloadbalancing:us-west-

Creation time: October 11, 2016 at 12:45:47 PM UTC-7

Hosted zone: Z1H1FLSHABSF5

DNS name: ApacheLoadBalancer-1272730108.us-west-2.elb.amazonaws.com (A Record)

VPC: vpc-b04225d4

Scheme: internet-facing

We then set up an ELB to front those 3 apache servers in the 3 different AZs

Apache Servers

ELB & Lambda Configured

Lambda Triggered

ELB logs in CloudWatch

Create a TargetGroup with the three Apache servers

NETWORK & SECURITY

Security Groups

Elastic IPs

Placement Groups

Key Pairs

Network Interfaces

LOAD BALANCING

Load Balancers

Target Groups

AUTO SCALING

Launch Configurations

Auto Scaling Groups

COMMANDS

Command History

Documents

Managed Instances

Activations

Create target group

Actions

Filter: Search

< 1 to 1 of 1 >

Name	Port	Protocol	VPC ID	Monitoring
LoadBalancerTargetGroup	80	HTTP	vpc-b04225d4	<input checked="" type="checkbox"/>

Edit

Registered instances

Instance ID	Name	Port	Availability Zone	Status
i-00255fd07ca832bd2	apache-2b	80	us-west-2b	healthy
i-0fa0adccbf720442	apache-2c	80	us-west-2c	healthy
i-0250179af2a546134	apache-2a	80	us-west-2a	healthy

Availability Zones

Availability Zone	Instance Count	Health?
-------------------	----------------	---------

We then add those 3 apache servers to our ELB

Apache Servers

ELB & Lambda Configured

Lambda Triggered

ELB logs in CloudWatch

Enable ELB log delivery to S3

EC2 Dashboard

Events

Tags

Reports

Limits

INSTANCES

Instances

Spot Requests

Reserved Instances

Scheduled Instances

Dedicated Hosts

AMIS

AMIs

Bundle Tasks

Elastic Block Store

Volumes

Snapshots

NETWORK & SECURITY

Security Groups

Elastic IPs

Placement Groups

Key Pairs

Network Interfaces

LOAD BALANCING

Load Balancers

Secret Resync

Create Load Balancer

Actions

Filter: apache

< 1 to 1 of 1 >

Name	DNS name	State	VPC ID	Availability Zones	Type	Created At
ApacheLoadBalancer	ApacheLoadBalancer-12727...	active	vpc-b04225d4	us-west-2b, us-west-2a...	application	October 11, 2018 at 12:

Edit load balancer attributes

Access Logs delivers detailed logs of all requests made to Elastic Load Balancing. The logs are stored in Amazon S3. See the [documentation](#) for more information.

Enable deletion protection

Idle timeout

Enable access logs

S3 location

s3:// apache-elb-logs

Example: S3bucket/prefix

Create this location for me

This location must exist in the same region as the load balancer.

Cancel Save

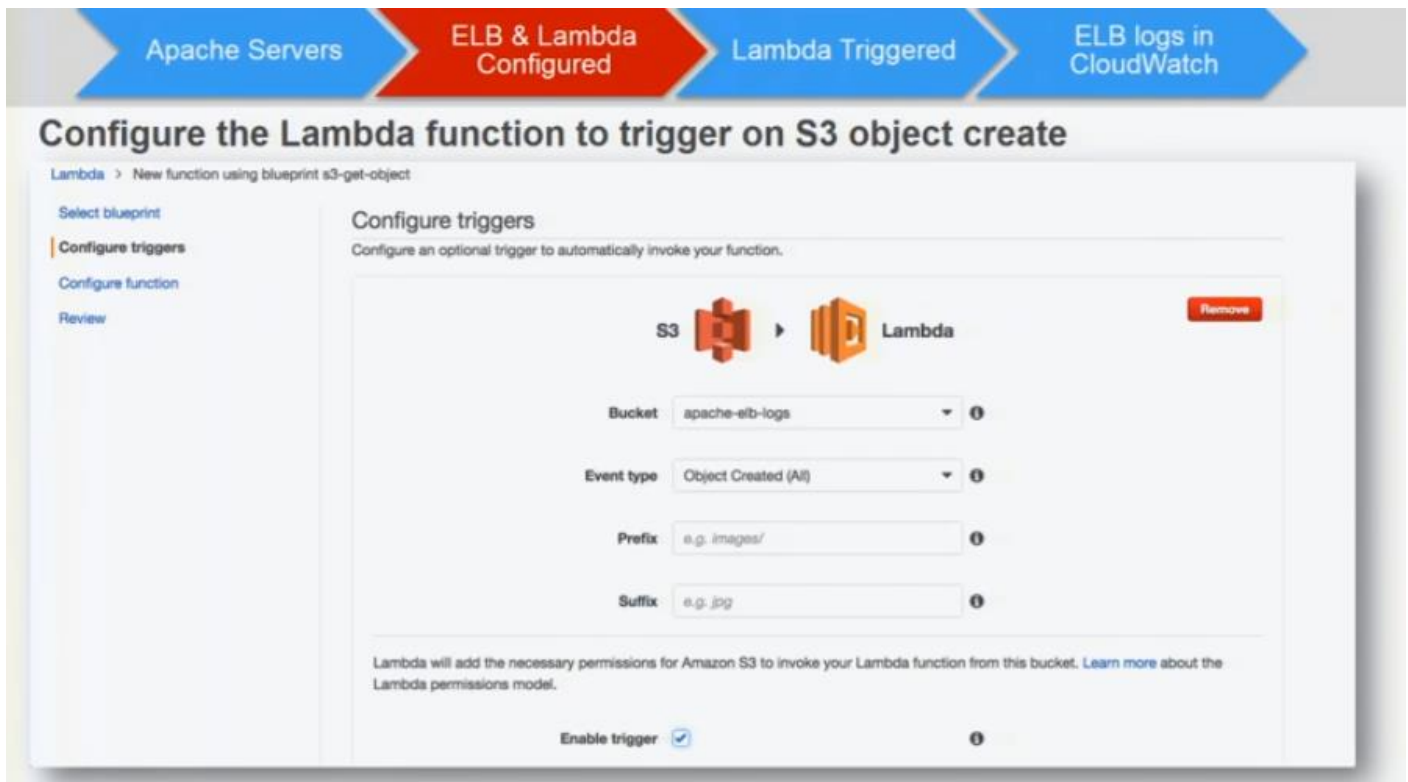
Deletion protection: Disabled

Idle timeout: 60 seconds

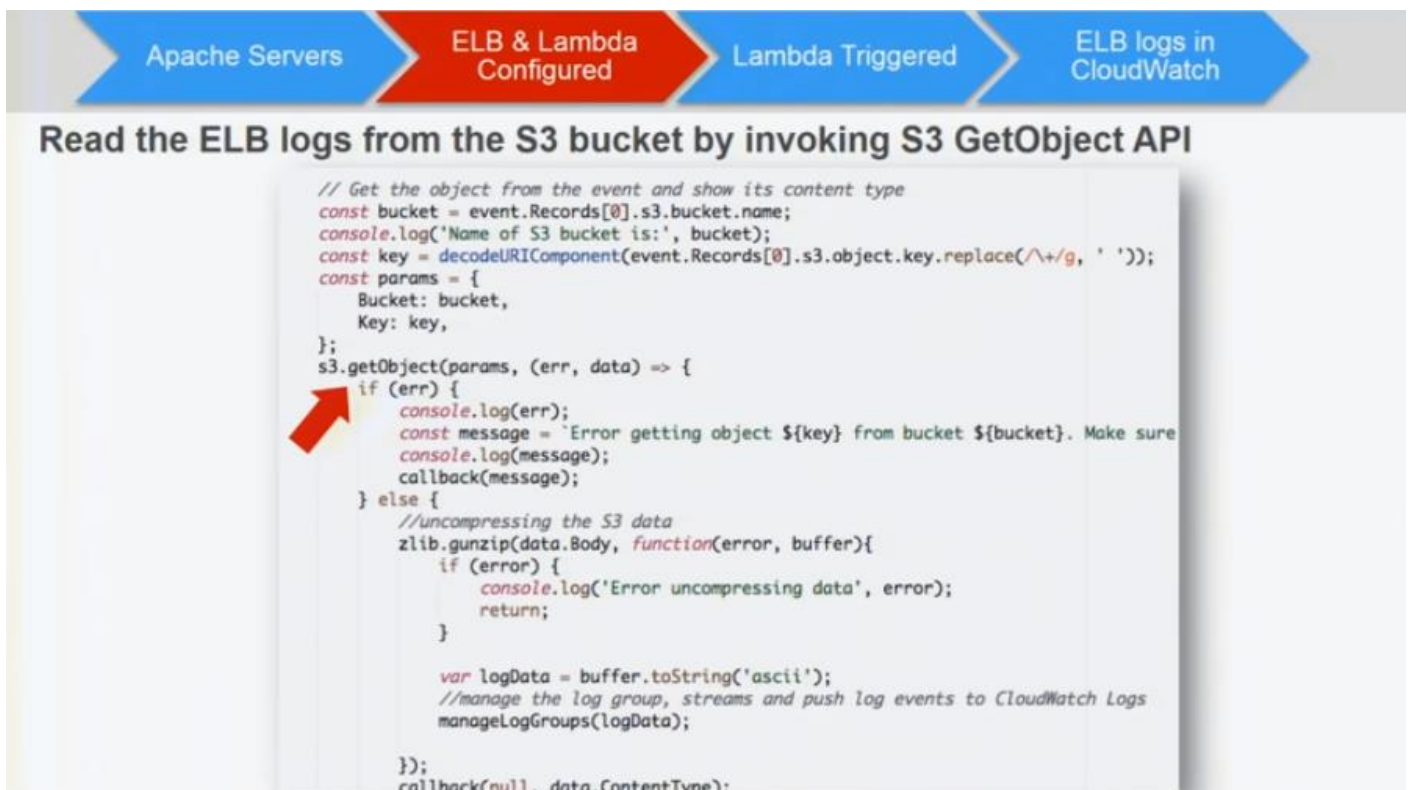
Access logs: S3 location apache-elb-logs /

Edit attributes

We then configure S3 delivery of the access logs to our S3 bucket



As those objects are being created, we have the option to set up a lambda function to trigger on the S3 object creation event



Our lambda function is going to get invoked when an object gets created in the S3 bucket, it then gets the object, decompresses it, break it up into pieces, then calls the AWS SDK's putLogEvents API for CloudWatch **CW** logs as below

Post the logs into CloudWatch Logs by invoking putLogEvents SDK API

```
//Put log events in CloudWatch Logs
function putLogEvents (sequenceToken, logData) {
  var putLogEventParams = {
    logEvents: [ {
      message: logData,
      timestamp: Date.now()
    } ],
    logGroupName: logGroupName,
    logStreamName: logStreamName
  }
  if (sequenceToken) {
    putLogEventParams['sequenceToken'] = sequenceToken;
  }

  cloudWatchLogs.putLogEvents (putLogEventParams, function (err, data) {
    if (err) {
      console.log('Error during put log events: ', err, err.stack);
      return;
    } else {
      console.log('Success in putting log events: ', data);
    }
  });
}
```

This **putLogEvents** API allows you to log from your apps directly into CW without first logging to a file.

ELB serving traffic



Apache2 Ubuntu Default Page

ubuntu

It works!

This is the default welcome page used to test the correct operation of the Apache2 server after installation on Ubuntu systems. It is based on the equivalent page on Debian, from which the Ubuntu Apache packaging is derived. If you can read this page, it means that the Apache HTTP server installed at this site is working properly. You should **replace this file** (located at `/var/www/html/index.html`) before continuing to operate your HTTP server.

If you are a normal user of this web site and don't know what this page is about, this probably means that the site is currently unavailable due to maintenance. If the problem persists, please contact the site's administrator.

Configuration Overview

Ubuntu's Apache2 default configuration is different from the upstream default configuration, and split into several files optimized for interaction with Ubuntu tools. The configuration system is **fully documented in [/usr/share/doc/apache2/README.Debian.gz](#)**. Refer to this for the full documentation. Documentation for the web server itself can be found by accessing the **manual** if the `apache2-doc` package was installed on this server.

The configuration layout for an Apache2 web server installation on Ubuntu systems is as follows:

```
/etc/apache2/
|-- apache2.conf
|   |-- ports.conf
|-- mods-enabled
|   |-- *.load
|   |-- *.conf
|-- conf-enabled
|   |-- *.conf
|-- sites-enabled
|   |-- *.conf
```

- `apache2.conf` is the main configuration file. It puts the pieces together by including all remaining configuration files when starting up the web server.
- `ports.conf` is always included from the main configuration file. It is used to determine the listening ports for incoming connections, and this file can be customized anytime.
- Configuration files in the `mods-enabled/`, `conf-enabled/` and `sites-enabled/` directories contain particular configuration snippets which manage modules, global configuration fragments, or virtual host configurations, respectively.
- They are activated by symlinking available configuration files from their respective `*-available/` counterparts. These should be managed by using our helpers `a2enmod`, `a2enconf`, `a2ensite`, `a2dissite`, and `a2dismod`. See their respective man pages for detailed instructions.

We can now test it out by visiting the ELK endpoint


```

2016-11-06
* 10:50:06 http 2016-11-06T10:48:47.113144Z app/ApacheLoadBalancer/cd3cfad23b36ea4 23.115.195.81:34810 172.31.40.221:80 0.001 0.001 0.000 200 200 123 11783 "GET http
http 2016-11-06T10:48:47.113144Z app/ApacheLoadBalancer/cd3cfad23b36ea4 23.115.195.81:34810 172.31.40.221:80 0.001 0.001 0.000 200 200 123 11783 "GET
http://apache.loadbalancer-1272730188.us-west-2.elb.amazonaws.com:80/ HTTP/1.0" Mozilla/5.0 (compatible; Googlebot/2.1; http://www.google.com/bot.html)" - -
arn:aws:elasticloadbalancing:us-west-2:643289538482:targetgroup/LoadBalancerTargetGroup/3e58fd867b34ac

* 11:00:06 http 2016-11-06T10:56:52.322702Z app/ApacheLoadBalancer/cd3cfad23b36ea4 200.215.162.246:50805 172.31.40.221:80 0.001 0.001 0.000 200 200 145 11788 "GET http
http 2016-11-06T10:56:52.322702Z app/ApacheLoadBalancer/cd3cfad23b36ea4 200.215.162.246:50805 172.31.40.221:80 0.001 0.001 0.000 200 200 145 11788 "GET
http://854.201.067.026:80/ HTTP/1.1" curl/7.17.1 (msis-unknown-linux-gnu) libcurl/7.17.1 OpenSSL/0.9.8i zlib/1.2.3" - - arn:aws:elasticloadbalancing:us-west-2:643289538482:targetgroup/LoadBalancerTargetGroup/3e58fd867b34ac

* 11:20:01 http 2016-11-06T11:17:07.183530Z app/ApacheLoadBalancer/cd3cfad23b36ea4 123.22.16.219:57198 172.31.40.220:80 0.001 0.001 0.000 200 200 123 11783 "GET http
http 2016-11-06T11:17:07.183530Z app/ApacheLoadBalancer/cd3cfad23b36ea4 123.22.16.219:57198 172.31.40.220:80 0.001 0.001 0.000 200 200 123 11783 "GET
http://apache.loadbalancer-1272730188.us-west-2.elb.amazonaws.com:80/ HTTP/1.0" Mozilla/5.0 (compatible; Googlebot/2.1; http://www.google.com/bot.html)" - -
arn:aws:elasticloadbalancing:us-west-2:643289538482:targetgroup/LoadBalancerTargetGroup/3e58fd867b34ac

* 11:45:08 http 2016-11-06T11:44:44.753834Z app/ApacheLoadBalancer/cd3cfad23b36ea4 70.53.229.49:60243 172.31.46.241:80 0.001 0.001 0.000 200 200 123 11783 "GET http://
http 2016-11-06T11:44:44.753834Z app/ApacheLoadBalancer/cd3cfad23b36ea4 70.53.229.49:60243 172.31.46.241:80 0.001 0.001 0.000 200 200 123 11783 "GET http://apache.loadbalancer-
1272730188.us-west-2.elb.amazonaws.com:80/ HTTP/1.0" Mozilla/5.0 (compatible; Googlebot/2.1; http://www.google.com/bot.html)" - - arn:aws:elasticloadbalancing:us-west-2:643289538482:targetgroup/LoadBalancerTargetGroup/3e58fd867b34ac

* 13:00:02 http 2016-11-06T12:58:38.333603Z app/ApacheLoadBalancer/cd3cfad23b36ea4 185.69.6.98:59838 172.31.40.221:80 0.001 0.001 0.000 200 200 127 11788 "GET http://
http 2016-11-06T12:58:38.333603Z app/ApacheLoadBalancer/cd3cfad23b36ea4 185.69.6.98:59838 172.31.40.221:80 0.001 0.001 0.000 200 200 127 11788 "GET http://852.024.283.111:80/
HTTP/1.1" curl/7.17.1 (msis-unknown-linux-gnu) libcurl/7.17.1 OpenSSL/0.9.8i zlib/1.2.3" - - arn:aws:elasticloadbalancing:us-west-2:643289538482:targetgroup/LoadBalancerTargetGroup/3e58fd867b34ac

```

Line Number	Protocol	Timestamp	Self	Client
1	http	2016-10-14T08:51:45.788677Z	app/apacheLoadBalancer/cd3c4fad23c36ae4	205.251
2	http	2016-10-14T08:51:47.804226Z	app/apacheLoadBalancer/cd3c4fad23c36ae4	205.251
3	http	2016-10-14T09:06:53.734455Z	app/apacheLoadBalancer/cd3c4fad23c36ae4	205.251

Apache Servers

ELB& Lambda
Configured

Lambda Triggered

ELB logs in
CloudWatch

Define a metric filter on the log group

Create Metric Filter and Assign a Metric

Filter for Log Group: apache-elb-logs

Log events that match the pattern you define are recorded to the metric that you specify. You can graph the metric and set alarms to notify you.

Filter Name: Filter Pattern:

Metric Details

Metric Namespace: [Create new namespace](#)Metric Name: Metric Value:

\$protocol \$timestamp \$elb \$client_port \$backend_port \$request_processing_time
\$backend_processing_time \$response_processing_time \$elb_status_code
\$backend_status_code \$received_bytes \$sent_bytes \$request \$user_agent \$ssl_cipher
\$ssl_protocol

Cancel

Previous

Create Filter

Apache Servers

ELB& Lambda
Configured

Lambda Triggered

ELB logs in
CloudWatch

View the metric graph for ELB latency

HighLatencyCount

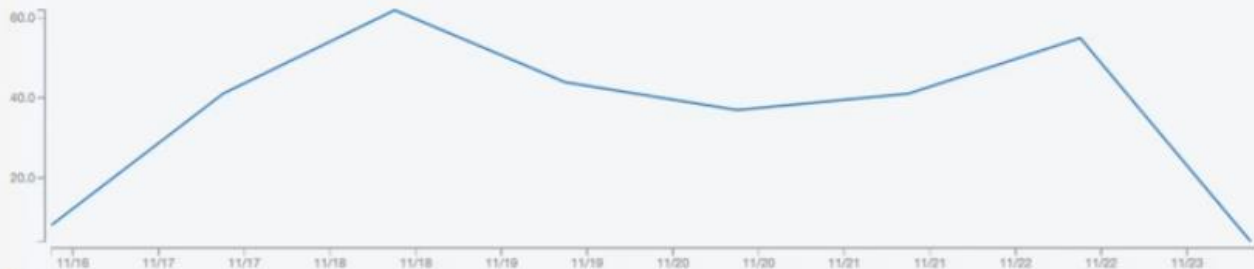
1h 3h 12h 1d 3d 1w custom ▾

Line ▾

Actions ▾

↺ ↻

?



■ ELBLatency

All metrics

Graphed metrics (1)

Graph options

Label	Namespace	Dimensions	Metric Name	Statistic	Period	Y Axis	Actions
ELBLatency	LogMetrics		HighLatencyCount	Sum	1 Day	< >	🔔 📄 ⌕

Key takeaways

1. S3 delivered log data from any source can be centralized into CloudWatch Logs using Lambda
2. You can search and extract metrics from those logs in near real time



Problem Statements

- When you get an alarm you want enough information to decide whether it needs immediate attention or not
- You want to customize the alarm text and format to your operational needs

Flow of Events



We can send generated CW error logs to trigger an alarm and published to an SNS topic, the SNS topic will then trigger an event notification to a lambda, the lambda function then pulls the information out of the alarm, use the information to search the log service for the SLA breach emitted from the log service and invoke SES with the logs to forward to the operator email in rich text format what to do

Metric Filter Defined

Alarm & Lambda Configured

Lambda Triggered

SES sends Email

Provision an EC2 instance with Tomcat running on it

Launch Instance

Connect

Actions

search : tomcat

Add filter

Name	elb-logs	Instance ID	Instance Type	Availability Zone	Instance State	Status
demo-tomcat-application		i-04df2c40bf1e2629f	t1.micro	us-west-2c	running	2/

Instance: i-04df2c40bf1e2629f (demo-tomcat-application)

Elastic IP: 52.37.119.7

Description

Status Checks

Monitoring

Tags

Instance ID

i-04df2c40bf1e2629f

Public DNS

Instance state

running

Public IP

Instance type

t1.micro

Elastic IPs

52.37.119.7

Private DNS

ip-172-31-10-66.us-west-2.compute.internal

Availability zone

us-west-2c

Private IPs

172.31.10.66

Security groups

awseb-e-awubna4swx-stack-AWSEBSecurityGroup

Metric Filter
DefinedAlarm & Lambda
Configured

Lambda Triggered

SES sends Email

CloudWatch agent sends EC2 instance logs to CloudWatch Logs

CloudWatch > Log Groups > /var/log/httpd/access_log > i-04df2c40bf1e2629f

Expand all

Row

Text

Filter events

all 30s 5m 1h 6h 1d 1w custom

Time (UTC +00:00)	Message
2016-11-07 22:11:22	113.253.124.253 - - [06/Nov/2016:22:11:21 +0000] "GET / HTTP/1.0" 401 994 "-" Mozilla/5.0 (compatible; Googlebot/2.1; +http://w
22:31:15	219.85.62.148 - - [06/Nov/2016:22:31:14 +0000] "GET / HTTP/1.0" 401 994 "-" Mozilla/5.0 (compatible; Googlebot/2.1; +http://w
2016-11-07 00:00:09	175.143.129.228 - - [07/Nov/2016:00:00:08 +0000] "GET / HTTP/1.0" 401 994 "-" Mozilla/5.0 (compatible; Googlebot/2.1; +http://h
03:58:20	54.159.5.179 - - [07/Nov/2016:03:58:20 +0000] "GET /api/json HTTP/1.1" 401 994 "-" python-requests/2.7.0 CPython/2.7.12 Linu
03:58:21	54.159.5.179 - - [07/Nov/2016:03:58:20 +0000] "x16/x03/x01" 405 1086 "-" "
03:58:24	54.85.103.120 - - [07/Nov/2016:03:58:23 +0000] "GET http://example.com/ HTTP/1.1" 401 994 "-" python-requests/2.9.1"
04:52:01	115.134.113.80 - - [07/Nov/2016:04:52:01 +0000] "GET / HTTP/1.0" 401 994 "-" Mozilla/5.0 (compatible; Googlebot/2.1; +http://w
05:41:26	91.197.234.20 - - [07/Nov/2016:05:41:26 +0000] "x03" 405 1086 "-" "
07:07:13	69.30.193.254 - - [07/Nov/2016:07:07:12 +0000] "GET / HTTP/1.1" 401 994 "-" Mozilla/4.0 (compatible; MSIE 8.0; Windows NT 6.
08:10:44	139.162.13.205 - - [07/Nov/2016:08:10:43 +0000] "GET / HTTP/1.1" 401 994 "-" Mozilla/5.0 (Windows NT 10.0; WOW64) AppleWe
09:03:49	220.202.123.178 - - [07/Nov/2016:09:03:48 +0000] "GET HTTP/1.1 HTTP/1.1" 400 303 "-" Mozilla/4.0 (compatible; MSIE 6.0; Winc
09:03:49	220.202.123.178 - - [07/Nov/2016:09:03:49 +0000] "GET /cgi-bin/test-cgi HTTP/1.1" 401 994 "-" Mozilla/4.0 (compatible; MSIE 6.0
09:21:23	63.141.250.155 - - [07/Nov/2016:09:21:23 +0000] "GET / HTTP/1.1" 401 994 "-" Mozilla/4.0 (compatible; MSIE 8.0; Windows NT 6
10:13:01	104.193.252.232 - - [07/Nov/2016:10:13:00 +0000] "GET / HTTP/1.1" 401 994 "-" Mozilla/5.0"
10:42:28	23.252.54.224 - - [07/Nov/2016:10:42:27 +0000] "GET / HTTP/1.1" 401 994 "-" Mozilla/5.0 (Macintosh; Intel Mac OS X 10_12_1) A
10:42:31	23.252.54.224 - - [07/Nov/2016:10:42:31 +0000] "GET /favicon.ico HTTP/1.1" 401 994 "http://ec2-52-37-119-7.us-west-2.compute

Metric Filter
DefinedAlarm & Lambda
Configured

Lambda Triggered

SES sends Email

Define a filter pattern to extract Unauthorized access attempts

Define Logs Metric Filter

Filter for Log Group: /var/log/httpd/access_log

You can use metric filters to monitor events in a log group as they are sent to CloudWatch Logs. You can monitor and count specific terms or extract values from log events and associate the results with a metric. [Learn more about pattern syntax.](#)

Filter Pattern

[Show examples](#)

Select Log Data to Test

0

Test Pattern

Clear

205.251.233.183 - - [11/Jul/2016:08:51:41 +0000] "GET / HTTP/1.1" 401 994 "-" Mozilla/5.0 (Macintosh; in
205.251.233.183 - - [11/Jul/2016:08:52:16 +0000] "GET / HTTP/1.1" 401 994 "-" Mozilla/5.0 (Macintosh; in
205.251.233.183 - - [11/Jul/2016:08:52:22 +0000] "GET / HTTP/1.1" 401 994 "-" Mozilla/5.0 (Macintosh; in
205.251.233.183 - - [11/Jul/2016:09:06:57 +0000] "GET / HTTP/1.1" 401 994 "-" Mozilla/5.0 (Macintosh; in
205.251.233.183 - - [11/Jul/2016:09:06:03 +0000] "GET /failattempt HTTP/1.1" 401 994 "-" Mozilla/5.0 (M
205.251.233.183 - - [11/Jul/2016:09:23:52 +0000] "GET / HTTP/1.1" 401 994 "-" Mozilla/5.0 (Macintosh; in
205.251.233.183 - - [11/Jul/2016:09:28:28 +0000] "GET / HTTP/1.1" 404 994 "-" Mozilla/5.0 (Macintosh; in

Results

Found 6 matches out of 7 event(s) in the sample log.

Line Number	Host	logName	User	Timestamp	Request
1	205.251.233.183	-	-	11/Jul/2016:08:51:41 +0000	GET / HTTP/1.1
2	205.251.233.183	-	-	11/Jul/2016:08:52:16 +0000	GET / HTTP/1.1
3	205.251.233.183	-	-	11/Jul/2016:08:52:22 +0000	GET / HTTP/1.1
4	205.251.233.183	-	-	11/Jul/2016:09:06:57 +0000	GET / HTTP/1.1
5	205.251.233.183	-	-	11/Jul/2016:09:06:03 +0000	GET /failattempt HTTP/1.1

Metric Filter
Defined

Alarm & Lambda
Configured

Lambda Triggered

SES sends Email

Define a metric filter on the log group

Create Metric Filter and Assign a Metric

Filter for Log Group: /var/log/httpd/access_log

Log events that match the pattern you define are recorded to the metric that you specify. You can graph the metric and set alarms to notify you.

Filter Name: UnauthorizedAccessFilter

Filter Pattern: [host, logName, user, timestamp, request, statusCode=401, size]

Metric Details

Metric Namespace: LogMetrics [Create new namespace](#)

Metric Name: UnauthorizedAccess

Metric Value: 1

\$host \$logName \$user \$timestamp \$request \$statusCode \$size

Cancel

Previous

Create Filter

Metric Filter
Defined

Alarm & Lambda
Configured

Lambda Triggered

SES sends Email

Define an alarm with a specific threshold for that metric

Modify Alarm

1. Select Metric 2. Define Alarm

Alarm Threshold

Provide the details and threshold for your alarm. Use the graph on the right to help set the appropriate threshold.

Name: UnauthorizedAccess

Description:

Whenever: UnauthorizedAccess

Is: \geq 5

for: 1 consecutive period(s)

Actions

Define what actions are taken when your alarm changes state.

Notification

Delete

Whenever this alarm: State is ALARM

Send notification to:

[New list](#) [Enter list](#)

Alarm Preview

This alarm will trigger when the blue line goes up to or above the red line for a duration of 5 minutes



Namespace: LogMetrics

Metric Name: UnauthorizedAccess

Period: 5 Minutes

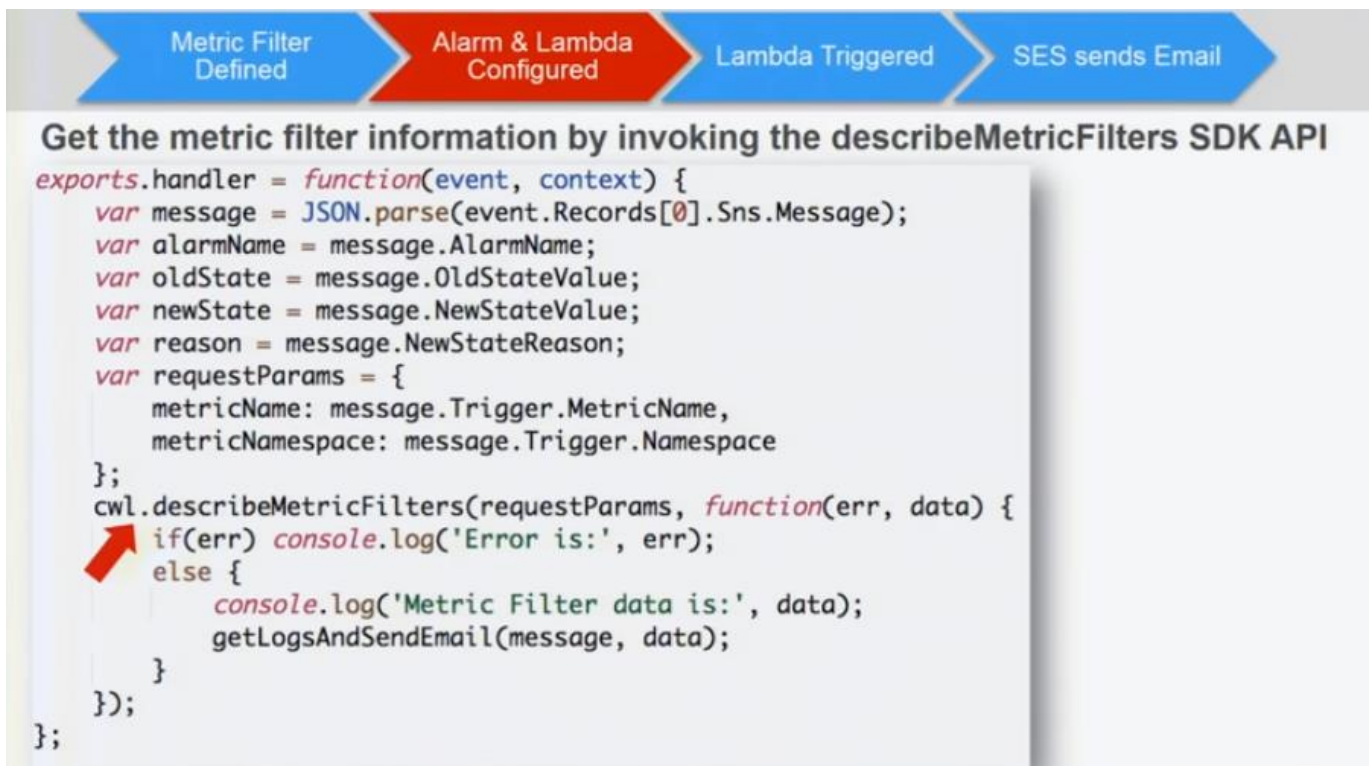
Statistic: Sum

Cancel

Previous

Next

Save Changes



Configure the lambda as above that takes the alarm notification from the SNS message, pulls out some data like the MetricName and Namespace, then called the **describeMetricFilters()** public API using the AWS SDK

Metric Filter
Defined

Alarm & Lambda
Configured

Lambda Triggered

SES sends Email

Get the relevant log data by invoking the filterLogEvents SDK API

```
function getLogsAndSendEmail(message, metricFilterData) {  
    var timestamp = Date.parse(message.StateChangeTime);  
    var offset = message.Trigger.Period * message.Trigger.EvaluationPeriods * 1000;  
    var metricFilter = metricFilterData.metricFilters[0];  
    var parameters = {  
        'logGroupName': metricFilter.logGroupName,  
        'filterPattern': metricFilter.filterPattern ? metricFilter.filterPattern : "",  
        'startTime': timestamp - offset,  
        'endTime': timestamp  
    };  
    cw.filterLogEvents(parameters, function (err, data) {  
        if (err) {  
            console.log('Filtering failure:', err);  
        } else {  
            console.log("==SENDING EMAIL==");  
            var email = ses.sendEmail(generateEmailContent(data, message), function (err, data) {  
                if (err) console.log(err);  
                else {  
                    console.log("==EMAIL SENT==");  
                    console.log(data);  
                }  
            });  
        }  
    });  
}
```

Then call the filterLogEvents() API with the parameters and data,

Metric Filter
Defined

Alarm & Lambda
Configured

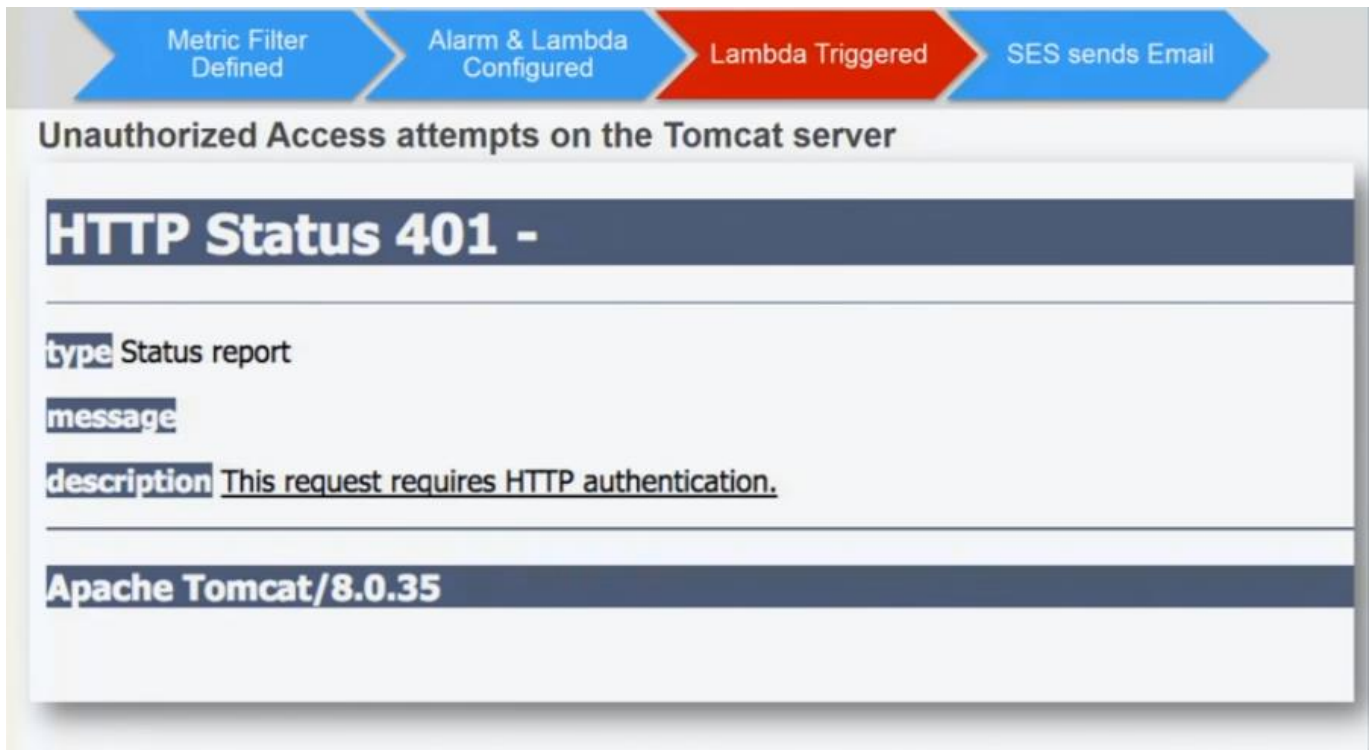
Lambda Triggered

SES sends Email

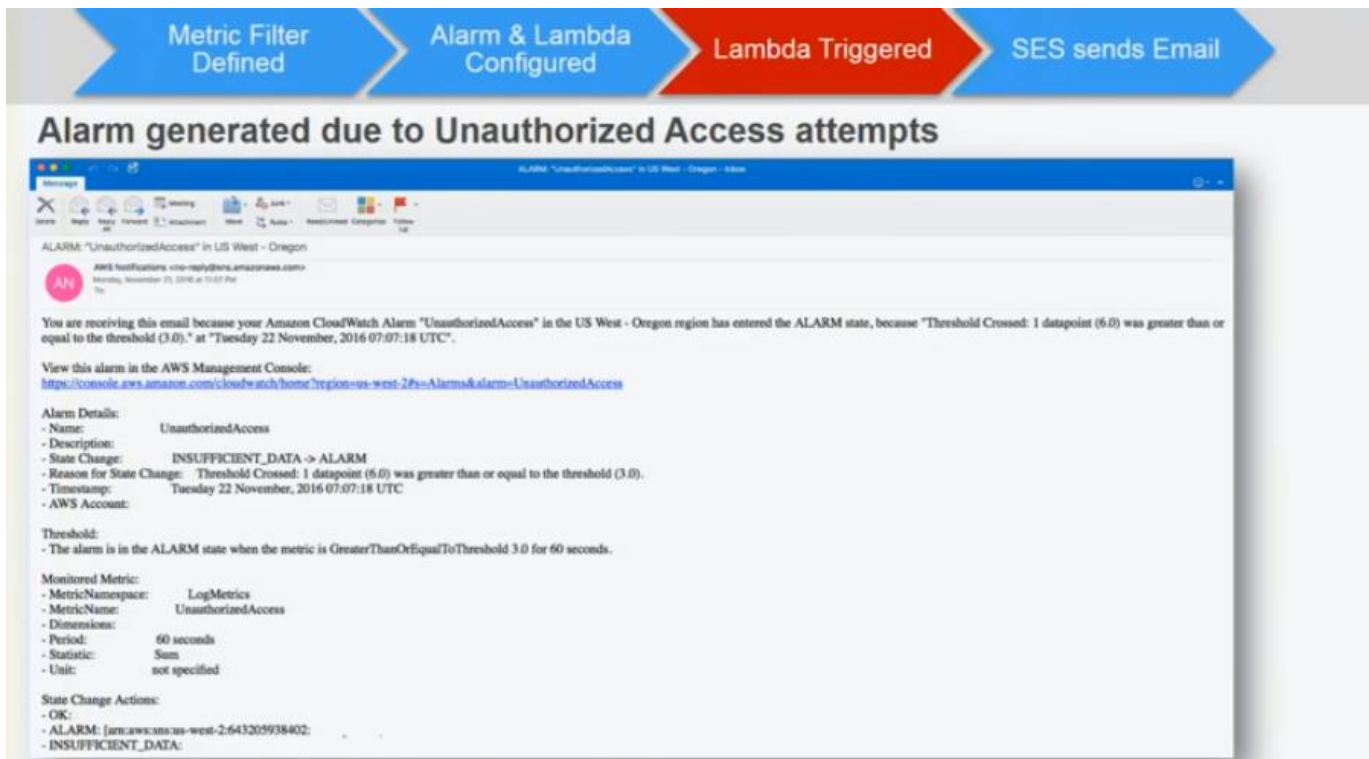
SES sends an email with the relevant log data

```
function generateEmailContent(data, message) {  
    var events = data.events;  
    console.log('Events are:', events);  
    var style = '<style> pre {color: red;} </style>';  
    var logData = '<br/>Logs:<br/>' + style;  
    for (var i in events) {  
        logData += '<pre>Instance: ' + JSON.stringify(events[i]['logStreamName']) + '</pre>';  
        logData += '<pre>Message: ' + JSON.stringify(events[i]['message']) + '</pre><br/>';  
    }  
  
    var date = new Date(message.StateChangeTime);  
    var text = 'Alarm Name: ' + '<b>' + message.AlarmName + '</b><br/>' +  
        'Runbook Details: <a href="http://wiki.mycompany.com/prodrunbook">Production Runbook</a><br/>' +  
        'Account ID: ' + message.AWSAccountId + '<br/>' +  
        'Region: ' + message.Region + '<br/>' +  
        'Alarm Time: ' + date.toString() + '<br/>' +  
        logData;  
    var subject = 'Details for Alarm - ' + message.AlarmName;  
    var emailContent = {  
        Destination: {  
            ToAddresses: ["Add destination email here"]  
        },  
        Message: {  
            Body: {  
                Html: {  
                    Data: text  
                }  
            },  
            Subject: {  
                Data: subject  
            }  
        },  
        Source: 'Add source email here'  
    };  
}
```

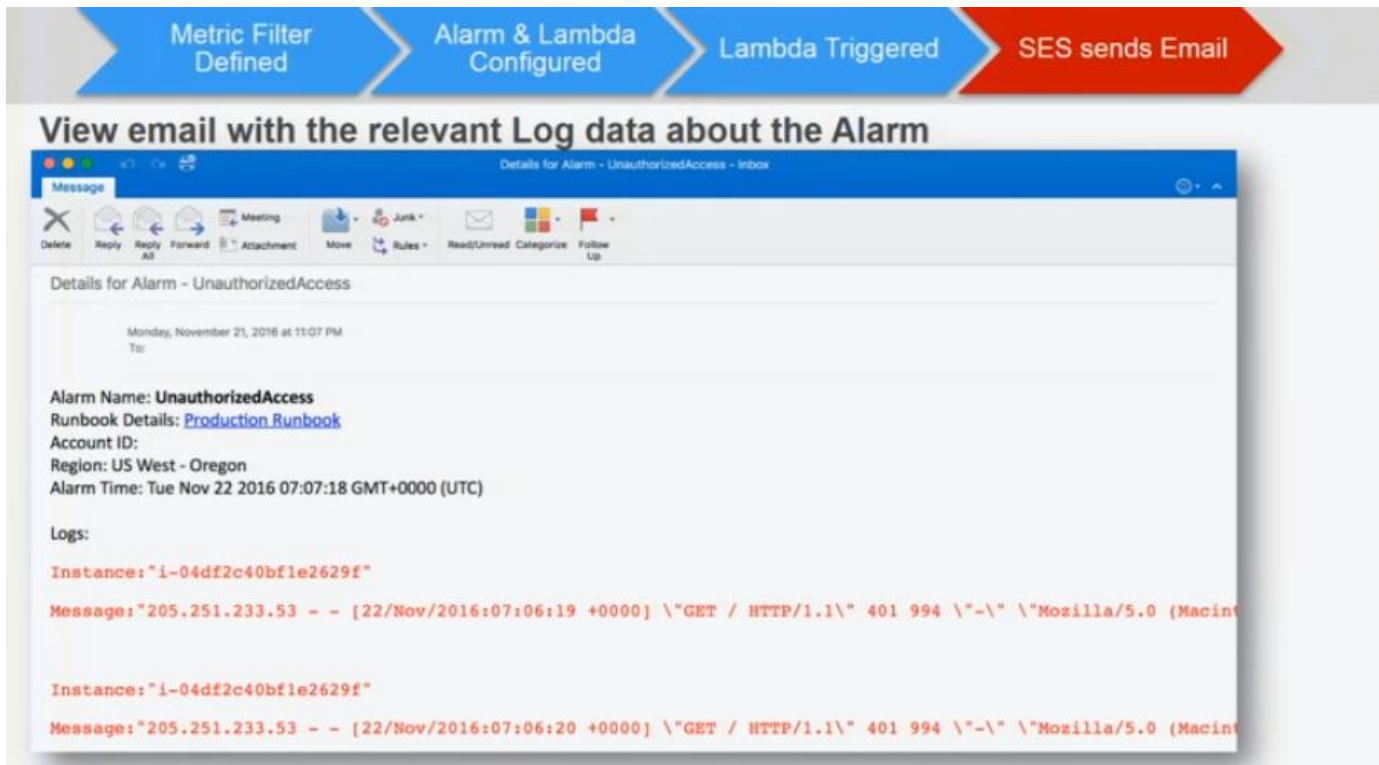
Send out the email with other details to the operator



Test it out



This is the alarm we will get



The email is much better with the details in it

Key takeaways

- Alarms can be customized to add specific details about the issue
- When you see a spike on a metric, you can also get the logs describing the issue triggering the alarm
- The Lambda function can be extended to add your specific information to the alarm

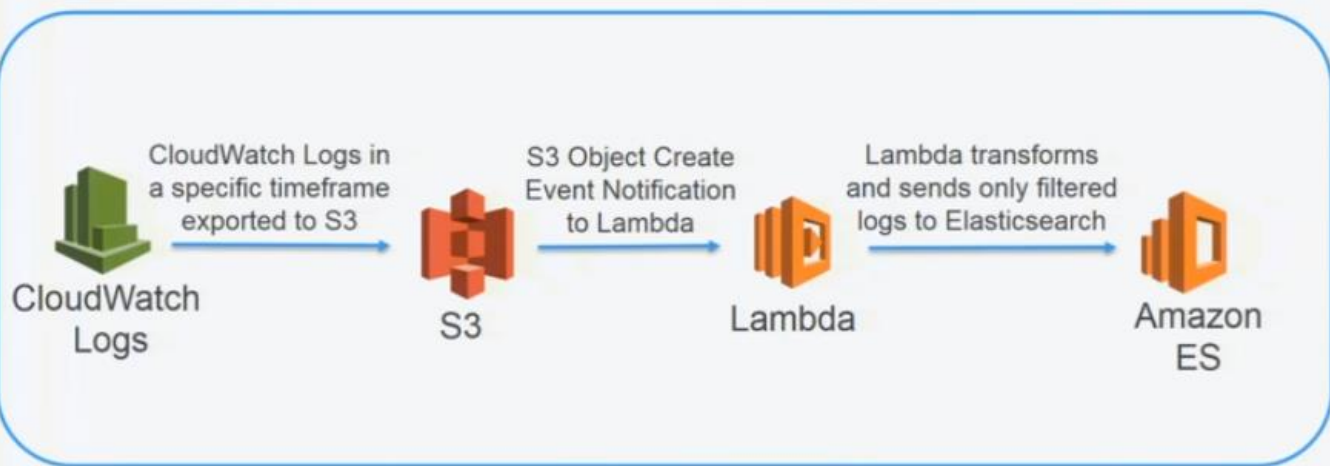


Problem Statements

- You want to do log analysis using Elasticsearch but don't want to leave the cluster running all the time
- You want to send data to Elasticsearch, but don't want to manage ongoing operations
- Build an on-demand Elasticsearch cluster from historical data

You can integrate with the Amazon Elasticsearch service, data can flow into it in real time and then do real time analysis on it without leaving the lambda running everytime. We might want to accumulate the historical CW logs data and then do searching on it when needed by standing up an Elasticsearch cluster

Flow of Events



You can also transform the logs by breaking them into smaller more indexable pieces before sending them into the ES endpoint.

The screenshot shows the AWS Management Console interface for configuring flow logging for a VPC. A navigation bar at the top indicates the process: "Logs exported to S3" (red), "Lambda Configured" (blue), "Logs to Elasticsearch" (blue), and "Visuals in Kibana" (blue). The main heading is "Configure flow logging for your VPC". Below this, a "Create Flow Log" dialog box is open, showing the following configuration:

- Resources:** eni-ba14b0c7
- Filter:** All
- Role:** flowlogsRole
- ARN:** arn:aws:iam::643205838402:role/flowlogsRole
- Destination Log Group:** vpc-flow-logs

The dialog box also includes a "Cancel" button and a "Create Flow Log" button. The background shows the "Create Network Interface" page with a table of network interfaces.

Name	Network Interface	Subnet ID	VPC ID	Zone	Security groups	Description	Instance ID	Status	Public IP	Primary private	Secondary private
	eni-ba14b0c7								52.40.89.187	172.31.25.200	

Logs exported to S3

Lambda Configured

Logs to Elasticsearch

Visuals in Kibana

View the VPC Flow logs in CloudWatch Logs

CloudWatch > Log Groups > vpc-flow-logs > eni-c38f7e82-all

Expand all ☐ Row ☒ Text ↺ ⚙ ⓘ

Filter events all 30s 5m 1h 6h 1d 1w custom ▾

Message

2016-11-04 08:39:26

2	643205938402	eni-c38f7e82	172.31.46.241	172.31.27.63	80	48917	6	6	12103	1478248766	1478248816	ACCEPT	OK
2	643205938402	eni-c38f7e82	172.31.27.63	172.31.46.241	48953	80	6	5	361	1478248816	1478248876	ACCEPT	OK
2	643205938402	eni-c38f7e82	79.46.115.16	172.31.46.241	24489	39987	17	1	139	1478248816	1478248876	REJECT	OK
2	643205938402	eni-c38f7e82	172.31.46.241	172.31.27.63	80	48939	6	6	12103	1478248816	1478248876	ACCEPT	OK
2	643205938402	eni-c38f7e82	172.31.40.238	172.31.46.241	41713	80	6	5	361	1478248816	1478248876	ACCEPT	OK
2	643205938402	eni-c38f7e82	172.31.46.241	172.31.40.238	80	41713	6	6	12103	1478248816	1478248876	ACCEPT	OK
2	643205938402	eni-c38f7e82	172.31.40.238	172.31.46.241	41689	80	6	5	361	1478248816	1478248876	ACCEPT	OK

Logs exported to S3

Lambda Configured

Logs to Elasticsearch

Visuals in Kibana

[Step 1: Define domain](#)
[Step 2: Configure cluster](#)
[Step 3: Set up access policy](#)
Step 4: Review

Review

Review the information below, and then choose **Confirm and create**.

Define domain

Elasticsearch domain name

demo-s3-es-vpc-flow-log

Edit

Elasticsearch version

2.3

Configure cluster

Instance type

m3.medium.elasticsearch (default)

Instance count

1

Dedicated master

Disabled

Zone awareness

Disabled

Storage type

Instance (default)

Start hour for the daily automated snapshot

0:00 Hour (UTC)

Advanced options

Elasticsearch parameters:
rest.action.multi.allow_explicit_index: true
indices.fielddata.cache.size: unbounded (default)

We create an ES domain, set it up with the number of nodes that we need, the types of data, etc

Logs exported to
S3

Lambda
Configured

Logs to
Elasticsearch

Visuals in Kibana

Read the VPC Flow Logs from the S3 bucket by invoking GetObject API

```
exports.handler = (event, context, callback) => {
  //console.log('Received event:', JSON.stringify(event, null, 2));

  // Get the object from the event and show its content type
  const bucket = event.Records[0].s3.bucket.name;
  console.log('The name of bucket is:', bucket);
  const key = decodeURIComponent(event.Records[0].s3.object.key.replace(/\+/g, ' '));
  console.log('The name of key is:', key);

  const params = {
    Bucket: bucket,
    Key: key,
  };
  s3.getObject(params, (err, data) => {
    if (err) {
      console.log(err);
      const message = `Error getting object ${key} from bucket ${bucket}. Make sure`;
      console.log(message);
      callback(message);
    } else {
      console.log('CONTENT TYPE:', data.ContentType);
      console.log('Reading the S3 data:');
      zlib.gunzip(data.Body, function (error, buffer) {
        if (error) {

```

We also set up our lambda function to trigger on the **s3ObjectCreate()** event to send the data to ES

Logs exported to
S3

Lambda
Configured

Logs to
Elasticsearch

Visuals in Kibana

Transform the VPC Flow logs into a JSON document for Elasticsearch

```
// index name format: cwl-YYYY.MM.DD
var indexName = [
  'cwl-' + timestamp.getUTCFullYear(),           // year
  ('0' + (timestamp.getUTCMonth() + 1)).slice(-2), // month
  ('0' + timestamp.getUTCDate()).slice(-2)       // day
].join('.');
var message = parts[2];

var source = buildSource(message, {});
source['@id'] = id;
source['@timestamp'] = new Date(1 * timestamp).toISOString();
source['@message'] = message;
source['@owner'] = payload.owner;
source['@log_group'] = bucket;
source['@log_stream'] = logStream;

var action = { "index": {} };
action.index._index = indexName;
action.index._type = bucket;
action.index._id = id;

bulkRequestBody += [
  JSON.stringify(action),
  JSON.stringify(source),
].join('\n') + '\n';
});
```

Logs exported to
S3

Lambda
Configured

Logs to
Elasticsearch

Visuals in Kibana

Ingest the logs into Elasticsearch by putting to its HTTP endpoint

```
function post(body, callback) {
  var requestParams = buildRequest(endpoint, body);

  var request = https.request(requestParams, function(response) {
    var responseBody = '';
    response.on('data', function(chunk) {
      responseBody += chunk;
    });
    response.on('end', function() {
      var info = JSON.parse(responseBody);
      var failedItems;
      var success;

      if (response.statusCode >= 200 && response.statusCode < 299) {
        failedItems = info.items.filter(function(x) {
          return x.index.status >= 300;
        });

        success = {
          "attemptedItems": info.items.length,
          "successfulItems": info.items.length - failedItems.length,
          "failedItems": failedItems.length
        };
      }
    });
  });
}
```

Logs exported to
S3

Lambda
Configured

Logs to
Elasticsearch

Visuals in Kibana

Export the VPC Flow logs to S3

CloudWatch > Log Groups

Create Metric Filter

Filter: Log Group Name

Log Groups

- ☐ /aws/kinesisfirehose/KinesisFirehose
- ☐ /aws/lambda/Customize
- ☐ /aws/lambda/Demo-S3
- ☐ /aws/lambda/Git-S3-E
- ☐ /aws/lambda/HelloWorld
- ☐ /aws/lambda/LambdaF
- ☐ /aws/lambda/LambdaF
- ☐ /aws/lambda/LogsToE
- ☐ /aws/lambda/S3-GetO
- ☐ /aws/lambda/Unauthor
- ☐ /var/log/dpkg.log
- ☐ /var/log/httpd/access
- ☐ /var/log/messages
- ☐ APIGateway

Export data to Amazon S3

Define data to export

From: 2016/11/03 00:00 UTC (GMT)

To: 2016/11/10 00:00 UTC (GMT)

Advanced +

Choose S3 bucket

Select account ☒ This account ☐ Another account

S3 bucket name: demo-rca-logs

Advanced +

Click **Export data** to export events between [2016/11/03 00:00] and [2016/11/10 00:00] from the vpc-flow-logs log group to the demo-rca-logs S3 bucket.

Cancel Export data

Subscriptions

- None
- None
- None
- None
- None
- None
- None
- None
- None
- None
- None
- None
- Lambda (UnauthorizedAccess)
- None
- None

Logs exported to
S3

Lambda
Configured

Logs to
Elasticsearch

Visuals in Kibana

VPC Flow logs exported to S3

Upload Create Folder Actions

Search by prefix

None Properties Transfers

All Buckets / demo-rca-logs / demo / 95d9a081-2e6a-4e75-90bc-48c2f26fa603

Name	Storage Class	Size	Last Modified
eni-100c5d4f-all	--	--	--
eni-1a1eb145-all	--	--	--
eni-1ba04144-all	--	--	--
eni-24045768-all	--	--	--
eni-43e6833e-all	--	--	--
eni-4a10010b-all	--	--	--
eni-4d0010c-all	--	--	--
eni-51d82910-all	--	--	--
eni-83e1c9fc-all	--	--	--
eni-93619a62-all	--	--	--
eni-b445ae5-all	--	--	--
eni-b645ae7-all	--	--	--
eni-b745ae8-all	--	--	--
eni-b945ae9-all	--	--	--
eni-ba14bcc7-all	--	--	--
eni-c087e62-all	--	--	--
eni-e795b698-all	--	--	--

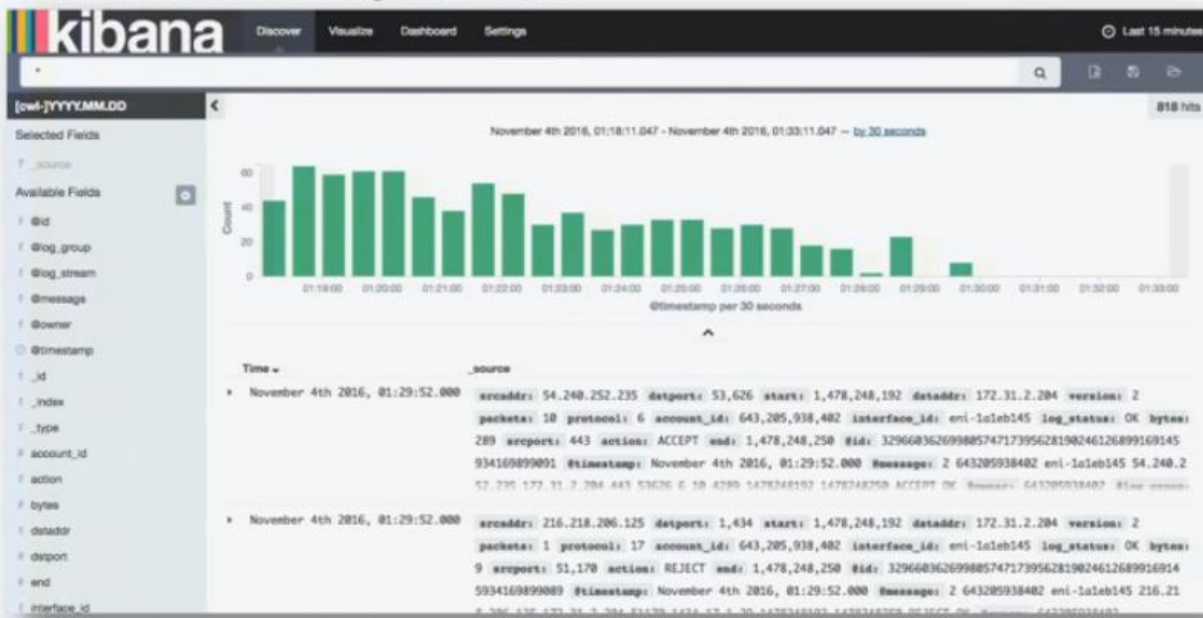
Logs exported to
S3

Lambda
Configured

Logs to
Elasticsearch

Visuals in Kibana

View the VPC Flow logs in Kibana



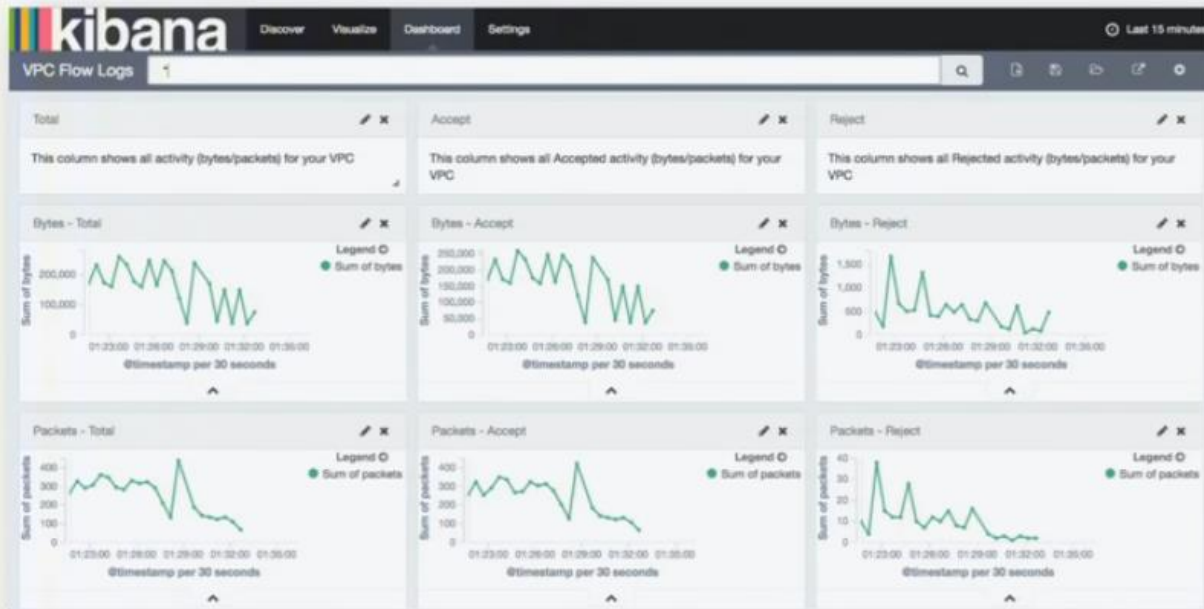
Logs exported to S3

Lambda Configured

Logs to Elasticsearch

Visuals in Kibana

View Different Dashboards in Kibana



Logs exported to S3

Lambda Configured

Logs to Elasticsearch

Visuals in Kibana

View Different Dashboards in Kibana



Key takeaways

- Send historical data within a timeframe to Elasticsearch on demand
- This reduces cost, burden of scalability, and operations time
- Troubleshooting gets easier because you have only limited and relevant data

Recap

- Monitoring is more important than ever, but still too hard
- CloudWatch is working to make monitoring easier
- CloudWatch Logs and Lambda are powerful tools to tailor your monitoring for your business needs

Useful Links

- CloudWatch Overview - <https://aws.amazon.com/cloudwatch/>
- Documentation - <https://aws.amazon.com/documentation/cloudwatch/>
- CloudWatch Blog - <https://aws.amazon.com/blogs/aws/category/amazon-cloud-watch/>
- Lambda functions used in the demo scenarios
 - Centralize - <https://github.com/awslabs/cloudwatch-logs-centralize-logs>
 - Customize - <https://github.com/awslabs/cloudwatch-logs-customize-alarms>
 - Analyze - <https://github.com/awslabs/cloudwatch-logs-analyze-data>



**Remember to complete
your evaluations!**