



SDD413

# Amazon S3 Deep Dive and Best Practices

Tim Hunt, Sr. Product Manager, Amazon S3

November 14, 2014 | Las Vegas, NV



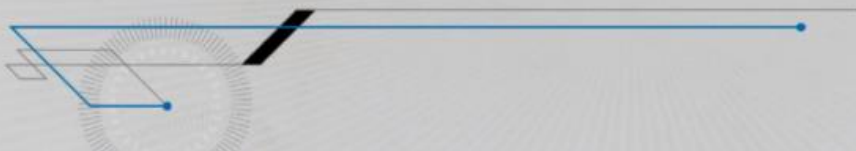
© 2014 Amazon.com, Inc. and its affiliates. All rights reserved. May not be copied, modified, or distributed in whole or in part without the express consent of Amazon.com, Inc.

Come learn about new and existing Amazon S3 features that can help you better protect your data, save on cost, and improve usability, security, and performance. We will cover a wide variety of Amazon S3 features and go into depth on several newer features with configuration and code snippets, so you can apply the learnings on your object storage workloads.

## Topics we will cover

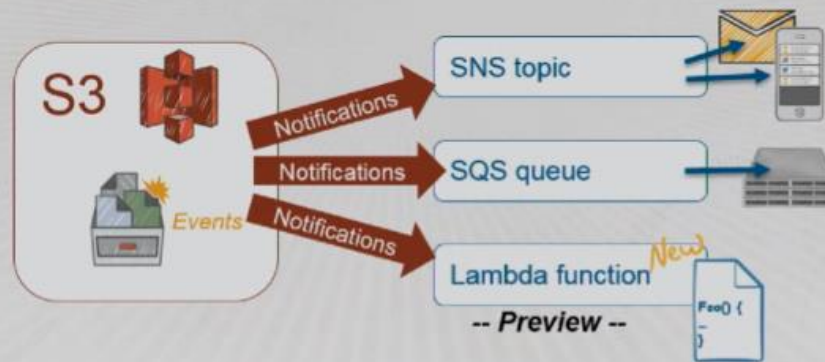
- Amazon S3 event notifications *New*
- Versioning + lifecycle policies
- Top tips for Amazon S3

## Amazon S3 event notifications



# Amazon S3 event notifications

Delivers notifications to Amazon SNS, Amazon SQS, or AWS Lambda when events occur in Amazon S3



Events are things that happen on your S3 buckets, some form of actions like a new object being created with a PUT. You can send the event notification to SNS for sending alerts to email addresses or HTTP endpoints, or to SQS to trigger a workflow, you can set up your EC2 worker fleet to pull those messages off the SQS queue and work on them. Or you can send the event to Lambda to run some code, this allows you take an event in S3 and do some processing in Lambda on-demand.

## What's in it for you?



Integration – A new surface on the Amazon S3 “building block” for event-based computing



Speed – typical time to send notifications is less than a second



Simplicity – Avoids proxies or polling to detect changes

This is event-based computing

# Now what can you do?

Customers have told us about powerful applications ...



Transcoding  
media files



Updating  
data stores



Processing  
data/log files



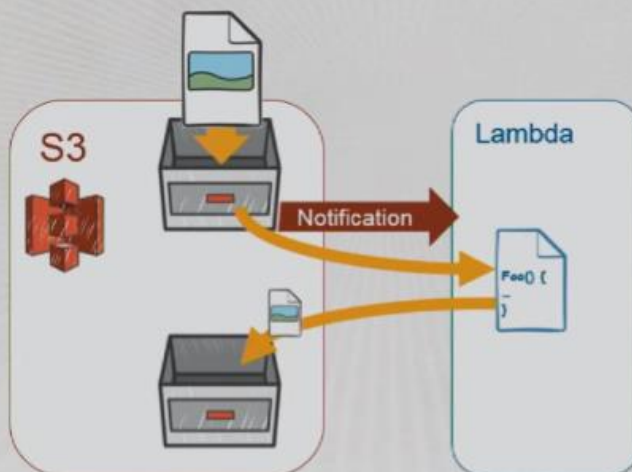
Object  
change alerts

... and we look forward to seeing what you create.

You can use notifications in any of the above ways like resizing and image file that arrives in your S3 bucket. You can process data files or log files as they arrive in S3 and forward the results to your DynamoDB table. You can also have your CloudWatch logs delivered to S3 and do something with the data as they come in.

## Event notification demo #1

Using Amazon S3 event notifications and a Lambda function to automatically generate thumbnails of new images



# Amazon S3 Console Screen Shot

▼ Events

Event Notifications enable you to send alerts or trigger workflows. Notifications can be sent via [Amazon Simple Notification Service \(SNS\)](#) or [Amazon Simple Queue Service \(SQS\)](#) or to a [Lambda function](#) (depending on the bucket location).

Name  ⓘ

Events  ⓘ

Send To ☒ SNS topic ☐ SQS queue ☐ Lambda function ⓘ

SNS topic  ▼

S3 must have permission to publish to the topic from this source bucket. See the [Developer Guide](#).

► Versioning

► Lifecycle

► Tags

We set this up for the S3 bucket

# Amazon S3 Console Screen Shot

▼ Events

Event Notifications enable you to send alerts or trigger workflows. Notifications can be sent via [Amazon Simple Notification Service \(SNS\)](#) or [Amazon Simple Queue Service \(SQS\)](#) or to a [Lambda function](#) (depending on the bucket location).

Name  ⓘ

Events  ⓘ

Send To ☒ SNS topic ☐ SQS queue ☐ Lambda function ⓘ

SNS topic  ▼

S3 must have permission to publish to the topic from this source bucket. See the [Developer Guide](#).

► Versioning

► Lifecycle

► Tags

- RRSObjectLost
- ObjectCreated (All)
- Put
- Post
- Copy
- CompleteMultipartUpload

You can choose the event that you want to create the notification for like PUTS, All, etc. then you choose where to send the event to like a Lambda and supply both the **arn** of the Lambda function and the invocation role, which is a way to control permission to the lambda function by giving S3 the authority to invoke that lambda.



# Amazon S3 Console Screen Shot

▼ Events

Event Notifications enable you to send alerts or trigger workflows. Notifications can be sent via [Amazon Simple Notification Service \(SNS\)](#) or [Amazon Simple Queue Service \(SQS\)](#) or to a [Lambda function](#) (depending on the bucket location).

Name  ⓘ

Events  ⓘ

Send To ☐ SNS topic ☐ SQS queue ☒ Lambda function ⓘ

AWS Lambda functions consist of custom code that you can execute in response to an event notification.

Enter the Amazon Resource Names (ARNs) of a Lambda function to run for this event and an invocation IAM role. You can visit the [Amazon Lambda Console](#) to create a Lambda function or look up its ARN. You can visit the [IAM Console](#) to create or look up an invocation role. See the [Developer Guide](#).

Lambda function ARN

Invocation role ARN  ⓘ

# Amazon S3 Console Screen Shot

► Static Website Hosting

► Logging

▼ Events

Event Notifications enable you to send alerts or trigger workflows. Notifications can be sent via [Amazon Simple Notification Service \(SNS\)](#) or [Amazon Simple Queue Service \(SQS\)](#) or to a [Lambda function](#) (depending on the bucket location).

Name	Event(s)	Type
myNotification	ObjectCreated (All)	<a href="#">CloudFunction</a> ✓✕

► Versioning

► Lifecycle

► Tags

► Requester Pays

We have now created the notification for this S3 bucket. We can create many notifications for a single bucket.

# Lambda function code

```
function handler(event, context) {  
  // read values from the event  
  var srcBucket = event.Records[0].s3.bucket.name;  
  var srcKey     = event.Records[0].s3.object.key;  
  var dstBucket = srcBucket + "resized";  
  var dstKey     = "resized-" + srcKey;  
  
  // sanity check: validate that source and destination are different buckets  
  if (srcBucket == dstBucket) {  
    return;  
  }  
  
  // make sure it's a jpg or png  
  var imageType = srcKey.match(/\.[^.]*/)[1];  
  if (imageType != "jpg" && imageType != "png") {  
    return;  
  }  
}
```

*Continued...*

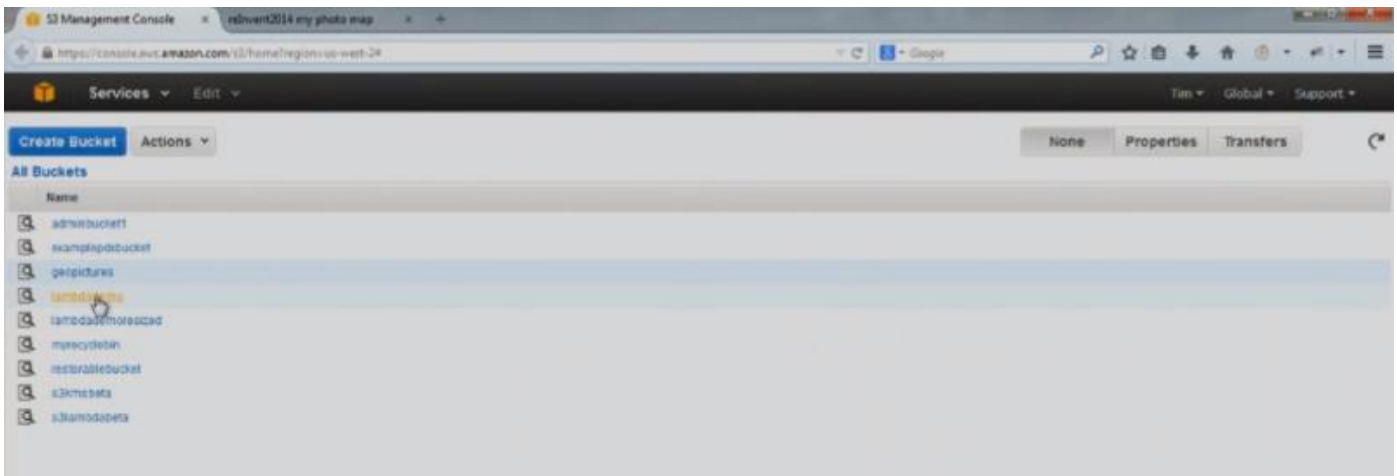
This is a NodeJS code for our lambda function

# Lambda function code – continued

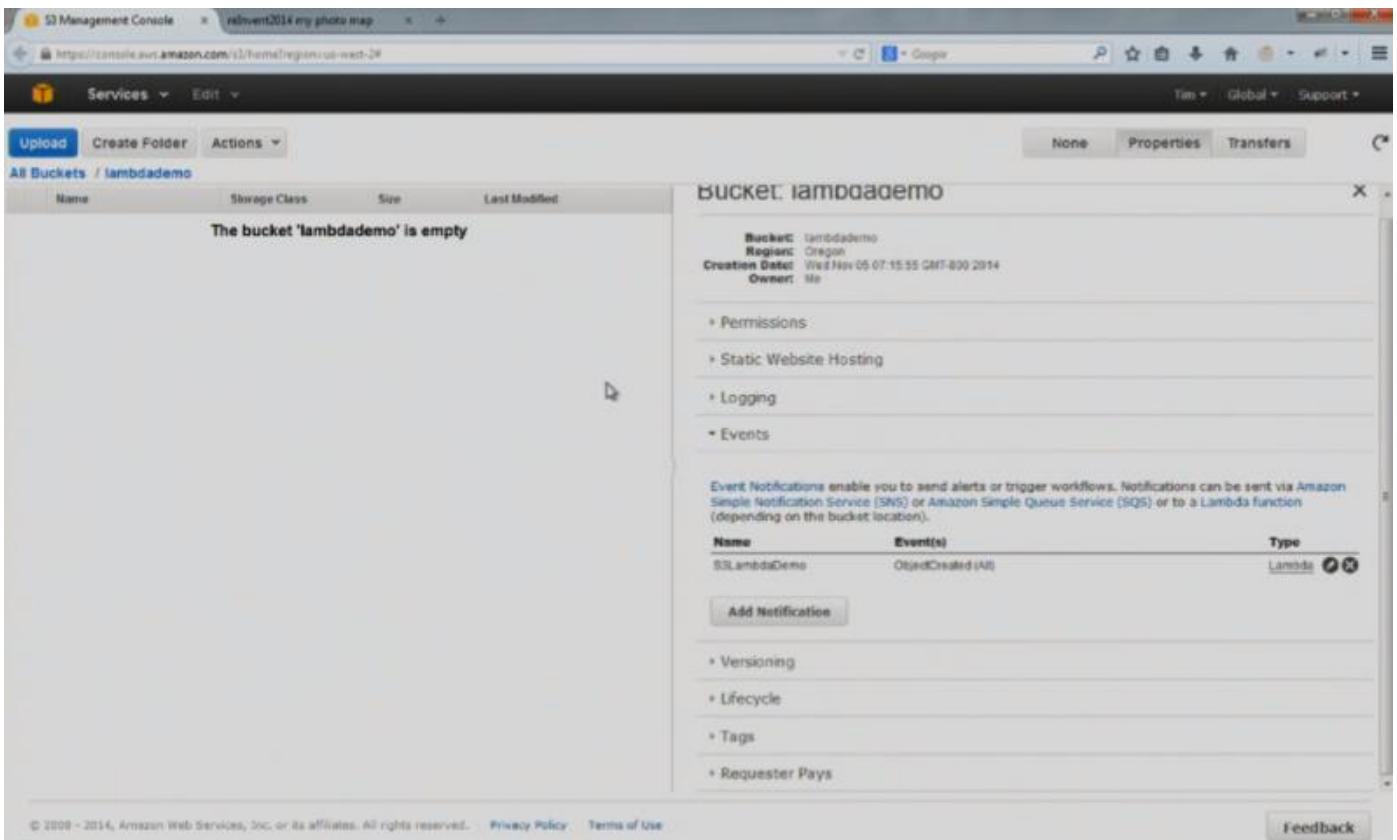
```
// download image from s3 into buffer  
function download(next) {  
  s3.getObject({Bucket: srcBucket, Key: srcKey}, next);  
},  
// generate the thumbnail  
function tranform(response, next) {  
  gm(response.Body).size(function(err, size) {  
    this.resize(width, height).toBuffer(imageType, function(err, buffer) {  
      if (err) { next(err);}  
      else {next(null, response.ContentType, buffer);}  
    });  
  });  
},  
// Put into S3  
function upload(data, next) {  
  s3.putObject({Bucket: dstBucket, Key: dstKey, Body: data, ContentType: contentType}, next);  
}
```

We are using the ImageMagick library that is available in Lambda to create the thumbnails

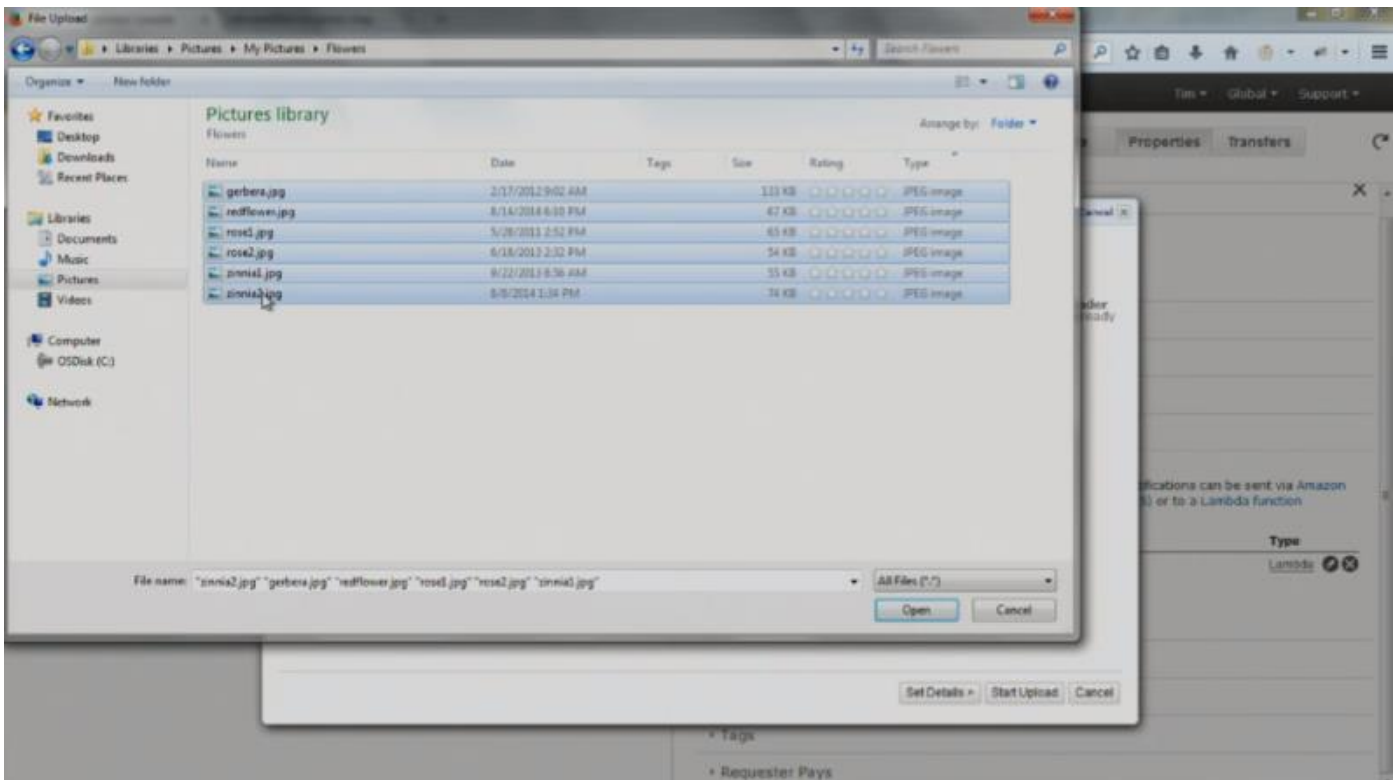
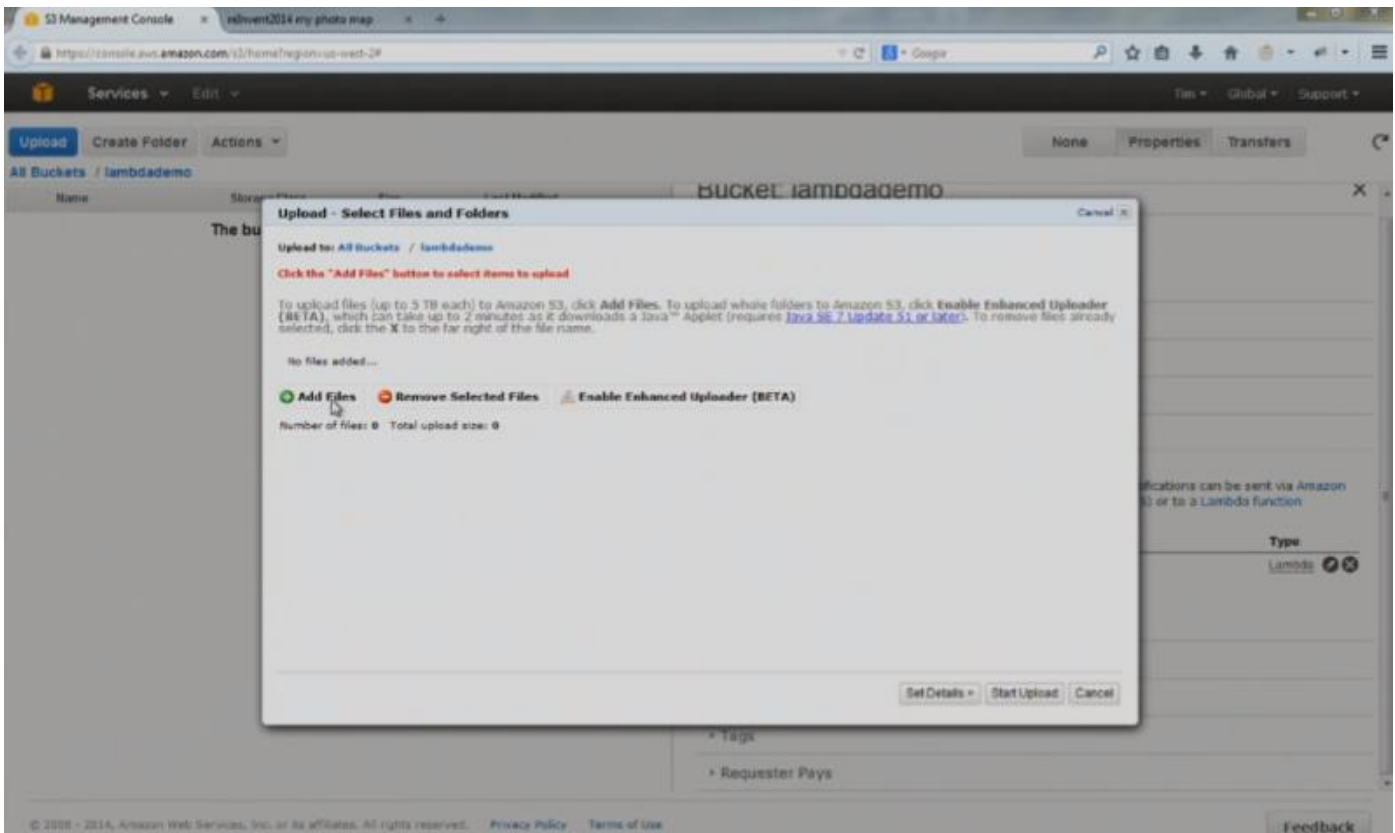
# Demo



This is the source S3 bucket

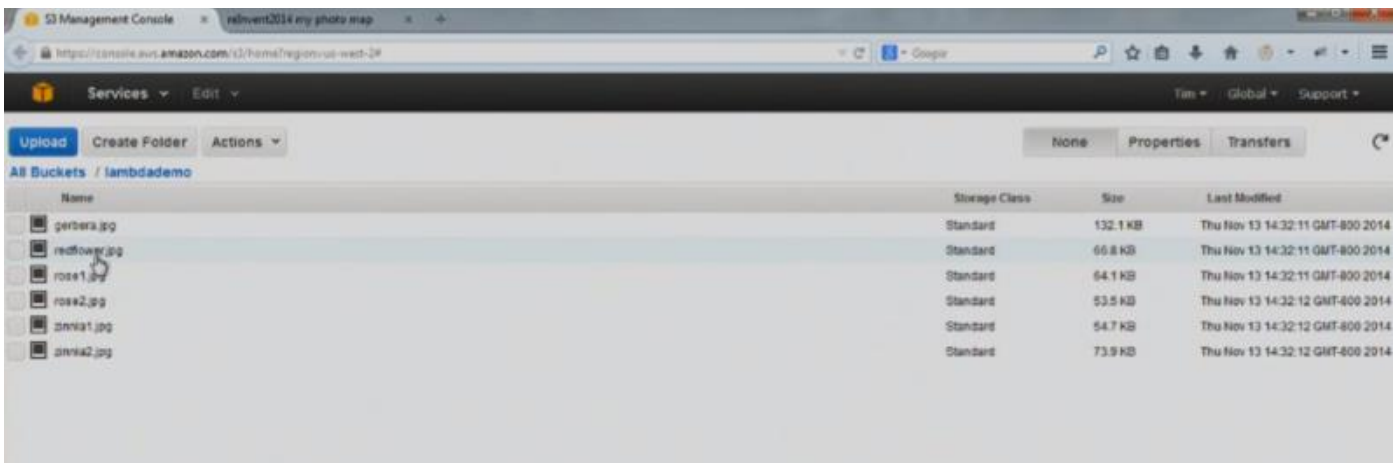
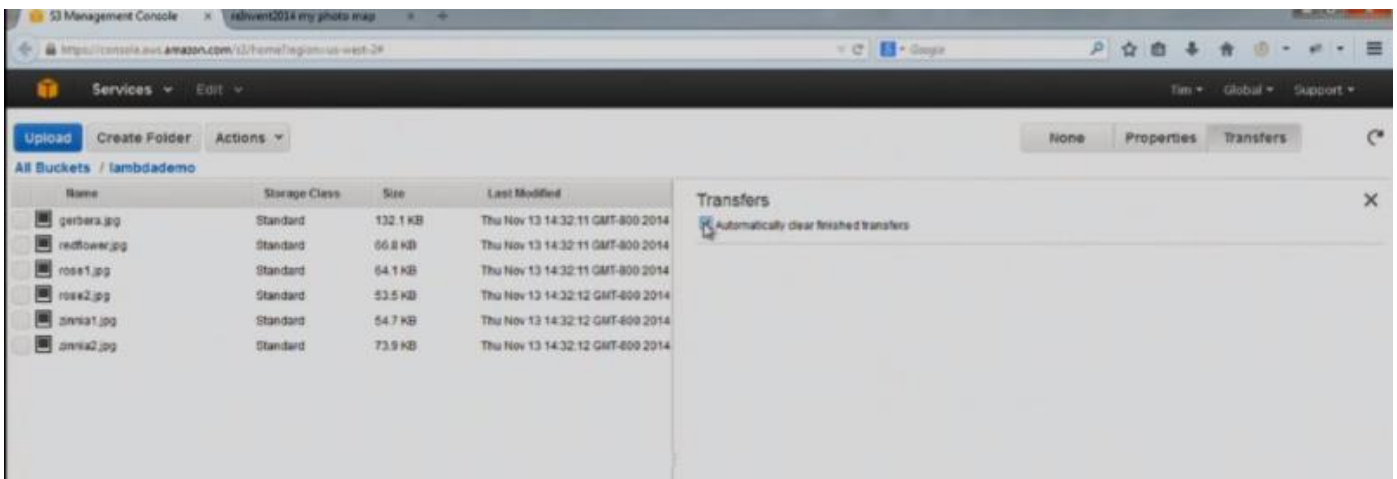
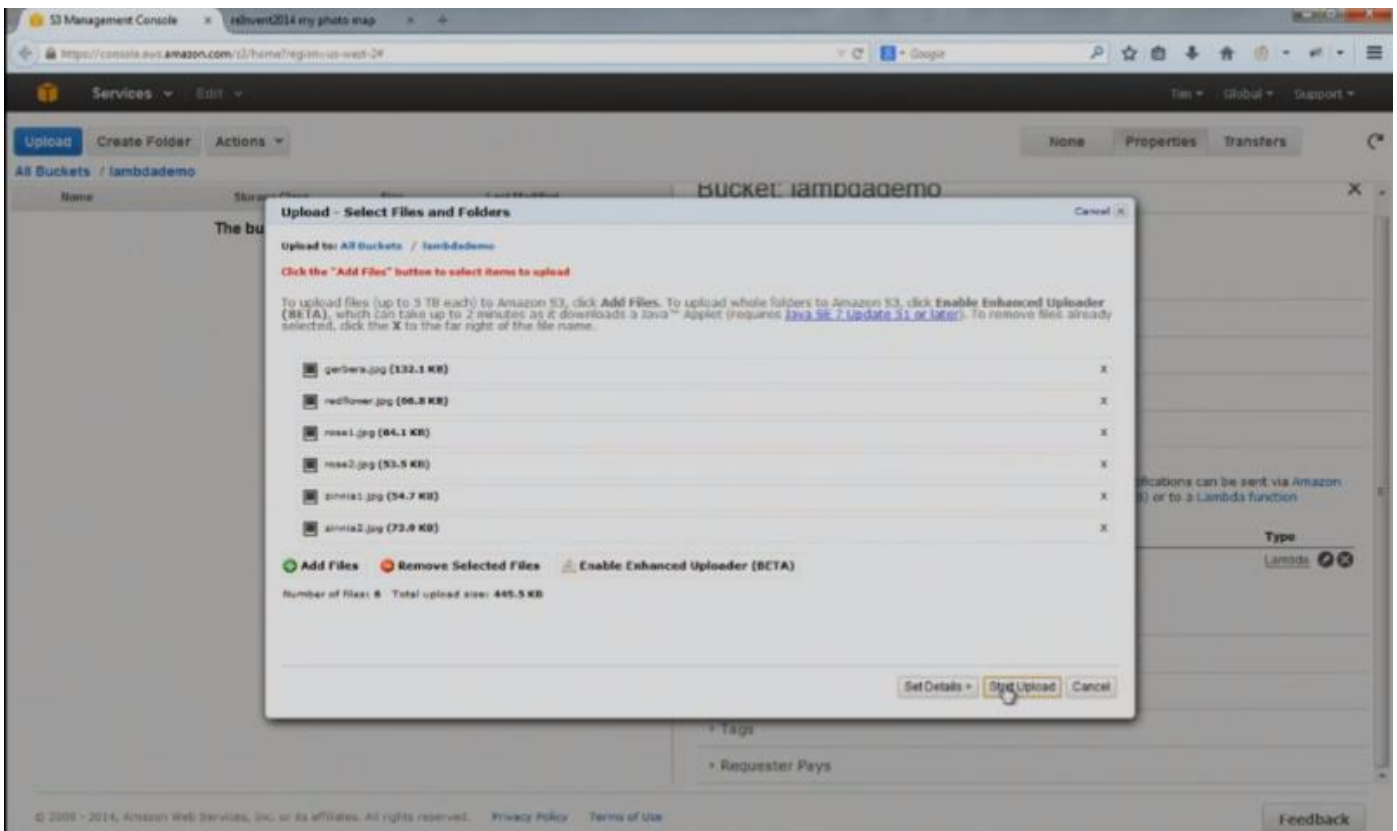


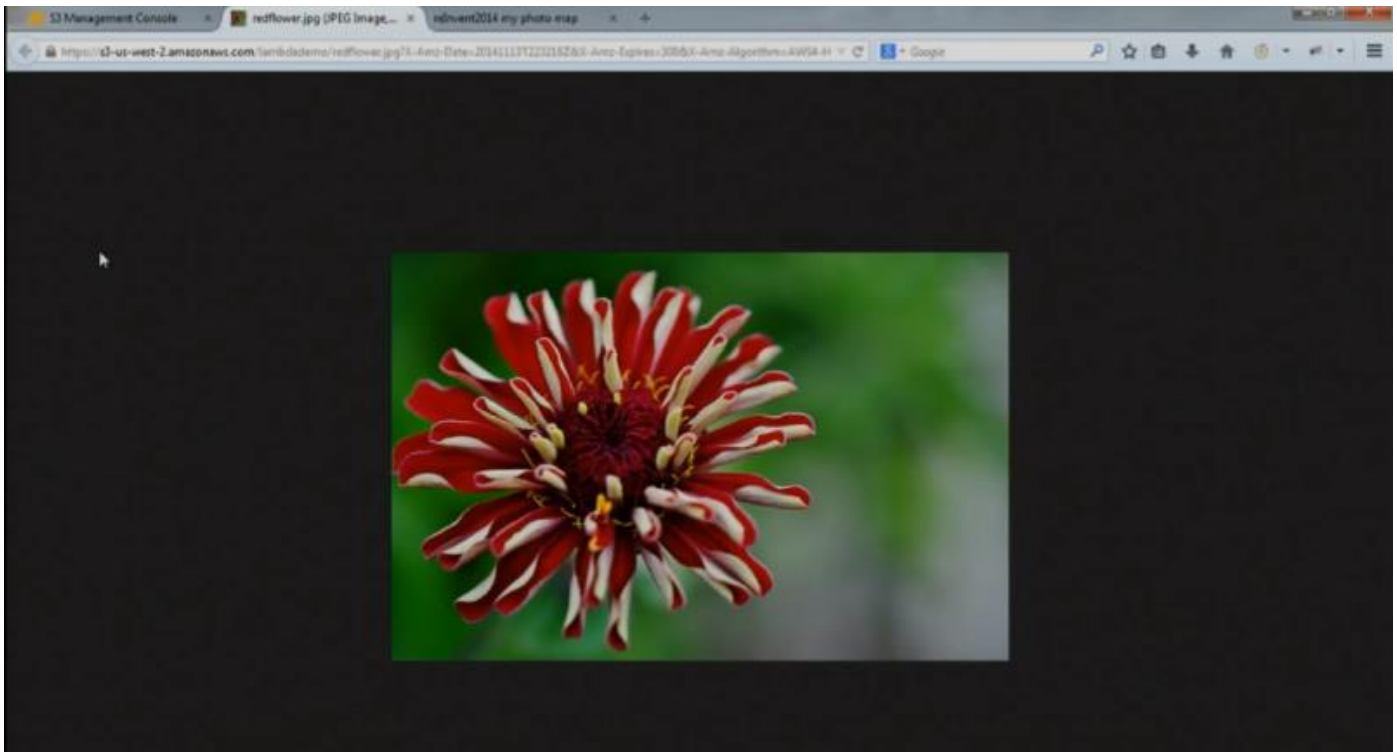
We have set up an ObjectCreated(All) notification for this S3 bucket to be sent to the S3LambdaDemo lambda function



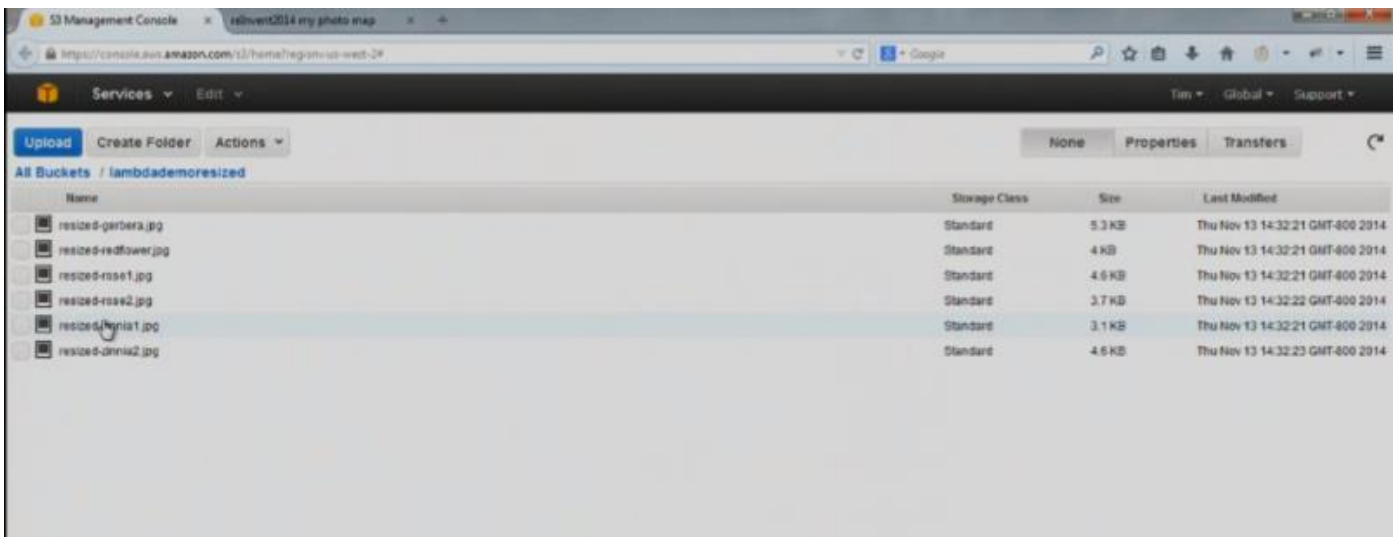
We can now upload some files into the source S3 bucket as above



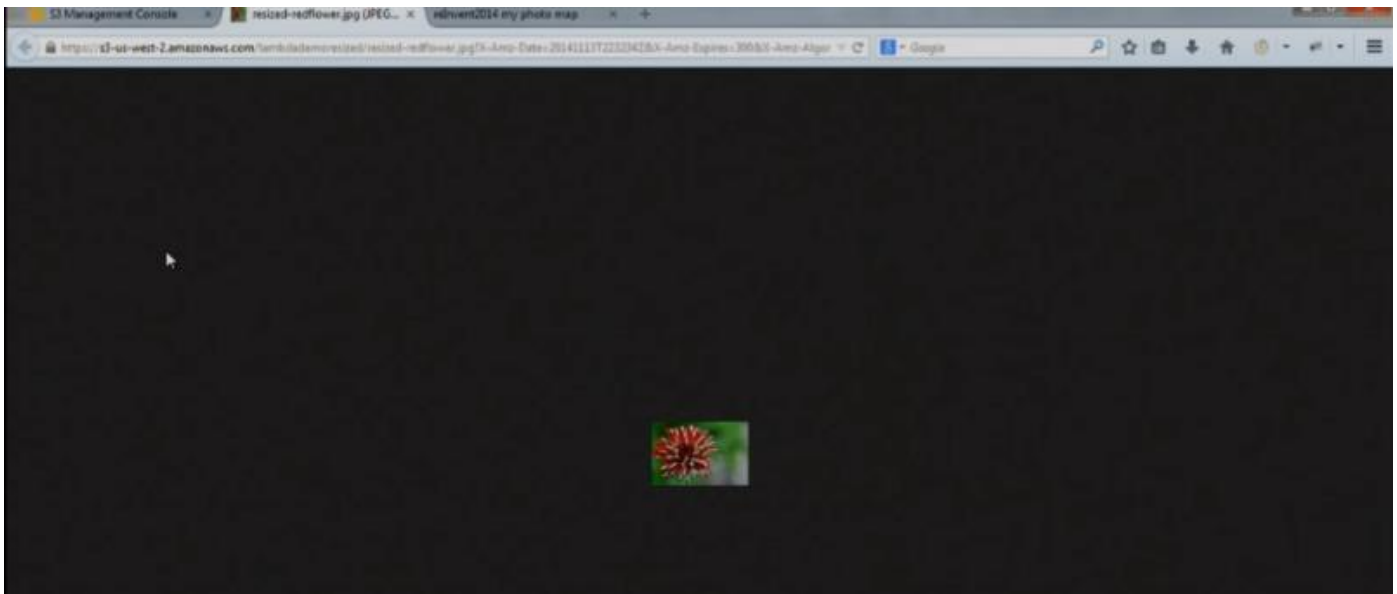




This is an example of one of the uploaded pictures that the Lambda will create a thumbnail for and save into the destination S3 bucket.



We can now see the created thumbnails images in our resized bucket, these were each created by the Lambda function using the S3 PutObject event notification as a trigger.



This is the equivalent thumbnail

## Event notification details

- No additional charge from Amazon S3; you pay only for associated use of Amazon SNS, Amazon SQS, Lambda
- Configured at the bucket level and supports different destinations for different events
- Today supports PUT, POST, COPY, MultiPartComplete, and RRSObjectLost events; more to come

## Event notification details, continued

- Highly reliable – designed for nine '9's with at least once delivery
- The destination topic, queue, or Lambda function must be in the same AWS region as the S3 bucket

## Sample notification JSON

```
{
  "Records": [
    {
      "eventVersion": "2.0",
      "eventSource": "aws:s3",
      "awsRegion": "us-west-2",
      "eventTime": "2014-11-06T18:02:44.551Z",
      "eventName": "ObjectCreated:Put",
      "userIdentity": {
        "principalId": "A1AM2HK7YZABCD"
      },
      "requestParameters": {
        "sourceIPAddress": "10.240.9.17"
      },
      "responseElements": {
        "x-amz-request-id": "0DAFE1024D7E1234",
        "x-amz-id-2": "YVbY/Q7Nst1Rao68w1z1VfJr4H8x3aaEzVgPbzN8J0WE19mkPa7pqwertyuiop+m"
      },
    }
  ]
}
```

*Continued...*

The content of the SNS notification that was sent to trigger the lambda function are in JSON format as above,



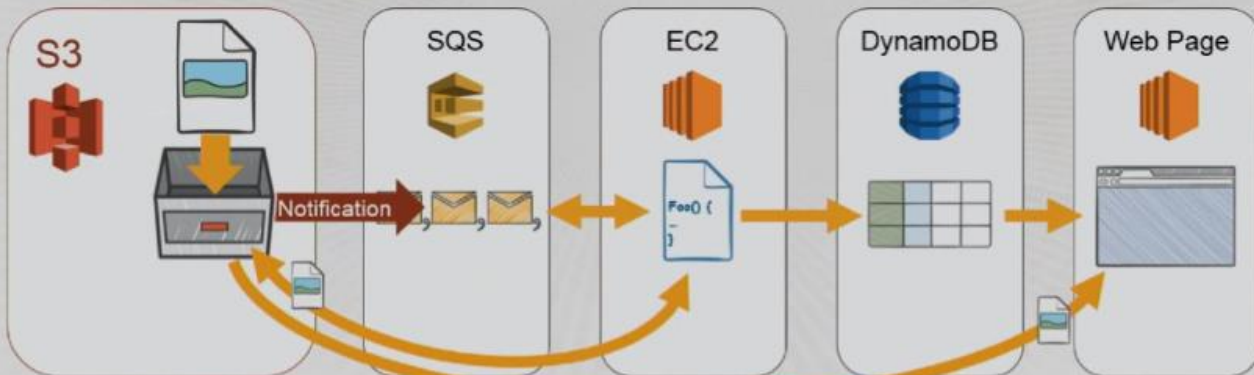
# Sample notification JSON - continued

```
"s3":{
  "s3SchemaVersion":"1.0",
  "configurationId":"lambdaeventdemo",
  "bucket":{
    "name":"lambdadedemo",
    "ownerIdentity":{
      "principalId":"A1AM2HK7YZABCD"
    },
    "arn":"arn:aws:s3:::lambdadedemo"
  },
  "object":{
    "key":"Geranium.jpg",
    "size":102010,
    "eTag":"e03cf0062f8e99a32dd4937a1b2c31b2"
  }
}
```

Some of the information pertains to the S3 bucket and the object. There is an issue with S3's eventual consistency model if you already have an image in the S3 folder by the same name that you want to replace. The notification can be faster than the eventual object replacement and your lambda might end up creating the thumbnail for the old image that was already in the S3 bucket. A way to mitigate this will be to compare the eTag that is in the Notification message with the name of the object in the S3 bucket or you can use versioning.

## Event notification demo #2

Using Amazon S3 event notifications with Amazon SQS and Java code to create thumbnails and record data in Amazon DynamoDB



We are going to use SQS as the message bus that gets the notifications as messages being stored on the queue. Then we have a Java application running in EC2 that will be pulling those messages off the SQS queue, then read the message to be able to go pull the object out of S3, it will then look at the data in the S3 object and put some of the data into DynamoDB and also creating the image thumbnail and putting it back in another S3 bucket. We are also going to have a web page that displays the information about things that are getting stored in DynamoDB.

# Main loop Java code

```
while (true) {  
    // Long poll for the messages  
    List<Message> messages = sqs.receiveMessage(receiveMessageRequest).getMessages();  
  
    for (Message message : messages) {  
  
        // Process the message  
        processMessage(message.getBody());  
  
        // Delete the message after we are done.  
        sqs.deleteMessage(myQueueUrl, message.getReceiptHandle());  
    }  
}
```

Authentication and access control not shown for brevity, but published best practices should be followed.

This is the main loop of the Java code, it starts by polling the SQS queue for messages

# processMessage Java code

```
private static void processMessage(String messageBody) {  
  
    // Parse the bucket and key from the event notification  
    S3EventNotificationRecord item = S3EventNotificationItem.parseJson(messageBody)  
        .getRecords().get(0);  
    String bucketName = item.getS3().getBucket().getName();  
    String objectName = item.getS3().getObject().getObjectKey();  
    String thumbnailFileName = thumbnailPrefix + objectName;  
  
    if (objectName.startsWith(thumbnailPrefix)) {  
        return; // skip thumbnails to avoid recursion  
    }  
  
    // Download the file from S3  
    s3.getObject(new GetObjectRequest(bucketName, objectName), new File(tempFileName));  
}
```

*Continued...*

Authentication and access control not shown for brevity, but published best practices should be followed.

It parses the notification message to pull out some of the details we want to store in the DynamoDB table for display.

# processMessage Java code - Continued

```
// Read the GPS info from the file
double[] gps = readGpsFromEXIF(tempFileName);

// Put the photo info into DynamoDB for public website display
putNewPhotoInDynamo (bucketName, objectName, gps, "http://" +
    sourceBucketName + ".s3.amazonaws.com/" + thumbnailFileName);

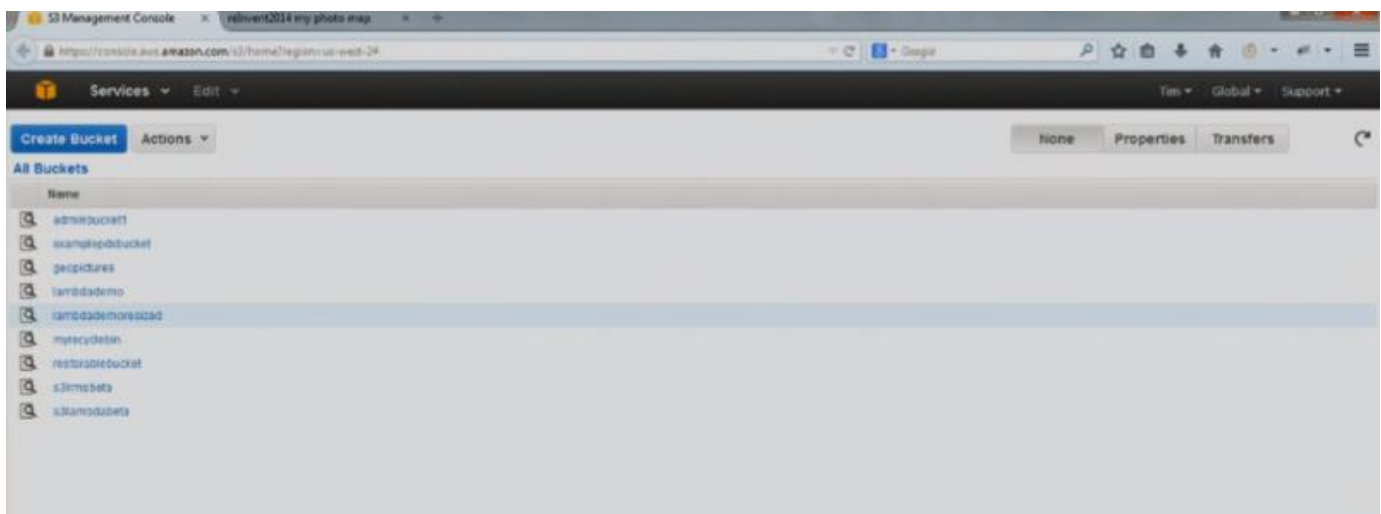
// Generate a thumbnail and put it into s3
generateThumbnail(tempFileName, thumbnailFileName);
s3.putObject(sourceBucketName, thumbnailFileName, new File(thumbnailFileName));

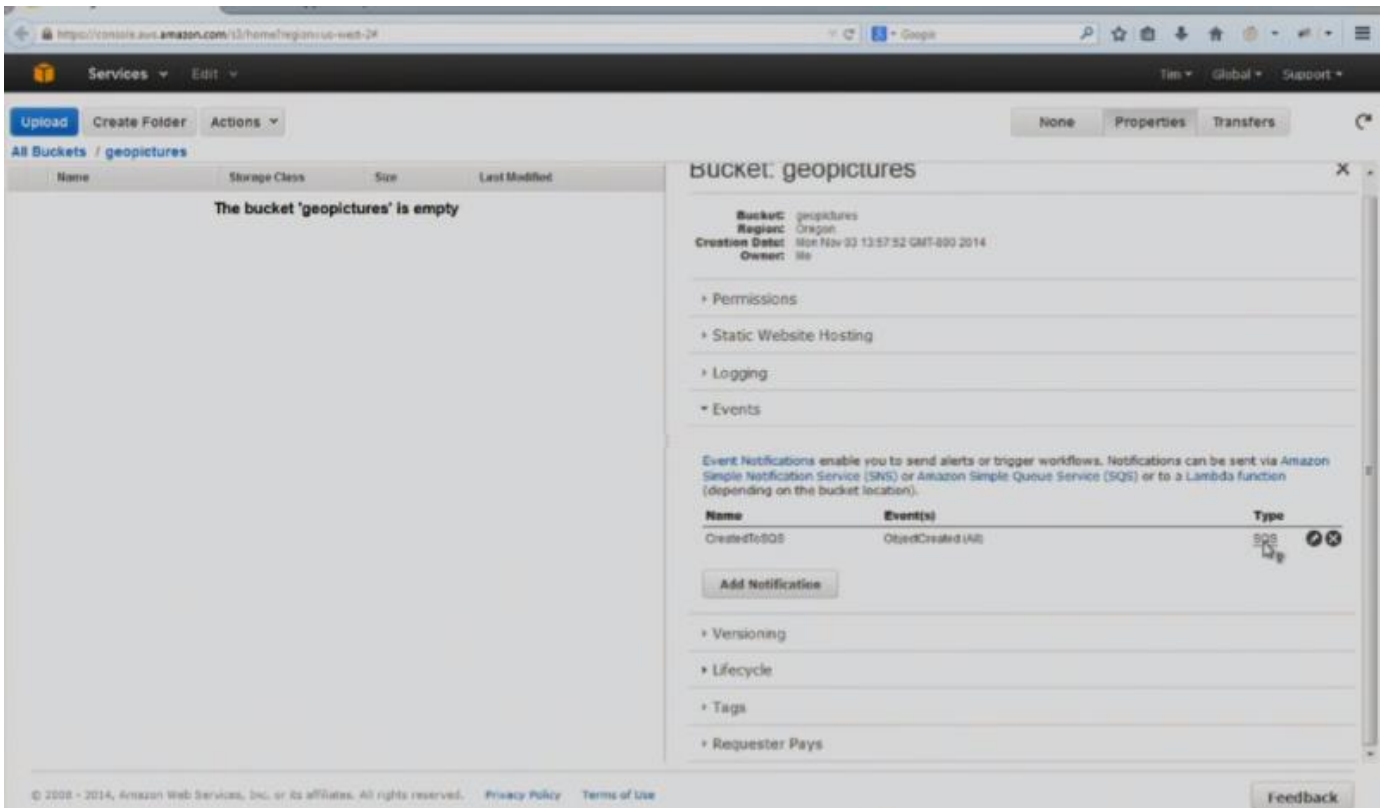
// Notify our web page to add it to the map
sqs.sendMessage(sqs.createQueue(realTimeQueueName).getQueueUrl(), key);
}
```

Authentication and access control not shown for brevity, but published best practices should be followed.

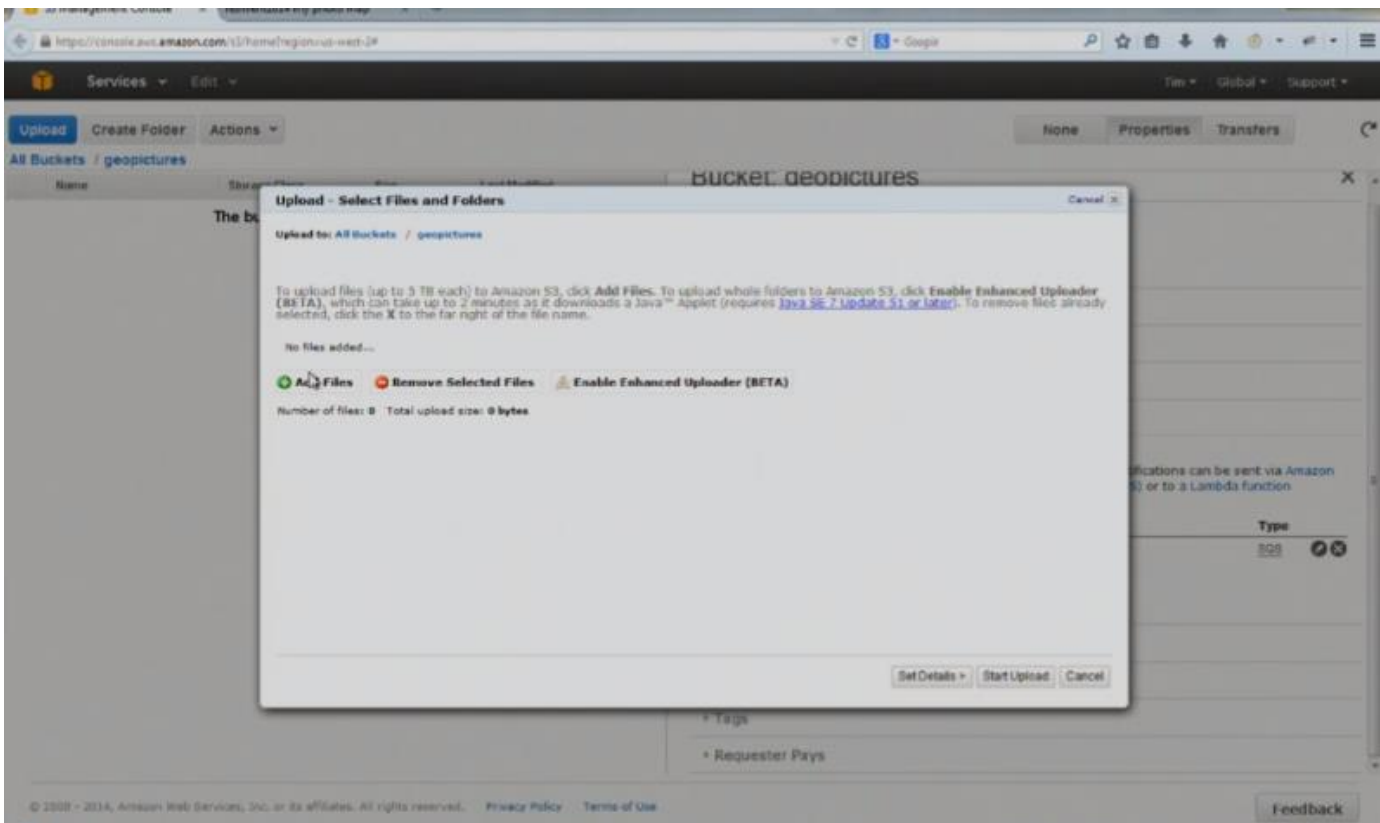
This 2<sup>nd</sup> SQS queue is used to trigger the web page that displays the results to the user so that the page knows there is a new object to display. Note that we are storing the thumbnail image in the same bucket as the original image

## Demo

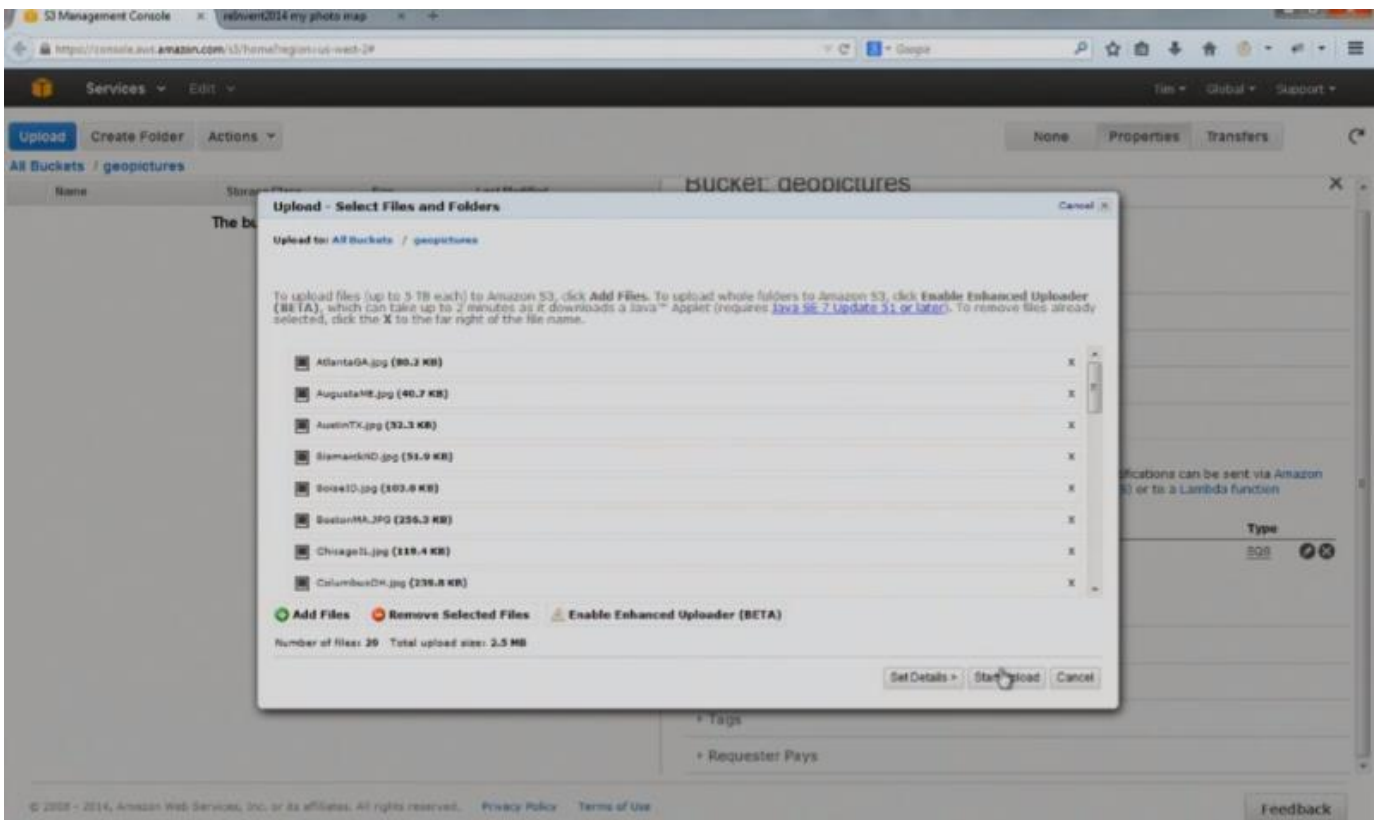
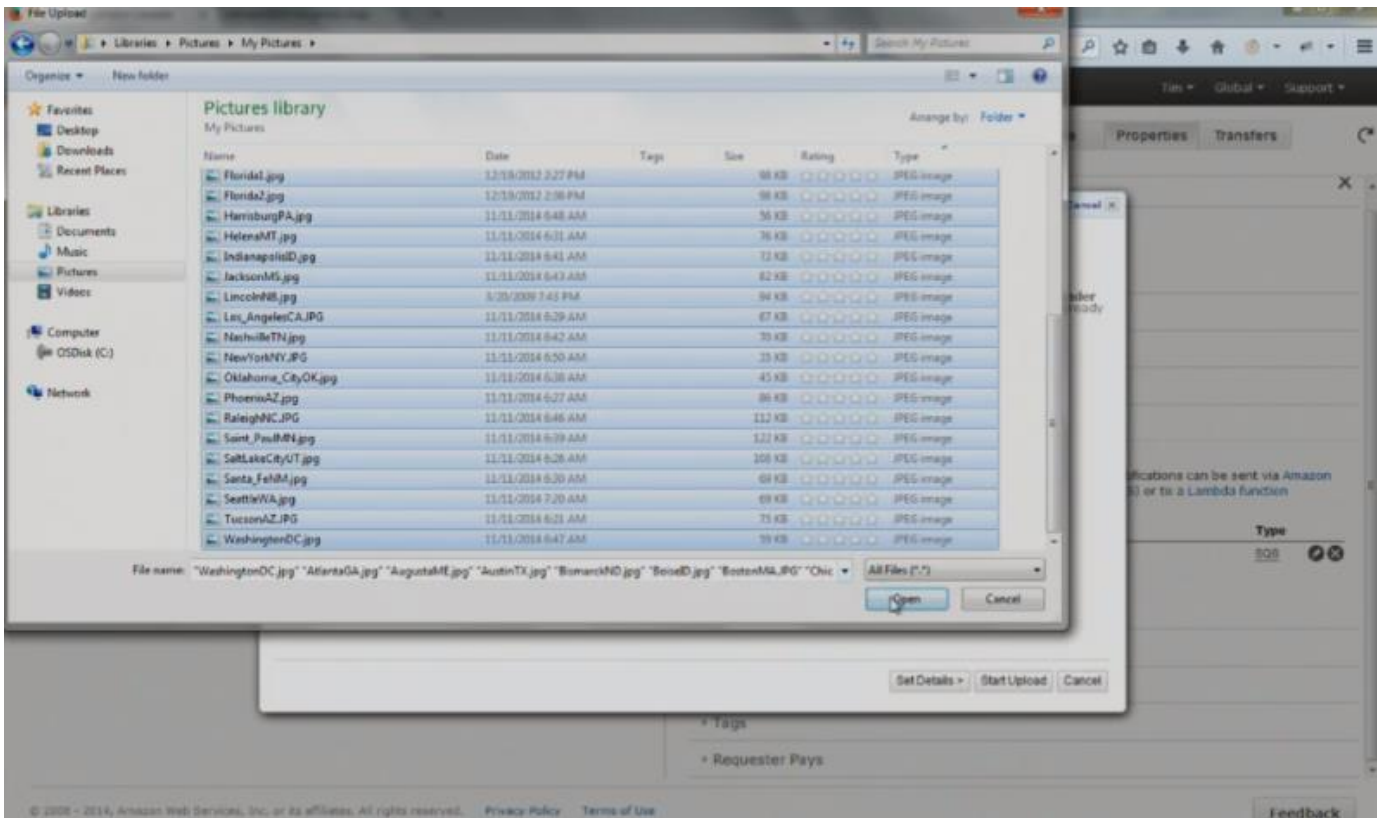




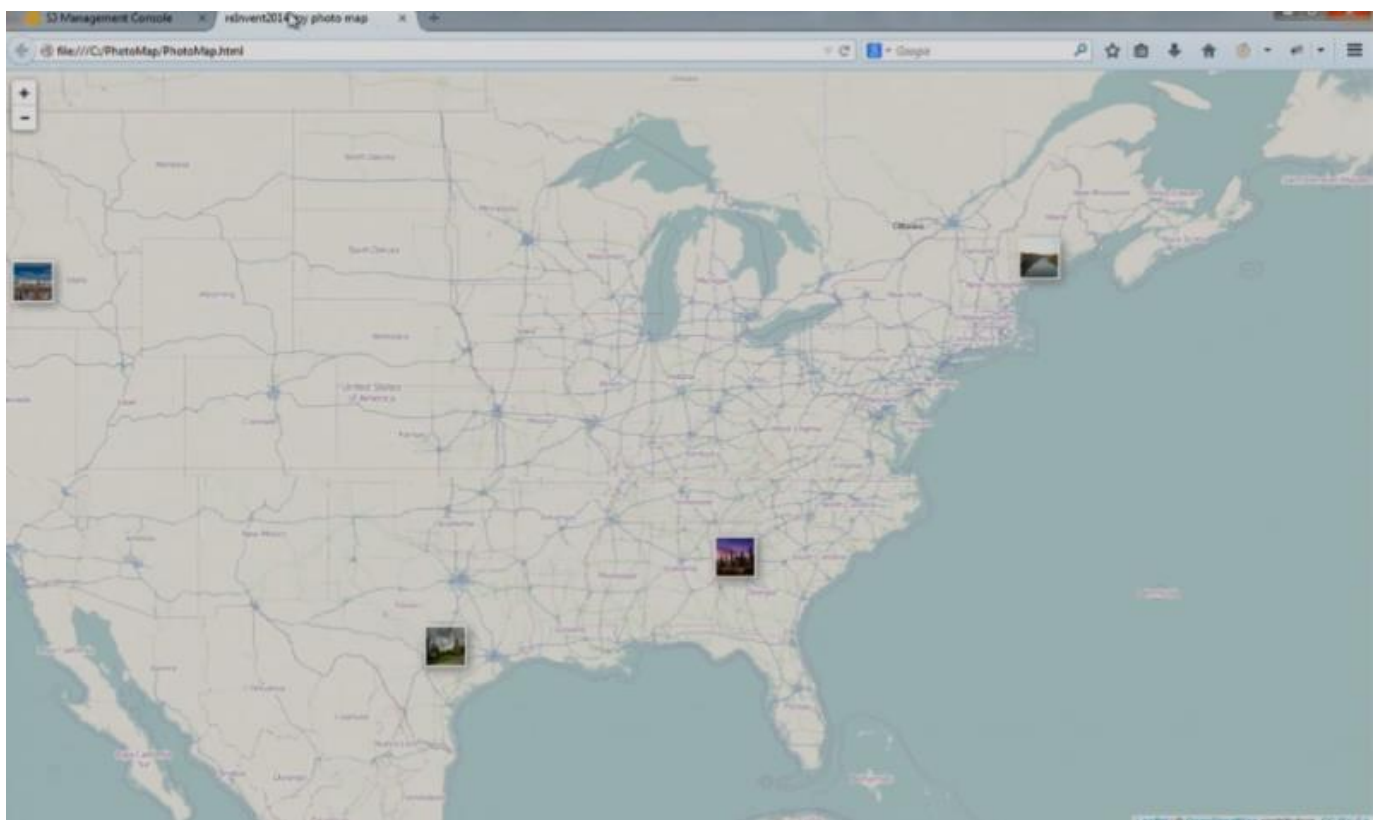
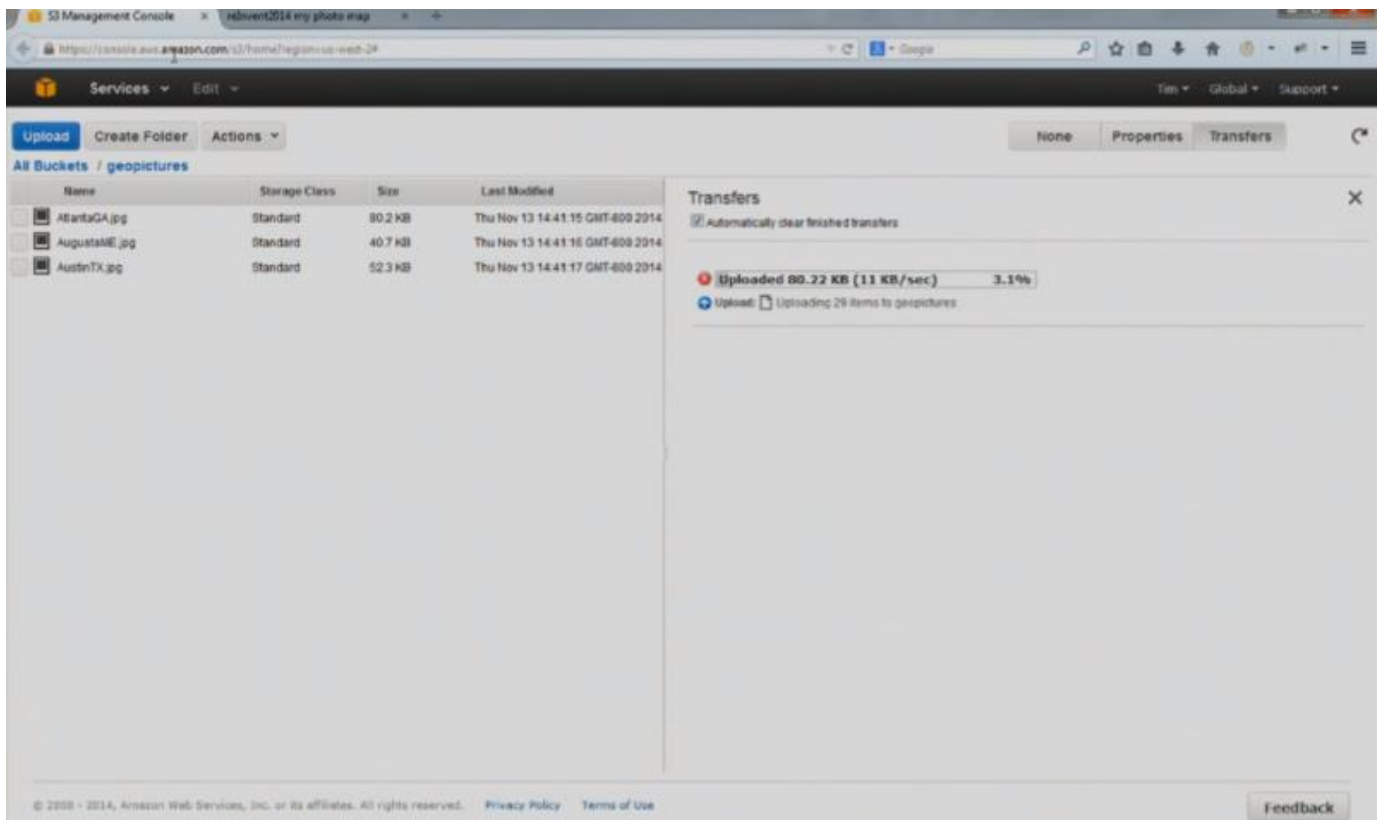
This is our source S3 bucket, it is set up to send notifications to the SQS queue

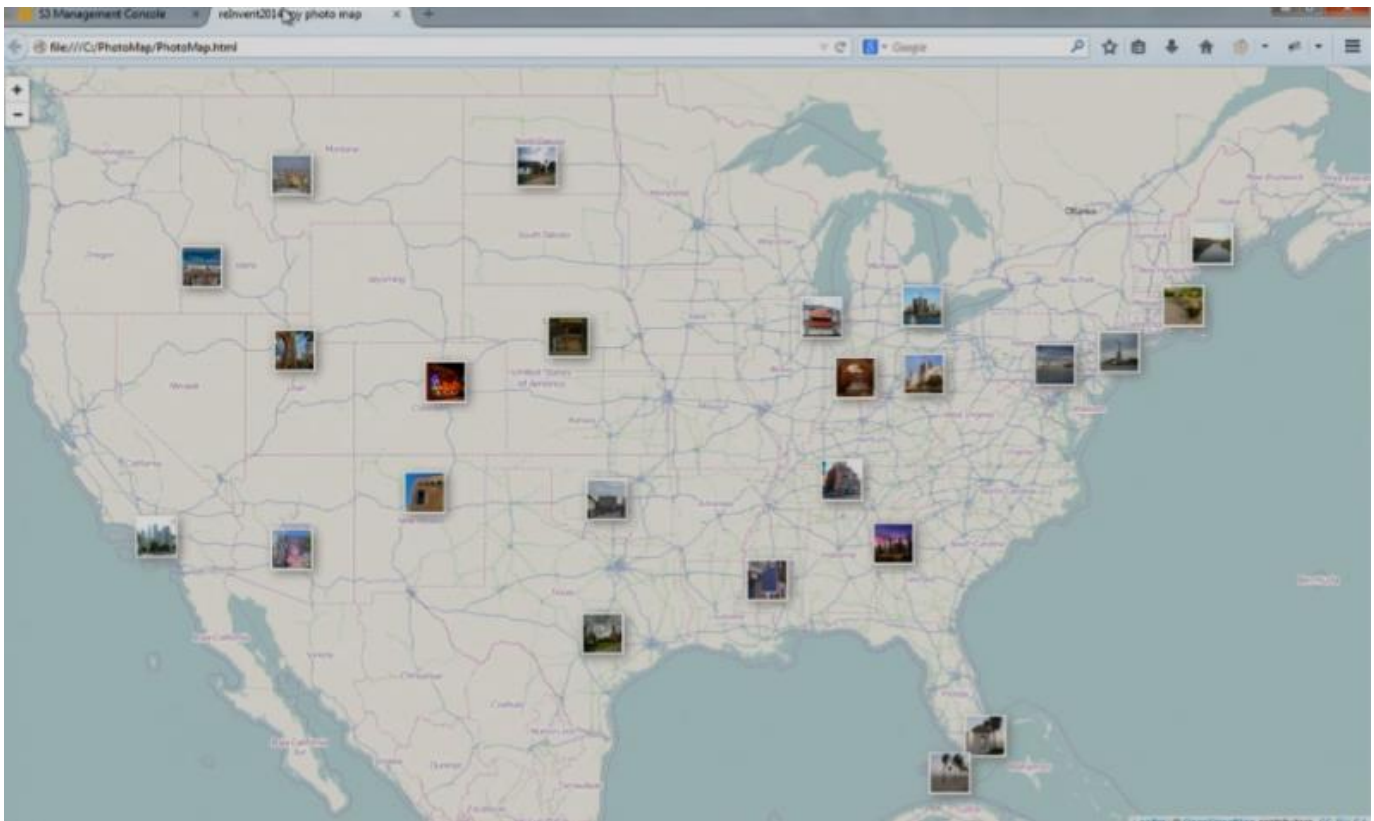




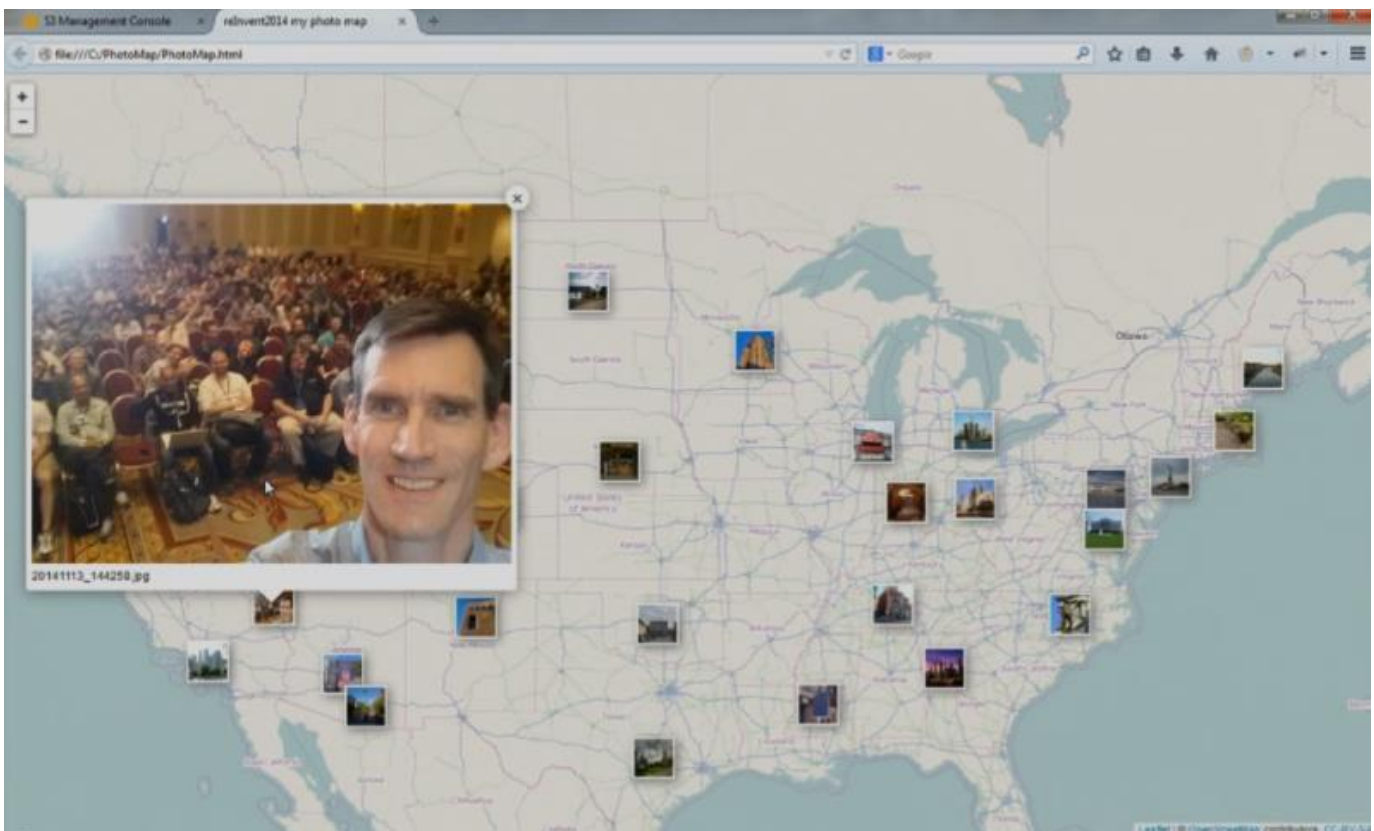


We then upload a bunch of files to the S3 bucket as above





The images are being displayed on the web page via the 2<sup>nd</sup> SQS message notifications as they become available and their thumbnails are being displayed



We can even do a real-time photo upload to the source bucket and see the thumbnail immediately as above



# Key takeaways for event notifications

- Enables event-based computing
- Destination choices of Amazon SNS, Amazon SQS, or Lambda
- No additional charge from Amazon S3

## Versioning + lifecycle policies



### Versioning

- Protects from accidental overwrites and deletes with no performance penalty
- Generates a new version with every upload
- Allows easily retrieval of deleted objects or roll back to previous versions
- Three states of an Amazon S3 bucket
  - Default – Un-versioned
  - Versioning-enabled
  - Versioning-suspended



# Lifecycle policies

- Provides automatic tiering to a different storage class and cost control
- Includes two possible actions:
  - Transition: archives to Amazon Glacier after a specified amount of time
  - Expiration: deletes objects after a specified amount of time
- Allows for actions to be combined – archive and then delete
- Supports lifecycle control at the prefix level

## Recycle bin with automatic cleaning



Versioning



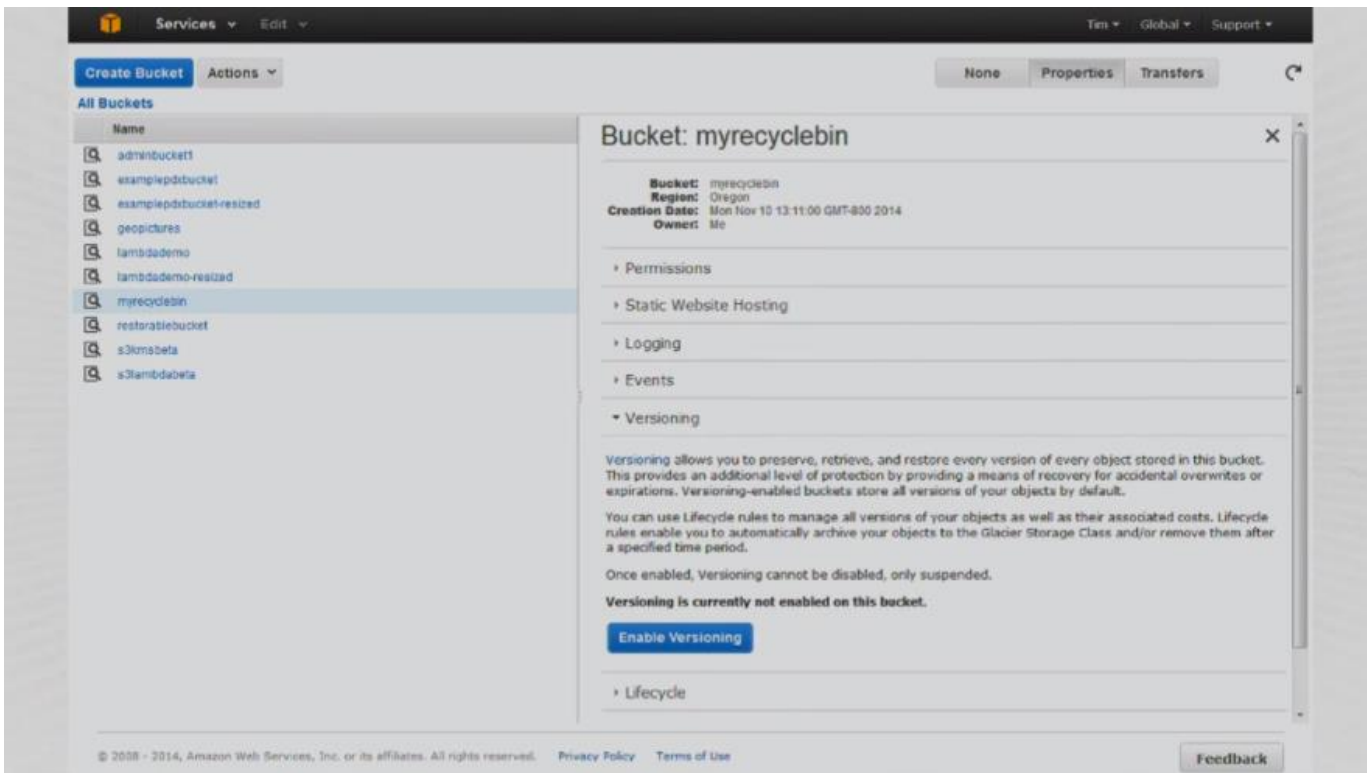
Recycle bin



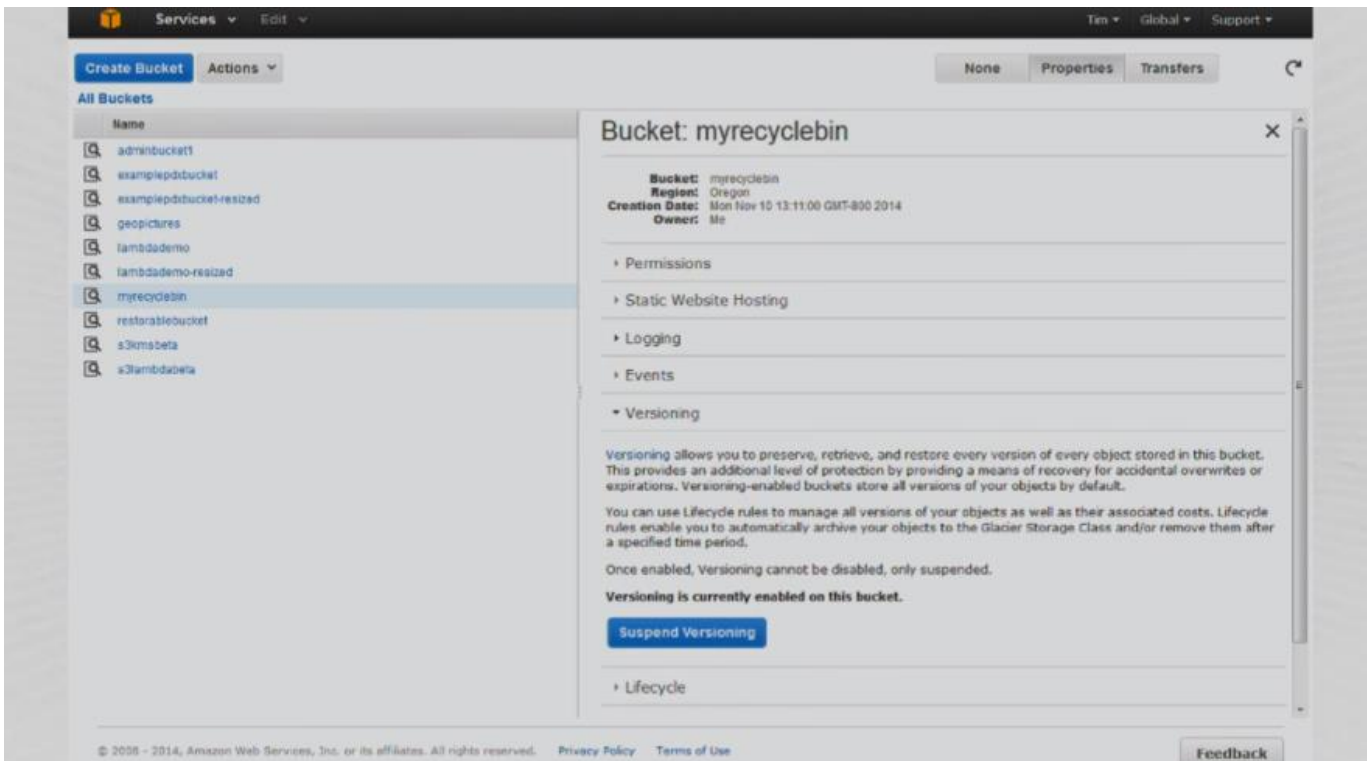
Lifecycle  
policies



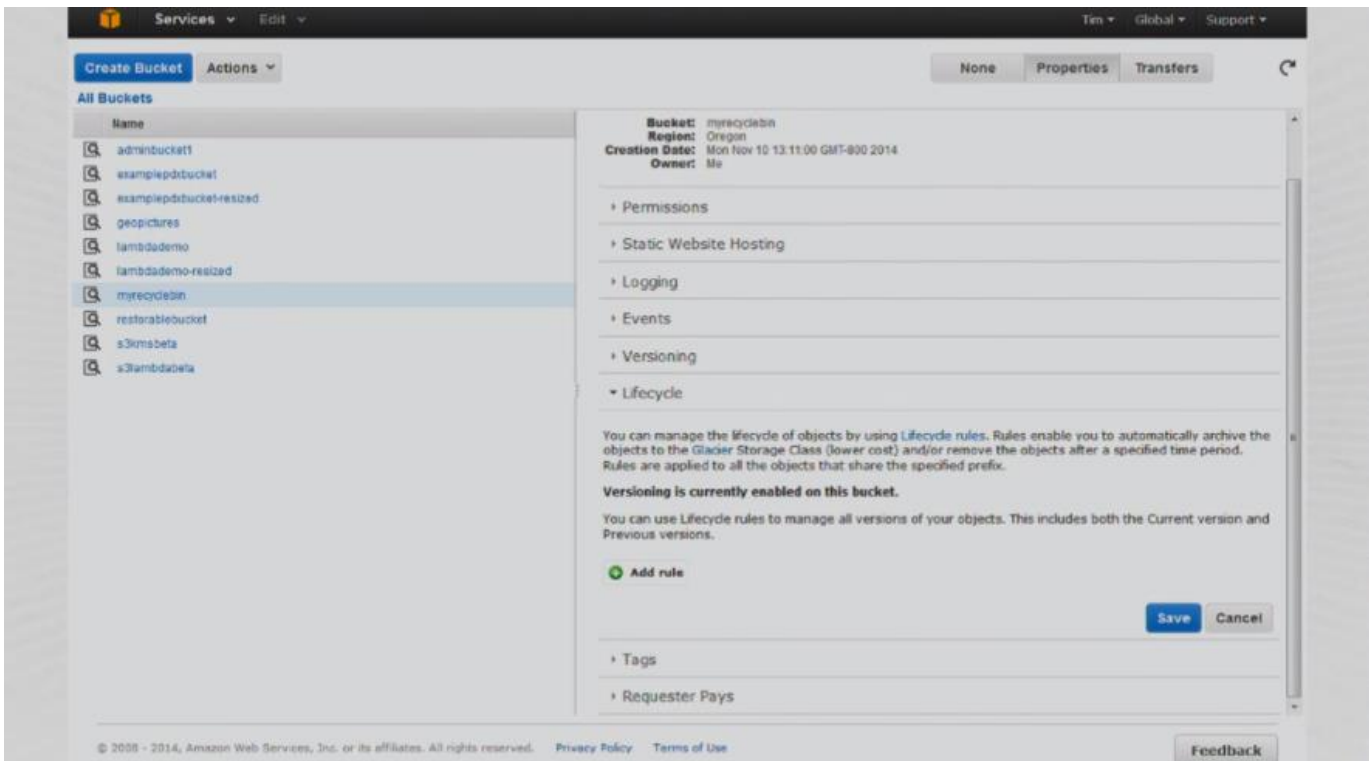
Automatic  
cleaning



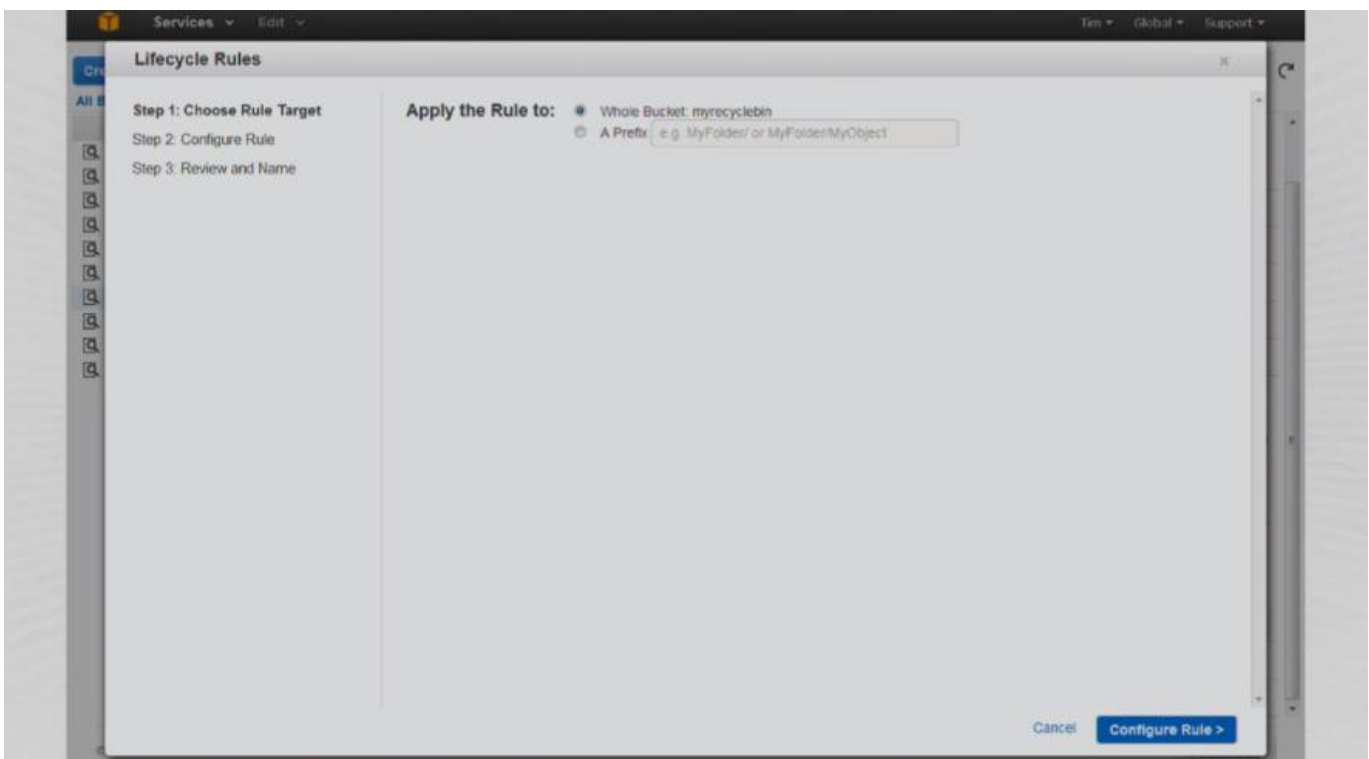
Let us now see how to set up this recycle bin functionality on your S3 buckets. You can simply click the Enable Versioning button to have it turn on for this bucket



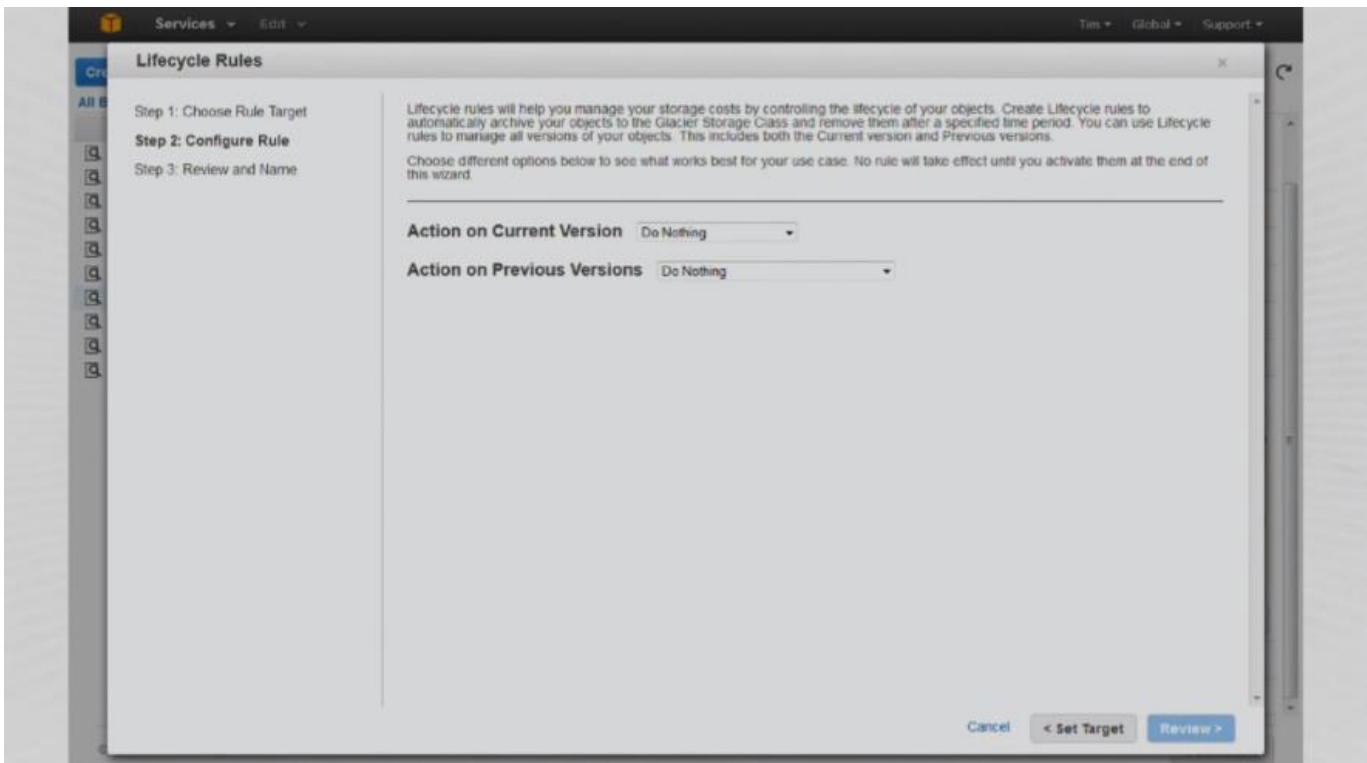
You can Suspend Versioning if you want later



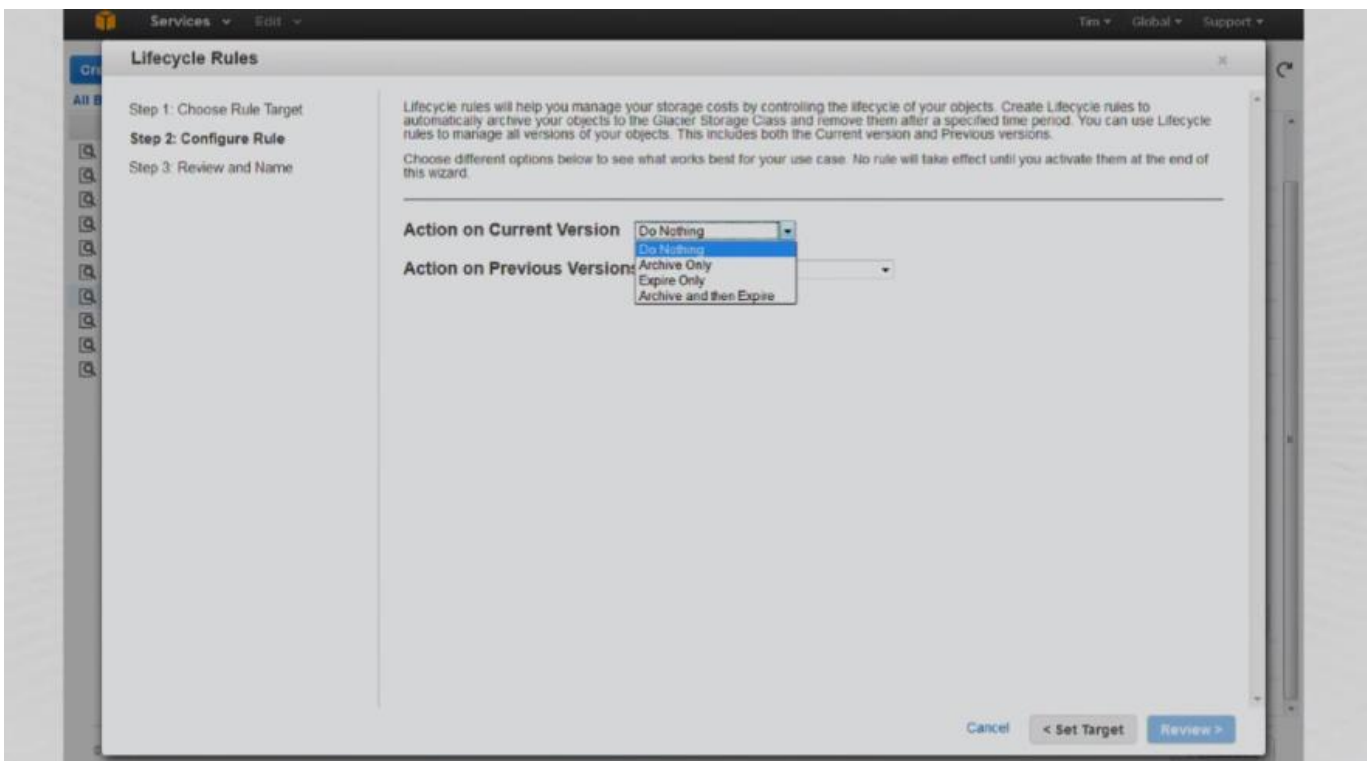
In the Lifecycle section for this bucket, we can turn on lifecycle rule for this bucket also by clicking the Add Rule link



This gives us a 3-step process to define a lifecycle policy for this S3 bucket. We decide if we want the rule to apply to the entire s3 bucket or just a prefix of it

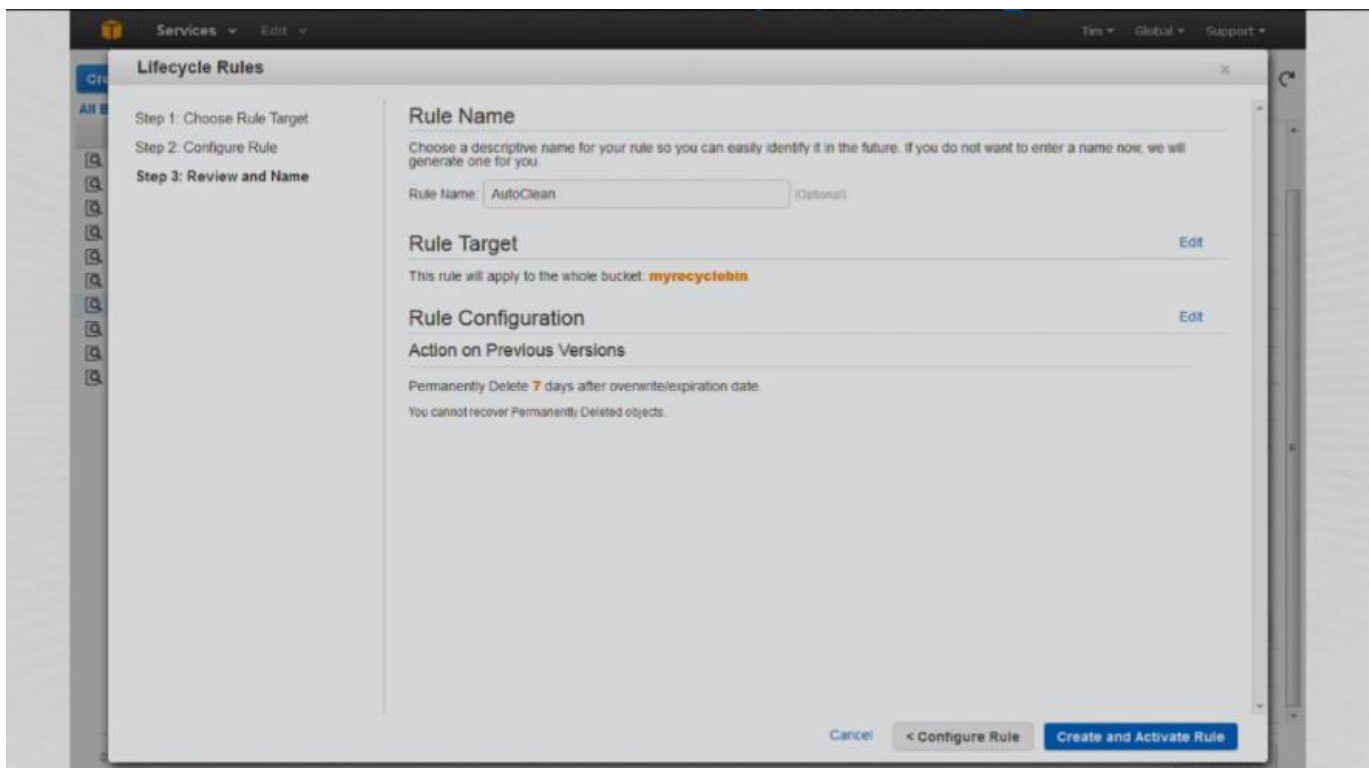
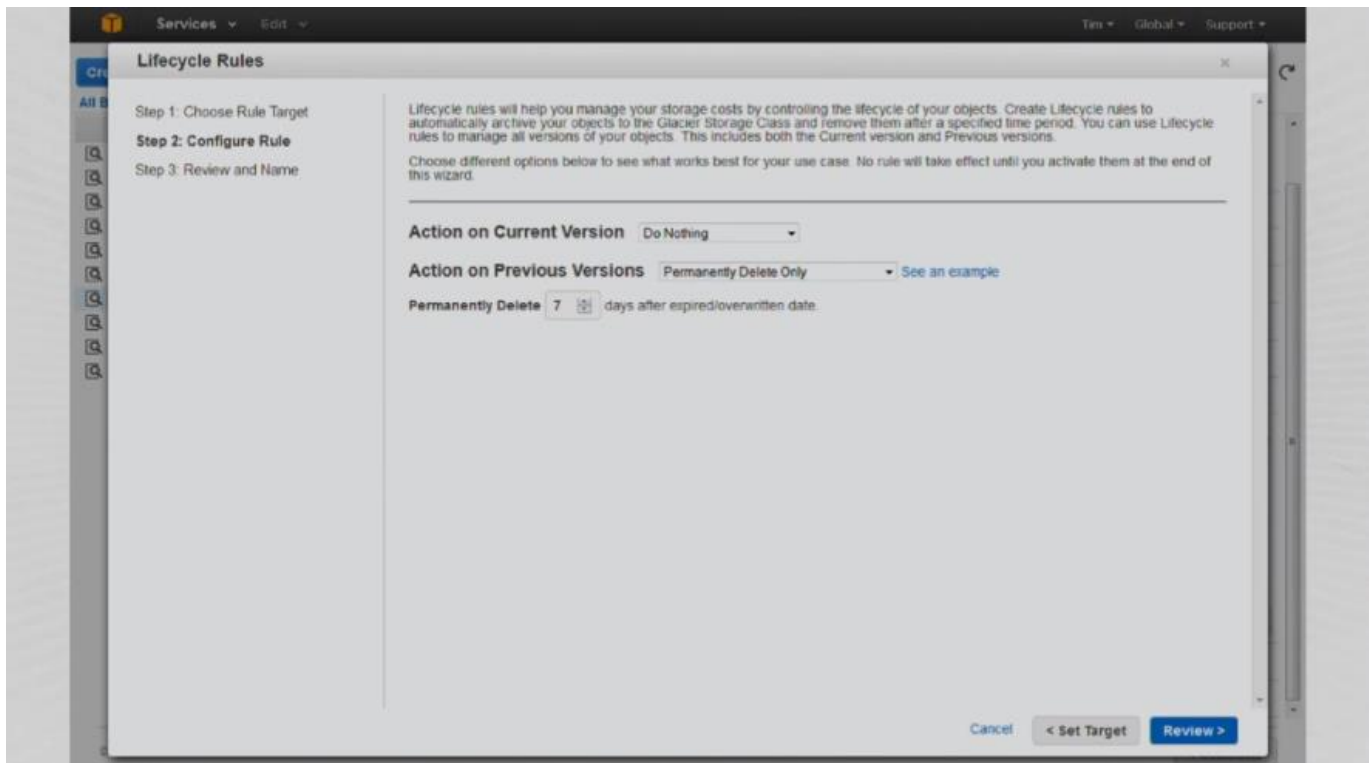


Then we chose what action we want to take both on the current version and on all the previous versions if they exist.

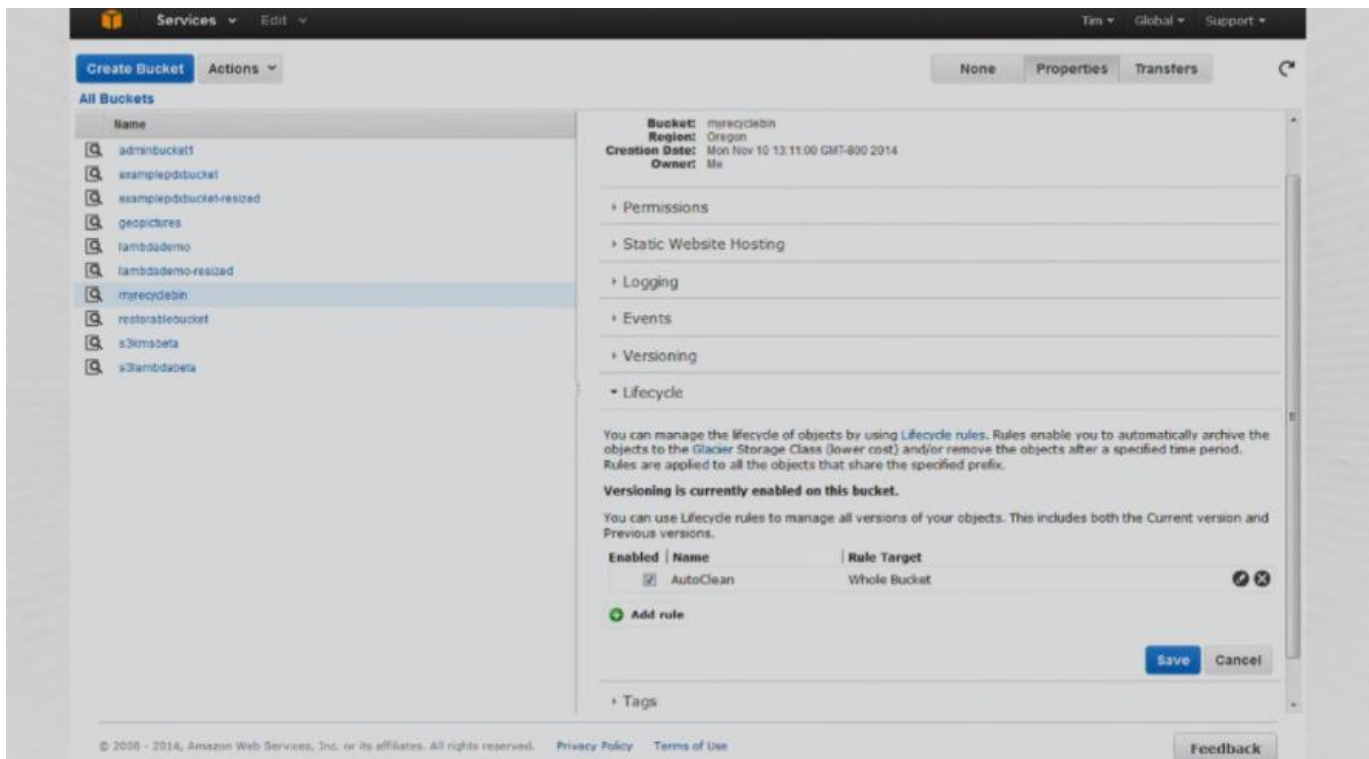


We have 4 choices listed above

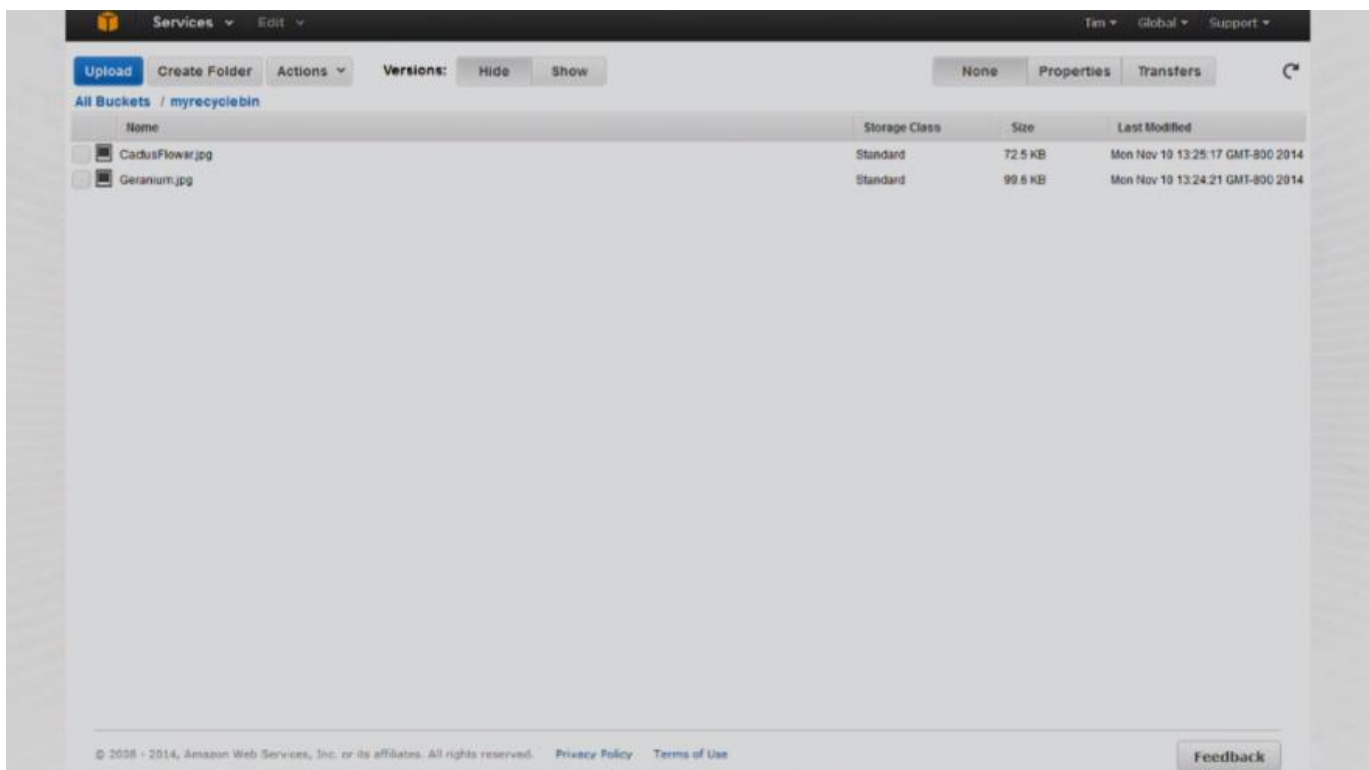




We then give the rule a name, review it, then save it



We now have a new lifecycle policy created for this S3 bucket.



We can see what versioning looks like in the console, we have a bucket with 2 objects



## Tip #1: 3 server-side encryption options

- SSE with Amazon S3 managed keys
  - “Check-the-box” to encrypt your data at rest
- SSE with customer provided keys
  - You manage your encryption keys and provide them for PUTs and GETS

New

- SSE with AWS Key Management Service managed keys
  - Keys managed centrally in AWS KMS with permissions and auditing of usage

## Tip #2: Detailed billing reports

- Provide object counts, storage GB, requests, and data transfer usage down to the bucket level
- Turn them on via the Preferences page in the Billing and Cost Management console
- Delivered to an Amazon S3 bucket you specify and downloadable from the Billing and Cost Management console



## Tip #3: Restricting deletes

- Bucket policies can restrict deletes
- For additional security, enable MFA (multi-factor authentication) delete, which requires additional authentication to:
  - Change the versioning state of your bucket
  - Permanently delete an object version
- MFA delete requires both your security credentials and a code from an approved authentication device



## Tip #4: Maximizing performance

- Use a key naming scheme with randomness at the beginning for high TPS
  - Most important if you regularly exceed 100 TPS on a bucket
  - Avoid starting with a date
  - Consider adding a hash or reversed timestamp (ssmmhhddmmyy)
- Multipart upload provides parallelism
  - Allows you to upload a single object as a set of parts
  - Enables pausing and resuming, and beginning before the total size is known
  - Encouraged for objects larger than 100MB; required above 5GB

# AWS re:Invent

Please give us your feedback on this session.  
Complete session evaluations and earn re:Invent swag.

SDD413

<http://bit.ly/awsevals>



Join the conversation on Twitter with **#reinvent**

© 2014 Amazon Web Services, Inc. and its affiliates. All rights reserved. May not be copied, modified, or distributed in whole or in part without the express consent of Amazon Web Services, Inc.