

CON 310

AWS re:INVENT

McDonald's – Moving to Containers

Thilina Gunasinghe and Manjeeva Silva

AWS
re:Invent



McDonald's is the world's largest restaurant company with 37,000 locations serving 64 million people per day. Using AWS, McDonalds built ***Home Delivery—a platform that integrates local restaurants with delivery partners such as UberEats.***

McDonald's built and launched the Home Delivery platform in less than four months using a microservices architecture running on Amazon Elastic Container Service, ***Amazon Elastic Container Registry, Application Load Balancer, Amazon ElastiCache, Amazon SQS, Amazon RDS, and Amazon S3.***

The ***cloud-native microservices architecture allows the platform to scale to 20,000 orders per second*** with less than 100-millisecond latency, and open APIs allow McDonald's to easily integrate with multiple global delivery partners. Using AWS also means the system provides McDonald's with a return on its investment, even for its average \$2–5 order value.

Contents

- About **McDonald's** and **Digital Acceleration**
- McDonald's **Home Delivery Platform**: Business Problem and Architecture
- Under the Covers: How we used ECS for **Scale, Speed, Security, DevOps and Monitoring**
- Final Thoughts and **Takeaways**

AWS
re:Invent



About McDonald's and Digital Acceleration

WORLD'S LARGEST RESTAURANT COMPANY



37K

RESTAURANTS



1.9M

PEOPLE
Working for
McDonald's
& Franchisees



120

COUNTRIES



64M+

CUSTOMERS
Served Every Day

VELOCITY ACCELERATORS

McDonald's uses scale as a competitive advantage to surpass the rising expectations of our customers across three growth initiatives.



digital




Seamless, personalized engagement with our customers when they are at home, on-the-go, or in a restaurants.

delivery



Bringing McDonald's directly to customers.

experience of the future



Elevating the McDonald's restaurant experience.

McDonald's Home Delivery Platform: Business Problem and Architecture



Introducing Home Delivery

digital & delivery



Bringing McDonald's to you



AWS
re:Invent



This journey starts when the customer picks a restaurant from a map on their phone to purchase their food from

Introducing Home Delivery

digital & delivery



Bringing McDonald's to you

AWS
re:Invent



aws

They then get shown the available menu at that particular restaurant

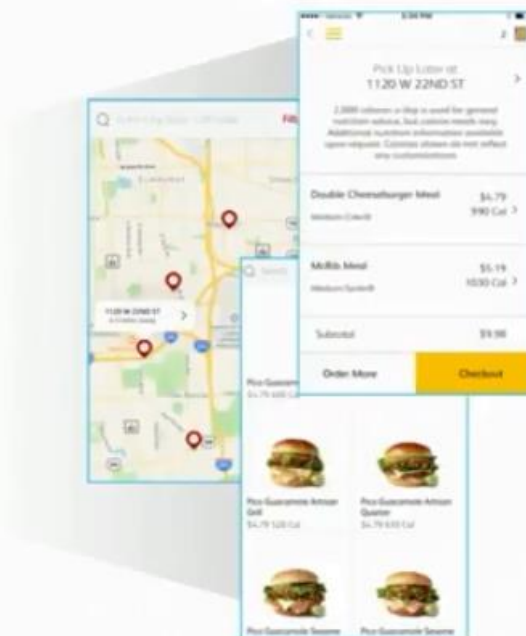
Introducing Home Delivery

digital & delivery



Bringing McDonald's to you

AWS
re:Invent



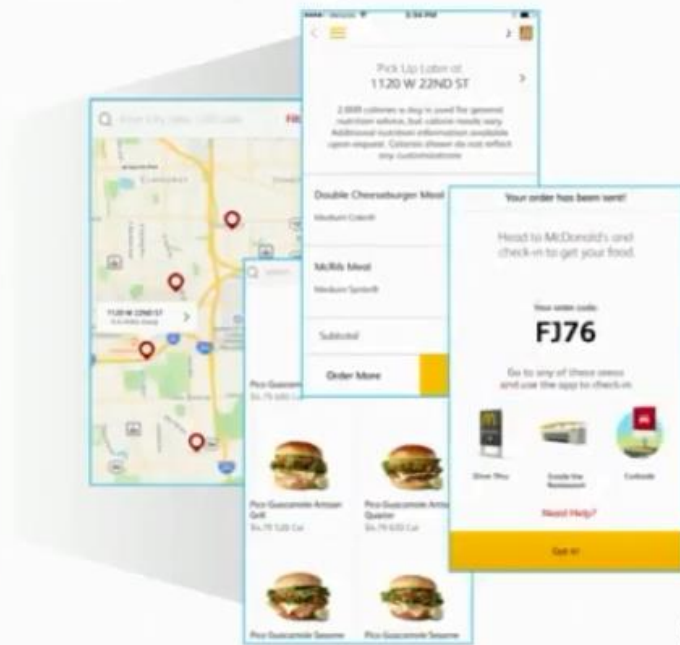
aws

The customer then completes their order in their basket

Introducing Home Delivery



AWS re:Invent



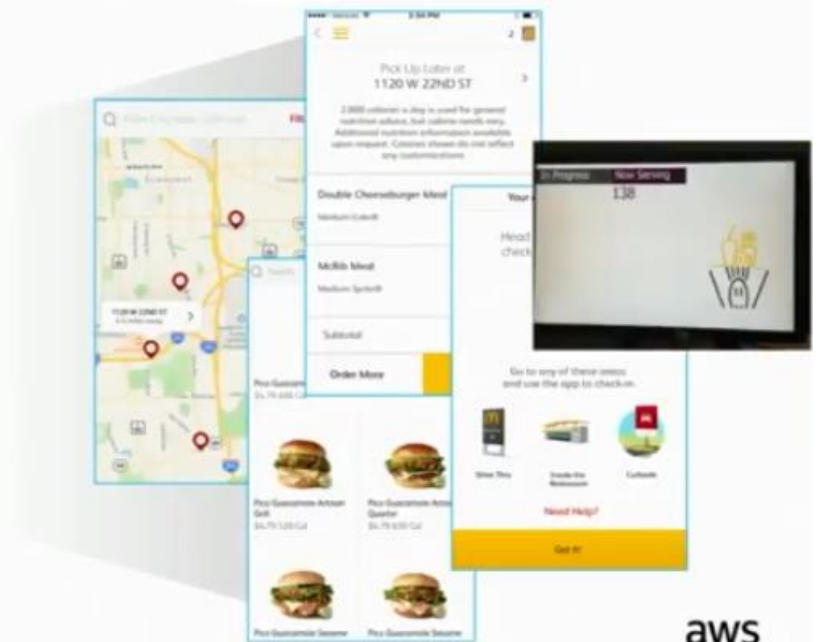
aws

At this point, order is complete and the delivery process begins

Introducing Home Delivery



AWS re:Invent



aws

The delivery driver or rider close to the restaurant gets a ping with the details for delivery




Critical Business Requirements

- **Speed to market**, quick turn around for features and functionality from concept to production
- **Scalability and reliability** targets of, 250K-500K orders per hour
- **Multi country support** and integration with multiple 3rd party food delivery partners
- **Cost sensitivity**, cost model based on low average check amounts

AWS
re:Invent



Home Delivery Architecture



AWS
re:Invent





Key Architecture Principles

- **Microservices**, with clean APIs, service models, isolation, independent data models and deployability.
- **Containers and orchestration**, for handling massive scale, reliability and speed to market requirements
- **PaaS**, based architecture model by leveraging AWS platform components such as ECS, SQS, RDS and ElastiCache
- **Synchronous and event based**, programming models based on requirements

Under the Covers: Using Amazon ECS for Scale, Speed, Security, DevOps and Monitoring

Why Amazon ECS?



Speed to market



Scalability and reliability



Security



DevOps – CI / CD



Monitoring

Speed to Market



- **4 months** from concept to production with 2 week dev iterations
- **Polyglot tech stack** ported to containers
 - Existing .net code was refactored and complied with .net core
 - Java was used where .net core is not supported
- **Simplified Amazon ECS deployment model** with easy integration to AWS services
- **Less time** was spend on **application tuning/testing** to achieve the scale and reliability requirements
- **DevOps** – Faster feedback loops on release iterations



Scalability and Reliability



- **Scale Targets:** Achieved the scale targets of 250k-500k orders per hour with below ~100ms response times
- **Autoscaling:** Used Amazon ECS “out-of-box” resource based autoscaling
- **Task Placement:** Task placement strategies and constrains were used to fine-tune and achieve container isolation with country / 3rd party scaling requirements
- **Scalability and Reliability Requirements (By Service):**
 1. {Service 1} Contains synchronous APIs with intense scalability and reliability requirements based on traffic burst patterns
 2. {Service 2} Requires more batch mode processing. Work load optimization is a critical requirement
 3. Some instances of {Service 3} will need isolation from others



Scalability and Reliability




```
{Service_1} Task Definition
"placementStrategy": [
  {
    "field": "attribute:ecs.availability-zone",
    "type": "spread"
  }
]

{Service_2} Task Definition
"placementStrategy": [
  {
    "field": "memory",
    "type": "binpack"
  }
]

{Service_3} Task Definition
"placementConstraints": [
  {
    "expression": "task:group == US",
    "type": "memberOf"
  }
]
```

re:Invent



Security



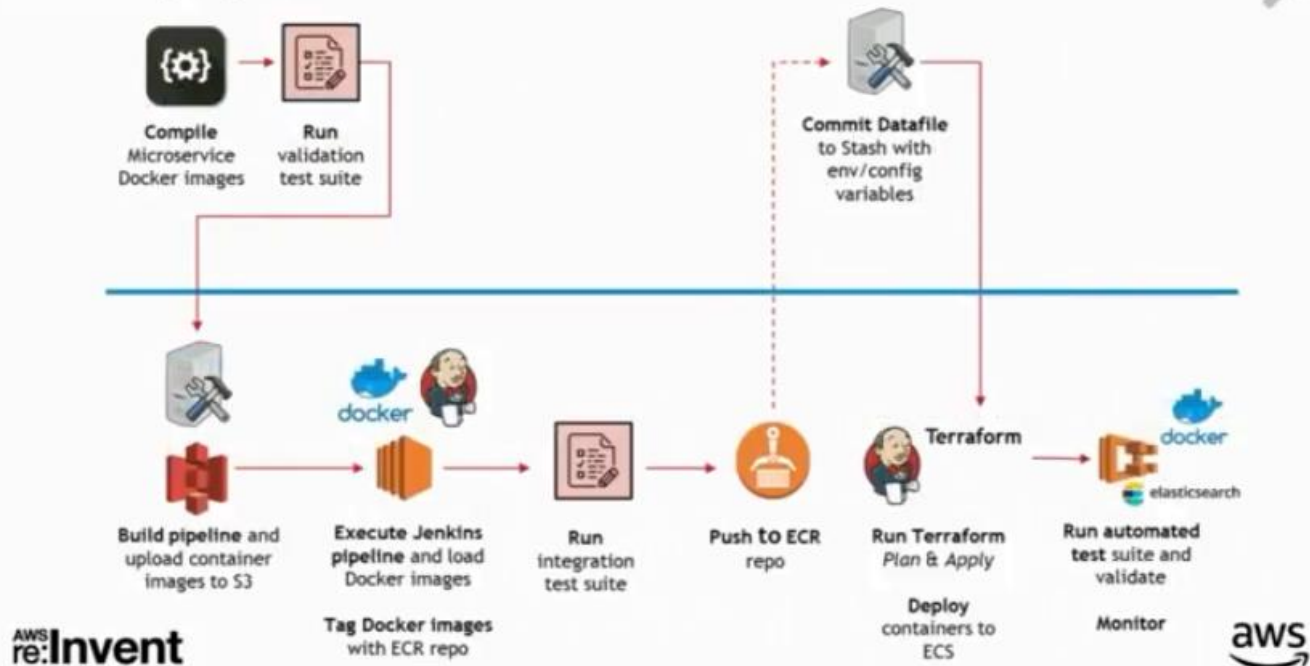
• Container security

- Container isolation through IAM policies and security groups
- Reduced container-to-container to communication (Reduce attack footprint).

• ECS instance security

- Automated AMI factory → start with ECS optimized AMIs → McD hardened gold AMIs → Project specific AMIs
- AMI factory pipeline listens to a SNS topic for obtaining updated AMIs

DevOps/CI&CD



Monitoring



- NewRelic agents configured to monitor the ECS instances, Containers and AWS PaaS components
- ELK stack configured for service logging and analytics



Major Technical Challenges

- “Out of memory error” due to containers not having access to *cgroup* file systems to get memory limits
 - **Solution:** Incorporated a new filesystem(*lxcfs*) to virtualized *cgroup* and virtualized views of */proc* files
- Docker containers are not honoring the ECS instance routing rules
 - **Solution:** Custom implementation of a Docker bridge

Final Thoughts

Final Thoughts and Key Takeaways

- A thought out **microservice architecture** is key for **scalability, reliability, and containerization**.
- **Massive scale** achievable (**north of 20k TPS under 100ms**) in a controlled manner **using auto-scale policies and task placement strategies**.
- **Moving to containers** simplified our development and deployment models and in turn provided **quicker dev/test iterations**.
- ECS out of the box integration and deployment models further **simplified our DevOps pipeline**.