

CMP218

AWS re:Invent

AWS Compute: What's New in Amazon EC2, Containers, and Serverless

Matt Garman, Vice President – Compute Services

November 29, 2017

AWS re:Invent



© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

Matt Garman, Vice President of AWS Compute Services, will introduce the latest innovations in the Compute space. At this session, we will be announcing new Compute capabilities, as well as insights into some of the underlying thinking of what makes the AWS Compute business unique. This session will cover new announcements around capabilities for EC2 instances, EC2 networking, EC2 Spot Instances, Amazon Lightsail, Containers and Serverless. Matt will also be joined by executives from our customers and partners, including GE CTO Chris Drumgoole, Heroku CEO Adam Gross, and Autodesk Chief of Product and Cloud Security Reeny Sondhi, who will share valuable success stories of how Amazon EC2 has helped their journey to digital transformation.

Thank you for your business!

Startup



Enterprise



Public Sector



SI/ISV



AWS re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.




AWS is the Market Leader for 7th consecutive year

AWS Positioned as a Leader in the Gartner Magic Quadrant for Cloud Infrastructure as a Service, Worldwide*

AWS is positioned
highest in execution
and furthest in
vision within the
Leaders Quadrant

Figure 1. Magic Quadrant for Cloud Infrastructure as a Service, Worldwide



*Gartner, Magic Quadrant for Cloud Infrastructure as a Service, Worldwide, Leong, Lydia, Sala, Raj, Lowery, Craig, Smith, Dennis, June 2017 G00315215


This graphic was published by Gartner, Inc. as part of a larger research document and should be evaluated in the context of the entire document. The Gartner document is available upon request from AWS: <http://www.gartner.com/doc/exprinfo?id=26202520&ctx=150518&ctx=26>

Gartner does not endorse any vendor, product or service depicted in its research publications, and does not advise technology users to select only those vendors with the highest ratings or other designation. Gartner research publications consist of the opinions of Gartner's research organization and should not be construed as statements of fact. Gartner disclaims all warranties, expressed or implied, with respect to this research, including any warranties of merchantability or fitness for a particular purpose.

AWS Positioned as a Leader in the IDC MarketScape: Worldwide Infrastructure as a Service 2017 Vendor Assessment*

"AWS has effectively
defined and led the core
offering portfolio in the
public cloud IaaS
market."

IDC MarketScape: Worldwide Infrastructure as a Service 2017



*SOURCE: "IDC MarketScape: Worldwide Infrastructure as a Service 2017 Vendor Assessment", by Deepak Mohan, Erik Berggren and Laura Dullos, September 2017 IDC # US43079816

IDC MarketScape vendor analysis model is designed to provide an overview of the competitive fitness of ICT suppliers in a given market. The research methodology utilizes a rigorous scoring methodology based on both qualitative and quantitative criteria that results in a single graphical illustration of each vendor's position within a given market. The Capabilities score measures vendor product, go-to-market and business execution in the short-term. The Strategy score measures alignment of vendor strategies with customer requirements in a 3-5 year timeframe. Vendor market share is represented by the size of the circles. Vendor year-over-year growth rate relative to the given market is indicated by a plus, neutral or minus next to the vendor name.

*"Amazon's inaugural
cloud conference draws
6,000 developers, IT pros"*

2012

The AWS Compute business,
5 years ago at 1st re:Invent

We had a single service focused on compute:



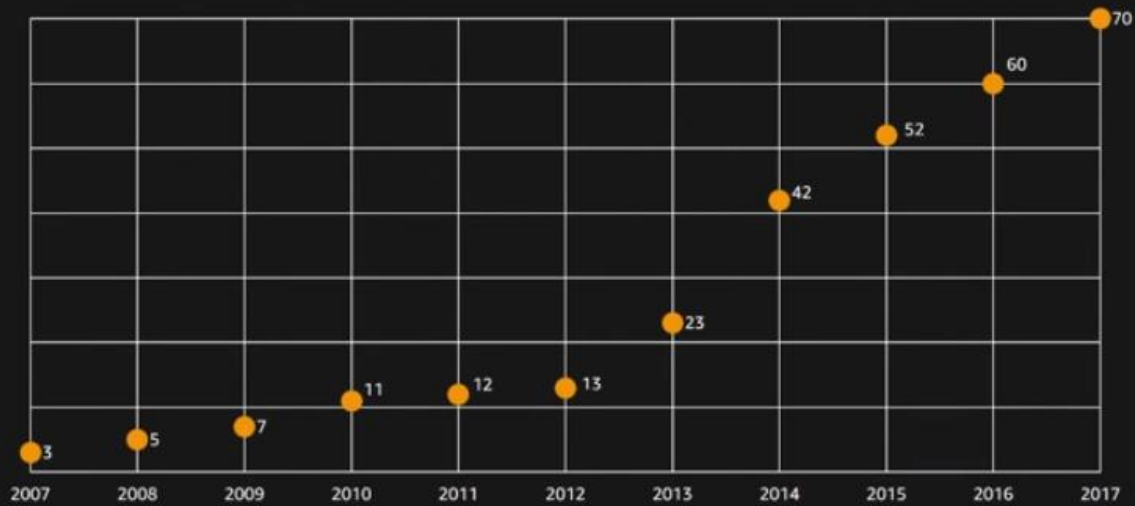
Amazon EC2

AWS
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



EC2 Instance Type Growth—Increased Choice



AWS
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



Compute Services are at the heart
of this amazing pace of innovative growth



Instances



Containers



Serverless

AWS
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.




The AWS Compute Business—Now



AWS
re:Invent



Compute services are at the heart of social good



AWS
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

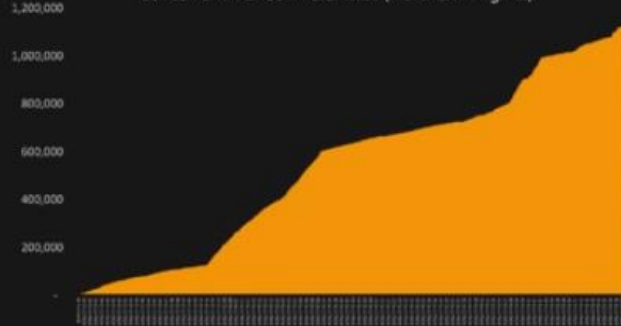


Cognoa is a startup that has a mobile app that uses AI to identify behaviors in children that are correlated with developmental issues in children

Compute services are at the heart of immense scale

CLEMSON
UNIVERSITY

Concurrent vCPUs in U.S. East (Northern Virginia)



Natural Language Processing on **1.1 million concurrent vCPUs**
in less than **24 hours** using EC2 Spot instances

AWS
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



Continuous Innovation Cloud

Continuous Innovation Cloud



AWS
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



Continuous Innovation Cloud



AWS
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



Strengthen Your Security Posture



Over 50 global
compliance
certifications &
accreditations



Benefit from AWS
industry leading
security teams 24/7,
365 days a year



Security infrastructure
built to satisfy military,
global banks, and other
high-sensitivity organizations



World class network
performance
and capabilities



"We work closely with AWS to develop a security model, which we believe enables us to operate more securely in the public cloud than we can in our own data centers."

Rob Alexander, CIO

AWS
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.





Access a Deep Set of Cloud Security Tools

Networking

Virtual private cloud
Isolated cloud resources
VPC Flow Logs

Web application firewall
Filter malicious
Web Traffic

Shield
DDoS protection

Certificate manager
Provision, manage,
and deploy SSL/TLS
certificates

Encryption

Key management service
Manage creation
and control of encryption
keys
Support for KMS

CloudHSM
Hardware-based
key storage

Server-side encryption
Flexible data
encryption options
Encrypted Boot & EBS
volumes

Identity & management

IAM
Manage user access
and encryption keys

SAML federation
SAML 2.0 support to
allow on-prem identity
integration

Directory service
Host and manage
Microsoft Active Directory

Organizations
Manage settings for
multiple accounts

Compliance

New compliances added,
especially government

Service catalog
Create and use
standardized products

Config
Track resource inventory
and changes

CloudTrail
Track user activity
and API usage

CloudWatch
Monitor resources
and applications

Inspector
Analyze application
security

Artifact
Self-service for AWS
compliance reports

AWS
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



Access a Deep Set of Cloud Security Tools

25% more
certifications
than other cloud providers



AWS
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.





10 years

of region & availability zone growth

AWS
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



Architected at the Core for High Availability



Legend:

- Regions
- Availability Zones
- New Regions
- New Availability Zones



NEW!

Spread Placement:

Reduce likelihood of correlated failure to customer workloads

AWS
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



Architected at the Core for High Availability



Legend:

- Regions
- Availability Zones
- New Regions
- New Availability Zones

SLA of **99.99%** availability



Spread Placement:

Reduce likelihood of correlated failure to customer workloads

AWS re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



Availability Expert Advice



Quick starts



Technical Account Managers



Whitepapers & best practices



Solution Architects



Webinars



SI Partner resources

Accenture, Capgemini, CloudMas, Cloudnexus, Cognizant, Infosys, Logicworks, Mobiquity, Nordcloud, NRI, Slatom, Wipro, and many more!

AWS re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



A world map with several location pins. Green pins are located in North America (USA and Canada), Europe, and Asia. Blue pins are located in South America, Africa, and Australia. The GE Digital logo is in the top right corner.

Cloud@GE

- Leverage global, multi-region model with both persistent and edge only connectivity.
- Multi-AZ application patterns as the norm.
- Deliver low RTO, low latency geo-distributed applications.

A horizontal collage of five images: a medical scanner, a large industrial turbine, a factory interior, a close-up of a turbine component, and a wind turbine.

GE applications has 9 layers of resiliency and are spread across AZs and Regions to improve resiliency.

Continuous Innovation Cloud



AWS
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



Flexible Access to Compute Power— Per-second Billing



On-demand

Pay for compute capacity by the second
with no long-term commitments



Per-second billing for EC2 Linux
instances & EBS volumes

AWS
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



EC2 Reserved Pricing— Flexibility & Savings



Flexible capacity
reservation
when used in a specific
Availability Zone



Convertible RIs
Change instance family,
OS, tenancy, & payment



Payment flexibility
3 upfront payment options
(all, partial, none)

NEW! 1 year convertible RIs

AWS
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



NEW

Flexible Access to Compute Power— EC2 Spot, Streamlined



Predictable prices, pause, & resume



NEW! Spot in China (11/20)

AWS
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



You can now use spot instances with the **run instance** command to launch a spot instance instantly, you can also hibernate your spot instances and resume them later



They use spot instances to parallelize their video re-rendering jobs

Joshua Burgin
GM, EC2 Spot

For more information

CMP202: Optimizing EC2 for Fun and Profit #bigsavings #newfeatures

- Wednesday, Nov 29, 3:15-4:15PM
Venetian, Level 2, Venetian F
- Thursday, Nov 30, 11:30-12:30PM
Venetian, Level 2, Titian 2204



AWS
re:Invent



Customers Using Containers on AWS



AWS re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



Flexible Access to Compute Power— Containers



Key advancements in 2017

Native service discovery

Task networking

Expanded region support



AWS
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



Flexible Access to Compute Power— Containers



Problem

Need to manage cluster of
instances used to launch containers

Flexible Access to Compute Power— Containers



Problem

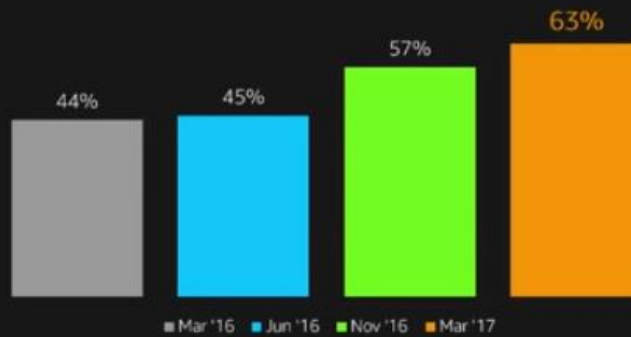


Need to manage cluster of
instances used to launch containers

NEW

Allows you to simply just run a container task without worrying about cluster management

Most Kubernetes Clusters Run on AWS



Cloud Native Computing Foundation Survey

NEW

Most Kubernetes Clusters Run on AWS



Cloud Native Computing Foundation Survey

Deepak Singh
GM, EC2 Container Services

For more information

CON201: Containers on AWS –
State of the Union

Thursday, Nov. 30, 12:15-1:15PM
Aria, Level 1, Pinyon 5



AWS
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



Customers Using AWS Lambda



AWS
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



Flexible Access to Compute Power— AWS Lambda



Key advancements in 2017

More accessible to more
developers

Make Lambda developers
more productive

Enable new application
patterns

24

Feature launches in 2017

2

New languages – Python 3, Node 6

All

Supported across all regions



ISO, PCI, HIPAA, and SOC compliances



AWS
Green
Grass



Lambda
@Edge



X-Ray

AWS
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



Tim Wagner GM, AWS Serverless Applications

For more information

SRV305: What's New in Serverless

Thursday, Nov. 30, 3:30-4:30PM
co-present w/ FICO (Jeet Kaul, VP)
Aria, Level 1, Pinyon 5

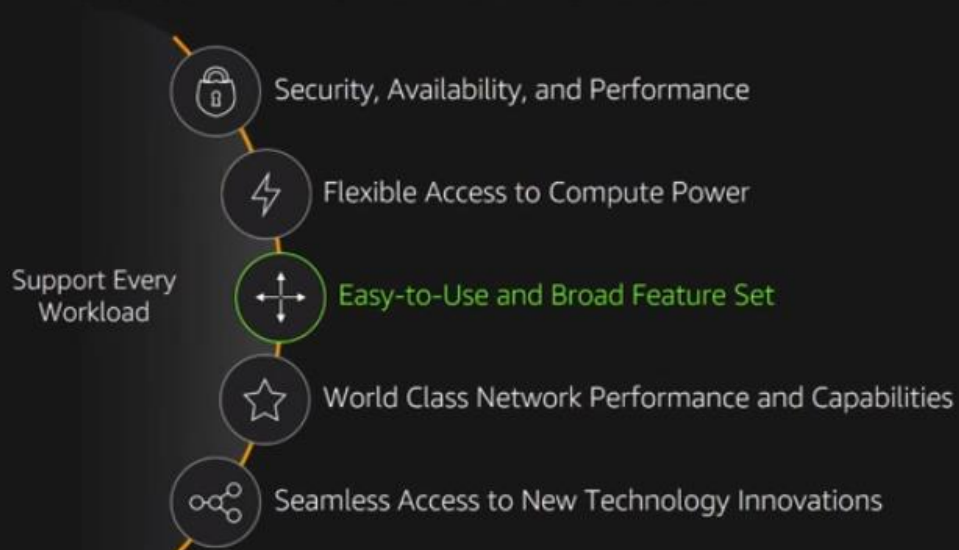


AWS
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



Continuous Innovation Cloud



AWS
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



Easy-to-Use with
Broad Feature Set—
Media & Entertainment



AWS
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



Easy-to-Use with Broad Feature Set—
HPC Innovations & Technology



Batch



HPC
Application
Performance



EnginFrame



AWS
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



Easy-to-Use with Broad Feature Set— Developers



Amazon Lightsail — easy to launch, easy to grow

Bundled compute,
storage, and networking

Low, predictable
prices

Intuitive
console

Access to AWS
services

Rapid innovation in 2017

Use worldwide, support Windows workloads, and scale applications easily



10 regions



8 languages



Windows Server



Block storage add-on

AWS
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



NEW

Introducing: Lightsail Load Balancing



Lightsail Load Balancer

Easy-to-use load balancing with free certificate management allows customers to create secure, high-availability web applications for a low, predictable price of \$18 per month

AWS
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



NEW


Expanding Features for Developers



Amazon Time Sync Service

Highly reliable service with redundant array of satellite and atomic clock sources

Available globally today!



AWS
re:Invent

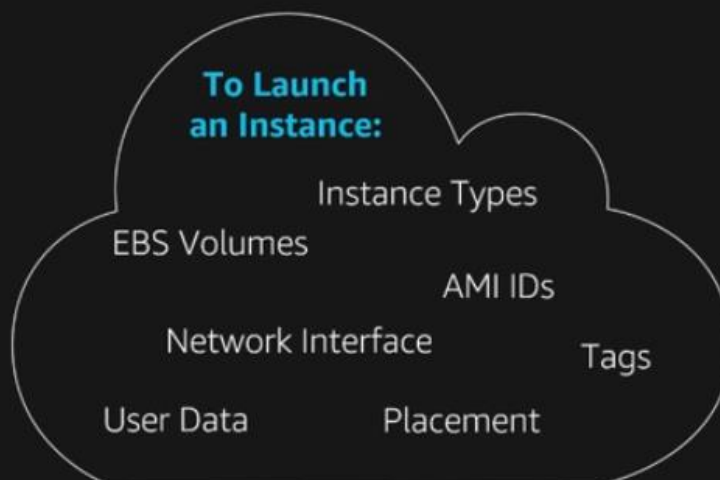
© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



Expanding Features for Developers



To Launch
an Instance:



AWS
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



NEW

Expanding Features for Developers

Launch Templates

To Launch an Instance **Now:**

Launch

Instance Types
EBS Volumes
Network Interface
User Data
Placement
Tags
AMI IDs

AWS re:Invent
© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

aws

You can now define all your launch time parameters in a launch time template that is versioned.

NEW

Expanding Features for Developers

Launch Templates

For more information

CMP324: Re:inventing EC2 instance launches with Launch Templates

Thursday November 30, 4:00–5:00pm
Venetian – Level 1 – Casanova 506

User Data
Placement

AWS re:Invent
© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

aws

Continuous Innovation Cloud




AWS
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



Continuous Innovation— On-instance Networking

Key highlights



AWS
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.




NEW

AWS PrivateLink



Share services privately
between VPCs and
on-premises networks

Secure. Scalable. Reliable.



Customers
and
Partners



Vanguard



TIBCO



APPDYNAMICS



Expedia



aqua



AWS
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



AUTODESK®

Reeny Sondhi
Chief of Product Security

Dave Brown Director, EC2 Networking

For more information

NET205: Networking State of the Union

Wednesday, Nov. 29, 3:15-4:15PM

Venetian, Level 3, Murano 3203



AWS
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



Continuous Innovation Cloud



AWS
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



Customer Obsessed



90%


of the items on the roadmap originate with customer requests and are designed to meet specific needs and requirements

AWS
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



Continuous Innovation— Instance Types



AWS
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



FPGAs—F1 Instance



edico  genome



 Children's Hospital
of Philadelphia

Guinness World Record for fastest
analysis of 1,000 Genomes, running
on Amazon EC2 FPGA instances

AWS
re:Invent



© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



High-performance Graphics with G3



HALLIBURTON



Seismic exploration and analytics for Oil & Gas

"The exploration and production models are increasingly complex with very large datasets, 3D and dynamic algorithms, security, and global reach... Amazon EC2 G3 instances enable Landmark to deliver value to our clients in ways that were not possible before."

—Chandra Yeleshwarapu
Global Head of Services and Cloud Landmark, Halliburton

AWS
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



NEW

New Technology Innovation—P3



SO MUCH DATA FOR MACHINE LEARNING & AI TO CONSUME!

By 2020,
there will be
44
zettabytes
of data¹



1.2 trillion
photos taken
in 2017²



5.2TB
per person
of data³



Over **100k** chatbots
today the chatbot
market is growing
at a compounded
annual growth
rate (CAGR)
of **37%**⁴



THE ANSWER: AMAZON EC2 P3 + MACHINE LEARNING & AI

Amazon EC2 P3s
are available across




Oregon, Virginia,
Dublin, and Tokyo

Up to **14X** faster
for Machine Learning
training than
previous generation



One Amazon EC2
P3 instance provides
1 petaflop of
compute performance



41k images/second
for image
recognition/
inference per EC2 P3



Tesla V100 GPUs
contains **5,120**
CUDA cores and
640 Tensor cores



AWS
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



NEW

Continuous Innovation—T2 Unlimited



Cost-effective
for most general-
purpose workloads



Generous baseline
performance,
automatically burst
to full-core



Now burst whenever
you want, for as long
as you want

AWS
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



NEW

Continuous Innovation—H1



H1

New Storage-optimized instance

Up to 16TB of locally
attached HDD storage

Up to 25 Gbps network
bandwidth with ENA



AWS
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

aws

NEW

Continuous Innovation— Purpose-built Virtualization Platform



C5

Next-generation, Compute
Optimized EC2 instance

M5

Next-generation, General
Purpose EC2 instance

25-50% price performance
improvement

Introducing our custom built all
new hypervisor – Nitro

AWS
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

aws

NEW

Continuous Innovation—EC2 Bare Metal



EC2 Bare Metal

Run bare metal workloads on EC2
with all the elasticity, security,
scale, and services of AWS



AWS
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

For public preview



Continuous Innovation Cloud

Support Every
Workload



Security, Availability, and Performance



Flexible Access to Compute Power



Easy-to-Use and Broad Feature Set



World Class Network Performance and Capabilities



Seamless Access to New Technology Innovations

AWS
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.




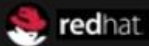


HEROKU

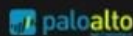
Adam Gross
CEO
SVP, Salesforce

Support Every Workload—Partners

Enterprise



Security




Networking



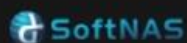
Databases



BI



Storage



Media



AWS
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



Support Every Workload—Every Industry

Gaming



IoT



Healthcare & Life Sciences



Financial Services



Media & Entertainment



Education



Retail



Nonprofits



Government



AWS
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

And many more!



Sandy Carter
VP, EC2

For more information

WIN201: EC2 Enterprise Workloads
State of the Union
Thursday, Nov. 30, 11:30-12:30PM
MGM, Level 1, Grand Ballroom 117



AWS
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



What's Next?

Don't expect us to slow down anytime soon!

Recap of What We're Announcing Today

Spot - Simplified Pricing Model
Spot on RunInstances
Instance Hibernation for Spot
Amazon Fargate
Amazon EKS
Spread Placement

Lightsail Load Balancer
Array Jobs
Amazon Time Sync Service
Launch Templates
PrivateLink for
Customers & Partners

T2 Unlimited
H1
C5
M5
EC2 Bare Metal Instances
P3
G3

Want to Learn More?



Compute Services
announcement

aws.amazon.com/new/reinvent



Sign up for a new
Lightsail account

amazonlightsail.com



Kickoff Spot
instances

aws.amazon.com/ec2/spot

FREE for the first month!

AWS
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



AWS re:Invent

THANK YOU!

AWS
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

