Serverless technologies let you build and scale applications and services rapidly without the need to provision or manage servers. In this session, we show you how to incorporate serverless concepts into your big data architectures. We explore the concepts behind and benefits of serverless architectures for big data, looking at design patterns to ingest, store, process, and visualize your data. Along the way, we explain when and how you can use serverless technologies to streamline data processing, minimize infrastructure management, and improve agility and robustness and share a reference architecture using a combination of cloud and open source technologies to solve your big data problems. Topics include: use cases and best practices for serverless big data applications; leveraging AWS technologies such as Amazon DynamoDB, Amazon S3, Amazon Kinesis, AWS Lambda, Amazon Athena, and Amazon EMR; and serverless ETL, event processing, ad hoc analysis, and real-time analytics.

Serverless Analytics Evolution...

Virtualized → Managed → Serverless

Provision servers

Configure Clusters

Run Analytics

Serverless characteristics

No servers to provision or manage

Scales with usage

Never pay for idle

Availability and fault tolerance built in

Serverless nicely fits into big data platforms

- Mix and Match Serverless, Managed, and Virtualized
- Leverage Services to easily
  - Rapidly ingest, categorize, and discover your data
  - Allow easy query and analysis of your data
  - Transform and Load data
  - Provide custom event based handlers
- Serverless allows you to focuses more analytics and not on infrastructure or servers

Serverless Compute

AWS Lambda

- **Run your code in the cloud - fully managed and highly-available**
- **Triggered through API or state changes in your setup**
- **Scales automatically to match the incoming event rate**
- **Node.js (JavaScript), Python, Java, and C#**
- **Charged per 100ms execution time**

Serverless Interactive Query Service

Amazon Athena

- Query directly from Amazon S3
- Use ANSI SQL
- Serverless
- Multiple Data Formats
- Pay per query

Serverless Catalog and ETL/ELT Service

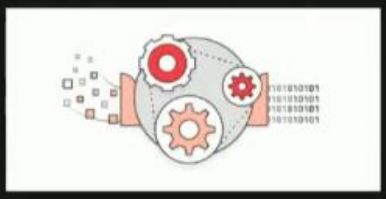

AWS Glue

Data Catalog

Crawl, Discover and Organize Data Integration with Managed and Serverless Analytics

Job Authoring

Serverless ETL – Pay for what you consume

Job Execution

Serverless Streaming Made Easy

Services make it easy to capture, deliver and process streams on AWS

Amazon Kinesis Streams

- For Technical Developers
- Build your own custom applications that process or analyze streaming data

Amazon Kinesis Firehose

- For all developers, data scientists
- Easily load massive volumes of streaming data into S3, Amazon Redshift and Amazon Elasticsearch

Amazon Kinesis Analytics

- For all developers, data scientists
- Easily analyze data streams using standard SQL queries

Applying Serverless to Big Data Applications?

Characteristics of a Big Data Applications

Collect Anything

Dive in Anywhere

Flexible Access

Future Proof

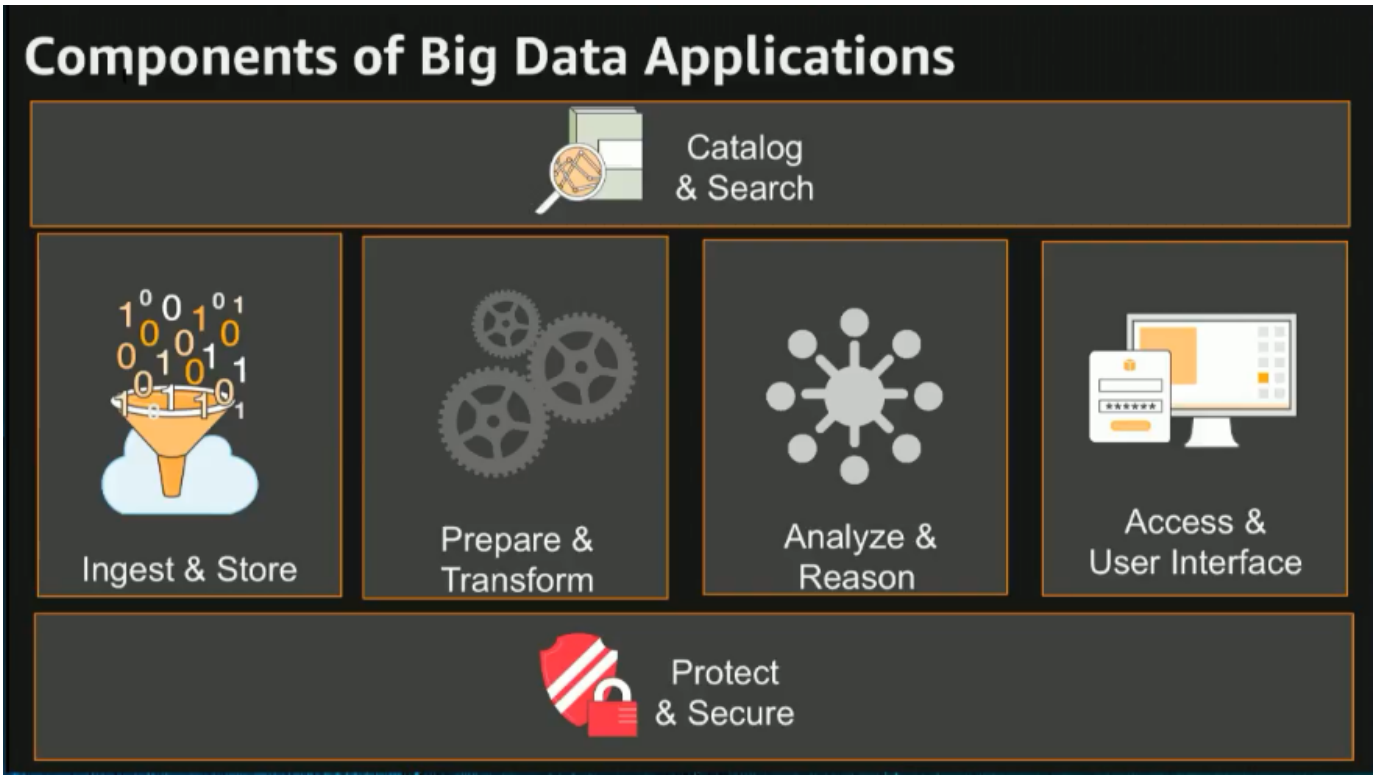

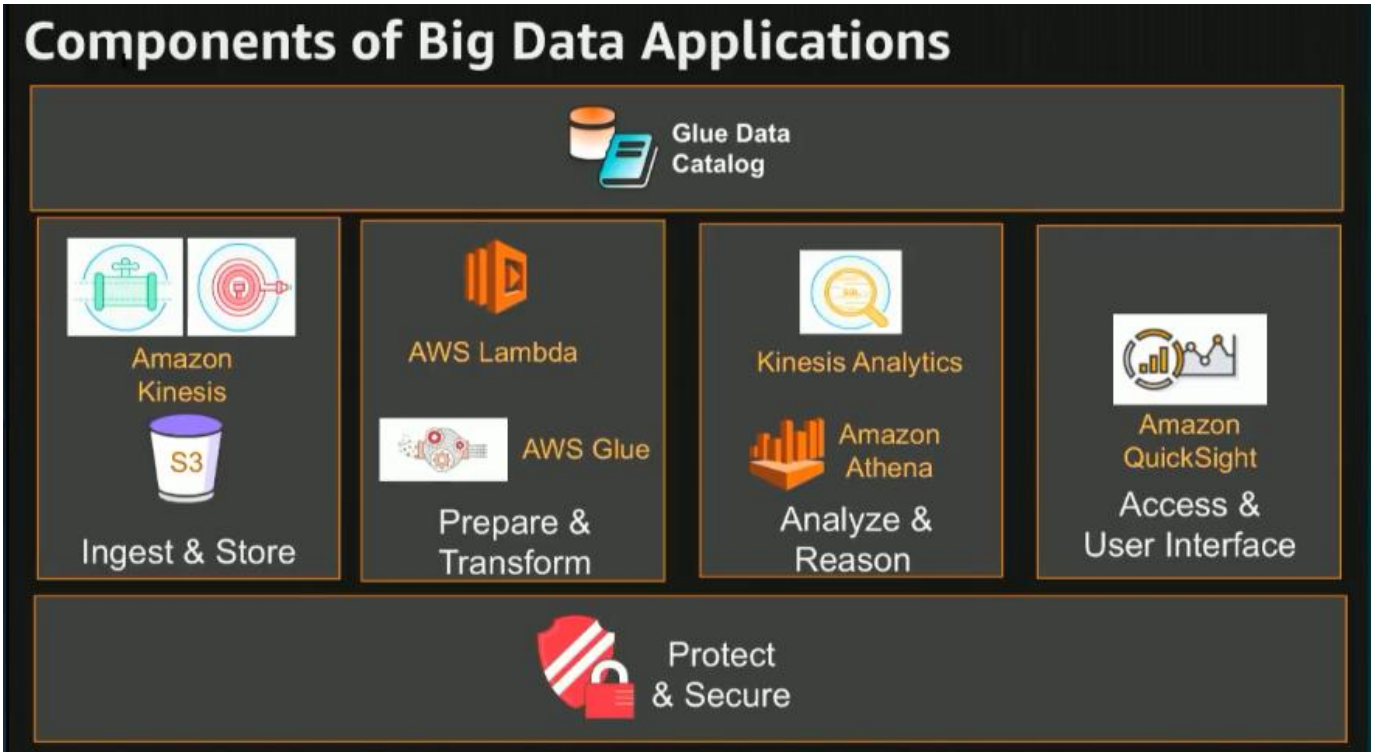

AWS Glue has the capability to be a metadata catalog that helps you bring in and maintain new dataset. ***Lambda*** allows you to write transformation logic for your data as they come in, ***Glue*** also allows you to write transformations of your data also. ***Kinesis Analytics*** provides SQL based access on top of data streams while Athena provides SQL based access on top of your data stored on ***S3***.

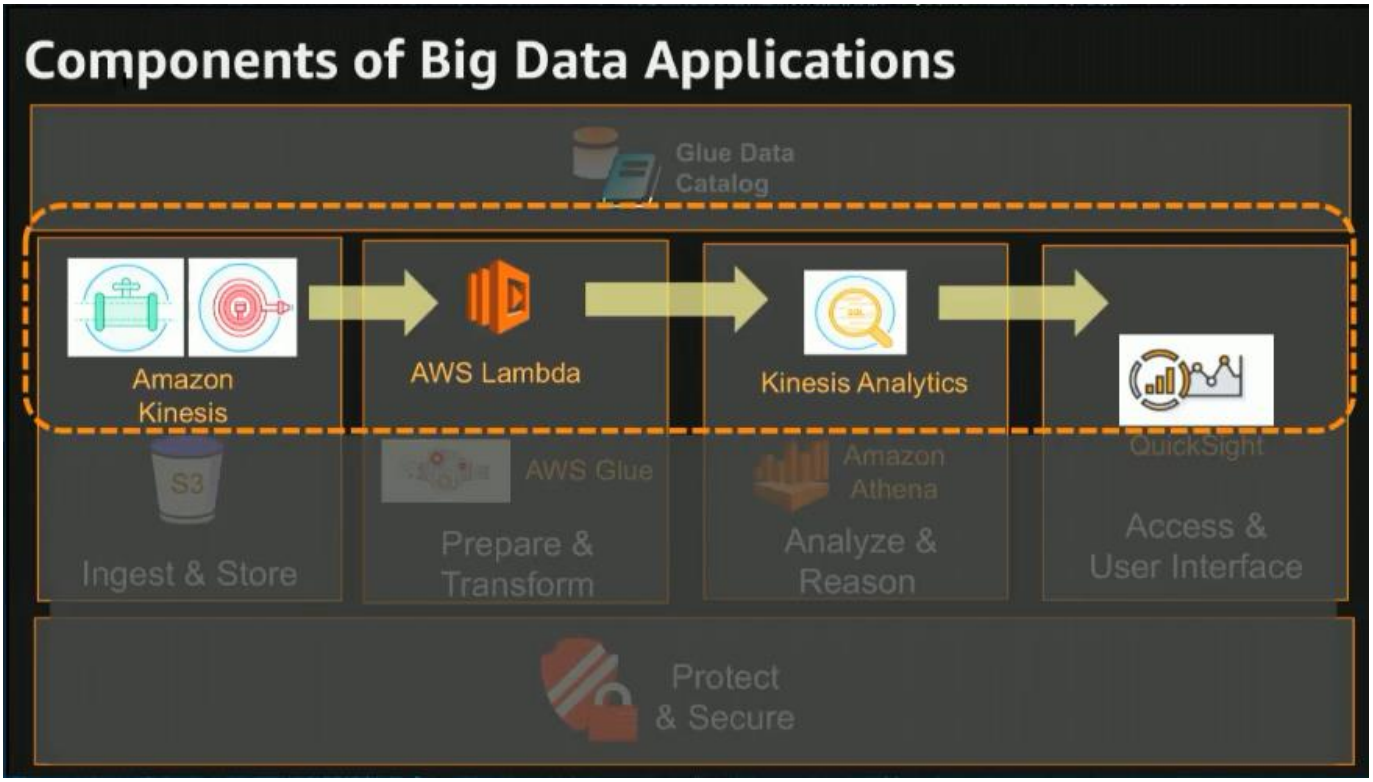

This is a very common real time analytical flow, it is very easy to setup to start capturing, transforming, and analyzing data

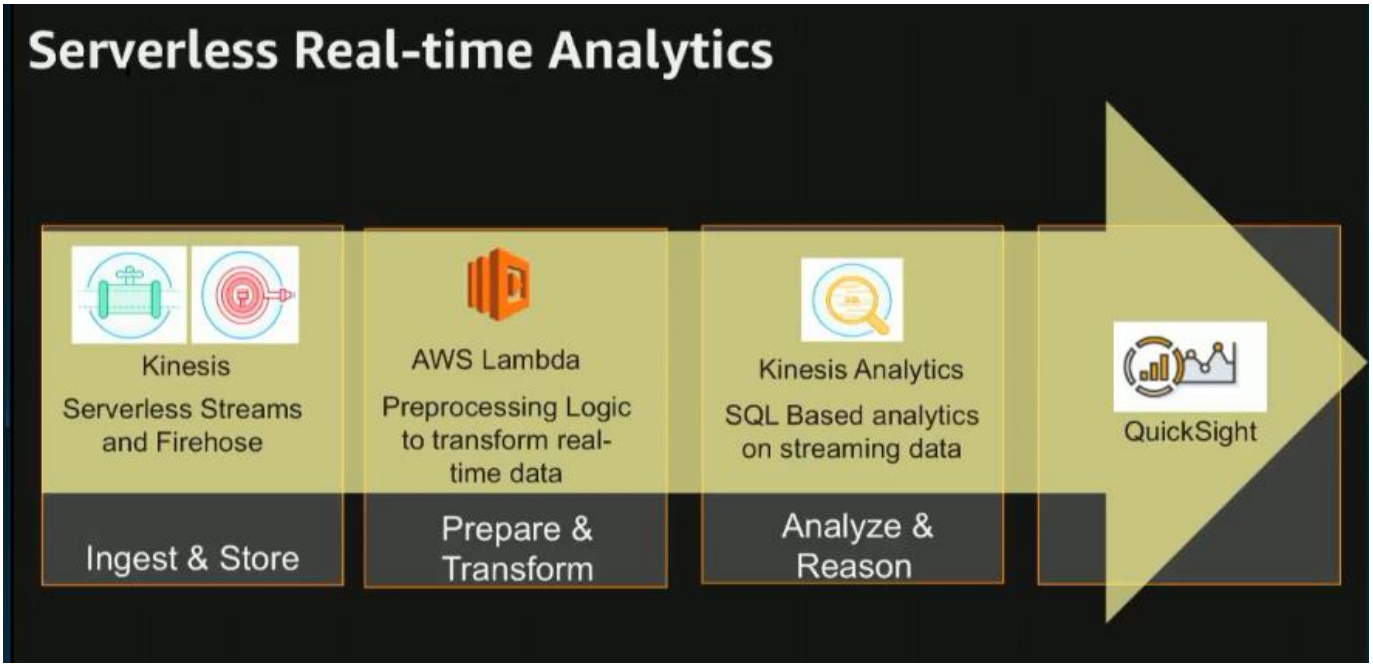

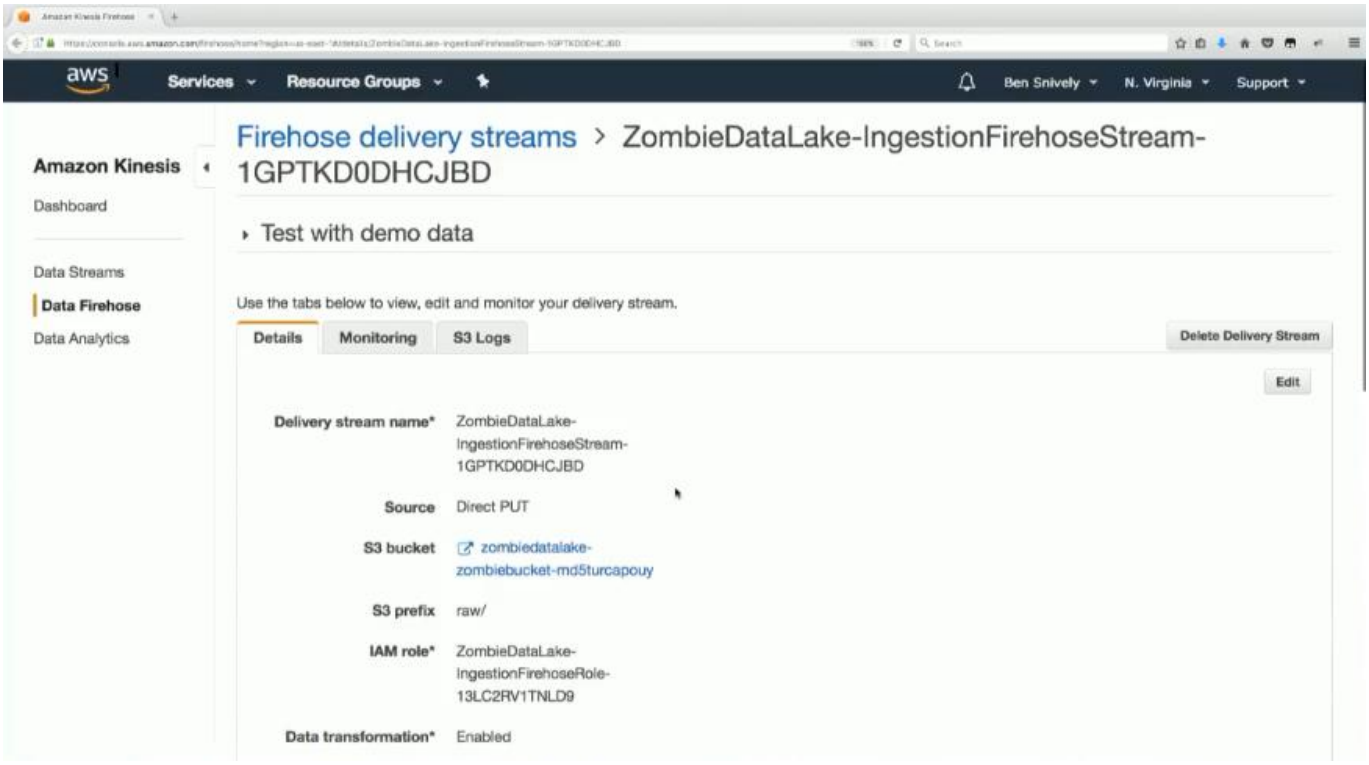

This is a Kinesis Firehose, we have specified an S3 bucket with a prefix ***raw/***. We want to capture app messages and analyzed the data

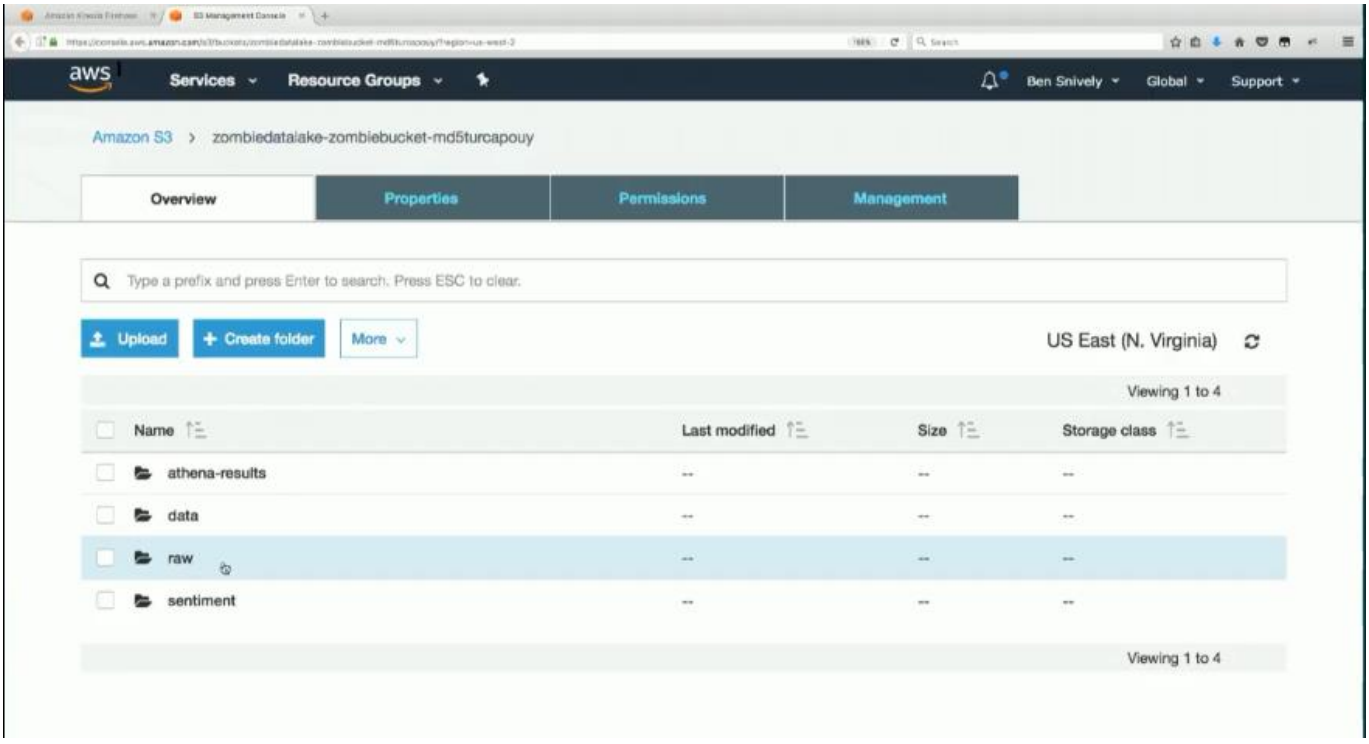

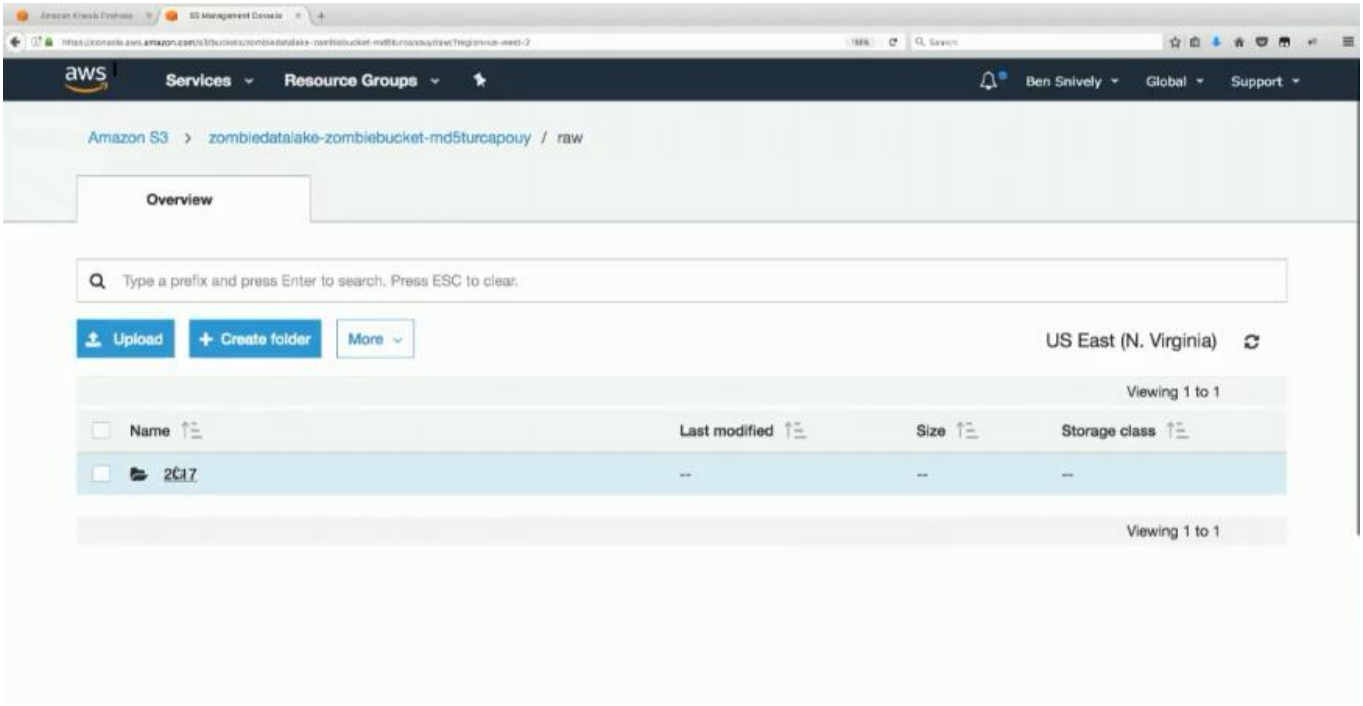
aws
Services
Resource Groups
Ben Snively
Global
Support
Amazon S3 > zombiedatalake-zombiebucket-md5turcapouy / raw
Overview
Type a prefix and press Enter to search. Press ESC to clear.
Upload
+ Create folder
More
US East (N. Virginia)
Viewing 1 to 1
Name
Last modified
Size
Storage class
2017
--
--
--
Viewing 1 to 1

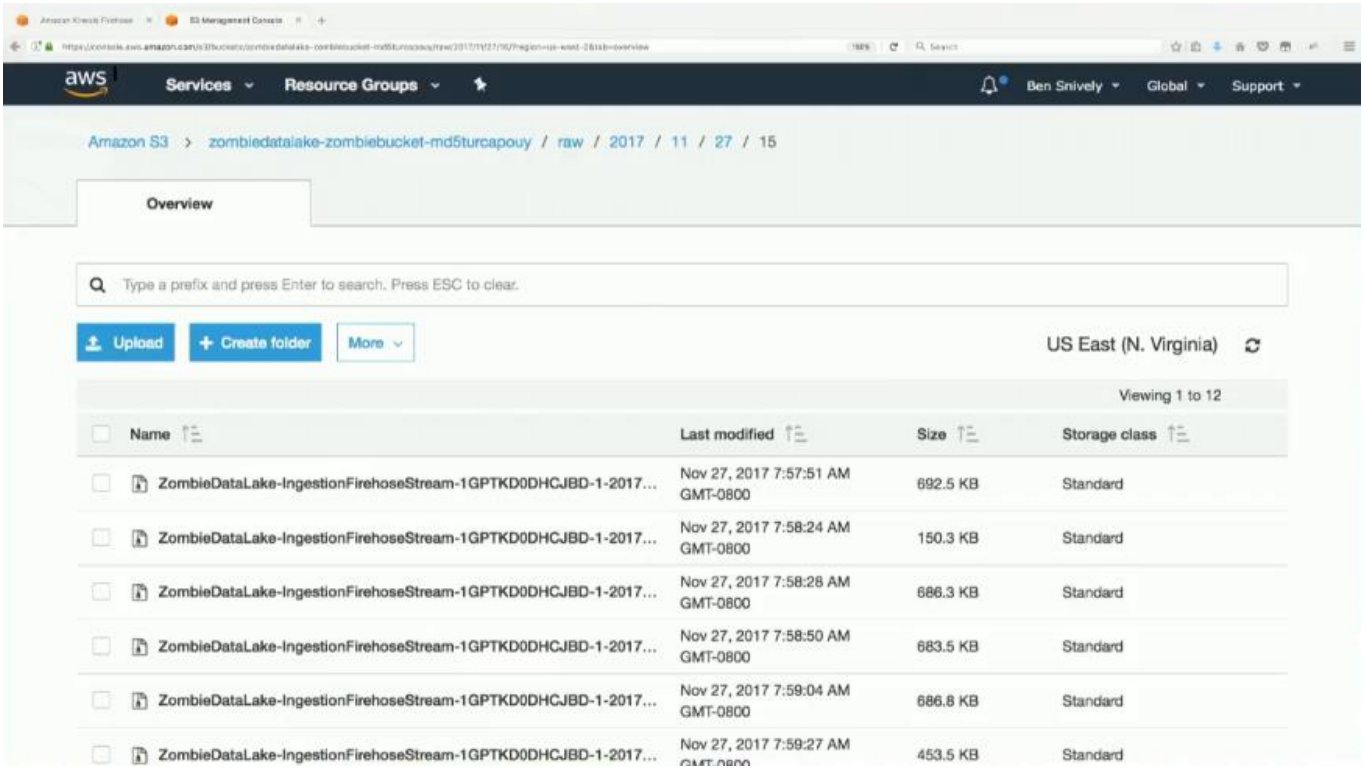
aws
Services
Resource Groups
Ben Snively
Global
Support
Amazon S3 > zombiedatalake-zombiebucket-md5turcapouy / raw / 2017 / 11 / 27 / 15
Overview
Type a prefix and press Enter to search. Press ESC to clear.
Upload
+ Create folder
More
US East (N. Virginia)
Viewing 1 to 12
Name
Last modified
Size
Storage class
ZombieDataLake-IngestionFirehoseStream-1GPTKD0DHCJBD-1-2017...
Nov 27, 2017 7:57:51 AM GMT-0800
692.5 KB
Standard
ZombieDataLake-IngestionFirehoseStream-1GPTKD0DHCJBD-1-2017...
Nov 27, 2017 7:58:24 AM GMT-0800
150.3 KB
Standard
ZombieDataLake-IngestionFirehoseStream-1GPTKD0DHCJBD-1-2017...
Nov 27, 2017 7:58:28 AM GMT-0800
686.3 KB
Standard
ZombieDataLake-IngestionFirehoseStream-1GPTKD0DHCJBD-1-2017...
Nov 27, 2017 7:58:50 AM GMT-0800
683.5 KB
Standard
ZombieDataLake-IngestionFirehoseStream-1GPTKD0DHCJBD-1-2017...
Nov 27, 2017 7:59:04 AM GMT-0800
686.8 KB
Standard
ZombieDataLake-IngestionFirehoseStream-1GPTKD0DHCJBD-1-2017...
Nov 27, 2017 7:59:27 AM
453.5 KB
Standard

```
{"sentiment_value": 0.0, "name": "oceane.batz", "sentiment": "neutral", "timestamp": 1511711811578, "longitude":
20.168331146240234, "country": "Albania", "latitude": 41.1533317565918, "message": "The group scavenged a surgical clinic",
"channel": "default"}
{"sentiment_value": 0.6124, "name": "mohammed.paucek", "sentiment": "positive", "timestamp": 1511711811679, "longitude":
20.168331146240234, "country": "Albania", "latitude": 41.1533317565918, "message": "The group stumbled upon the nearby urgent
care", "channel": "default"}
{"sentiment_value": 0.0, "name": "rodrick.russel", "sentiment": "neutral", "timestamp": 1511711811779, "longitude":
20.168331146240234, "country": "Albania", "latitude": 41.1533317565918, "message": "I discovered a mellow Sugar", "channel":
"default"}
{"sentiment_value": 0.6124, "name": "scottie.hirthe", "sentiment": "positive", "timestamp": 1511711811879, "longitude":
20.168331146240234, "country": "Albania", "latitude": 41.1533317565918, "message": "I hunted down the local urgent care",
"channel": "default"}
{"sentiment_value": 0.0, "name": "friedrich.larson", "sentiment": "neutral", "timestamp": 1511711811980, "longitude":
20.168331146240234, "country": "Albania", "latitude": 41.1533317565918, "message": "This morning the group came upon a naval
doctor's office", "channel": "default"}
{"sentiment_value": 0.0, "name": "evert.satterfield", "sentiment": "neutral", "timestamp": 1511711812085, "longitude":
20.168331146240234, "country": "Albania", "latitude": 41.1533317565918, "message": "The group signed a truce with a golden
swing, reported a municipal clam", "channel": "default"}
{"sentiment_value": 0.0, "name": "willard.bogan", "sentiment": "neutral", "timestamp": 1511711812185, "longitude":
20.168331146240234, "country": "Albania", "latitude": 41.1533317565918, "message": "This morning i spoke to the crib, reported
a famous elephant", "channel": "default"}
{"sentiment_value": 0.4588, "name": "sheldon.cummerata", "sentiment": "positive", "timestamp": 1511711812285, "longitude":
20.168331146240234, "country": "Albania", "latitude": 41.1533317565918, "message": "I ate the sweet Oysters", "channel":
"default"}
{"sentiment_value": 0.0, "name": "felicita.turcotte", "sentiment": "neutral", "timestamp": 1511711812474, "longitude":
20.168331146240234, "country": "Albania", "latitude": 41.1533317565918, "message": "I was converted into a red stove",
"channel": "default"}
{"sentiment_value": -0.34, "name": "madie.schowalter", "sentiment": "negative", "timestamp": 1511711812574, "longitude":
20.168331146240234, "country": "Albania", "latitude": 41.1533317565918, "message": "I plans to cease fire against a swing",
"channel": "default"}
{"sentiment_value": 0.0, "name": "josie.schamberger", "sentiment": "neutral", "timestamp": 1511711812674, "longitude":
20.168331146240234, "country": "Albania", "latitude": 41.1533317565918, "message": "I scavenged a makeshift doctor's office",
"channel": "default"}
{"sentiment_value": 0.0, "name": "lorine.stanton", "sentiment": "neutral", "timestamp": 1511711812775, "longitude":
20.168331146240234, "country": "Albania", "latitude": 41.1533317565918, "message": "I drank the juicy Loquats", "channel":
"default"}
{"sentiment_value": 0.0, "name": "coby.greenholt", "sentiment": "neutral", "timestamp": 1511711812875, "longitude":
```


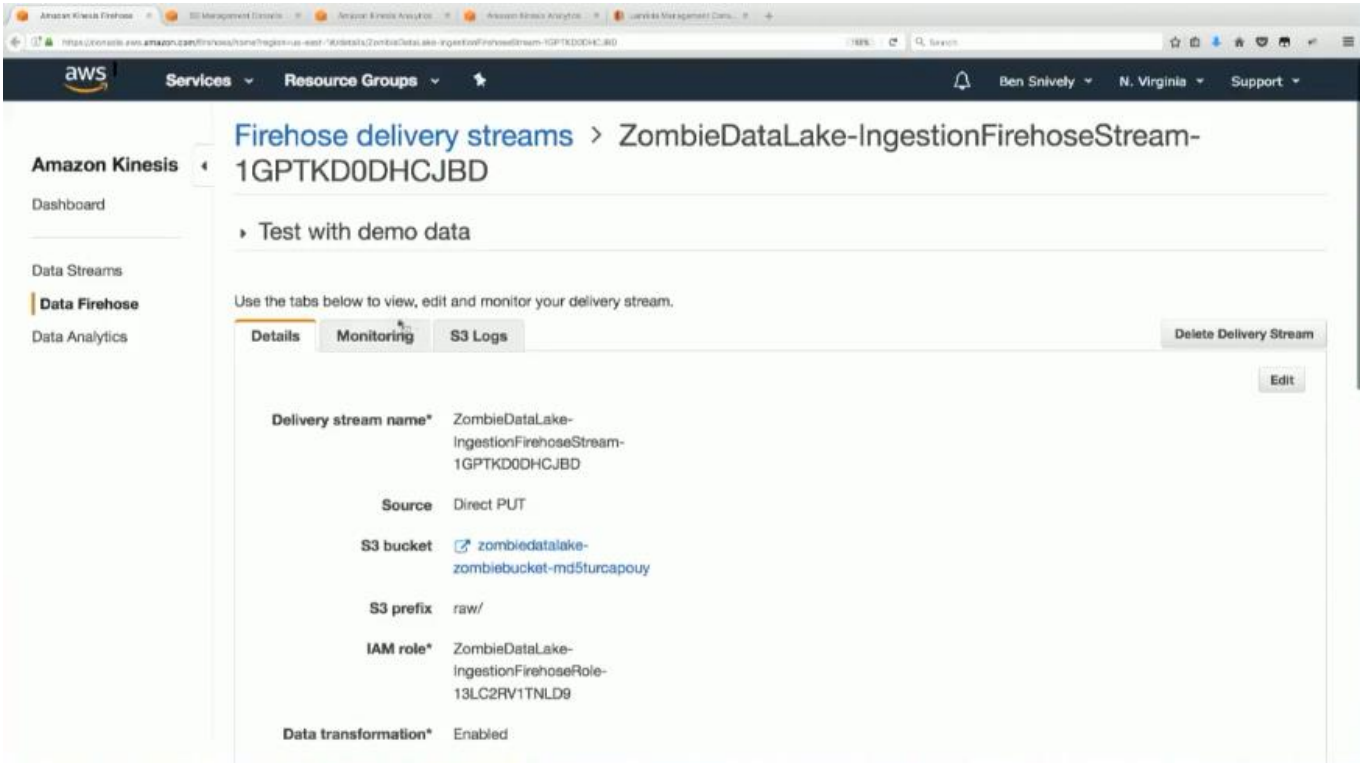

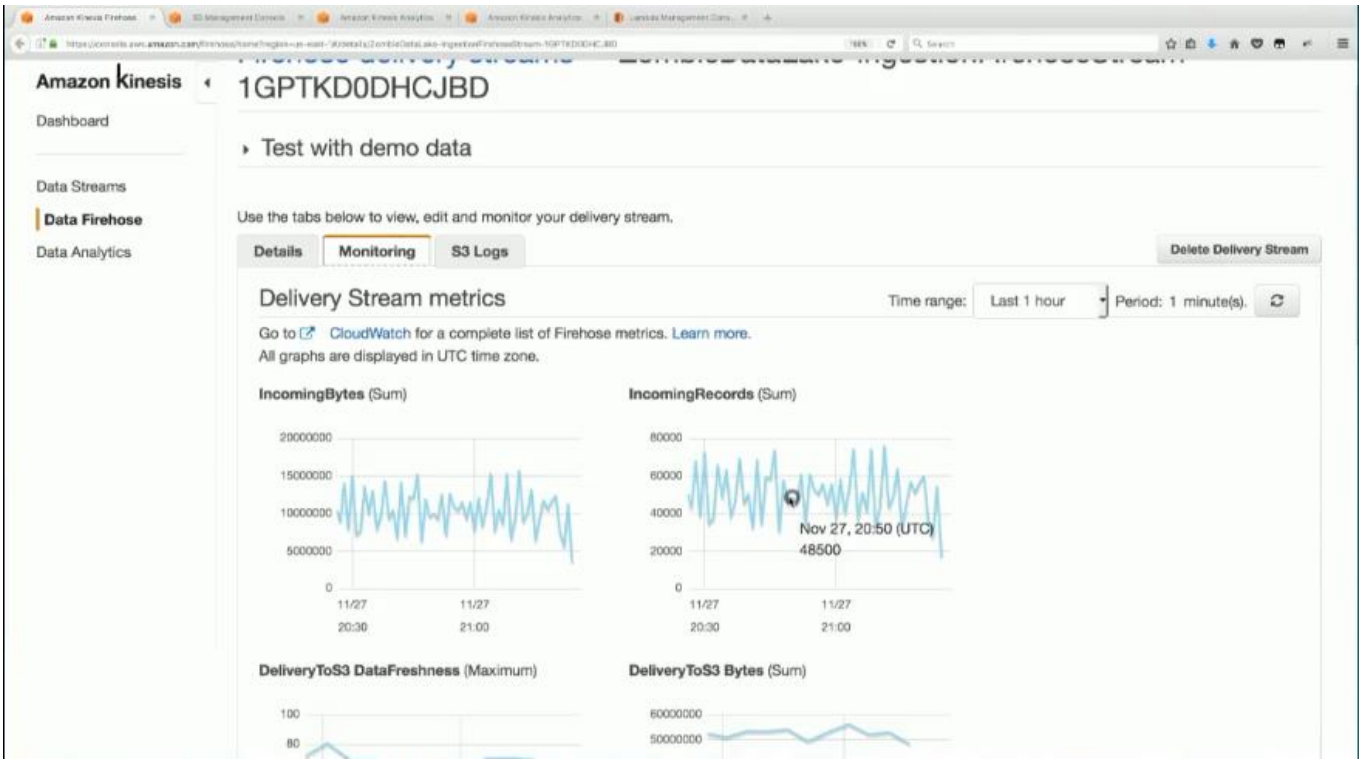

We can easily query this data using Athena or Kinesis Analytics as shown below

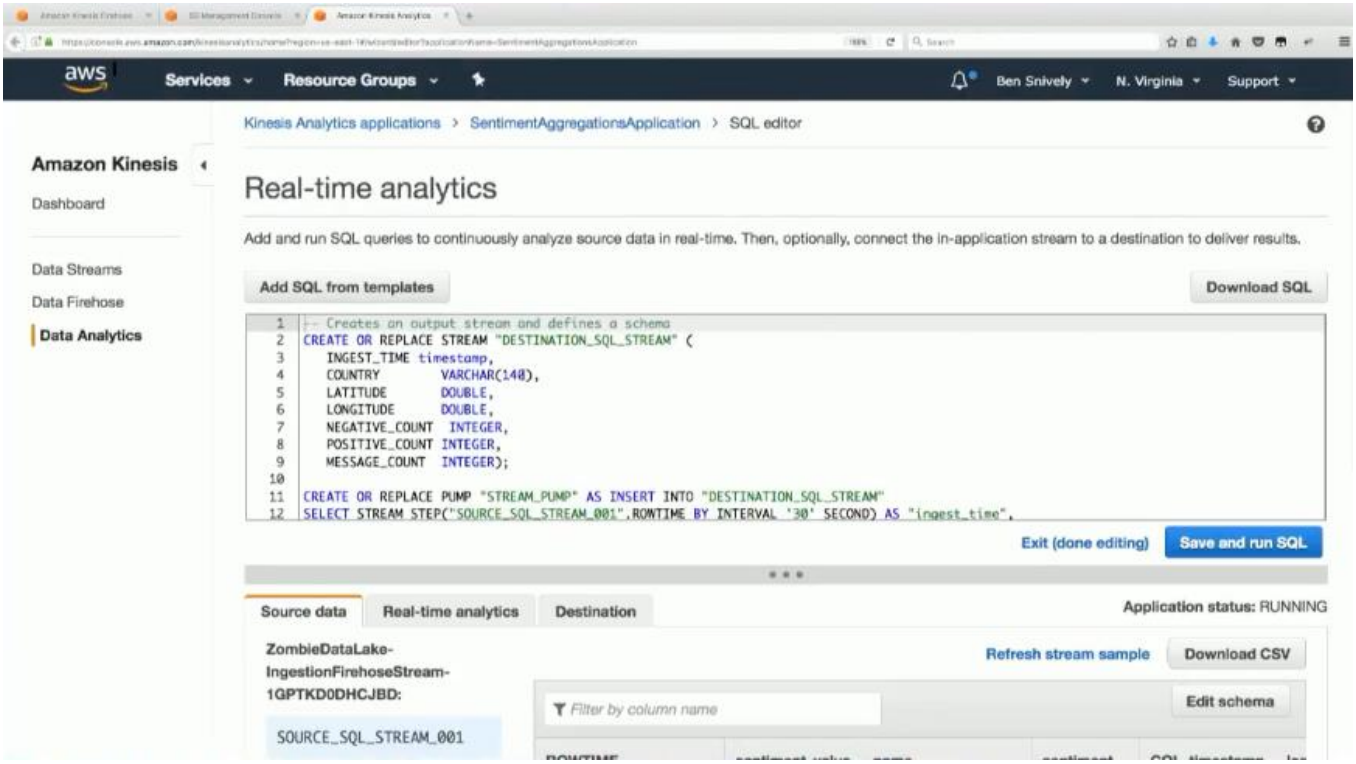

We write a SQL query, we have a SQL statement that does our real time analytics and start getting results

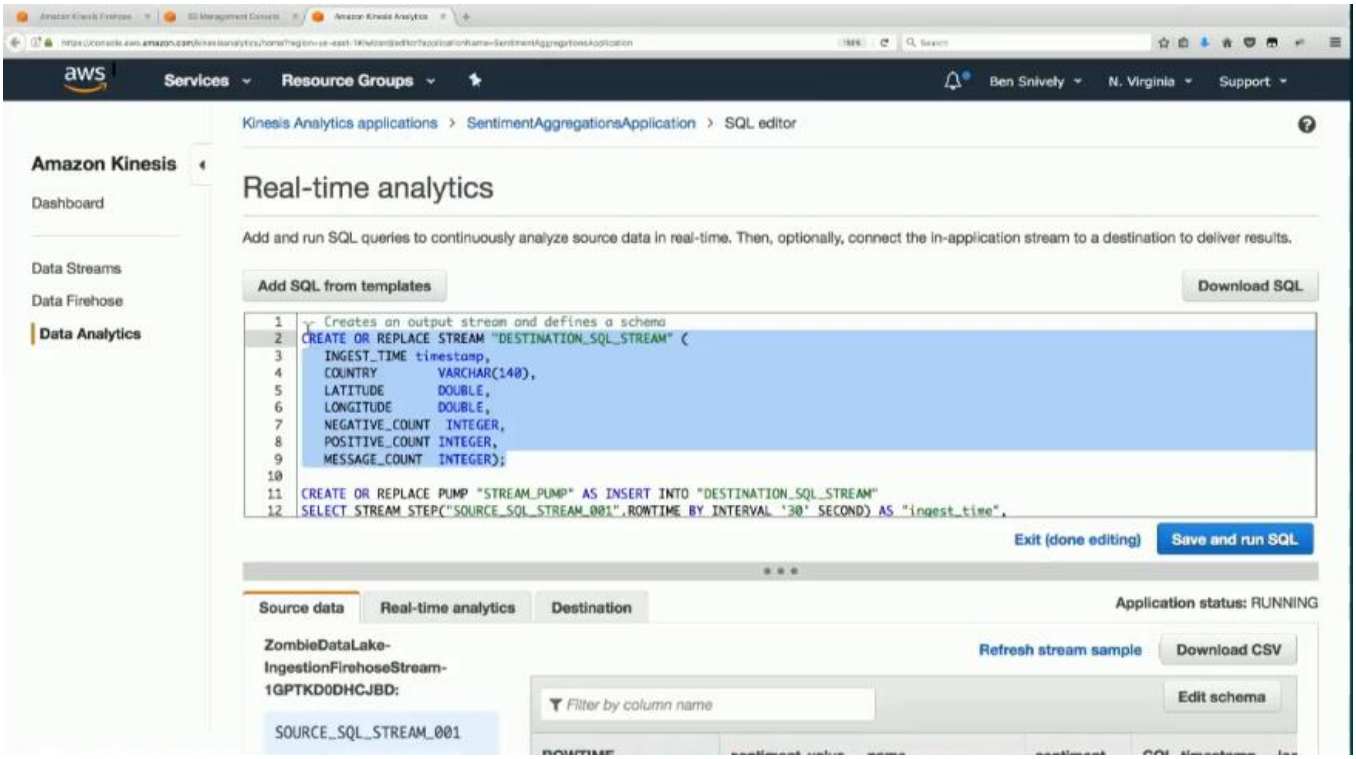

This defines the output using a sliding window

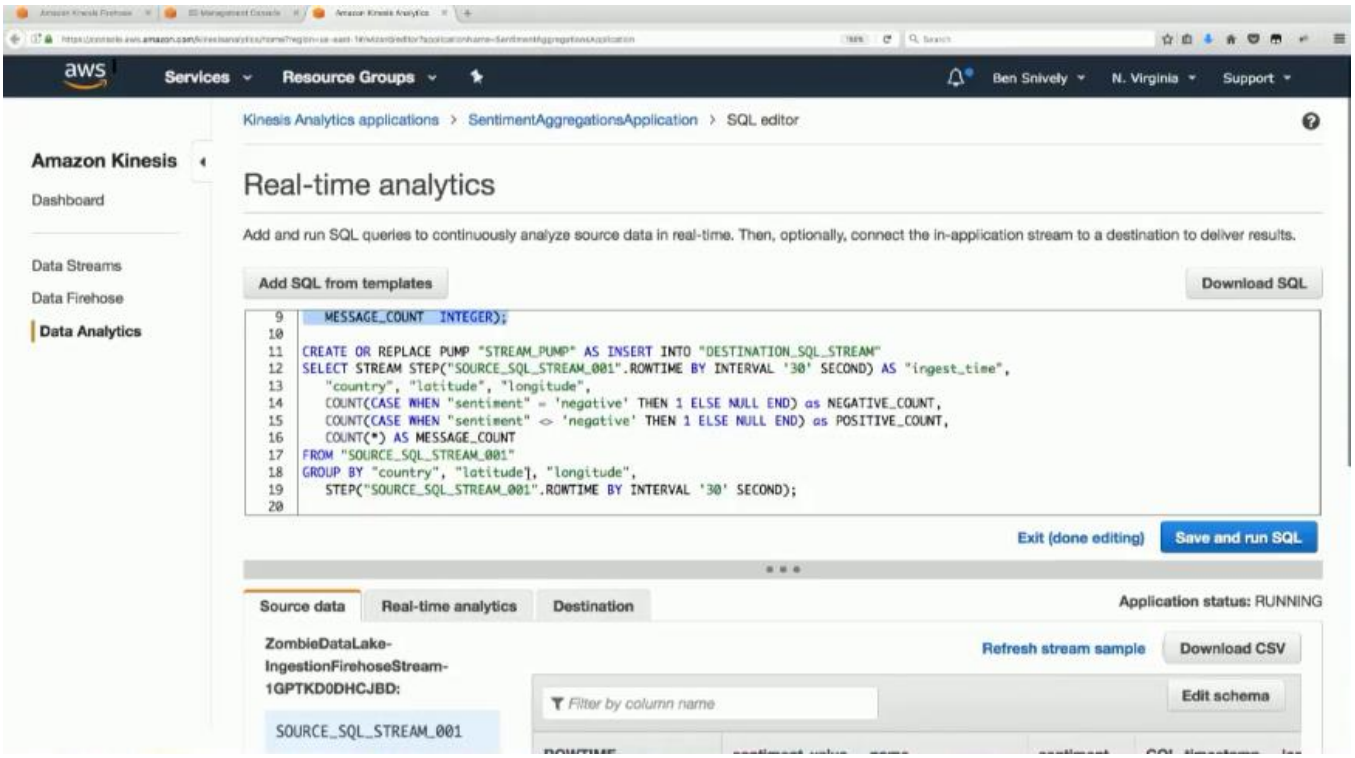

We are now generating the data into the stream, this is doing a 30 second window of data

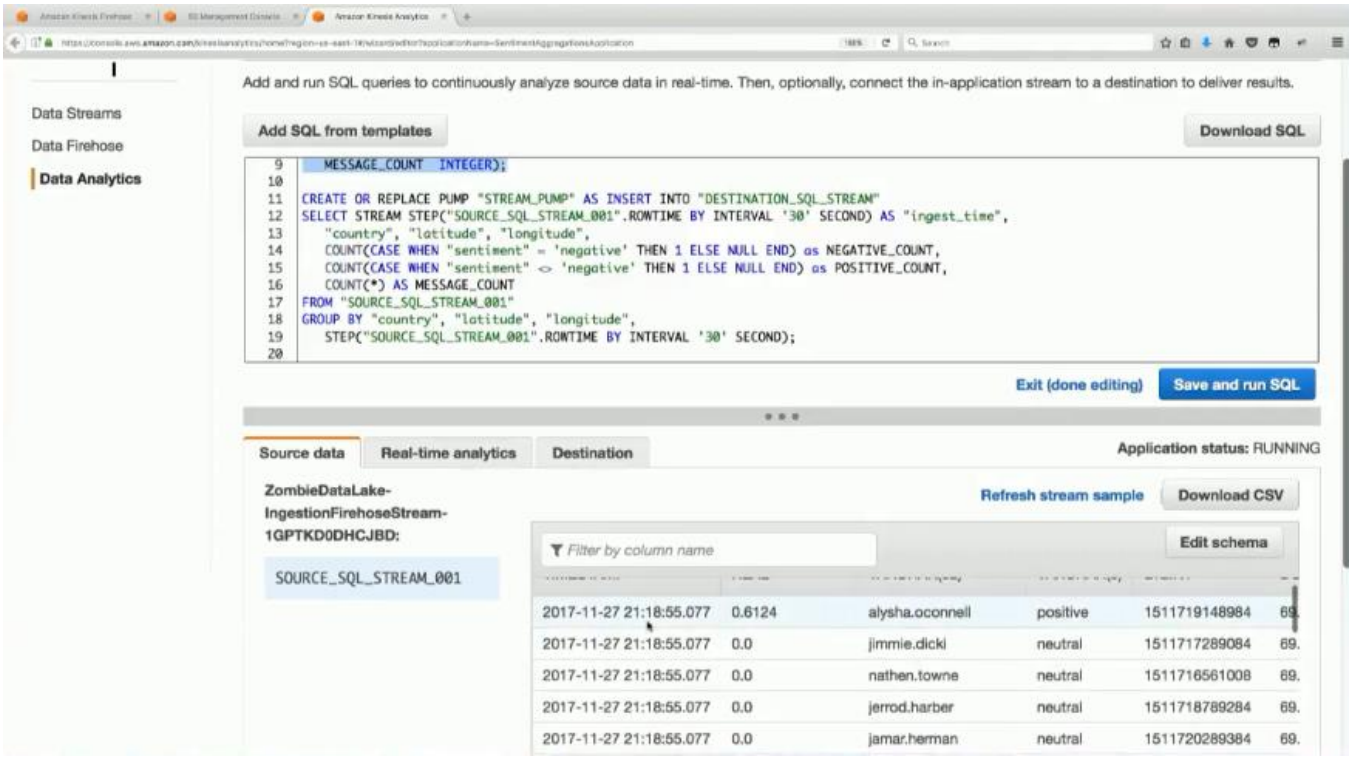
Data Streams
Data Firehose
Data Analytics
Add and run SQL queries to continuously analyze source data in real-time. Then, optionally, connect the in-application stream to a destination to deliver results.
Add SQL from templates
Download SQL
MESSAGE_COUNT INTEGER);
CREATE OR REPLACE PUMP "STREAM_PUMP" AS INSERT INTO "DESTINATION_SQL_STREAM"
SELECT STREAM STEP("SOURCE_SQL_STREAM_001".ROWTIME BY INTERVAL '30' SECOND) AS "ingest_time",
"country", "latitude", "longitude",
COUNT(CASE WHEN "sentiment" = 'negative' THEN 1 ELSE NULL END) as NEGATIVE_COUNT,
COUNT(CASE WHEN "sentiment" <> 'negative' THEN 1 ELSE NULL END) as POSITIVE_COUNT,
COUNT(*) AS MESSAGE_COUNT
FROM "SOURCE_SQL_STREAM_001"
GROUP BY "country", "latitude", "longitude",
STEP("SOURCE_SQL_STREAM_001".ROWTIME BY INTERVAL '30' SECOND);
Exit (done editing)
Save and run SQL
Source data
Real-time analytics
Destination
Application status: RUNNING
ZombieDataLake-IngestionFirehoseStream-1GPTKD0DHCJBD:
SOURCE_SQL_STREAM_001
Refresh stream sample
Download CSV
Edit schema
Filter by column name
2017-11-27 21:18:55.077 0.6124 alysha.oconnell positive 1511719148984
2017-11-27 21:18:55.077 0.0 jimmie.dicki neutral 1511717289084
2017-11-27 21:18:55.077 0.0 nathen.towne neutral 1511716561008
2017-11-27 21:18:55.077 0.0 jerrod.harber neutral 1511718789284
2017-11-27 21:18:55.077 0.0 jamar.herman neutral 1511720289384

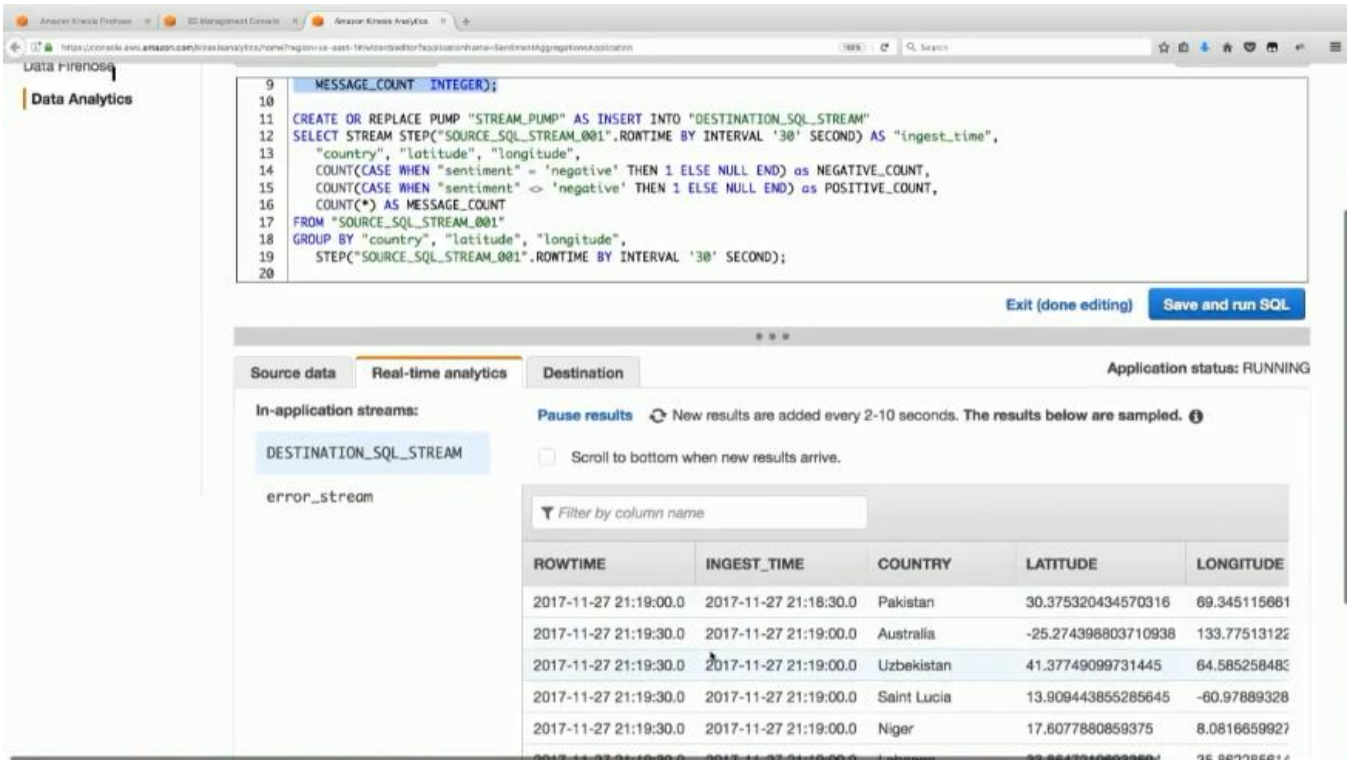
Data Firehose
Data Analytics
MESSAGE_COUNT INTEGER);
CREATE OR REPLACE PUMP "STREAM_PUMP" AS INSERT INTO "DESTINATION_SQL_STREAM"
SELECT STREAM STEP("SOURCE_SQL_STREAM_001".ROWTIME BY INTERVAL '30' SECOND) AS "ingest_time",
"country", "latitude", "longitude",
COUNT(CASE WHEN "sentiment" = 'negative' THEN 1 ELSE NULL END) as NEGATIVE_COUNT,
COUNT(CASE WHEN "sentiment" <> 'negative' THEN 1 ELSE NULL END) as POSITIVE_COUNT,
COUNT(*) AS MESSAGE_COUNT
FROM "SOURCE_SQL_STREAM_001"
GROUP BY "country", "latitude", "longitude",
STEP("SOURCE_SQL_STREAM_001".ROWTIME BY INTERVAL '30' SECOND);
Exit (done editing)
Save and run SQL
Source data
Real-time analytics
Destination
Application status: RUNNING
In-application streams:
DESTINATION_SQL_STREAM
error_stream
Pause results
New results are added every 2-10 seconds. The results below are sampled.
Scroll to bottom when new results arrive.
Filter by column name
ROWTIME INGEST_TIME COUNTRY LATITUDE LONGITUDE
2017-11-27 21:19:00.0 2017-11-27 21:18:30.0 Pakistan 30.375320434570316 69.345115661
2017-11-27 21:19:30.0 2017-11-27 21:19:00.0 Australia -25.274398803710938 133.77513122
2017-11-27 21:19:30.0 2017-11-27 21:19:00.0 Uzbekistan 41.37749099731445 64.58525848
2017-11-27 21:19:30.0 2017-11-27 21:19:00.0 Saint Lucia 13.909443855285645 -60.97889328
2017-11-27 21:19:30.0 2017-11-27 21:19:00.0 Niger 17.6077880859375 8.081665992

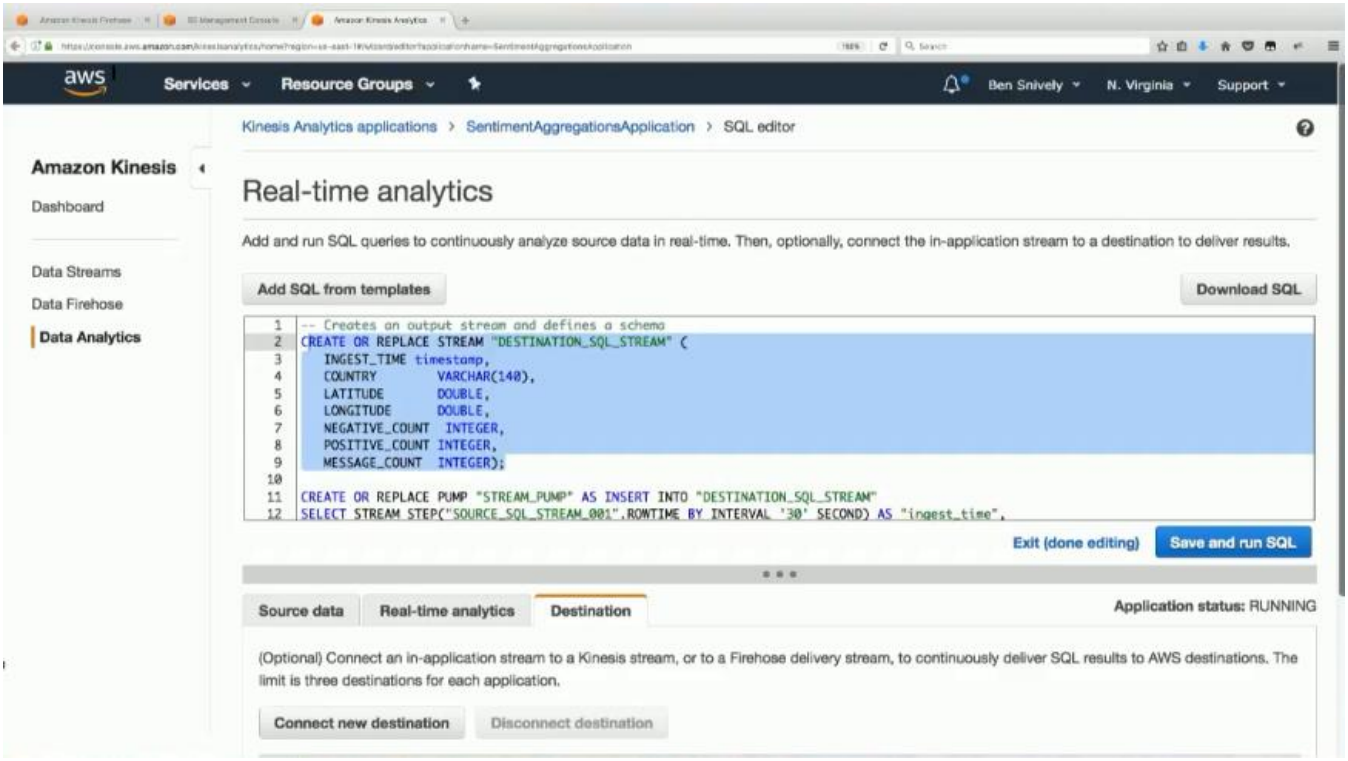

We can specify a destination like S3 for the analyzed result

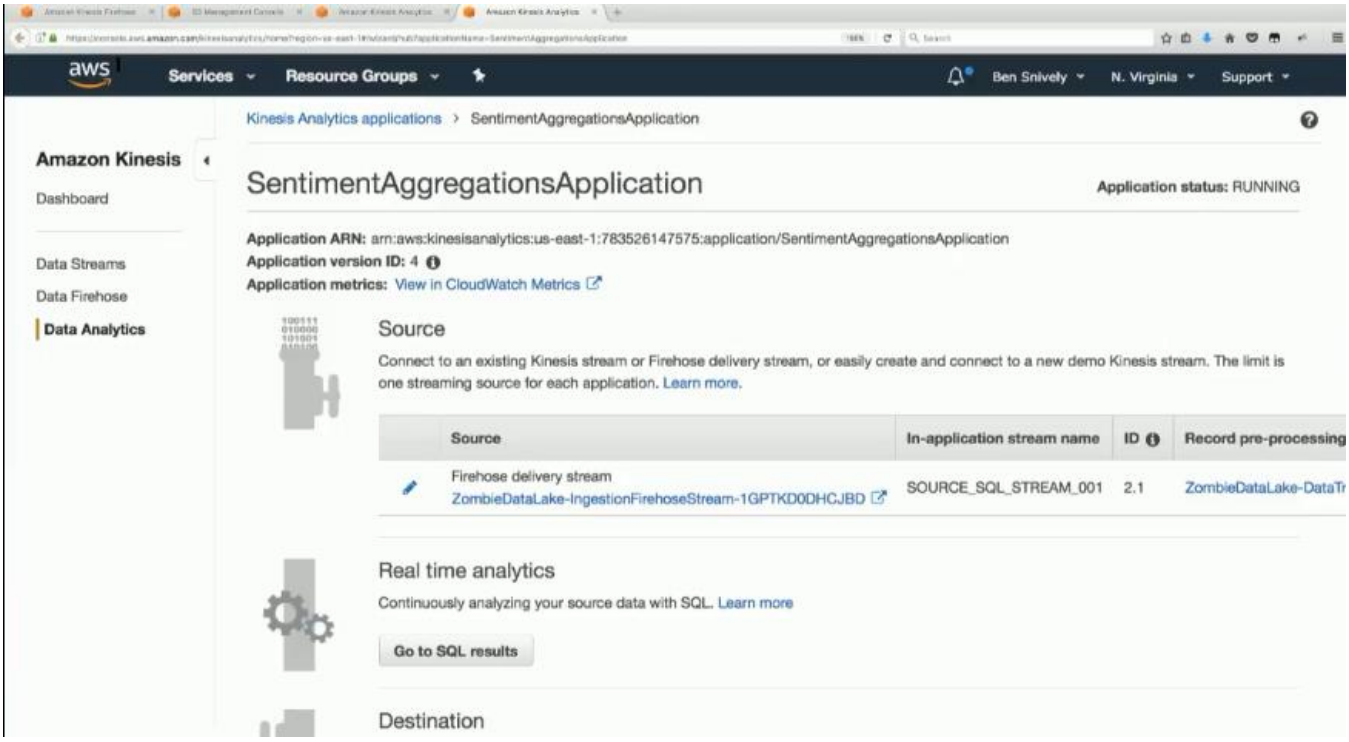

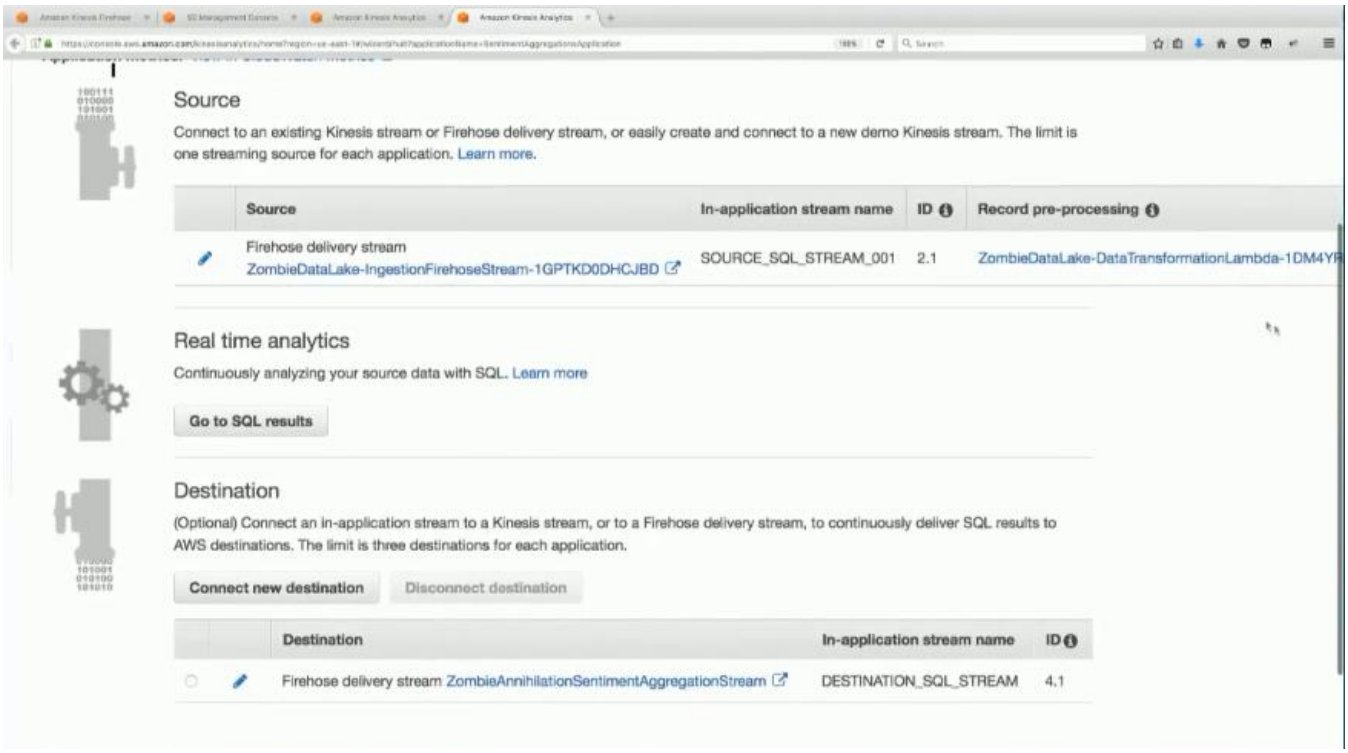
Source
Connect to an existing Kinesis stream or Firehose delivery stream, or easily create and connect to a new demo Kinesis stream. The limit is one streaming source for each application. Learn more.
Source
In-application stream name
ID
Record pre-processing
Firehose delivery stream
ZombieDataLake-IngestionFirehoseStream-1GPTKD0DHCJBD
SOURCE_SQL_STREAM_001
2.1
ZombieDataLake-DataTransformationLambda-1DM4YR
Real time analytics
Continuously analyzing your source data with SQL. Learn more
Go to SQL results
Destination
(Optional) Connect an in-application stream to a Kinesis stream, or to a Firehose delivery stream, to continuously deliver SQL results to AWS destinations. The limit is three destinations for each application.
Connect new destination
Disconnect destination
Destination
In-application stream name
ID
Firehose delivery stream ZombieAnnihilationSentimentAggregationStream
DESTINATION_SQL_STREAM
4.1

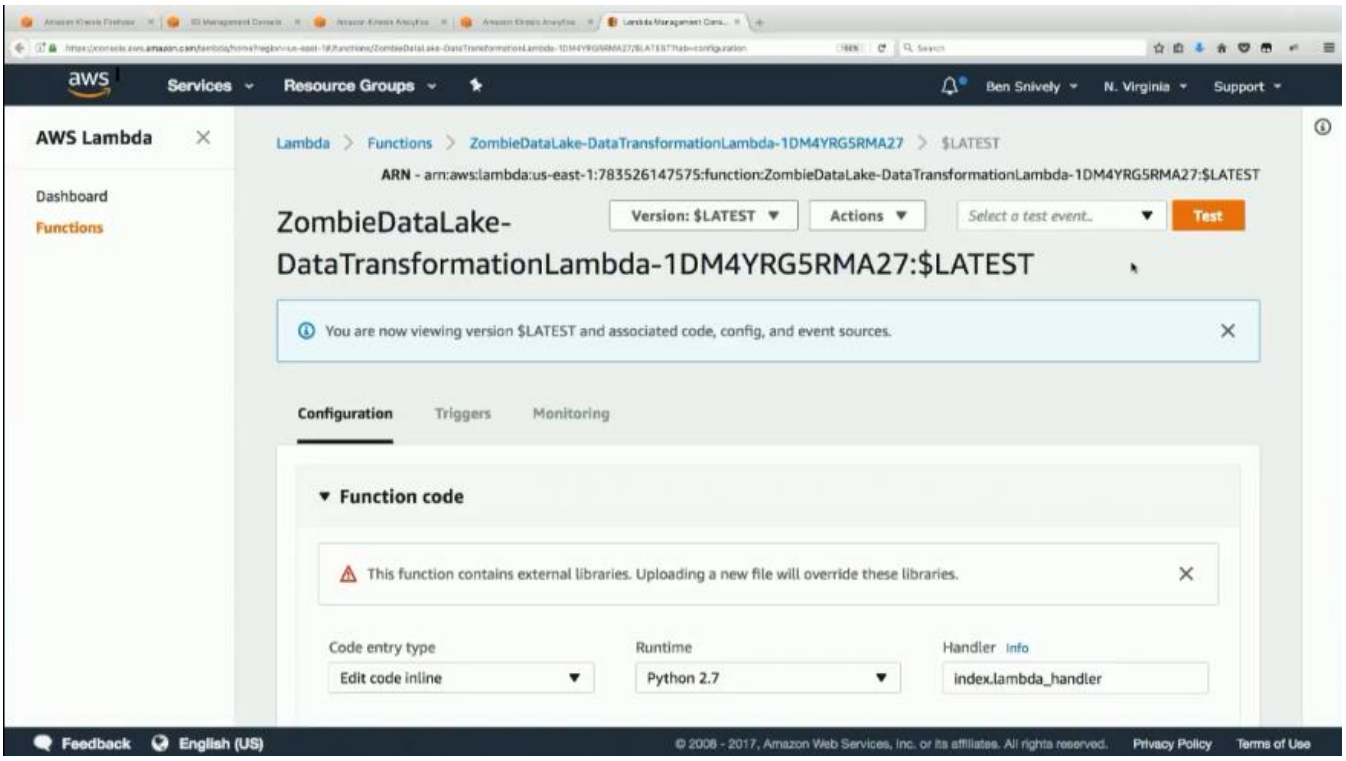
aws
Services
Resource Groups
Ben Snively
N. Virginia
Support
AWS Lambda
Dashboard
Functions
Lambda > Functions > ZombieDataLake-DataTransformationLambda-1DM4YRG5RMA27 > $LATEST
ARN - arn:aws:lambda:us-east-1:783526147575:function:ZombieDataLake-DataTransformationLambda-1DM4YRG5RMA27:$LATEST
ZombieDataLake-DataTransformationLambda-1DM4YRG5RMA27:$LATEST
Version: $LATEST
Actions
Select a test event..
Test
You are now viewing version $LATEST and associated code, config, and event sources.
Configuration
Triggers
Monitoring
Function code
This function contains external libraries. Uploading a new file will override these libraries.
Code entry type
Edit code inline
Runtime
Python 2.7
Handler Info
index.lambda_handler
Feedback
English (US)
© 2008 - 2017, Amazon Web Services, Inc. or its affiliates. All rights reserved.
Privacy Policy
Terms of Use

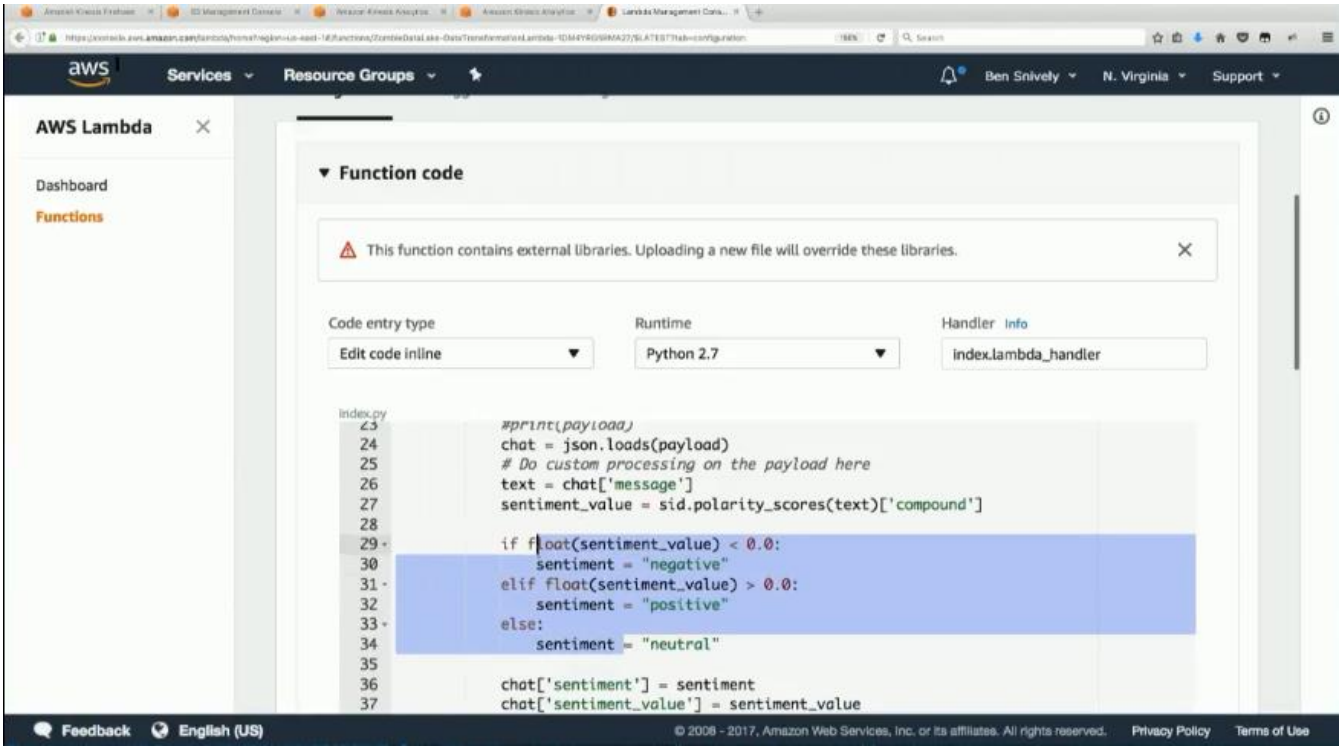

We can run some custom lambda code to transform or enrich the data coming into the stream before storing it

```
2017-11-27 16:25:00.000,Fiji,-16.5781936645078,179.41441345214844,1,249,250
2017-11-27 16:25:00.000,India,20.59368324279785,78.96288299560547,5,245,250
2017-11-27 16:25:00.000,Japan,36.2048225402832,138.2529296875,5,245,250
2017-11-27 16:25:00.000,Saint Lucia,13.909443855285645,-60.978893280029304,5,245,250
2017-11-27 16:25:00.000,Turkey,38.9637451171875,35.24332046508789,2,248,250
2017-11-27 16:25:00.000,Burkina Faso,12.238332748413088,-1.5615930557250977,2,248,250
2017-11-27 16:25:00.000,Cambodia,12.565678596496584,104.990966796875,1,249,250
2017-11-27 16:25:00.000,Spain,40.46366882324219,-3.74921989440918,2,248,250
2017-11-27 16:25:00.000,Australia,-25.274398803710938,133.77513122558594,0,250,250
2017-11-27 16:25:00.000,Niger,17.6077880859375,8.081665992736816,0,250,250
2017-11-27 16:25:00.000,Uruguay,-32.52277755737305,-55.76583480834961,0,250,250
2017-11-27 16:25:00.000,Cayman Islands,19.513469696044922,-80.56695556640625,1,249,250
2017-11-27 16:25:00.000,El Salvador,13.79418468475342,-88.89653015136719,10,240,250
2017-11-27 16:25:00.000,Ethiopia,9.145000457763672,40.48967361450195,1,249,250
2017-11-27 16:25:00.000,Costa Rica,9.748916625976562,-83.75342559814453,2,248,250
2017-11-27 16:25:00.000,Bhutan,27.514162063598633,90.43360137939452,0,250,250
2017-11-27 16:25:00.000,Martinique,14.641528129577637,-61.024173736572266,7,243,250
2017-11-27 16:25:00.000,Egypt,26.820552825927738,30.80249786376953,1,249,250
2017-11-27 16:25:00.000,Nepal,28.39485740661621,84.12400817871094,4,246,250
2017-11-27 16:25:00.000,Togo,8.619543075561523,0.8247820138931273,3,247,250
2017-11-27 16:25:00.000,Portugal,39.399871826171875,-8.224453926086426,1,249,250
2017-11-27 16:25:00.000,Chile,-35.675148010253906,-71.54296875,0,250,250
2017-11-27 16:25:00.000,Bahrain,25.9304141998291,50.63777160644531,4,246,250
2017-11-27 16:25:00.000,Madagascar,-18.76694679260254,46.86910629272461,0,250,250
2017-11-27 16:25:00.000,Greece,39.0742073059082,21.82431221008301,1,249,250
2017-11-27 16:25:00.000,Mauritius,-20.348403930664066,57.55215072631836,10,240,250
2017-11-27 16:25:00.000,Montserrat,16.74249839782715,-62.1873664855957,6,244,250
2017-11-27 16:25:00.000,Mozambique,-18.66569519042969,35.529563903808594,0,250,250
2017-11-27 16:25:00.000,Norway,60.472023010253906,8.46894645690918,0,250,250
2017-11-27 16:25:00.000,Venezuela,6.423749923706056,-66.58972930908203,0,250,250
2017-11-27 16:25:00.000,Saint Kitts and ,17.35782241821289,-62.78299713134766,6,244,250
```

This is what the analyzed result looks like in S3

Lambda Pre-processing

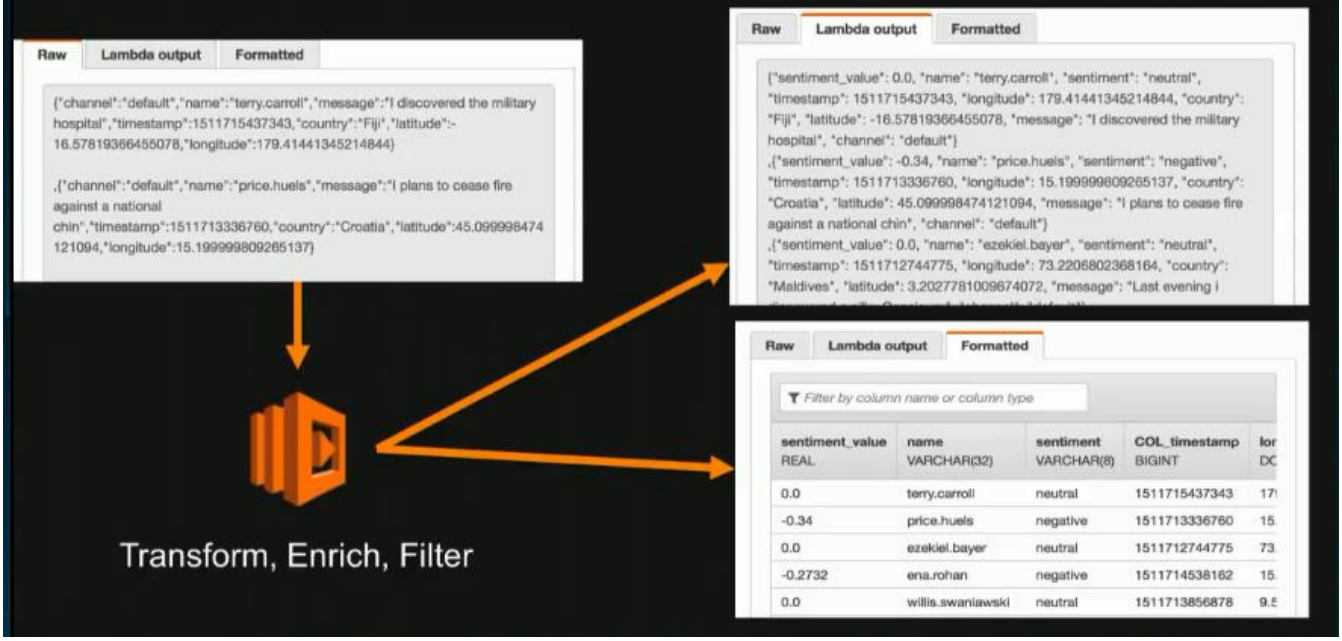

Amazon Kinesis Analytics - SQL

In this example:

19 Lines of SQL = Serverless Realtime Analytics

```
11 CREATE OR REPLACE PUMP "STREAM_PUMP" AS INSERT INTO "DESTINATION_SQL_STREAM"
12 SELECT STREAM STEP("SOURCE_SQL_STREAM_001".ROWTIME BY INTERVAL '30' SECOND) AS "ingest_time",
13    "country", "latitude", "longitude",
14    COUNT(CASE WHEN "sentiment" = 'negative' THEN 1 ELSE NULL END) as NEGATIVE_COUNT,
15    COUNT(CASE WHEN "sentiment" <> 'negative' THEN 1 ELSE NULL END) as POSITIVE_COUNT,
16    COUNT(*) AS MESSAGE_COUNT
17 FROM "SOURCE_SQL_STREAM_001"
18 GROUP BY "country", "latitude", "longitude",
19    STEP("SOURCE_SQL_STREAM_001".ROWTIME BY INTERVAL '30' SECOND);
```

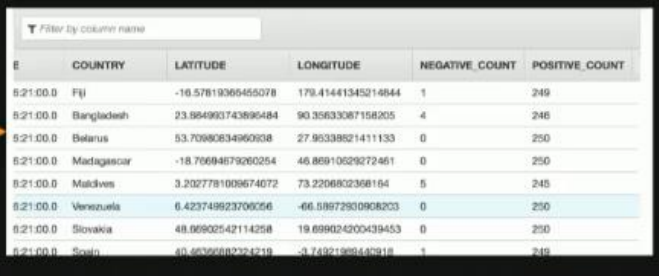

Larger Picture of what we showed:

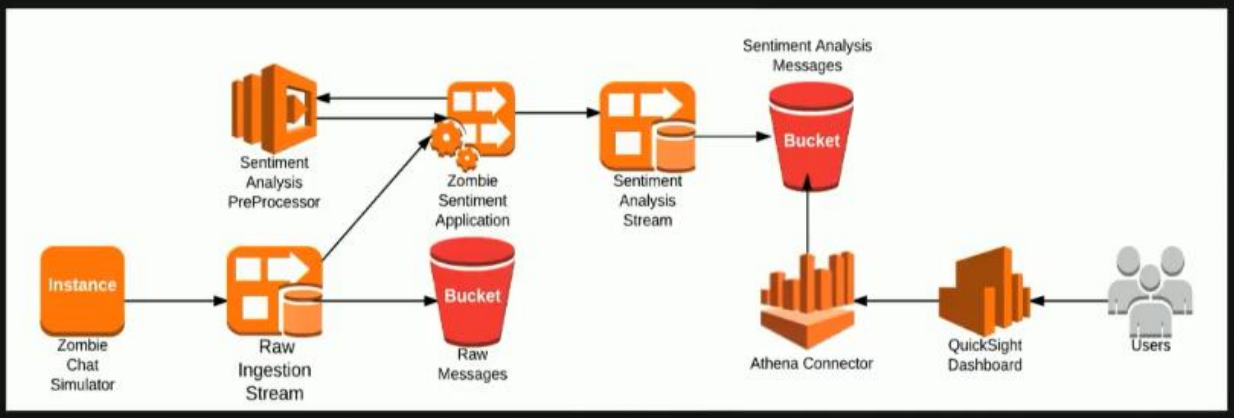

Fitting into existing Real-time Analytics

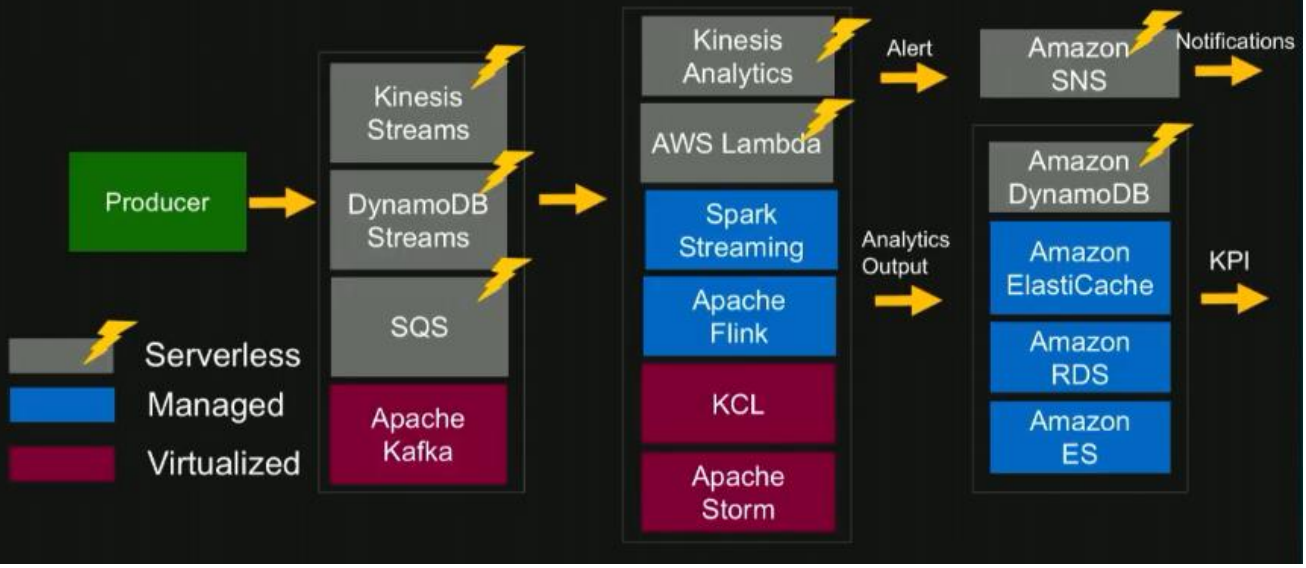

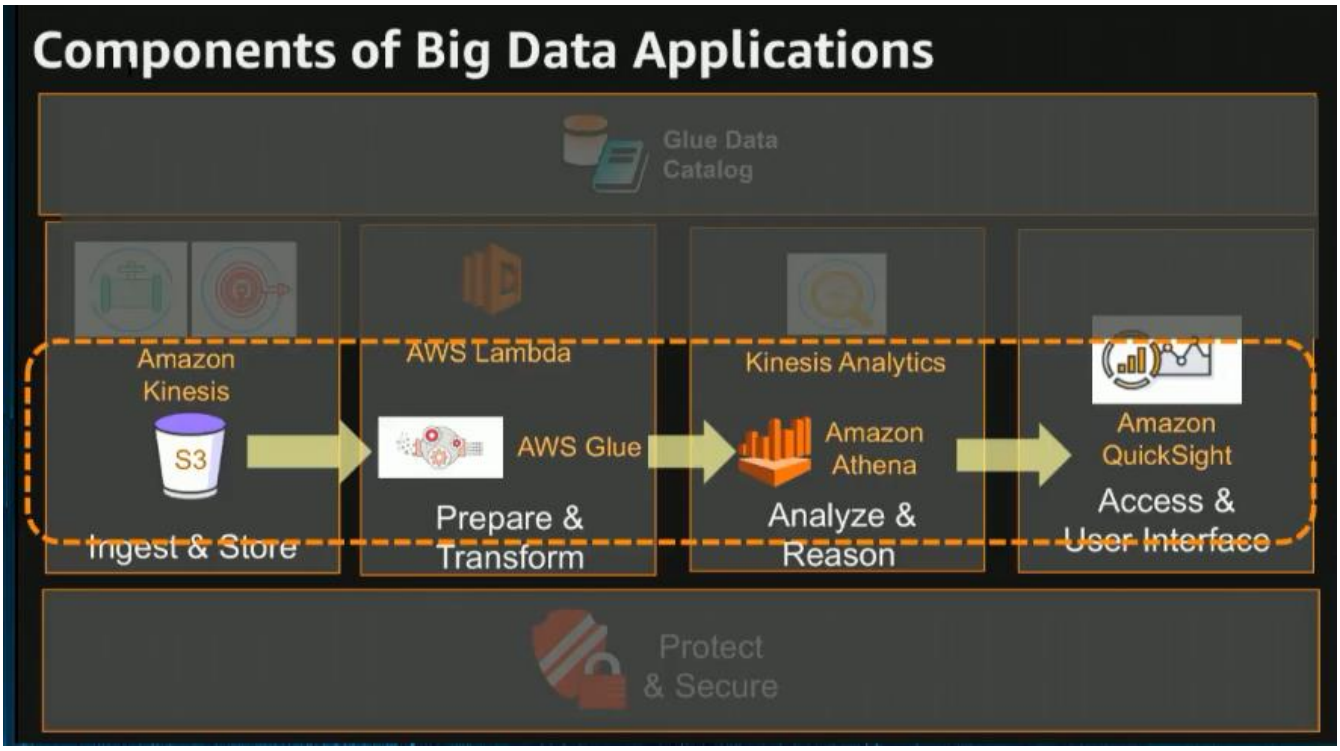

Let us now do some interactive analytics on some of our data at rest in S3.

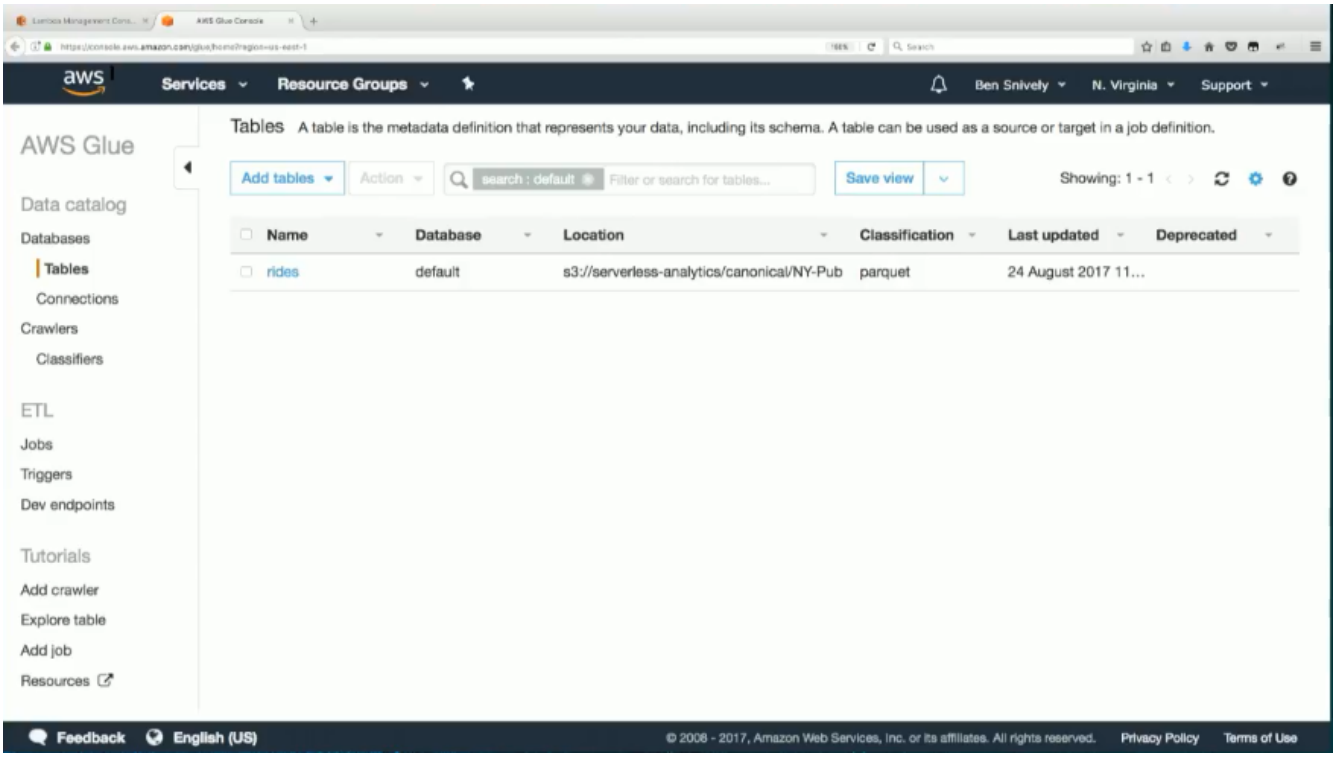

We are starting with the Glue catalog to help describe our data on S3

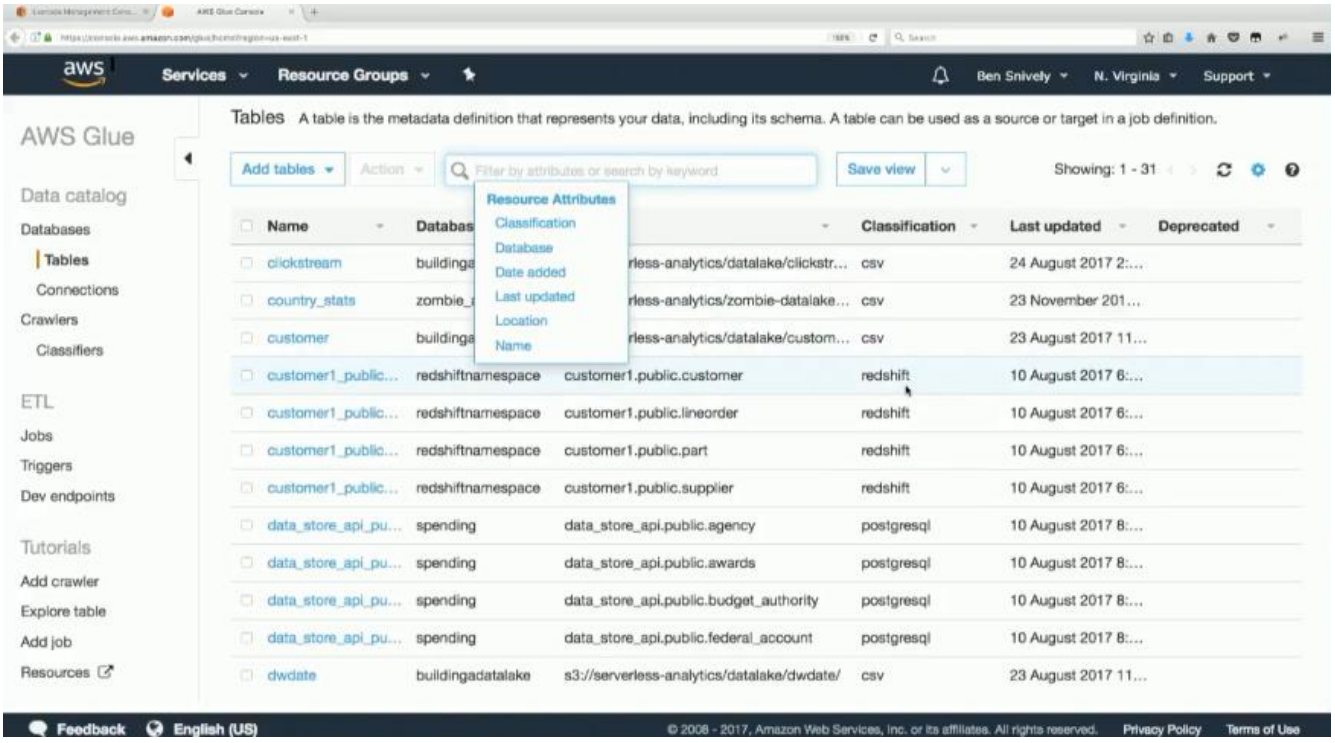

There are actually several data type available in the S3 location, we can filter the data by type

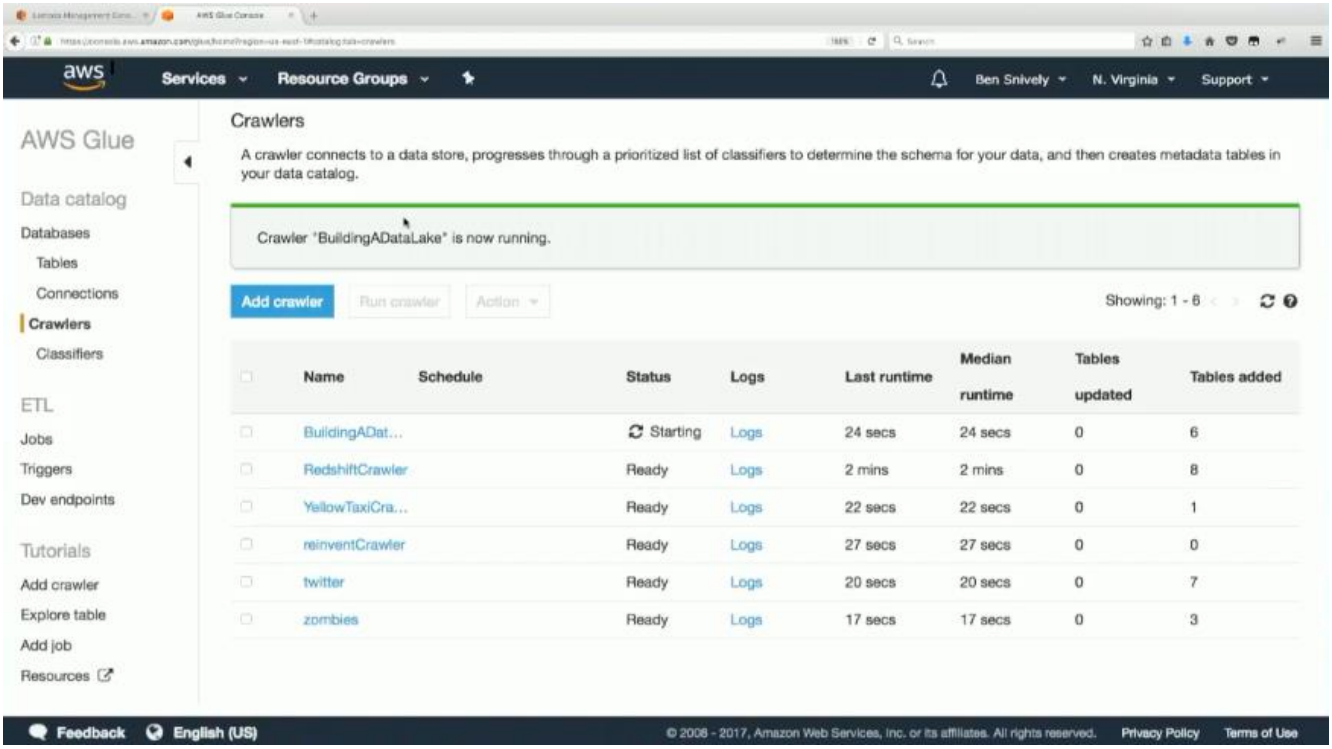

We can use Glue to crawl a new S3 location, the crawlers run with an IAM role that we set up for the crawler to allow it crawl our data on S3 to discover all the new tables in there

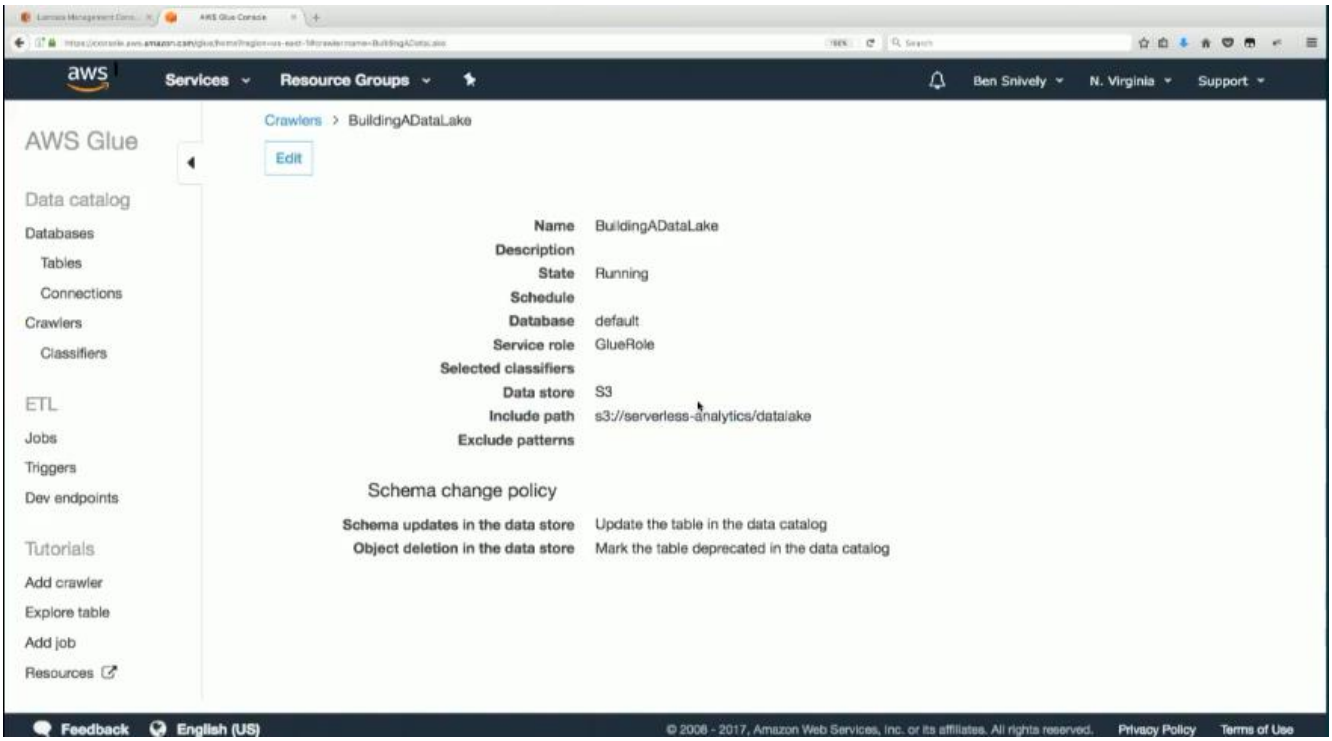

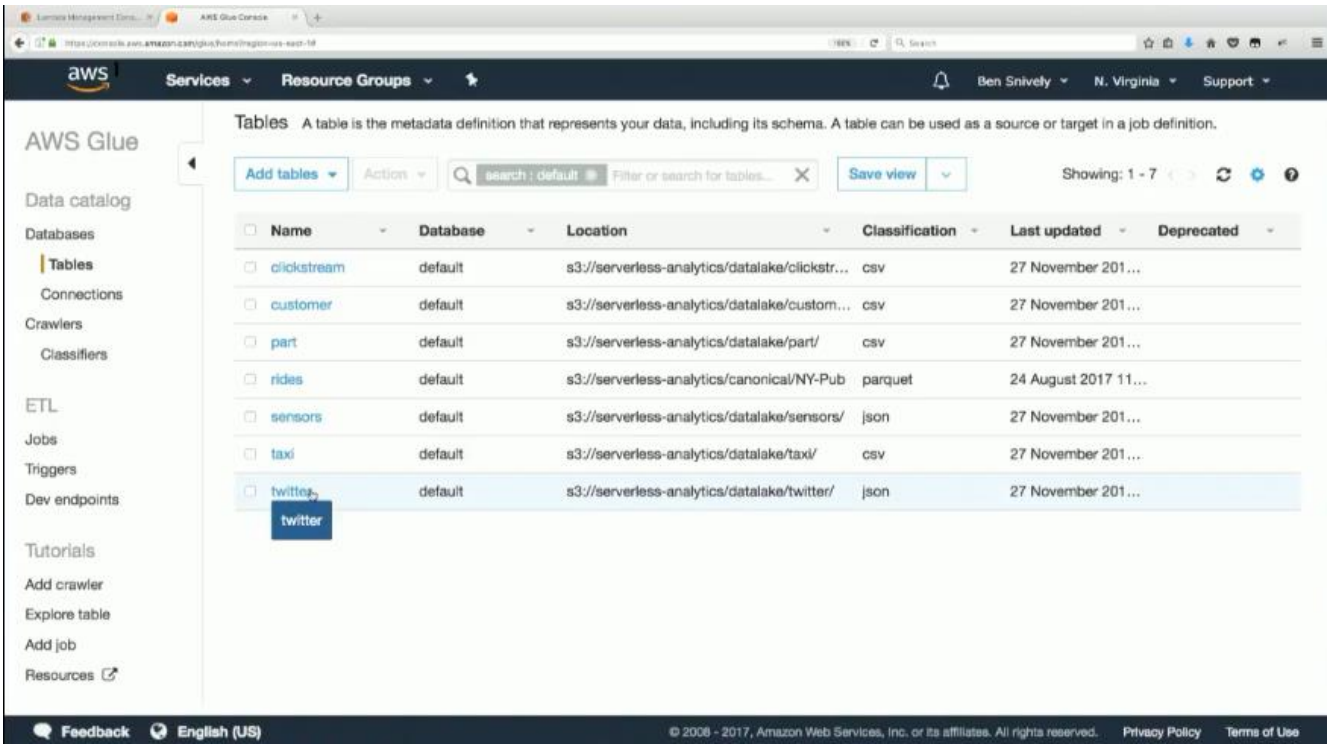

We now have several tables created for us by the crawler with the data sets discovered

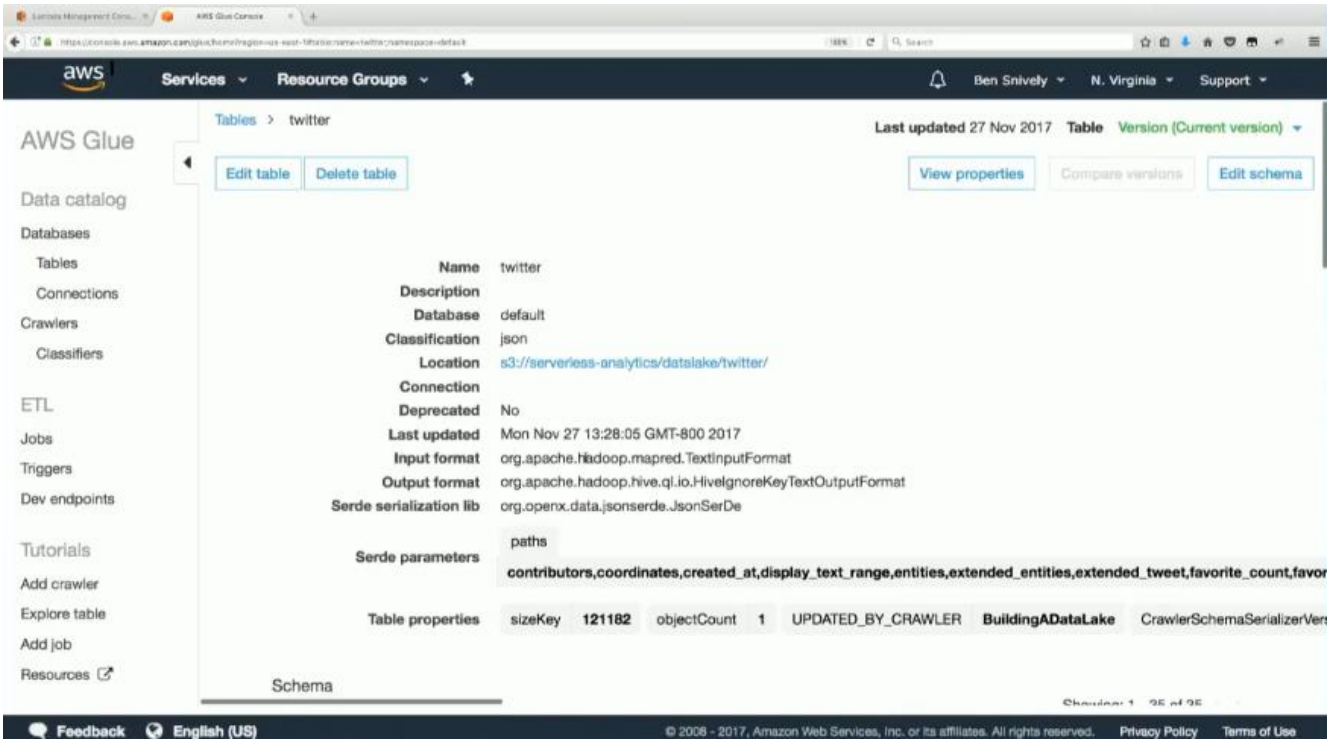

This is the twitter data it discovered, we can now use this in our analytics ad queries

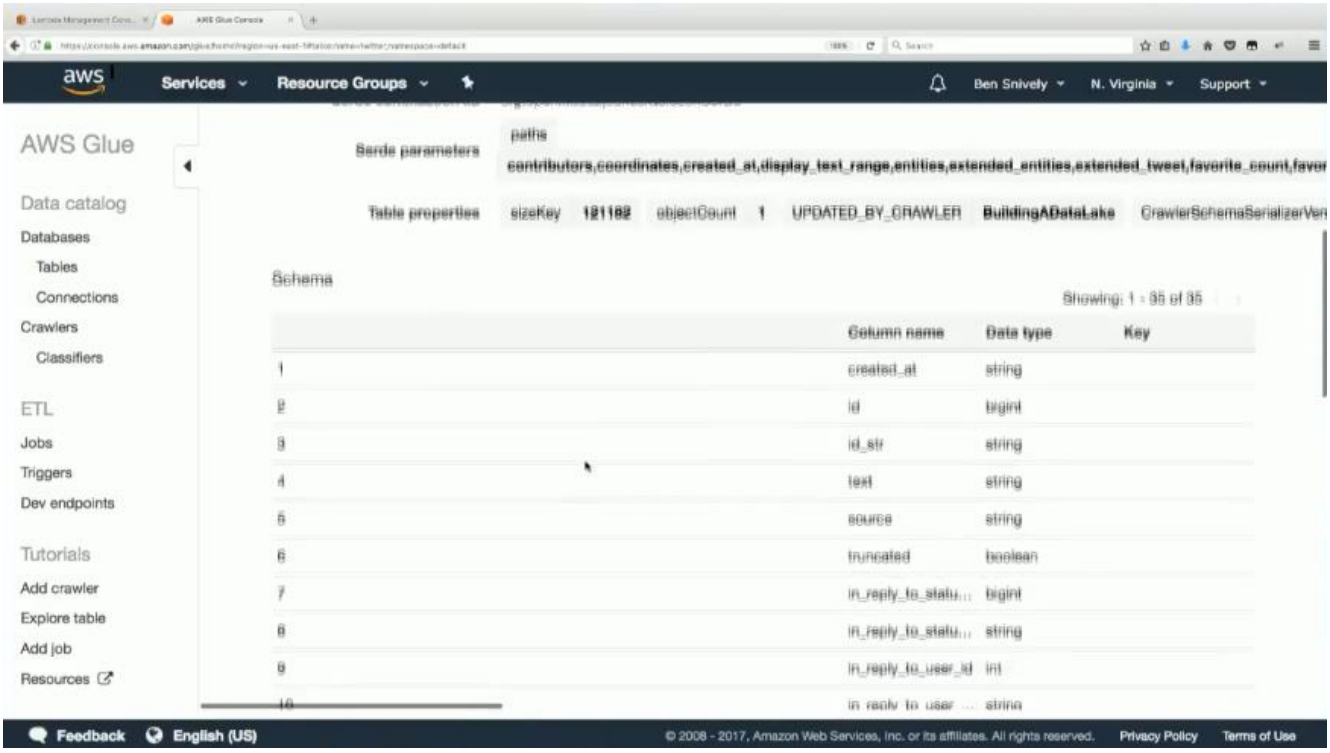

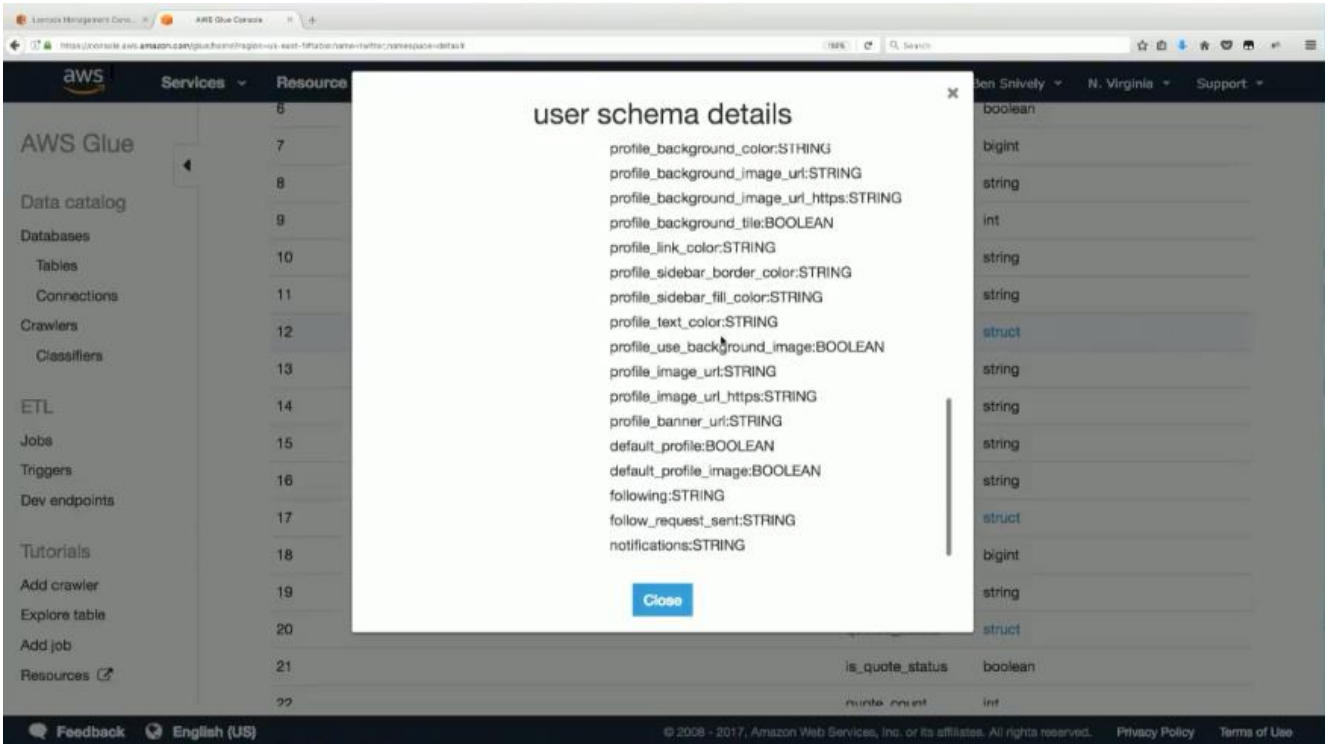

It also discovered the JSON that can represent the twitter data, we can modify this schema as we like

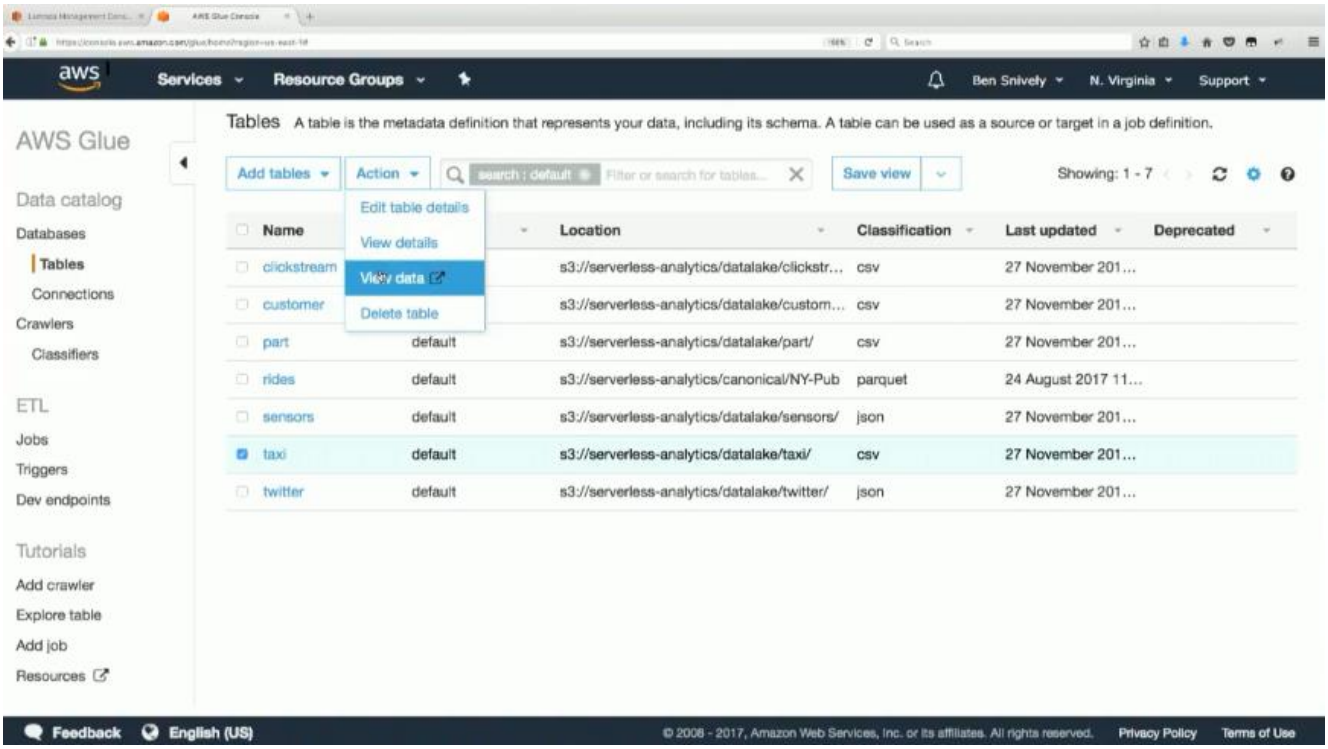

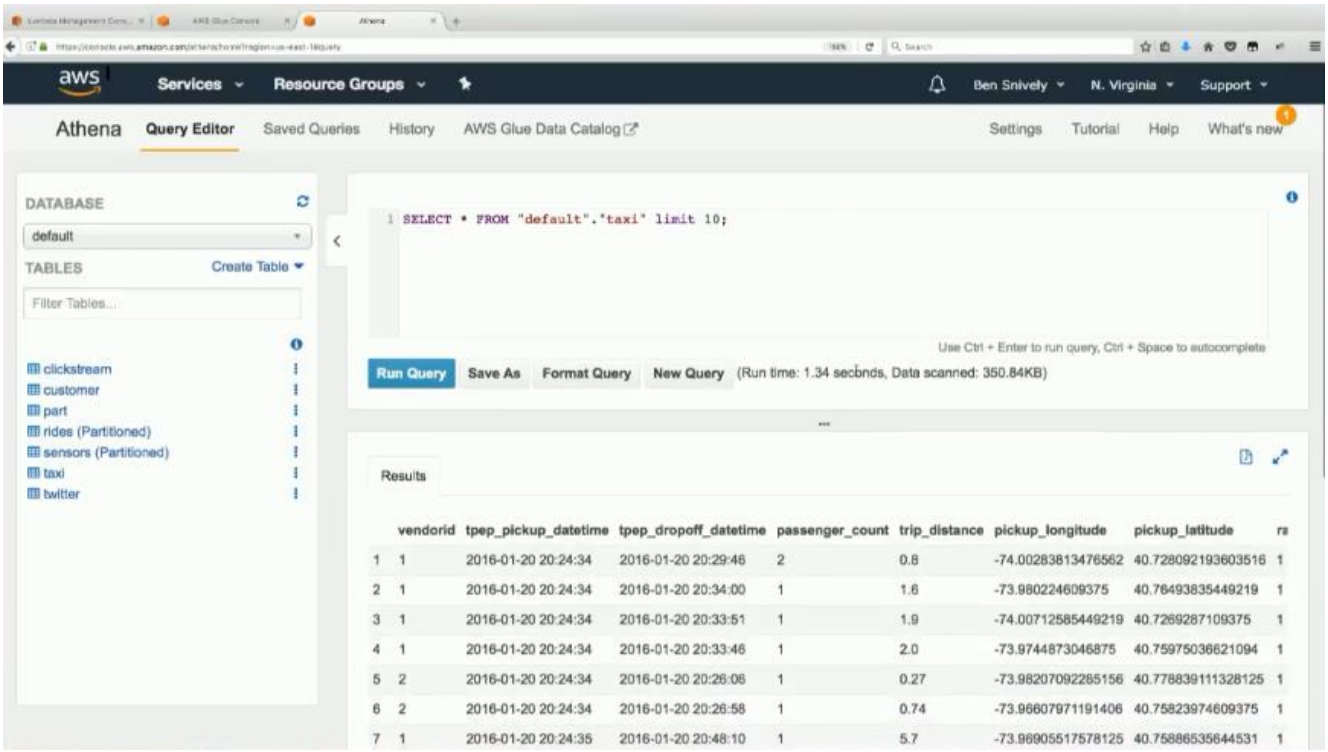

Now we can use Athena to start querying the data using SQL to analyze this dataset using SELECT commands.

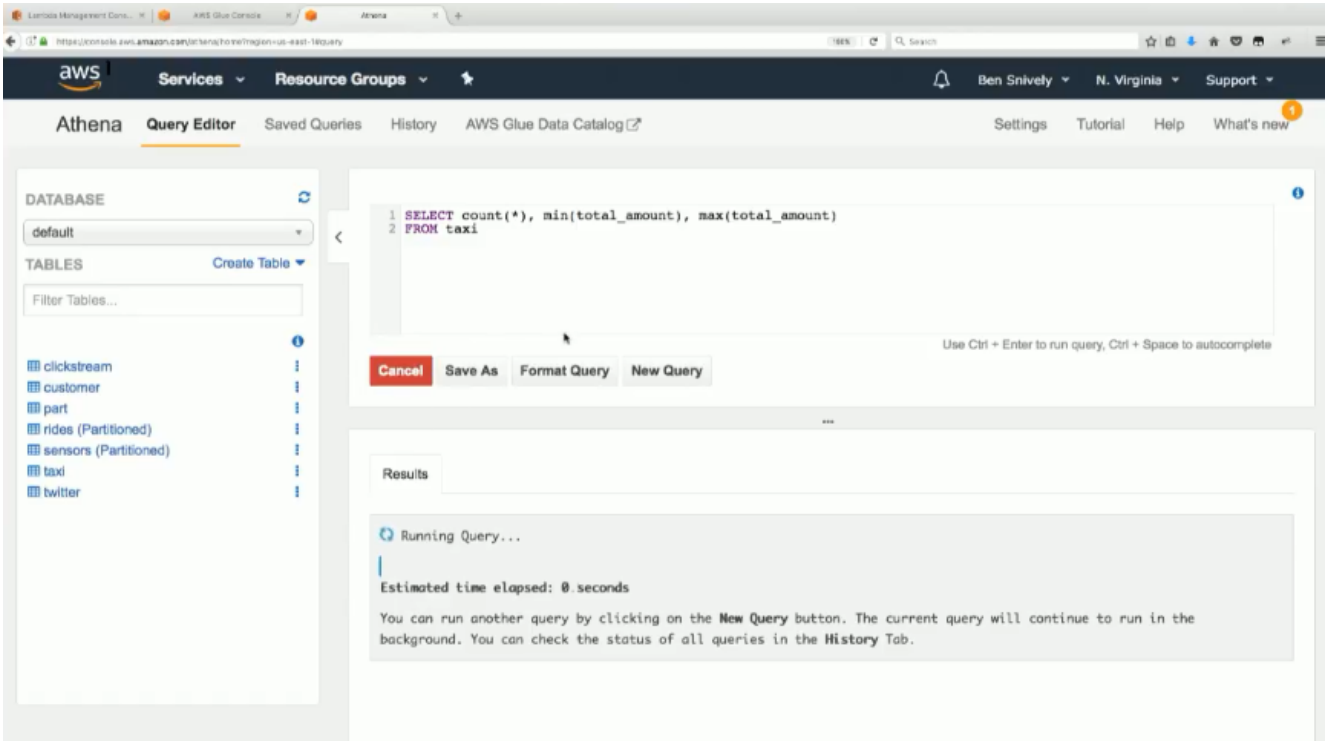

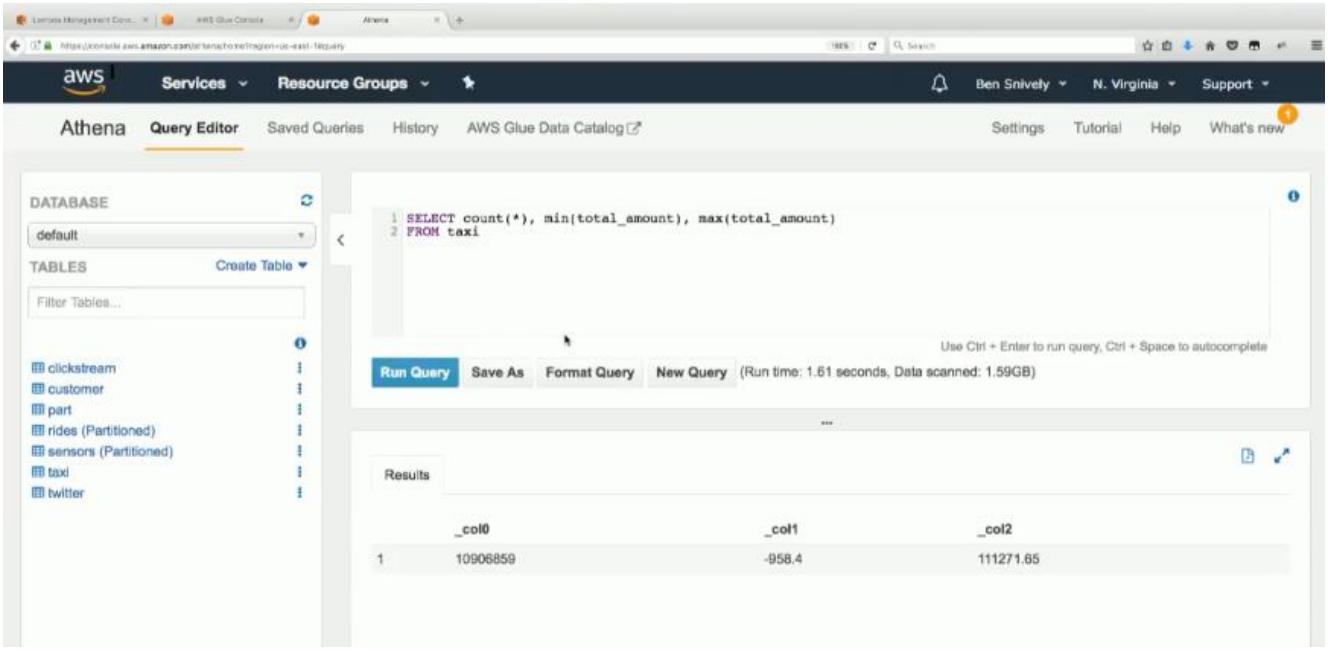

We can see that we have over 10 million records available for this month

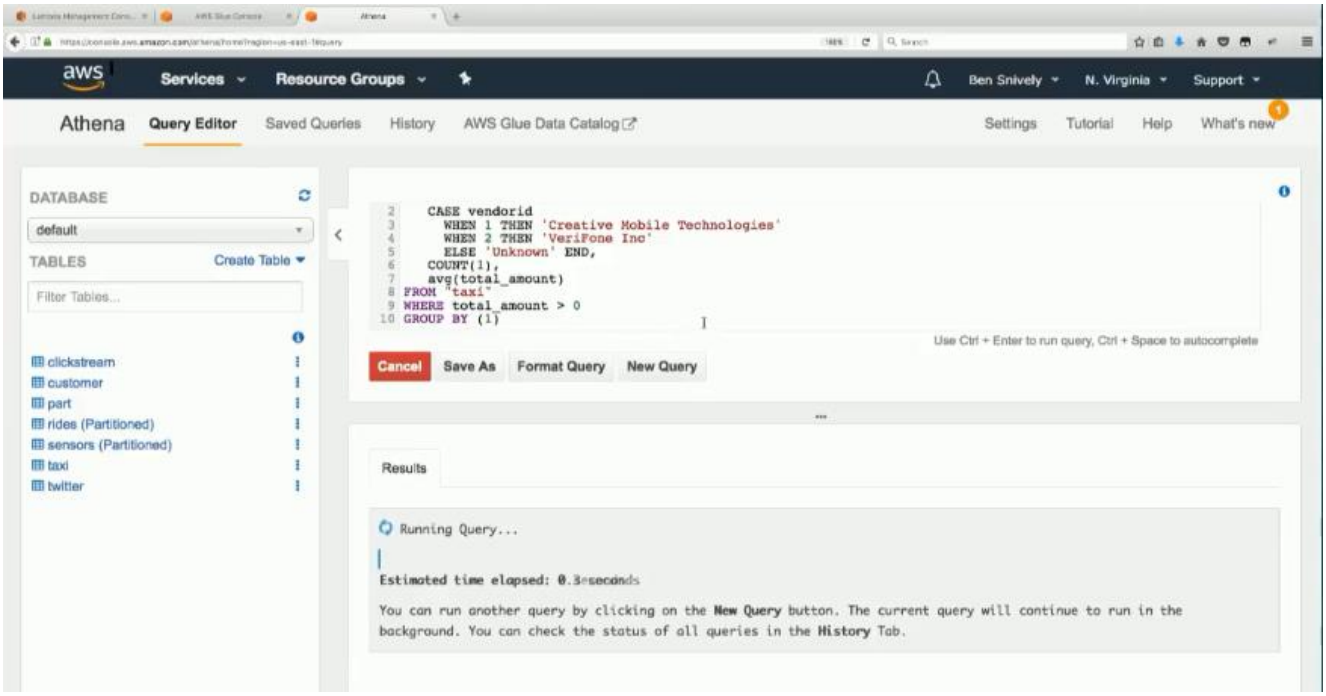

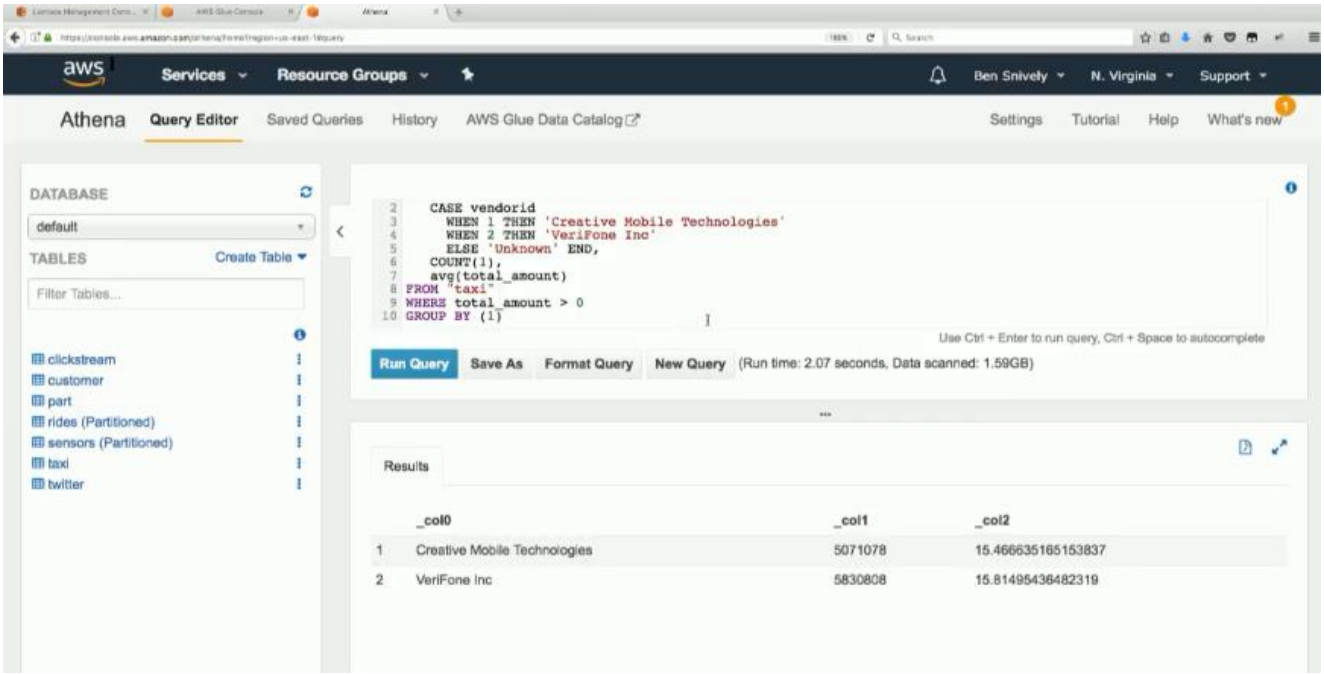

The results come back very fast

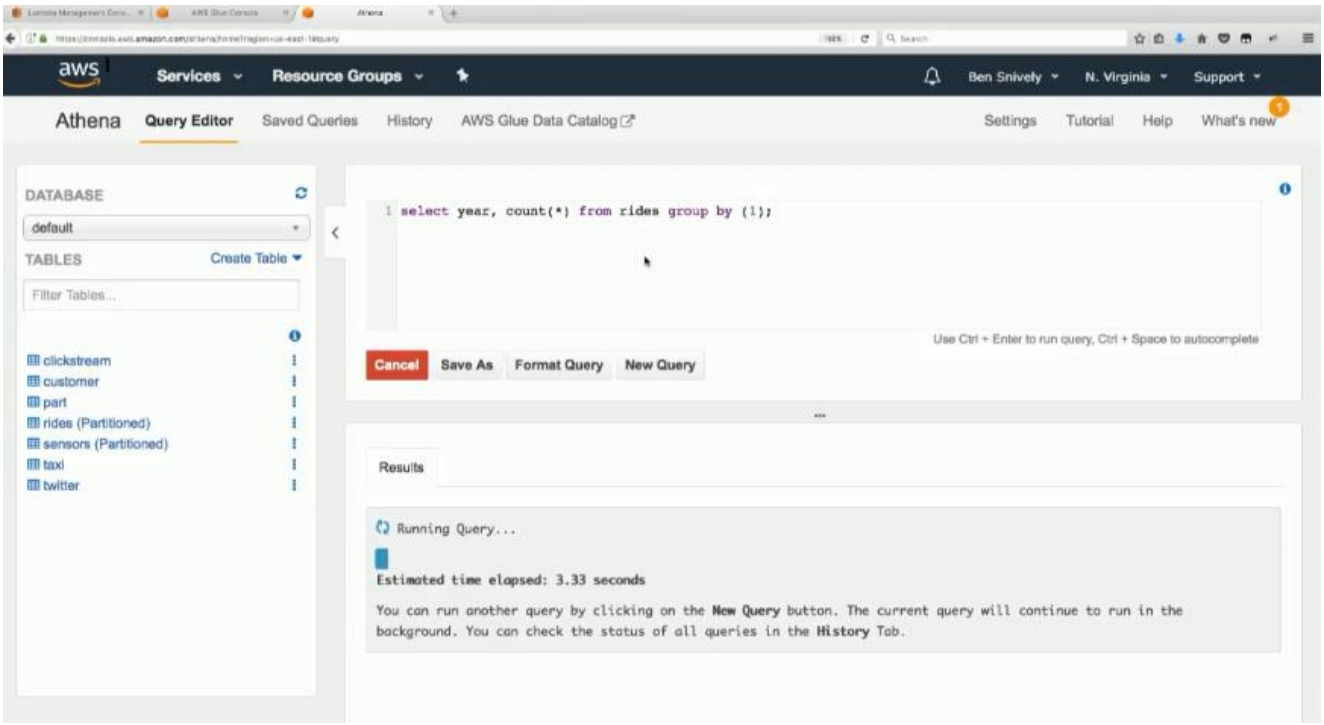

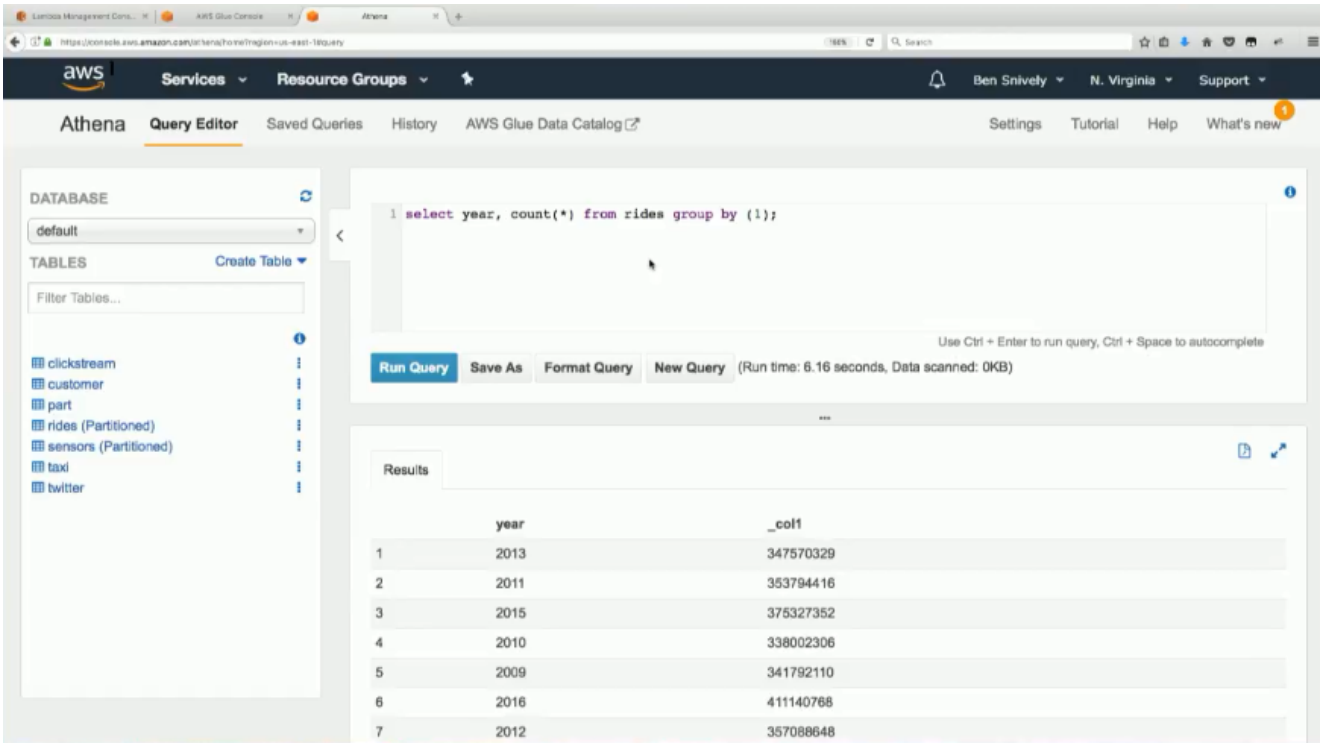

	year	_col1
1	2013	347570329
2	2011	353794416
3	2015	375327352
4	2010	338002306
5	2009	341792110
6	2016	411140768
7	2012	357088648

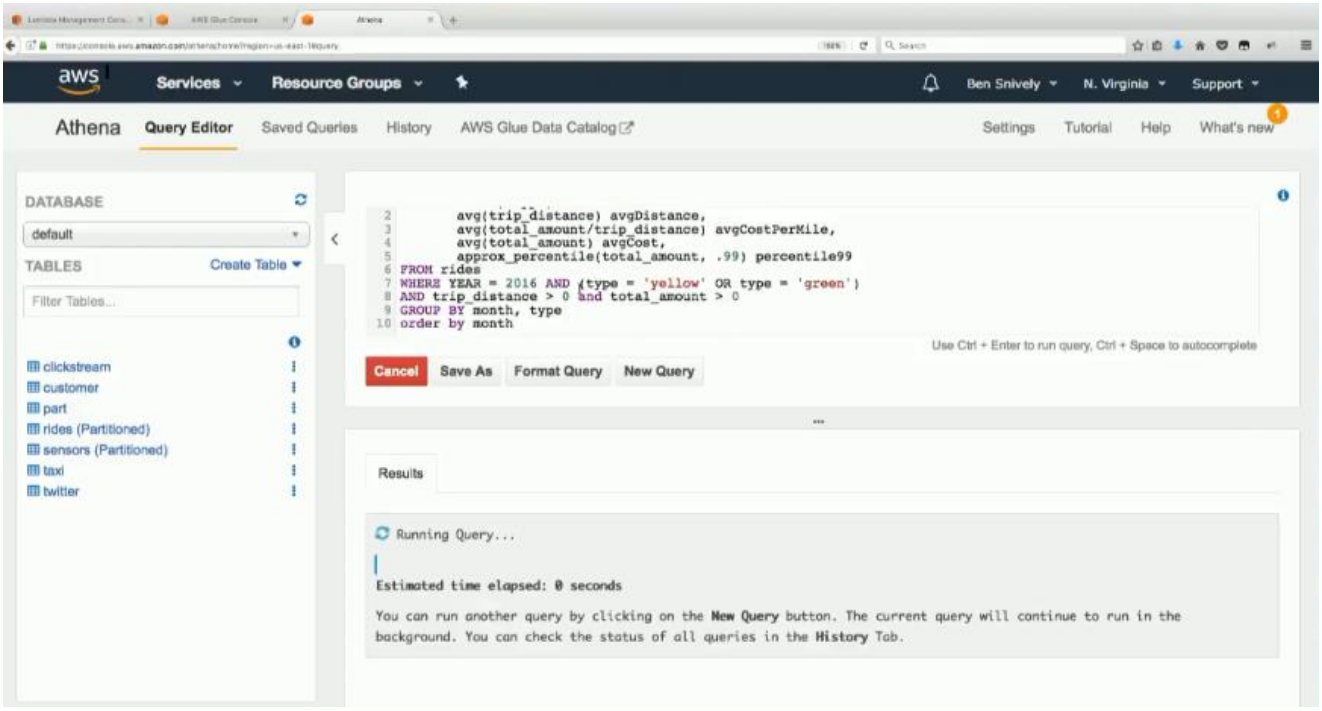

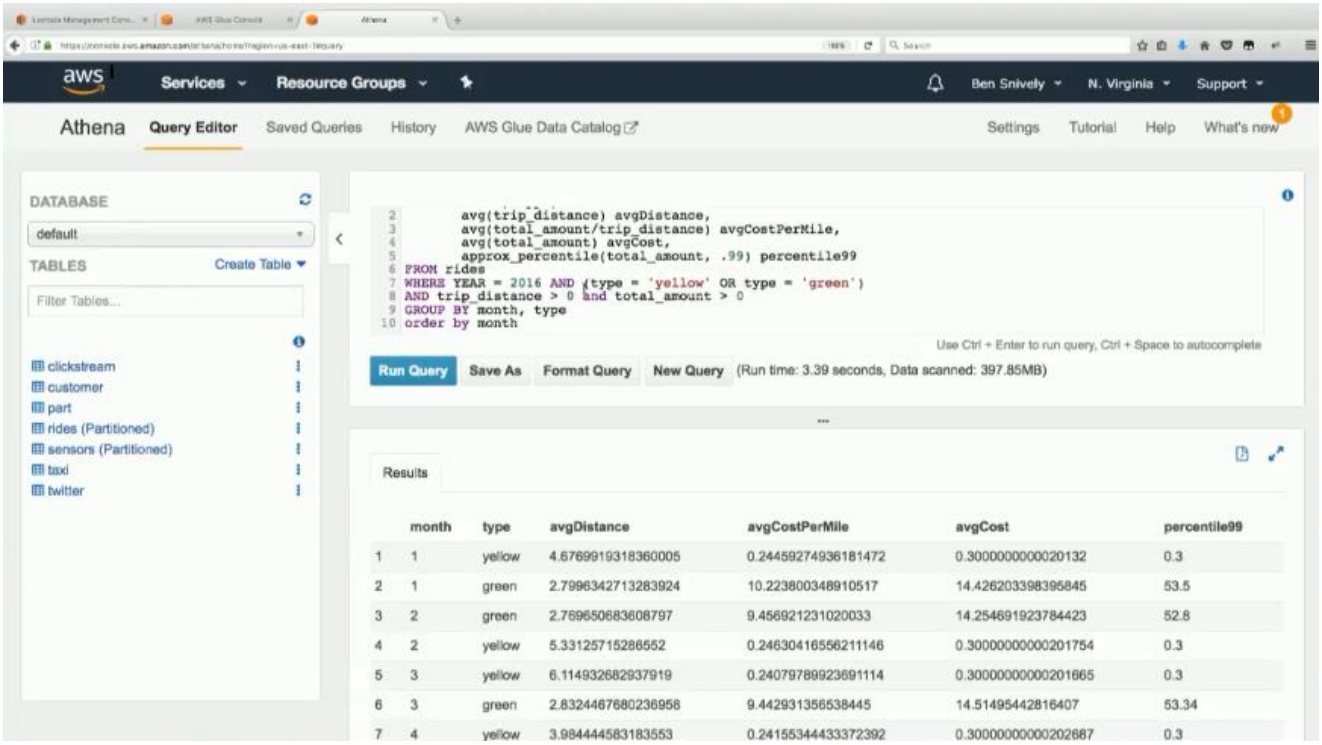

	month	type	avgDistance	avgCostPerMile	avgCost	percentile99
1	1	yellow	4.6769919318360005	0.24459274936181472	0.3000000000020132	0.3
2	1	green	2.7996342713283924	10.223800348910517	14.426203398395845	53.5
3	2	green	2.769650683608797	9.456921231020033	14.254691923784423	52.8
4	2	yellow	5.33125715286552	0.24630416556211146	0.30000000000201754	0.3
5	3	yellow	6.114932682937919	0.24079789923691114	0.30000000000201665	0.3
6	3	green	2.8324467680236958	9.442931356538445	14.51495442816407	53.34
7	4	yellow	3.984444583183553	0.24155344433372392	0.30000000000202687	0.3

Demonstration

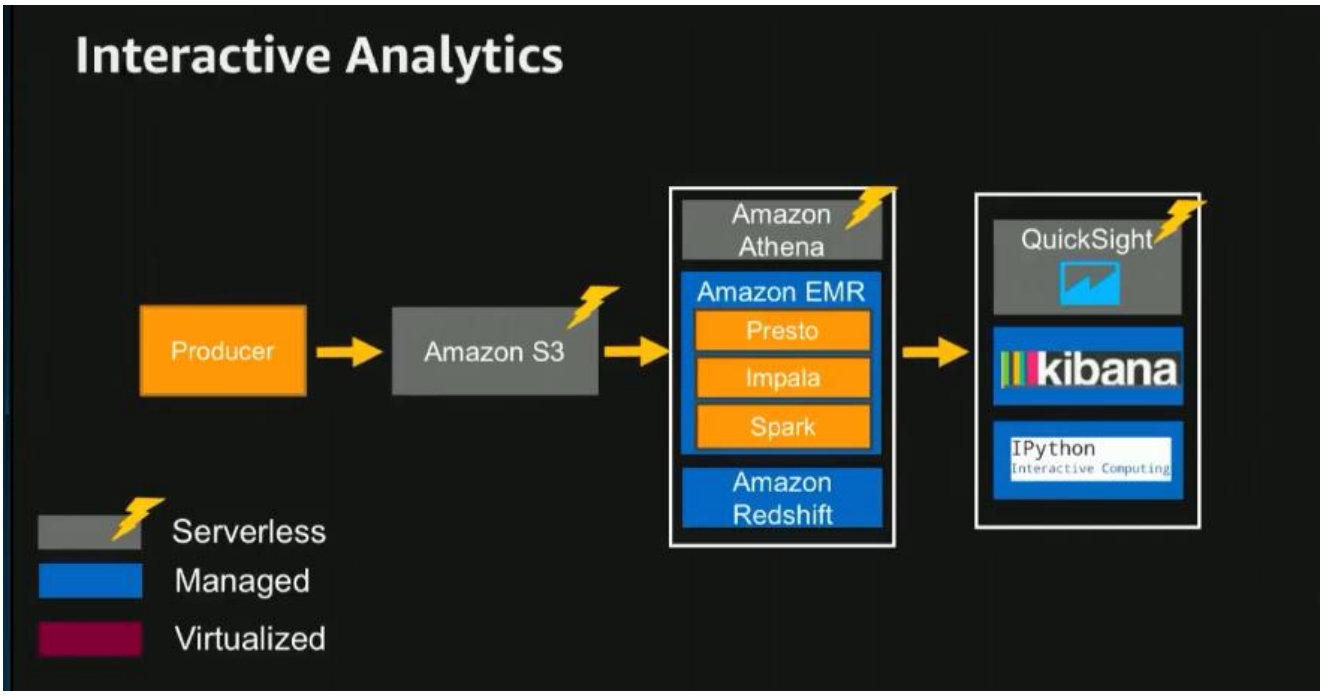
Interactive Analytics
Producer
Amazon S3
Amazon Athena
Amazon EMR
Presto
Impala
Spark
Amazon Redshift
QuickSight
kibana
IPython
Interactive Computing
Serverless
Managed
Virtualized

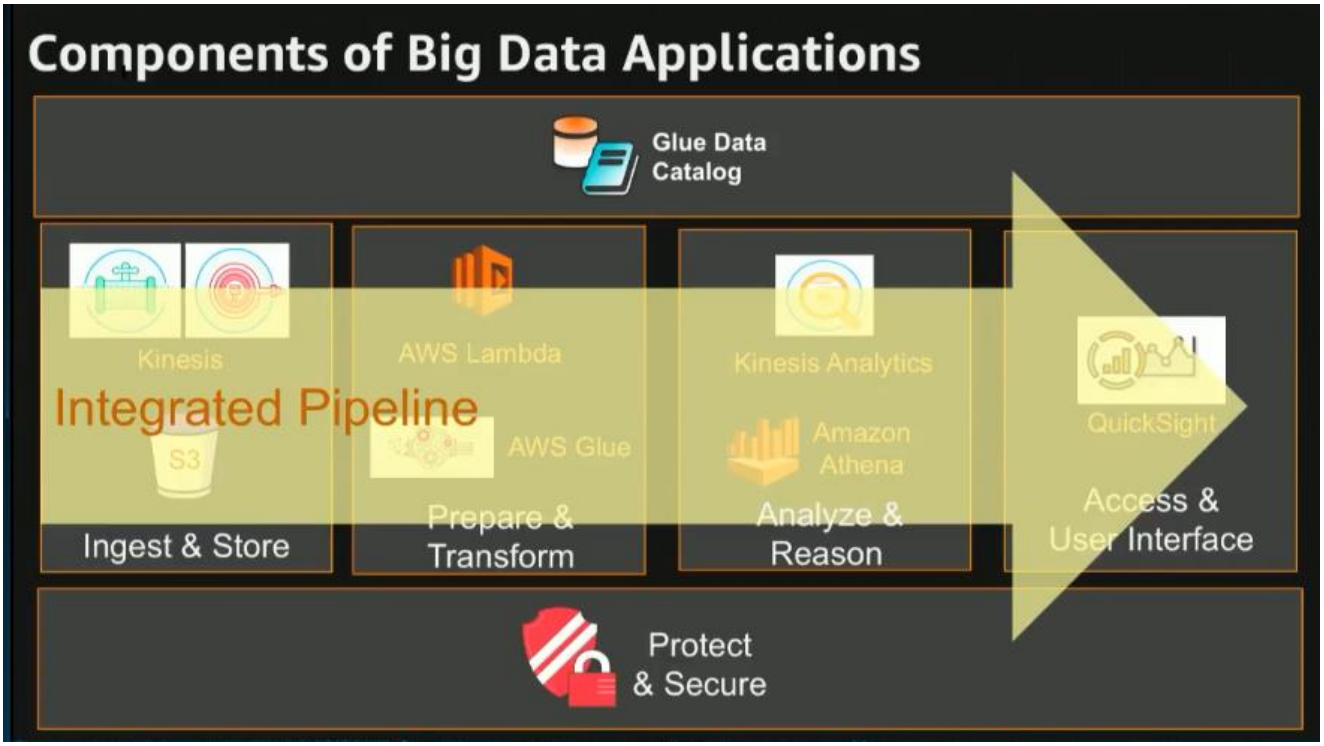
Components of Big Data Applications
Glue Data Catalog
Kinesis
AWS Lambda
Kinesis Analytics
QuickSight
Integrated Pipeline
S3
AWS Glue
Amazon Athena
Ingest & Store
Prepare & Transform
Analyze & Reason
Access & User Interface
Protect & Secure

Data Lake reference architecture

The Right tool for the Right Job

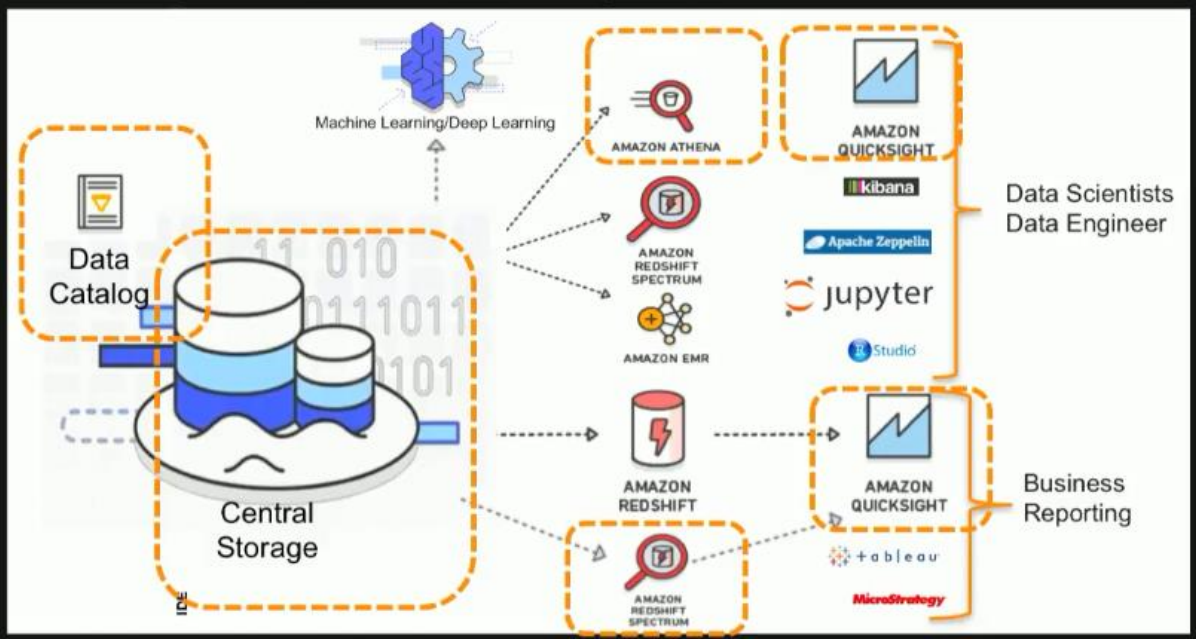

What about existing Hadoop Clusters?

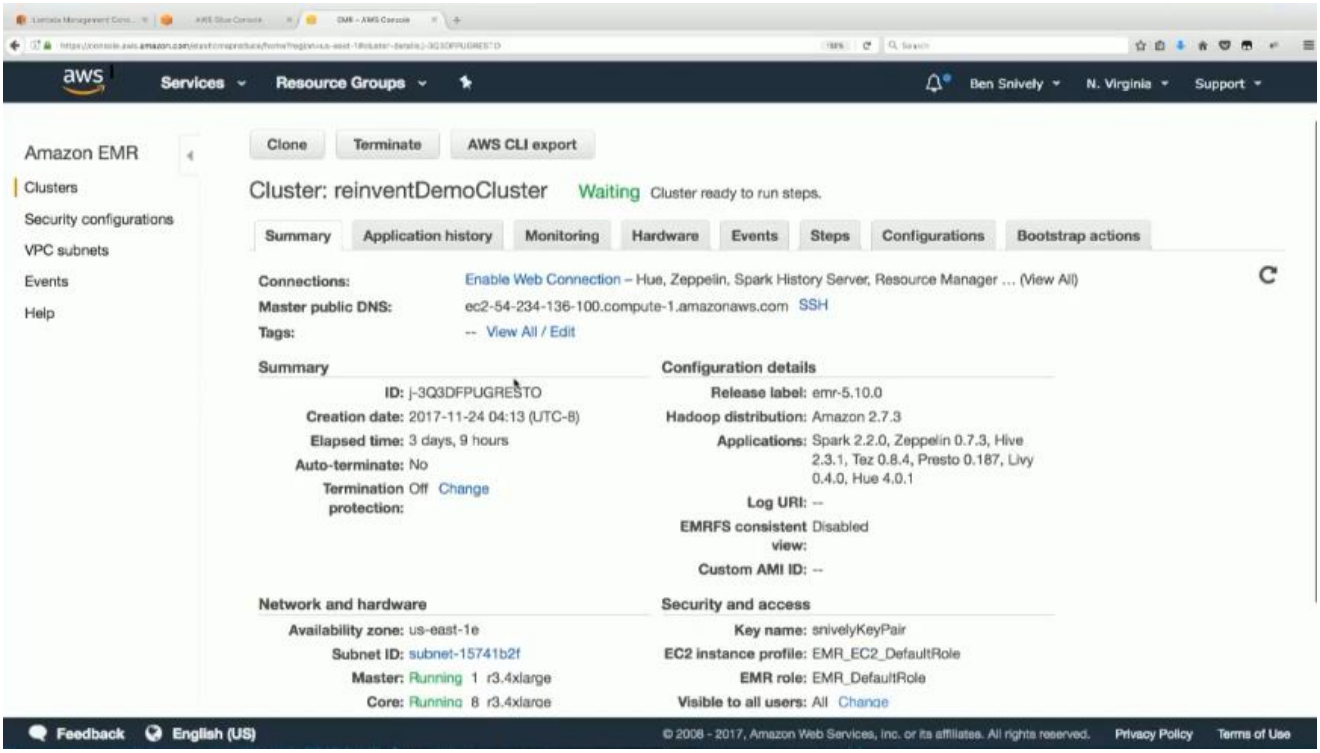

We have our EMR cluster running here, it has several of the ecosystem applications on it, we also specified that it use the Glue data catalog.

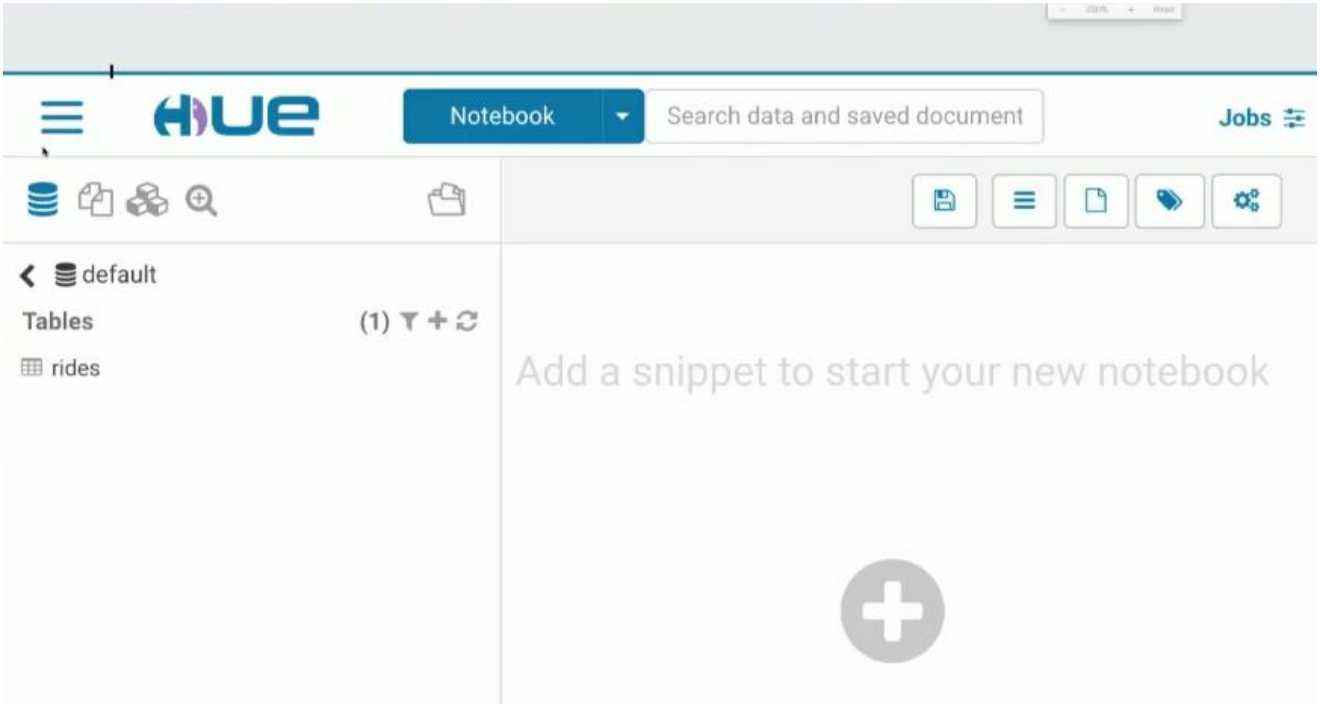

Hue is an interactive web interface for interacting with your Hadoop cluster.

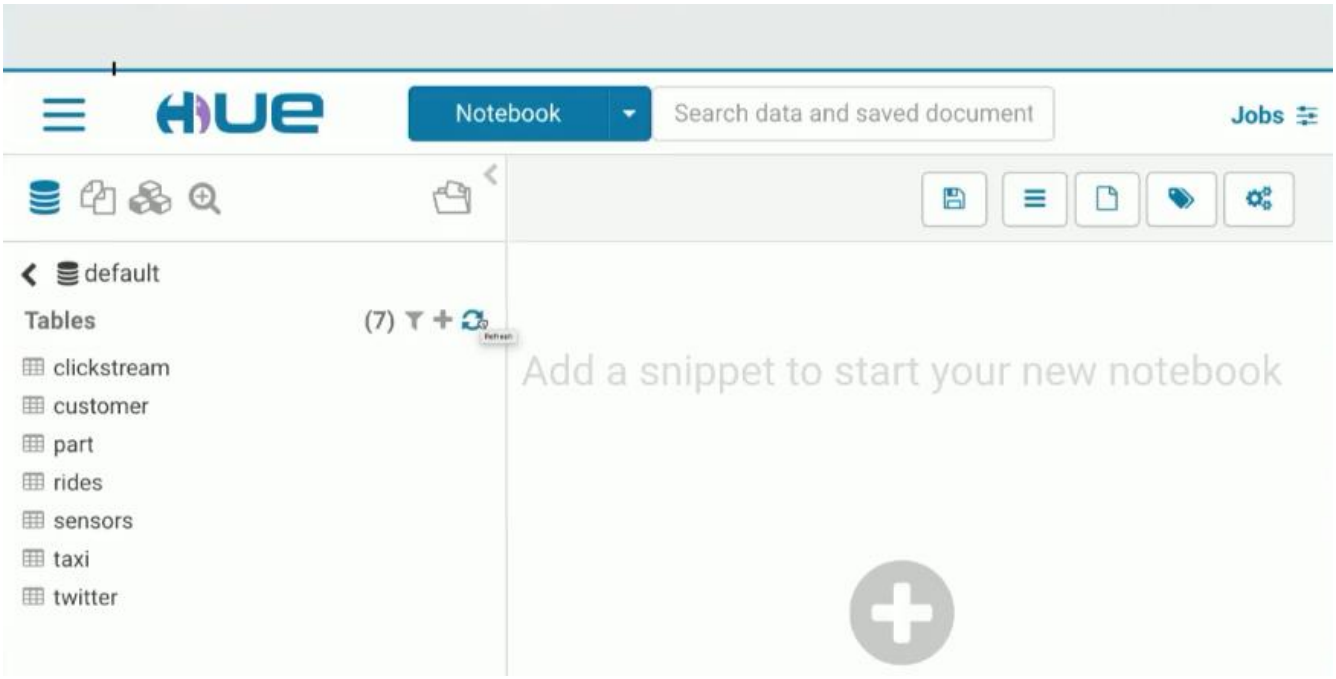

These are the data tables we discovered with the Glue crawler

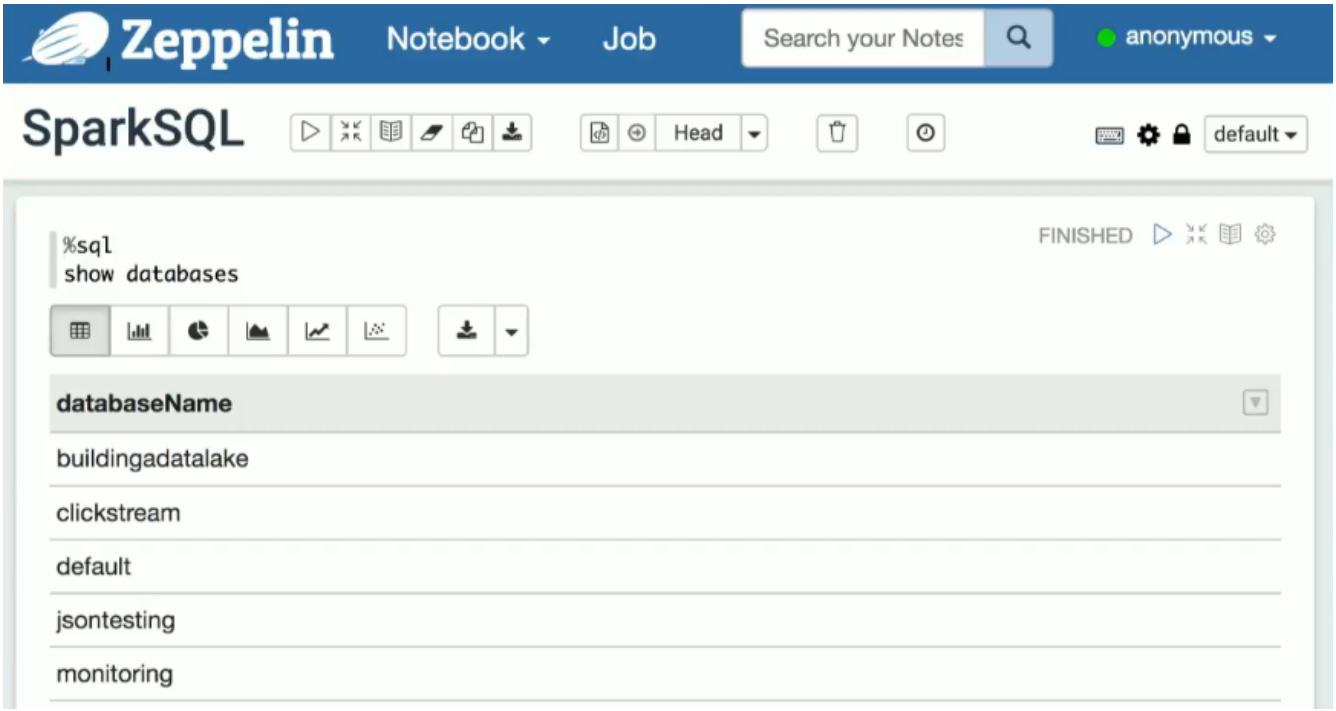

We can also have our data scientist using Zeppelin notebooks running on the EMR cluster look at the data

Zeppelin Notebook Job Search your Notes anonymous

SparkSQL Head default

Took 0 sec. Last updated by anonymous at November 27 2017, 11:26:03 AM.

Run this paragraph (Shift+Enter)

FINISHED

```
%sql
show tables
```

database	tableName	isTemporary
default	clickstream	false
default	customer	false
default	part	false
default	rides	false
default	sensors	false

What about existing Hadoop Clusters?

The Right tool for the Right Job

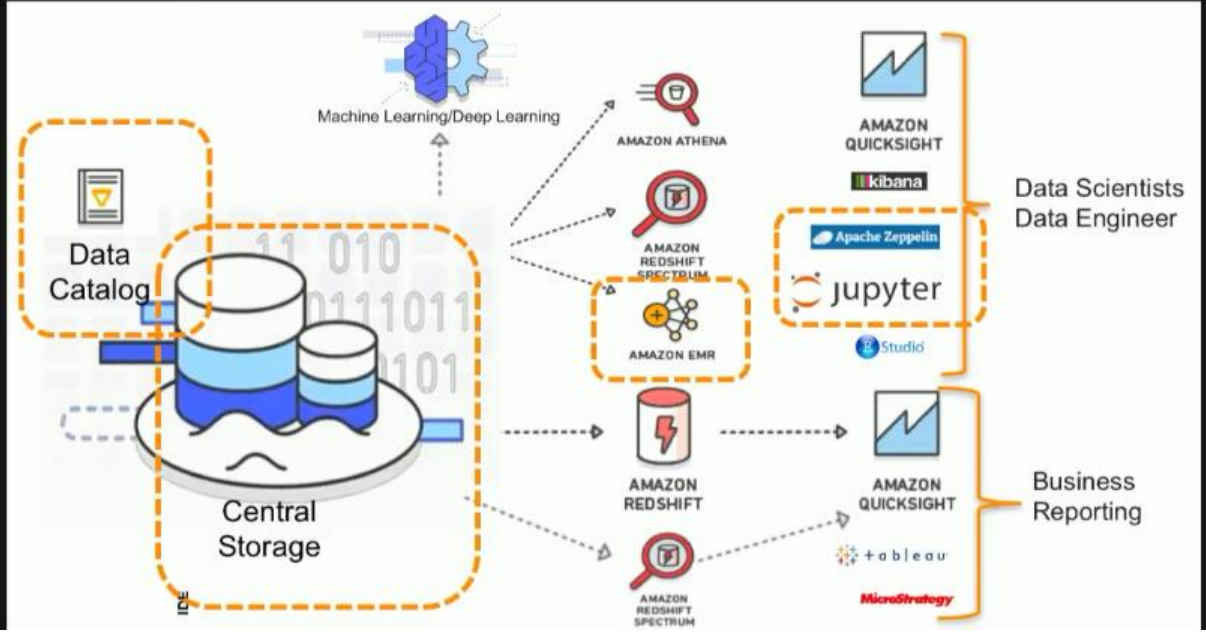

Serverless nicely fits into big data platforms

- Mix and Match Serverless, Managed, and Virtualized Services
- Leverage Services to easily
  - Rapidly ingest, categorize, and discover your data
  - Allow easy query and analysis of your data
  - Transform and Load data
  - Provide custom event based handlers
- Serverless allows you to focuses more analytics and not on infrastructure or servers
- Pay only for what you use