

AWS re:Invent

Best Practices for Orchestrating AWS Lambda Workloads

Matthew Williams
Evangelist @ Datadog

AWS re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



Technovangelist

I'm Matt Williams and I am an Evangelist at Datadog and an organizer for DevOps Days Boston. This site is a place where I write about the stuff that interests me.

296 Posts

November 01, 2017

CSV Lookup with Typinator for Working with Repetitive Forms

Datadog just had it's third customer summit this week in Austin Texas. I gave two workshops there to help customers get up to speed on Datadog, and...

October 28, 2017

Deciding I Need To Use AWS CodeBuild

If you at least partially live in the Amazon Web Services (AWS) world, then you know that re:INVENT is coming up in just a few weeks. That's the big...

July 31, 2017

Moved Technovangelist to a New Site

This weekend I built a new site. Why? Not sure. I have lots of reasons, but none are so strong that they are the main reason to do it. I moved from...

December 18, 2016

Reading The Next 93 - A goal for 2017

Books that my wife and I have checked out right now that I am eager to read (.pt) Books we own and I want to reread (.ps) Books that we own but L...

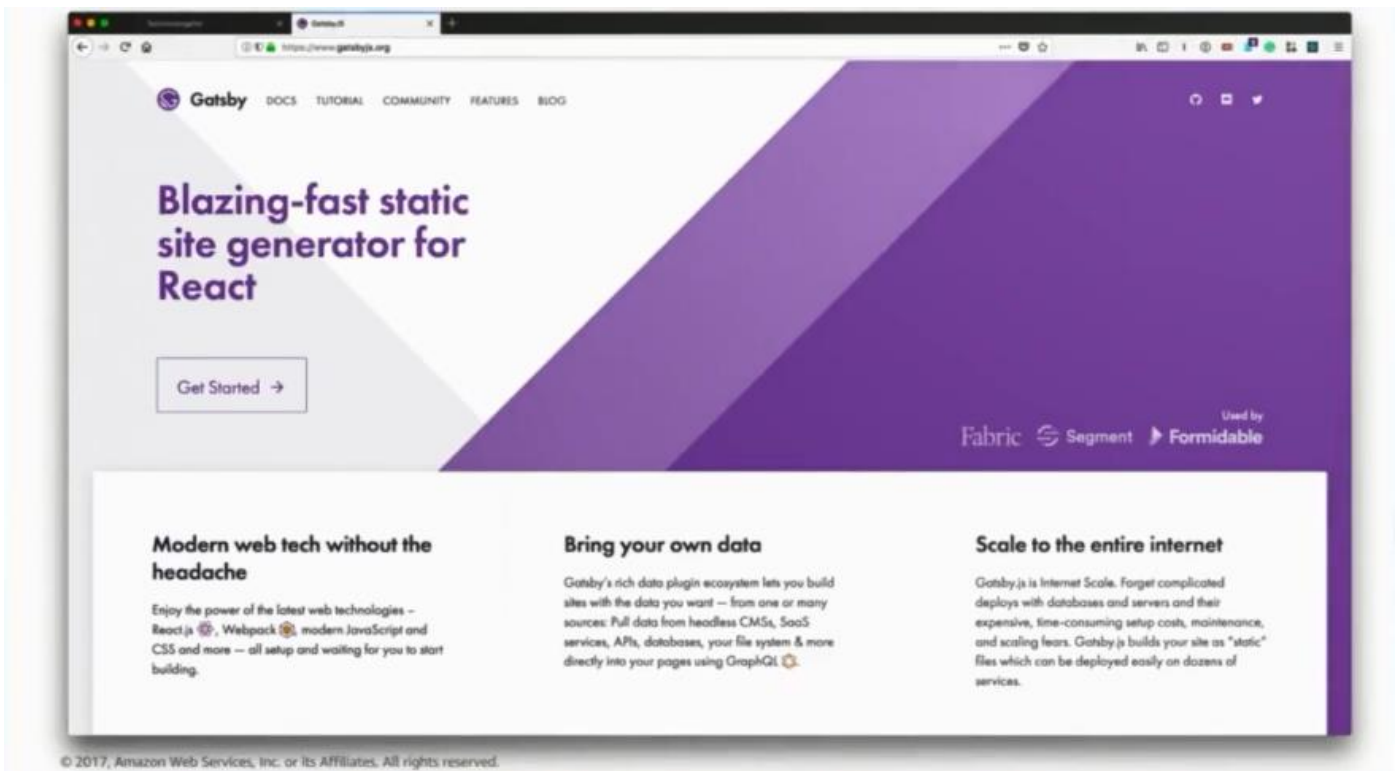
November 15, 2015

Podcasts You Need To Listen To: 99 Percent Invisible

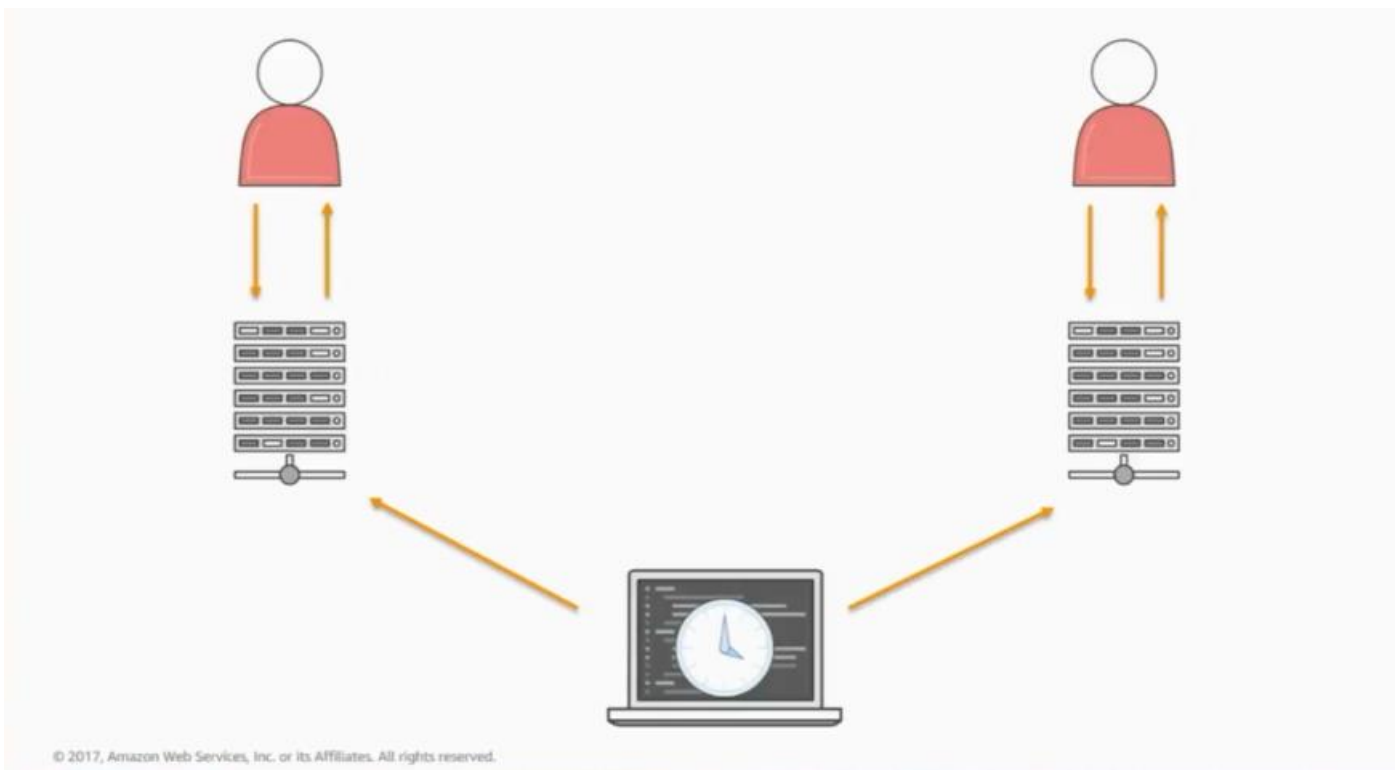
January 28, 2015

Go build and install on OSX in SublimeText3


© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



Now using GatsbyJS hosted on S3 for building out my blog that converts our JS, CSS, HTML assets into a static site



We don't need to keep a fixed EC2 instance when



<http://sf.technovangelist.com>

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

We are going to be talking about 2 step functions that are used to build our static sites

Who is Datadog

SaaS-based Monitoring & Analytics
Infrastructure, APM, Logs
Open Source Agent
Trillions of data
points per day

Booth 109 in Aria
Booth 1021 in Venetian
We are hiring!!

AWS re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



aws

Who is Datadog

SaaS-based Monitoring & Analytics

Infrastructure, APM, Logs

Open Source Agent

Trillions of data
points per day

Booth 109 in Aria

Booth 1021 in Venetian

We are hiring!!

**AWS
re:Invent**

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



Last year I talked about AWS Lambda...






...but AWS Lambda != Amazon EC2

...but AWS Lambda != application

...applications are complicated

...Lambdas should respect the single responsibility principle...



Other things learned over the last year
Serverless is about automation

Serverless is about event-driven computing

Scalability is different

we can create many serverless functions

but how do we orchestrate them?

Event/Message/Database as Traffic Cop



Amazon Kinesis




Amazon SQS



Amazon DynamoDB

The concept of the Traffic Cop is something that a lambda can report back to after it does its work. Then the next lambda can look into the Traffic Cop to determine whether it should do its work or not.

<https://github.com/Nordstrom/hello-retail>



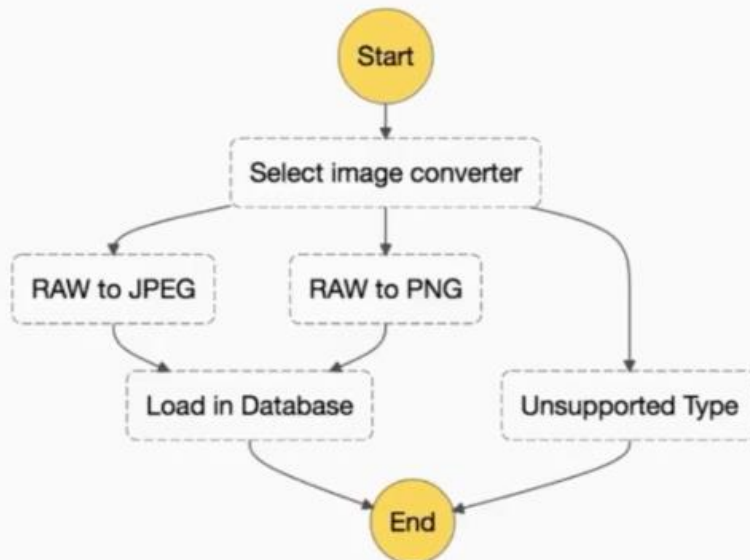
AWS re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



This is a very good example of the Traffic Cop concept. You can try this out using the GitHub code and link above

Enter AWS Step Functions



AWS re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



Now we can orchestrate several lambdas that each do their individual work and then hand off the result to another lambda function using Step Functions

a step function is a state machine

but what about SWF?

(Amazon Simple Workflow Service)

AWS Step Function Concepts

- State Machine
 - Transitions
 - Executions
 - States
 - Task
 - Choice
 - Succeed/Fail
 - Pass
 - Wait
 - Parallel
- You define the State Machine with the Amazon States Language
- A diagram will be generated based on that structure

AWS Step Function States: Task


Resource *	Type *
ResultPath	Next
Retry	End
Catch	Comment
TimeoutSeconds	InputPath
HeartbeatSeconds	OutputPath

A task can be a Lambda or an 'activity'

The Type* is the task and the Resource* is the ARN for that lambda or activity.

State Example: Task

```
1 {  
2   "Comment": "A simple task example",  
3   "StartAt": "The Task",  
4   "States": {  
5     "The Task": {  
6       "Type": "Task",  
7       "Resource": "arn:aws:lambda:REGION:ACCOUNT_ID:function:AwesomeTask",  
8       "End": true  
9     }  
10  }  
11 }
```



AWS Step Function States: Choice

Choices

Default

Type

Next

End


Comment

InputPath

OutputPath

State Example: Choice

```
10 "Choice": {  
11   "Type": "Choice",  
12   "Choices": [  
13     {  
14       "Variable": "$.foo",  
15       "NumericEquals": 1,  
16       "Next": "FirstTask"  
17     },  
18     {  
19       "Variable": "$.foo",  
20       "NumericEquals": 2,  
21       "Next": "SecondTask"  
22     }  
23   ]  
24 },
```




AWS Step Function States: Succeed / Fail

Cause
Error

Type
Next
End
Comment
InputPath
OutputPath

State Example: Succeed / Fail

```
7  "Choices": [  
8    {  
9      "Variable": "$.foo",  
10     "NumericGreaterThan": 0,  
11     "Next": "Positive"  
12   }  
13 ],  
14 "Default": "Negative"  
15 },  
16  
17 "Positive": {  
18   "Type": "Succeed"  
19 },  
20  
21 "Negative": {  
22   "Type": "Fail",  
23   "Error": "DefaultStateError",  
24   "Cause": "Its not positive!"  
25 }
```



AWS re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.




AWS Step Function States: Pass

Result
ResultPath

Type
Next
End
Comment
InputPath
OutputPath

State Example: Pass

```
1 {  
2   "Comment": "A simple pass",  
3   "StartAt": "FirstTask",  
4   "States": {  
5     "FirstTask": {  
6       "Type": "Task",  
7       "Resource": "arn:aws:lambda:REGION:ACCOUNT_ID:function:AwesomeTask",  
8       "Next": "Pass"  
9     },  
10    "Pass": {  
11      "Type": "Pass",  
12      "Result": "Wow, this was exciting",  
13      "ResultPath": "ExtraDetail",  
14      "End": true  
15    }  
16  }  
17 }  
18
```



AWS re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



AWS Step Function States: Wait

Seconds

Timestamp

SecondsPath

TimestampPath

Type

Next

End

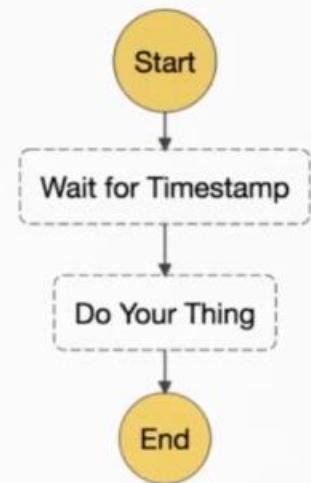
Comment

InputPath

OutputPath

State Example: Wait

```
1 {  
2   "Comment": "A simple wait example",  
3   "StartAt": "Wait for Timestamp",  
4   "States": {  
5     "Wait for Timestamp": {  
6       "Type": "Wait",  
7       "TimestampPath": "$.trigger_date",  
8       "Next": "Do Your Thing"  
9     },  
10    "Do Your Thing": {  
11      "Type": "Task",  
12      "Resource": "arn:aws:lambda:REGION:ACCOUNT_ID:function:DoTheThing",  
13      "End": true  
14    }  
15  }  
16 }
```



AWS re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



AWS Step Function States: Parallel

Branches

ResultPath

Retry

Catch

Type

Next

End


Comment

InputPath

OutputPath

State Example: Parallel

```
5  "Parallel": {
6    "Type": "Parallel",
7    "Next": "Final State",
8    "Branches": [
9      {
10       "StartAt": "Left Brain",
11       "States": {
12         "Left Brain": {
13           "Type": "Wait",
14           "Seconds": 2,
15           "Next": "Sequence"
16         },
17         "Sequence": {
18           "Type": "Task",
19           "Resource": "arn:aws:lambda:REGION:ACCOUNT_ID:function:sequence",
20           "End": true
21         }
22       }
23     },
24     {
25       "StartAt": "Right Brain",
26       "States": {
27         "Right Brain": {
28           "Type": "Pass",
29           "Next": "Visualize"
```



AWS re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



These are the basics of using step functions

Many ways to build Step Functions

I used Serverless Framework

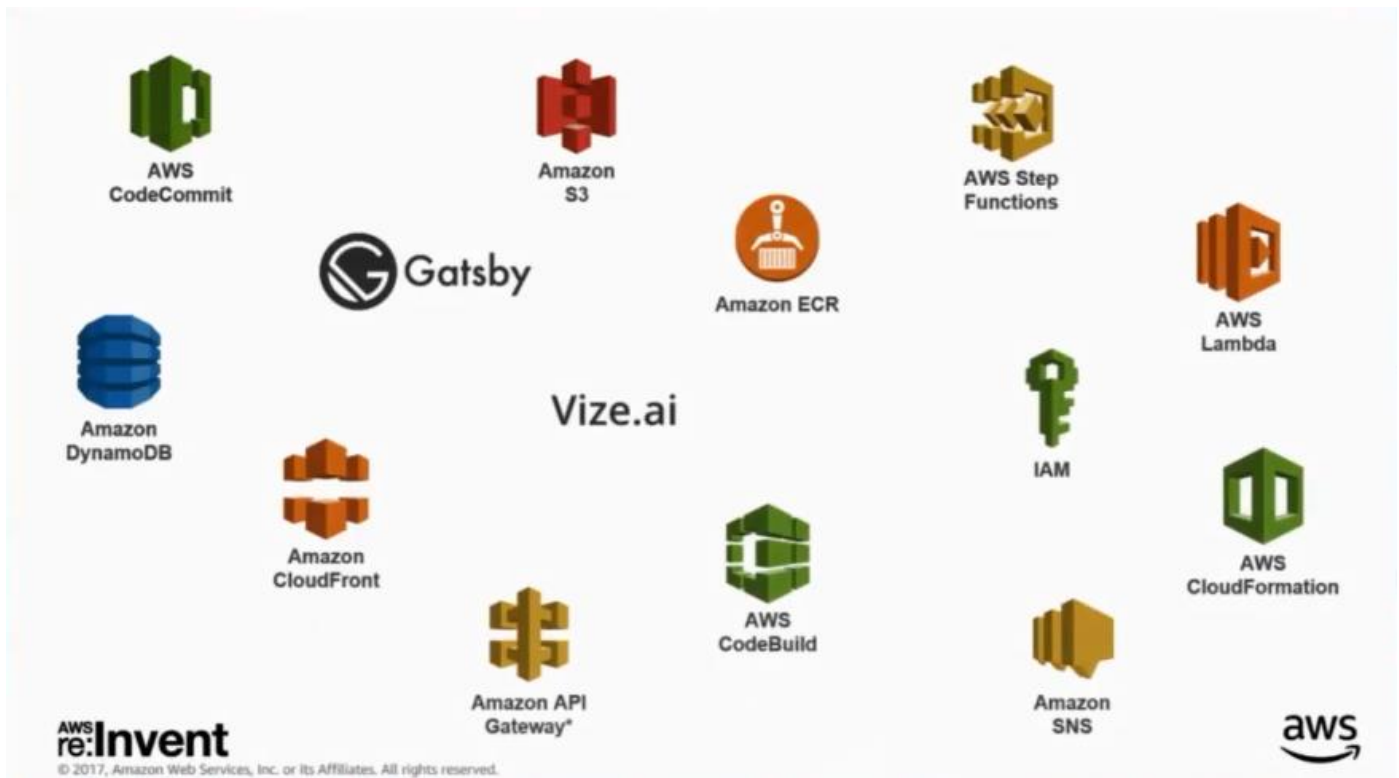


The **Serverless framework** has an add-on called **Serverless Step Functions** that make it really easy to build step functions and their integrations in a single file

Benefits of Serverless Framework

Define everything in a single file:

- IAM Role Statements
- Lambda Function Handlers
- Lambda Function Triggers
- Step Function States
- CloudFormation for other things



serverless.yml

```
service:
  name: mattw-reinvent2017-build-website

plugins:
  - serverless-pseudo-parameters
  - serverless-step-functions
  - serverless-webpack
custom:
  repoName: reinvent2017-website
```

serverless.yml cont'd

```
provider:
  name: aws
  runtime: nodejs6.10
  stage: ${opt:stage, 'dev'}
  region: ${opt:region, 'us-east-1'}
  iamRoleStatements:
    - Effect: "Allow"
      Action:
        - "states:StartExecution"
      Resource:
        - ${self:resources.Outputs.MyStateMachine.Value}
    - Effect: "Allow"
      Action:
        - "codecommit:GetCommit"
        - "codecommit:GetDifferences"
      Resource:
        - arn:aws:codecommit:${self:provider.region}:${AWS::AccountId}:${...custom.repoName}
```

AWS
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



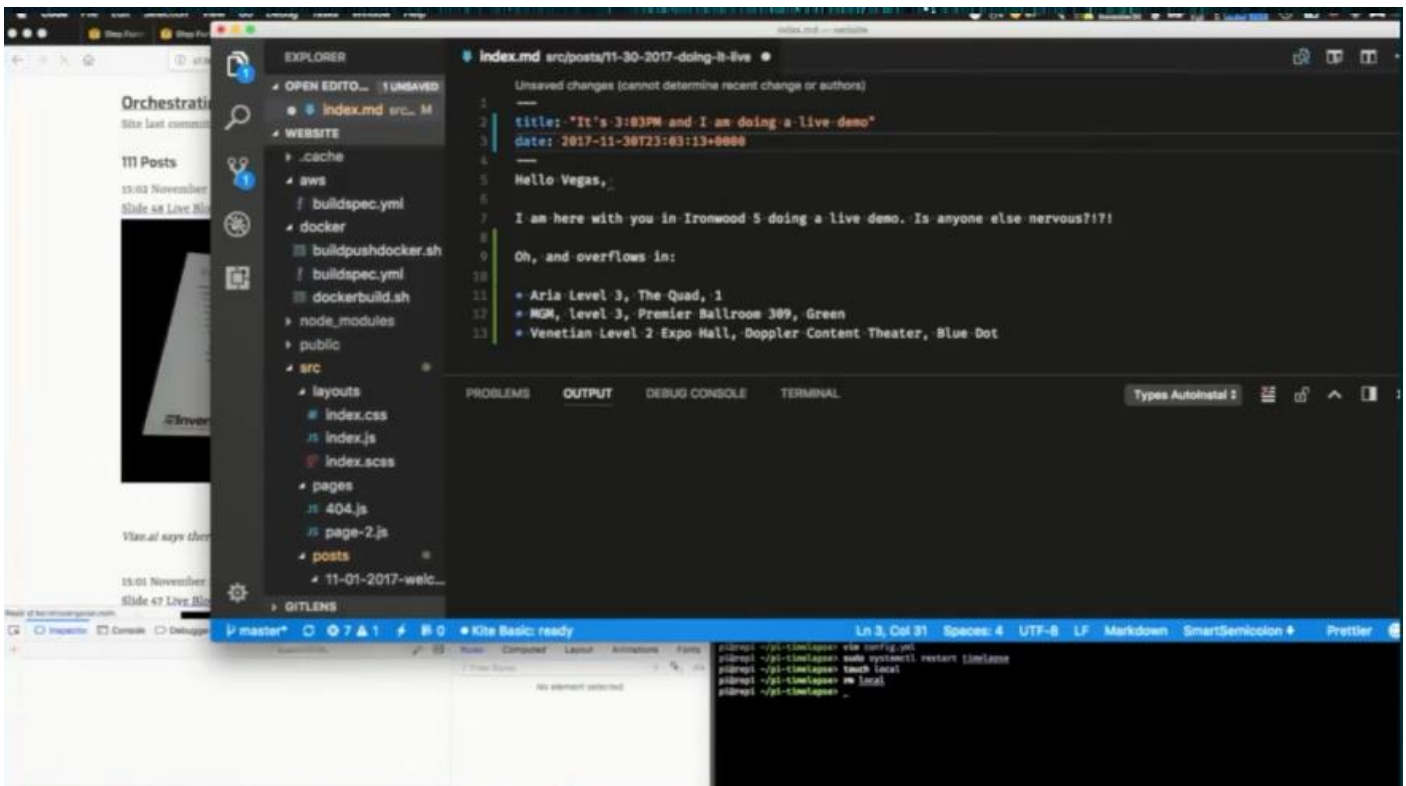
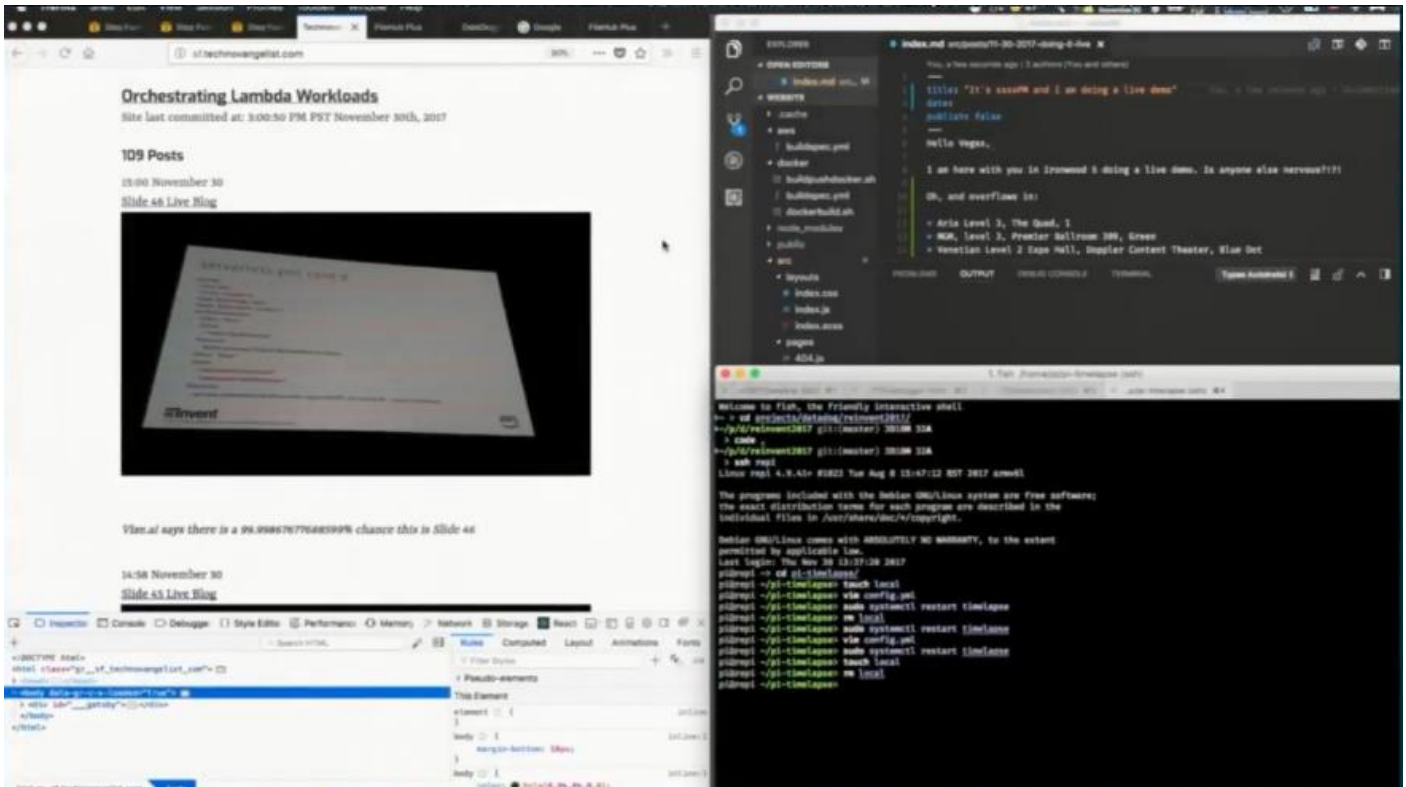
serverless.yml cont'd


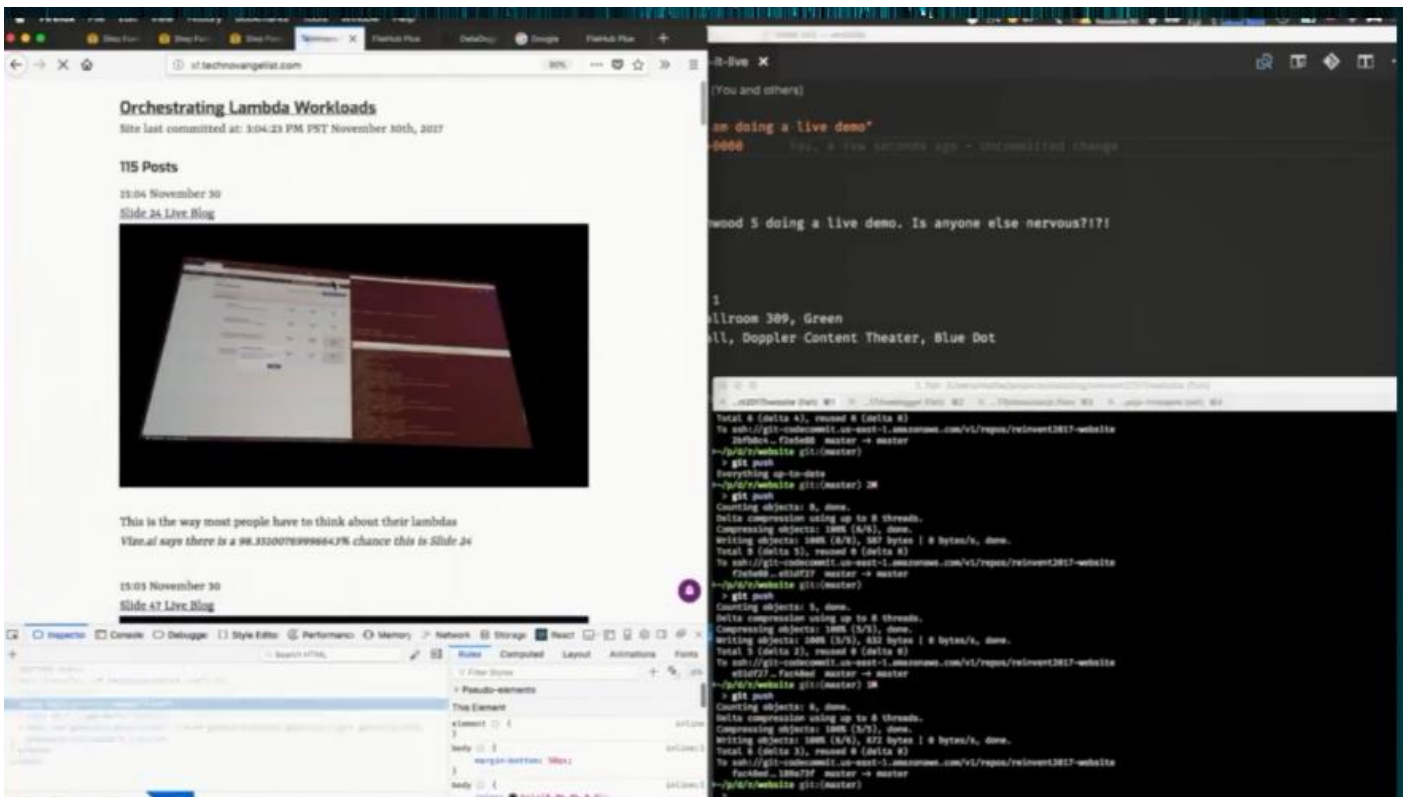
```
functions:
  buildSite:
    handler: handler.build
  whatChanged:
    handler: handler.whatChanged
  startStepFunction:
    handler: handler.startStepFunction
  events:
    - sns: arn:aws:sns:${...provider.region}:${AWS::AccountId}:mattw-reinvent-websitebuild
  environment:
    statemachine_arn: ${self:resources.Outputs.MyStateMachine.Value}
```

serverless.yml cont'd

```
stepFunctions:
  stateMachines:
    FirstStateMachine:
      name: myStateMachine
      definition:
        StartAt: WhatChanged
        States:
          WhatChanged:
            Type: Task
            Resource: arn:aws:lambda:${AWS::Region}:${AWS::AccountId}:
            function: ${self:service}-${self:provider.stage}-whatChanged
            Next: BuildDockerOrNot
```

Demo





When we committed code into **CodeCommit** it triggered off an event to an **SNS** topic, we already have a **lambda** waiting for that SNS topic to run

```

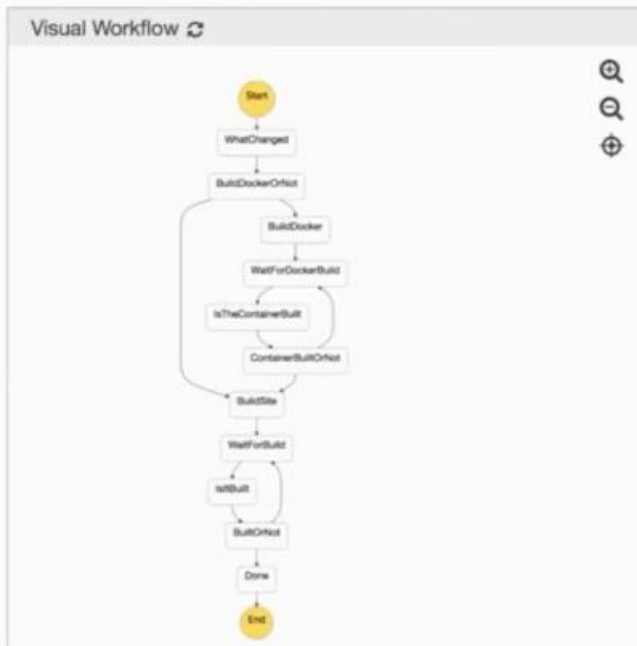
startStepFunction:
  handler: handler.startStepFunction
  events:
    - sns: arn:aws:sns:...mattw-reinvent-websitebuild
  environment:
    statemachine_arn: ${self:reso...puts.MyStateMachine.Value}

```

This is our definition for this **startStepFunction** step in our serverless.yml function. The handler called handler.startStepFunction is waiting for an SNS topic, we have also given it an environment variable

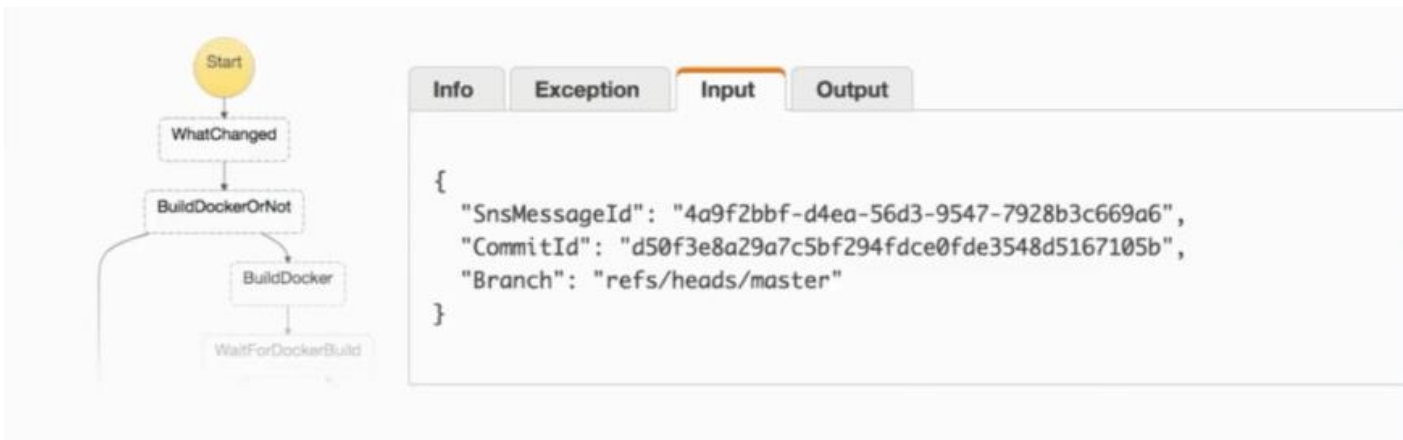
StartStepFunction
 Get Step Function ARN
 Get SNS details
 Start SF Execution

The **startStepFunction** basically gets the step function **arn** from **serverless**, then it gets the SNS details, and then starts a step function execution.

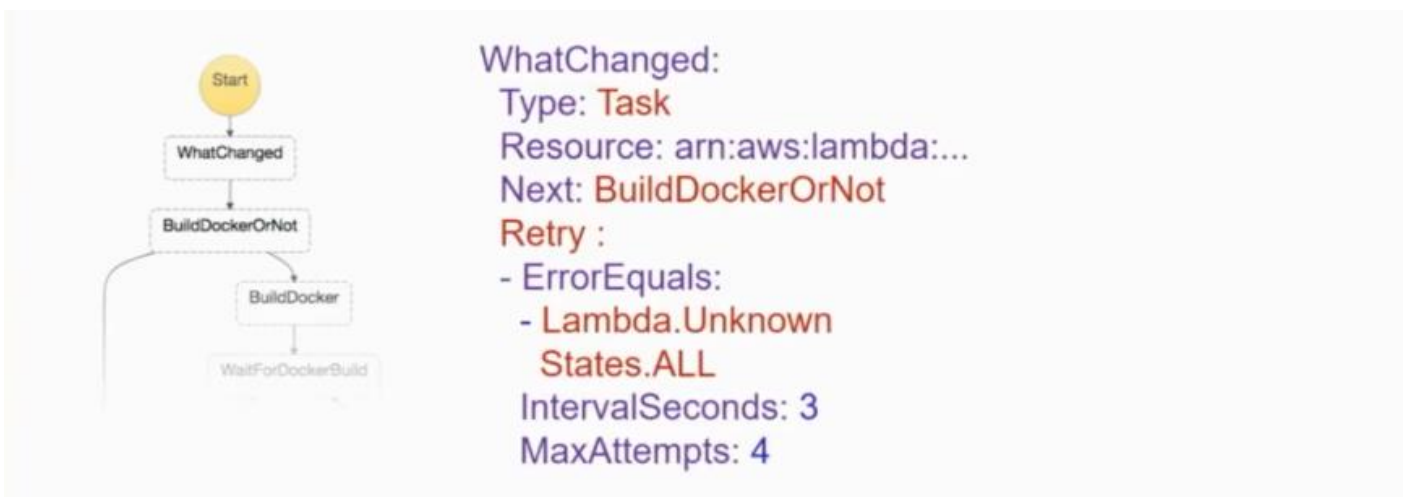


1. What changed in the repo?
2. If docker-related, build docker
 1. Tell CodeBuild to rebuild docker image
 2. Wait for it
 3. Push to Amazon ECS repo
3. Use the Docker image to build the website
4. Wait for it
5. Send Matt an SMS

This is the step function SF that gets started. The SF is asking what and doing the steps above. It will ask what changed, if it's a file related to Docker then it will build a new image using CodeBuild.



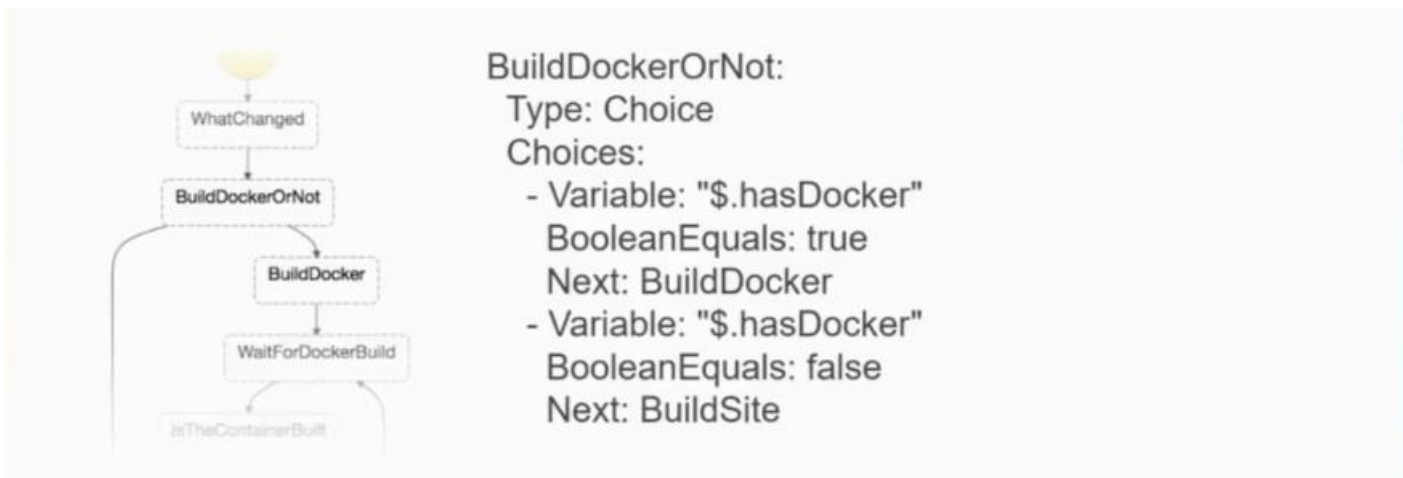
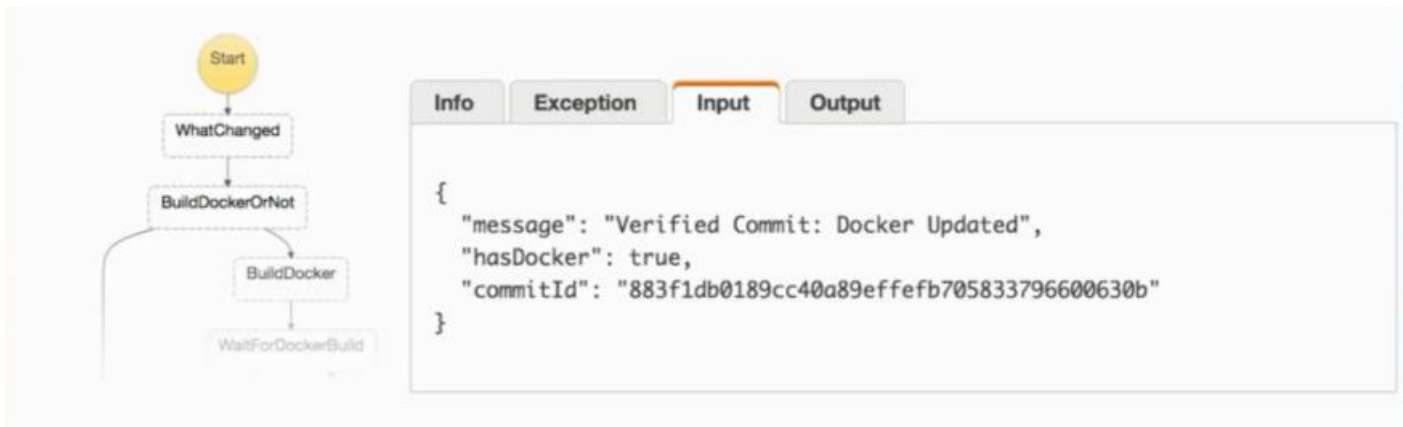
This is the type of interface that you use when working with step functions, you can do some debugging and see information related to any step/task/state. The Input tab shows you what JSON data is going into that particular task, and the Output tab shows you the output produced from that particular task.



This is the **whatChanged** task in the **serverless.yml** file, we also have 4 retries



This is the **pseudocode** for the **whatChanged** task



Note that we needed to add the needed IAM role permissions for this in the **serverless.yml** file. This will trigger a **CodeBuild** step as below

phases:

install:

commands:

- nohup /usr/local/bin/dockerd -G dockremap --host=unix:///...
- timeout -t 15 sh -c "until docker info; do echo .; sleep 1; done"

pre_build:

commands:

- echo Logging in to Amazon ECR...
- \$(aws ecr get-login --no-include-email --region us-east-1)

build:

commands:

- echo Build started on `date`
- echo Building the Docker image...
- docker build -t mattw-stepfunction-demo:latest .
- docker tag mattw-stepfunction-demo:latest ...dkr.ecr.us-east-1.amazona...


post_build:

commands:

- echo Build completed on `date`
- echo Pushing the Docker image...
- docker push ...dkr.ecr.us-east-1.amazonaws.com/mattw-stepfunc...

aws
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.




WaitForDockerBuild:

Type: Wait

Seconds: 20

Next: IsTheContainerBuilt




IsTheContainerBuilt:

Type: Task

Resource: arn:aws:lambda...


Next: ContainerBuiltOrNot




Info Exception **Input** Output

```

{
  "message": "The container has not been built",
  "containerBuildId": "buildDockerImage:9f9600d1-a028-406d-bb6f-7053f693!",
  "containerBuildDone": false,
  "commitId": "883f1db0189cc40a89effefb705833796600630b"
}
  
```




IsTheContainerBuilt
Get build info **from** CodeBuild
Check **if** its done



Info Exception **Input** **Output**

```

{
  "message": "The container has been built",
  "containerBuildId": "buildDockerImage:9f9600d1-a028-406d-bb6f-7053f693!",
  "containerBuildDone": true,
  "commitId": "883f1db0189cc40a89effefb705833796600630b"
}
  
```



ContainerBuiltOrNot:

Type: **Choice**

Choices:

- Variable: **"\$.containerBuildDone"**
BooleanEquals: **true**
Next: **BuildSite**
- Variable: **"\$.containerBuildDone"**
BooleanEquals: **false**
Next: **WaitForDockerBuild**


```

phases:
  build:
    commands:
      - echo "***** In Build *****"
      - aws s3 sync s3://mattw-reinvent2017-rawsitepages src/posts
      - export GATSBY_BUILDTIME=$(tail -1 .git/logs/HEAD | awk '{print $5}')
      - ln -s /backupmod/node_modules ./node_modules
      - gatsby build
      - echo "***** Build Complete *****"
  post_build:
    commands:
      - echo "***** In Post Build *****"
      - aws s3 sync public s3://mybucket --delete --acl public-read
      - echo "***** Post Build Complete *****"
# - command
artifacts:
  files:
    - public/**/*

```

AWS re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



The **buildspec.yml** file for the **Gatsby** site above is a little different because we don't need to build a **Docker** image. We are generating the **markdown** files and then building the Gatsby site before doing a sync back to the **S3** bucket hosting our site



Here is a view of the workflow we have been seeing in **Step Functions**, we start by taking pictures using a **RaspberryPi**. We went to **Vize.ai** service website and uploaded about **6000 images** having many different rotations of the slides that I am going to be presenting at the conference today and trained the network to identify what slide number is presented to it with very high accuracy. We are then using this service to take real time image of the presentation slide at the conference and sending it to the Vize.ai service endpoint to get what slide we are on and get the slide text from **DynamoDB** to post on our website for live blog updates. We can also use the **Amazon Rekognition** service instead of Vize.ai.

How about monitoring?

We monitor Lambda via Logs

```
MONITORING|unix_epoch_timestamp|value|count|  
            my.metric.name|#tag1:value,tag2  
  
console.log(monitoringstring);
```

We do monitoring in 2 ways, using **CloudWatch metrics** where we are collecting a lot of data about the different lambdas execution times, how many failures are occurring, etc. we are also collecting similar metrics from the step functions. But we can also do custom metrics, this is done inside your lambda function where you simply create and write a line to your log file as above.

What are the Lambda Metrics

```
aws.lambda.duration, min, max, sum  
aws.lambda.errors  
aws.lambda.invocations  
aws.lambda.throttles  
aws.lambda.iterator_age
```

Here are the metrics that we grab from CloudWatch every 15 secs,

What are the Step Function Metrics

```
aws.states.execution_time (.maximum, .minimum)  
aws.states.executions_aborted, failed, started, succeeded, timed_out  
aws.states.lambda_functions_scheduled, started, succeeded, timed_out  
aws.states.execution_throttled  
aws.states.lambda_function_time (.maximum, .minimum)  
aws.states.lambda_function_run_time (.maximum, .minimum)  
aws.states.lambda_function_schedule_time (.maximum, .minimum)  
aws.states.lambda_functions_started
```

And here are the Step Function metrics that we also collect every 15 secs and kept for a 15 months duration.

Summary

Step Functions make AWS Lambda orchestration easy

Lambda and Step Functions can be an interesting part of your infrastructure

Monitor your Lambda functions to ensure peak performance



Source can be found at:

<https://github.com/DataDog/mattw-reinvent2017-demo>

AWS re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



GitHub repository view for **DataDog / mattw-reinvent2017-demo**

Content for demo created for Matt Williams' re:Invent 2017 demo

13 commits 1 branch 0 releases 1 contributor

Branch: master New pull request Create new file Upload files Find file Clone or download

technovangelist add some readmes Latest commit c8abcf8 on Dec 4, 2017

.vscode	valid package.jsons	2 months ago
liveblogger	add some readmes	a month ago
pushTestImage	add some readmes	a month ago
sitebuilderjs	add some readmes	a month ago
slideimageprep	add some readmes	a month ago
slidenotes	add some readmes	a month ago
website	It works!!	2 months ago
.gitignore	Updates . session is done and it went really well	2 months ago
LICENSE	Create LICENSE	2 months ago
LICENSE-3rdparty.csv	Updates . session is done and it went really well	2 months ago
README.md	add some readmes	a month ago

README.md



AWS re:Invent

THANK YOU!

Matthew Williams - @technovangelist - Datadog

AWS
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

