

# Building Microservices with Event Sourcing and CQRS

Michael Ploed - innoQ

@bitboss



springone *EXPL*



Unless otherwise indicated, these slides are © 2012-2015 Pivotal Software, Inc. and licensed under a Creative Commons Attribution-NonCommercial license: <http://creativecommons.org/licenses/by-nc/3.0/>


**Let's review the  
classical old  
school N-Tier  
architecture**



## Characteristics

1

*We read and write data through the same layers*



2

*We use the same model for read and write access*



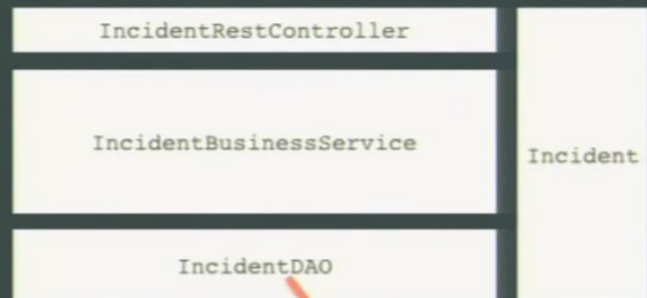
3

We use coarse grained deployment units that combine read and write code




4

We change datasets directly



ID	USER_ID	DATE	TEXT
1	23423	11.03.2014	Mouse is broken
2	67454	12.03.2014	EMail not working
3	93729	12.03.2014	Office license
...	...	...	...

We've done that  
for ages already and  
it's proven



Many applications will  
run smooth and fine  
with this kind of  
approach

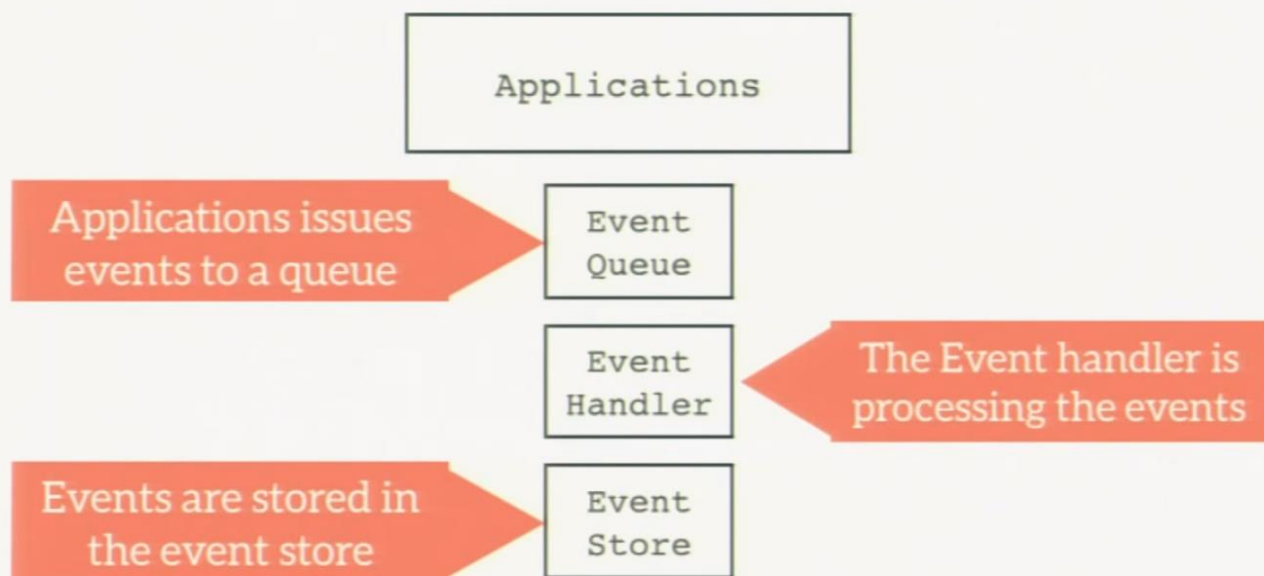
# However there are drawbacks to this kind of architecture

- 1 *The data model is a compromise*
- 2 *You can't scale read and write independently*
- 3 *No data history, no snapshots, no replay*
- 4 *Tendency to a monolithic approach*



**Event Sourcing is an architectural pattern in which the state of the application is being determined by a sequence of events**

### Building Blocks



**The sequence of events in the queue  
is called event stream**



## **Event Stream Example**

### **IncidentCreatedEvent**

```
incidentNumber: 1  
userNumber: 23423  
timestamp: 11.03.2014 12:33:33  
text: „Mouse broken“  
status: „open“
```


### **IncidentTextChangedEvent**

```
incidentNumber: 1  
text: „Left button of mouse broken“
```

### **IncidentClosedEvent**

```
incidentNumber: 1  
solution: „Mouse replaced“  
status: „closed“
```

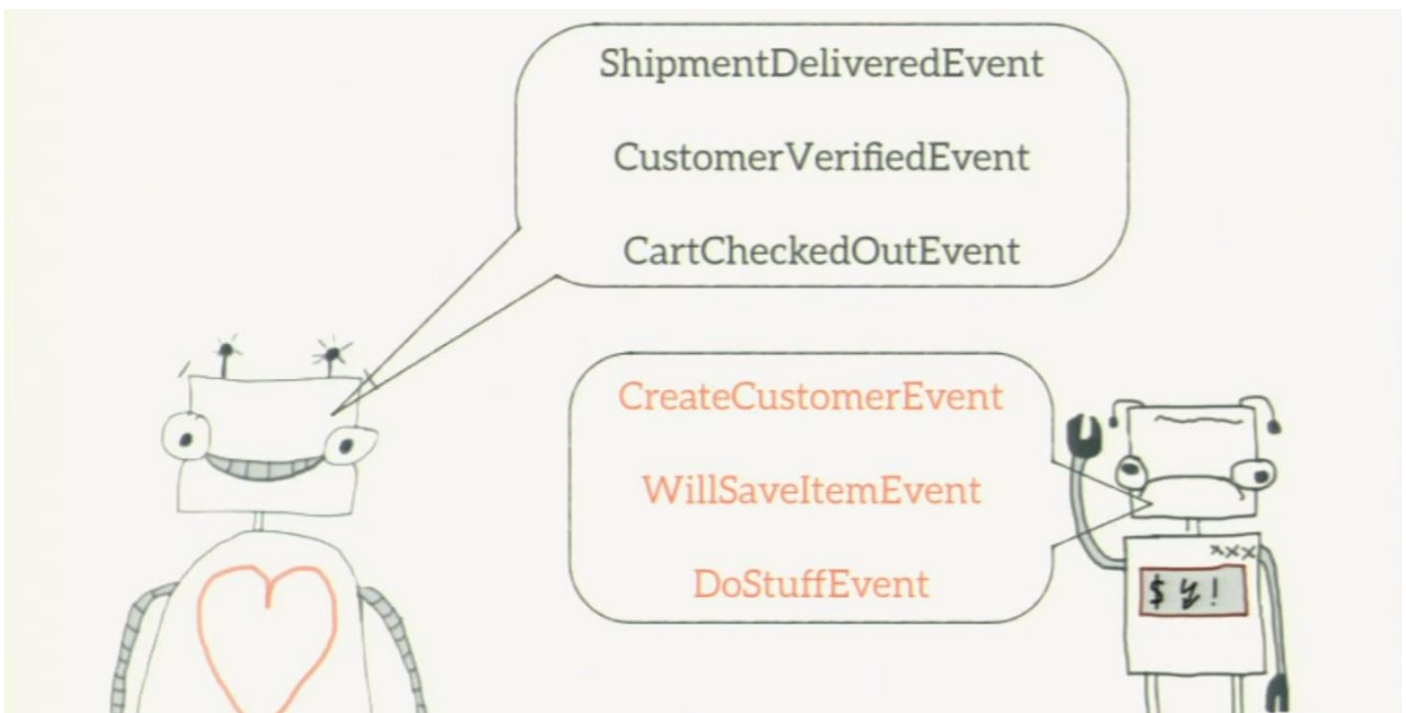
**An event is something  
that happened in the past**



The names of the events are part  
of the

Ubiquitous Language

D D D



The bad event names in red are command-style event names that are bad in DDD!



## Code Example

```
public class CustomerVerifiedEvent {  
    private String eventId;  
    private Date eventDate;  
    private CustomerNumber customerNumber;  
    private String comment;  
  
    public CustomerVerifiedEvent(CustomerNumber custNum,  
                                String comment) {  
        this.customerNumber = cusNum;  
        this.comment = comment;  
        this.eventDate = new Date();  
    }  
}
```

Scope your events based on  
Aggregates

DDD

How fine grained should we model our events? Aggregates in DDD are a collection of entities bounded to a specific business model. There is always a root entity that other entities are based on, the root entity is usually the aggregate that we can use to generate events.



An Event is always immutable



There is no deletion of events

*A delete is  
just another  
event*

#### **IncidentCreatedEvent**

```
incidentNumber: 1  
userNumber: 23423  
timestamp: 11.03.2014 12:23:23  
text: „Mouse is broken defekt“  
status: „open“
```

#### **IncidentChangedEvent**

```
incidentNumber: 1  
text: „Maus ist Kaputt“
```


#### **IncidentRemovedEvent**

```
incidentNumber: 1
```

**The event bus is usually  
implemented by a message  
broker**



**Let's reuse the ESB  
from the failed SOA  
project**



# NO



*Prefer dumb pipes  
with smart  
endpoints as a  
suitable message  
broker architecture*

1

*Complete rebuild is possible*

2

*Temporal Queries*

3


*Event Replay*

Well known examples


=

Version Control Systems  
or  
Database Transaction Logs





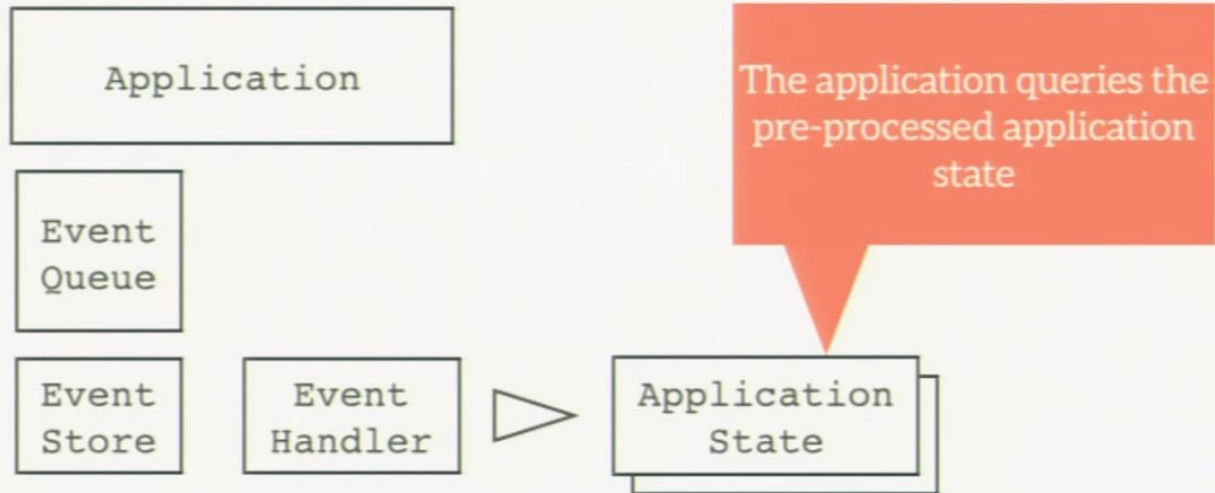
# **The Event Store has a very high business value**



Aren't there performance issues attached to this kind of data storage?

YES!

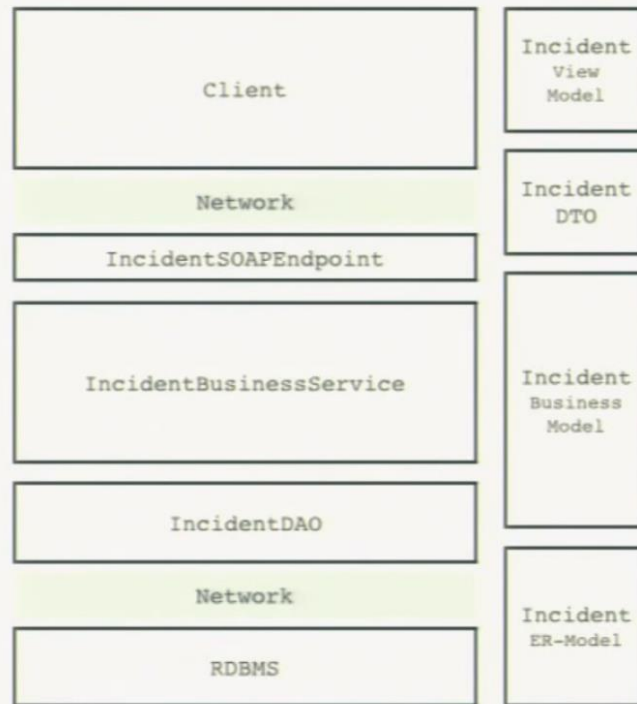
## Think about application state



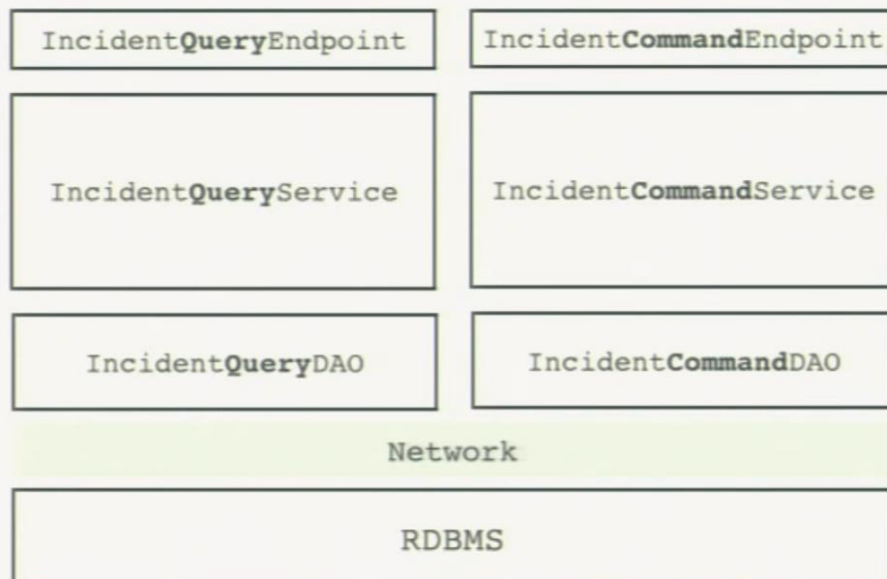
# CQRS

**Command**  
**Query**  
**Responsibility**  
**Separation**

Command are Write operations, Queries and Read operations. We need to separate our Read and Write operations.



*Basically the idea behind CQRS is simple*



We split up the Read and Write parts of our application into separate logic different from the above. Because we are still working with the same data model

# Code Example

# Classic Interface

```
public interface IncidentManagementService {  
    Incident saveIncident(Incident i);  
    void updateIncident(Incident i);  
    List<Incident> retrieveBySeverity(Severity s);  
    Incident retrieveById(Long id);  
}
```

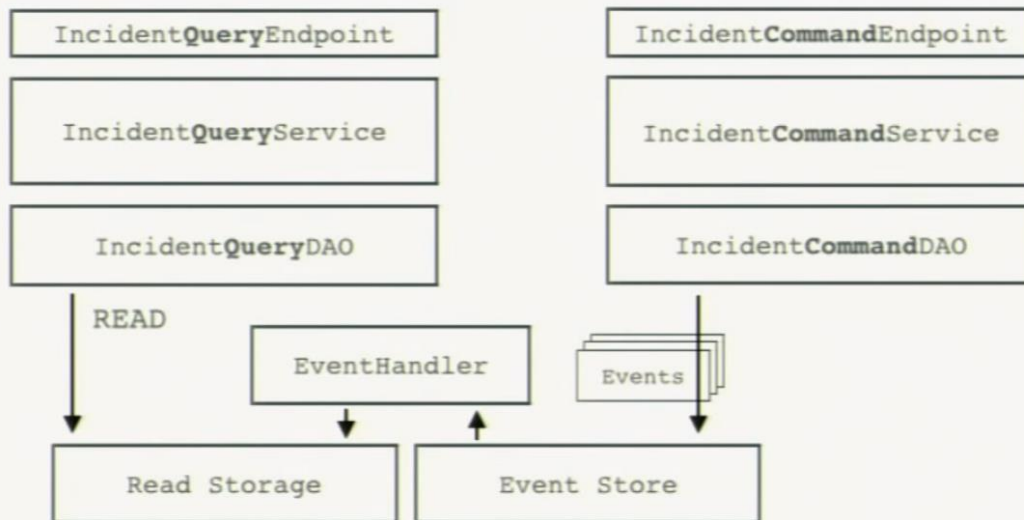
# CQRS-ified Interfaces

```
public interface IncidentManagementQueryService {  
    List<Incident> retrieveBySeverity(Severity s);  
    Incident retrieveById(Long id);  
}
```

```
public interface IncidentManagementCommandService {  
    Incident saveIncident(Incident i);  
    void updateIncident(Incident i);  
}
```

Split it up into a query service and a command service as above

# Event Sourcing & CQRS



The Command part issues events to the events store where we have an event handler that is able to derive the application state into a separate Read storage that you query against.

1

*Individual scalability and deployment options*

2

*Technological freedom of choice for command, query and event handler code*

3

*Excellent Fit for Bounded Context (Domain Driven Design)*

A bounded context for a microservice might also include many data aggregates, bind your bounded context of your microservices come from aggregates in your business domain



*Event Sourcing and CQRS are interesting architectural options. However there are various challenges, that have to be taken care of*

1

Consistency

2

Validation

3

Parallel Updates

Can we use a RDBMS as our events data store? Where do we validate data about our aggregates in event sourcing? What about parallel updates or processing of events? An event can be a serialized JSON data, generally a data format that can be easily stored like JSON, XML, or even string.

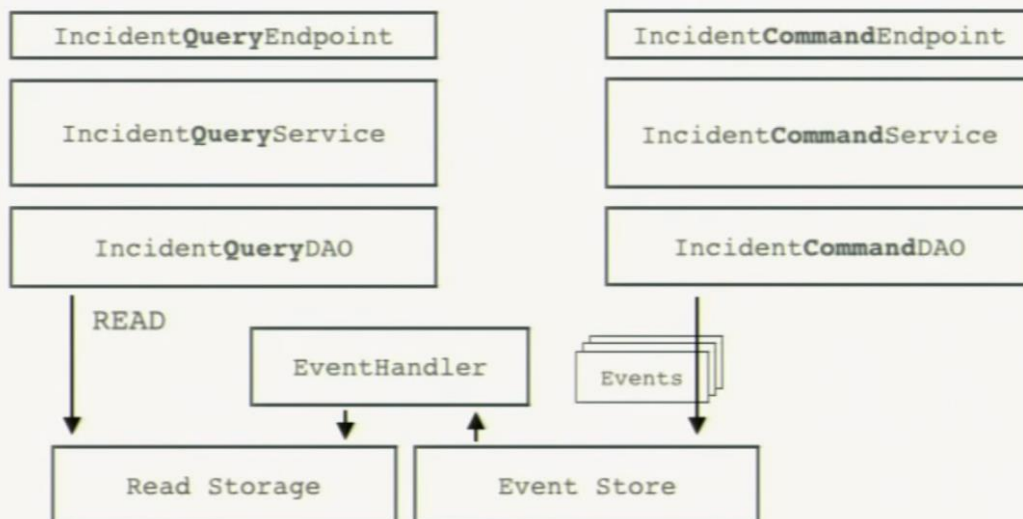
**YES**

In terms of consistency, the answer is YES



*Systems based on  
CQRS and Event  
Sourcing are  
mostly eventually  
consistent*

## *Eventual Consistency*



# Eventual Consistency



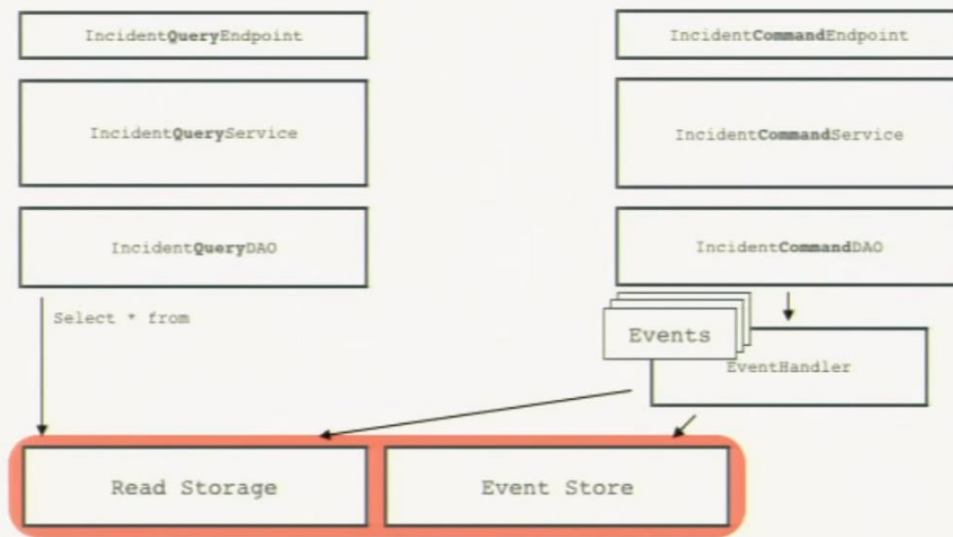
There are consistency risks at the Events level, the EventHandler level, or the Read Storage if we get issues. We can lose event being stored in the if the event store goes does. From the **CAP Theorem**, Consistency, Availability and Partition Tolerance, **mostly is that we need to pick between Consistency over Availability**.

# BUT



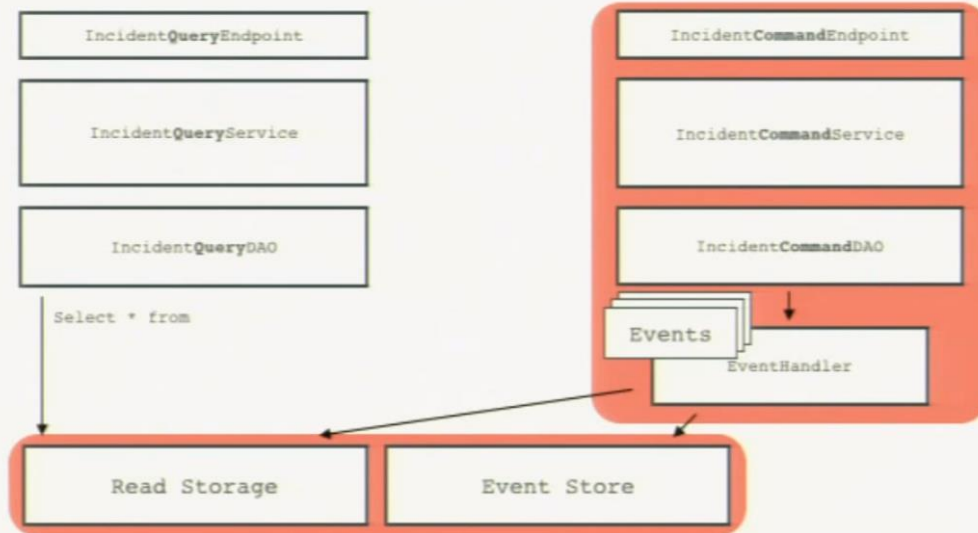
You can build a  
fully consistent  
system which  
follows Event  
Sourcing  
principles

# Full Consistency



We can go ahead and place the Event Store and the Read Storage in a RDBMS on the same node

# Full Consistency



Then we could add the event to both the Event Store and the Read Storage in a single transaction, this will then provide very high transactional guarantee and consistent, but this system will not scale very well for Availability!!!

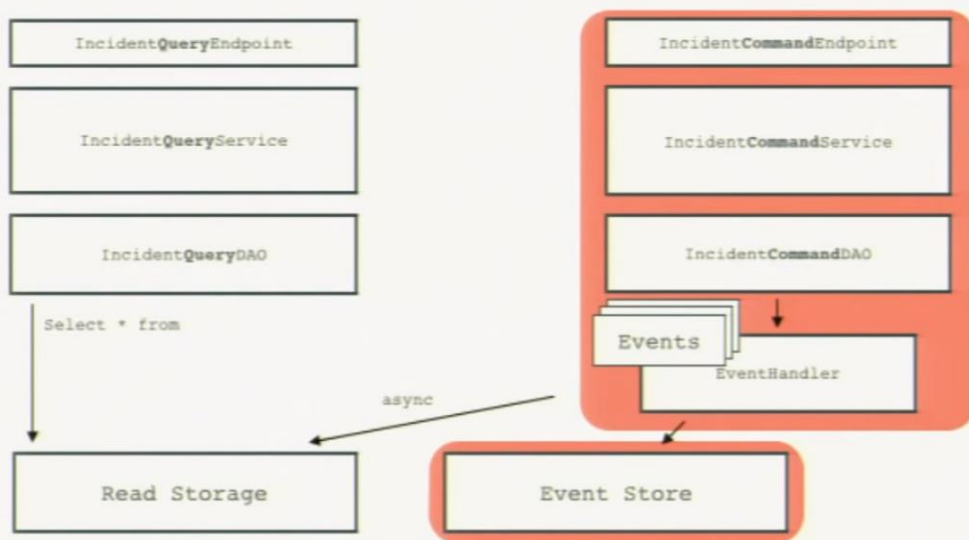
# Your business domain drives the level of consistency not technology

## Deeper Insight

# DDD

You need to **consider the business value of failed transactions in your system** as well as the **business and regulatory requirements** for your use case **before making a decision regarding the level of data consistency** you want.

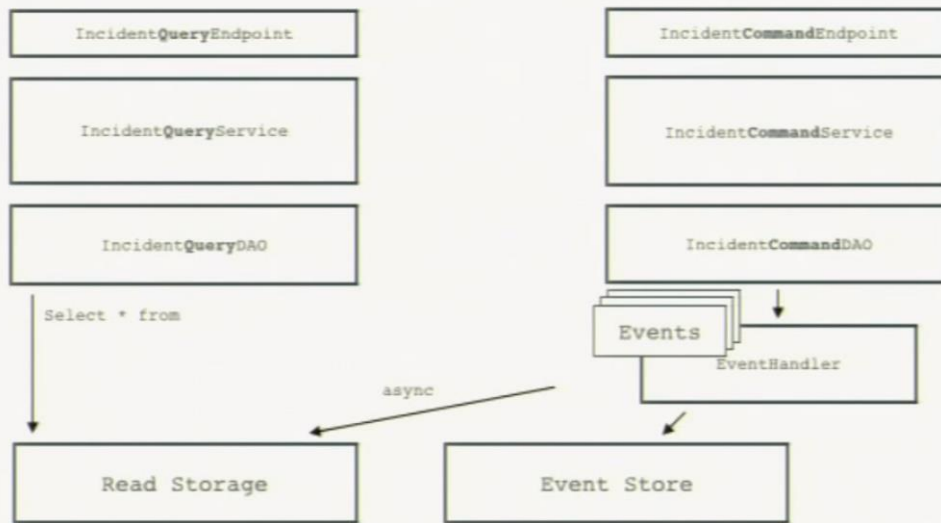
## Increased (but still eventual) consistency



You could decide to add events to the Event store only and do an async storage to the Read Storage as above



# Increased (but still eventual) consistency



We can use updates to the read storage more infrequently

**There is no  
standard solution**

1

Consistency

2

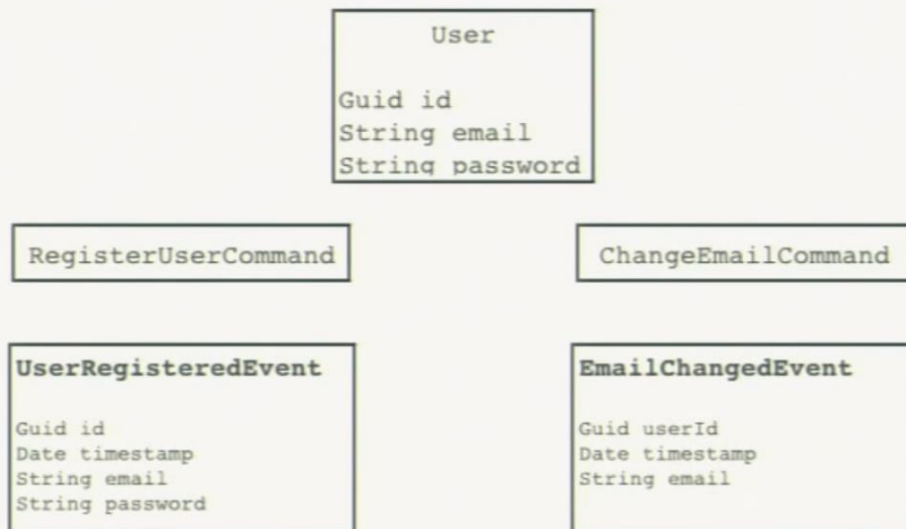
Validation

3

Parallel Updates

Validation is easier because we can do things like null checks for some required fields. Uniqueness validations are different

# Example Domain



We have 2 commands based on the User object as above,



Then we have the requirement above for our event source system. We have a risk of having wrong validation if we do not update our Read Storage immediately after an event occurred, we might not be able to guarantee email uniqueness. We might have to sacrifice some availability in order to increase the consistency level in this use case.



*How high is the  
probability that a  
validation fails*

*Which data is required  
for the validation*

*Where is the required  
data stored*



*What is the business  
impact of a failed  
validation that is not  
recognized due to  
eventual consistency  
and how high is the  
probability of failure*

Your business domain drives the level  
of consistency  
*Deeper Insight*

**D D D**

1

*Validate from Event Store*


2

*Validate from Read Store*

3

*Perform Validation in Event Handler*

We should strongly validate from the Read Store and update it based on the criticality of the business operation



***Never validate  
from the event  
store***

This will create problems in high volume, high transaction systems

1

Consistency

2

Validation

3

Parallel Updates

## Example Domain

User

Guid id  
String email  
String password

RegisterUserCommand

ChangeEmailCommand


**UserRegisteredEvent**

Guid id  
Date timestamp  
String email  
String password

**EmailChangedEvent**

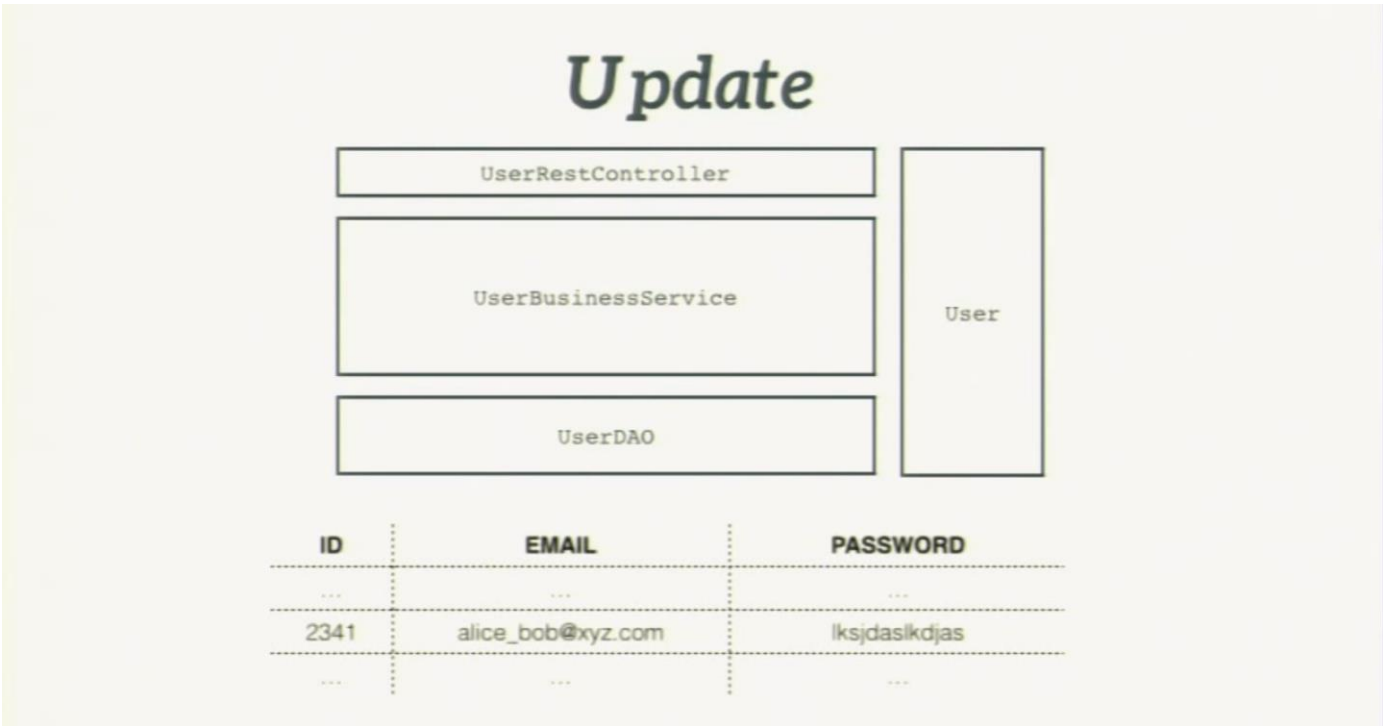
Guid userId  
Date timestamp  
String email



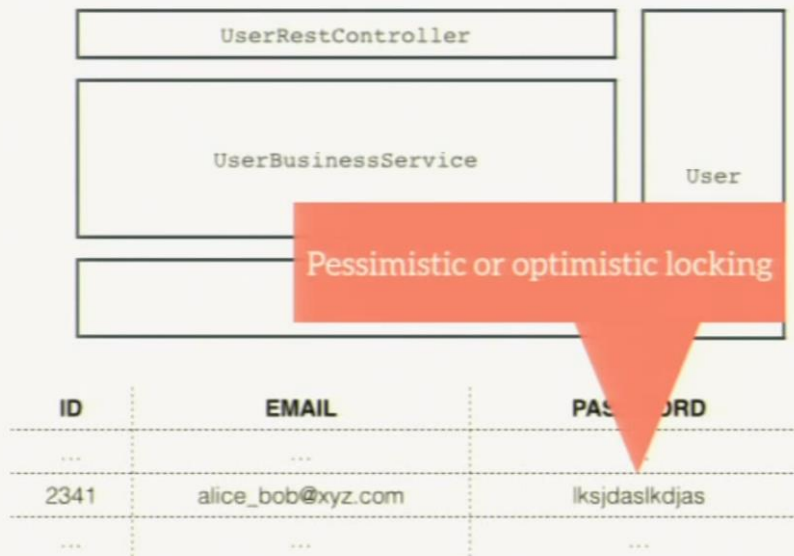




What would we do  
in a  
„classic old school  
architecture“



# Update



We usually issue the classical 'SELECT FOR' update command while locking that field

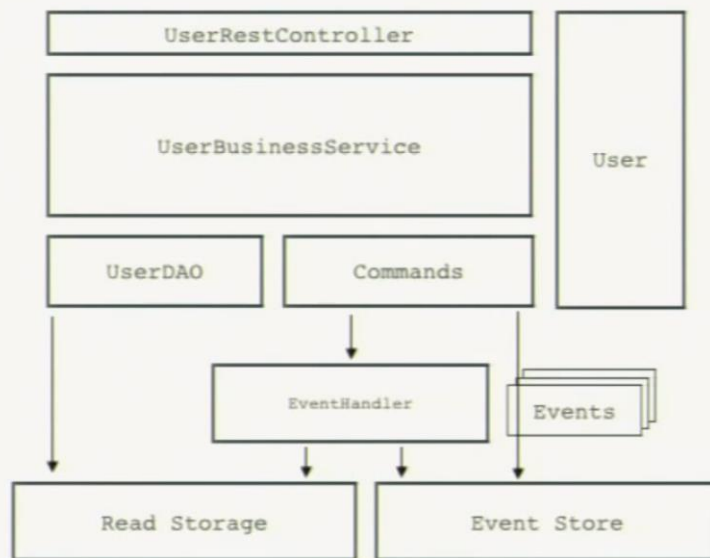
Your business domain drives the  
locking quality  
Deeper Insight

D D D

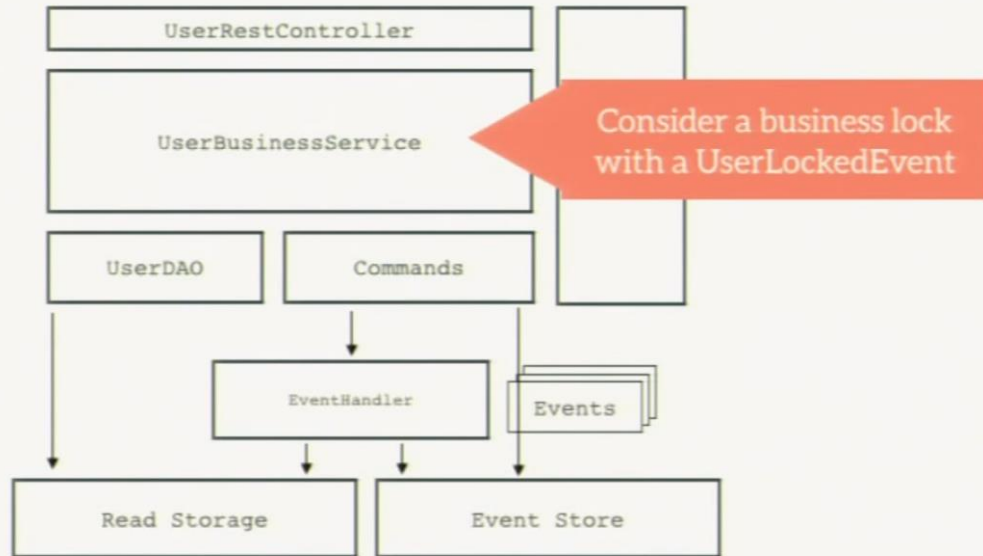


**Pessimistic  
locking on a  
data-level will  
hardly work in  
event sourcing  
architectures**


**Where to „pessimistically“ lock?**



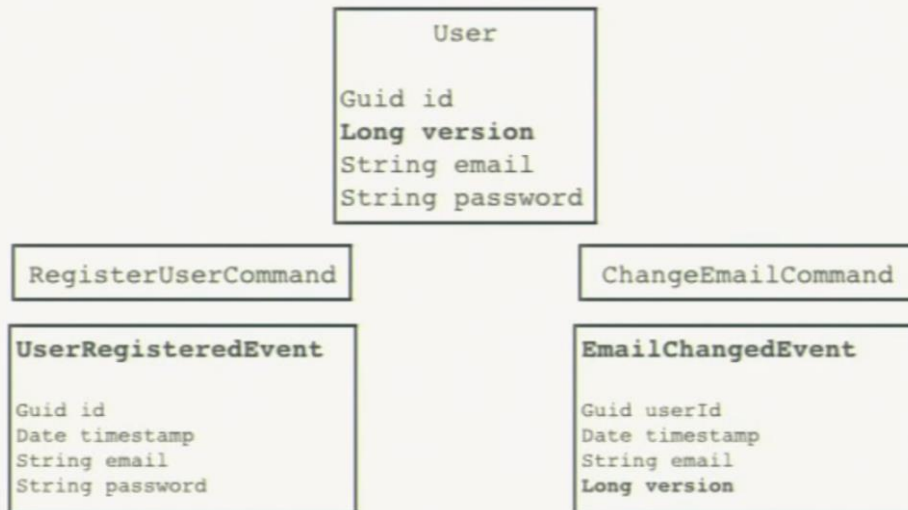
## Where to „pessimistically“ lock?



This 'pessimistic lock' can be permitted in a mortgage application business scenario




# Introduce a version field for the domain entity

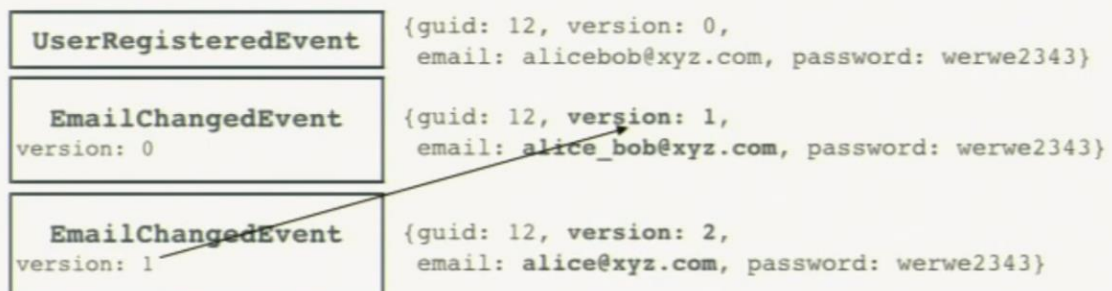


This will be an 'optimistic locking' scenario

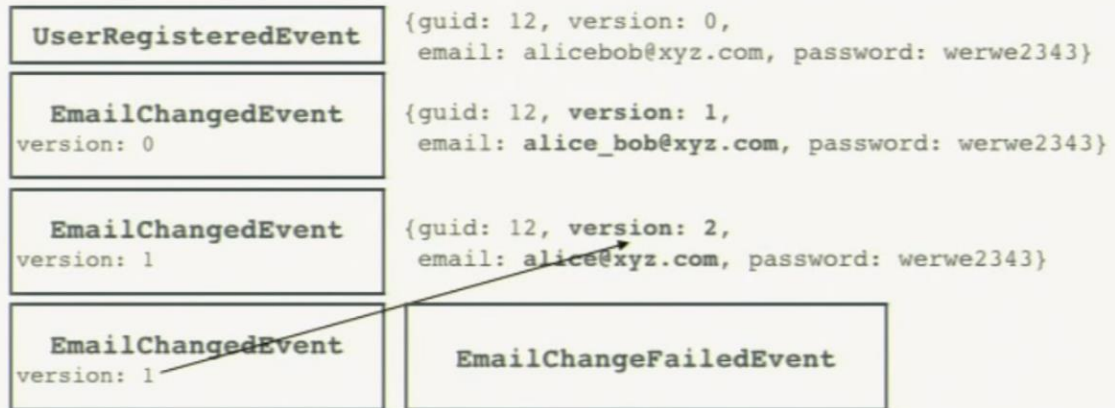
## Each writing event increases the version



## Each writing event increases the version



## Each writing event increases the version



This will be detected when we update our Read Store and we can raise another event to notify the user to go and change their email again

Also here:  
you should be as  
consistent as  
your domain  
requires