Financial institutions today must manage multiple data types from a wide variety of sources. Among these various data types, archive data presents a particular challenge: it is invisible to much of the organization and not easily leveraged by the lines of business for analytics, insight, and product innovation. Faced with massive volumes of archive data, Financial Services organizations are finding that delivering insights in a timely manner requires a data storage and analytics solution with more agility and flexibility than traditional data management systems can provide. In this session, we will discuss a design pattern that (1) brings this data into a highly available, lower-cost queryable archive within AWS than you currently have and (2) migrates that data to a data lake that the entire organization can use to extract insight and drive innovation. We will walk through a strategy that addresses the following topics: storing archive data in compressed, cost-effective, and readily available formats; creating lifecycle policies to archive older data sets and make them easily accessible; fully utilizing the features of object storage to enrich the data lake; and applying AWS analytics tools to gather business insights.

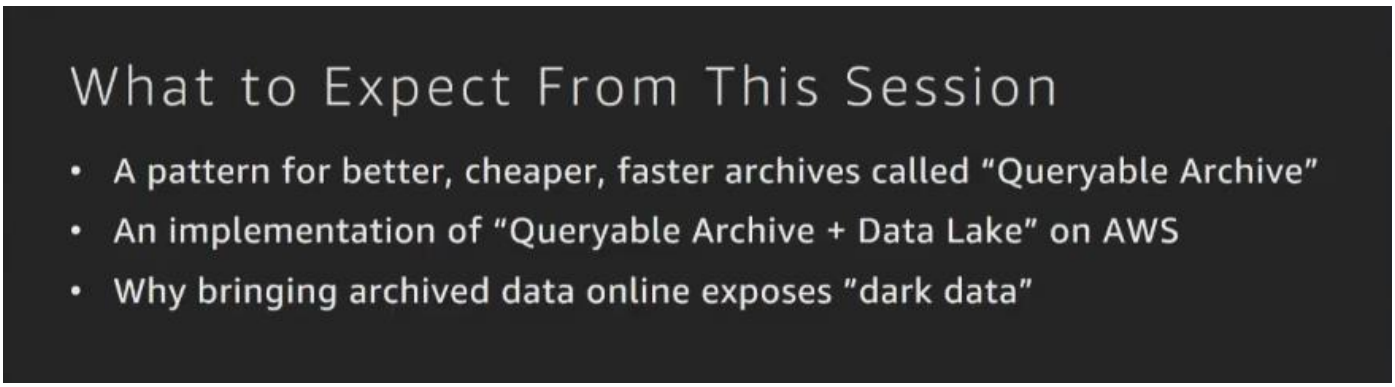

Agenda

- Why We Archive
- Building A "Queryable Archive + Data Lake"
- Demo of A "Queryable Archive + Data Lake"
- Benefits and Costs
- Unleashing "Dark Data"
- Next Steps

Why We Archive

Archive For Performance

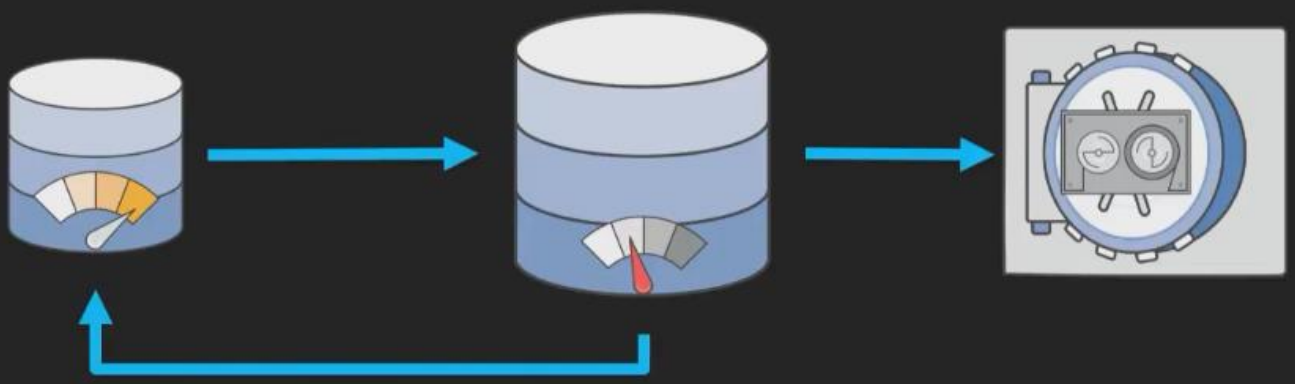

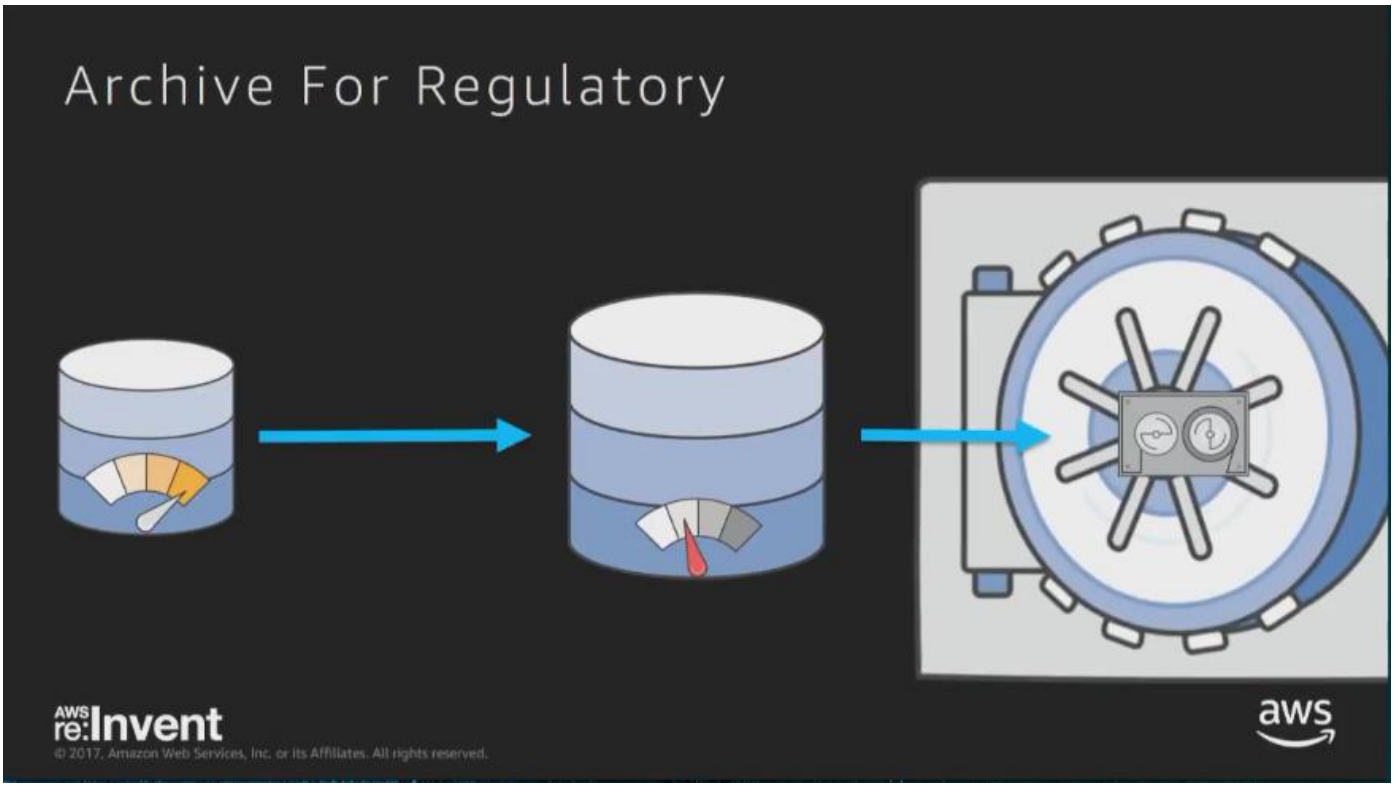

We now have way more storage attached to compute which is not what we intended during design phase

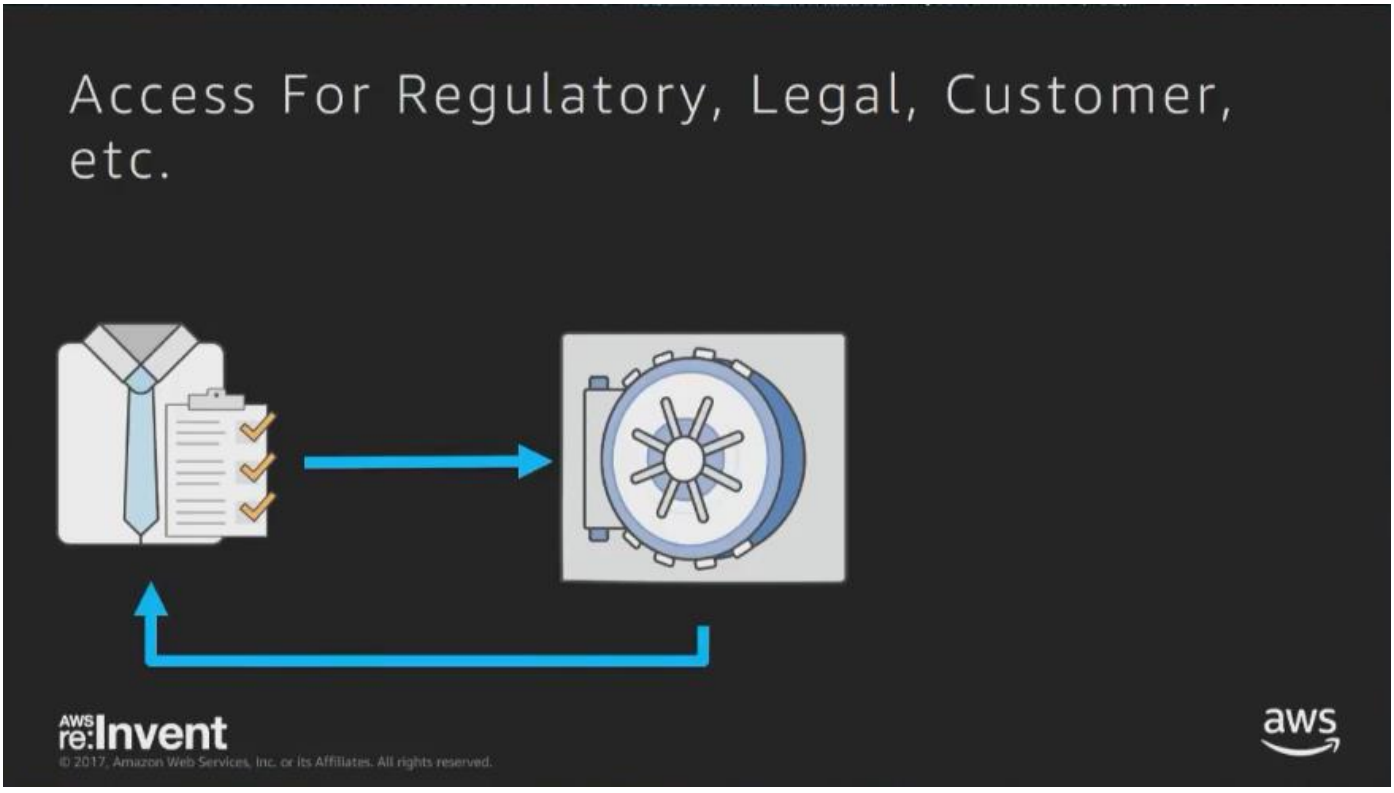

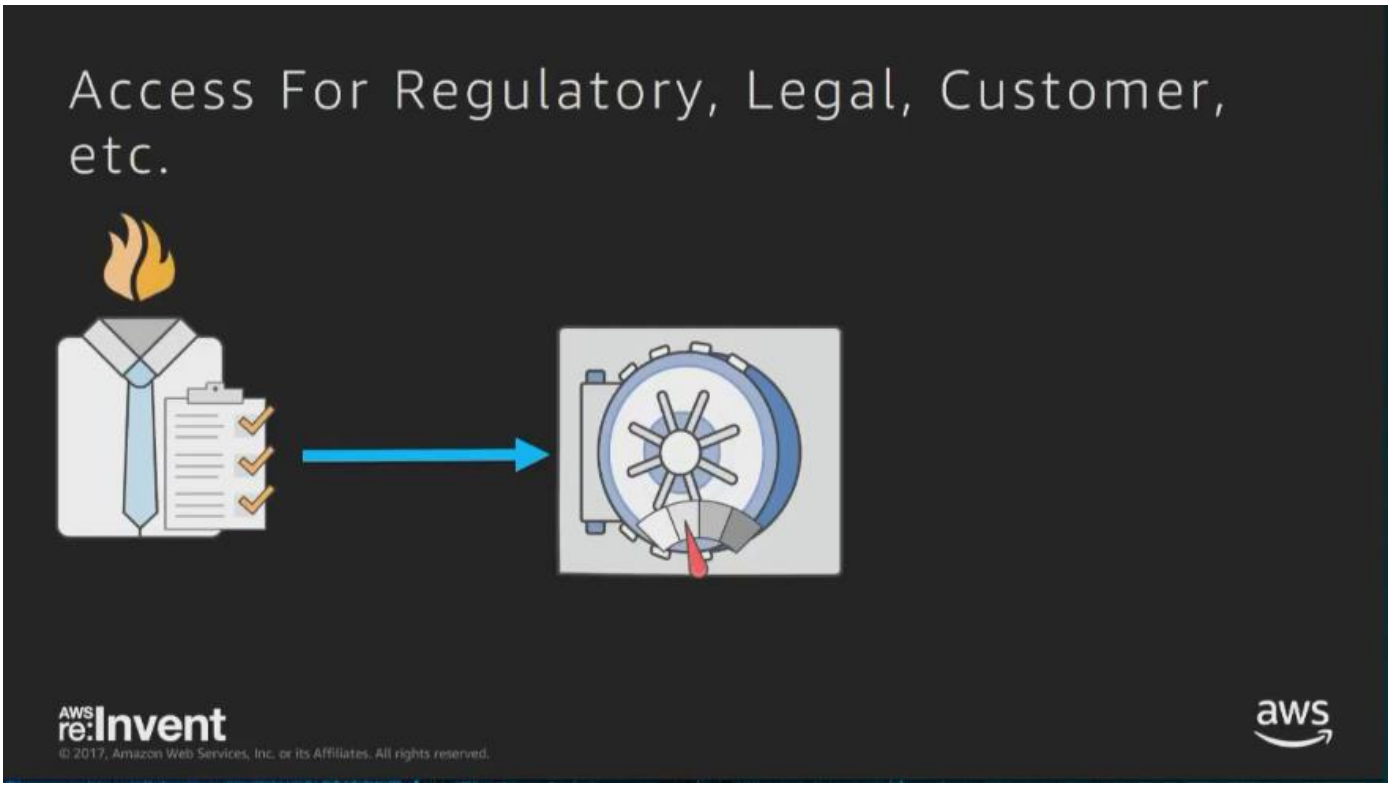

we now have an end user who needs access to data and can’t get it

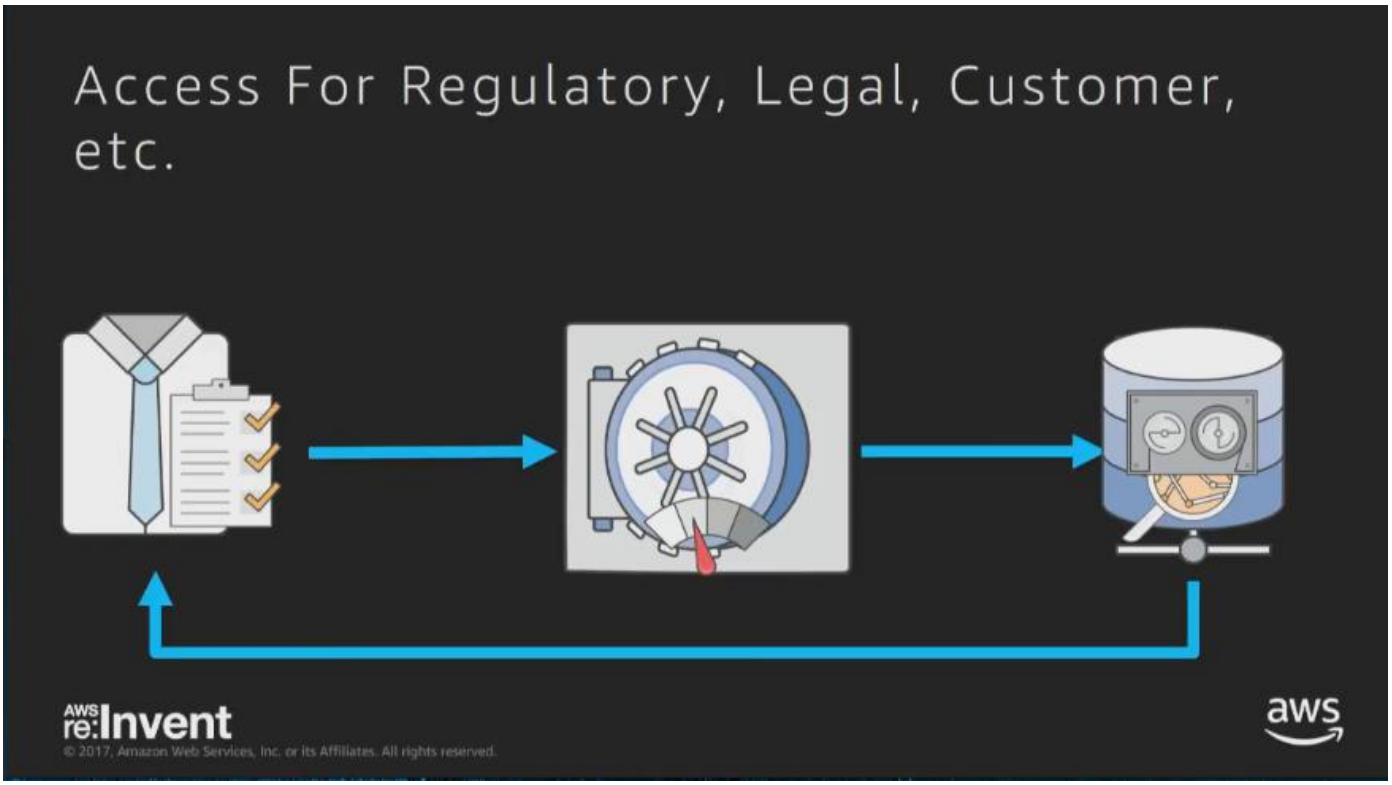

We have to set up spare database and copy the archives into it, then get the date out for the customer

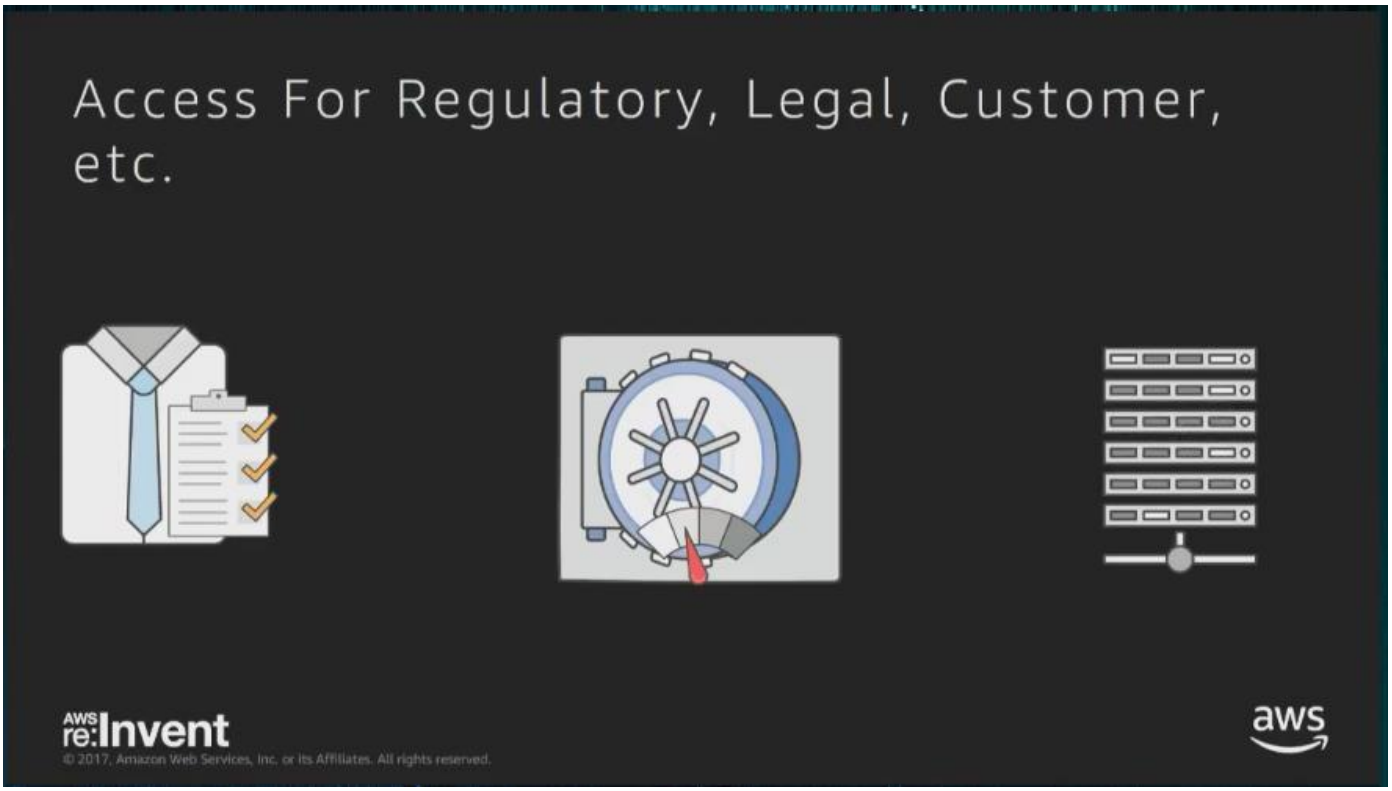

We will then tear everything down again

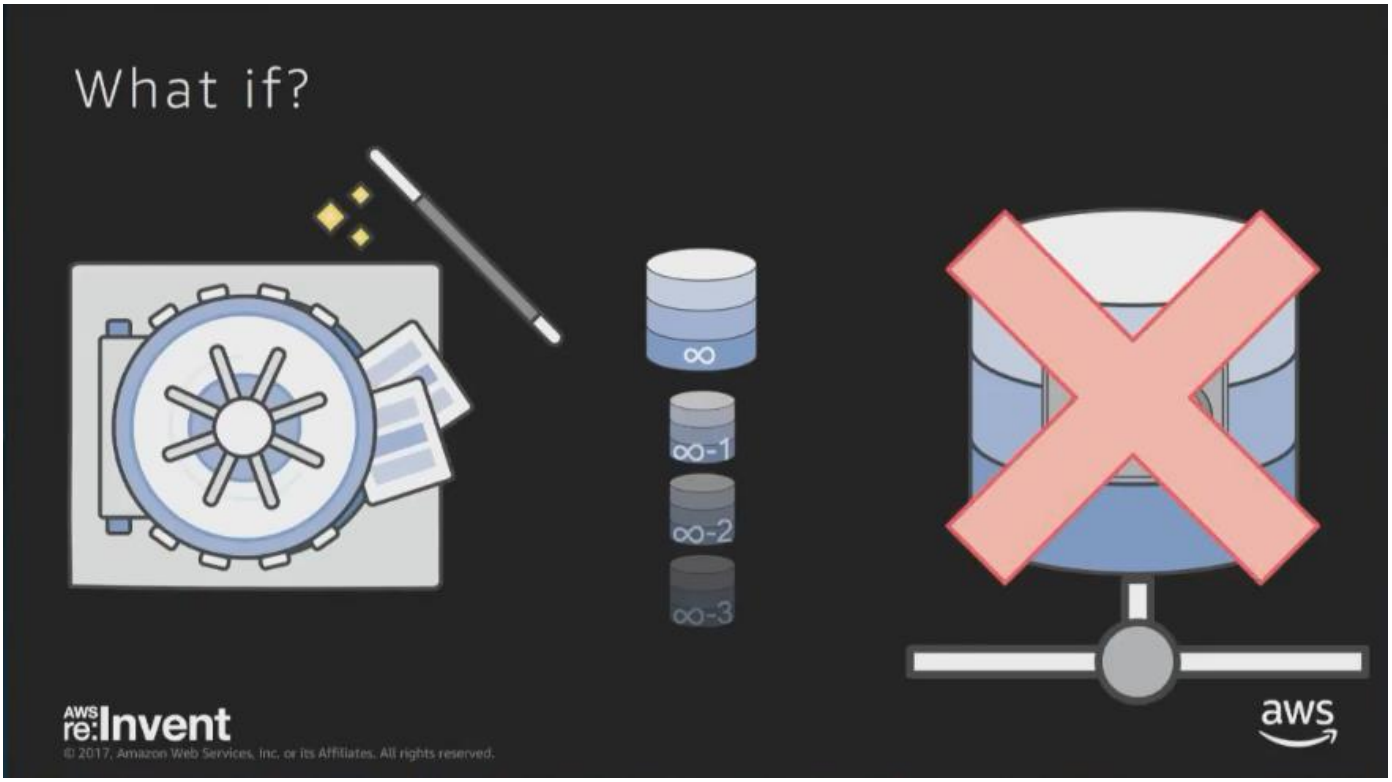

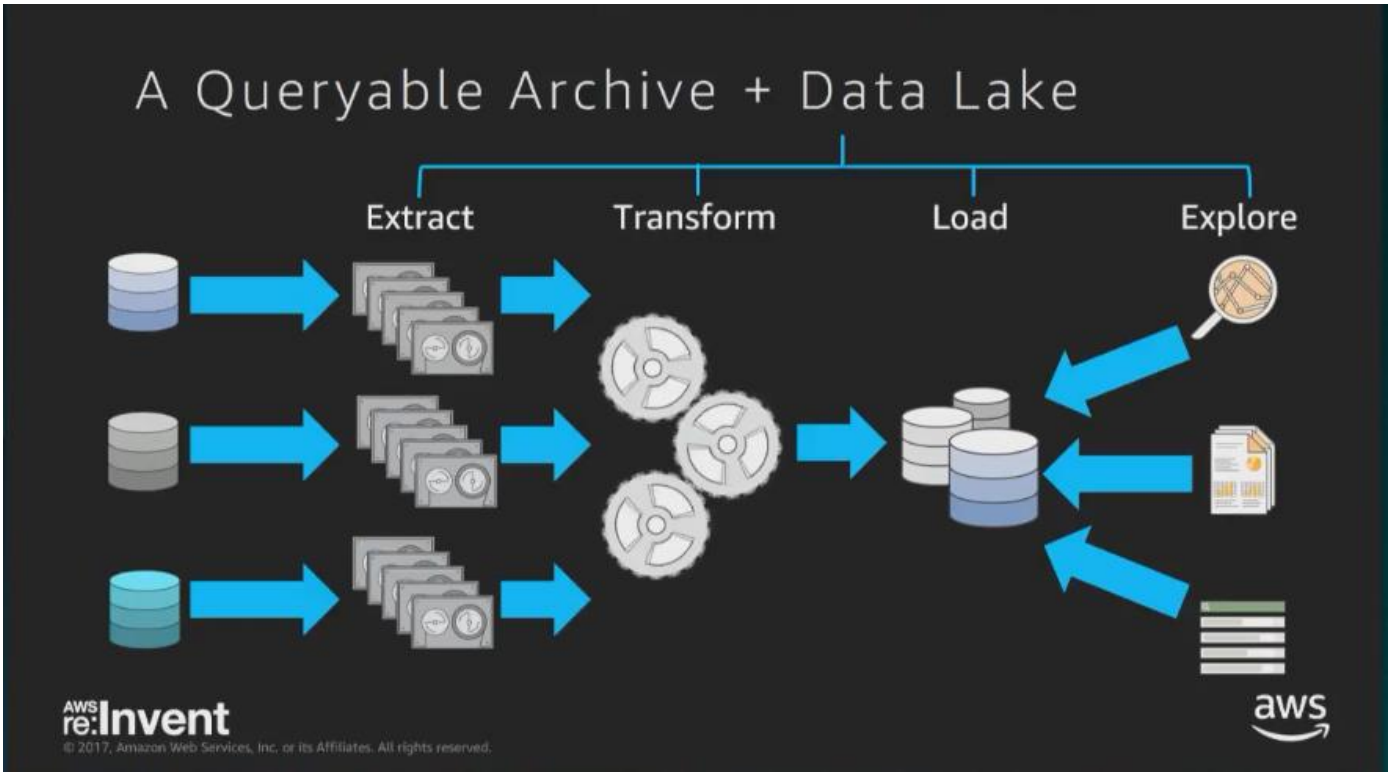

We start out with all our historical financial datasets that are sitting in our databases in data centers today. Once the data is loaded into the specific AWS database services, we have a variety of tools to start exploring that data

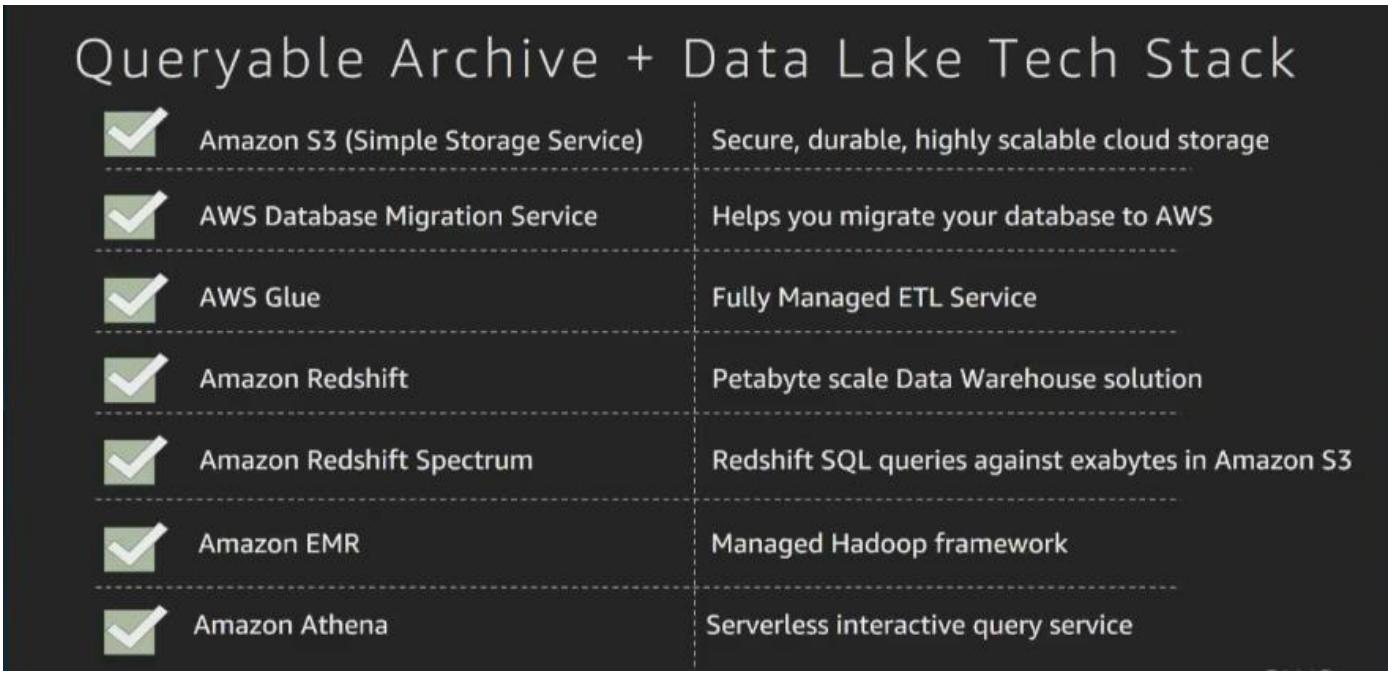

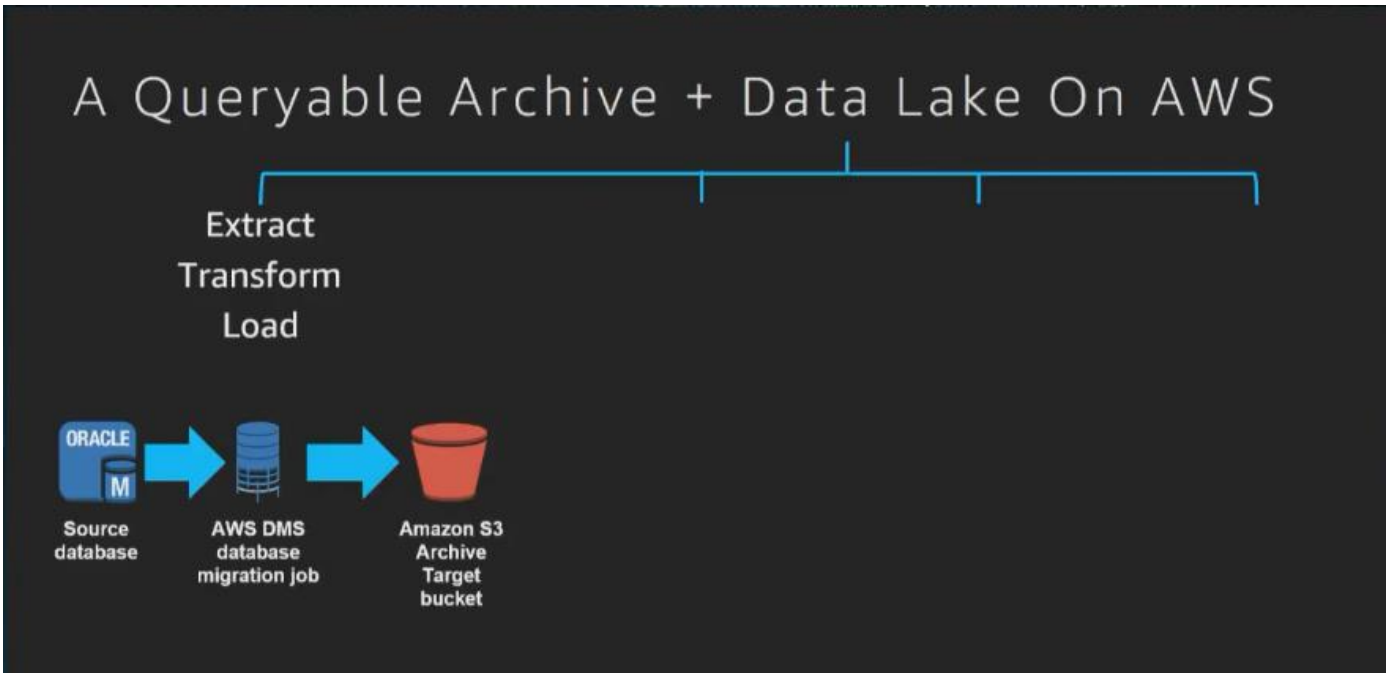

We use the AWS Data Migration Service and now have the data in S3 in a compressed CSV format.

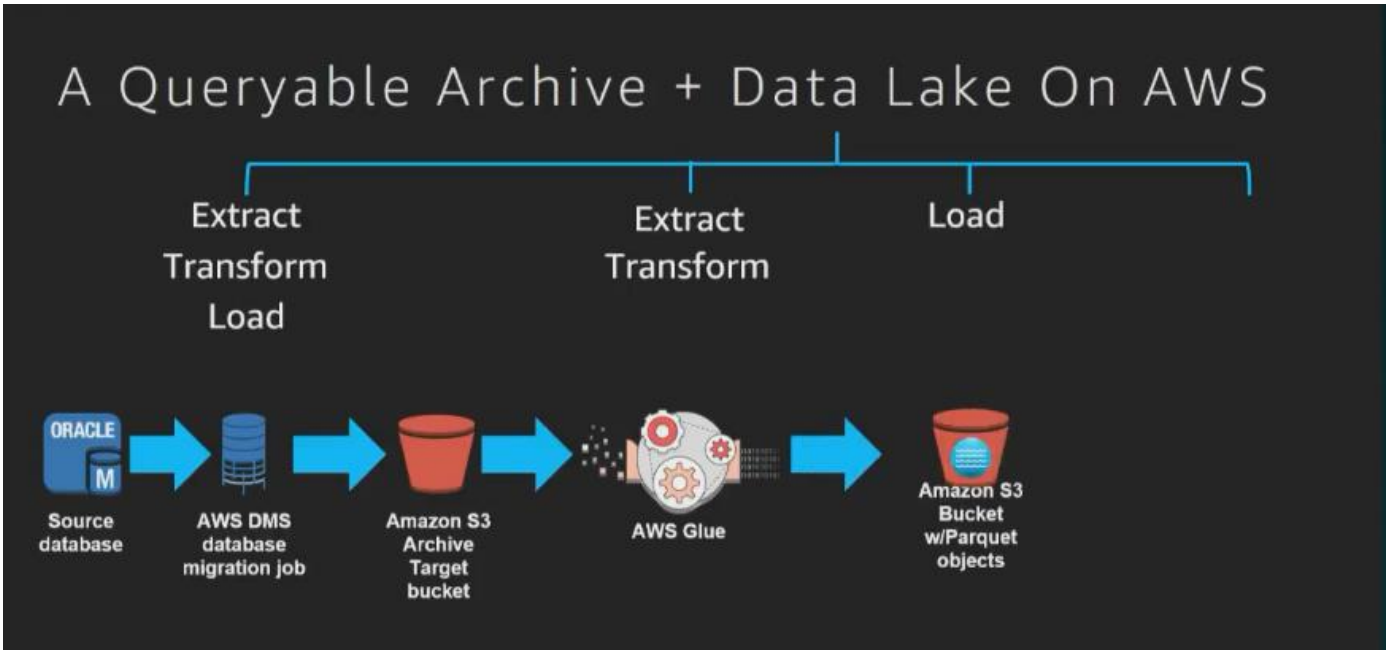

We now again use AWS Glue to transform the data into our enterprise data format and store the data in S3 in the Parquet object format

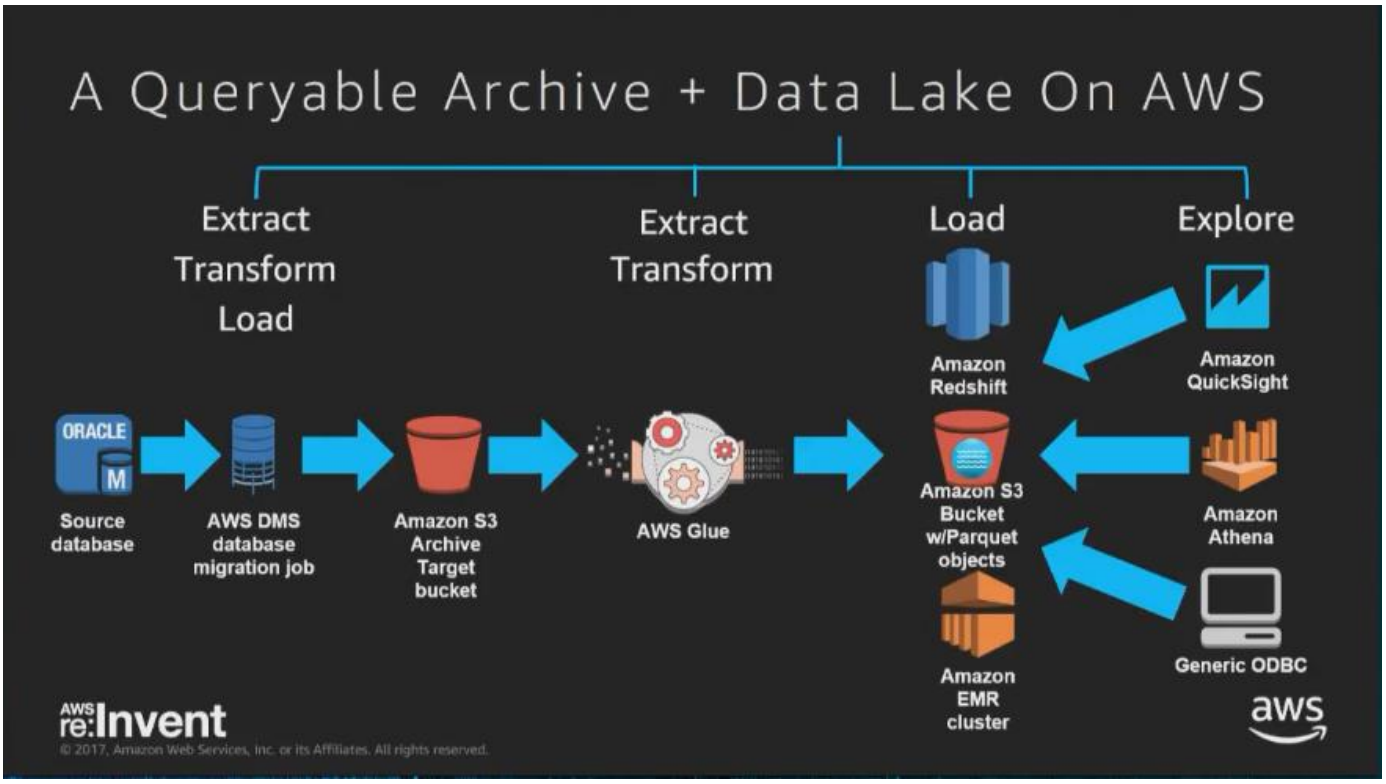

We can now access and explore the data

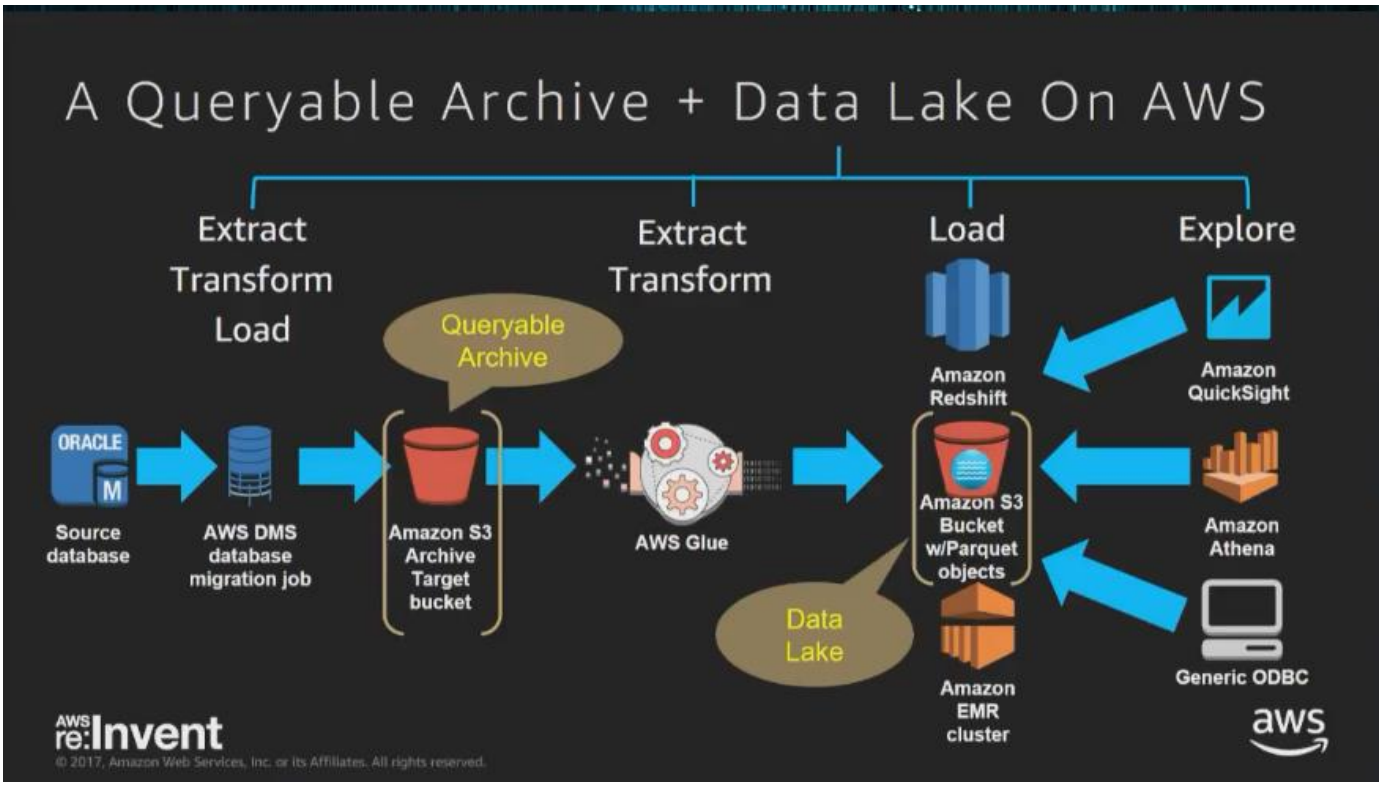

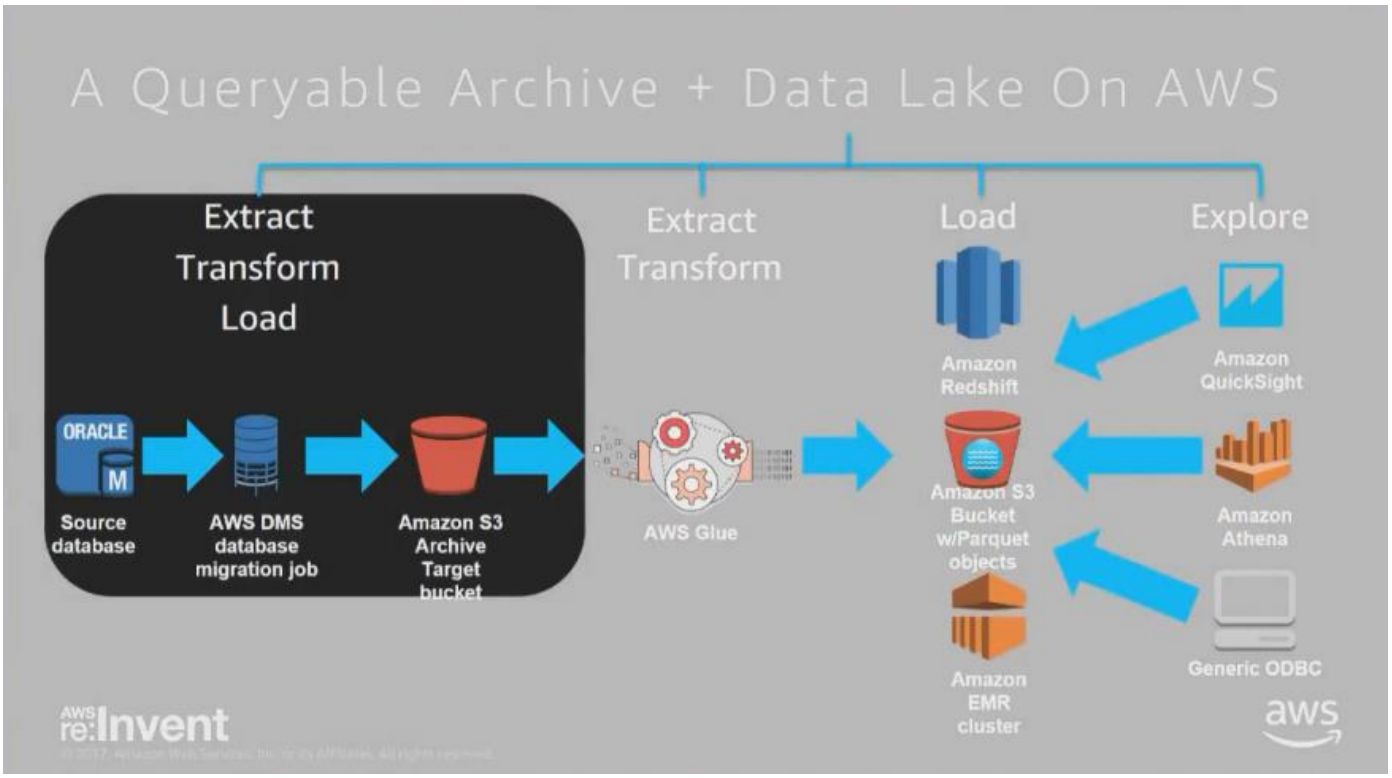
A Queryable Archive + Data Lake On AWS
Extract
Transform
Load
Extract
Transform
Load
Explore
ORACLE
Source database
AWS DMS database migration job
Amazon S3 Archive Target bucket
AWS Glue
Amazon Redshift
Amazon S3 Bucket w/Parquet objects
Amazon EMR cluster
Amazon QuickSight
Amazon Athena
Generic ODBC
aws re:Invent
aws

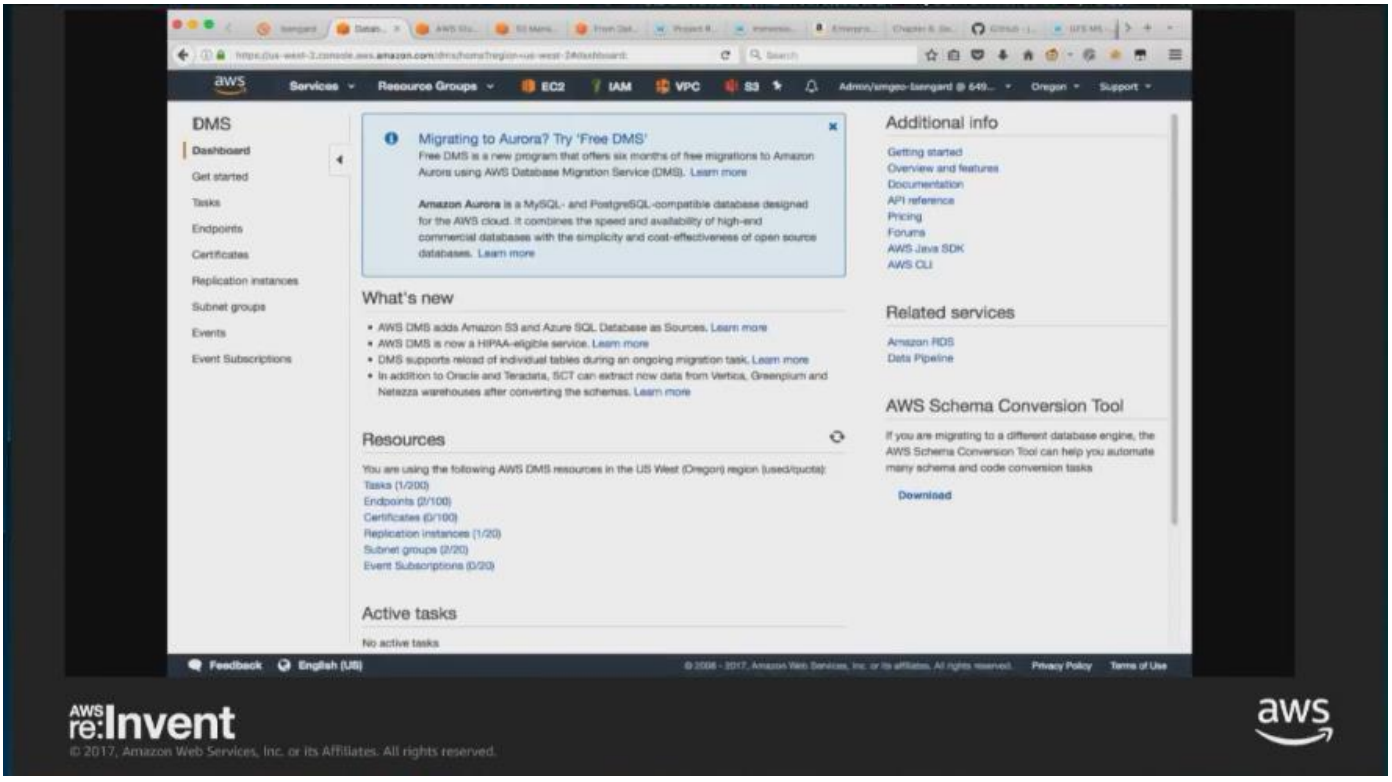
DMS
Dashboard
Get started
Tasks
Endpoints
Certificates
Replication instances
Subnet groups
Events
Event Subscriptions
Migrating to Aurora? Try 'Free DMS'
Free DMS is a new program that offers six months of free migrations to Amazon Aurora using AWS Database Migration Service (DMS). Learn more
Amazon Aurora is a MySQL- and PostgreSQL-compatible database designed for the AWS cloud. It combines the speed and availability of high-end commercial databases with the simplicity and cost-effectiveness of open source databases. Learn more
What's new
AWS DMS adds Amazon S3 and Azure SQL Database as Sources. Learn more
AWS DMS is now a HIPAA-eligible service. Learn more
DMS supports reload of individual tables during an ongoing migration task. Learn more
In addition to Oracle and Teradata, SCT can extract now data from Vertica, Greenplum and Netezza warehouses after converting the schemas. Learn more
Resources
You are using the following AWS DMS resources in the US West (Oregon) region (used/quota):
Tasks (1/200)
Endpoints (2/100)
Certificates (0/100)
Replication instances (1/20)
Subnet groups (2/20)
Event Subscriptions (0/20)
Active tasks
No active tasks
Additional info
Getting started
Overview and features
Documentation
API reference
Pricing
Forums
AWS Java SDK
AWS CLI
Related services
Amazon RDS
Data Pipeline
AWS Schema Conversion Tool
If you are migrating to a different database engine, the AWS Schema Conversion Tool can help you automate many schema and code conversion tasks
Download
Feedback
English (US)
Privacy Policy
Terms of Use
aws re:Invent
© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.
aws

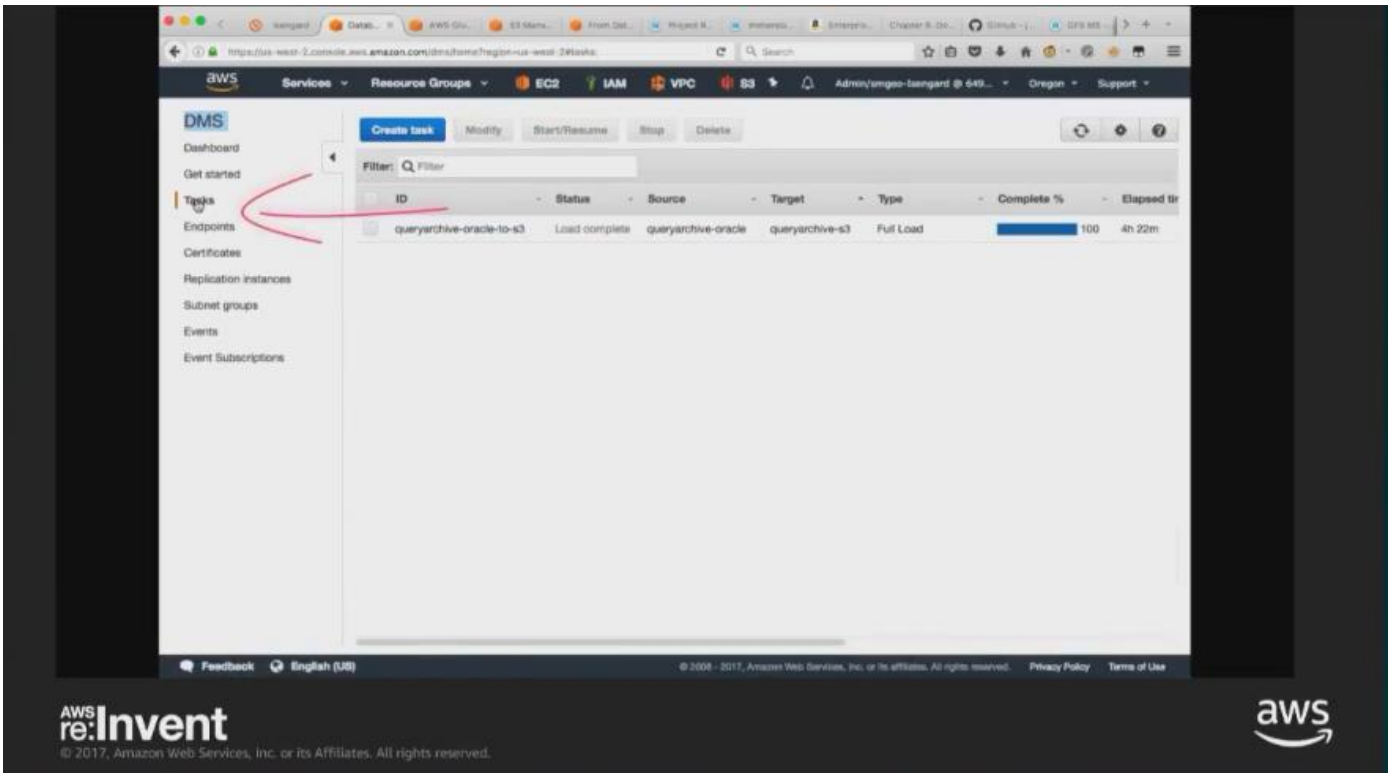

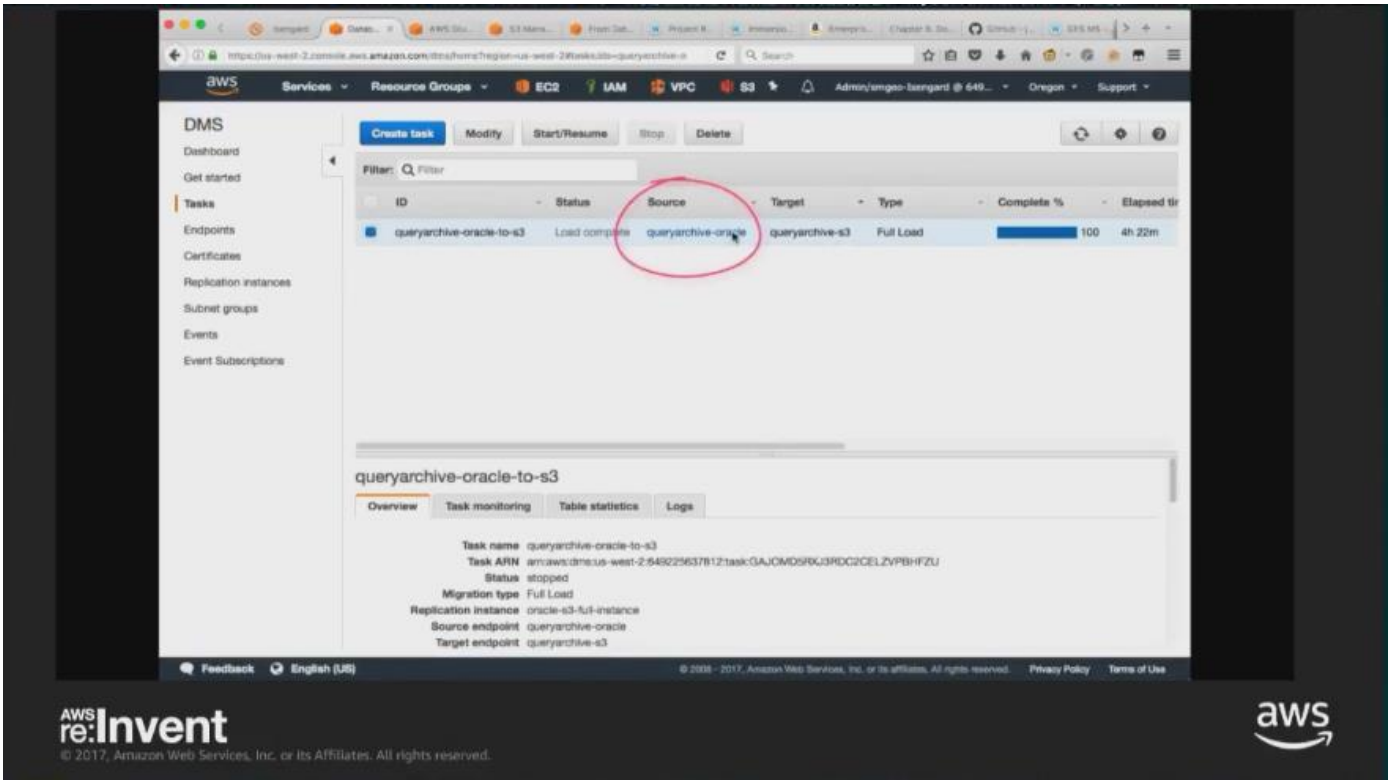

A Task is a batch job that moves data from one point to another

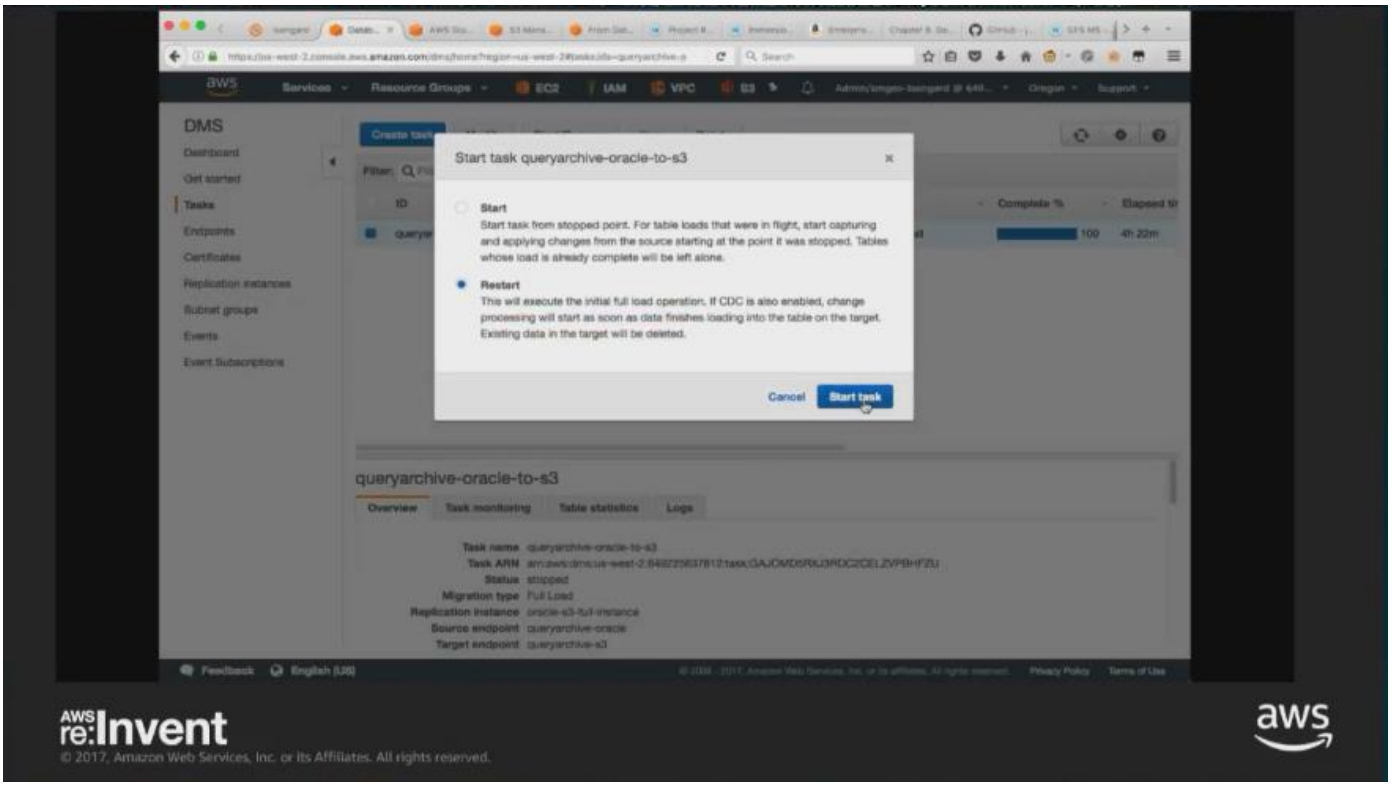

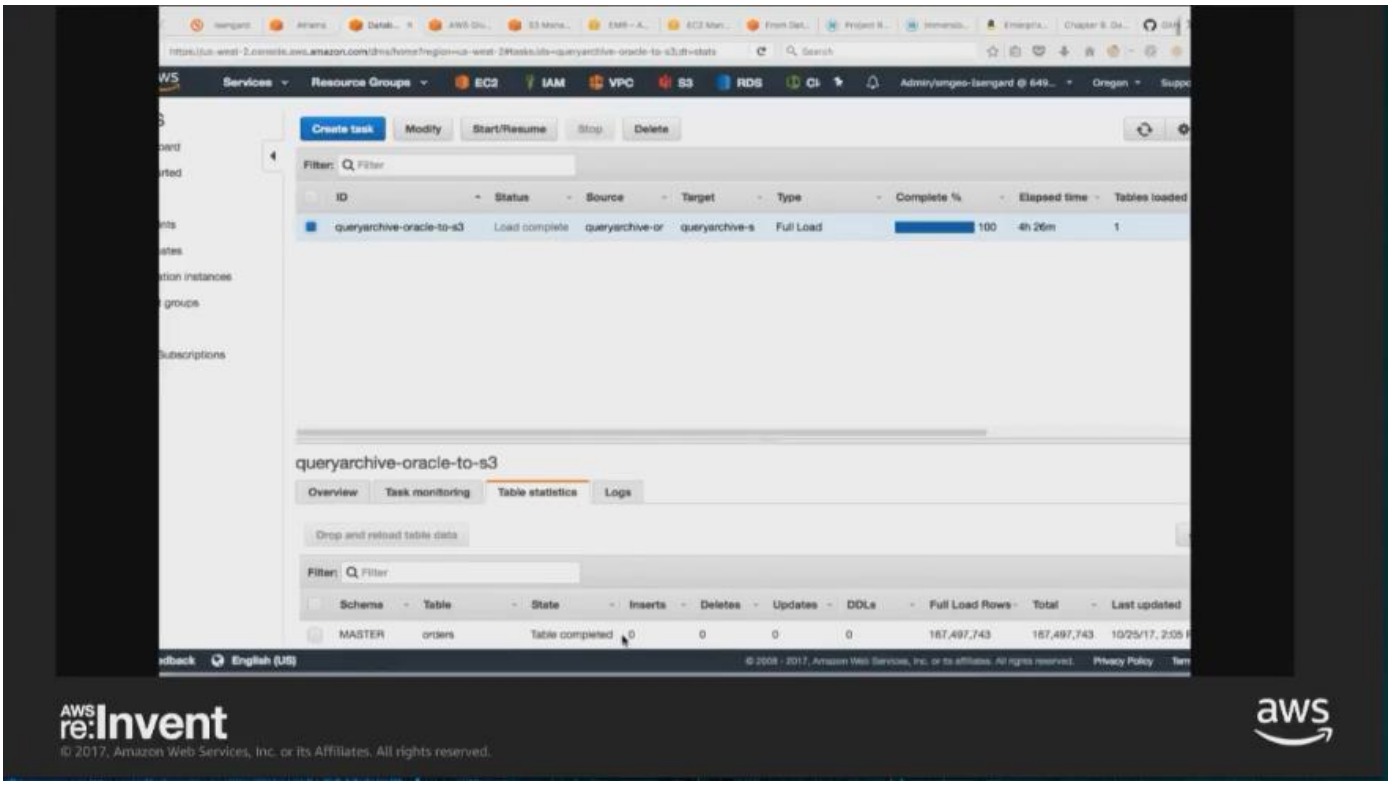

The job is done 4h 26m later and we have copied 167,497,743 rows from the Oracle database to S3 bucket in compressed CSV format. The dataset contains orders and order executions by a trading desk.

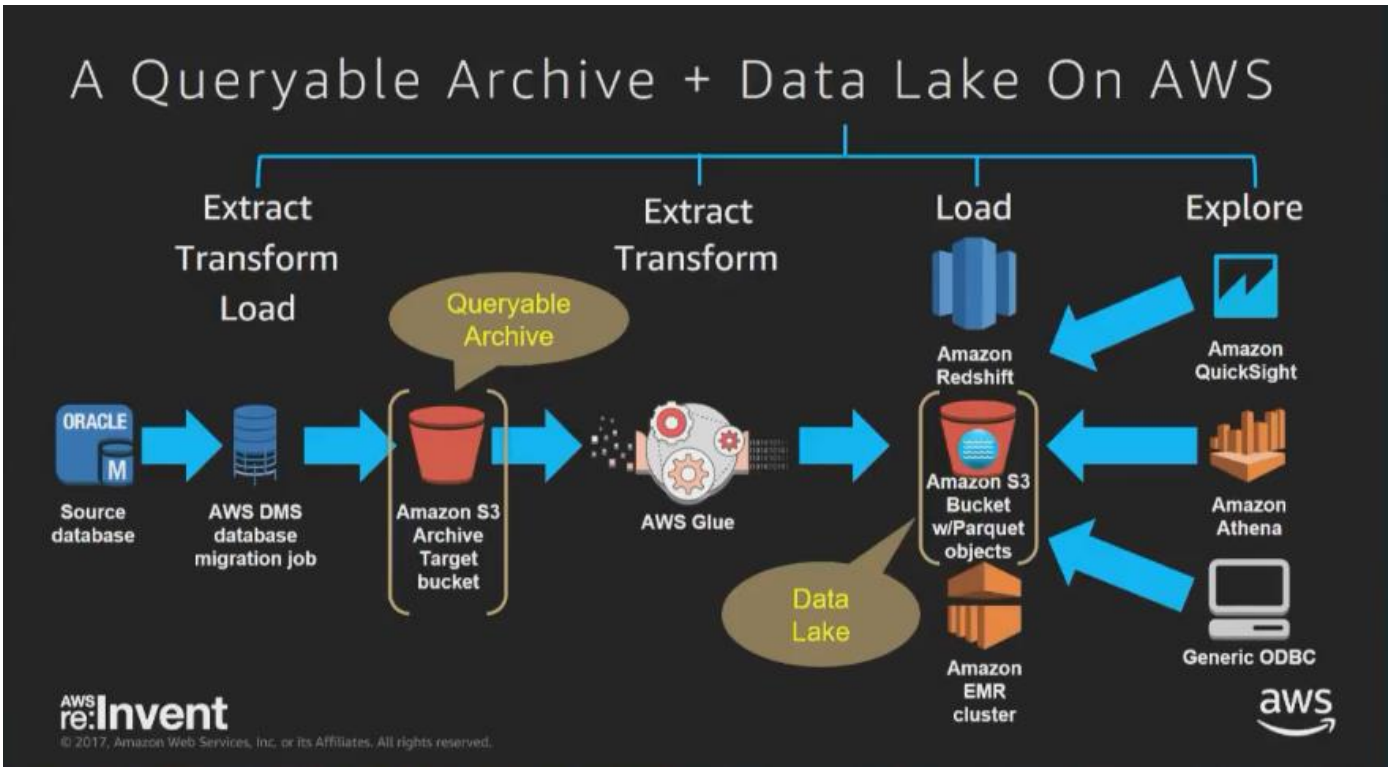

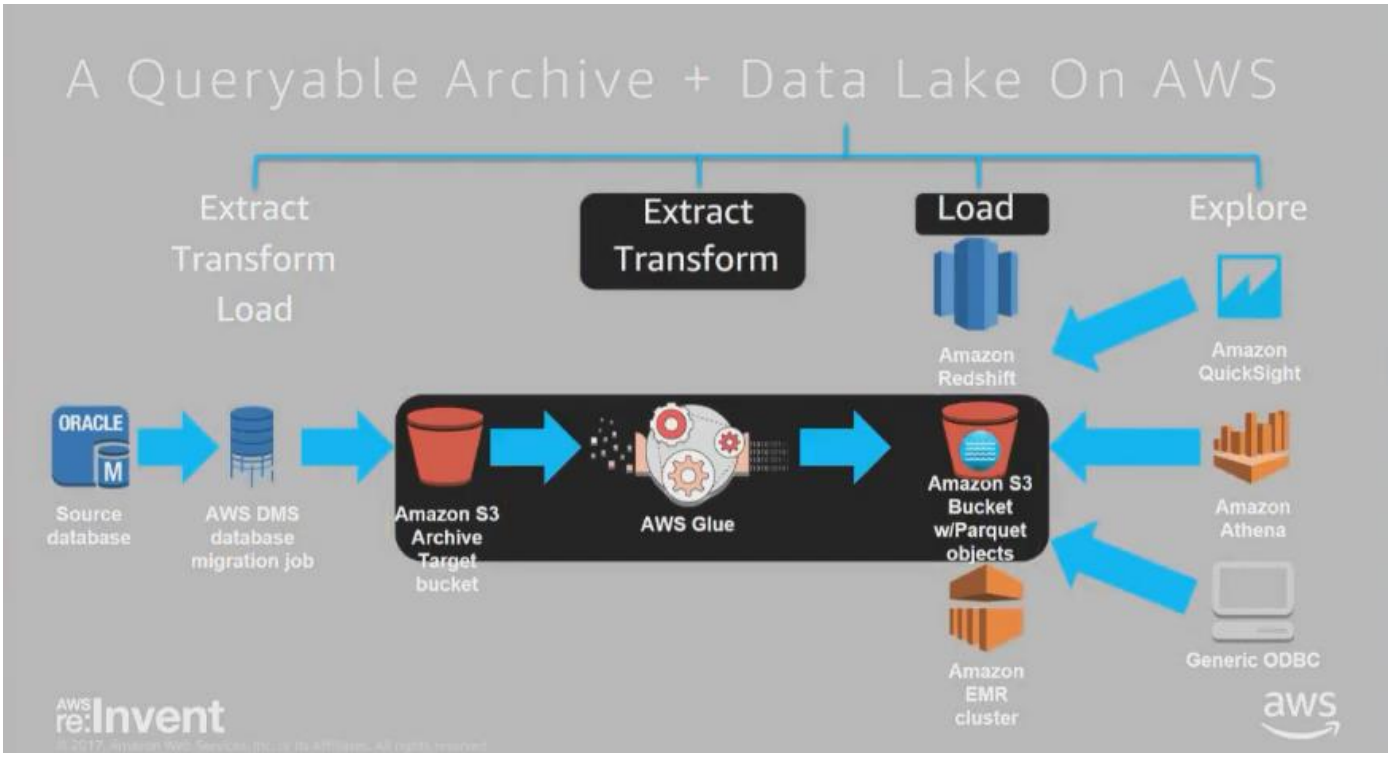

We now want to transform the compressed CSV data into the enterprise data lake format for the data model in Parquet objects. You will be transforming and enriching the data along the way

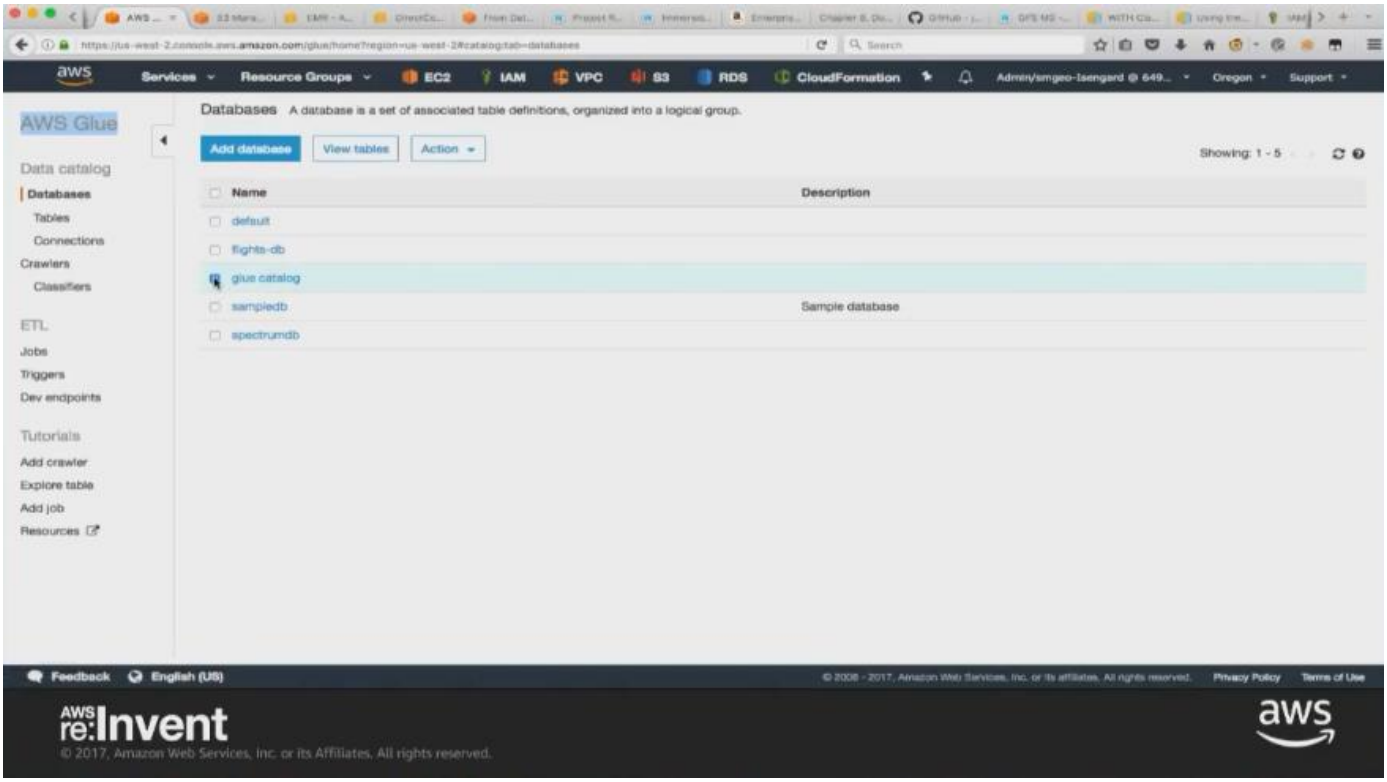

We are going to be using AWS Glue for this batch job,

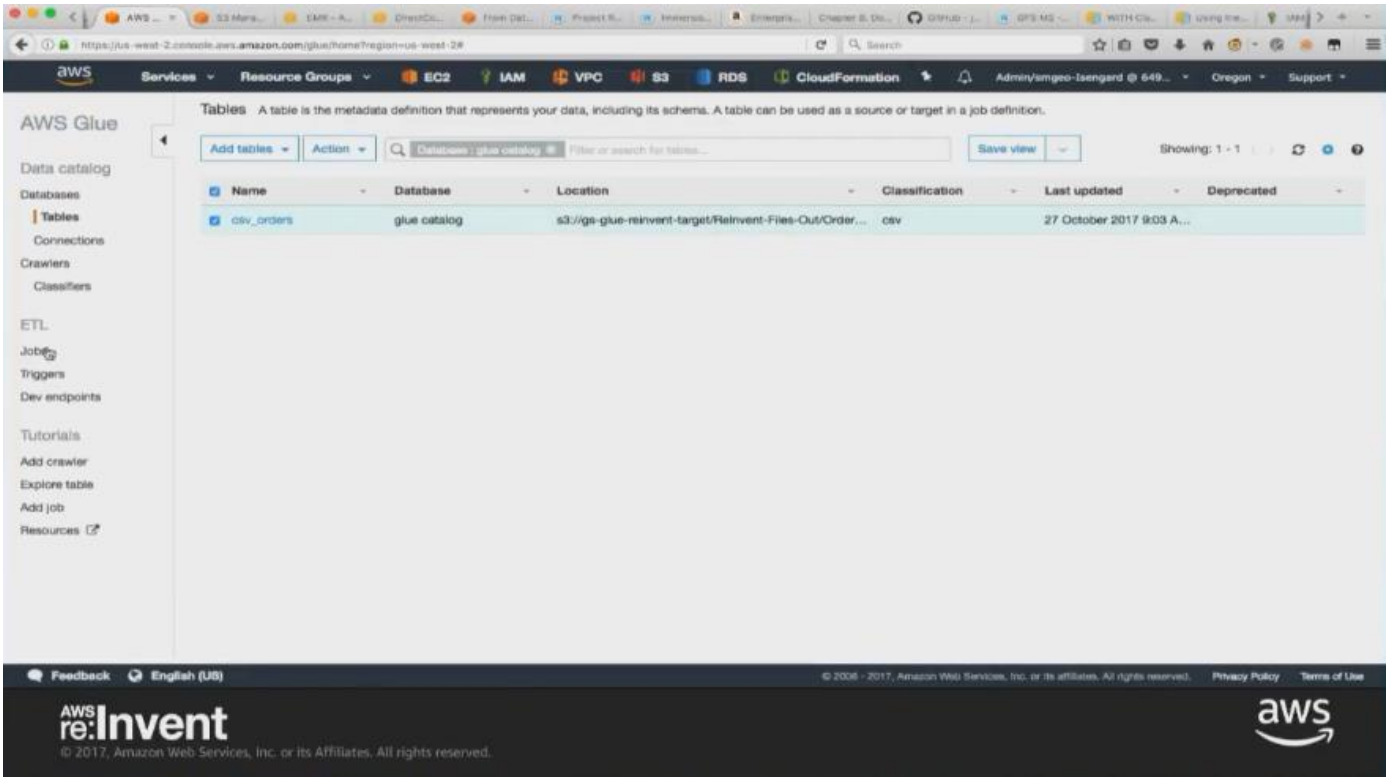

Within the glue_catalog database, we have the table containing our compressed CSV files

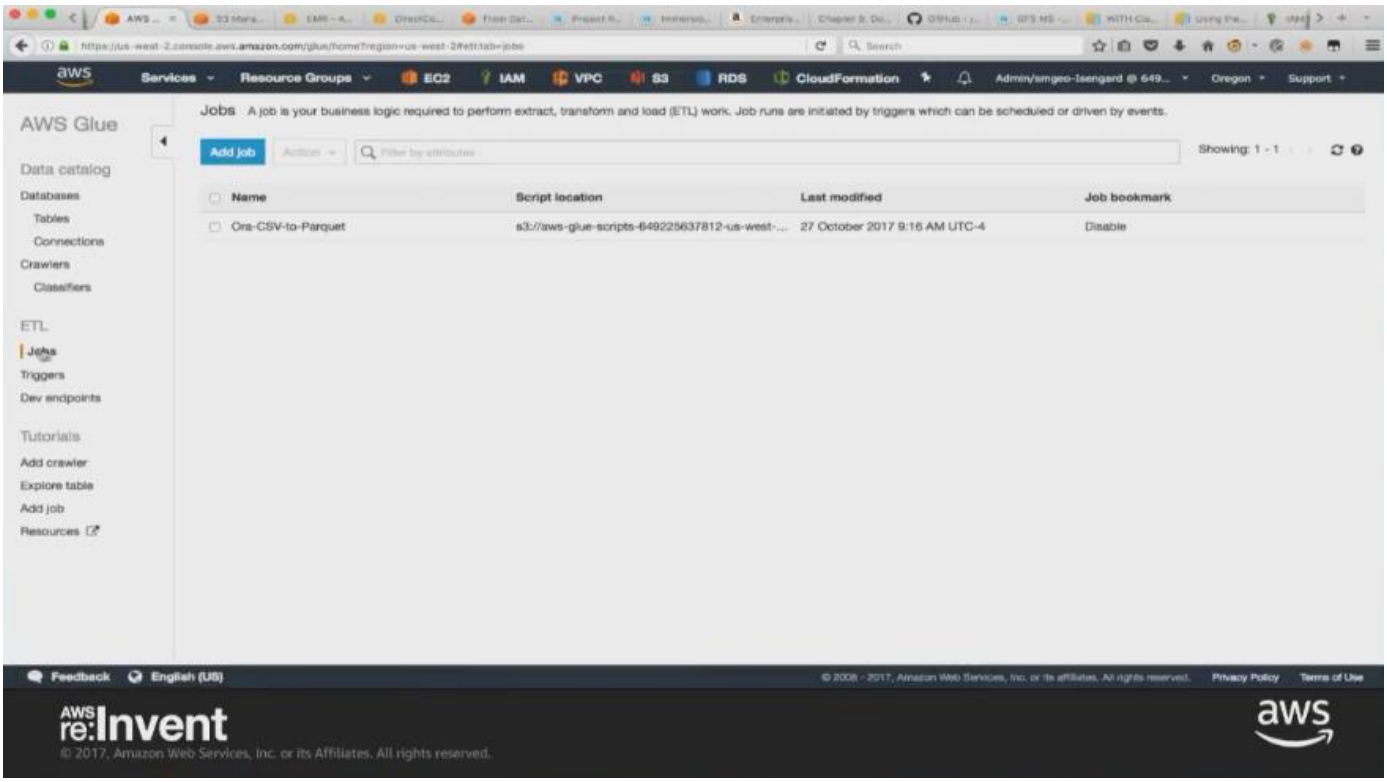
Services
Resource Groups
EC2
IAM
VPC
S3
RDS
CloudFormation
Oregon
Support
AWS Glue
Jobs A job is your business logic required to perform extract, transform and load (ETL) work. Job runs are initiated by triggers which can be scheduled or driven by events.
Add job
Action
Showing: 1 - 1
Name
Script location
Last modified
Job bookmark
Ora-CSV-to-Parquet
s3://aws-glue-scripts-649225637812-us-west-...
27 October 2017 9:16 AM UTC-4
Disable
Data catalog
Databases
Tables
Connections
Crawlers
Classifiers
ETL
Jobs
Triggers
Dev endpoints
Tutorials
Add crawler
Explore table
Add job
Resources
Feedback
English (US)
Privacy Policy
Terms of Use
re:Invent
© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

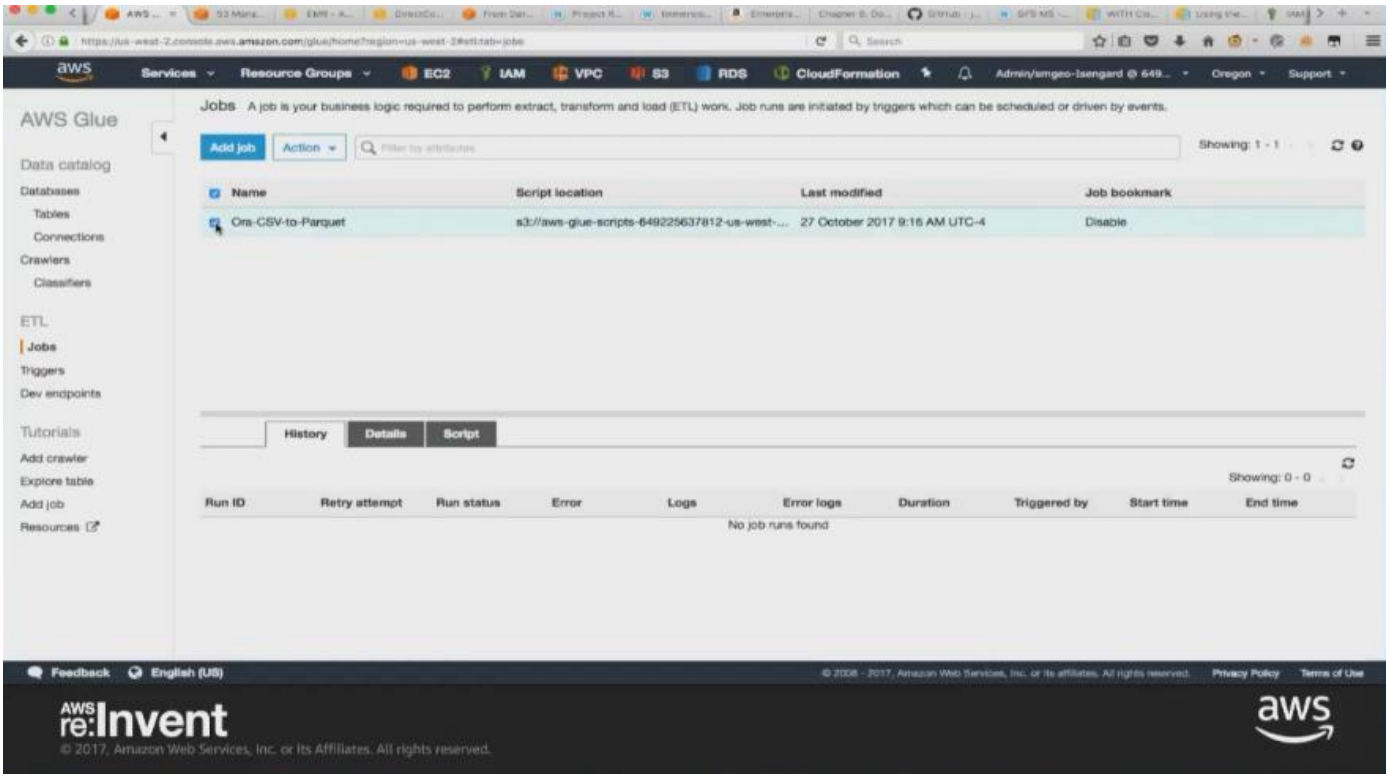
Services
Resource Groups
EC2
IAM
VPC
S3
RDS
CloudFormation
Oregon
Support
AWS Glue
Jobs A job is your business logic required to perform extract, transform and load (ETL) work. Job runs are initiated by triggers which can be scheduled or driven by events.
Add job
Action
Showing: 1 - 1
Name
Script location
Last modified
Job bookmark
Ora-CSV-to-Parquet
s3://aws-glue-scripts-649225637812-us-west-...
27 October 2017 9:16 AM UTC-4
Disable
History
Details
Script
Showing: 0 - 0
Run ID
Retry attempt
Run status
Error
Logs
Error logs
Duration
Triggered by
Start time
End time
No job runs found
Data catalog
Databases
Tables
Connections
Crawlers
Classifiers
ETL
Jobs
Triggers
Dev endpoints
Tutorials
Add crawler
Explore table
Add job
Resources
Feedback
English (US)
Privacy Policy
Terms of Use
re:Invent
© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

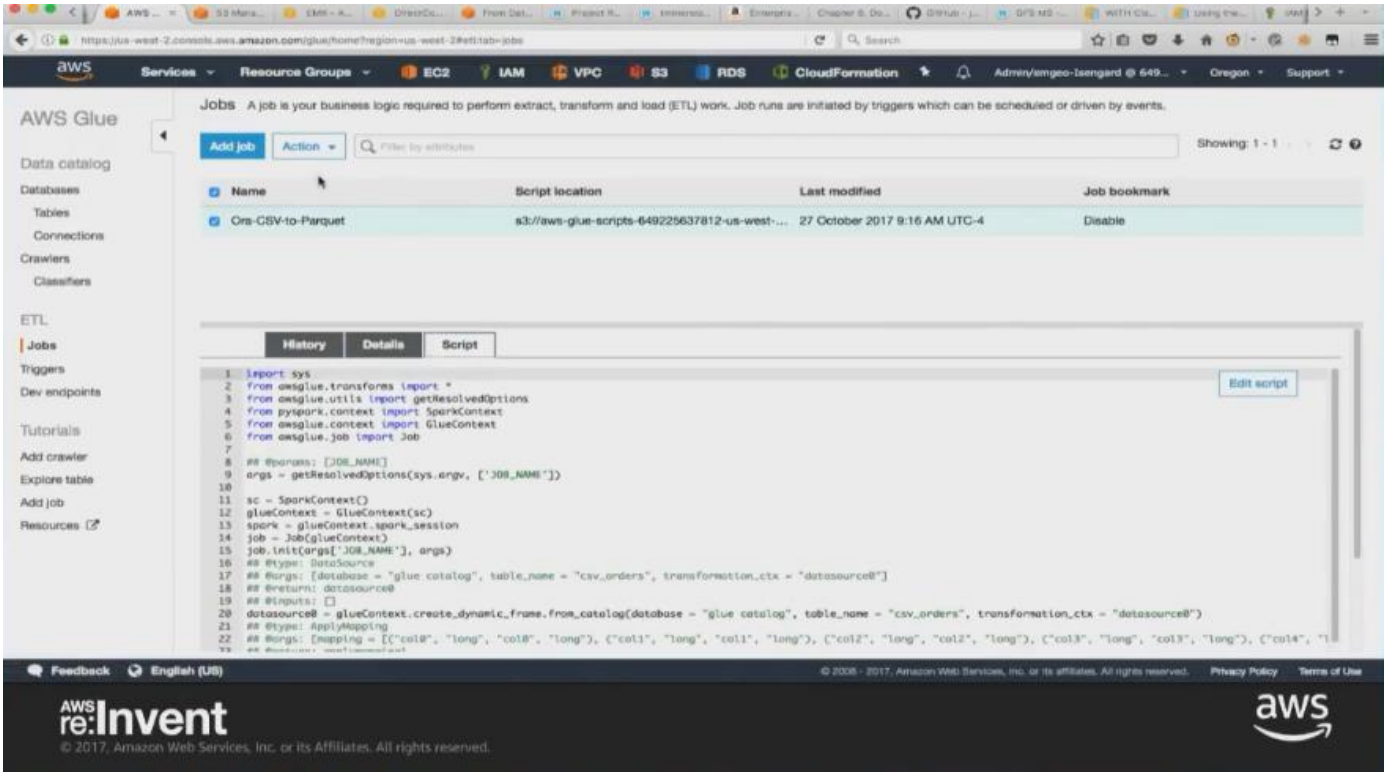

AWS Glue generate python code for the CSV_to_Parquet job for us

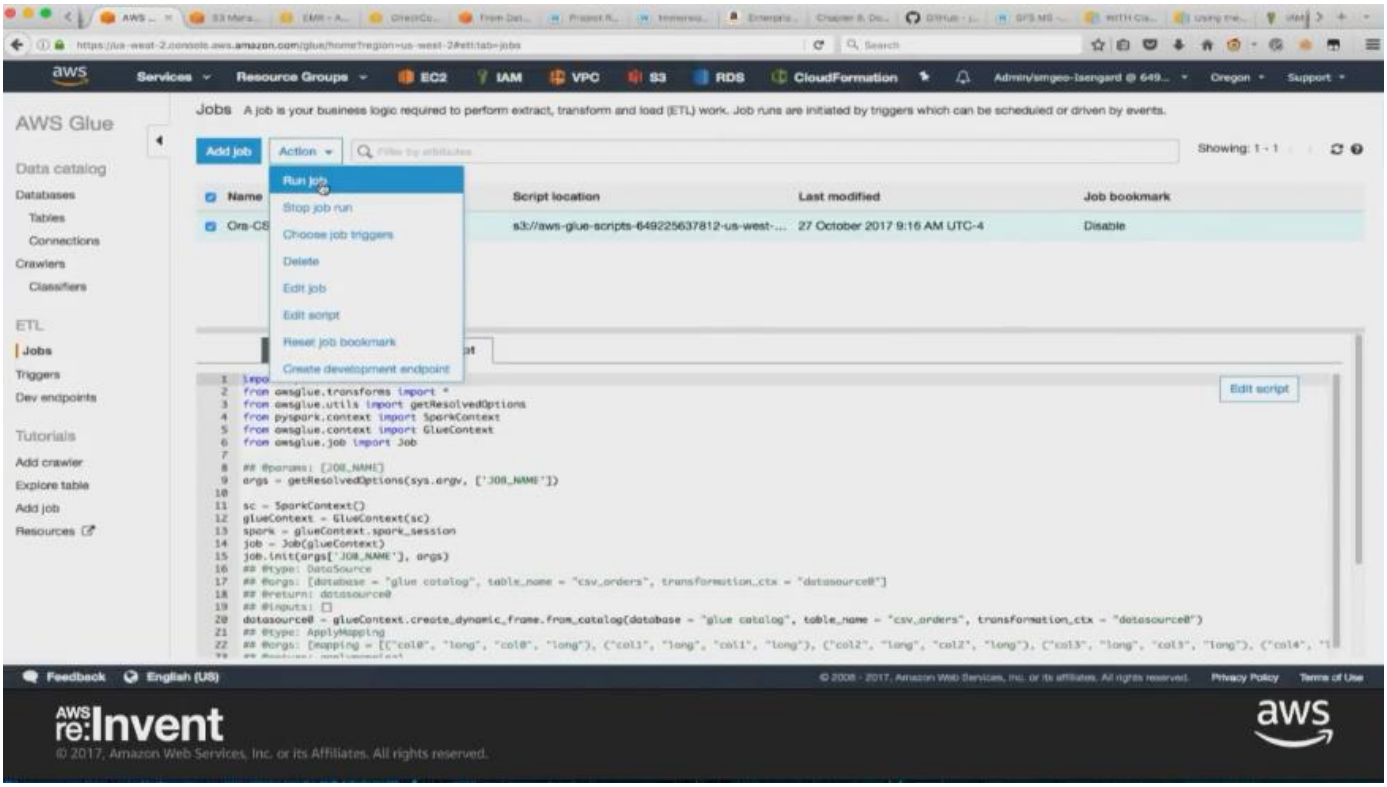

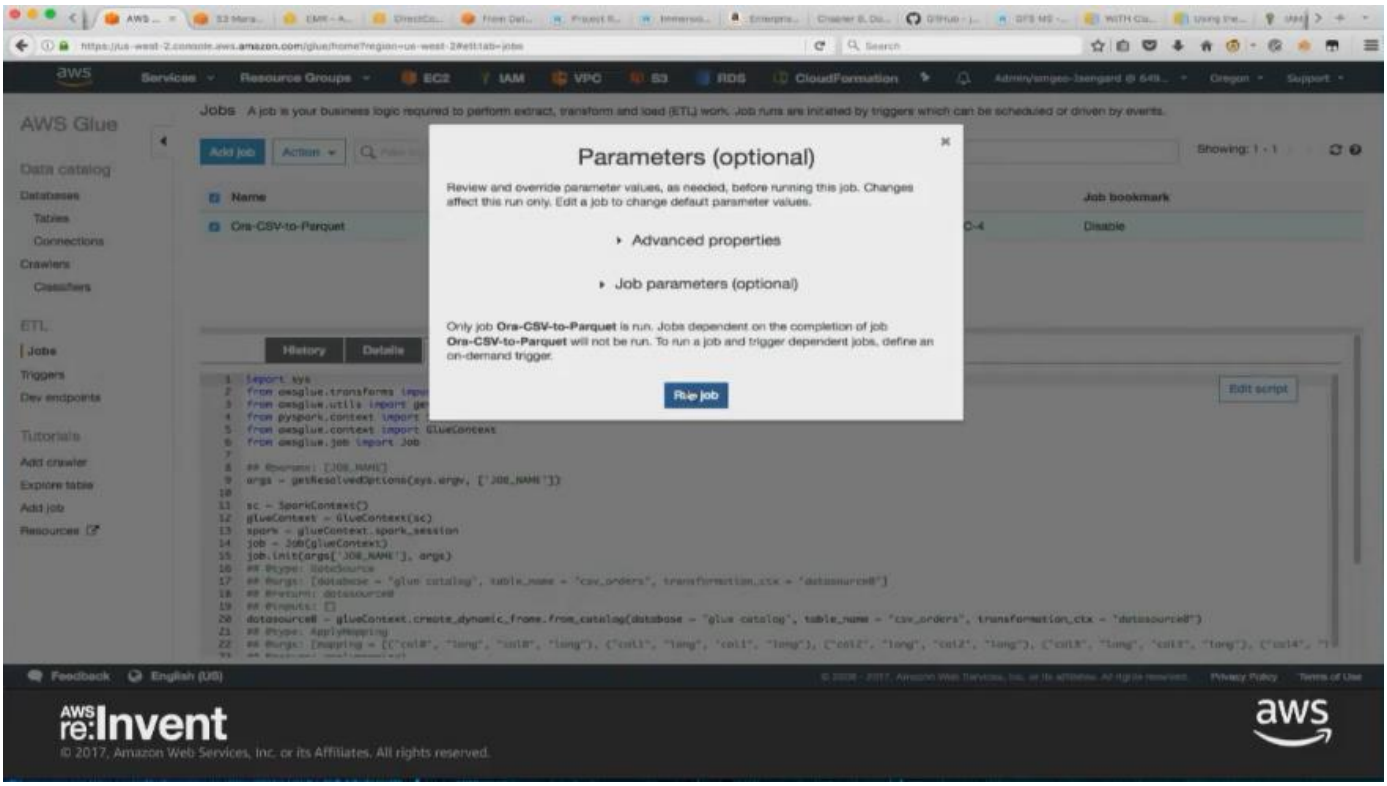
Parameters (optional)
Review and override parameter values, as needed, before running this job. Changes affect this run only. Edit a job to change default parameter values.
Advanced properties
Job parameters (optional)
Only job Ora-CSV-to-Parquet is run. Jobs dependent on the completion of job Ora-CSV-to-Parquet will not be run. To run a job and trigger dependent jobs, define an on-demand trigger.
Run job
AWS Glue
Jobs
Add job
Action
Name
Ora-CSV-to-Parquet
Job bookmark
Disable
Edit script
Feedback
English (US)
aws re:Invent
© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.
aws

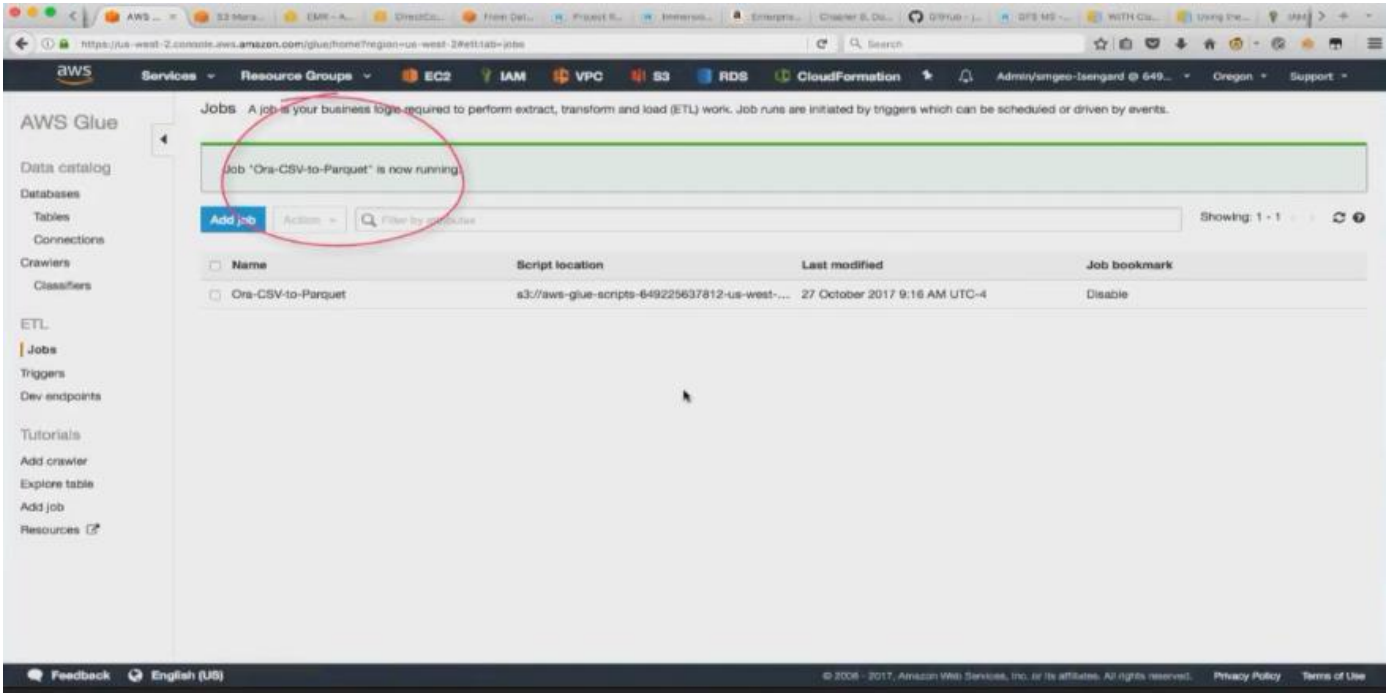
Services
Resource Groups
EC2
IAM
VPC
S3
RDS
CloudFormation
Oregon
Support
AWS Glue
Jobs A job is your business logic required to perform extract, transform and load (ETL) work. Job runs are initiated by triggers which can be scheduled or driven by events.
Job "Ora-CSV-to-Parquet" is now running
Add job
Action
Showing: 1 - 1
Name
Script location
Last modified
Job bookmark
Ora-CSV-to-Parquet
s3://aws-glue-scripts-649225637812-us-west-...
27 October 2017 9:16 AM UTC-4
Disable
Data catalog
Databases
Tables
Connections
Crawlers
Classifiers
ETL
Jobs
Triggers
Dev endpoints
Tutorials
Add crawler
Explore table
Add job
Resources
Feedback
English (US)
Privacy Policy
Terms of Use

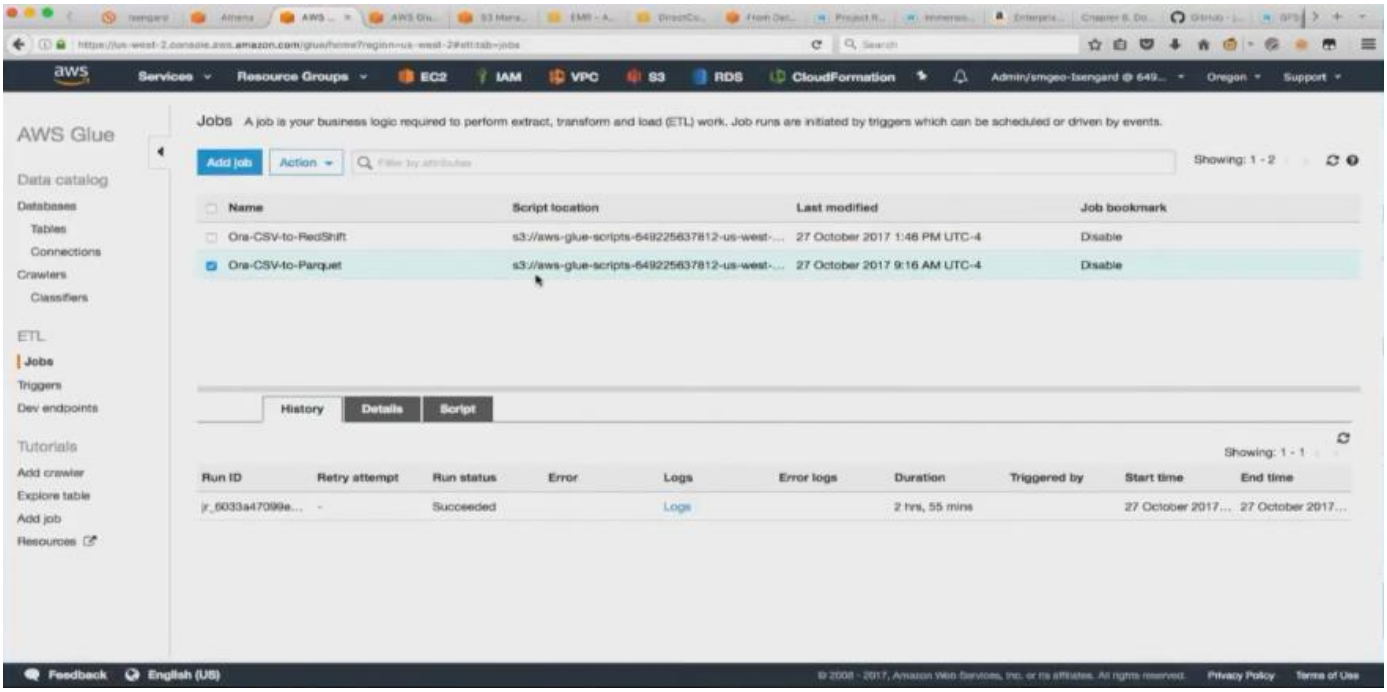

The job finishes after 2h 55mins and we have loaded all 167,497,743 rows into our data lake.

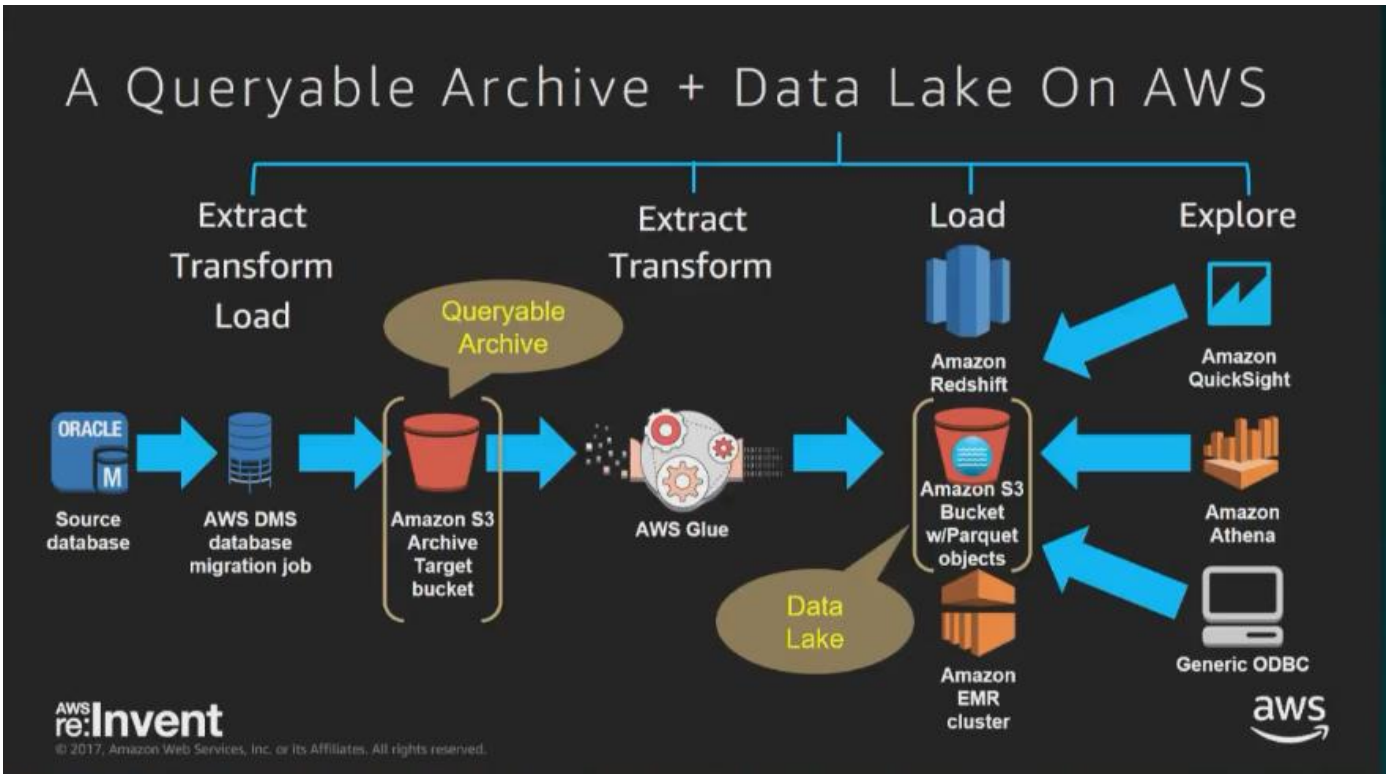

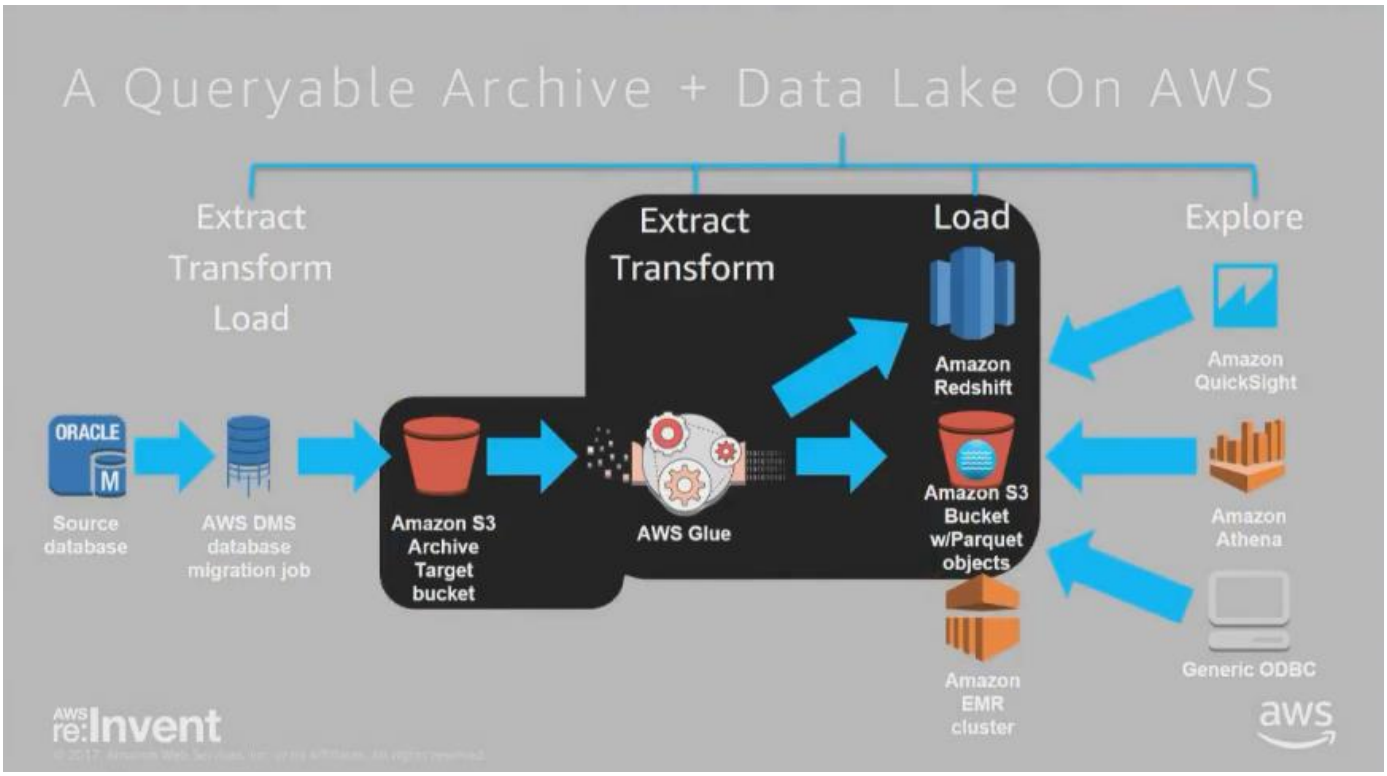

Now we have data in 2 locations.

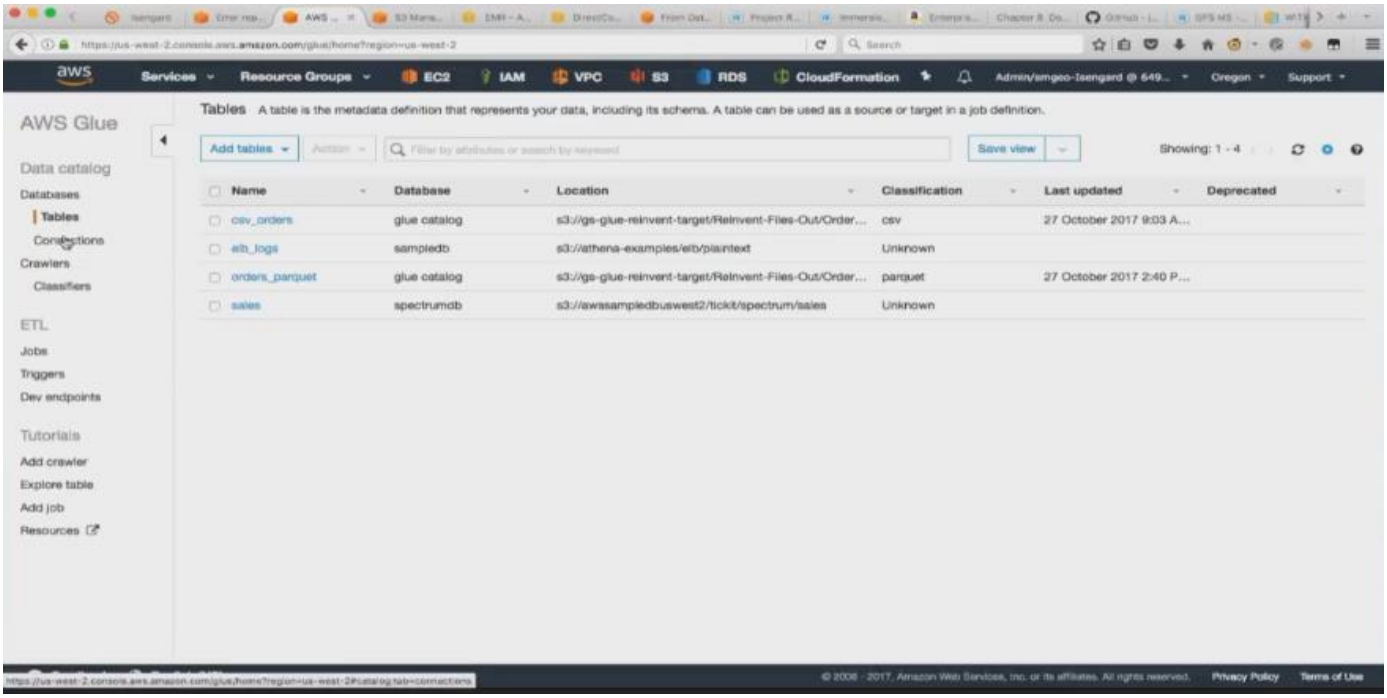

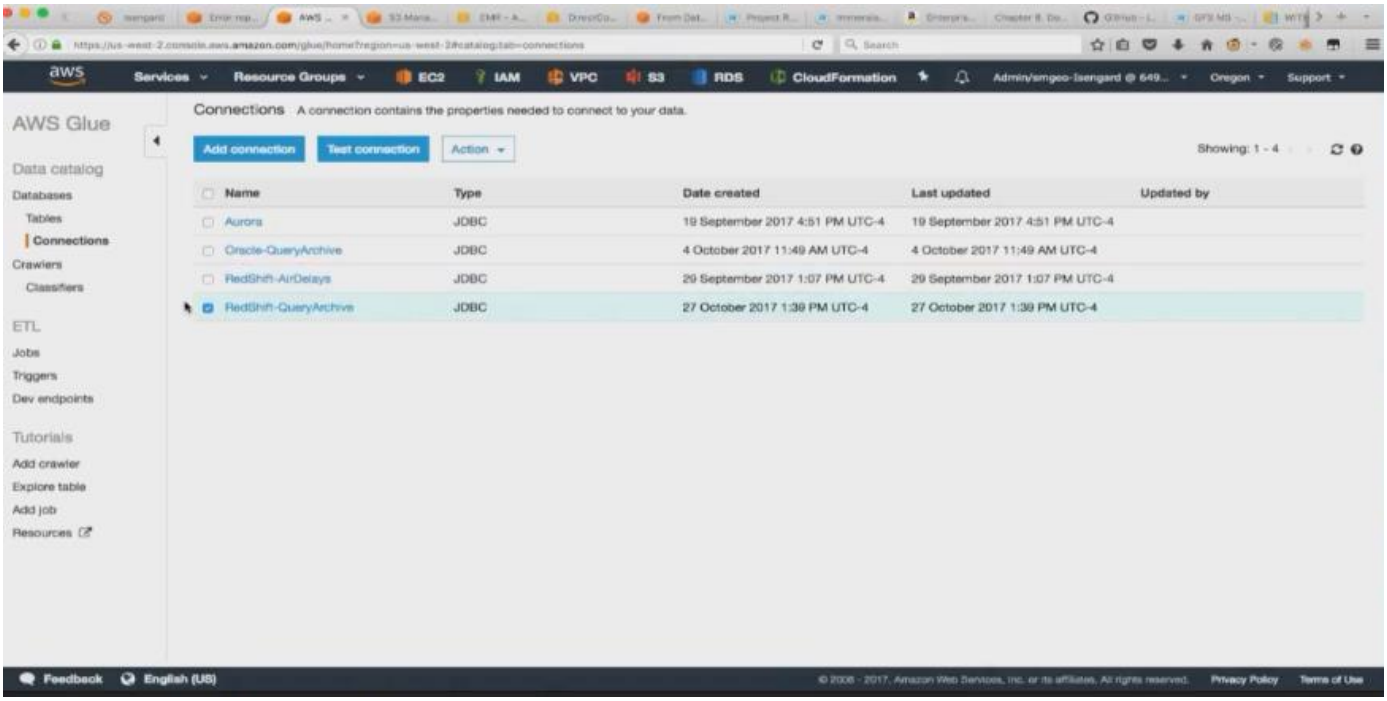

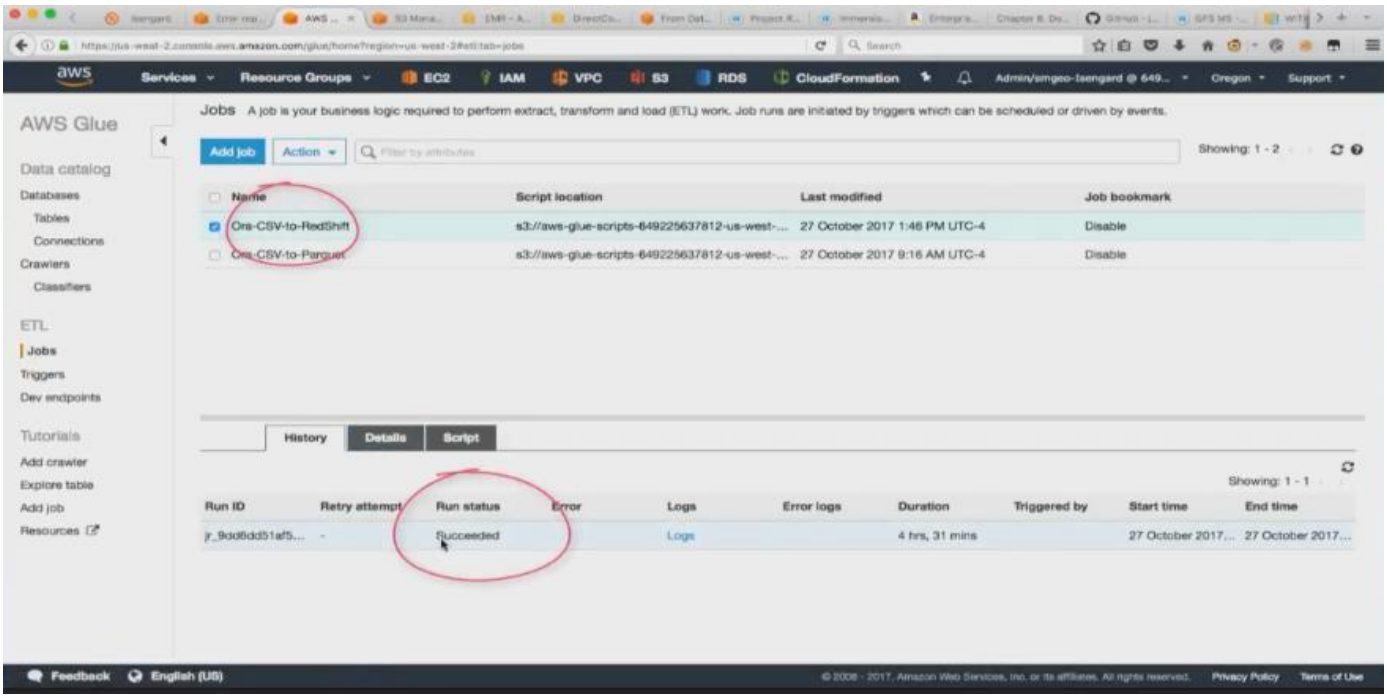

We also load data into Redshift

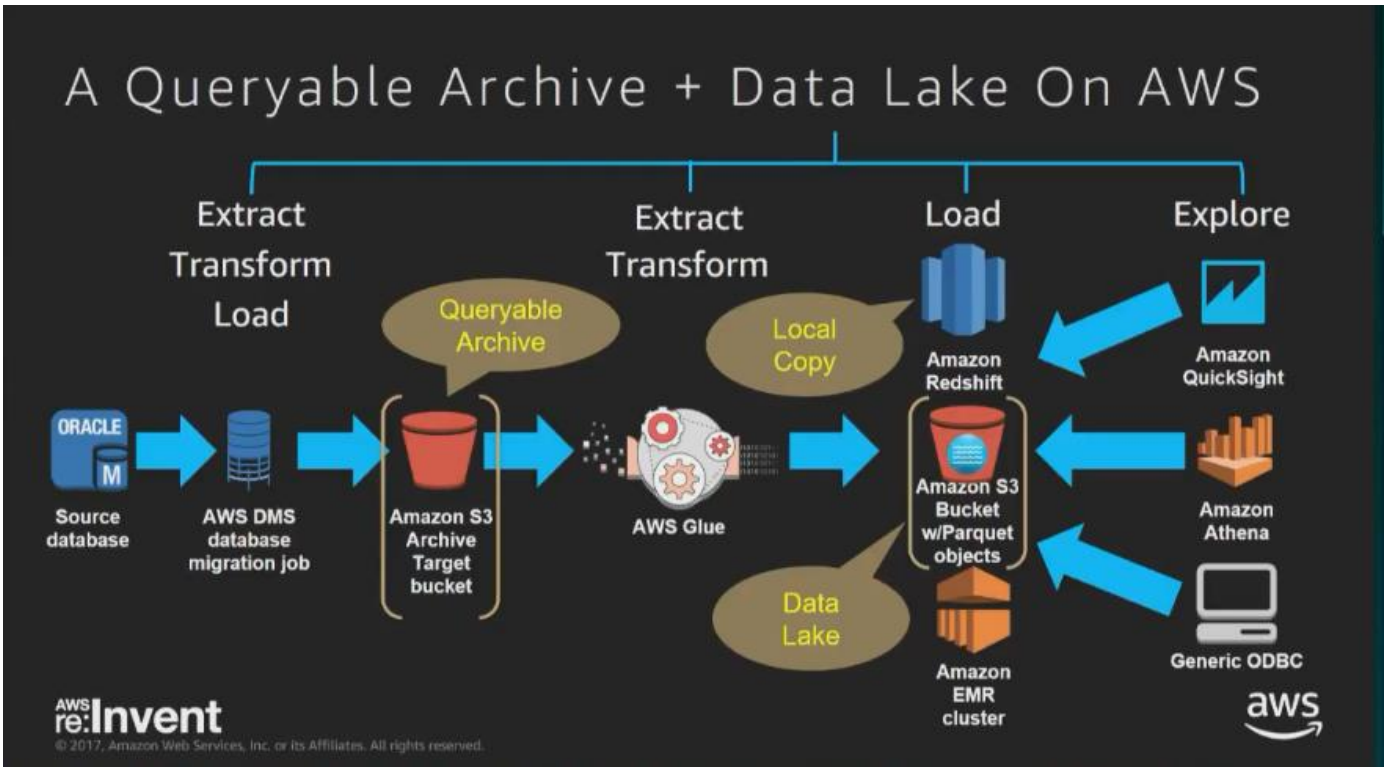
A Queryable Archive + Data Lake On AWS
Extract Transform Load
Extract Transform
Load
Explore
Queryable Archive
Local Copy
Data Lake
Source database
AWS DMS database migration job
Amazon S3 Archive Target bucket
AWS Glue
Amazon Redshift
Amazon S3 Bucket w/Parquet objects
Amazon EMR cluster
Amazon QuickSight
Amazon Athena
Generic ODBC
aws re:Invent
© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.
aws

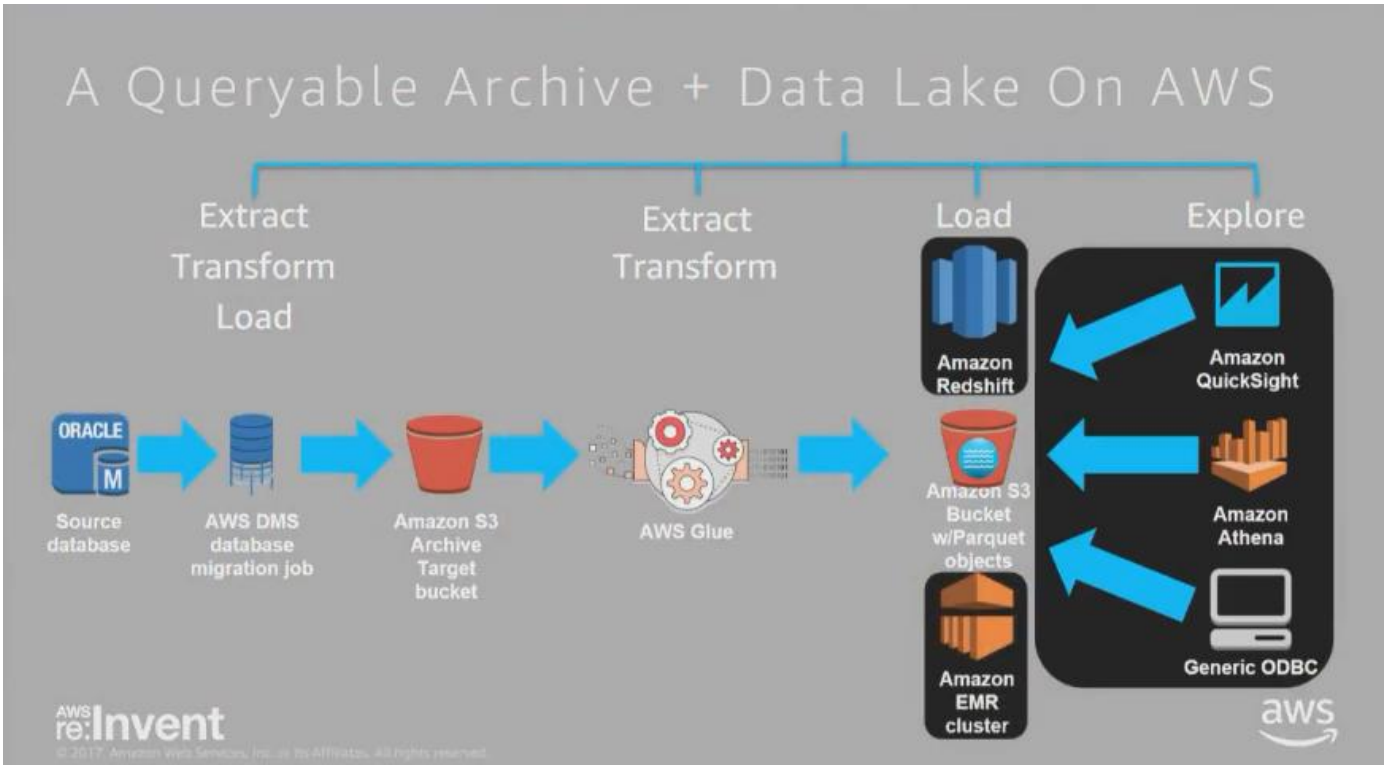
A Queryable Archive + Data Lake On AWS
Extract Transform Load
Extract Transform
Load
Explore
Source database
AWS DMS database migration job
Amazon S3 Archive Target bucket
AWS Glue
Amazon Redshift
Amazon S3 Bucket w/Parquet objects
Amazon EMR cluster
Amazon QuickSight
Amazon Athena
Generic ODBC
aws re:Invent
aws

```
f45c89bd7fe5:~ smgeo$ aws emr create-cluster \
> --termination-protected \
> --applications Name=Hadoop Name=Hive Name=Pig Name=Hue Name=Spark Name=Presto \
> --tags 'name=queryarchive' \
> --ec2-attributes '{"KeyName":"MyEC2Key","InstanceProfile":"EMR_EC2_DefaultRole","SubnetId":"subnet-17b0fa73","EmrManagedSlaveSecurityGroup":"sg-ea673093","EmrManagedMasterSecurityGroup":"sg-eb673092"}' \
> --service-role EMR_DefaultRole \
> --enable-debugging \
> --release-label emr-5.9.0 \
> --log-uri 's3n://aws-logs-649225637812-us-west-2/elasticmapreduce/' \
> --name 'QueryArchive Cluster Glue' \
> --instance-groups '[{"InstanceCount":1,"InstanceGroupType":"MASTER","InstanceType":"m3.xlarge","Name":"Master - 1"},{"InstanceCount":2,"InstanceGroupType":"CORE","InstanceType":"r3.2xlarge","Name":"Core - 2"}]' \
> --region us-west-2 \
> --configurations file://glueConfiguration.json
{
    "ClusterId": "j-3JJX71R6GHGVP"
}
f45c89bd7fe5:~ smgeo$
```

```
f45c89bd7fe5:~ smgeo$ aws emr create-cluster \
> --termination-protected \
> --applications Name=Hadoop Name=Hive Name=Pig Name=Hue Name=Spark Name=Presto \
> --tags 'name=queryarchive' \
> --ec2-attributes '{"KeyName":"MyEC2Key","InstanceProfile":"EMR_EC2_DefaultRole","SubnetId":"subnet-17b0fa73","EmrManagedSlaveSecurityGroup":"sg-ea673093","EmrManagedMasterSecurityGroup":"sg-eb673092"}' \
> --service-role EMR_DefaultRole \
> --enable-debugging \
> --release-label emr-5.9.0 \
> --log-uri 's3n://aws-logs-649225637812-us-west-2/elasticmapreduce/' \
> --name 'QueryArchive Cluster Glue' \
> --instance-groups '[{"InstanceCount":1,"InstanceGroupType":"MASTER","InstanceType":"m3.xlarge","Name":"Master - 1"},{"InstanceCount":2,"InstanceGroupType":"CORE","InstanceType":"r3.2xlarge","Name":"Core - 2"}]' \
> --region us-west-2 \
> --configurations file://glueConfiguration.json
{
    "ClusterId": "j-3JJX71R6GHGVP"
}
f45c89bd7fe5:~ smgeo$ more glueConfiguration.json
[
  {
    "Classification": "hive-site",
    "Properties": {
      "hive.metastore.client.factory.class": "com.amazonaws.glue.catalog.metastore.AWSGlueDataCatalogHiveClientFactory"
    }
  },
{
    "Classification": "spark-hive-site",
    "Properties": {
      "hive.metastore.client.factory.class": "com.amazonaws.glue.catalog.metastore.AWSGlueDataCatalogHiveClientFactory"
    }
  },
{
    "Classification": "spark-log4j",
    "Properties": {
        "log4j.rootCategory": "WARN, console"
        }
  }
]

f45c89bd7fe5:~ smgeo$
```


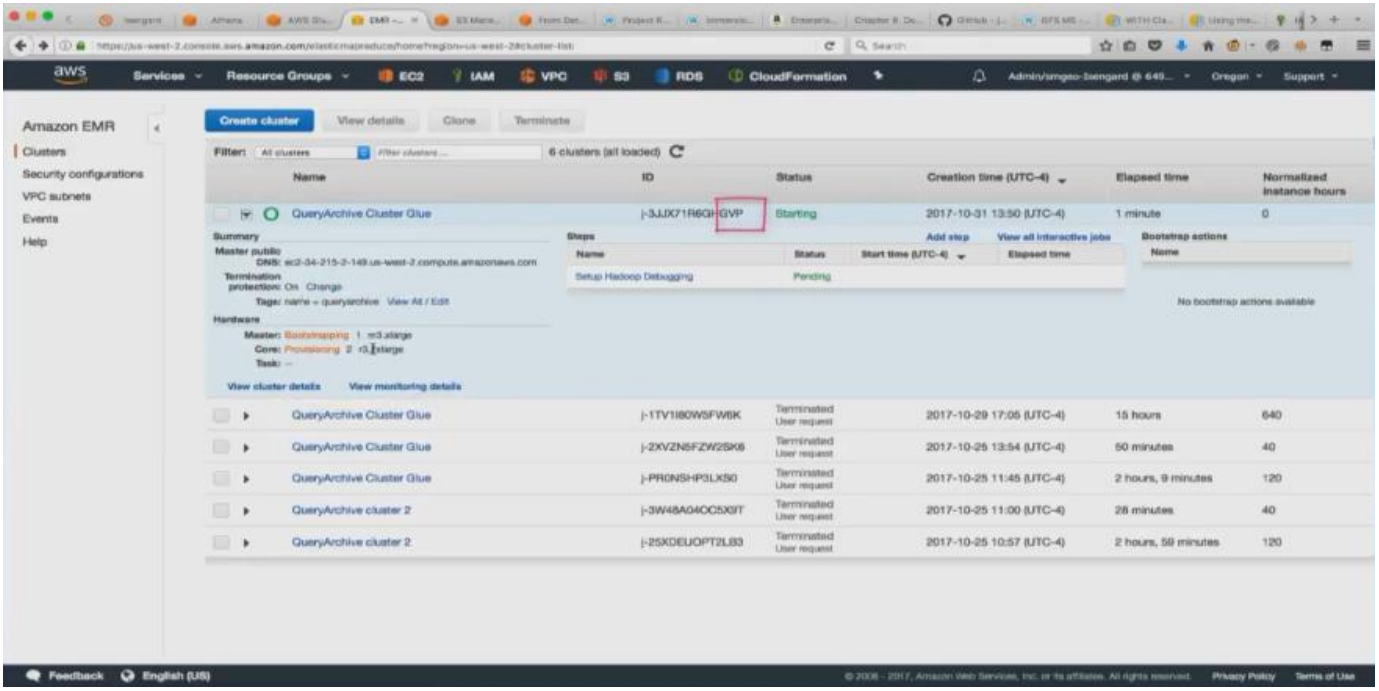

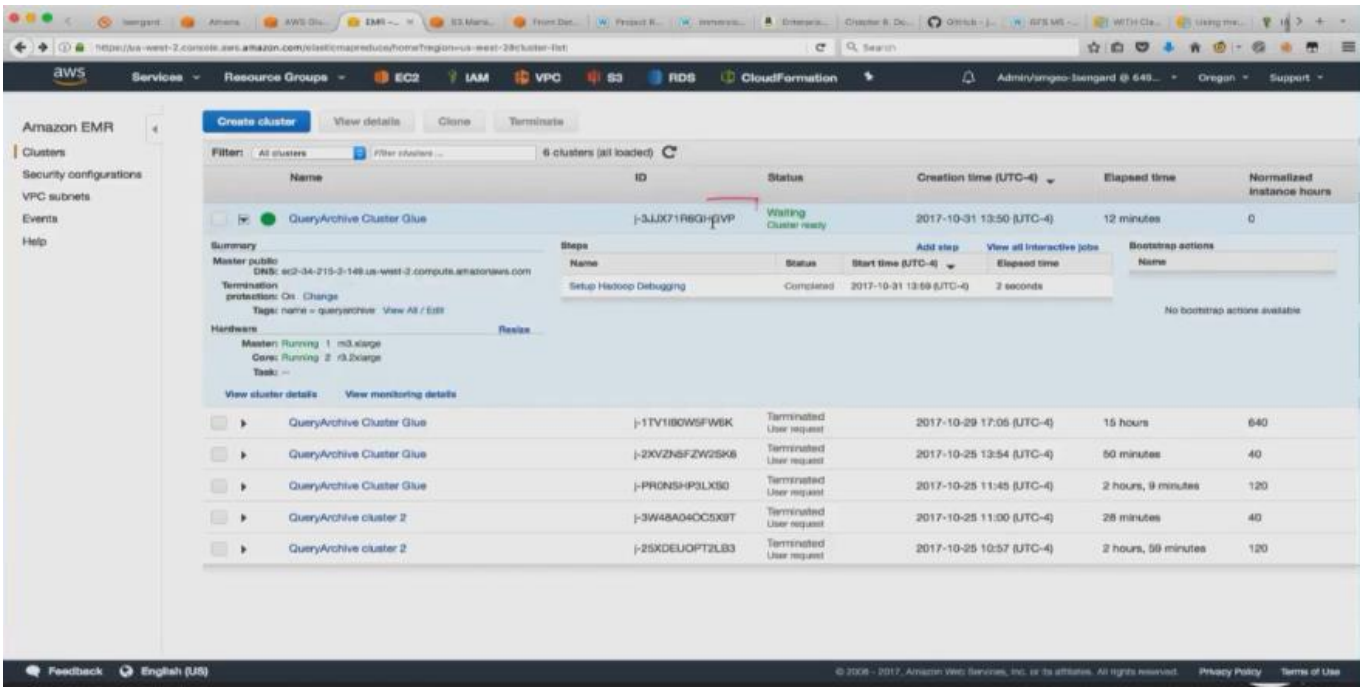

Benefits

- Single golden source of truth
- Right sizing compute for workload
- Increased availability

A Queryable Archive + Data Lake On AWS

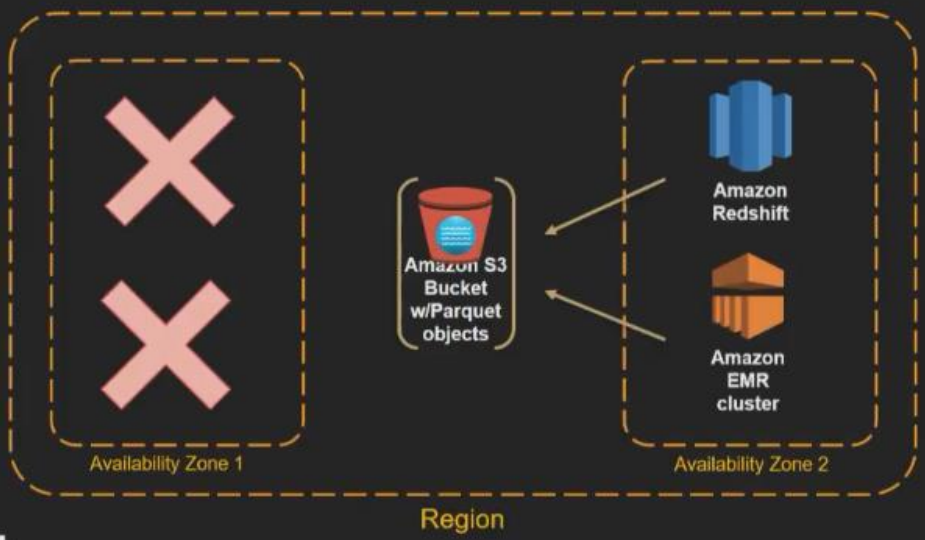

Storage Efficiencies

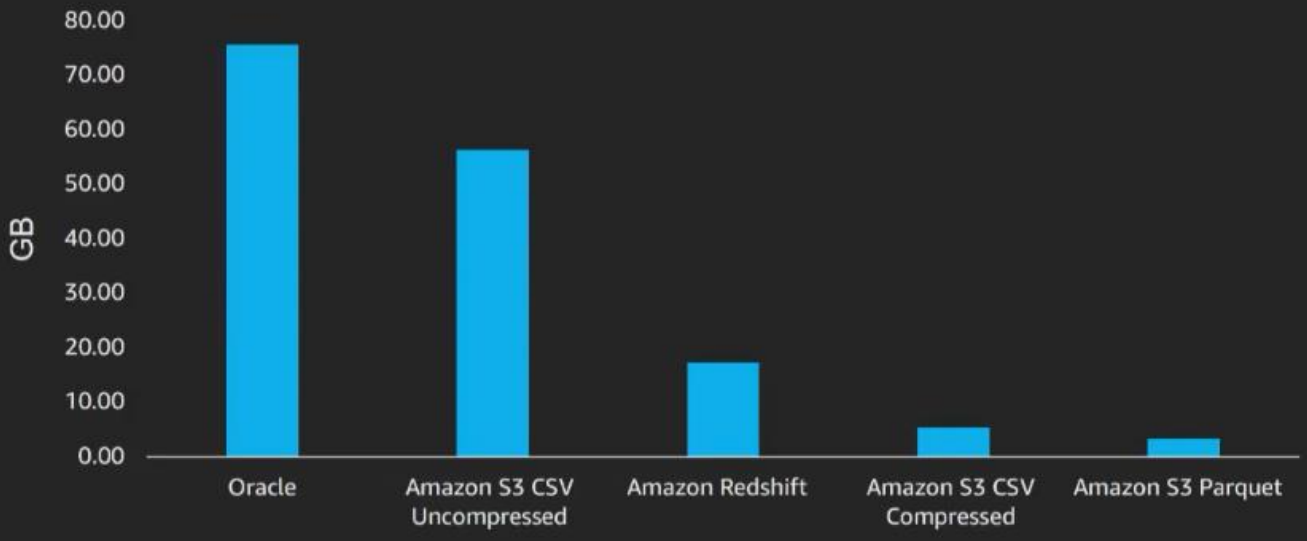

Compute Efficiencies

Storage Format	Size in GB	Size as a percent of largest store*	Costs Per Month
RDS Oracle	75.70	100.00%	$1449.78 / $706.69
Amazon S3 CSV Uncompressed	56.40	74.50%	$1.30
Amazon Redshift	17.30	22.85%	$366.00
Amazon S3 CSV Compressed	5.43	7.17%	$0.13
Amazon S3 Parquet	3.40	4.49%	$0.08

Data Efficiencies

Storage Tier	Monthly Storage Pricing per PB
Amazon S3	$22,583.30
Amazon S3-IA	$13,107.20
Amazon Glacier	$4,194.31

Unleashing "Dark Data"

What is Dark Data?

"Buried within raw information generated in mind-boggling volumes by transactional systems . . . are critical strategic, customer, and operational insights that, once illuminated by analytics, can validate or clarify assumptions, inform decision making, and help chart new paths to the future."

Kambies, T., Roma, P., Mittal, N., & Sharma, S. K. (2017, February 7). Dark analytics: Illuminating opportunities hidden within unstructured data. Retrieved October 16, 2017, from https://dupress.deloitte.com/dup-us-en/focus/tech-trends/2017/dark-data-analyzing-unstructured-data.html

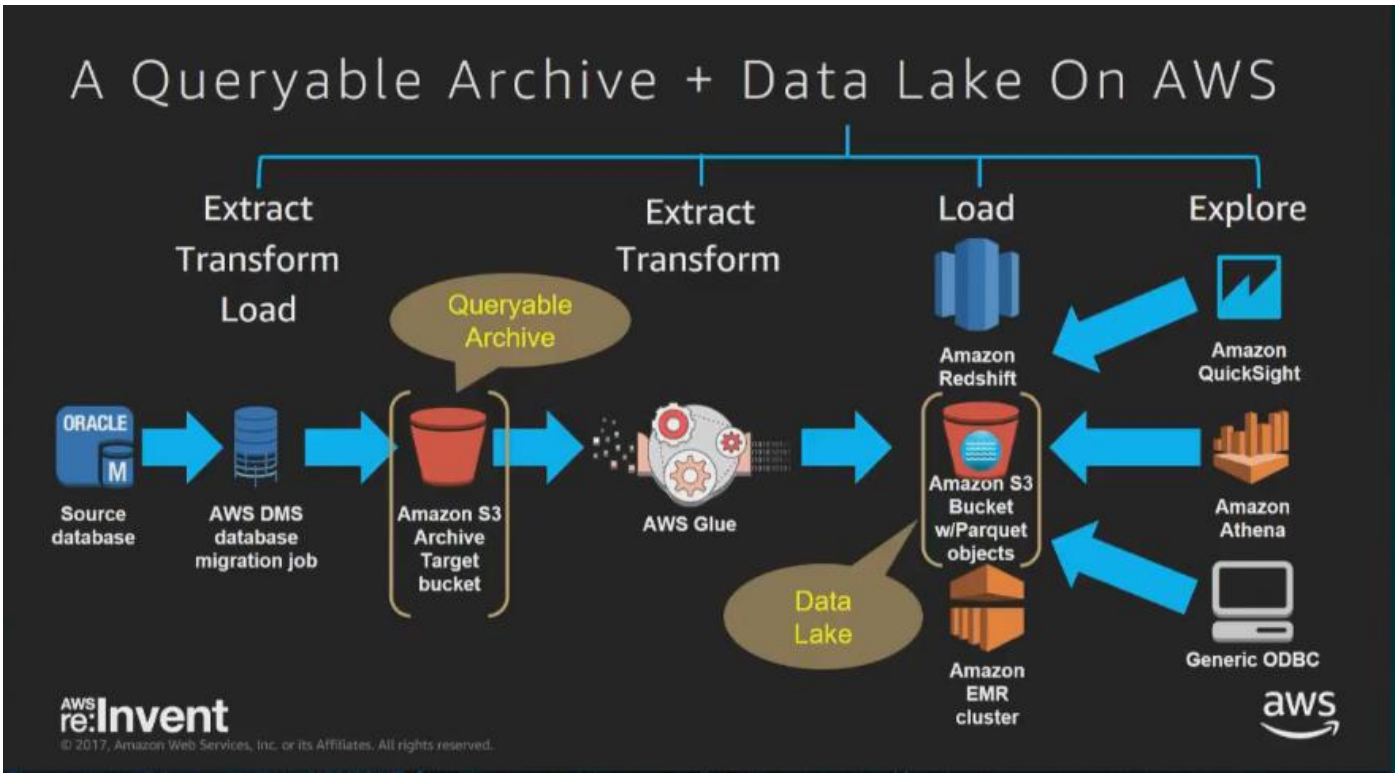
A Queryable Archive + Data Lake On AWS
Extract
Transform
Load
Extract
Transform
Load
Explore
Queryable
Archive
Amazon
Redshift
Amazon
QuickSight
ORACLE
Source
database
AWS DMS
database
migration job
Amazon S3
Archive
Target
bucket
AWS Glue
Amazon S3
Bucket
w/Parquet
objects
Amazon
Athena
Data
Lake
Amazon
EMR
cluster
Generic ODBC
aws
re:Invent
© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

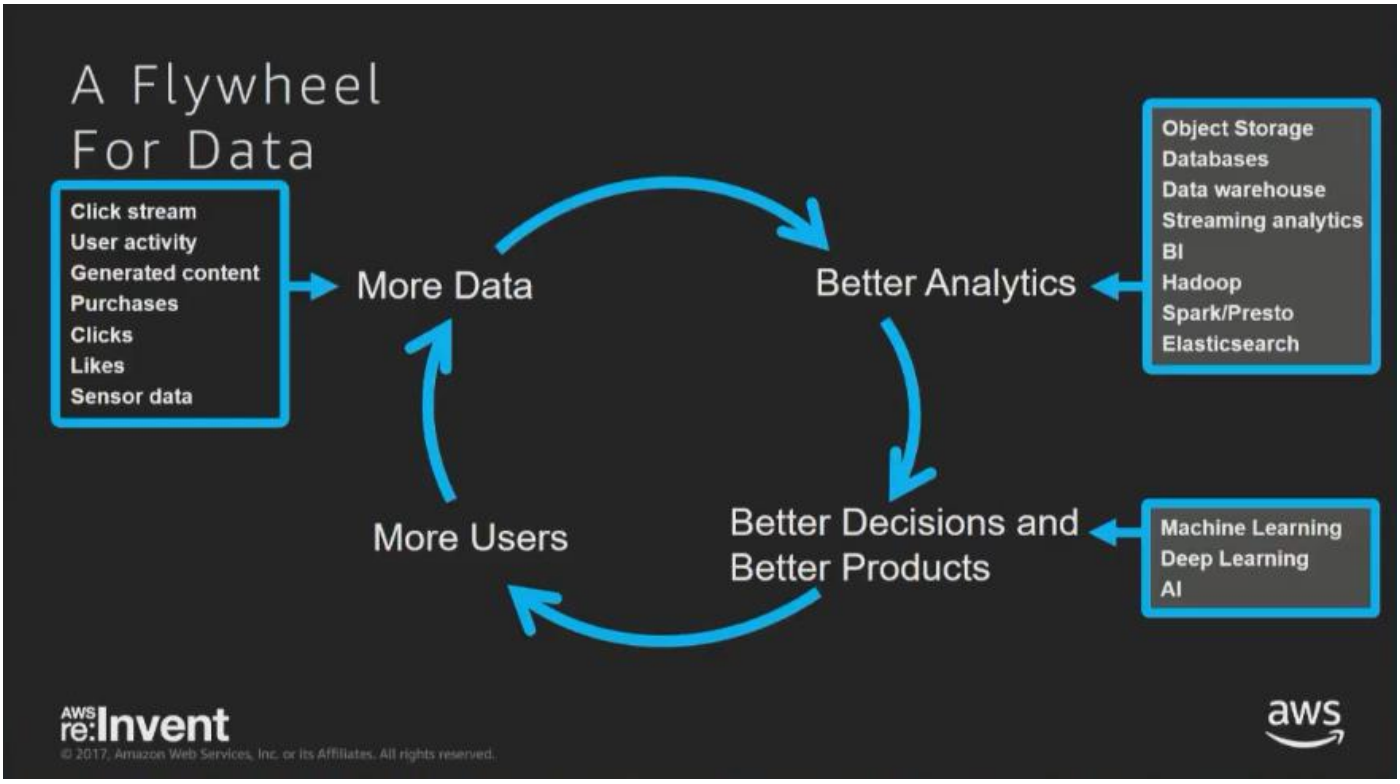
A Flywheel
For Data
Click stream
User activity
Generated content
Purchases
Clicks
Likes
Sensor data
More Data
Better Analytics
Object Storage
Databases
Data warehouse
Streaming analytics
BI
Hadoop
Spark/Presto
Elasticsearch
More Users
Better Decisions and
Better Products
Machine Learning
Deep Learning
AI
aws
re:Invent
© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.


Next Steps

What to Expect From This Session

- ✓ A pattern for better, cheaper, faster archives called "Queryable Archive"
- ✓ An implementation of "Queryable Archive + Data Lake" on AWS
- ✓ Why bringing archived data online exposes "dark data"

What you can do next?

- Go back and determine the cost to your company of your archive data stores
- What would it cost if you moved them into AWS with this solution?
  - Estimate AWS cost with the simple monthly calculator: https://calculator.s3.amazonaws.com/index.html
- Ask yourself, can this increase my organization's efficiency?
- Go build it!

AWS re:Invent

aws