

ATC303

# AWS re:INVENT Cache Me If You Can

Minimizing Latency While Optimizing Cost  
Through Advanced Caching Strategies

Markus Ostertag, VP Engineering, Team Internet, @osterjour  
Constantin Gonzalez, Principal Solutions Architect, AWS, @zalez

November 27, 2017

AWS  
re:Invent


© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



From CloudFront to ElastiCache to DynamoDB Accelerator (DAX), this is your one-stop shop for learning how to apply caching methods to your AdTech workload: What data to cache and why? What are common side effects and pitfalls when caching? What is negative caching and how can it help you maximize your cache hit rate? How to use DynamoDB Accelerator in practice? How can you ensure that data always stays current in your cache? These and many more topics will be discussed in depth during this talk and we'll share lessons learned from Team Internet, the leading provider in domain monetization.

## What you'll get out of this session

- Speed!
- Best practices
- Architecture patterns
- Real-world examples from Team Internet
- Lower cost



AWS  
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



This session is all about getting more speed into your applications.

# AdTech

- Publishers, advertisers
- Bidding
- Limited time
- Hard cutoff
- Typical: x00 ms

long live  
bacon

AWS  
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

aws

Photo by [Antonio Barrero](#) on [Unsplash](#)

It's all about speed!

Three ways to become faster

1. Increase rate

2. Parallelize

3. Do less

Complex

Complex

Easy!

To get speed into your AdTech applications and services, you can increase the rate at which you are running your transactions to minimize time taken, you can parallelize your application, or use caching

Three ways to become faster

1. Increase rate

2. Parallelize

3. Do less ← Caching

Caching

- Do time-consuming stuff **once**
- Re-use results **many times**

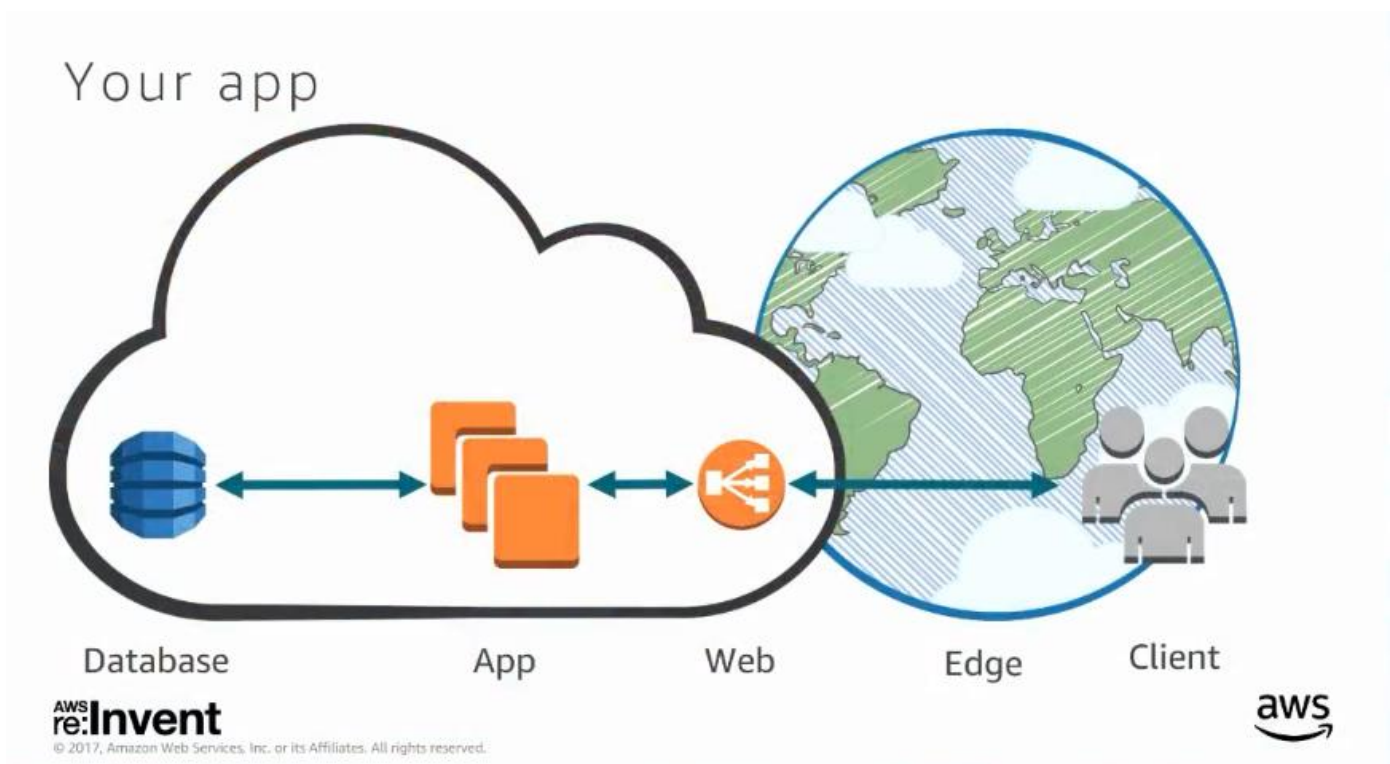
Because memory is **cheaper and faster**  
than CPU

You can get your transactions computed and then you save the results, it's a lot easier to retrieve those results instead of computing them again






Using caching at each of these layers will save you time and money



This is what most applications look like,

## Edge caching



AWS  
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



**Edge caching** is your first opportunity to become faster and add some speed to your applications using CloudFront

## Amazon CloudFront

- 107 Edge locations
- 55 cities
- 24 countries
- Static and **dynamic content**
- Optimized last-mile connectivity
- Lower cost than directly from region
- PCI DSS compliant



AWS  
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



Markus Ostertag

VP of Engineering, Team Internet AG

# Who is Team Internet?

- Domain monetization business
- 35 people
- HQ in Munich, Germany
- Tech focused



© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



## TONIC.



# The challenge

- Need to support clients worldwide
- Need latency of <300 ms
- Need 100% consistency at the DB level

Should we go multi-region?  
Or is there an alternative?




AWS  
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

aws

## Multiple regions



AWS  
re:Invent

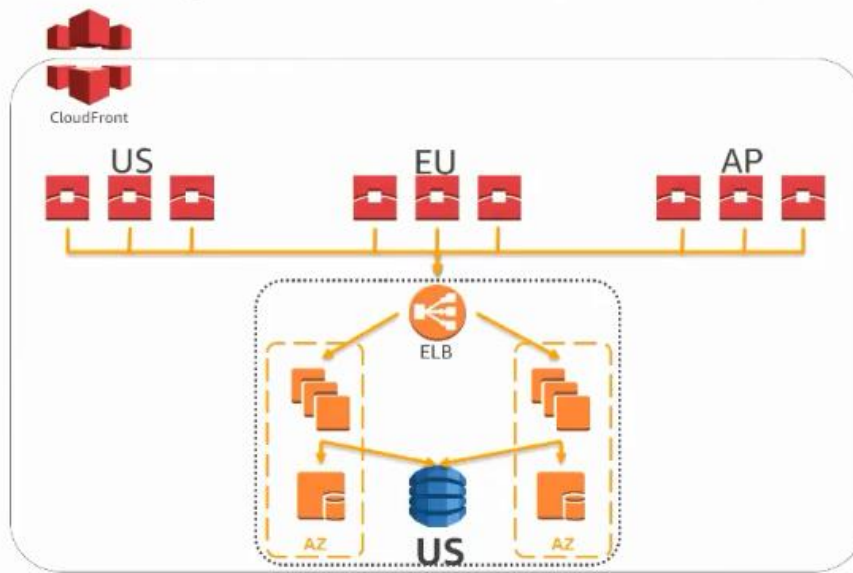
© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

aws

This was the architecture we had in mind before consulting AWS



# CloudFront dynamic content acceleration



**AWS re:Invent**

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



This is what was suggested to us for the real time bidding system architecture, CloudFront will help optimize the latency to our customers, and we only work out of only the US region and having CloudFront be the overlay for us around the world and help optimize the last mile to our customers.

## Lambda@Edge

Available in all  
Amazon CloudFront  
edge locations

- Low-latency request/response customization
- Supports viewer and origin events



**AWS re:Invent**


© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.




You can now modify the content at the edge by using lambdas at specific points in the request/response path as below




# CloudFront triggers for Lambda@Edge functions




# CloudFront triggers for Lambda@Edge functions



# CloudFront triggers for Lambda@Edge functions



# CloudFront triggers for Lambda@Edge functions



# Lambda@Edge use cases



Content  
customization



Visitor  
validation



A/B testing

## Edge caching recap

- Use CloudFront for reducing last-mile latency
- Almost always a win!
- Can save cost, too
- Consider Lambda@Edge



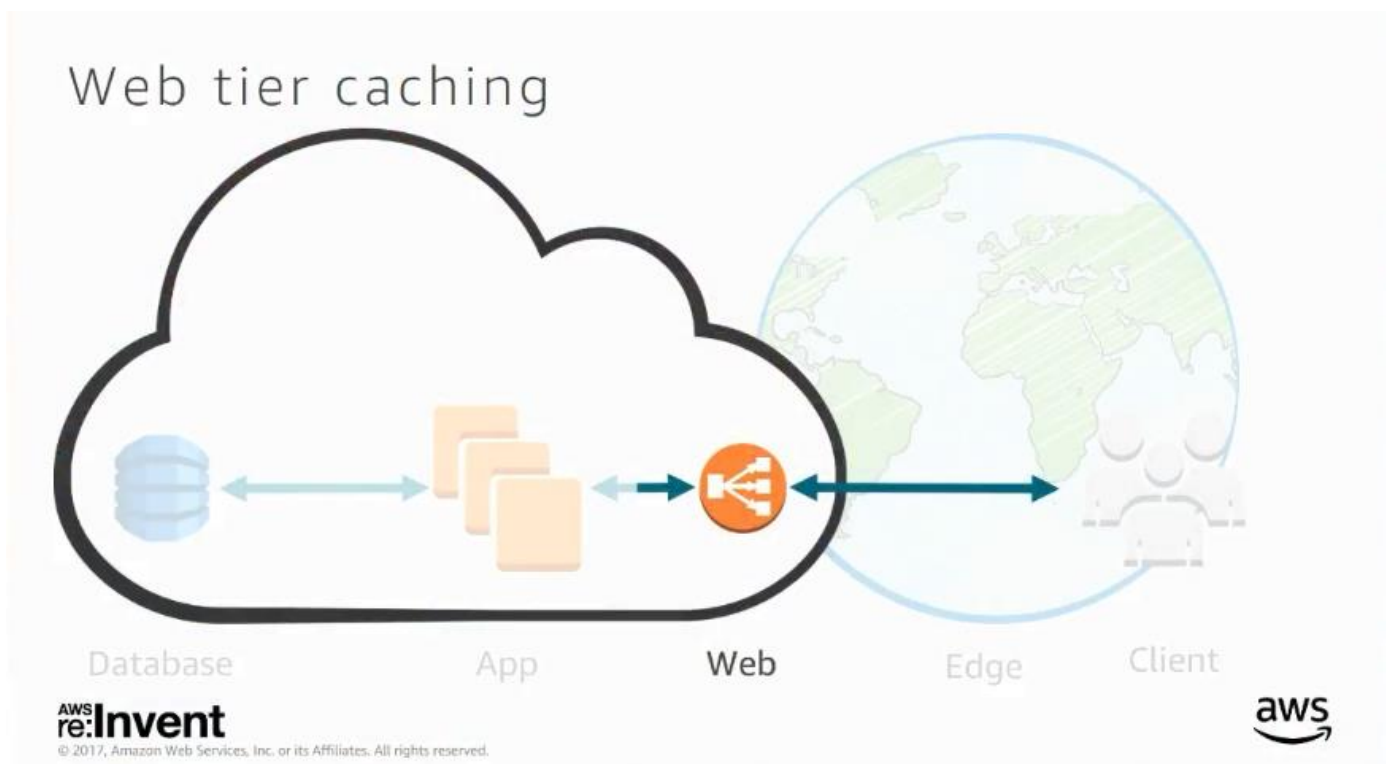




- Edge ✓
- Web tier
- App tier
- Database


**AWS re:Invent**  
© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

**aws**  
Photo by [Annie Spratt](#) on [Unsplash](#)



This is the bit where you deliver content to your users

## Web tier caching



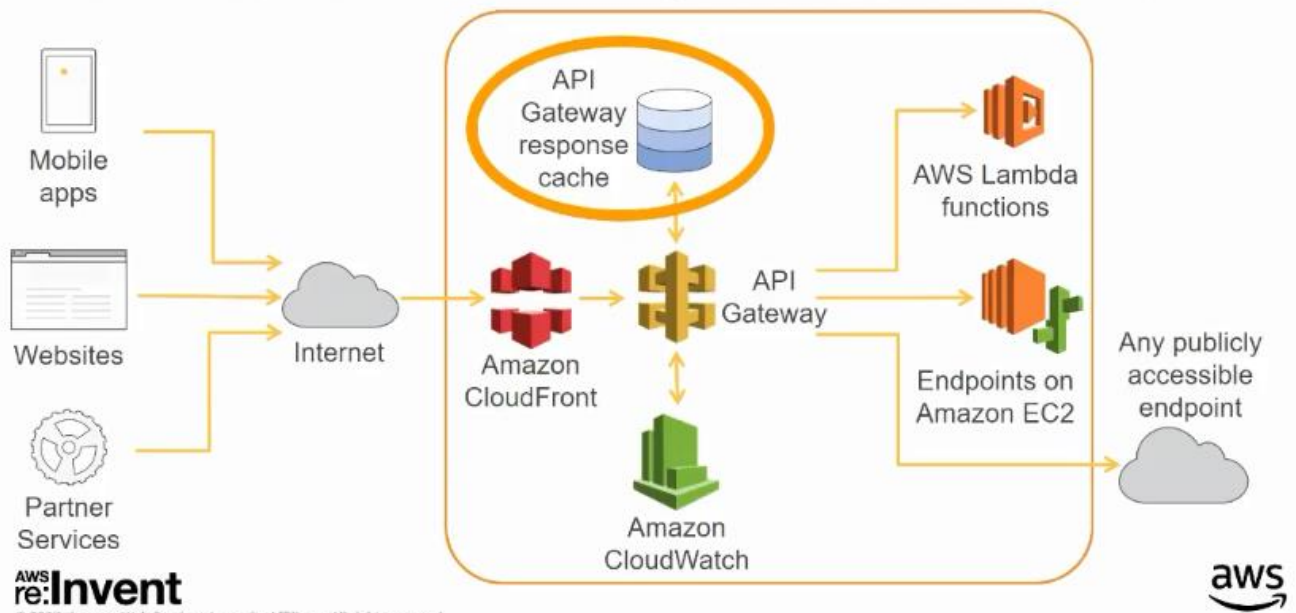
You can put a cache here before it hits your application, you can decide what to cache upfront by trying to cache your HTTP responses in-memory at the web layer. This means that you should use EC2 instances that have more memory

## Popular web cache solutions

- Varnish
- Nginx
- Apache with mod\_cache/mod\_proxy
- Squid
- Language/framework caches (APC, Zend)
- Cache in-memory
- Choose high RAM instances, like the R4 family




# Amazon API Gateway: Serverless APIs



You can add an in-memory cache layer from an API Gateway on top of your HTTP stack to be used for caching some of your responses

## Web cache tips

- Cache all static content
- Cache logged out users
- Look for heavy hitters in your logs
- Monitor hit/miss ratios
- Choose TTLs wisely:
  - Don't affect deployments
  - Even small TTLs (for example, 60 seconds) help





## Web caching recap

- Web caching makes sense, even if you already use CloudFront
- Amazon API Gateway comes with built-in caching: Easy
- Choose TTLs wisely



- Edge
- Web tier
- App tier
- Database




aws  
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

aws

Photo by [Annie Spratt](#) on [Unsplash](#)

## App tier caching



You can also put a cache on top of your applications that are running on some EC2 servers to help reduce latency

## What to cache


- Sessions
- Results
- Aggregations
- Templates
- Environments
- Configurations



In short: **Everything!**

# Everything counts (in larger amounts)

- X0,000–X00,000 requests
- Per second
- Even microseconds add up!
- Everything worth caching
- Even for very short times
- 1 ms @ 10K rps  
=> 10 app seconds per real seconds saved  
=> 7120 Instance h per month

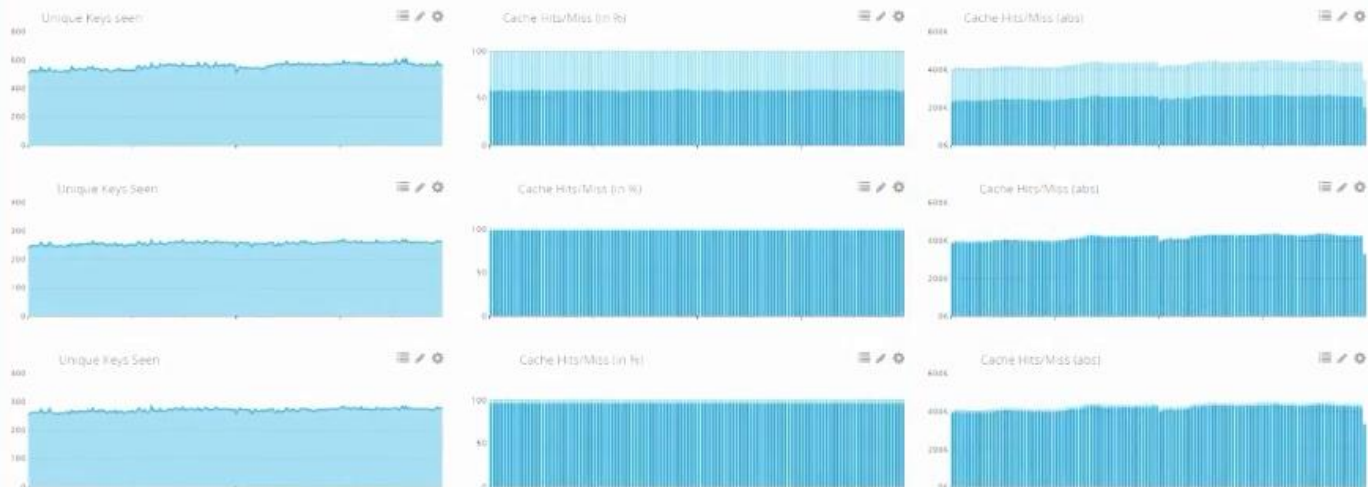


**AWS re:Invent**

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



## Monitor everything



It is important to know what you can cache and how you can cache. Above is one of our data dashboards that helps us figure out what to cache. We are caching domains, advertisers, budgets, etc. We are using 3 main columns like *unique keys seen in a specific timeframe* that tells us how much stuff we are caching in-memory, the **cache hit/miss ratio (number and percentage)** tells us how well we are using the cache we have and how much of the overall requests we are serving from the cache.



## Think 80/20

- Find your heavy hitters
- Likely a very small percentage
- Have them in memory, all the time
- It pays off to do special things for them




Vilfredo Pareto, 1848–1923  
Image: sie-ase.org/

You need to cache things specially for your heavy hitters to improve responses to them,

## RAM

- Every instance type has RAM
- Use it!
- Duplicate data is good
- Pre-load popular data
- Leverage file system cache
- Use your language's caching framework



You need to monitor RAM usage, if you are not using a lot of your RAM then you can cache things in-memory to use spare RAM. You can use idle time to pre-compute things into your cache

# Add more RAM with Amazon ElastiCache



Amazon  
ElastiCache

In-memory key-value store

High-performance

Redis and Memcached

Fully managed; zero admin

Highly available and reliable

Hardened by Amazon



redis

Memcached




AWS  
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

**ElastiCache is a RAM-as-a-service option** that you can use if your instance does not have enough RAM

## Memcached—fast caching

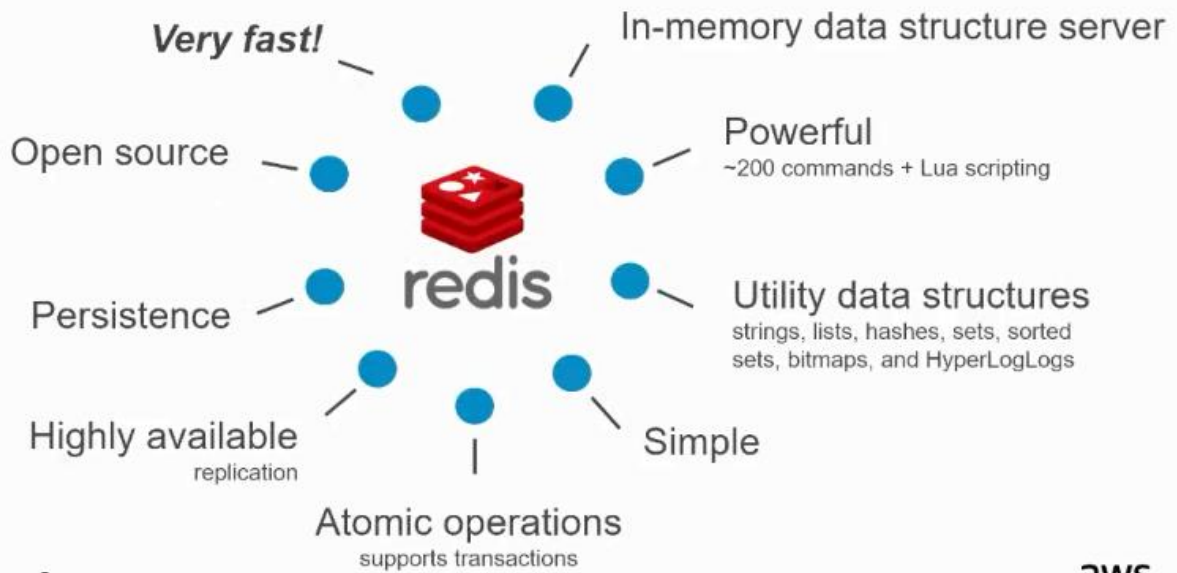


AWS  
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



# Redis—the in-memory leader



**AWS re:Invent**

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



## App caching recap

- Monitor everything!
- Find the 20% of users with 80% impact
- Cache everything, in RAM!
- Consider using ElastiCache







- Edge ✓
- Web tier ✓
- App tier ✓
- Database


AWS  
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

aws

Photo by [Annie Spratt](#) on [Unsplash](#)

## Database caching



Database

App

Web

Edge


Client

AWS  
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

aws

## Caching the traditional way



**AWS re:Invent**

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



We can make our application talk to the cache first when trying to get data before we talk to our database

## Cache invalidation

*"There are only two hard things in Computer Science: cache invalidation and naming things." –Phil Karlton*

Two ways to cache:

- Time to Live (TTL) with invalidation
- Keep the cache in sync all the time

We can put a TTL on the key in the cache we are using like Redis or Memcached so that the cache can invalidate the cache object after TTL elapses. The 2<sup>nd</sup> option is to always keep the cache in sync using synchronous write.

# Synchronous writes



Sync/check in the app with after-write return values

## After-write return with DynamoDB

```
var updateDynamo = function(value, key, callback) {  
  var params = {  
    Key: {key: {"s": key}},  
    TableName: "dynamodb_table",  
    AttributeUpdates: {  
      Action: "ADD", value: {N: JSON.stringify(value)}}  
    },  
    ReturnValues: "UPDATED_NEW"  
  };  
  dynamodb.updateItem(params, callback);  
};
```

This is a call against DynamoDB, this will make DynamoDB give us the newly written data back

## Compare after-write values

```
var calls = [
  function(callback) {
    updateDynamo(value, key, function(err, result) {
      callback(err, result.Attributes["val"].N);
    });
  },
  function(callback) { redis.incrbyfloat(key, value, callback); }
];
async.parallel(calls, function(err, result) {
  if(Math.abs(parseFloat(result[0]) - parseFloat(result[1])) > 0)
    logger.error("diffs @ the result"); // do something against it!
  callback(err);
});
```

We can also do the same thing with Redis in parallel to when we do the write to DynamoDB. This allows us to know if our cache is out of sync with the Database.

## Uncoupled writes/invalidation



Sync/updates via AWS Lambda (uncoupled)

Downside: Small delay for cache updates

In case we don't want to implement changes in our apps, we can use a lambda function to make sure that our cache is in sync with the database. Our app can continue to write to the database while it can read first from the cache.




# Amazon DynamoDB Accelerator (DAX)

DynamoDB API compatible:

```
// If you have this...  
var dynamodb = new AWS.DynamoDB();  
// change it to  
var dynamodb = new AmazonDaxClient({  
    endpoints: ["your-dax.amazonaws.com"],  
    region: "us-east-1"  
});
```

Multiple tables with one DAX cluster possible!



**AWS re:Invent**

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

We can instead talk to DAX client (instead of the DynamoDB client) all the time and not worry about the cache anymore, DAX will take care of the caching and the TTL.

## Performance of DAX

Without DAX



Average ~5000–6000  $\mu$ s  
Consistent performance  
No warming phase  
Detailed metrics per request


With DAX



Average ~400–450  $\mu$ s  
Very consistent performance  
Warming phase (average ~4500  $\mu$ s on cold keys)  
No metrics per request


# Cache everything

Without negative caching



Cache hit ratio 25%–30%

With negative caching



Cache hit ratio 89%–95%


Hit

Miss

Negative result caching is when we are not saving the 'no data found' result from DynamoDB for some queries, then we keep on asking DynamoDB about those same queries over and over again. You need to cache your negative result caching scenarios.

## Database caching recap

- Cache everything, even negative results!
- Consider cache auto-update with Lambda
- Can combine app with database cache
- Use DAX





## Things to do:

- Check out Amazon CloudFront/API Gateway
- Monitor everything
- Understand your top 20% of users
- Consider adding a web cache
- Explore app caches with Amazon ElastiCache
- Use DAX
- Cache everything!

# Speed on!



AWS  
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

aws

Photo by Robin Pierre on Unsplash

Liked this talk? Meet us again!

## ARC303–Running Lean Architectures: How to Optimize for Cost Efficiency

Tuesday, November 29, 2:30 p.m.

Venetian, Level 2, Venetian E

AWS  
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

aws