In this session, we explore the new Network Load Balancer that was launched as part of the Elastic Load Balancing service, which can load balance any kind of TCP traffic. This offers customers a high-performance, scalable, low-cost load balancer that can handle millions of requests per second with very low latencies, while maintaining high levels of performance. Come and learn more about this new Network Load Balancer.

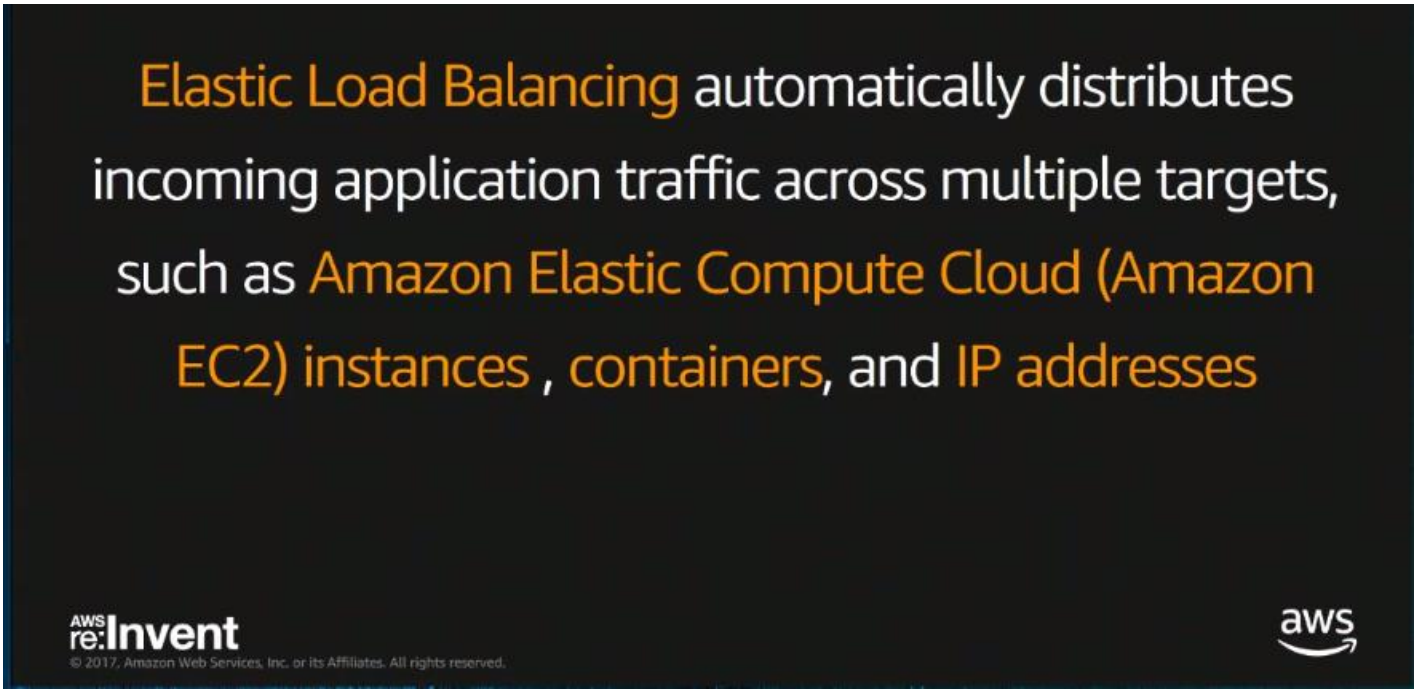

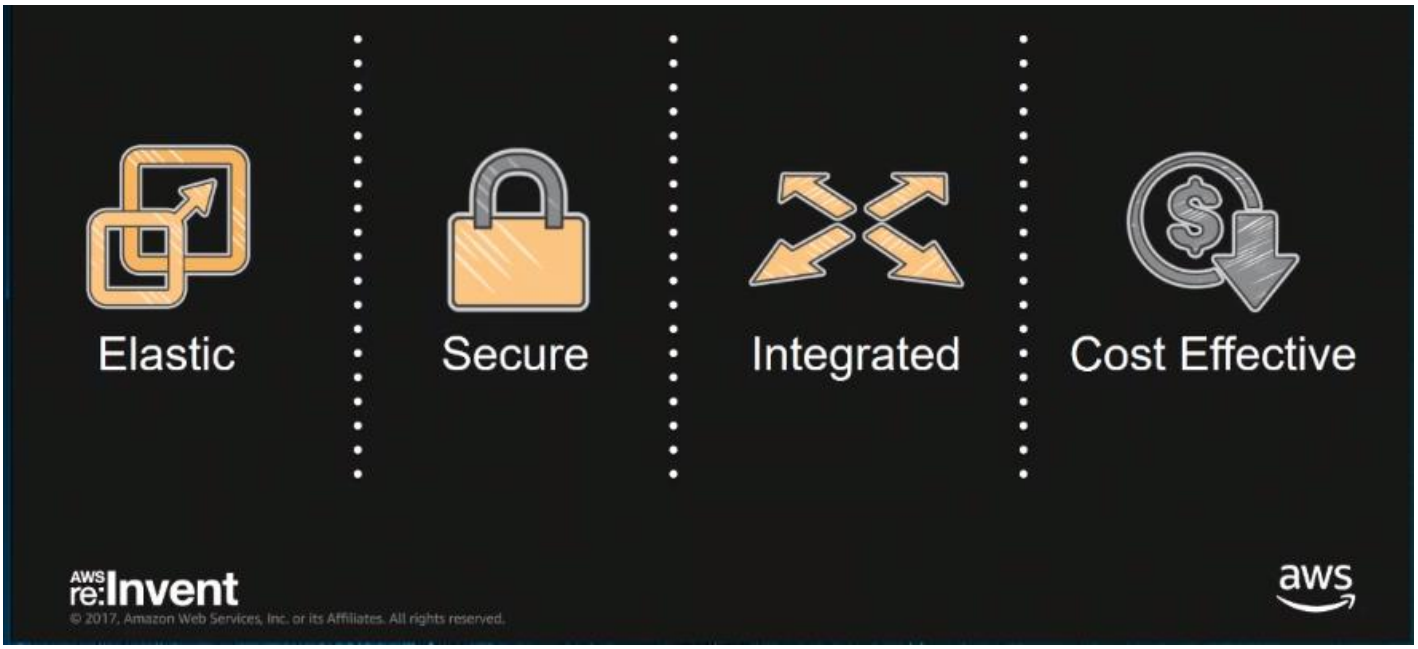

These are the advantages of the ELB, you simply add the targets to the ELB and you get all the above benefits. ELB uses best in class ciphers and protocols to provide security to your LB and targets. ELB is also integrated with several AWS services like ECS, EKS, CF, Route53, etc

This instance doesn't scale for instant spike in traffic and is not secure

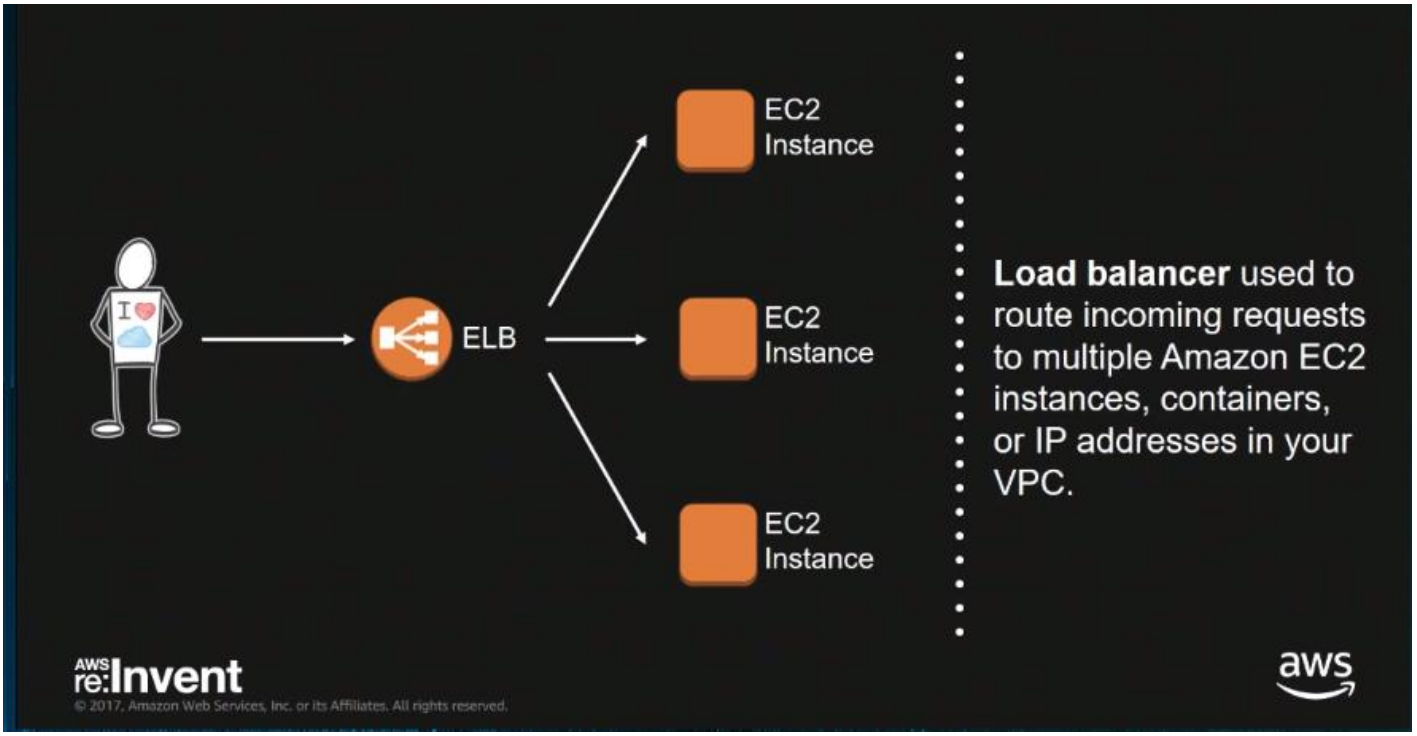

This is what is recommended when deploying an application in EC2

Layer 4 (network)

Supports TCP

Incoming client connection bound to server connection.

No header modification.

Source IP is preserved in the header or Proxy Protocol prepends source and destination IP and ports to request

Layer 7 (application)

Supports HTTP and HTTPS.

Connection terminated at the load balancer and pooled to the server.

Headers may be modified.

X-Forwarded-For header contains client IP address.

These are the 2 types of LB available, it depends on what layer the LB runs in the OS stack

The Elastic Load Balancing Family

Application Load Balancer

HTTP & HTTPS (VPC)

Network Load Balancer

TCP Workloads (VPC)

Classic Load Balancer

Previous Generation
for HTTP, HTTPS, TCP
(Classic Network)

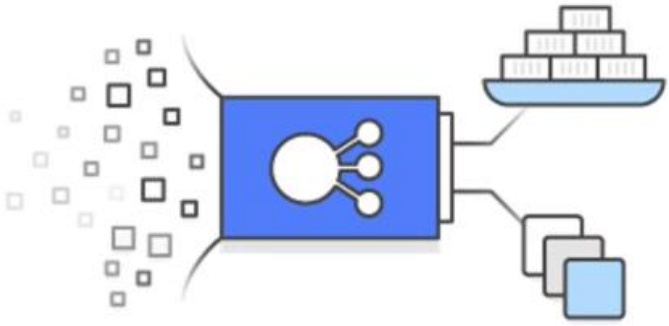

Network Load Balancer (NLB)

aws re:Invent

aws

	Application Load Balancer	Network Load Balancer	Classic Load Balancer
Protocol	HTTP, HTTPS, HTTP/2	TCP	TCP, SSL, HTTP, HTTPS
SSL offloading	✓		✓
IP as Target	✓	✓	
Path-based routing, Host-based routing	✓		
Static IP		✓	
WebSockets	✓	✓	
Container Support	✓	✓	

aws re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

aws

Static IP means a single IP per AZ even in volatile traffic, it can be used for whitelisting valid IPs that we want to allow communication with.

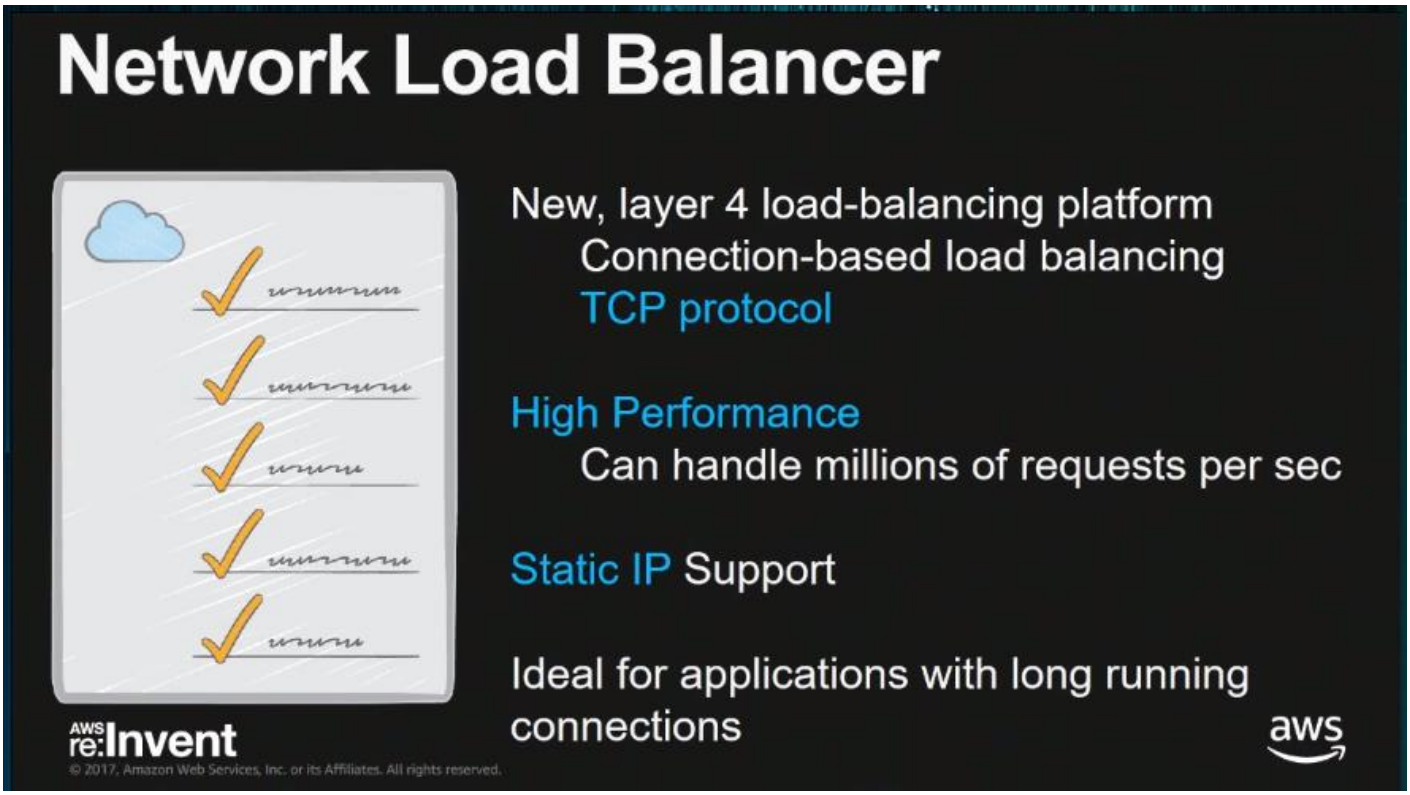

NLB performs well with no cold starting and in very volatile traffic patterns, NLB also supports using static IPs per AZ so that you have a single IP being used as you scale and can whitelist in your firewalls and other devices sitting in front of your LBs.

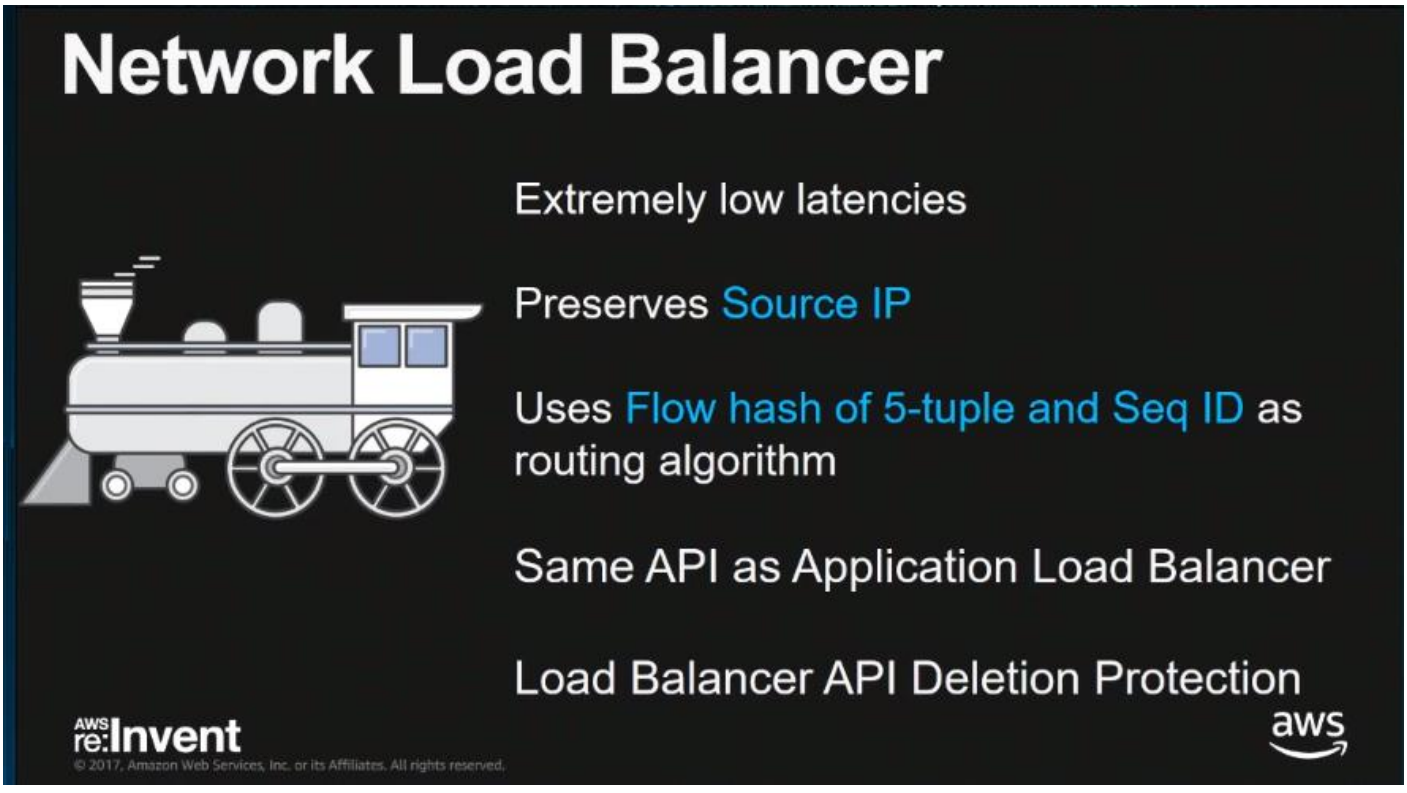

NLB uses the same set of new APIs as the ALB

NLB Performance

Can handle Millions of rps out of box and volatile traffic patterns

Test Setup

New out of box NLB with 3 Availability Zones in US-EAST-1

100 c4.xlarge clients running Apache bench
Backend fleet with 75 c4.2xlarge servers
Get requests
Response – 1KB static page
1,000 concurrent connections/client

NLB Performance

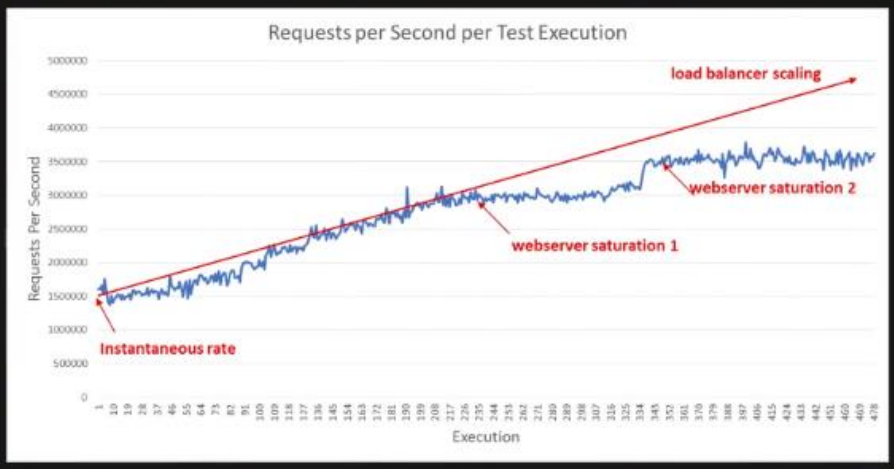

Bees with machine guns, executed in a loop

bees attack --url '<NLB-URL>' --number 10000000 --concurrent 100000 –keepalive

Performance Graph shows no errors and content was served fine

The NLB can linearly scale to tens of millions of requests per second

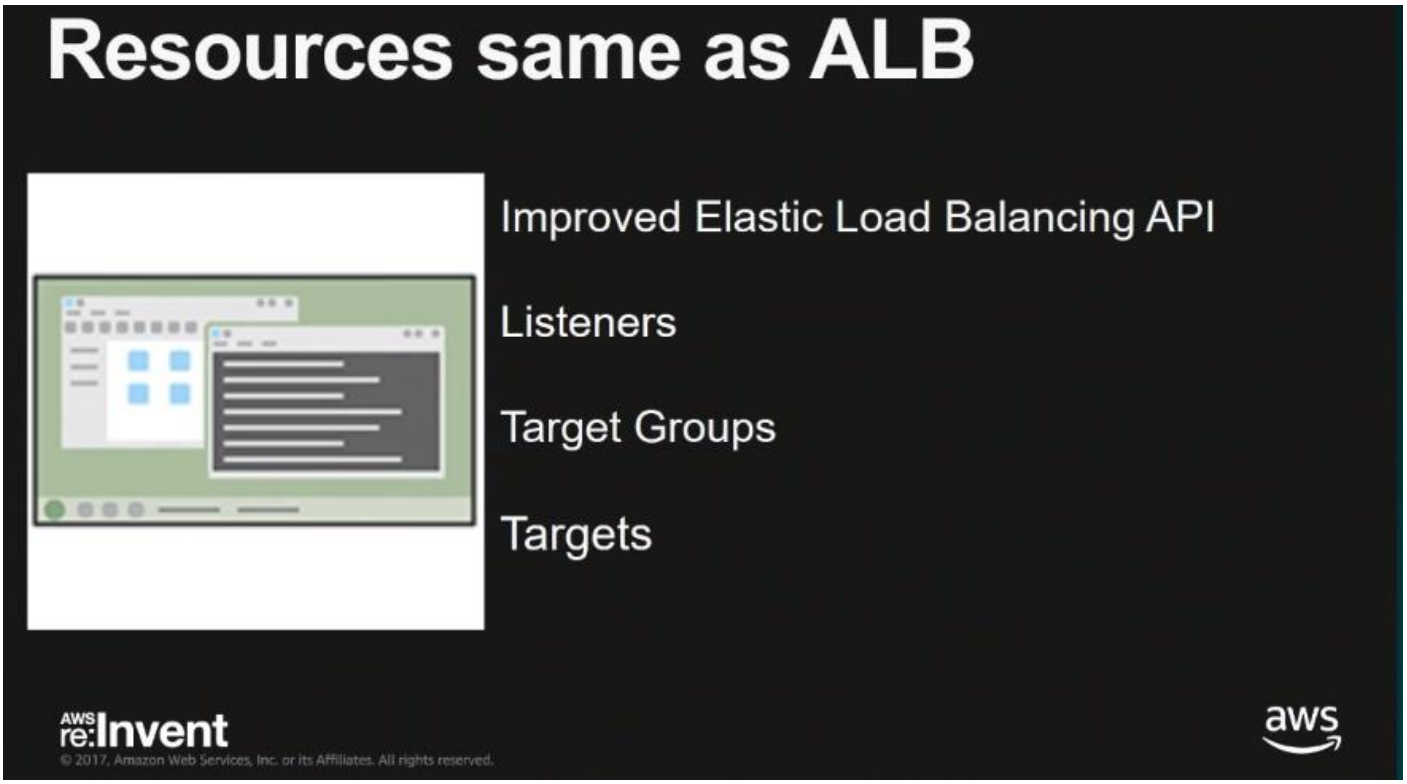

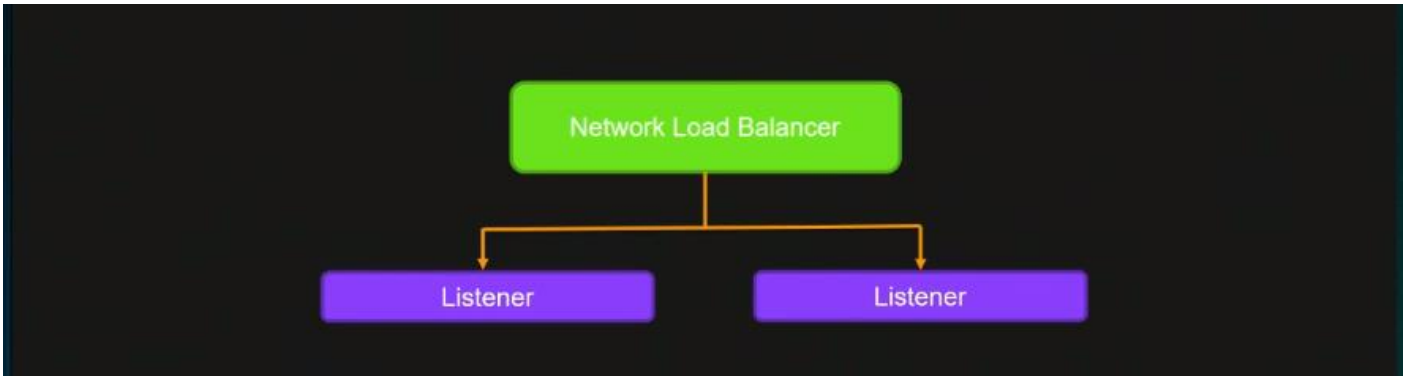

We start by attaching listeners to the NLB as above

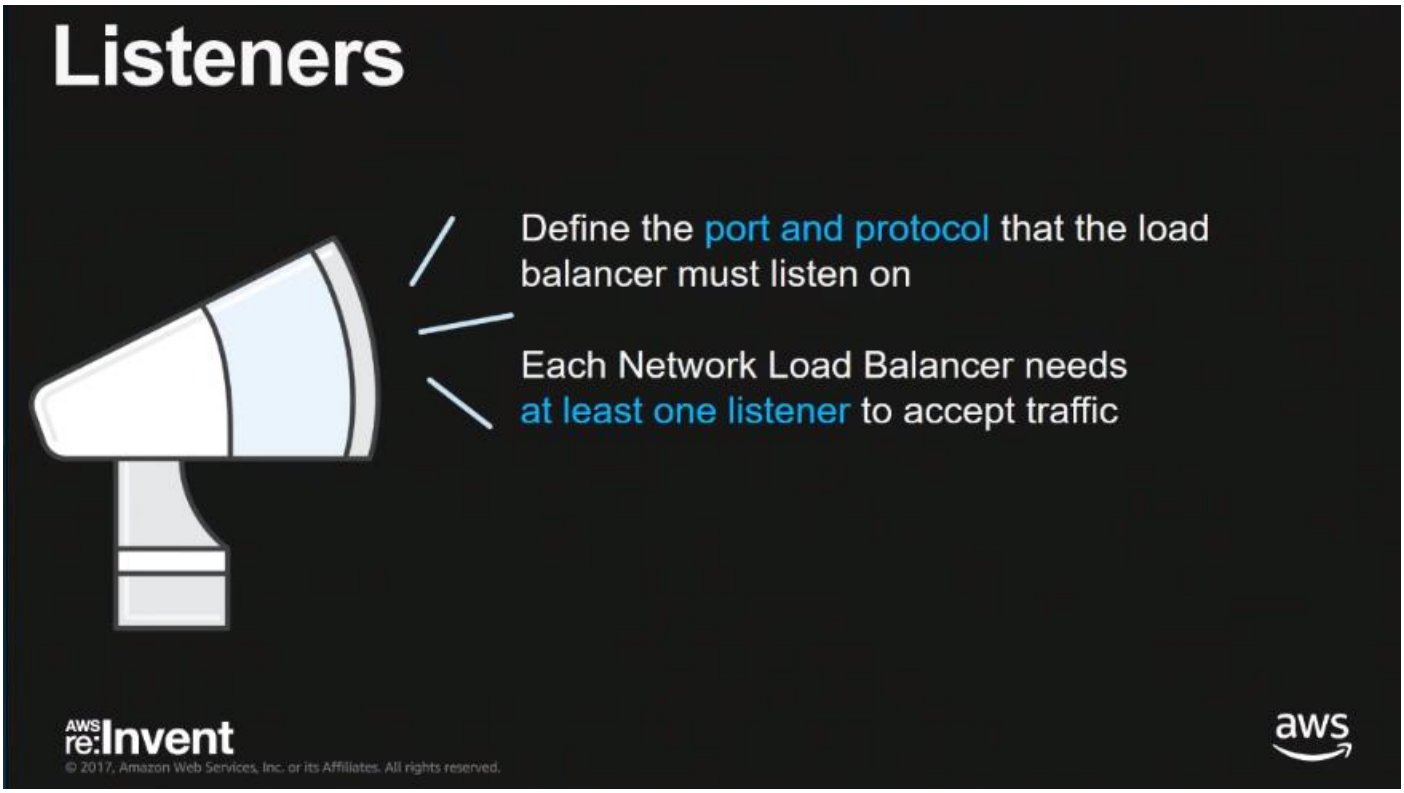

You need to have at least 1 listener before you start sending traffic through

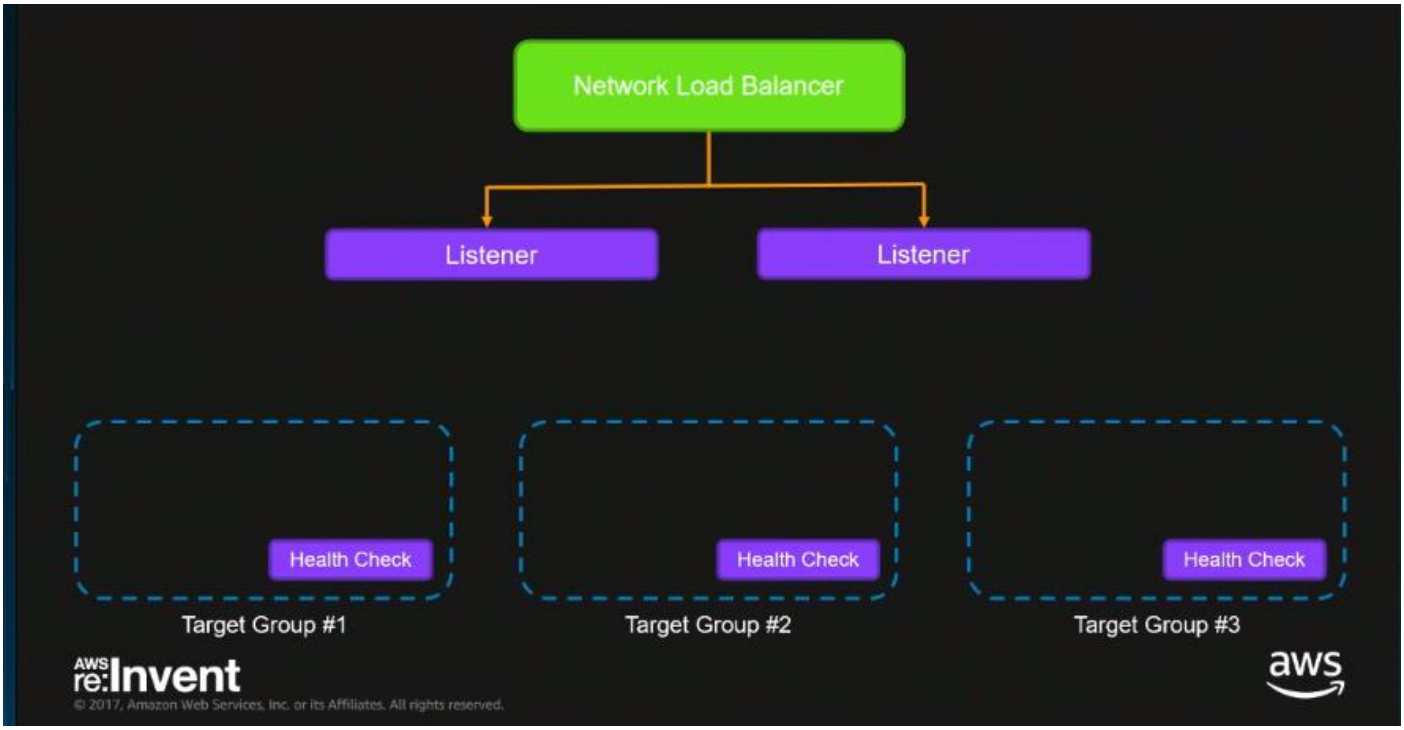

We then create or identify the target groups that we want to have the listener work with

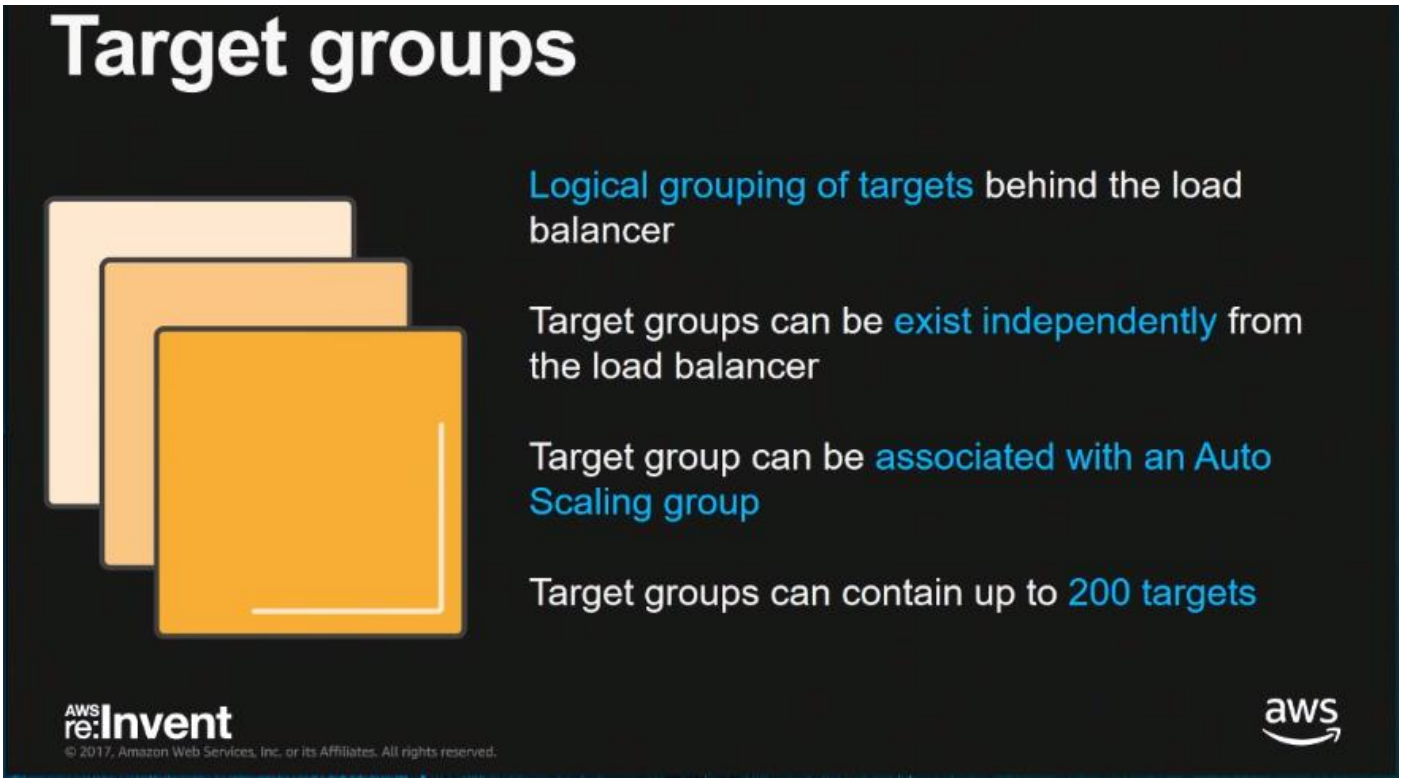

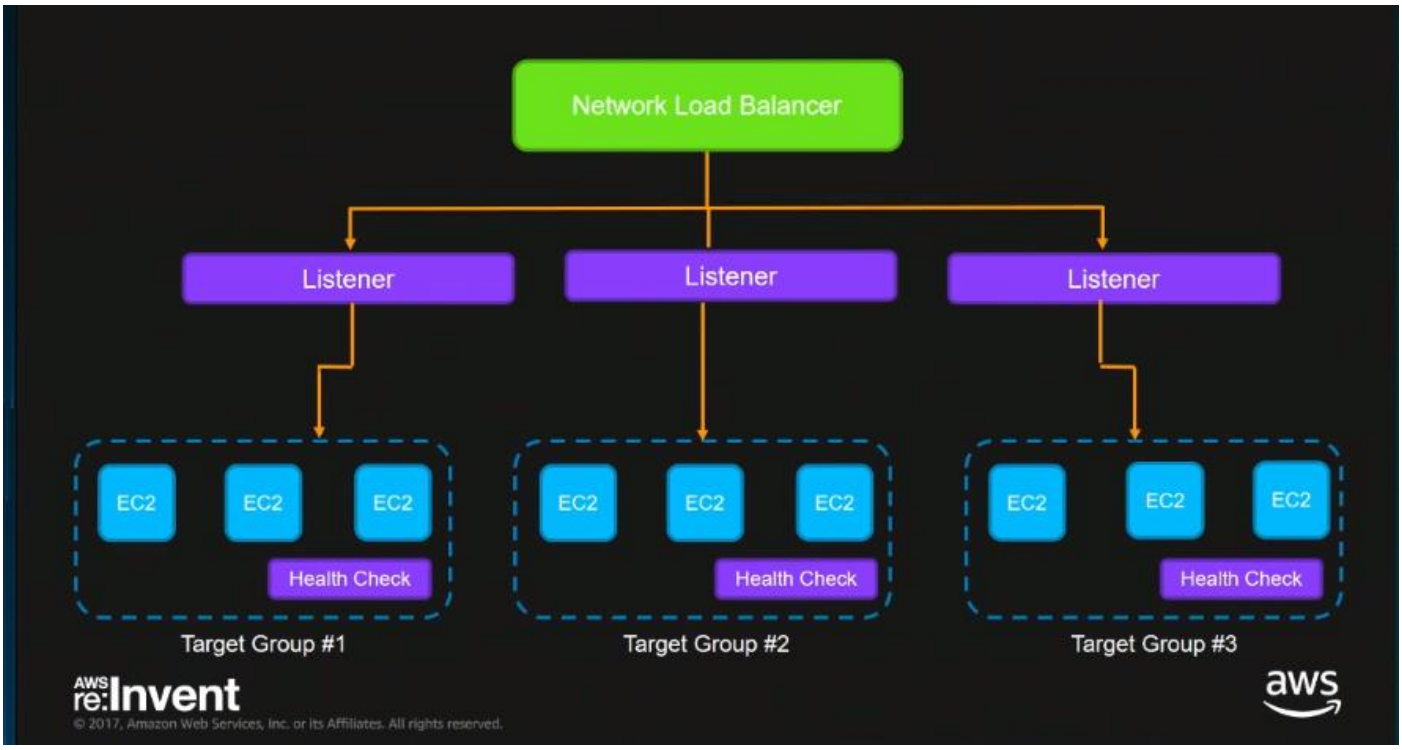

We then configure the health checks on the target groups to monitor the targets

Targets

Support for Amazon EC2 instances, Amazon ECS containers, and IP Addresses.

Amazon EC2 instances can be registered with the same target group using multiple ports for Containers

A single target can be registered with multiple target groups within the same load balancer

Targets can be IP Addresses both accessible within your VPC or via AWS Direct Connect

IP as a Target

Use IP address from the load balancer's VPC CIDR for targets within load balancer's VPC

Use IP address from the RFC 1918 and RFC 6598 range for targets located outside the load balancer's VPC such as on-premises targets reachable over AWS Direct Connect (10.0.0.0/8, 172.16.0.0/12,192.168.0.0/16 and 100.64.0.0/10)

ECS integration

NLB is fully integrated with Amazon EC2 Container Service (Amazon ECS)

Amazon ECS will automatically register tasks with the load balancer using a dynamic port mapping

Can also be used with other container technologies

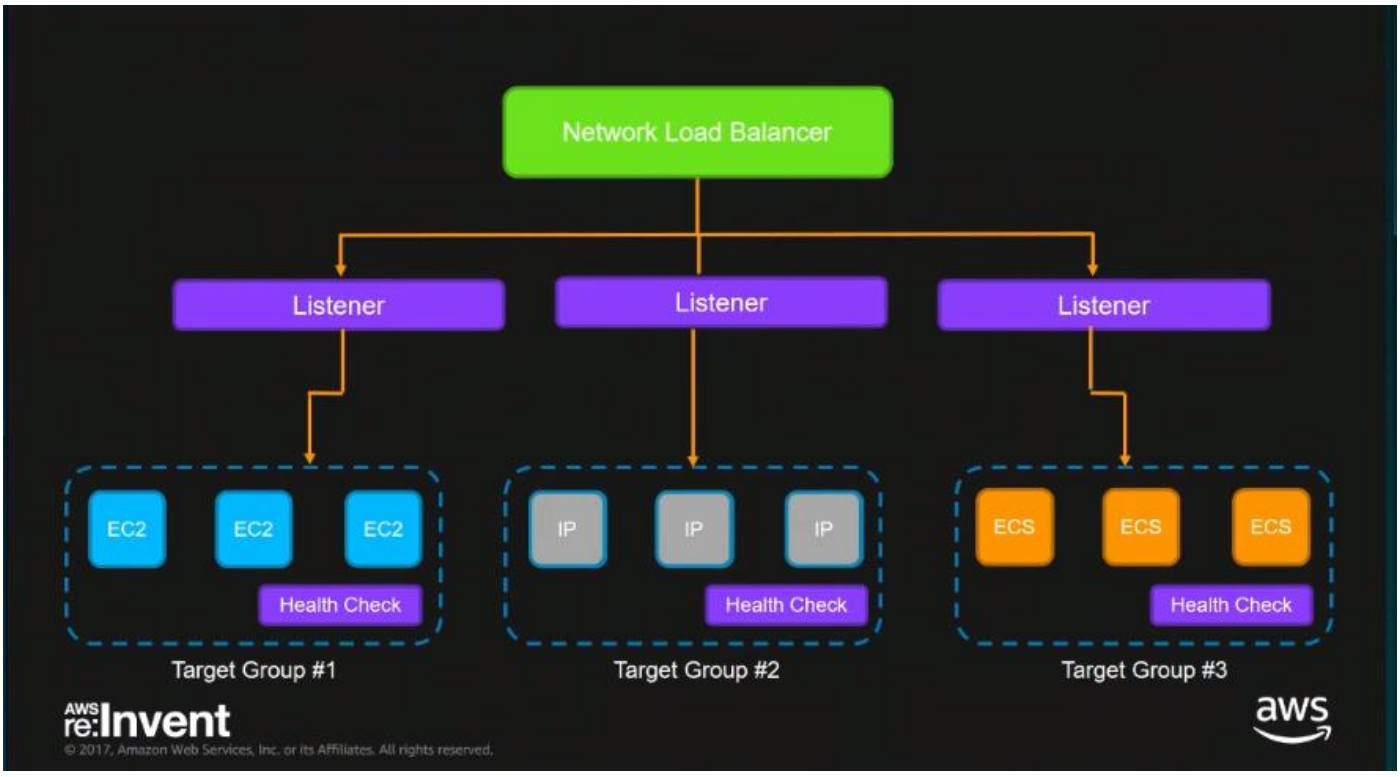

Let us now build up this architecture

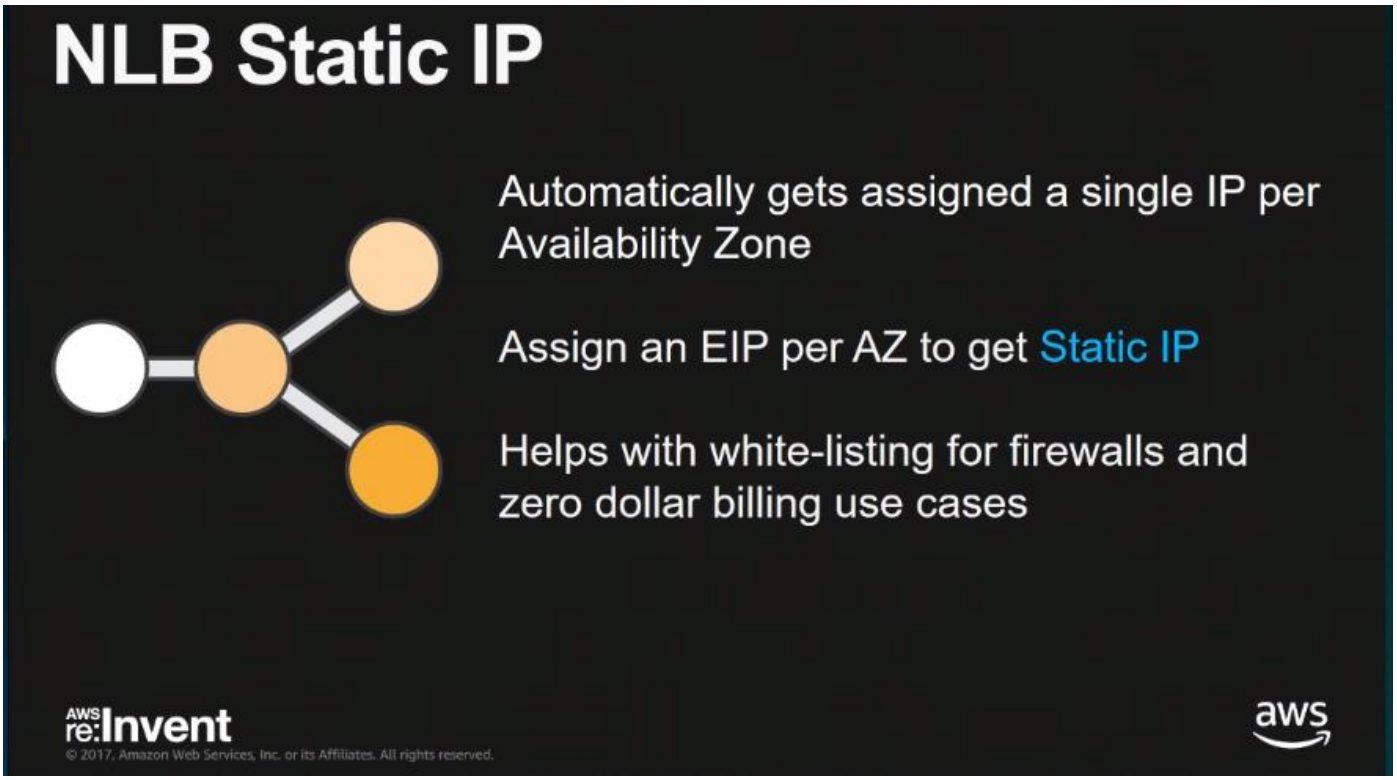

With NLB, you get a single IP per AZ when you create the NLB. If you do not want these IPs to change for the entire life of the LB, you can pick an EIP from your own EIP pool and assign it to the NLB. Static IP use has many use cases, the most common use case is in firewalls where you want to whitelist specific IPs to go across your firewalls.

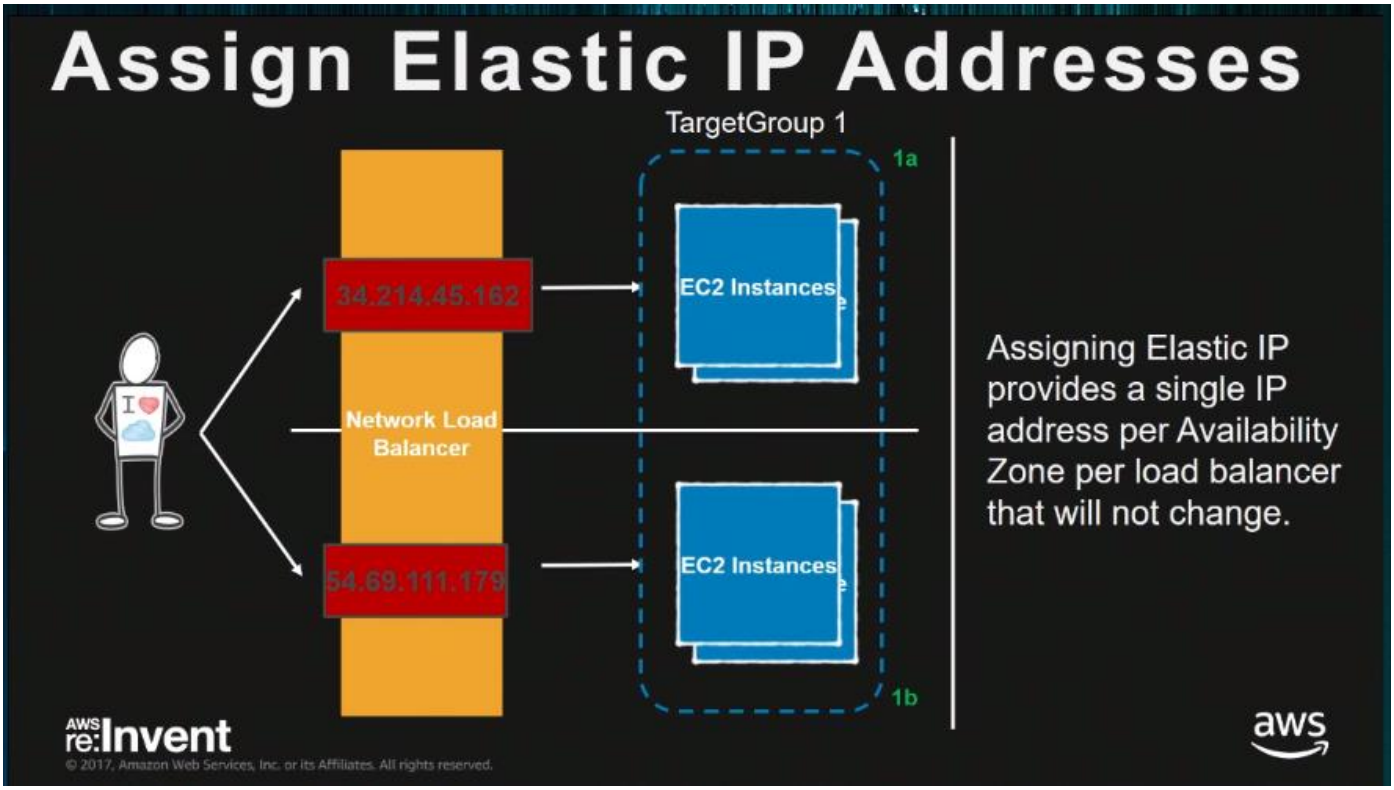

You can assign your own EIP per AZ to the NLB as above

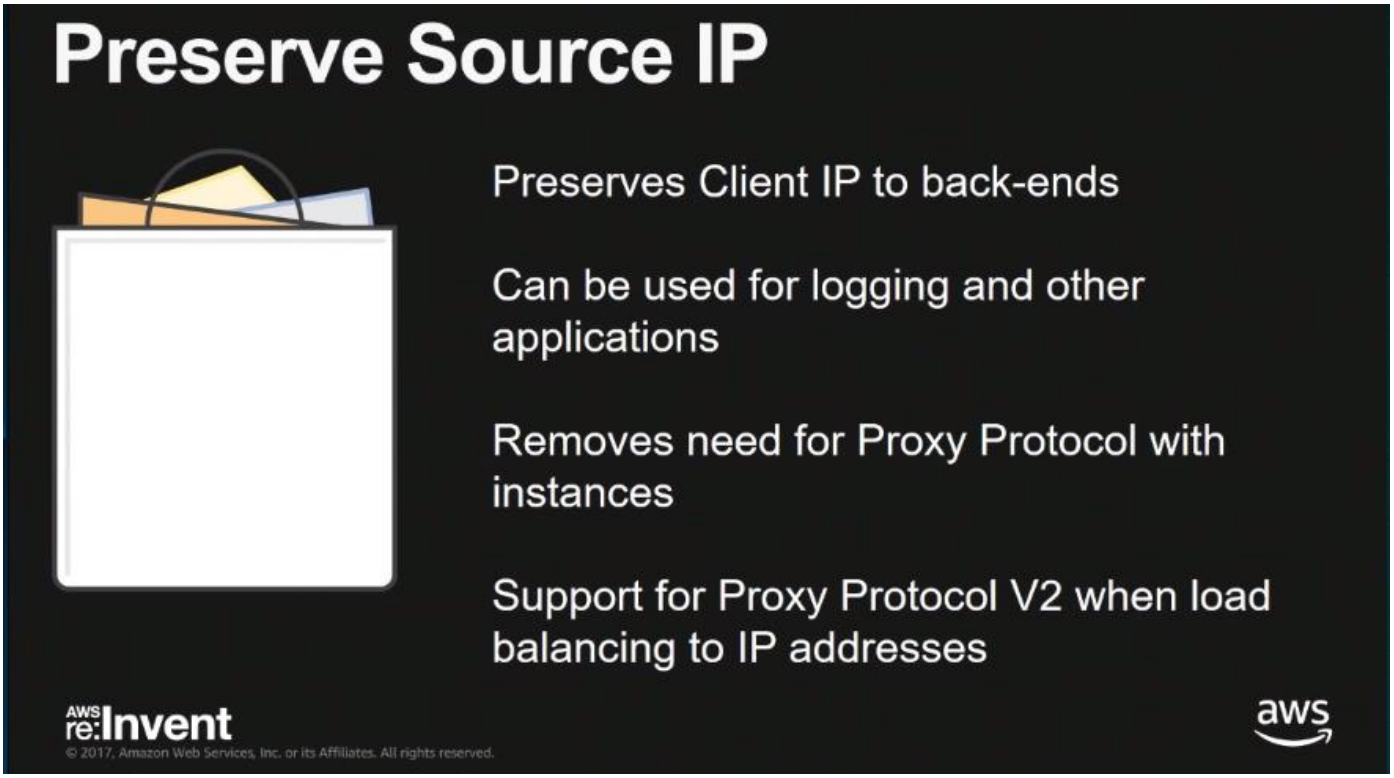

NLB does not terminate the connection to the LB itself and acts as a pass through, the NLB is able to pass the source IP all the way to your backend instances unlike the CLB. You can now set up SGs on your backend targets to allow access to certain CIDR blocks.

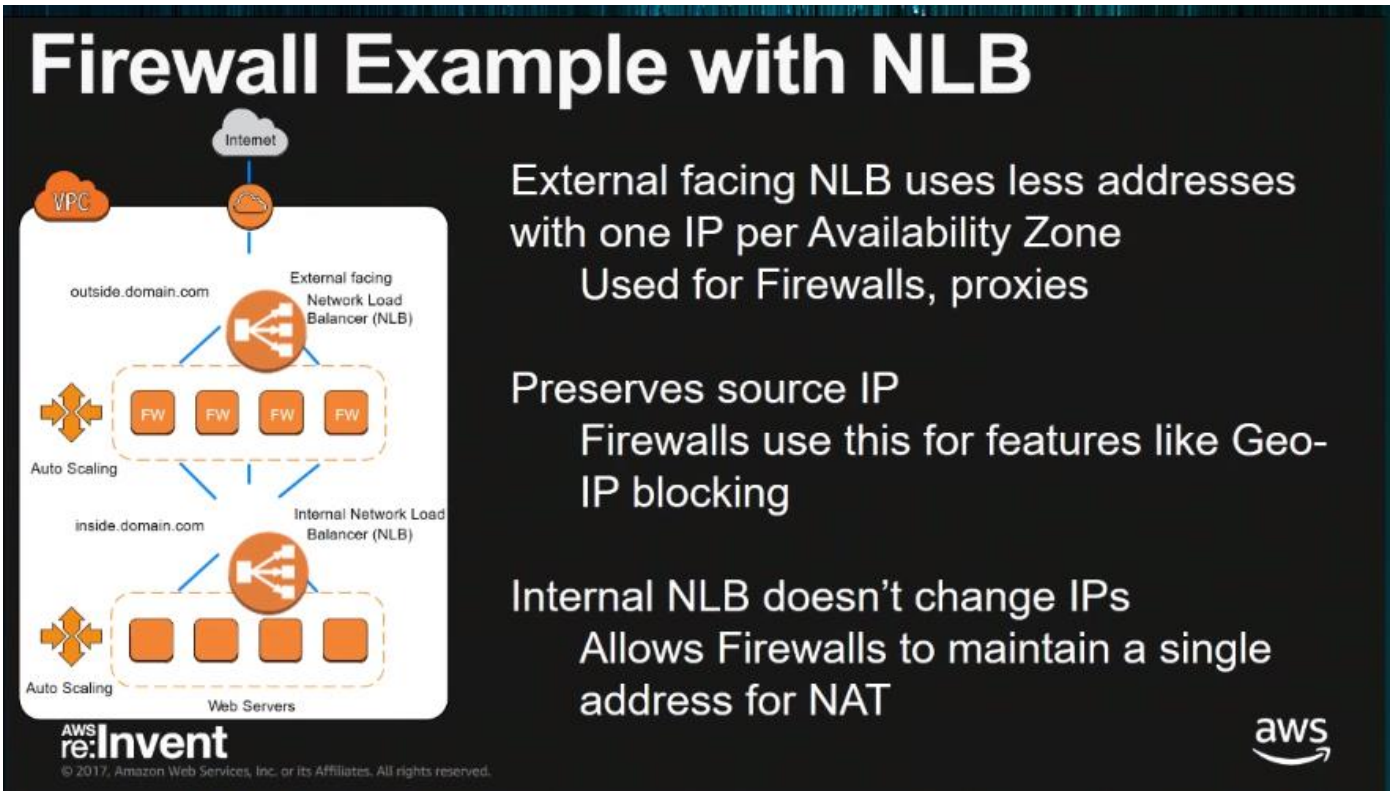

Above is a real-life example where we have an external LB that is fronting a firewall, we then have an internal LB that is fronting a set of web servers. Because the external LB can pass through the source IPs of the requests to the firewalls, you can use the firewalls to do different things like geo-IP blocking. You can also use the sing static IP in NATs in your network.

Health checks allow for traffic to be shifted away from failed instances

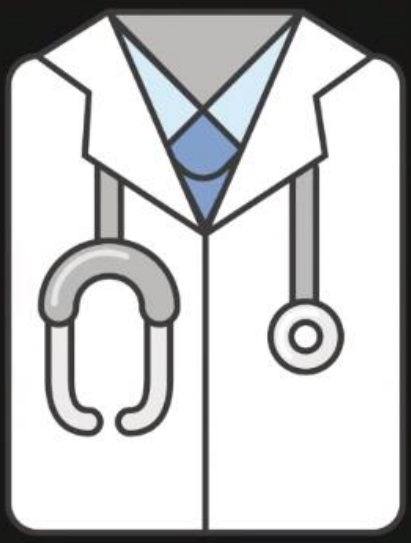

Health Checks

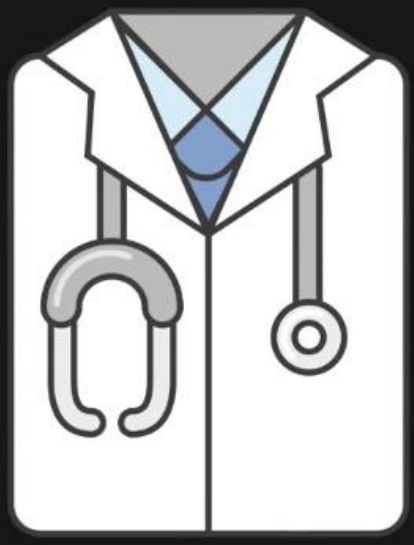

Supports both Network and Application Target health checks

Network health checks
- Based on overall response of your target to normal traffic
- Will fail unresponsive targets in millisec

Application level health checks
- HTTP, HTTPS, and TCP HC
- Customize frequency, failure thresholds

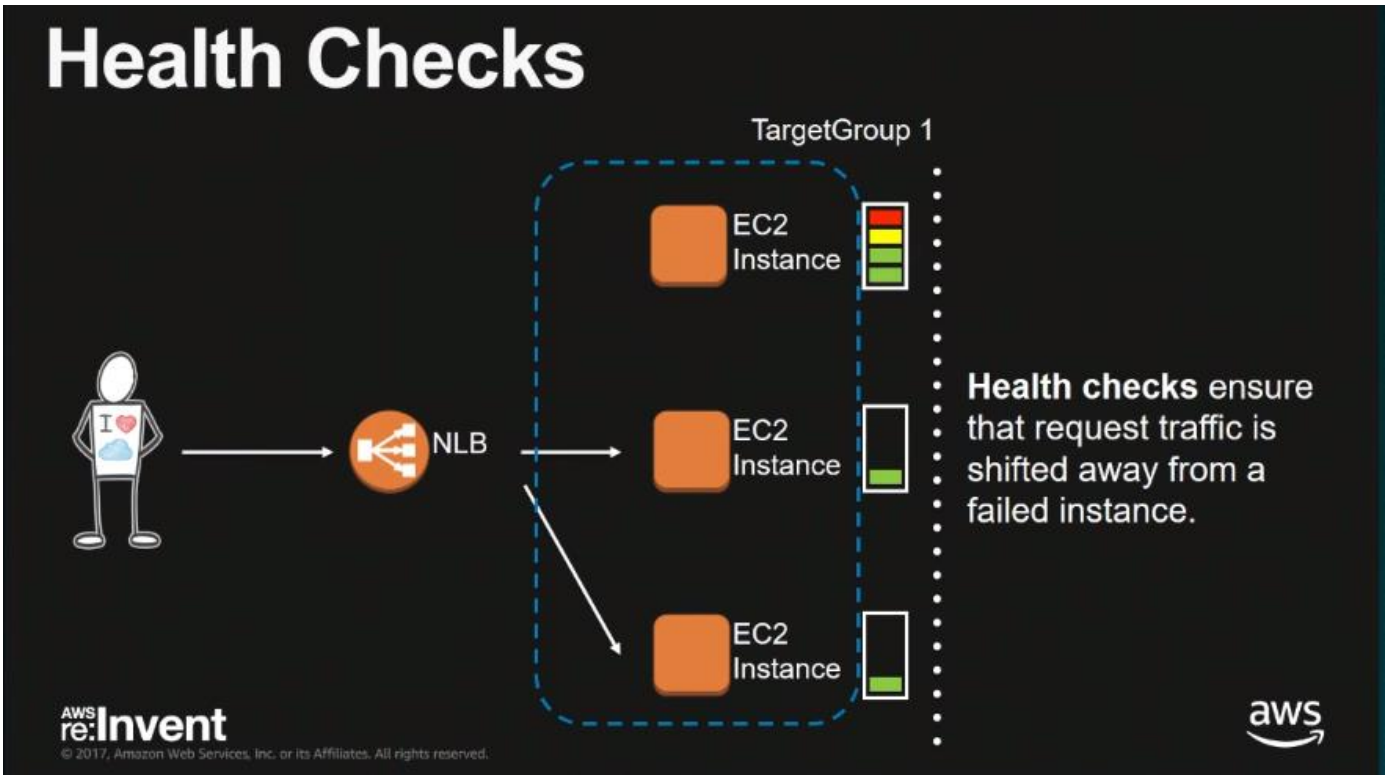

There is a command ***describe-target-health*** that can be used to see the health status of the target

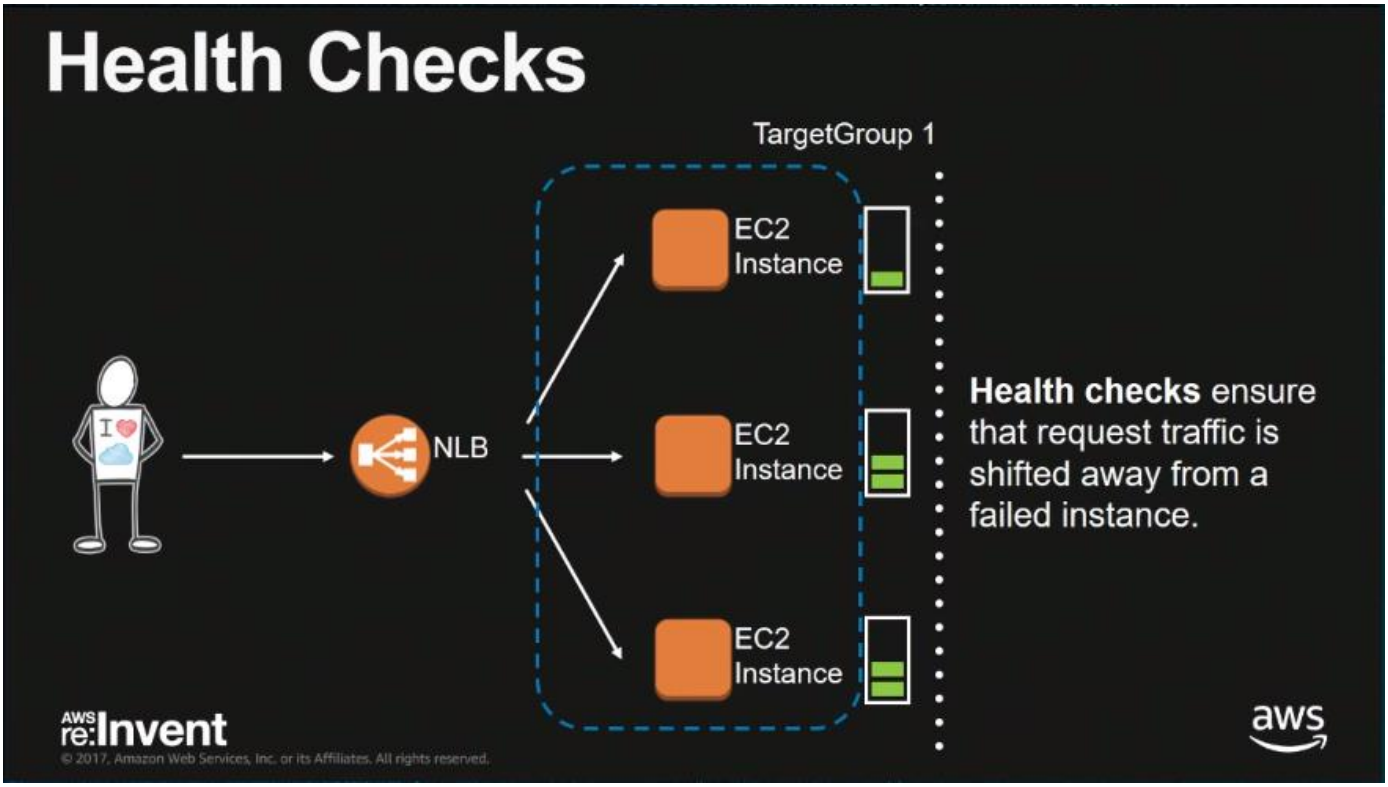

Now the load gets evenly distributed among the targets again

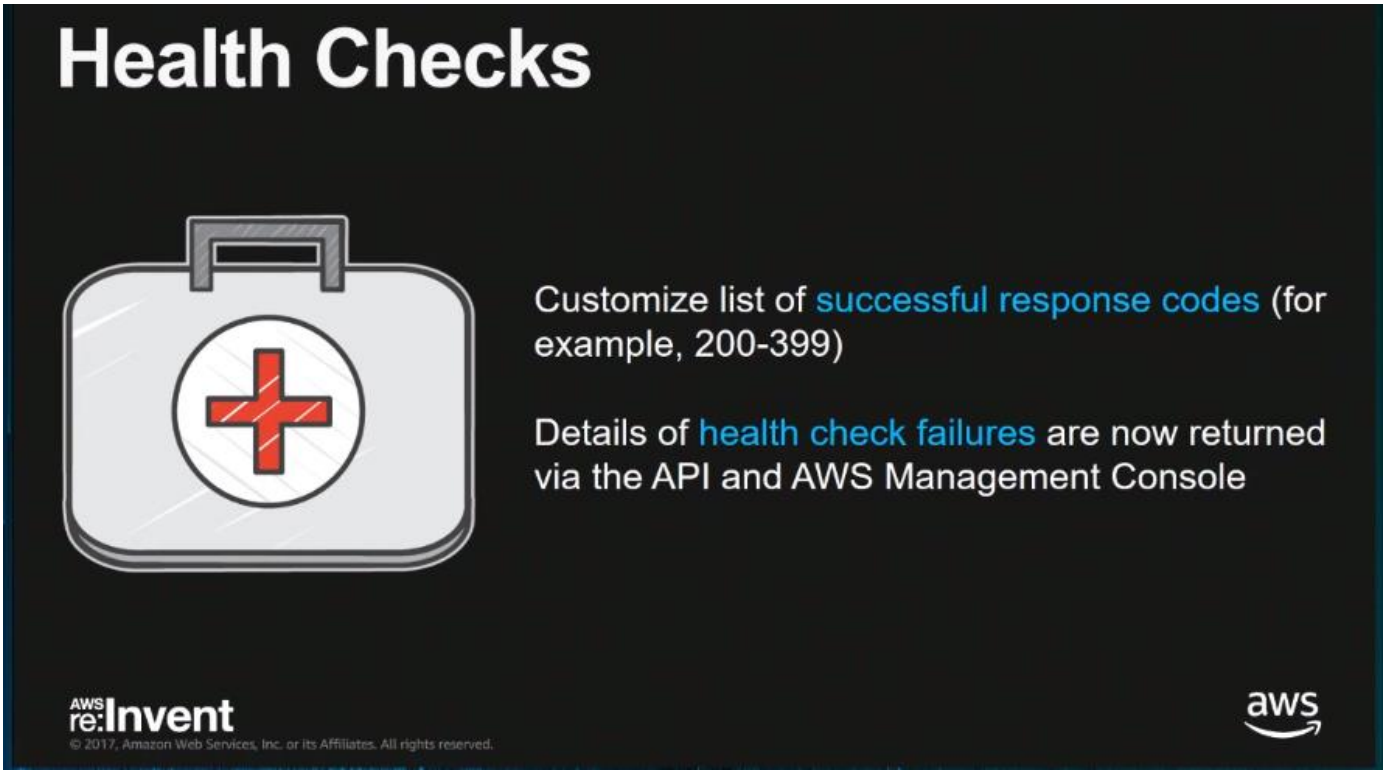

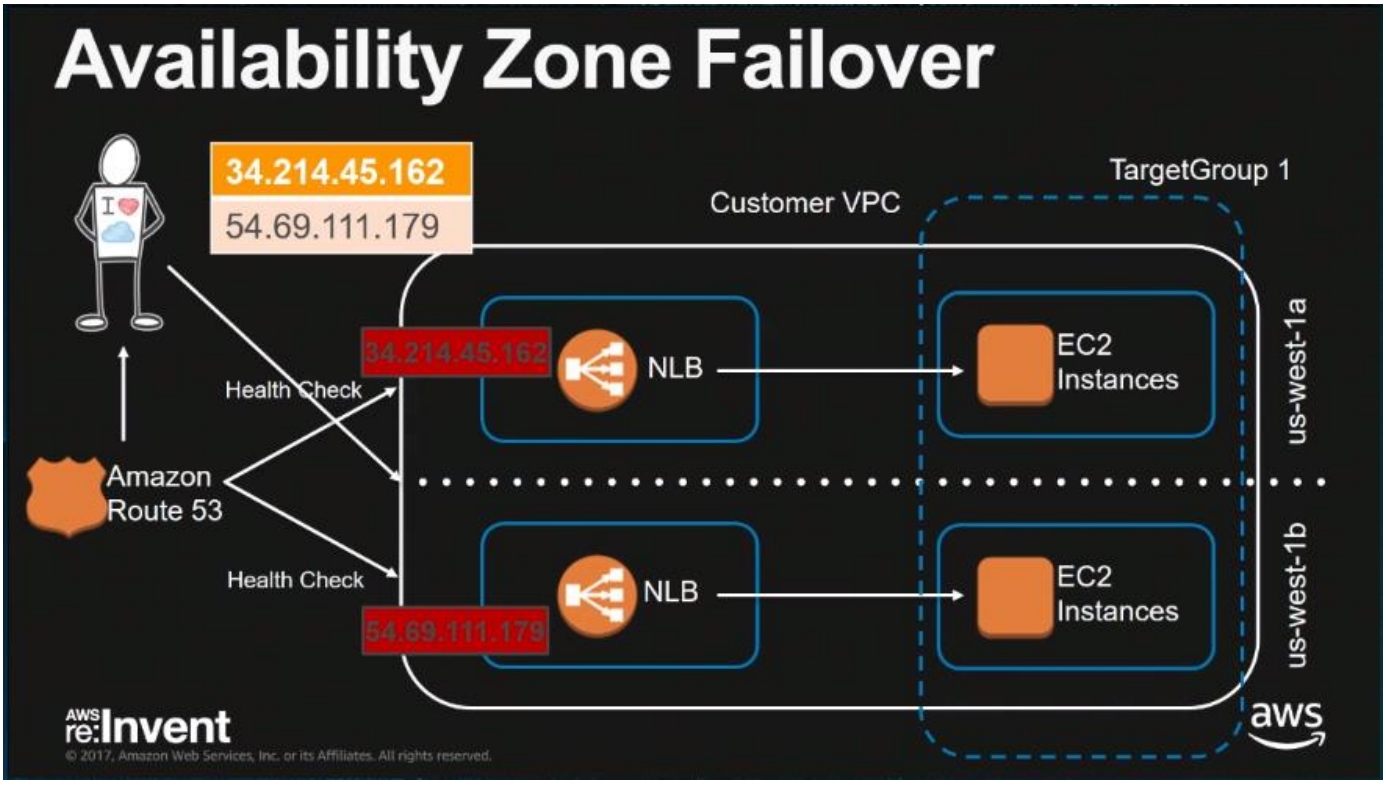

We have the NLB in 2 different AZs with different static address EIPs for the AZ NLB endpoints, the Route53 is resolving the DNS name of the NLB ***www.examplenlb.com*** to those 2 Static IP addresses.

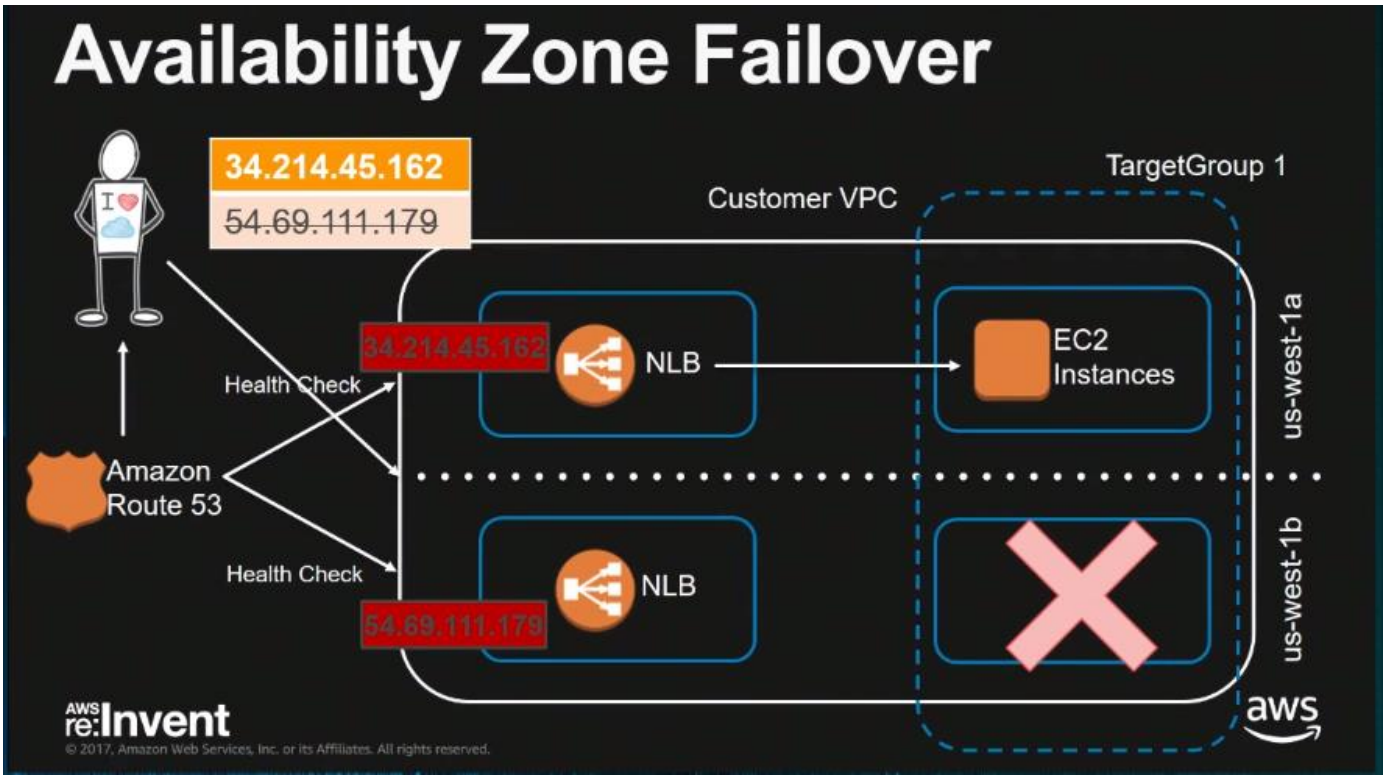

If the backends in one AZ get unhealthy, Route53 will automatically start resolving to only the NLB in the healthy AZ. This is an easy DNS failover.

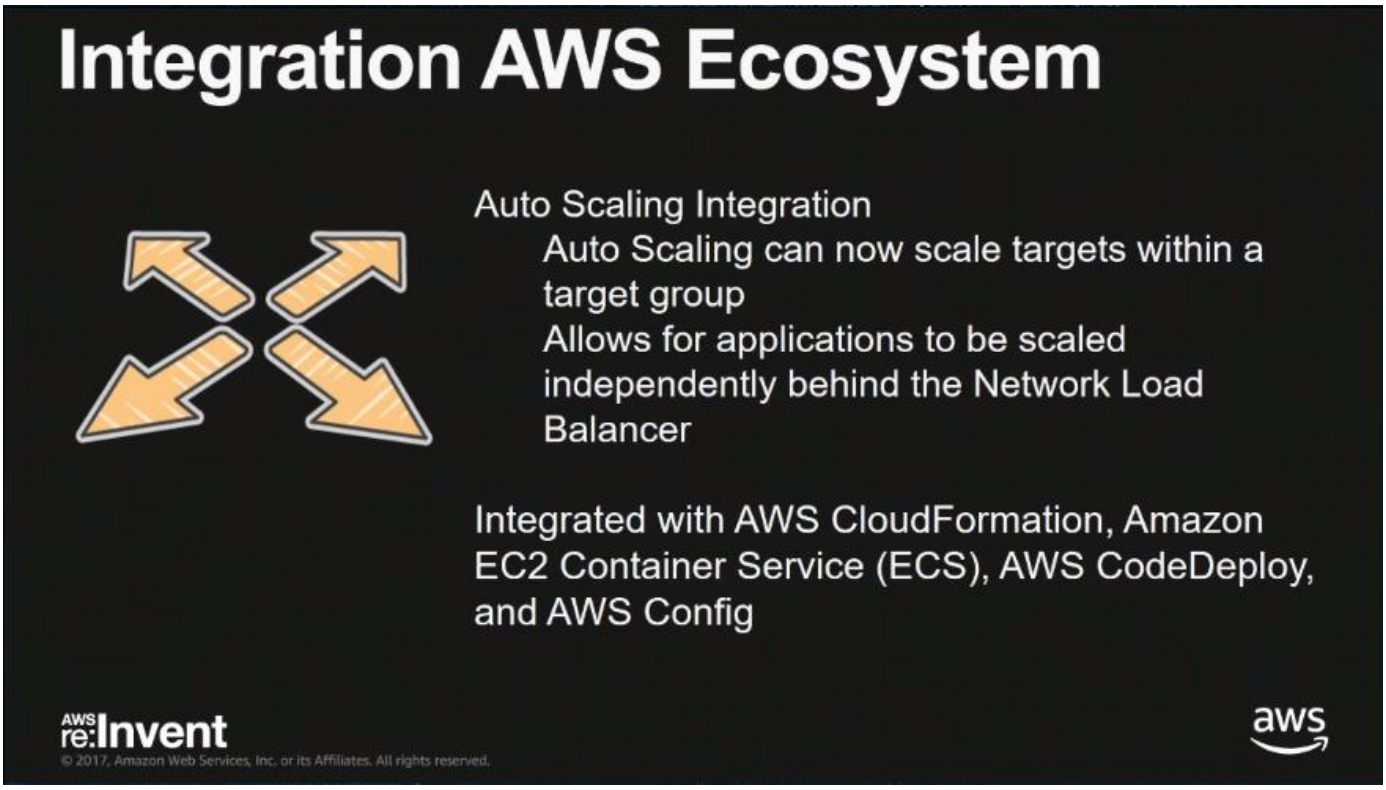

With ECS integration, you can use dynamic port mapping along with your containers and have the NLB front those applications.

Amazon CloudWatch metrics

Amazon CloudWatch metrics provided for each load balancer.

Provide detailed insight into traffic and capacity, errors and back-end health for the Network Load Balancer

Amazon CloudWatch alarms can be configured to notify or take action should any metric go outside the acceptable range.

All metrics provided at the 1-minute granularity.

Traffic and Capacity Metrics

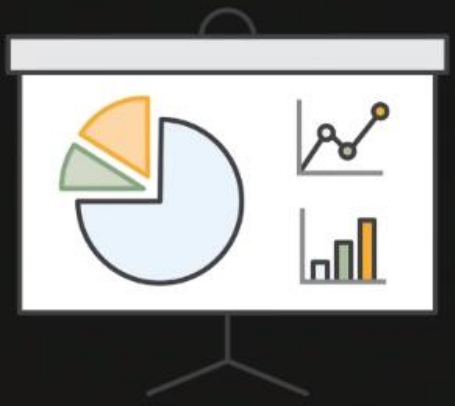

ActiveFlowCount - total number of concurrent TCP flows (or connections) from clients to targets

NewFlowCount - total number of new TCP flows (or connections) established from clients to targets

ProcessedBytes - total number of bytes processed by the load balancer

ResetCounts

TCPClientResetCount – number of reset (RST) packets sent from a client to a target

TCPELBResetCount – number of reset (RST) packets generated by the load balancer

TCPTargetResetCount- number of reset (RST) packets sent from a target to a client

Backend Health

HealthyHostCount – number of targets that are considered healthy

UnHealthyHostCount – number of targets that are considered unhealthy

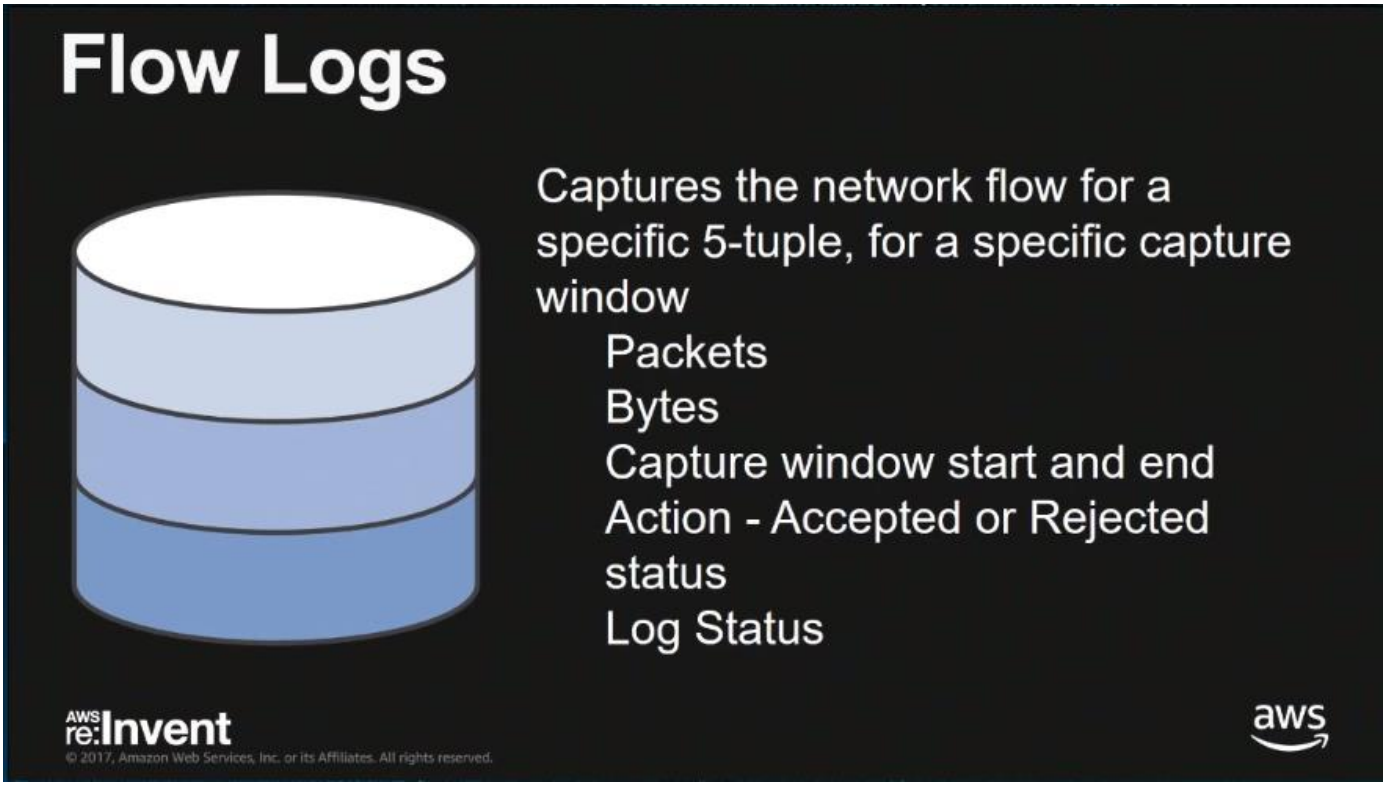

Let us see how the NLB API works

We have an EC2 instance running that is running a bash script that will create an NLB and the various resources like targets, target groups, and a listener. This will help us see the API structure that can be used for creating an NLB.

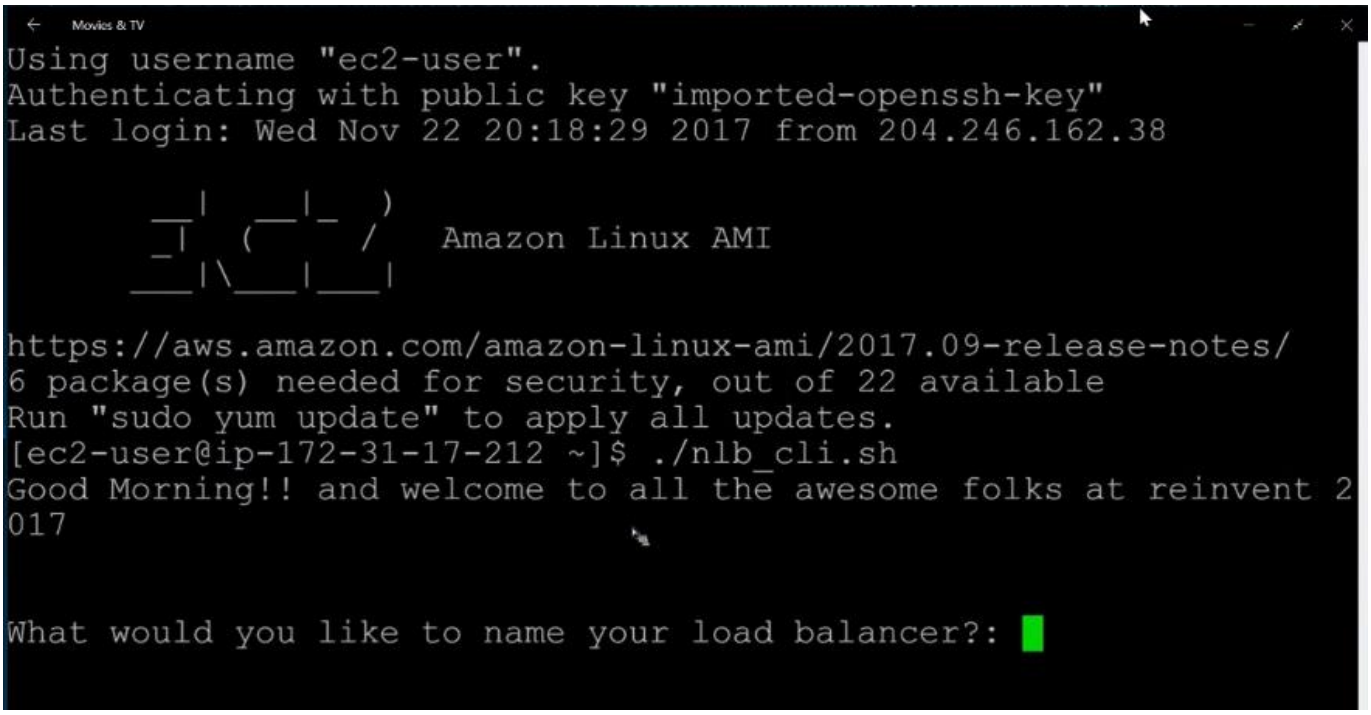

The script asks what name you want to give your NLB

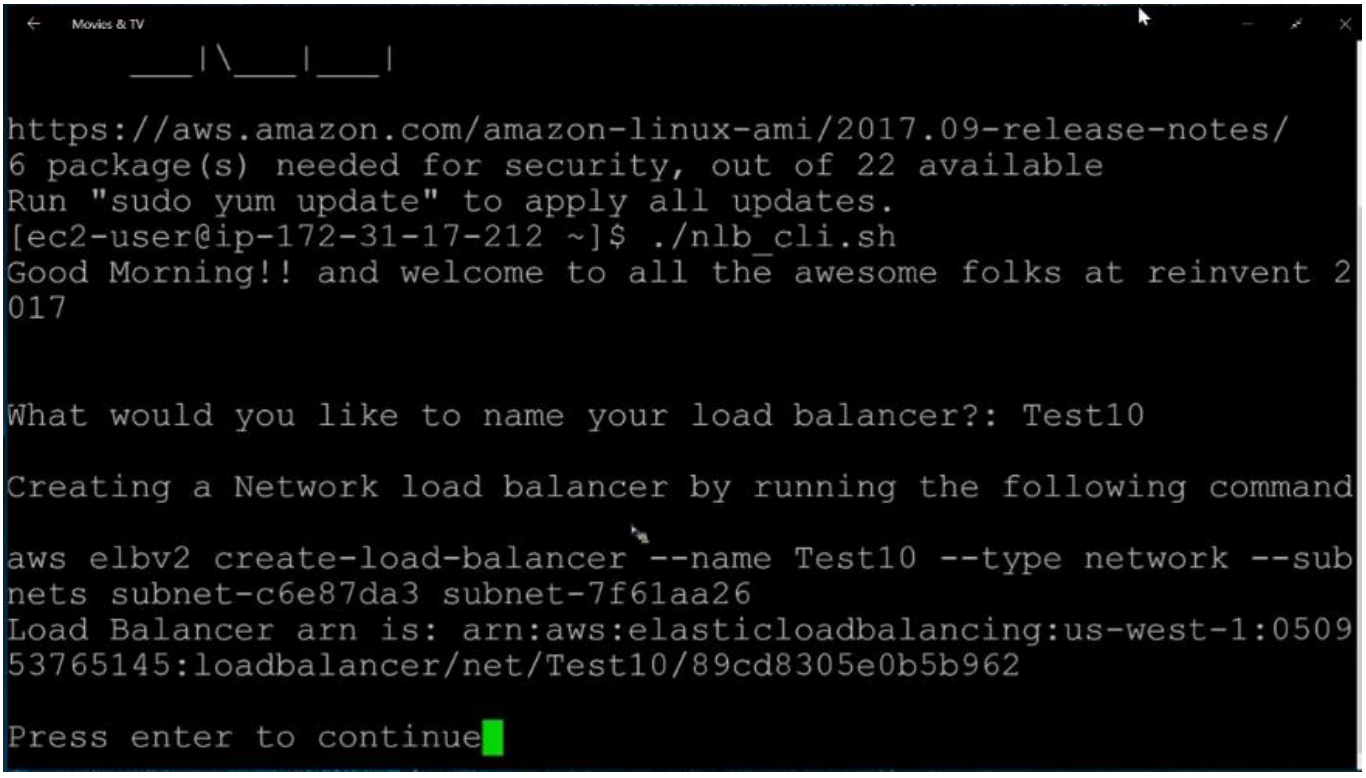

We choose the name ***Test10***, then the script issues the ***$ aws elbv2 create-load-balancer - -name Test10 - -type network - -subnets subnet-c6e87da3 subnet-7f61aa26*** command to create the NLB. We then get the NLB's ***arn*** value back after it has been created successfully. We have used the 2 subnets that we created in 2 different AZs for HA.

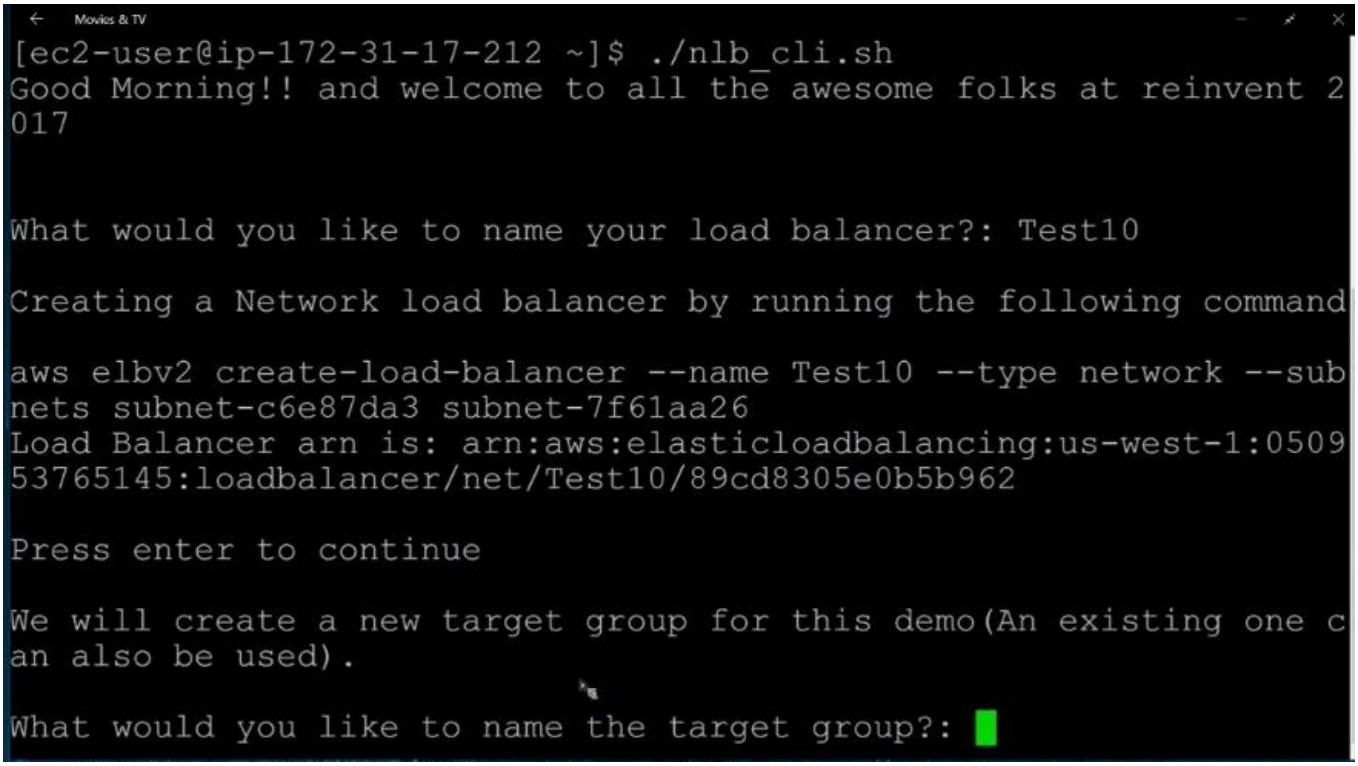

The script then asks if we want to create a new target group to use with a lister for this NLB.

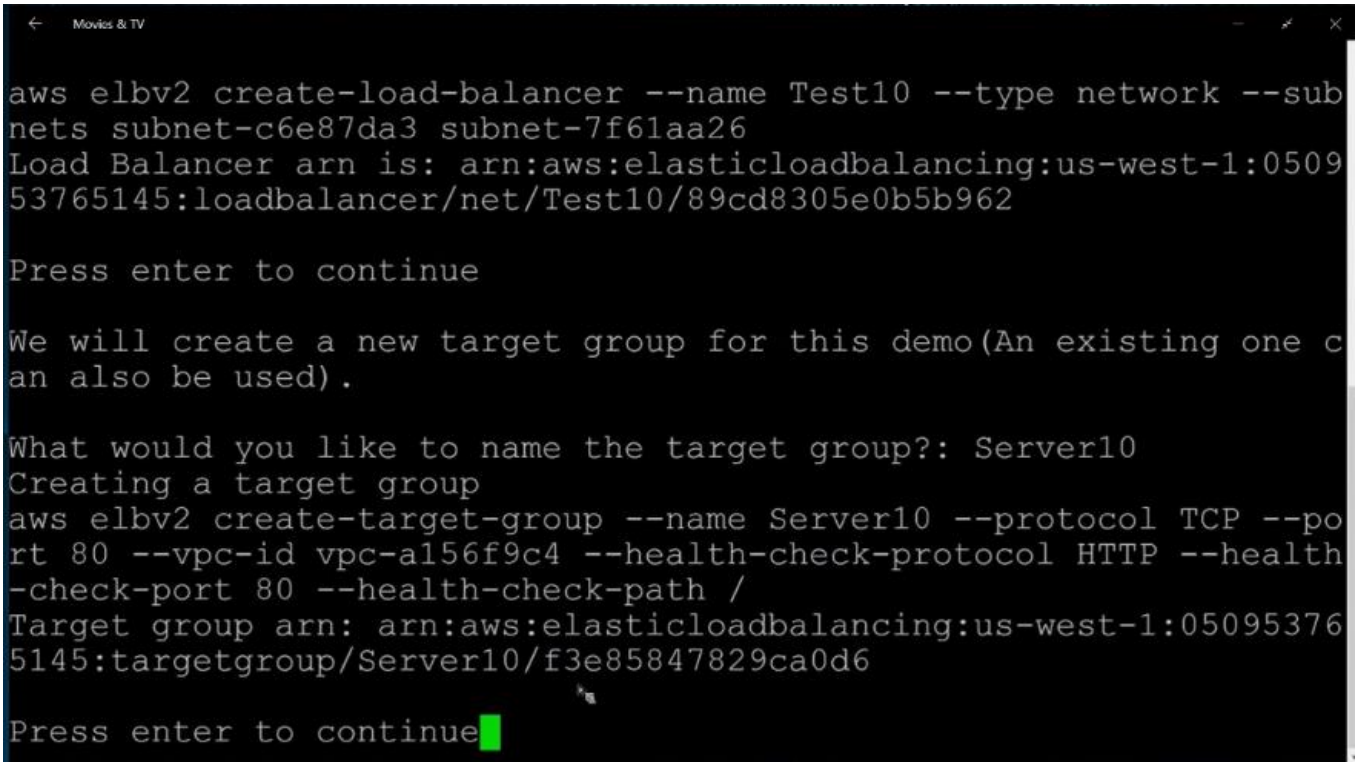

We choose the target group name Server10 and the script then creates a new target group for this NLB with the ***$ aws elbv2 create-target-group - -name Server10 - -protocol TCP - -port 80 - -vpc-id vpc-a156f9c4 - -health-check-protocol HTTP - -health-check-port 80 - -health-check-path /*** command. It then returns back to us the newly created target group's ***arn*** value.

Now it is going to ask that we ***register targets that we want to use with this NLB*** by showing us the list of all available targets that we can select to add to this NLB.

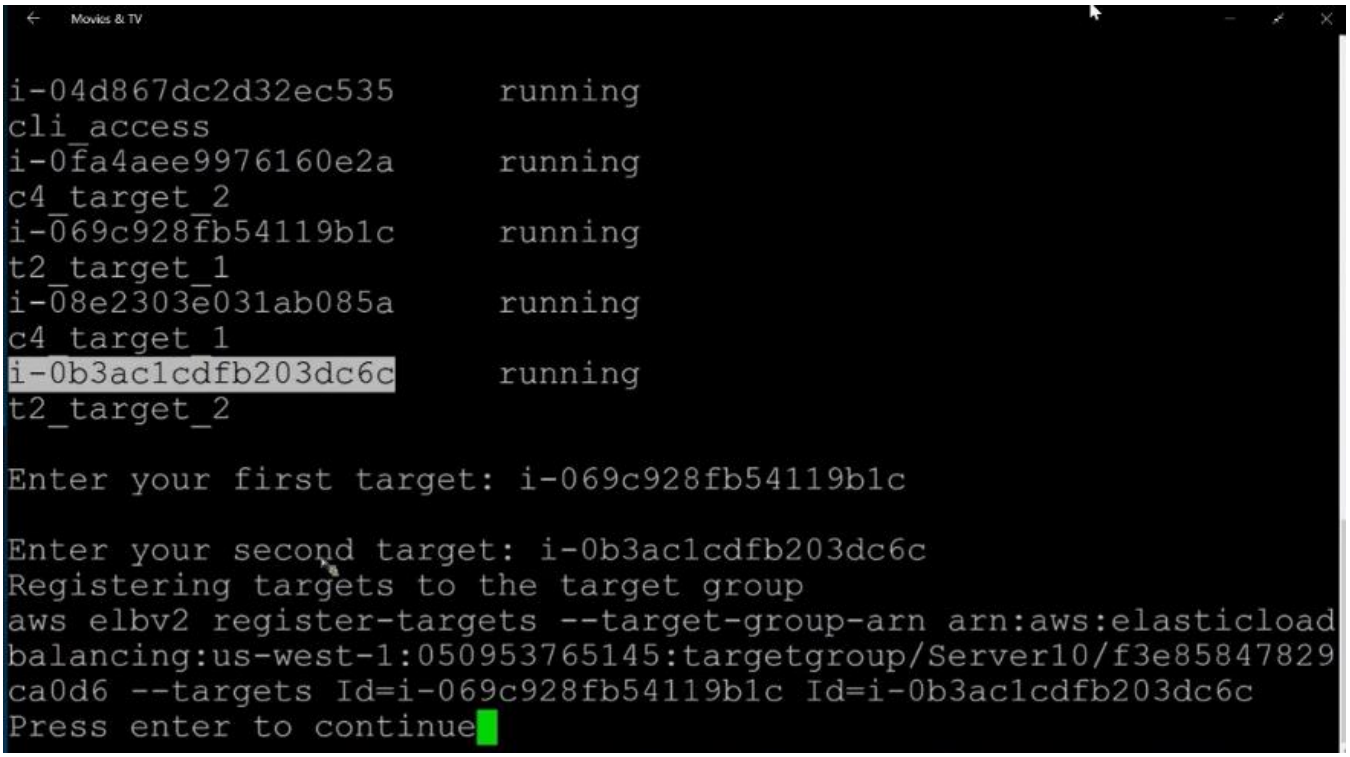

We select 2 target group to have them registered with this NLB as target groups, the script the issues the ***$ aws elbv2 register-targets - -target-group-arn arn:aws:elasticloadbalancing:us-west-1:0509673673676:targetgroup/Server10/f3e85847829ca0d6 - -targets Id=<target-group 1> Id=<target-group 2>*** command to have the 2 target groups registered as target groups for this NLB.

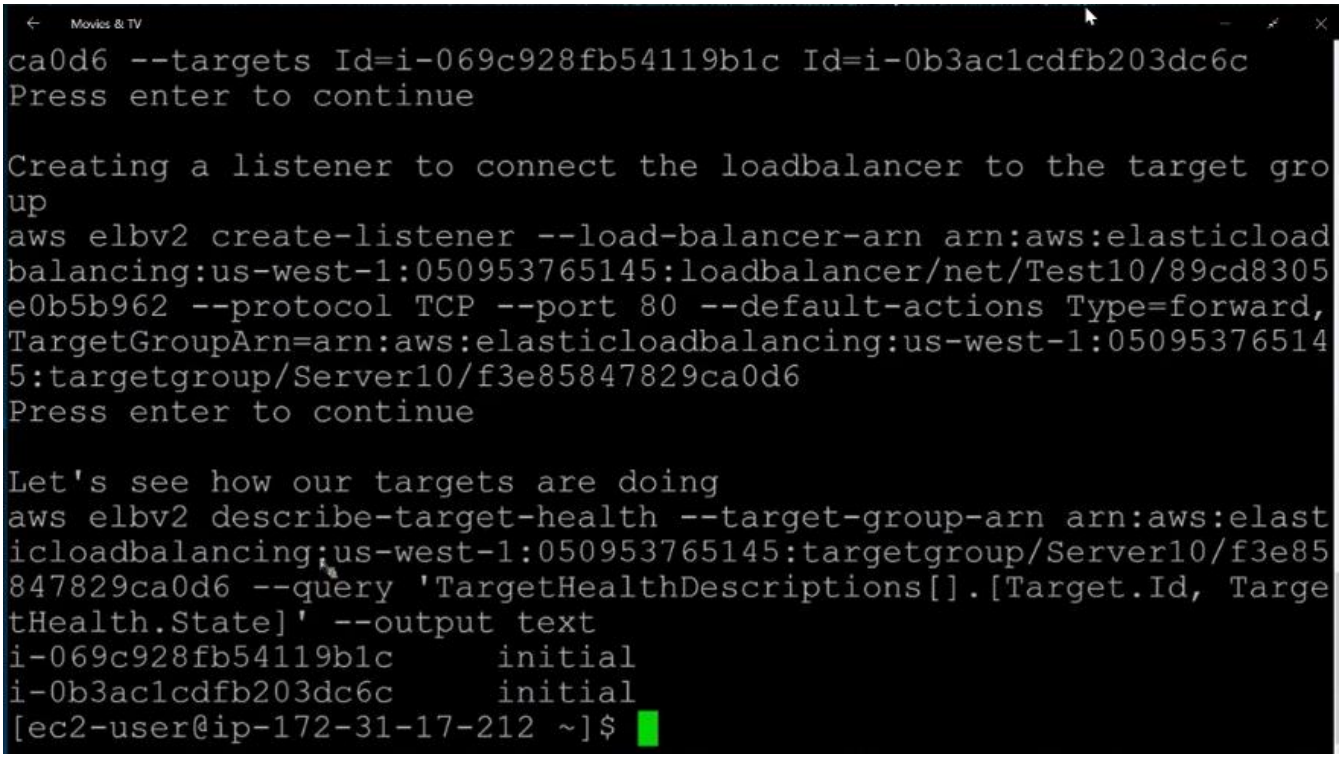

The script then goes and automatically creates a listener for us using the ***$ aws elbv2 create-listener - -load-balancer-arn arn:aws:elasticloadbalancing:us-west-1:0509673673676:loadbalancer/net/Test10/89cd8305e0b5b962 - -protocol TCP - -port 80 - -default-actions Type=forward,TargetGroupArn=arn:aws:elasticloadbalancing:us-west-1: 0509673673676:targetgroup/Server10/f3e85847829ca0d6*** command. We are using the default action type of 'forward' so that any connection that comes into our NLB at port 80 with the TCP protocol, will be forwarded to the target group called Server10. This then uses the 5-tuple hash algorithm and the sequenceId to pick the specific backend to send the request to within the Server10 target group.

We can then use the script to check the status of our targets using the ***$ aws elbv2 describe-target-health - -target-group-arn arn:aws:elasticloadbalancing:us-west-1: 0509673673676:targetgroup/Server10/ f3e85847829ca0d6 – query 'TargetHealthDescriptions[].[target.Id, TargetHealth.State]' - -output text*** command as above. We get the result back that the 2 target groups are been initialized and the healthchecks are still in progress for the targets in that particular target group.

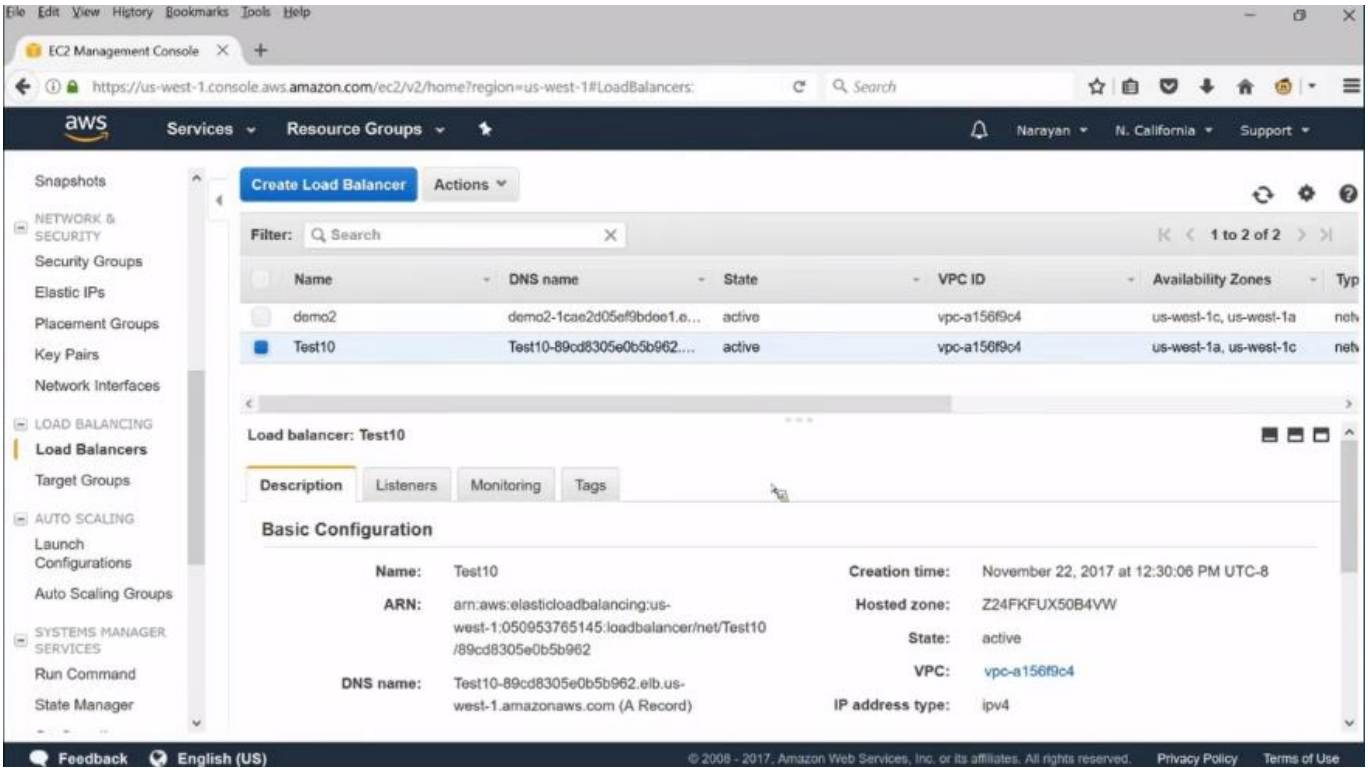

We can now go to the Load Balancers section of the EC2 service console and refresh the console, we can see that the Test10 NLB has been created

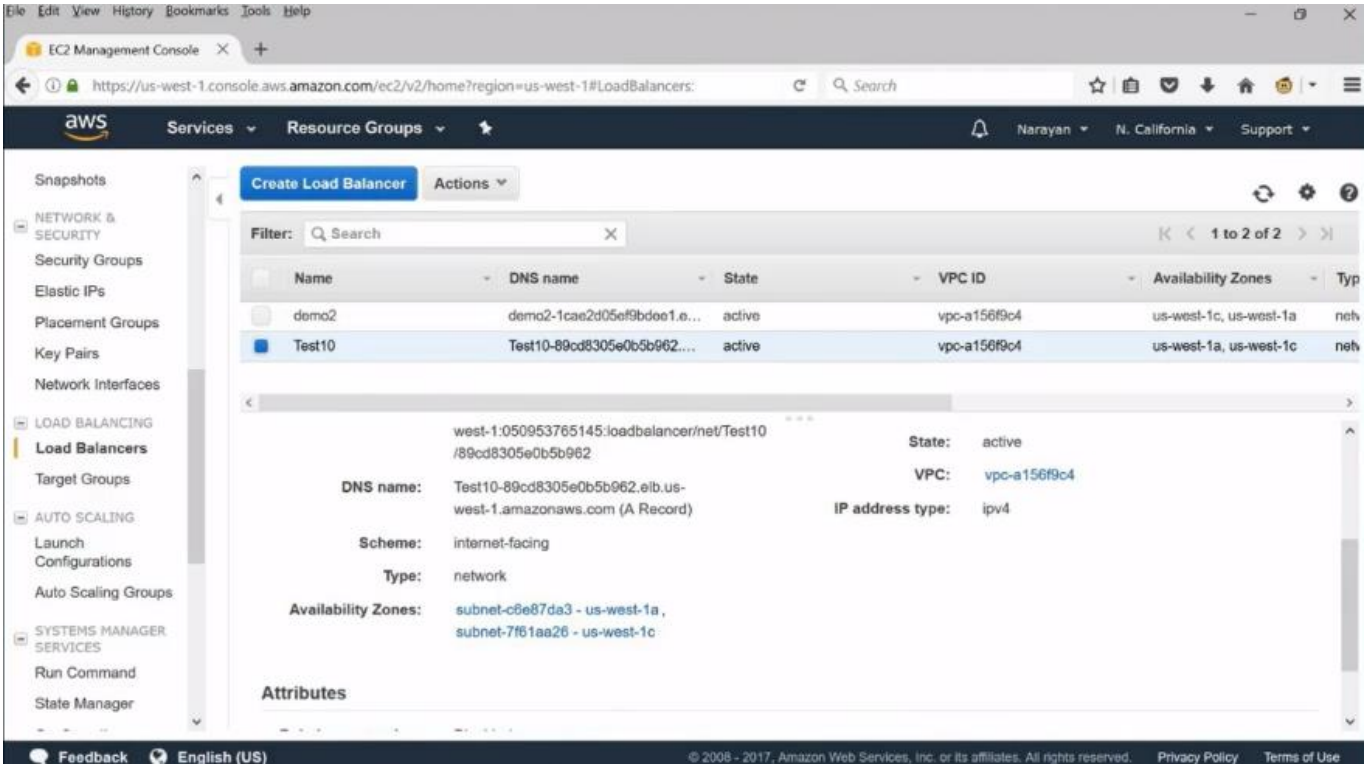

We see that it is a network type LB, internet-facing with 2 subnets in 2 different AZs,

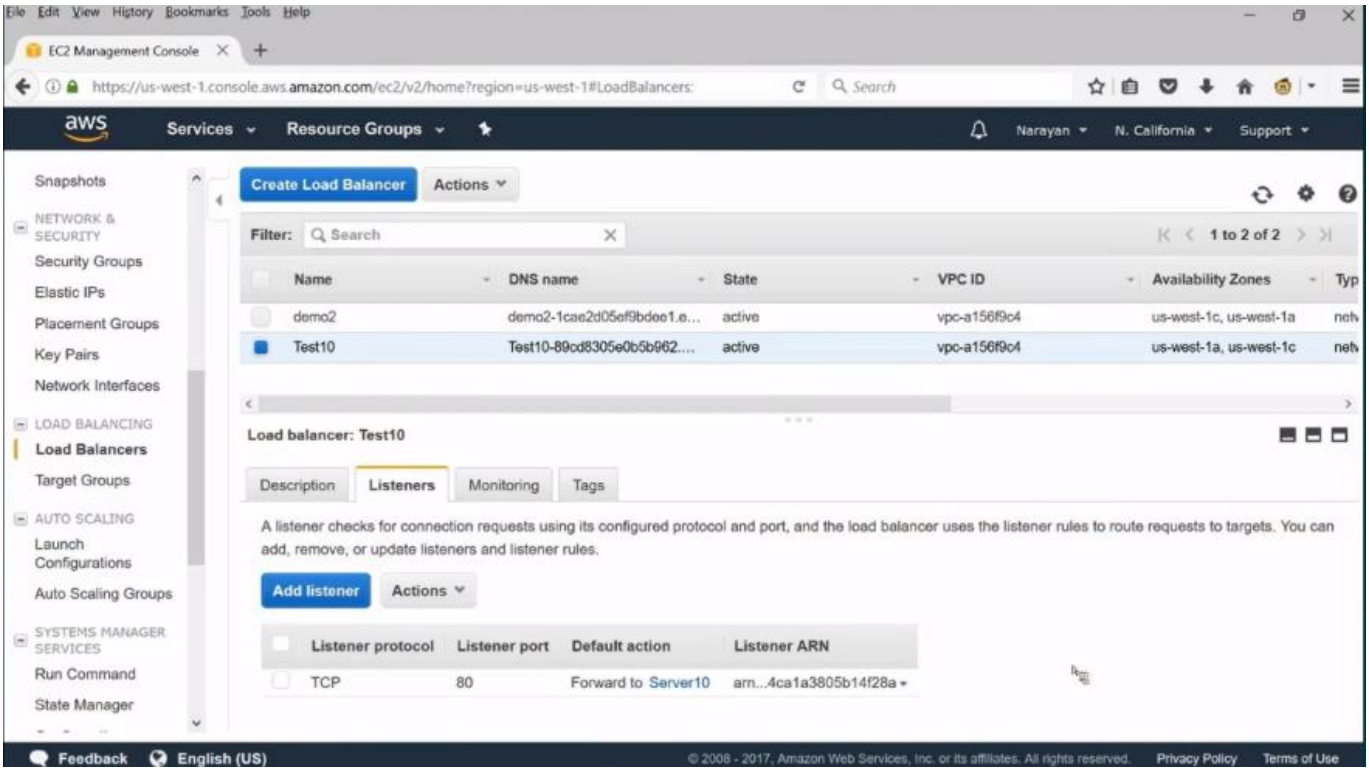

For this Test10 NLB, we also have a port 80 listener with TCP protocol that has the default action to forward to the target group called Server10.

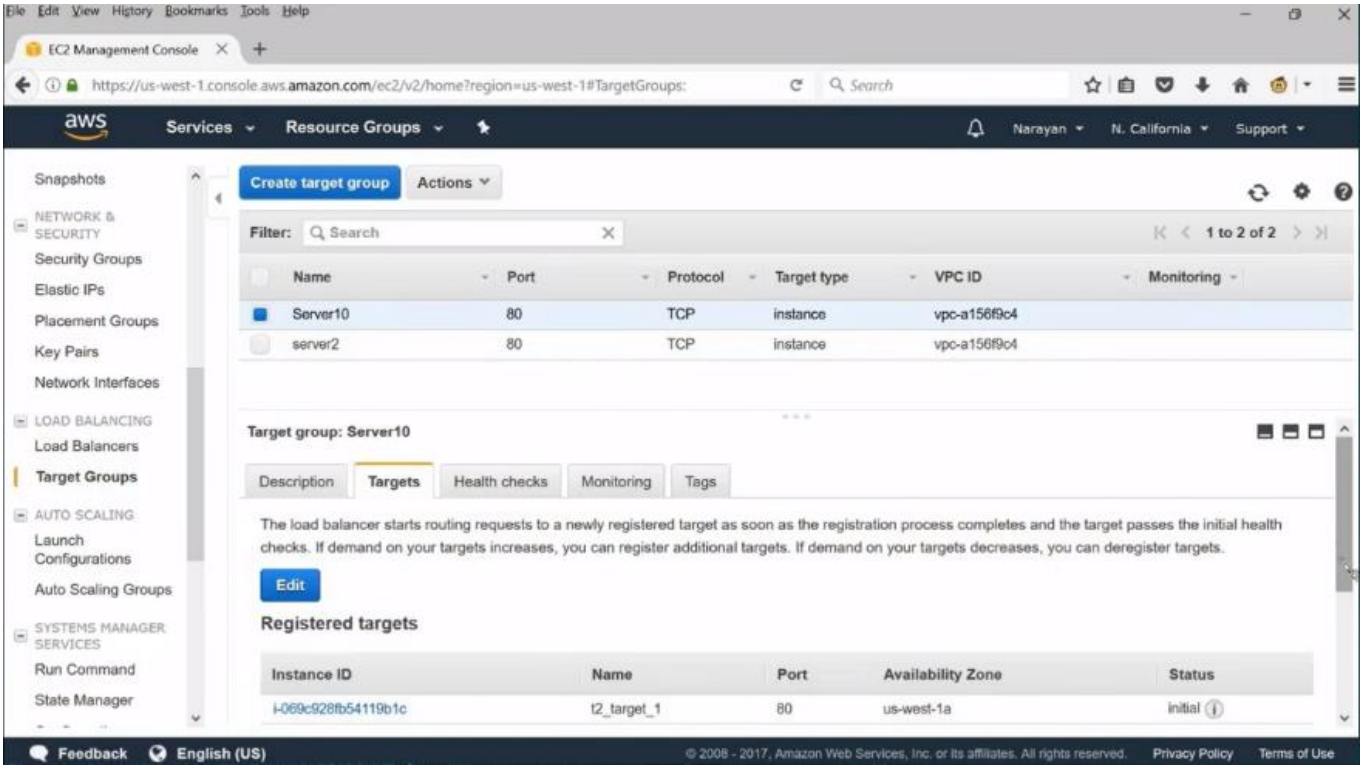

We can choose the Target Groups section to see the Server10 target group as above.

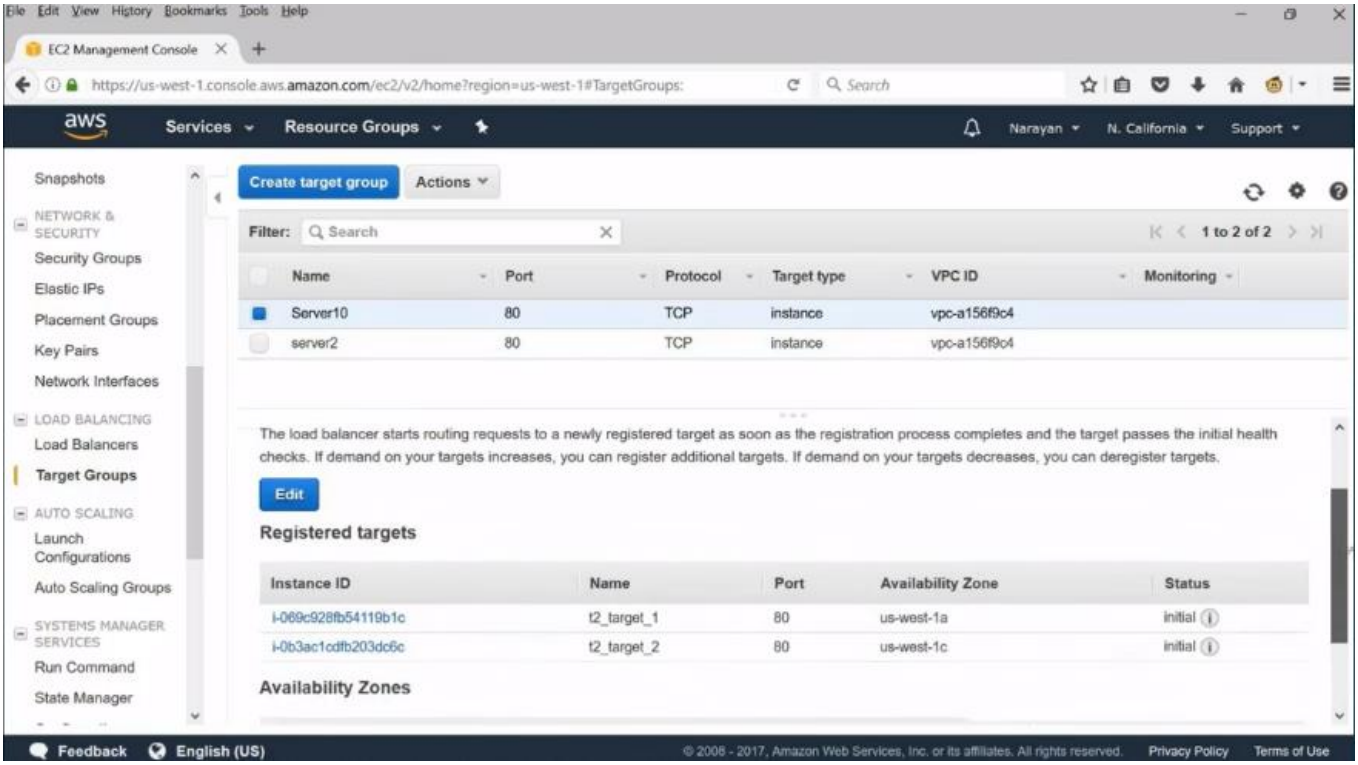

We can see the 2 registered targets that we created, one in each AZ. They both have their status as saying the health checks are still being initialized also.

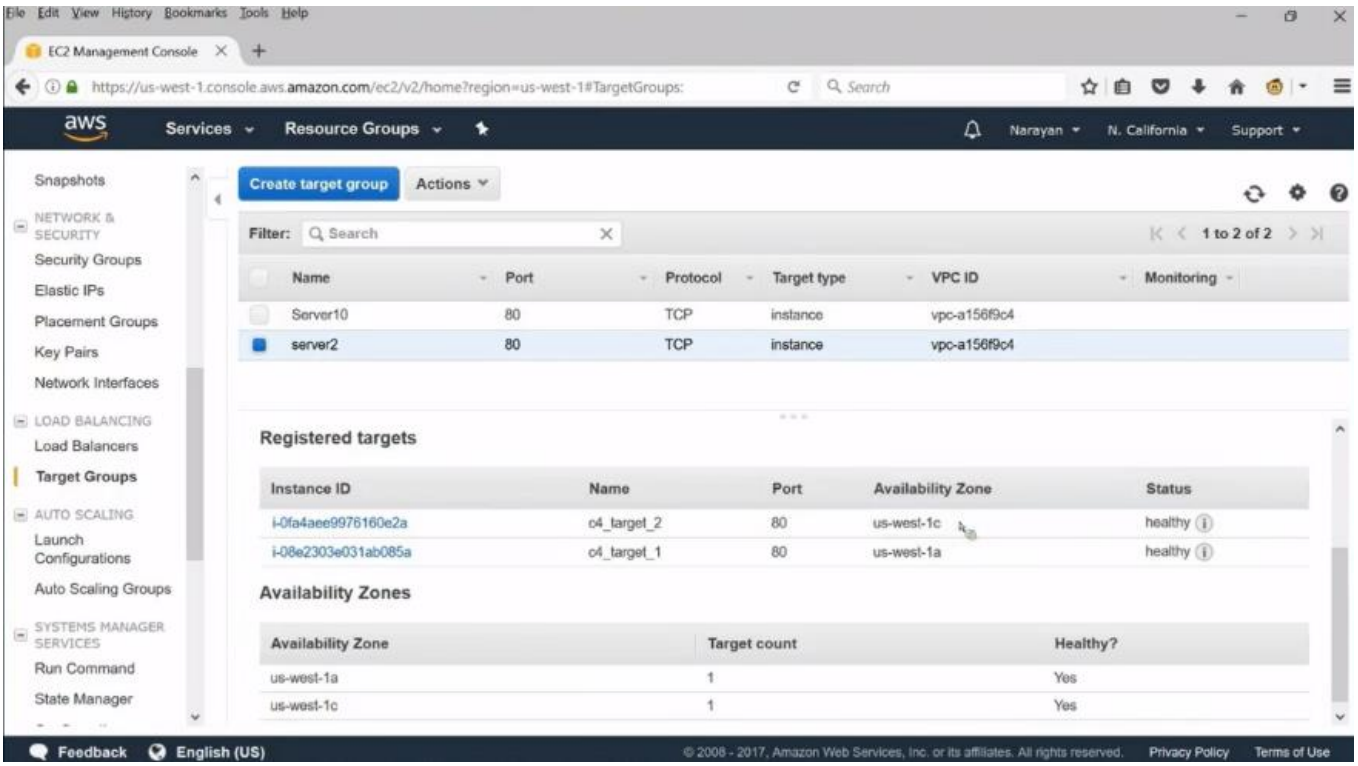

This is what will happen once the targets are completed and in a healthy state, their status will go from ‘initial’ to ‘healthy’ as seen above

Network Load Balancer pricing

With the Network Load Balancer, you only pay for what you use. You are charged for each hour or partial hour your Network load balancer is running and the number of Load Balancer Capacity Units (LCU) used per hour

- $0.0225 per Network Load Balancer-hour (or partial hour) (US-EAST-1)
- $0.006 per LCU-hour (or partial hour) (US-East-1)

Hourly charge is 10% less expensive than Classic Load Balancer; Data Processing charge is 25% less expensive than Classic and Application Load Balancer; reducing the cost for virtually all of our customers

Load balancer capacity units

An LCU measures the dimensions on which the Network Load Balancer processes your traffic (averaged over an hour). The three dimensions measured are:

- New connections: up to 800 new connections per second
- Active connections: up to 100,000 active connections
- Bandwidth: Up to 2.22 mbps (1 GB per hour)

You are charged only on the dimension with the highest usage over the hour.

Migrating to Network Load Balancer

Migration is as simple as creating a new Network Load Balancer, registering targets and updating DNS to point at the new CNAME.

Classic Load Balancer to Network Load Balancer migration utility:
https://github.com/aws/elastic-load-balancing-tools

AWS re:Invent

aws

When should I use Network Load Balancer?

	Application Load Balancer	Network Load Balancer	Classic Load Balancer
Protocol	HTTP, HTTPS,HTTP/2	TCP	TCP, SSL, HTTP, HTTPS
SSL offloading	✓		✓
IP as Target	✓	✓	
Path-based routing, Host-based routing	✓		
Static IP		✓	
WebSockets	✓	✓	
Container Support	✓	✓	

AWS re:Invent

aws

For TCP in VPC, use Network Load Balancer.

For all other use cases in VPC , use Application Load Balancer

For Classic networking, use Classic Load Balancer

aws

LOGGLY

Cloud-based log management
Founded in 2009
Based in San Francisco
10,000+ customers
Startups to Fortune 500

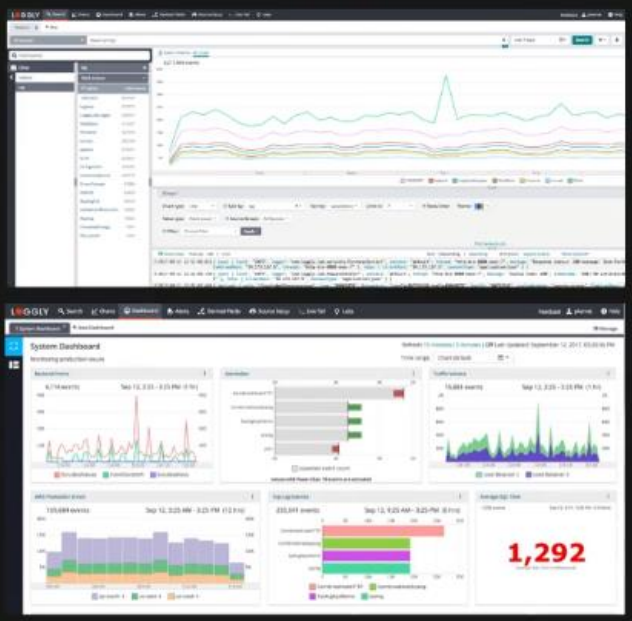

LOG MANAGEMENT IS A BIG DATA PROBLEM

- Massive incoming event stream
- Fundamentally multi-tenant
- Scalable framework for analysis
- Near real-time indexing
- Near real-time search
- Time-series index management
- Logs are TLDR
- Traffic is unpredictable

Big data services require a very resilient architecture.

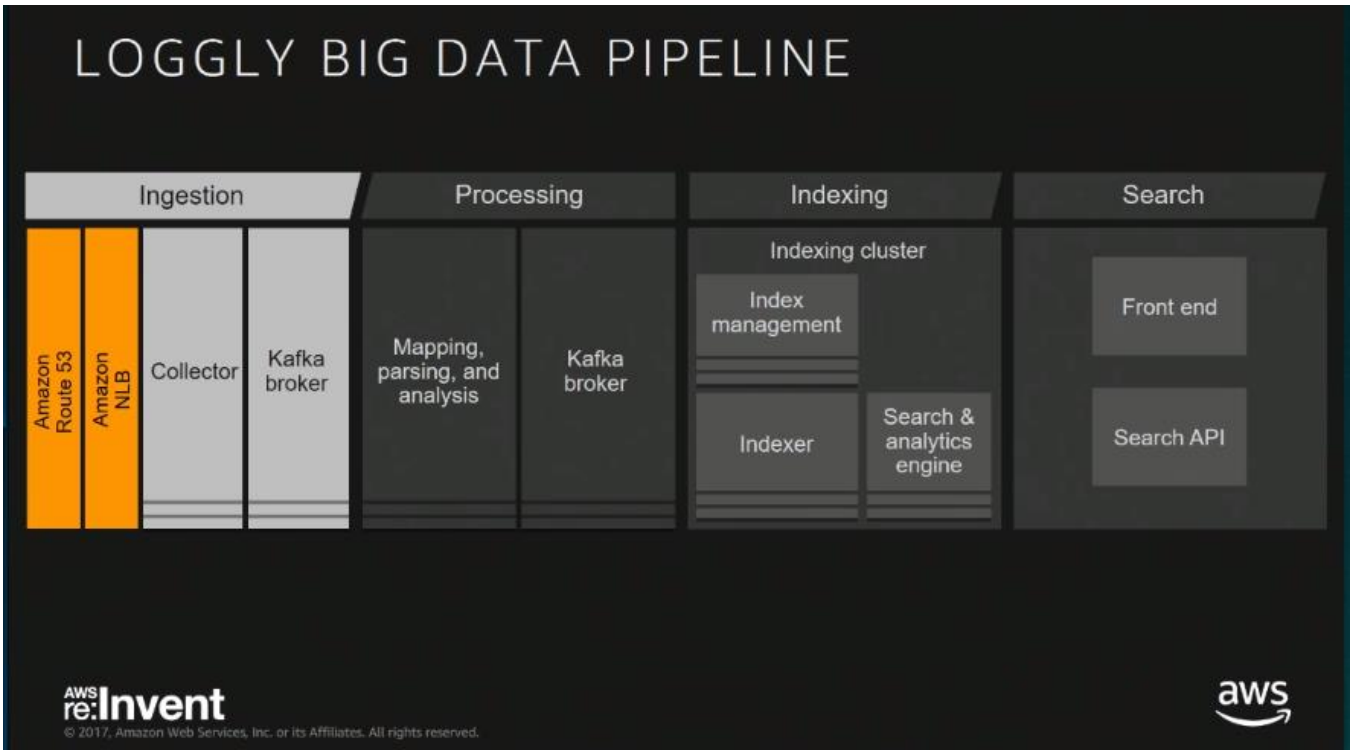

Loggly uses AWS route53 and NLB products to load balance data ingestion traffic, this is the entry point for customer logs into the pipeline. NLB evenly distributes the incoming data to a fleet of collectors where data validation is done, then pass on the data for processing, indexing, and make search available for customers. Routing and distributing loads to healthy collectors in the pipeline is one of the most important services.

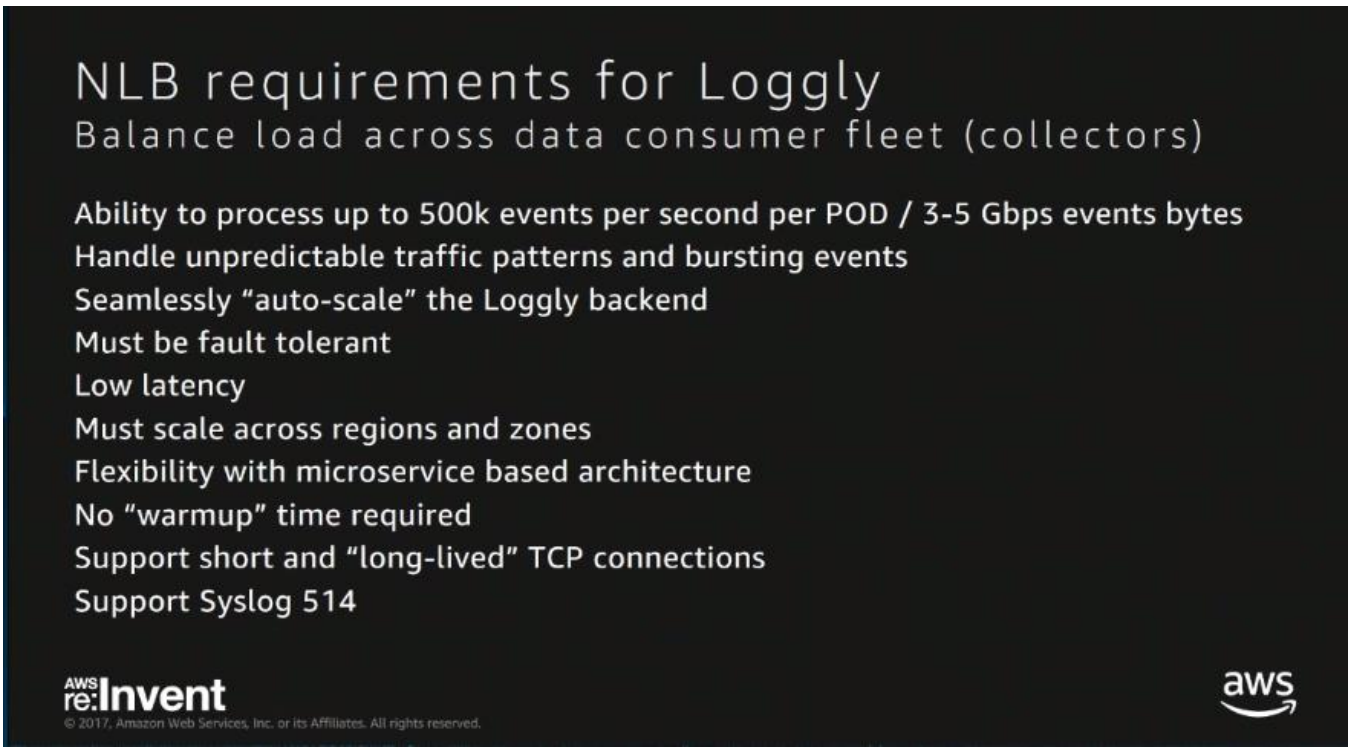

This system can adapt to different customer streaming profiles

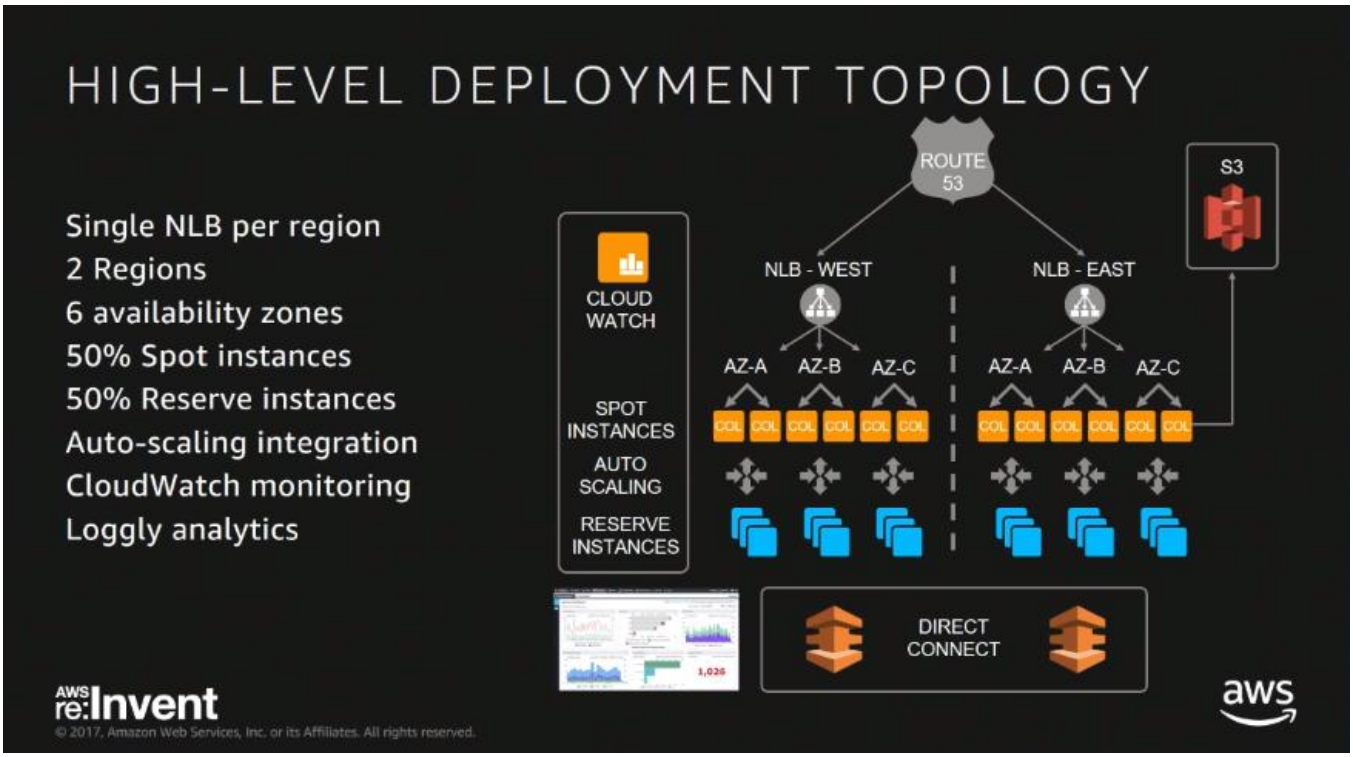

Loggly operates 2 NLBs, one in the west and the other in the east. We have 50% spot fleet instances upfront supported by another 50% spot fleet in the back to balance risk and cost for this topology. We use CloudWatch and ***Loggly Management Analytics*** to provide critical insights into our ***Ingestion Pipeline*** and Consumer Endpoint Health.

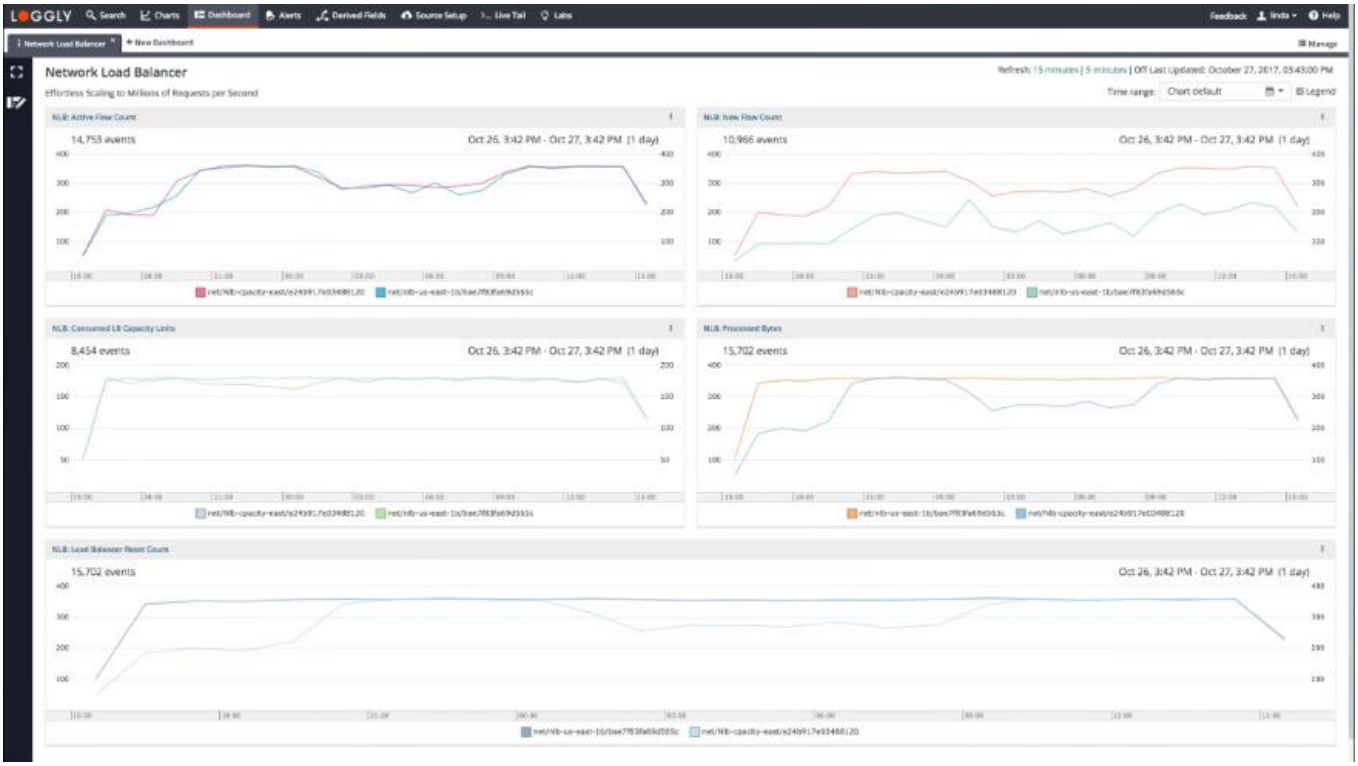

This is one of the view used by operations and network engineers to check the health of incoming data flow, and the critical data patterns

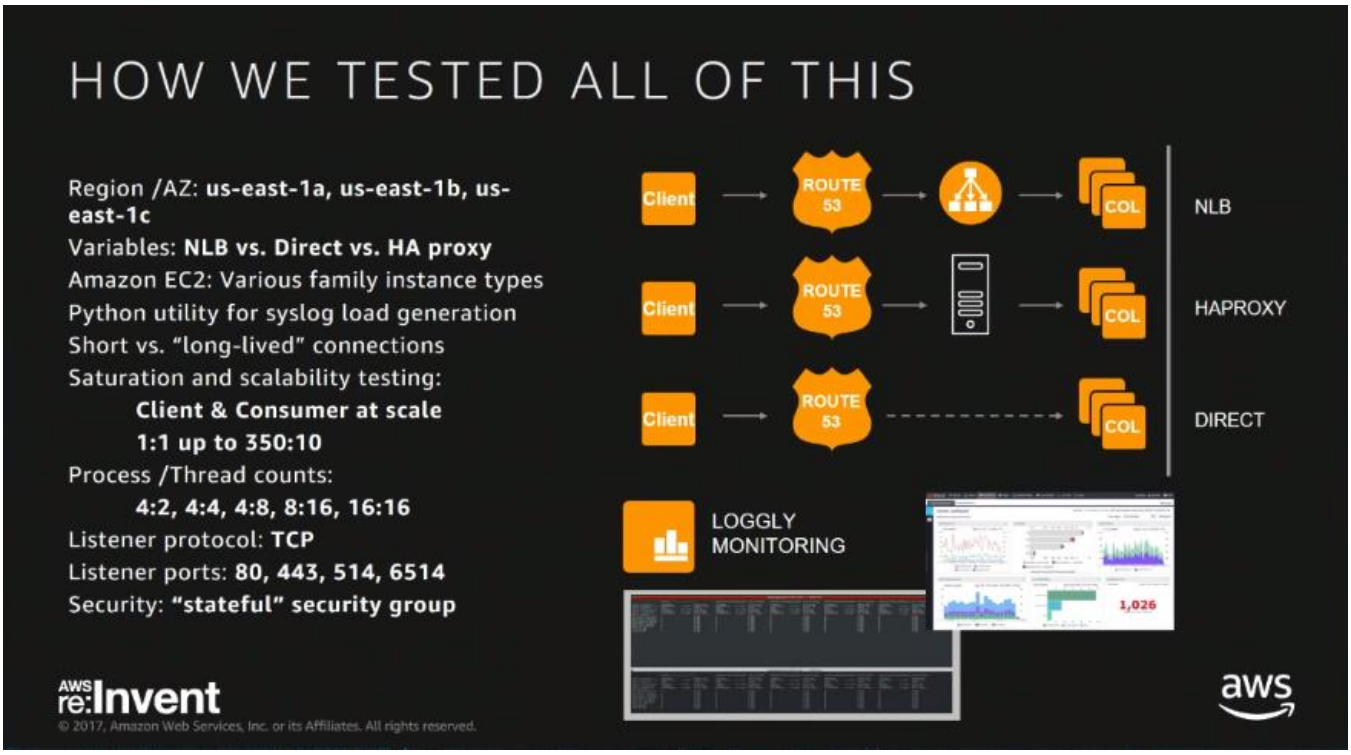

To test all these during the PoC phase. We started with historical review of all the incoming traffic volume data, throughput and performance benchmarks, then we used these historical data to establish specific NLB acceptance criteria that needs to be met. Our use case scope was to test both syslog long-lived connections and short-lived HTTP connections. We also need a balanced comparative testing perspective, so we setup 3 test harness variables Route53, NLB, and a generic HAProxy setup for sanity testing.

We started with 1 client sending simulated load to 1 connector through the NLB to establish a baseline, this establish the direct baseline for the HAProxy and the Direct testing. We then ramped up to 100s of clients sending loads to 10s of collectors, adjusting processing and threads to simulate customer profiles. The test results verified that the NLB can consistently manage load as the direct testing could with negligible differences.