





The pace of technology innovation is relentless, especially at AWS. Designing and building new system architectures is a balancing act between using established, production-ready technologies while maintaining the ability to evolve and take advantage of new features and innovations as they become available. In this session, learn how Amazon Game Studios built a flexible analytics pipeline on AWS for their team battle sport game, Breakaway, that provided value on day one, but was built with the future in mind. We will discuss the challenges we faced and the solution we built for ingesting, storing and analyzing gameplay telemetry and dive deep into the technical architecture using AWS many services including Amazon Kinesis, Amazon S3, and Amazon Redshift. This session will focus on game analytics as a specific use case, but will emphasize an overarching focus on designing an architectural flexibility that is relevant to any system.



# Analytics

Measure



Understand



Improve



Engagement



Monetization




Data-driven

**AWS**  
**re:Invent**

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



## Analytics: Scientific method



**AWS**  
**re:Invent**


© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



## Analytics: Level design



## Analytics: Heatmaps



Fewer Deaths


More Deaths

**AWS**  
**re:Invent**

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



## Analytics: Heatmaps




**AWS**  
**re:Invent**

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.


**aws**

## Analytics: Heatmaps







## Analytics: Heatmaps



## Analytics: Heatmaps




## Flexibility: Background




Back in 2015...

## Flexibility: Background



## Flexibility: Motivation

**flex·i·bil·i·ty** (\, flek-sə-'bi-lə-tē\) (*noun*)



***"Characterized by a ready capability to adapt to new, different, or changing requirements"***

- Ambiguous requirements
- Changing requirements
- Evolving tech landscape
- The "awesome prototype conundrum"

**AWS re:Invent**

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.




# Architecture




Onward to the architecture...

## Architecture: High-level



## Produce: Characteristics

### Telemetry events



**AWS re:Invent**

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



## Produce: Format





## Produce: Sample event


```
{  
  "event_version": "1.0",  
  "event_id": "4e96de3b-2bb5-4aca-b631-f3827317d90a",  
  "event_timestamp": 1505491200685,  
  "event_type": "player_death",  
  
  "app_name": "game_name",  
  "app_version": "1.0",  
  "client_id": "b2228ae9-10a3-4dc6-bfda-53591d34d065",  
  
  "level_id": "map_name",  
  "position_x": 78.35,  
  "position_y": 39.192  
}
```

**AWS**  
**re:Invent**

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



## Produce: Data sources



**AWS**  
**re:Invent**

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



## Produce: Servers



### Authoritative source of

- Gameplay
- Performance (server)



Trusted (mostly)

## Produce: Clients

### Authoritative source of

- Engagement
- Performance (client)
- Gameplay (local)



Untrusted



**AWS**  
**re:Invent**


© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.




# Architecture: High-level




## Ingest



## Ingest



# Ingest: Consumers



AWS  
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



# Ingest: Producers




AWS  
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



# Ingest: Producers




**AWS re:Invent**

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



# Ingest: Producers



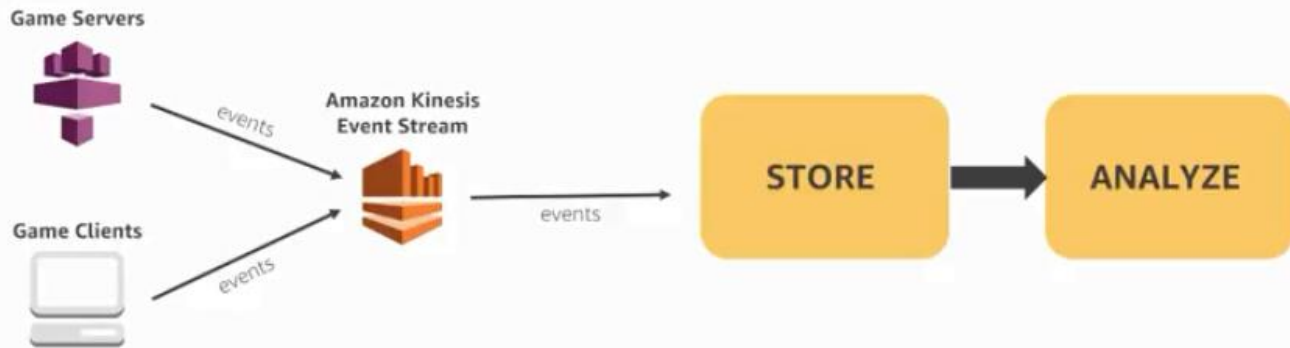
**AWS re:Invent**

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.





## Architecture: High-level



## Architecture: High-level



## Store




**Amazon S3**




**Cold data**

|            |                 |
|------------|-----------------|
| Retention  | Years           |
| Access     | Slow            |
| Turnaround | N/A             |
| Structure  | Semi-structured |
| Duplicates | OK              |

## Store: Cold data



## Store: Cold data




**AWS re:Invent**

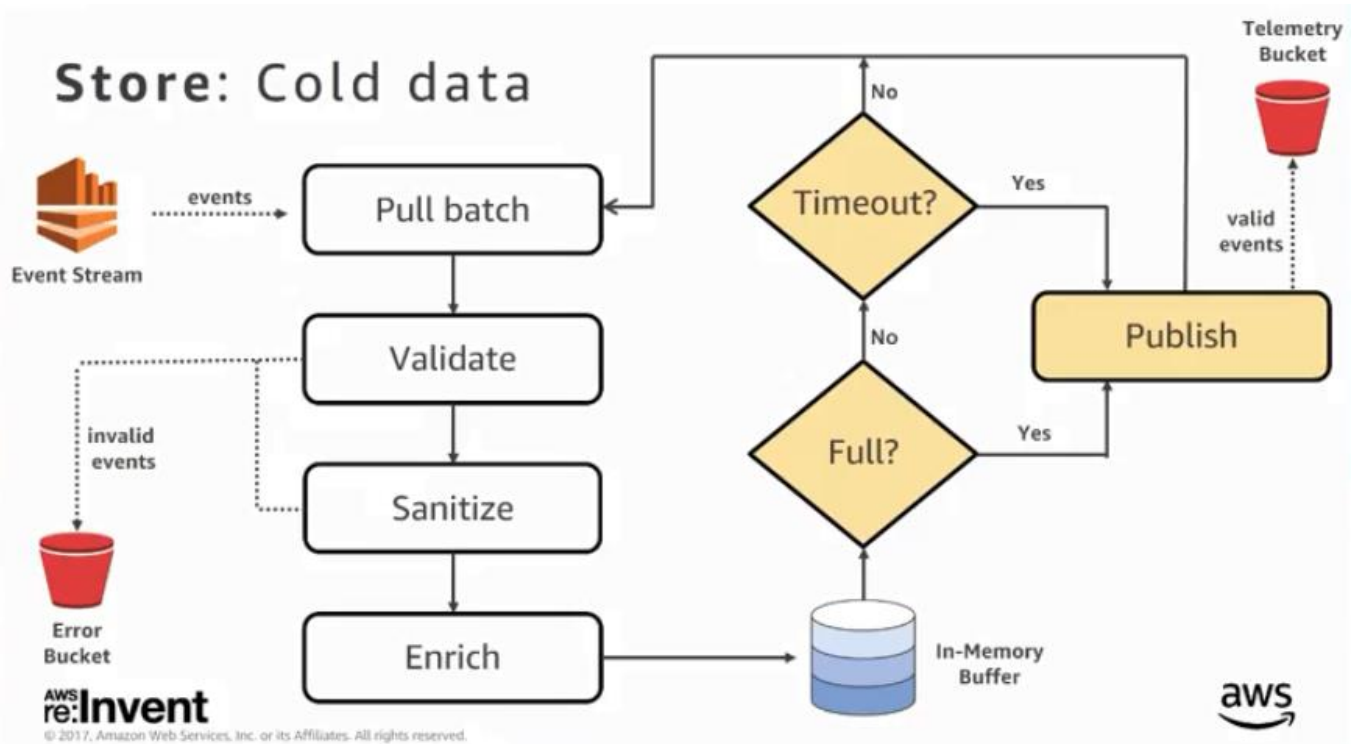
© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



## Store: Cold data



This is what the S3 App is doing.



We are using 100MB buffer size to batch up the data into a single JSON file and store it up in S3

## Store: Cold data



**AWS re:Invent**

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



## Store



**Amazon  
Redshift**



**Warm data**

|            |            |
|------------|------------|
| Retention  | 6 months   |
| Access     | Fast       |
| Turnaround | 1 hour     |
| Structure  | Structured |
| Duplicates | No         |

**AWS re:Invent**

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



Newer data like the last 6 months is more relevant to our gaming scenario, we also want structured data, filtering out duplicates. We use Redshift also because it is SQL compatible

## Store: Warm data

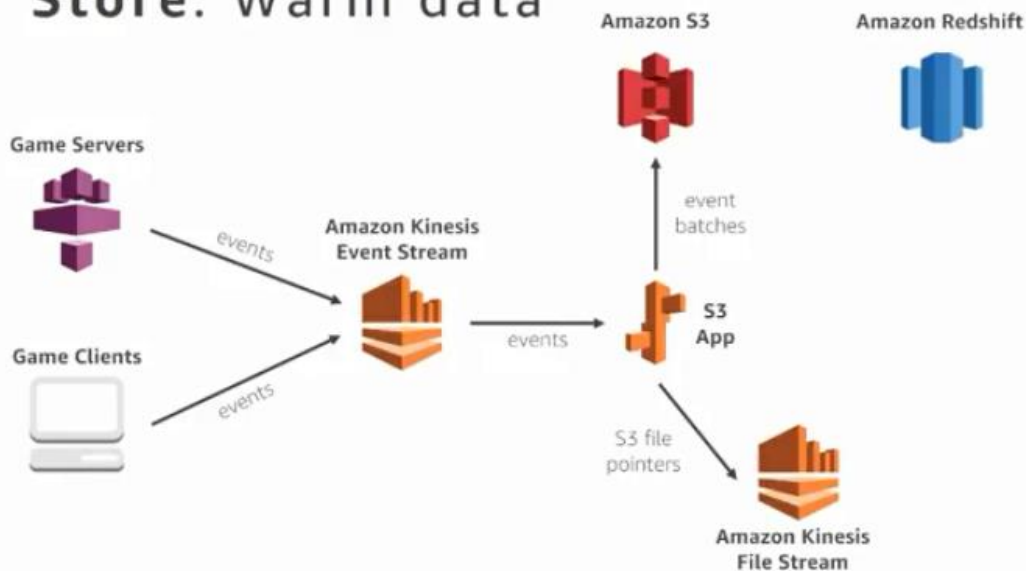


AWS  
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

aws

## Store: Warm data



AWS  
re:Invent


© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

aws

Every time we send one of the 100MB batched JSON file to S3, we also send a secondary file pointer data to another Kinesis stream



## Store: Warm data



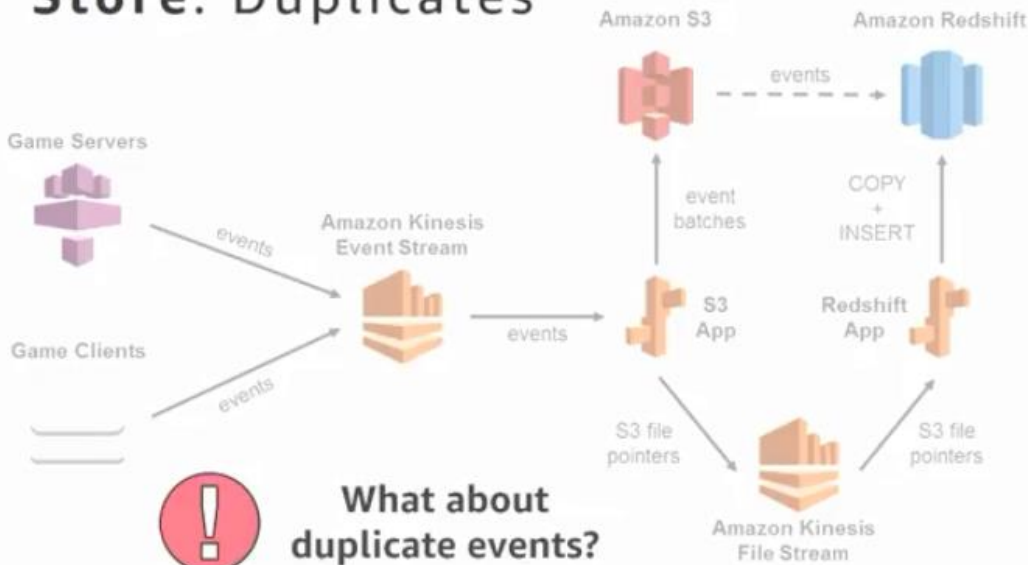
AWS  
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



We then have another Elastic Beanstalk application that collects data from the 2<sup>nd</sup> Kinesis stream, buffer the file pointers data up in-memory, when it deems the data large enough, it will push the data into our Redshift cluster using the Redshift COPY command to copy all the data up from S3 using the file pointers.

## Store: Duplicates

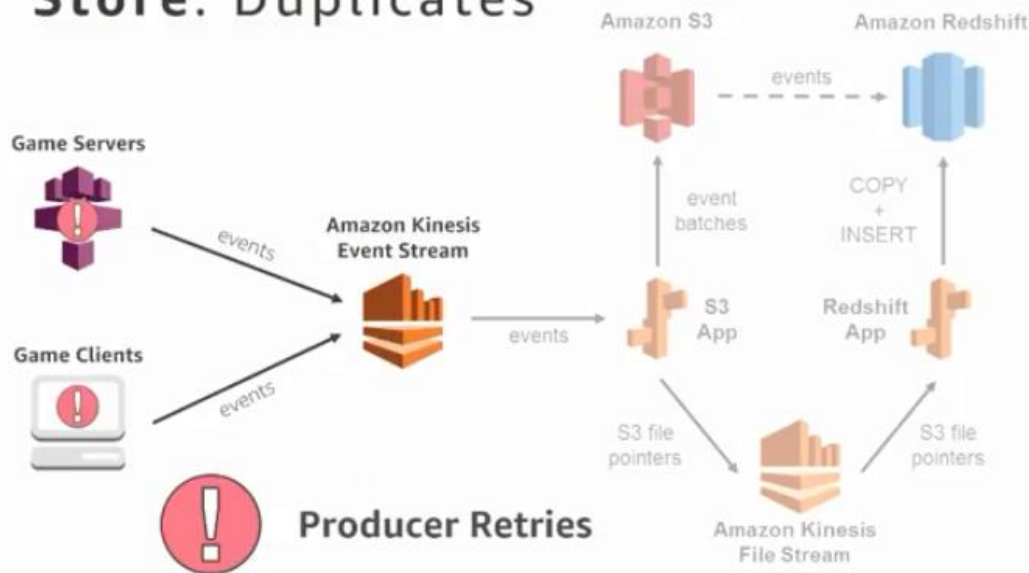


AWS  
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



## Store: Duplicates



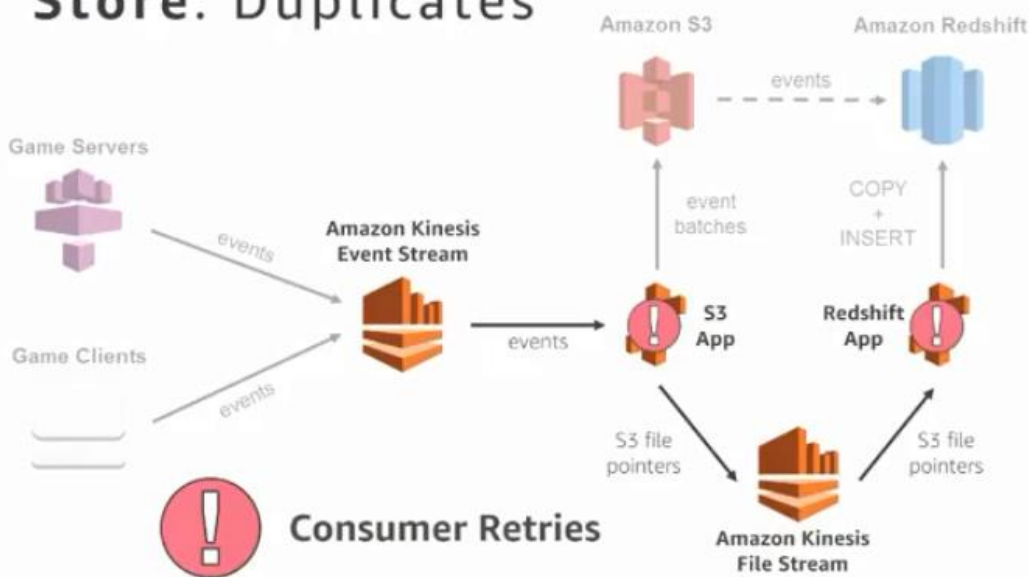
AWS  
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



We get duplicate event data from Producer retries when network is flaky

## Store: Duplicates



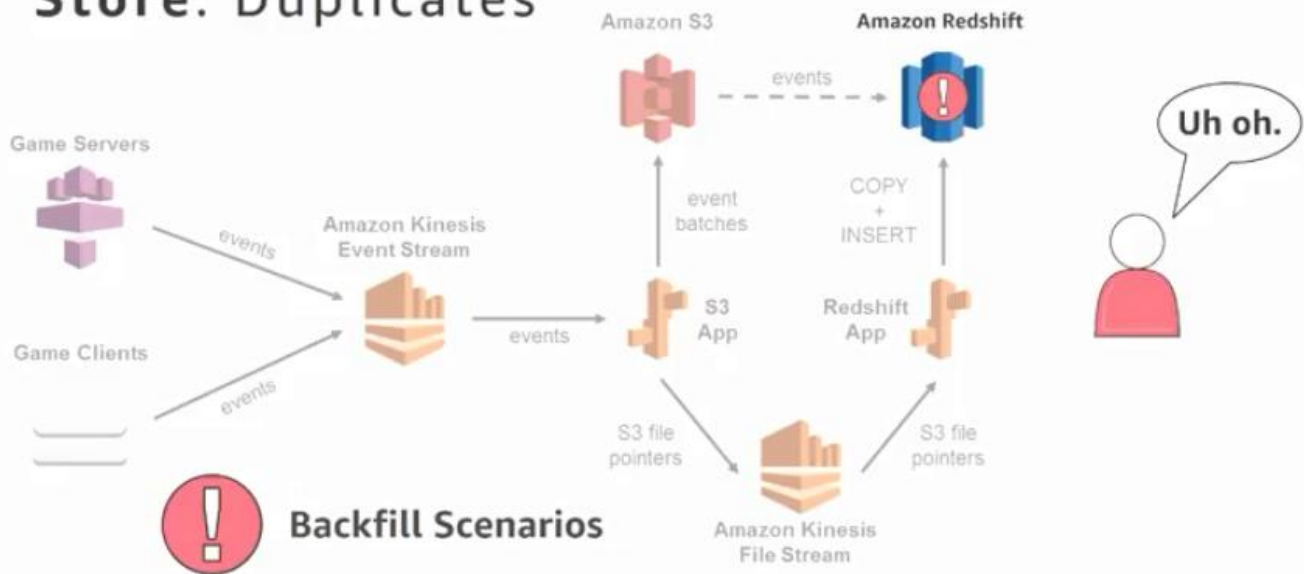
AWS  
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



Consumer retries can also create duplicate data being generated and stored

## Store: Duplicates



AWS re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



Duplicates can also be caused due to human errors

## Store: Dedupe



Amazon S3



AWS re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



## Store: Dedupe



**Pull S3 file pointers from Amazon Kinesis**



Once the Redshift elastic beanstalk app deems it has enough file pointers batched up in-memory, it's going to get the files from S3

## Store: Dedupe

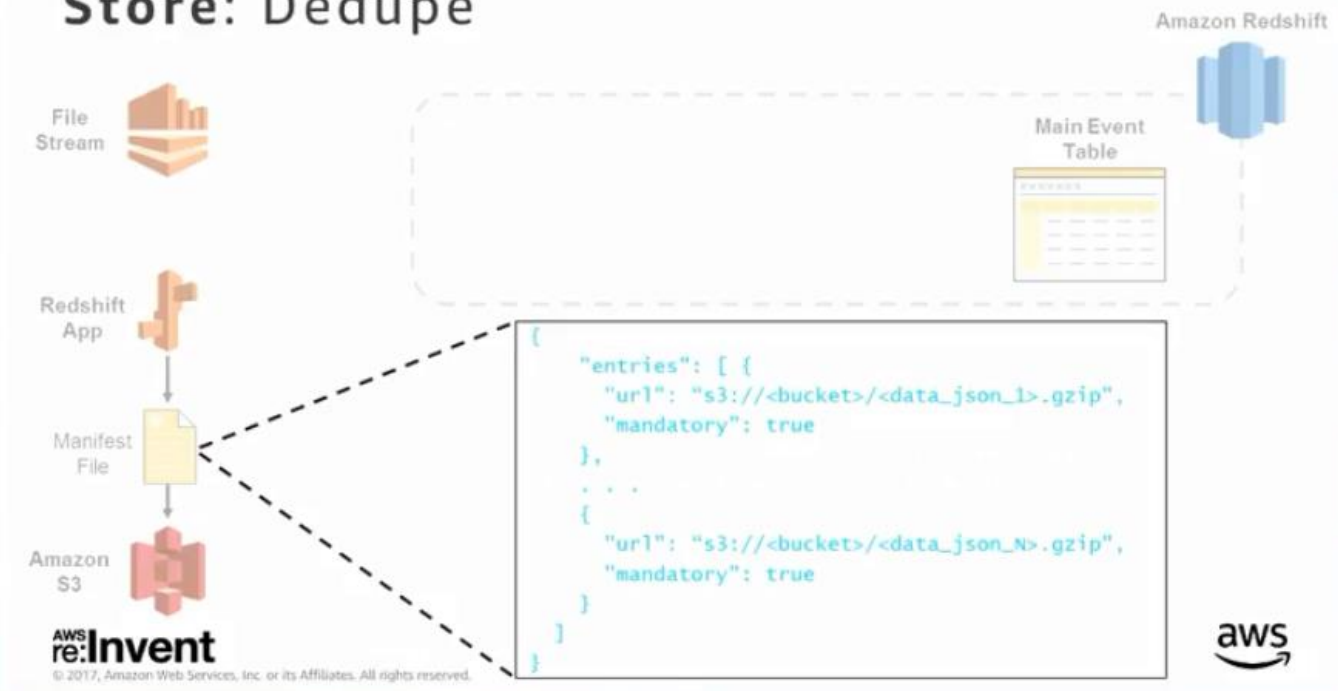


**Write a manifest file to Amazon S3**



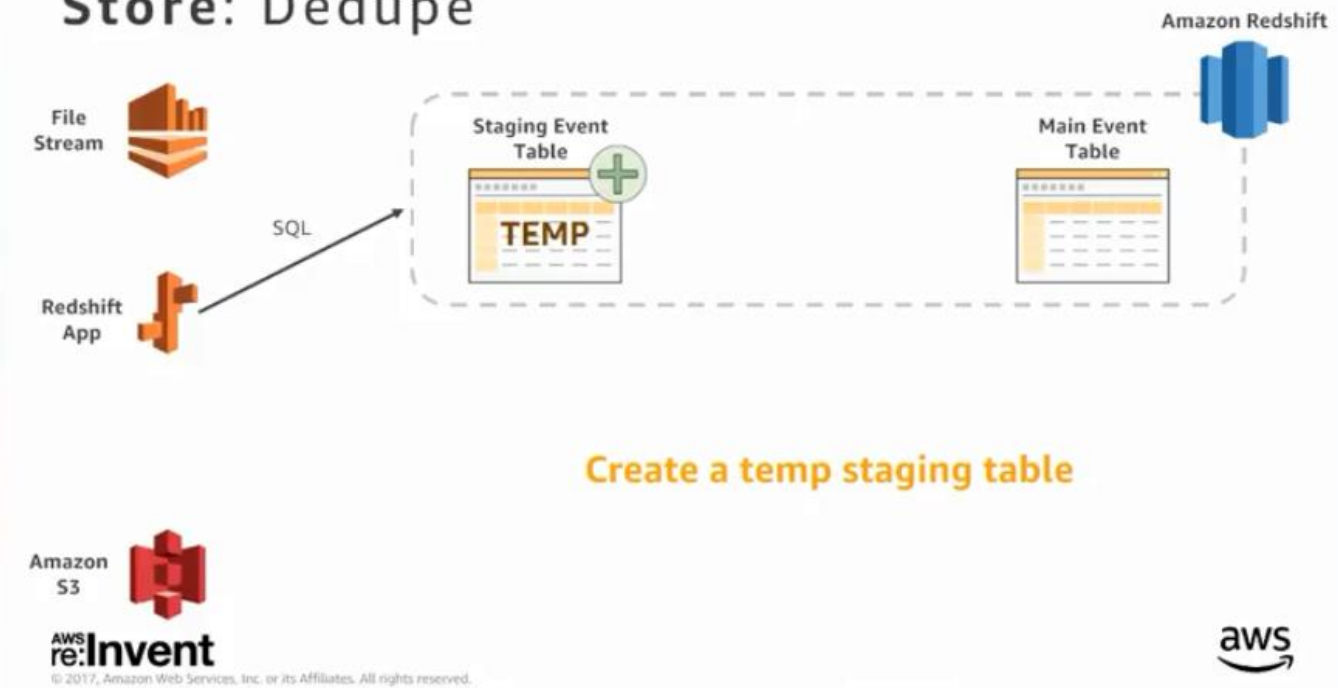
It then writes up all the file pointer data that it has in its memory and push them into a single manifest file

## Store: Dedupe



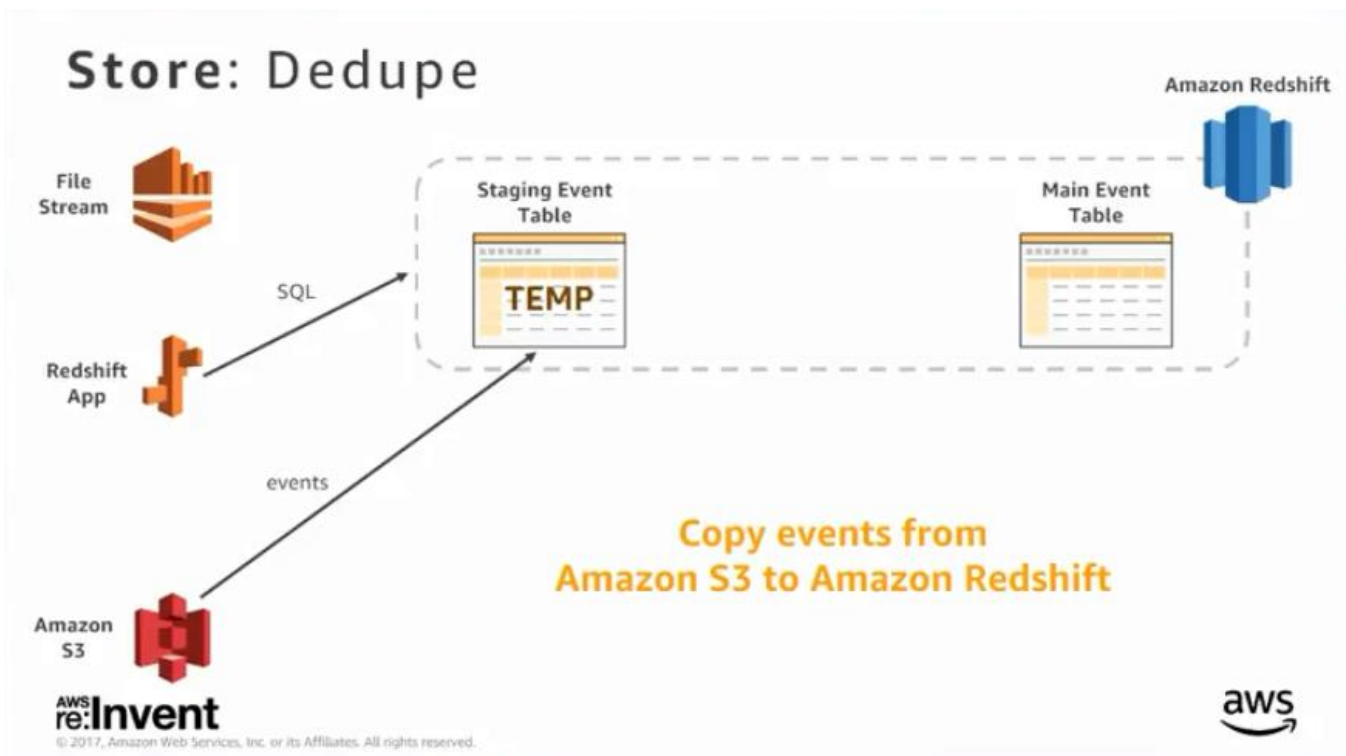
Manifests in Redshift are a set of S3 URIs or pointers combined with a mandatory entry that tells it whether or not this load should fail if that file is not there.

## Store: Dedupe

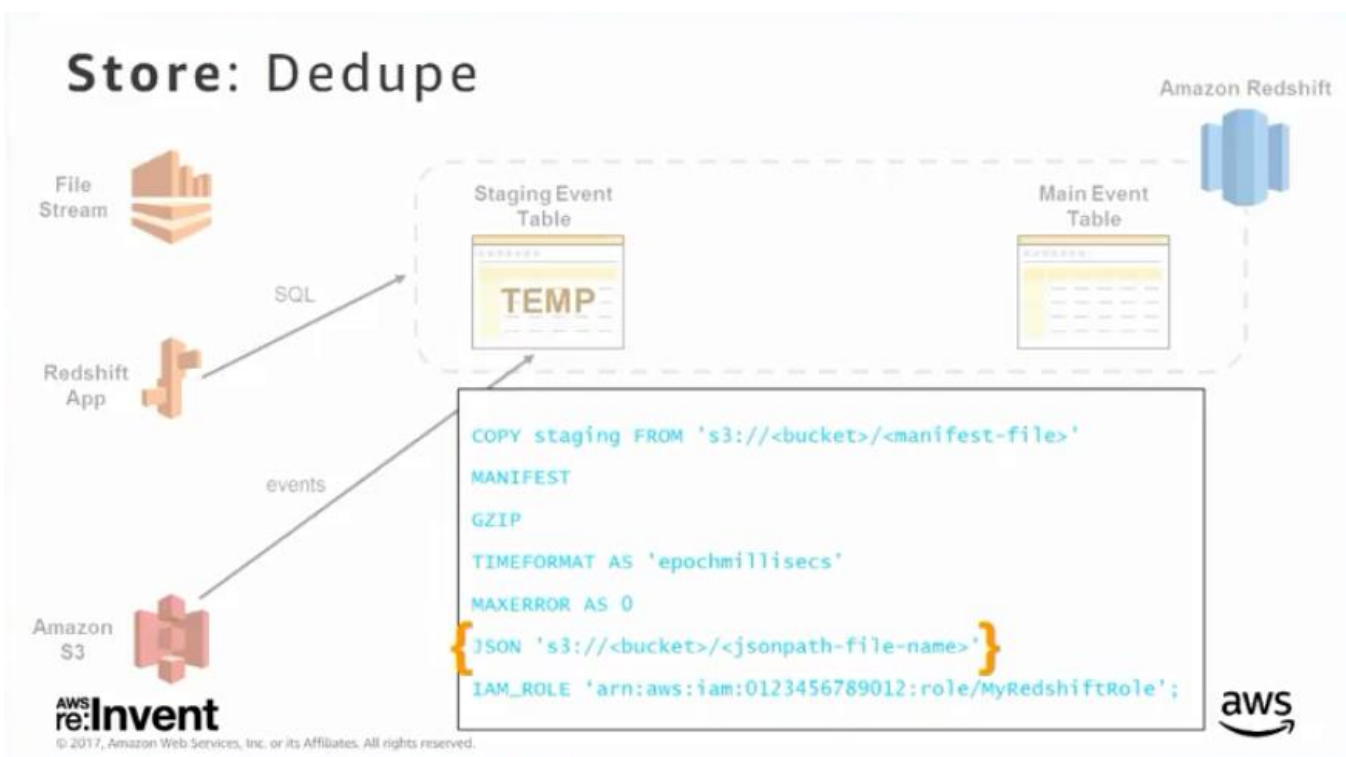


We create an empty temporary table in Redshift to put the data into





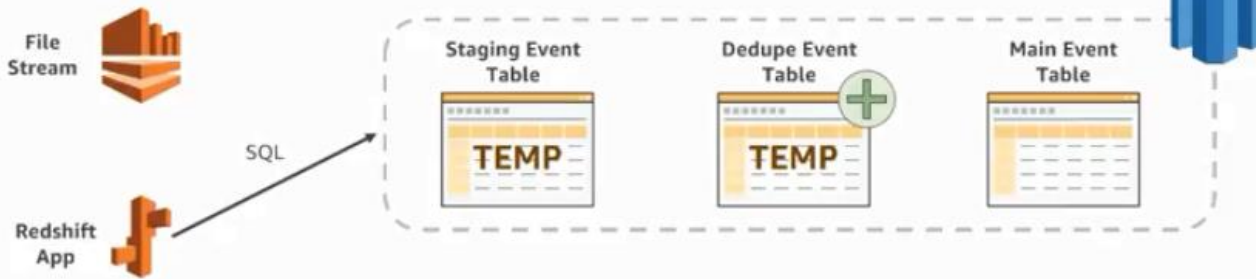
We then use the Redshift COPY command to get all the data out from S3



This is what a Redshift COPY command looks like, we tell Redshift we are loading JSON data. The **jsonpath** lets you map the fields on your input JSON onto database columns, it also lets you choose and pick only the fields that you care about mapping into database columns and also lets you order the database columns. This helps us decouple our database schema from our Redshift schema

## Store: Dedupe

Amazon Redshift



Create a temp dedupe table

Amazon  
S3

**AWS re:Invent**

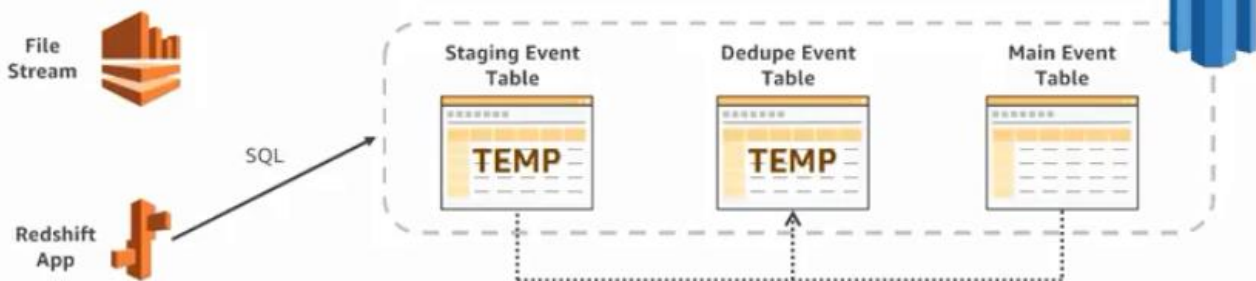
© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

aws

We now have all the events in the staging table including duplicates if present

## Store: Dedupe

Amazon Redshift



Insert only new events into dedupe table

Amazon  
S3

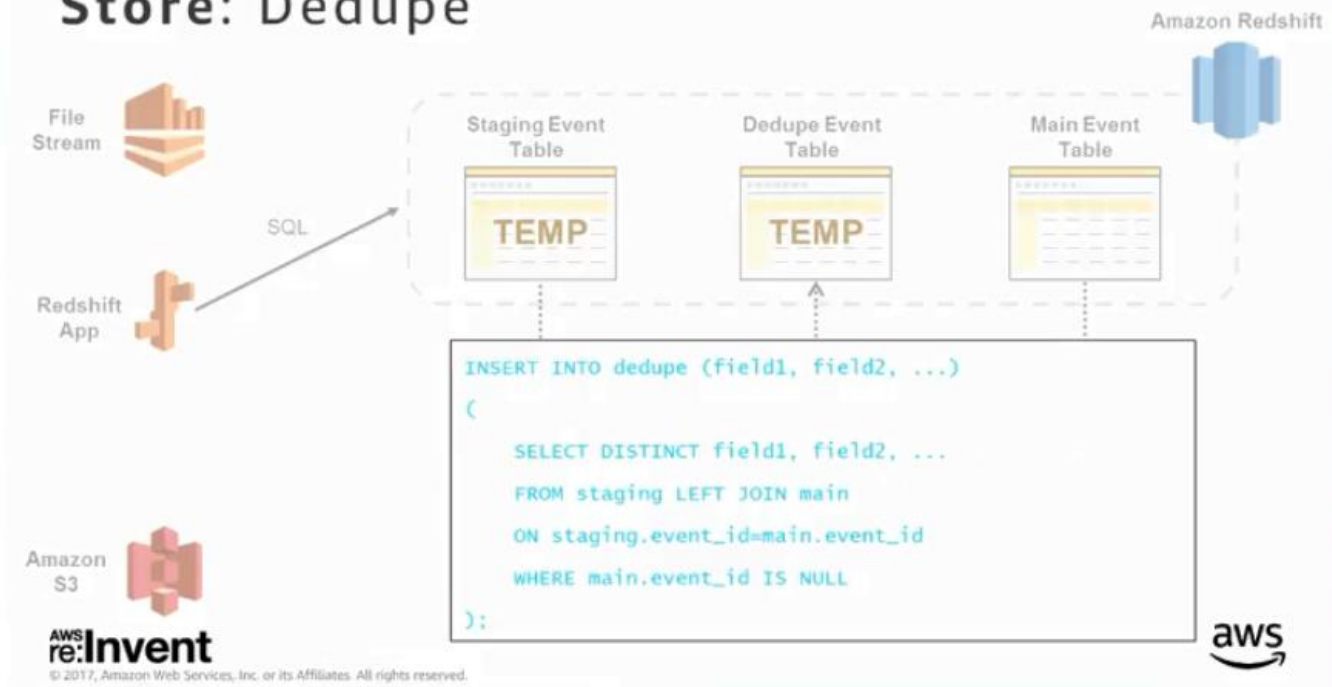
**AWS re:Invent**

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

aws

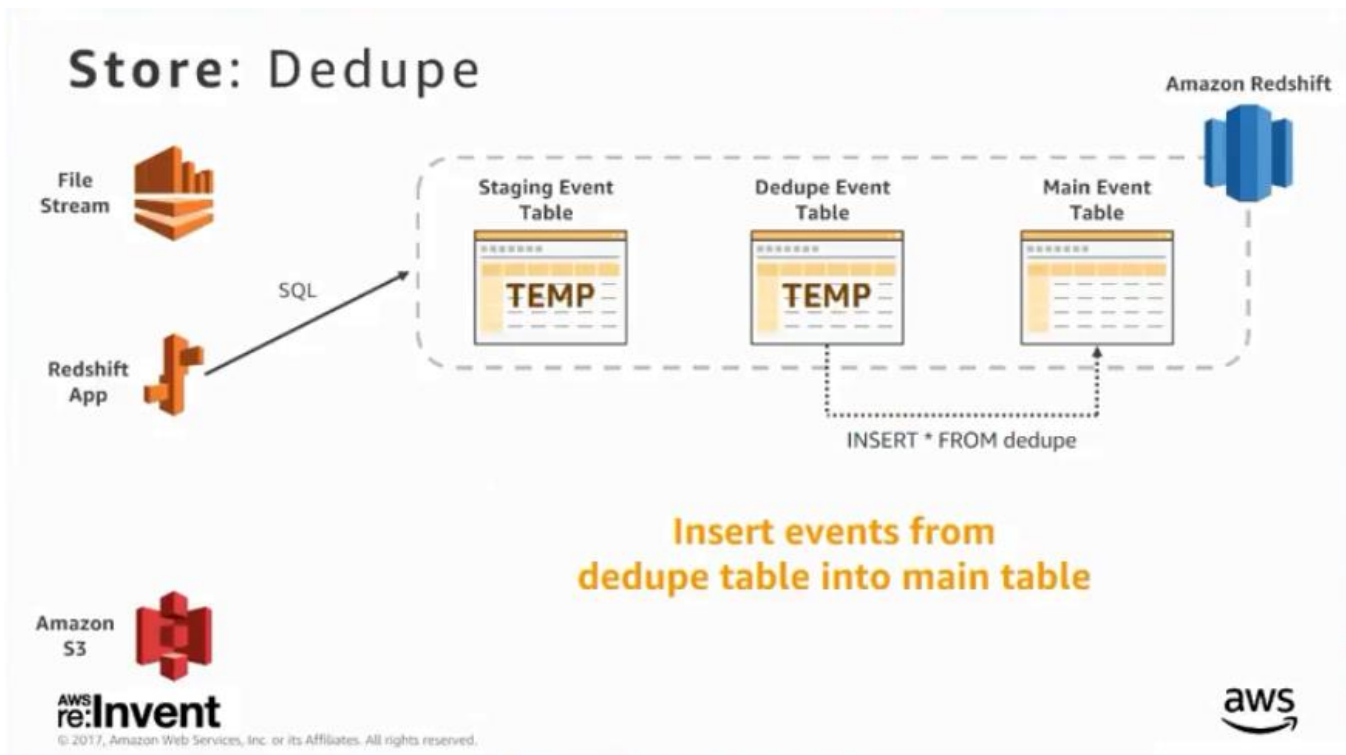
We the create a secondary table called the dedupe table, we now do some dedupe magic by taking everything in the staging table and left join it with our existing set of events in the main table and put only the unique data into the secondary dedupe table

## Store: Dedupe




This looks like above

## Store: Dedupe



We get unique event data in the dedupe table and then do a simple merge back to the Main table

## Store: Dedupe



Drop temp tables

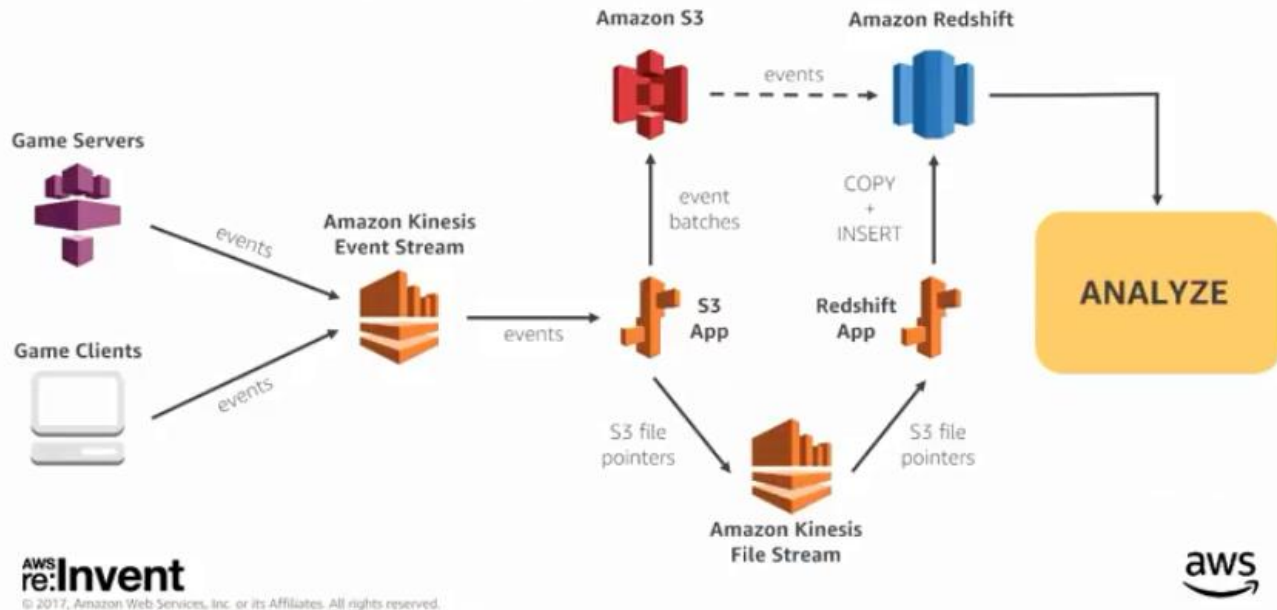


We then drop the temp tables when done

## Store: Warm data

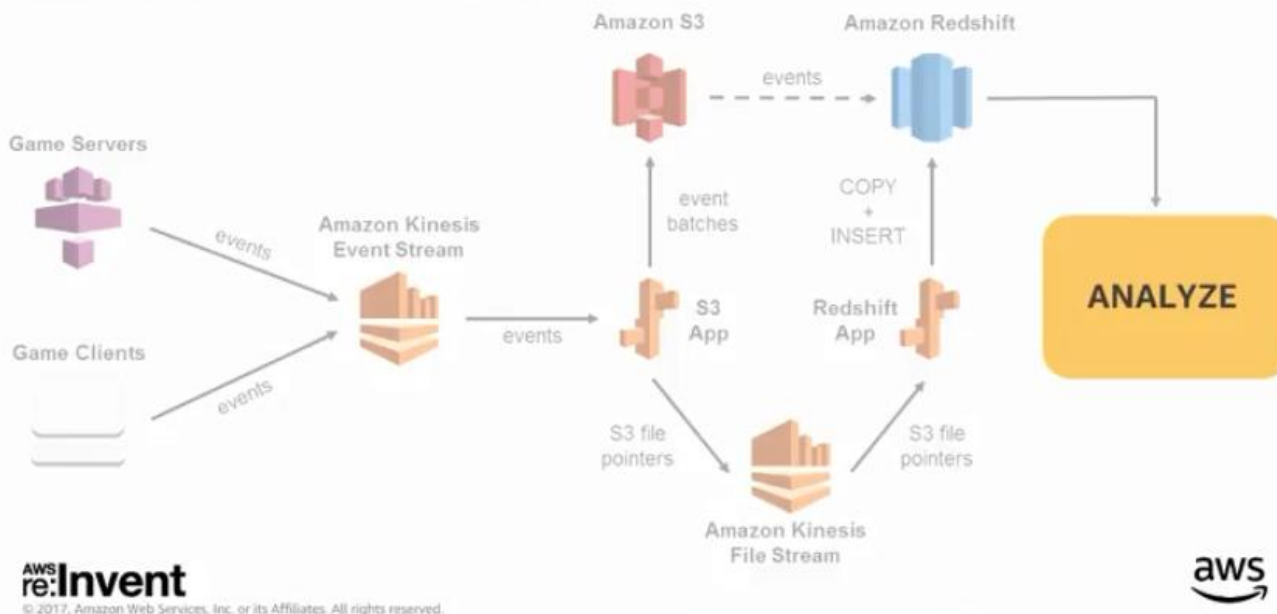


# Architecture: High-level



We are back to our architecture now, we have our cold data in S3 and our warm data in Redshift.

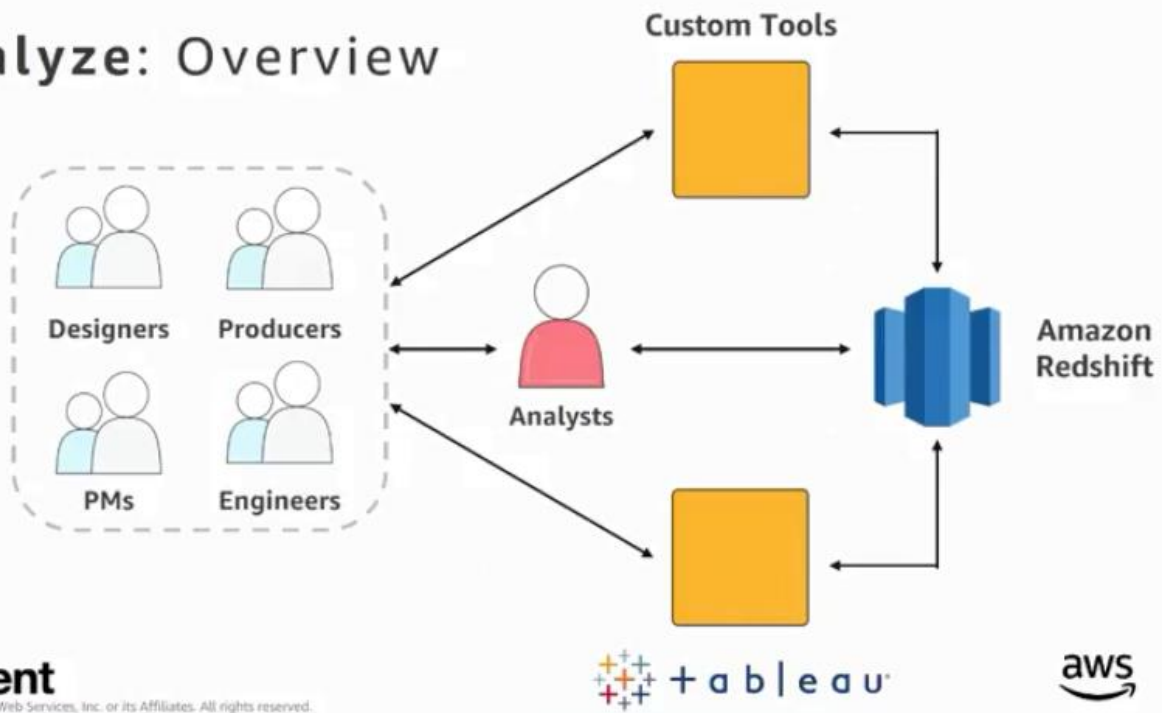
# Architecture: High-level



We are now going to analyze the data in Redshift




## Analyze: Overview



We have our custom tools that we built like our heatmap generator that can be run as a self-service, we have our warm data in S3, we have our Analysts that run queries against S3, and we use Tableau a lot for visualization from the market store

## Analyze: Heatmaps



We now want to generate the heatmap using a python script

# Analyze: Heatmaps with Python

```
import pg8000, pandas, numpy, matplotlib.pyplot as plt

# Connect to DB
conn = pg8000.connect(user='user', password='p@ssw0rd', host=host_name, port=5439, database='analytics')
cursor = conn.cursor()

# Run aggregation query
cursor.execute('''SELECT FLOOR(position_x), FLOOR(position_y), COUNT(*)
                  FROM game_events
                  WHERE level_id='map_name' AND event_type='player_death'
                  GROUP BY 1, 2''')
deaths = cursor.fetchall()

# Convert (x,y) world space to image pixel space

# Generate hexbin plot
df = pandas.DataFrame.from_records(deaths, columns=['pos_x', 'pos_y', 'deaths'])
df.plot.hexbin(x='pos_x', y='pos_y', c='deaths', reduce_c_function=numpy.max, gridsize=25, alpha=0.5, linewidth=0)

# Draw image and plot
plt.imshow(x=plt.imread('map_background_image.png'), zorder=0)
plt.show()
```

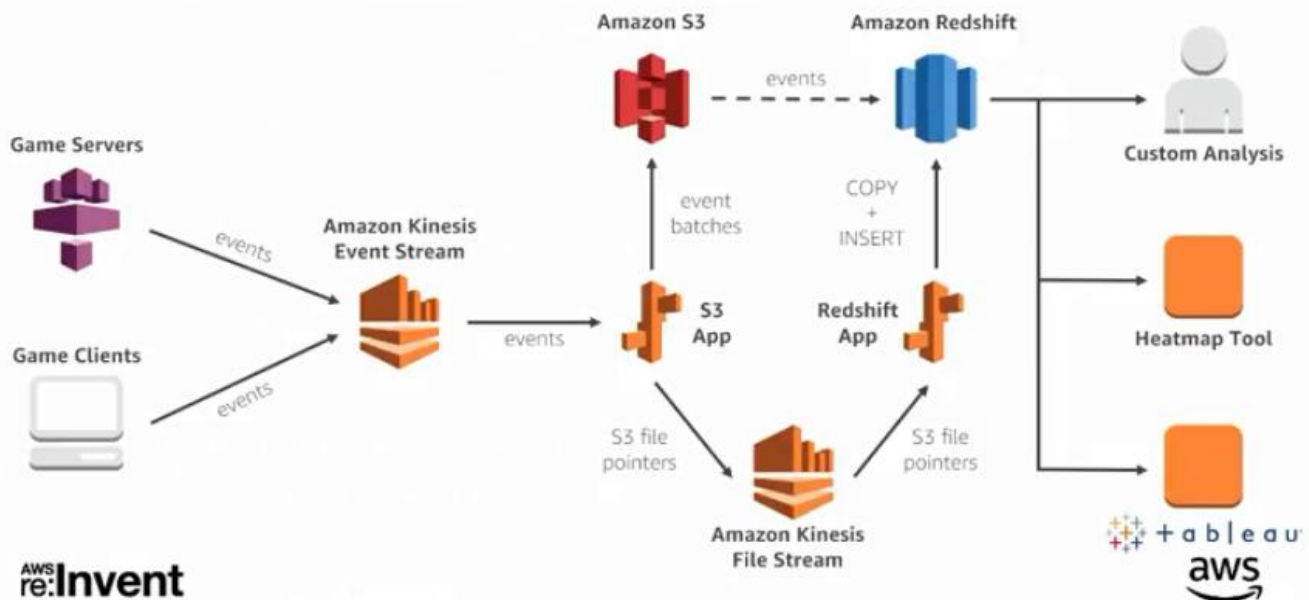
AWS  
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



We are using a library called pg8000 to connect to the Redshift database. Redshift can crunch data aggregations quickly. We also use the pandas dataframe data format and plot the data

## Architecture: High-level



AWS  
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

We have now finished our 4-stage architecture

## Architecture



Cold data



Warm data



Heatmaps



**WE'RE DONE!**

**AWS re:Invent**

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



## Evolve



**Amazon  
Elasticsearch  
Service**



**Hot data**

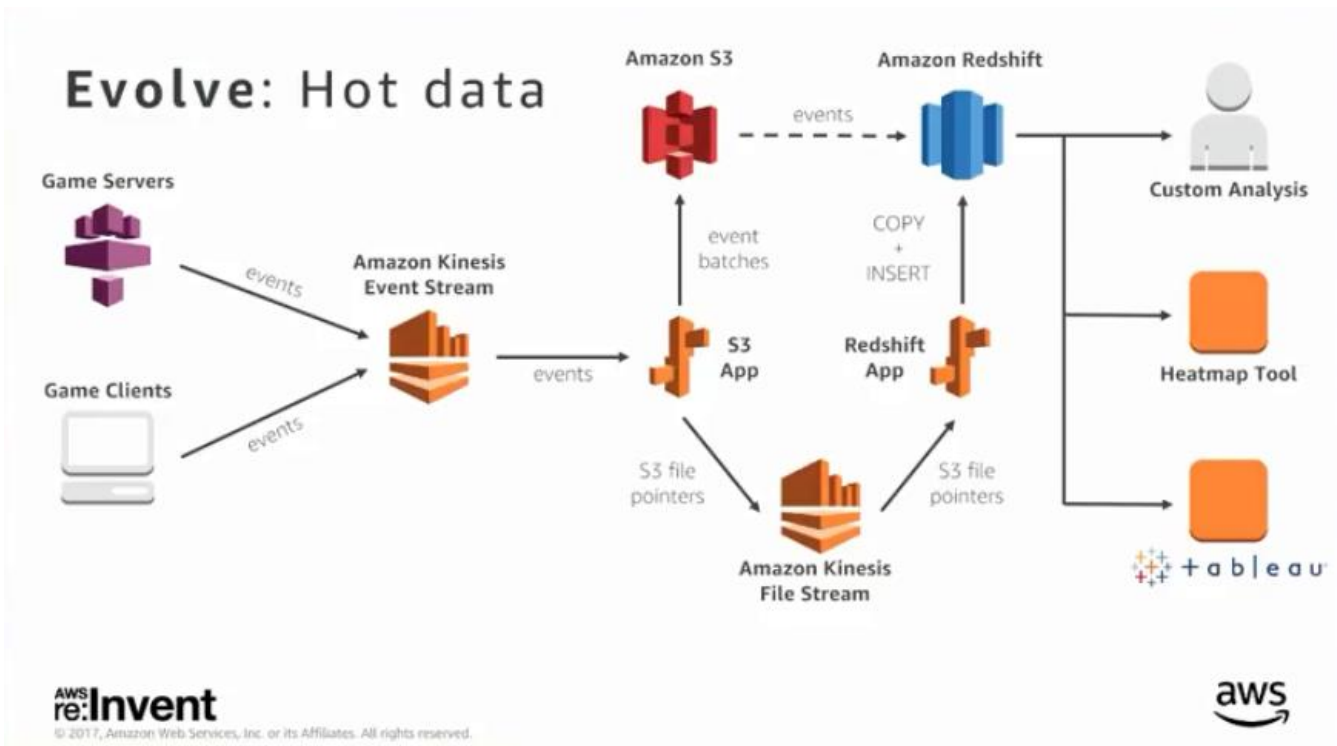
|            |            |
|------------|------------|
| Retention  | 7 days     |
| Access     | Very fast  |
| Turnaround | 5 minutes  |
| Structure  | Structured |
| Duplicates | No         |

**AWS re:Invent**

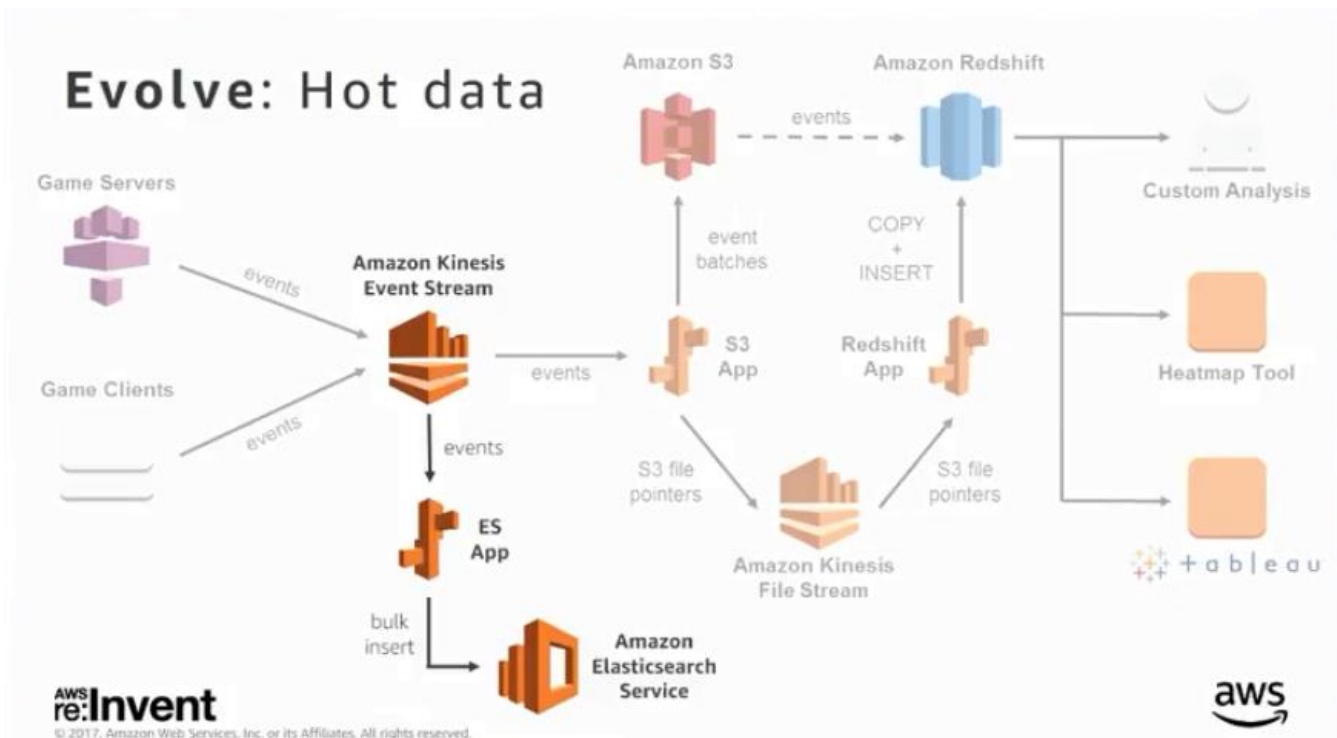
© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



Hot data is because 'most recent is most hot', we are keeping a week data only for hot data and we use Elasticsearch, it is very fast and does cool stuff with time series data, has plugins like kibana



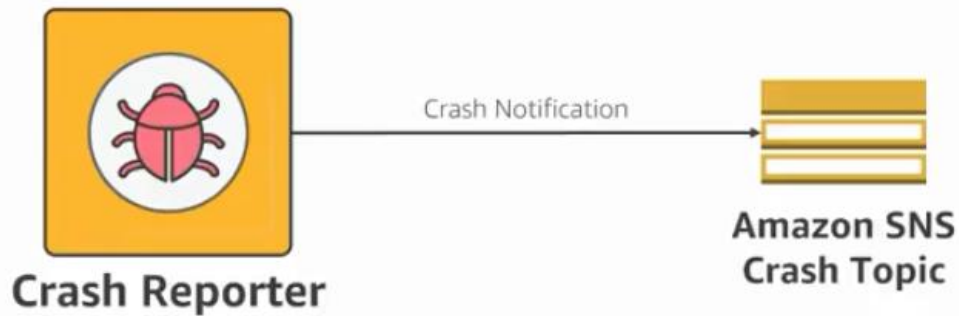
Where do we add Elasticsearch?



We hooked it up to the Kinesis stream using an Elasticsearch App (ES App) as above. We also use the ES App to drop some events that we don't want stored in Elasticsearch like combat events

# Evolve: Crashes

What about other data producers?

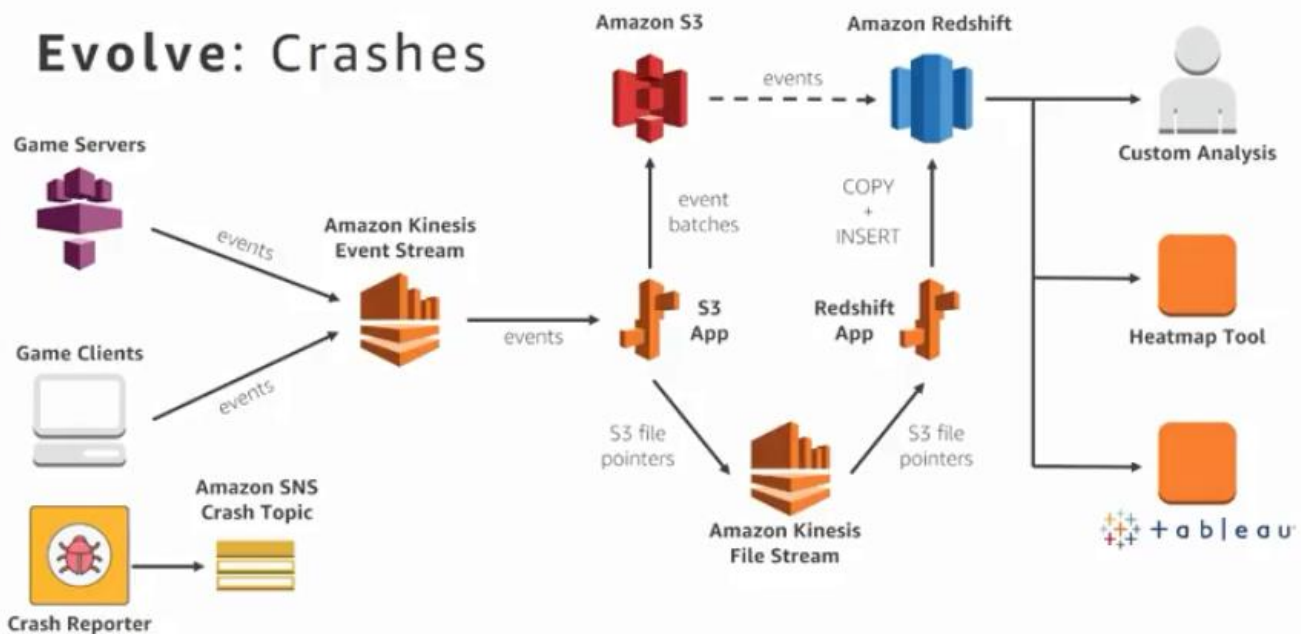


AWS  
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



# Evolve: Crashes

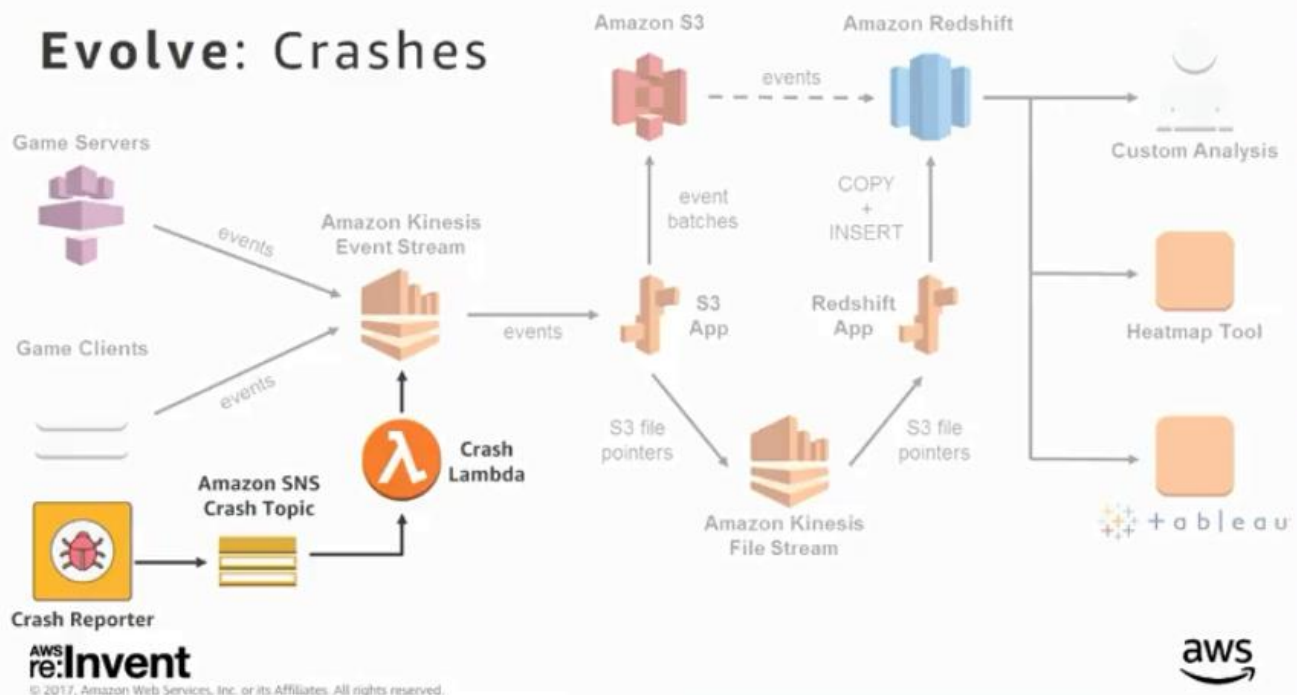


AWS  
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



## Evolve: Crashes



## Evolve: New tech

What about new AWS features and services?



Amazon  
Kinesis  
Analytics



Amazon  
Redshift  
Spectrum



Amazon  
QuickSight



Amazon  
Athena


AWS  
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

aws




## Evolve: New tech



We want to evaluate new services now available to refine our architecture

## Evolve: New tech



# Takeaways

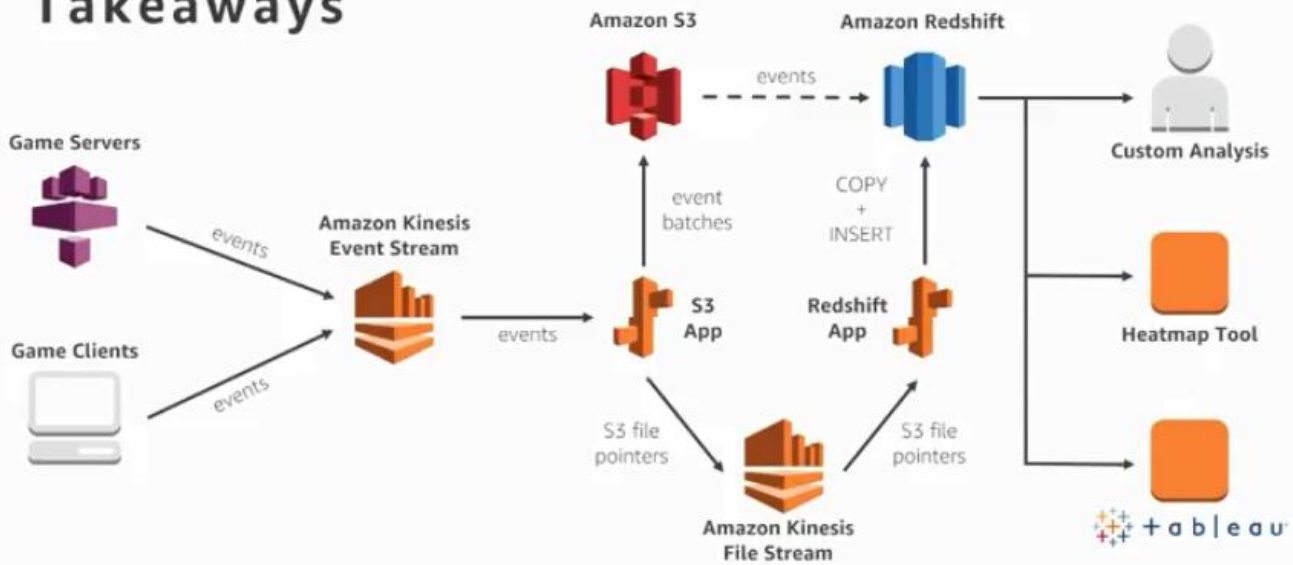


**AWS re:Invent**

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



# Takeaways

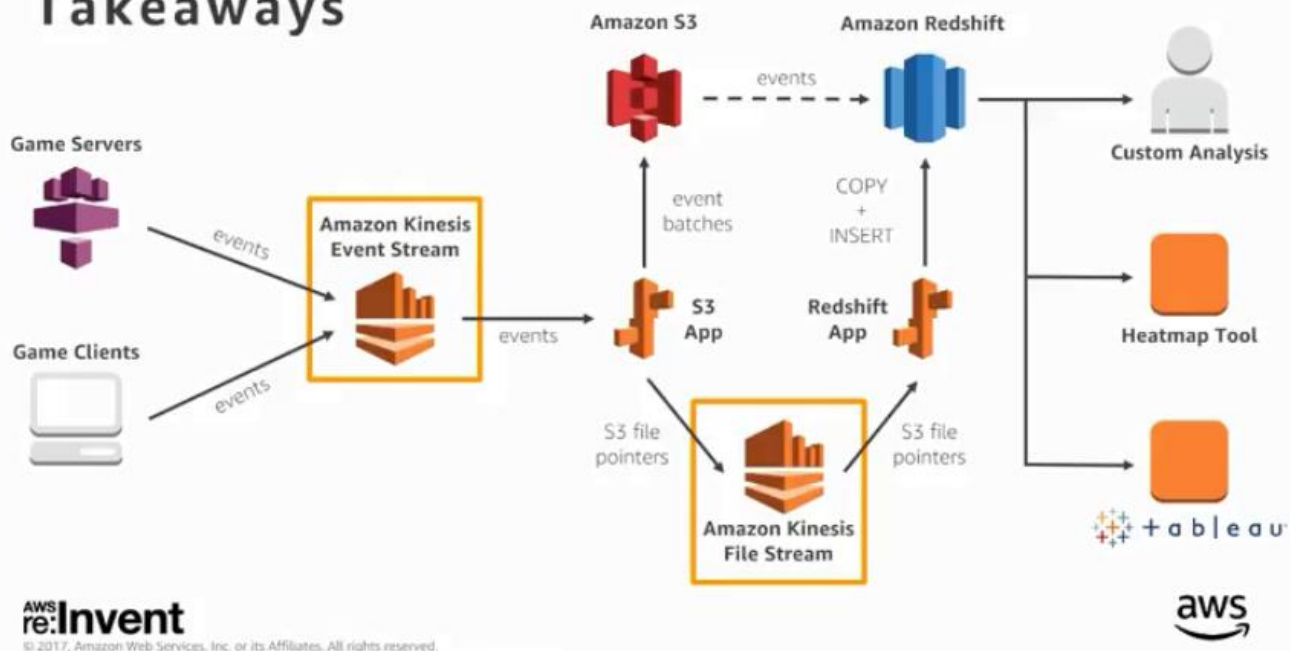


**AWS re:Invent**

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

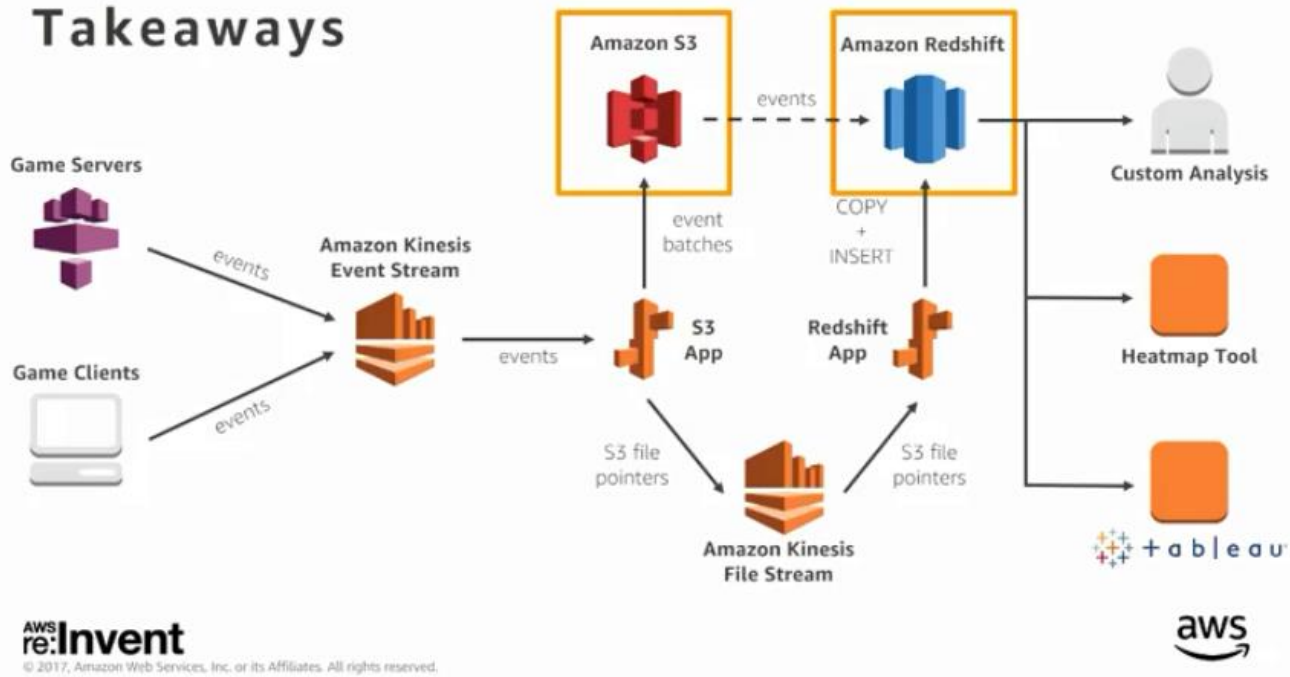


## Takeaways



Those Kinesis streams really don't need to be there from a pure functionality point of view, but they help down the line

## Takeaways



Putting your data into data stores with fan out capability is recommended because they can enable you do many other things in future

# Takeaways



- Assume things are going to **change**
- **Abstract** and decouple wherever possible
- Bias toward **fan-out** (no data dead ends)
- Run **experiments**, be agile, learn, and be curious

**AWS**  
**re:Invent**

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



# Takeaways



Amazon Game Studios

<https://games.amazon.com>

AWS for Gaming

<https://aws.amazon.com/gaming/>

Gaming Analytics Pipeline - Solution

<https://aws.amazon.com/answers/big-data/gaming-analytics-pipeline/>

Gaming Analytics Pipeline - Source Code

<https://github.com/aws-labs/gaming-analytics-pipeline>

**AWS**  
**re:Invent**

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



**AWS**  
**re:Invent**

Thank you!

**AWS**  
**re:Invent**

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

