

DEV328

DevOps Lessons from **Coursera**: Site Performance, Reliability, and Developer Productivity

Lewis Chung + Bryan Kane
Frontend Infrastructure @ Coursera

November 28, 2017

AWS
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



Intro
Coursera
overview

AWS CodeBuild
Optimizing
application builds

Amazon ECS
Isolating front-end
applications

Application
Load Balancer
Enhancing staging
environments

We envision a world where
anyone, anywhere
can transform their life
by accessing the world's
best learning experience.

AWS
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



AWS CodeBuild
Optimizing
application builds

Amazon ECS
Isolating front-end
applications


Application



30+ million
learners



150+
partners




2,200+
courses

AWS CodeBuild
Optimizing
application builds

Amazon ECS
Isolating front-end
applications

Application



70+
engineers

Intro
Coursera
overview

AWS CodeBuild
Optimizing
application builds

Amazon ECS
Isolating front-end
applications

Application
Load Balancer
Enhancing staging
environments

Keep Coursera as **fast** and
reliable as possible.

Empower Coursera's engineers to
be as **productive** as possible.

Adding build flexibility and reducing costs with
AWS CodeBuild

Intro
Coursera
overview

AWS CodeBuild
Optimizing
application builds

Amazon ECS
Isolating front-end
applications

Application
Load Balancer
Enhancing staging
environments



Productivity *increases* as time
between commit and deploy *decreases*

Intro
Coursera
overview

AWS CodeBuild
Optimizing
application builds

backend



Backend has always been developed as microservices

Intro
Coursera
overview

AWS CodeBuild
Optimizing
application builds

Amazon ECS
Isolating front-end
applications

Application
Load Balancer
Enhancing staging
environments



frontend



AWS re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



But the frontend was still one big JS monolith



Coursera in 2013

- commit to deploy in less than 5 minutes
- commit and deploy happen together
- regression debugging is easy w/ granular deploys

build time: **5 minutes**



build time: **15 minutes**



build time: 30 minutes

coursera

Intro
Coursera
overview

AWS CodeBuild
Optimizing
application builds

Amazon ECS
Isolating front-end
applications

**Application
Load Balancer**
Enhancing staging
environments

Amazon ECS
Isolating front-end
applications

**Application
Load Balancer**
Enhancing staging
environments

AWS CodeBuild
Optimizing
application builds

Amazon ECS
Isolating front-end
applications

Coursera in 2016

- builds take over 30 minutes
- developers forget to deploy
- debugging regressions becomes more difficult



Say we have 5 feature commits at once that are ready for testing and deployment





If 3rd commit has a problem



We can't deploy them all, we need to break them out and deploy one by one. So, we decided to break our frontend monolith up



Intro
Coursera
overview

AWS CodeBuild
Optimizing
application builds

Amazon ECS
Isolating front-end
applications

Application
Load Balancer
Enhancing staging
environments



Intro
Coursera
overview

AWS CodeBuild
Optimizing
application builds

Amazon ECS
Isolating front-end
applications

Application
Load Balancer
Enhancing staging
environments

Breaking up the monolith

- determine what applications were affected by a code change
- build individual applications in the codebase, not the whole

Intro
Coursera
overview

AWS CodeBuild
Optimizing
application builds

Amazon ECS
Isolating front-end
applications

Application
Load Balancer
Enhancing staging
environments





We now build mSes that are affected by a feature commit



Except for features that affect all mSes like the page header



We kept a pool of 8 workers on EC2 that monitors our build queue


Coursera overview

AWS CodeBuild
Optimizing application builds

Amazon ECS
Isolating front-end applications

Application Load Balancer

5 min



When we have 50 builds in the build queue

Coursera overview

AWS CodeBuild
Optimizing application builds

Amazon ECS
Isolating front-end applications

Application Load Balancer

5 min

5 min


5 min

5 min

5 min

5 min

5 min



Optimal full build:
35 minutes


Coursera overview

AWS CodeBuild
Optimizing application builds

Amazon ECS
Isolating front-end applications

Application Load Balancer

5 min



This is what we actually want, building everything in about 5 minutes

Optimizing application builds

Amazon ECS
Isolating front-end applications

N workers

overview

AWS CodeBuild
Optimizing application builds

Amazon ECS
Isolating front-end applications

Application Load Balancer

Automatically scaling Jenkins?

Too slow for
bursty build jobs.

Can we automatically scale our Jenkins worker pool quickly enough to do this?

overview

AWS CodeBuild
Optimizing application builds

Amazon ECS
Isolating front-end applications

Overprovision Jenkins?

Not cost effective

AWS CodeBuild
Optimizing application builds

Amazon ECS
Isolating front-end applications

Multiple commits
might land minutes from
each other...

Intro
Coursera
overview

AWS CodeBuild
Optimizing application builds

Amazon ECS
Isolating front-end applications

Application Load Balancer
Enhancing staging environments


commit 13a0bd
Author: Alice
Date: **Mon Oct 16 2:45:00 2017 -0700**

Update page header with new styles...



commit 13a0bd
Author: Bob
Date: **Mon Oct 16 2:47:00 2017 -0700**

Add links to footer...



commit 13a0bd
Author: Carol
Date: **Mon Oct 16 2:49:00 2017 -0700**

Update button component to support the new transition required for feature...



This scenario shows that we have over 100 builds queued up in the build system and 8 workers is not enough for 5 mins

Can we do better?



AWS CodeBuild



AWS CodeBuild

What is CodeBuild?


- on-demand build environment
- run your builds inside docker containers
- near-infinite elastic capacity -- scale up and down as you need

Compute instance type	Memory (GB)	vCPU	Price per build minute (\$)
small	3	2	0.005
medium	7	4	0.010
large	15	8	0.020

AWS CodeBuild
Optimizing application builds

Amazon ECS
Isolating front-end applications

Application




Coursera overview

AWS CodeBuild
Optimizing application builds

Amazon ECS
Isolating front-end applications

Application

5 min



Now we can go from nothing being built to everything being built in 5 min...we can!

Intro
Coursera overview

AWS CodeBuild
Optimizing application builds

Amazon ECS
Isolating front-end applications

Application Load Balancer
Enhancing staging environments

commit 13a0bd
Author: Alice
Date: **Mon Oct 16 2:45:00 2017 -0700**

Update page header with new styles...



commit 13a0bd
Author: Bob
Date: **Mon Oct 16 2:45:30 2017 -0700**

Add links to footer...




commit 13a0bd
Author: Carol
Date: **Mon Oct 16 2:46:00 2017 -0700**

Update button component to support the new transition required for feature...



- Intro
Coursera
overview
- AWS CodeBuild**
Optimizing
application builds
- Amazon ECS
Isolating front-end
applications
- Application
Load Balancer
Enhancing staging
environments

~5 min



We can now get this done too


- AWS CodeBuild**
Optimizing
application builds
- Amazon ECS
Isolating front-end
applications

Migrating from Jenkins to CodeBuild was easy

We were able to take our build scripts that we are running on Jenkins and make a few minor modifications to get it running on CodeBuild to make it compliant. Check the documentation

- Intro
Coursera
overview
- AWS CodeBuild**
Optimizing
application builds
- Amazon ECS
Isolating front-end
applications
- Application
Load Balancer
Enhancing staging
environments

Total time to build:



slowest individual build time

Tips and Tricks

Intro
Coursera

Intro
Coursera
overview

AWS CodeBuild
Optimizing
application builds

Amazon ECS
Isolating front-end
applications

**Application
Load Balancer**
Enhancing staging
environments

Build Visualization



Supercharge Logs

Intro
Coursera
overview

AWS CodeBuild
Optimizing
application builds

Amazon ECS
Isolating front-end
applications

**Application
Load Balancer**
Enhancing staging
environments

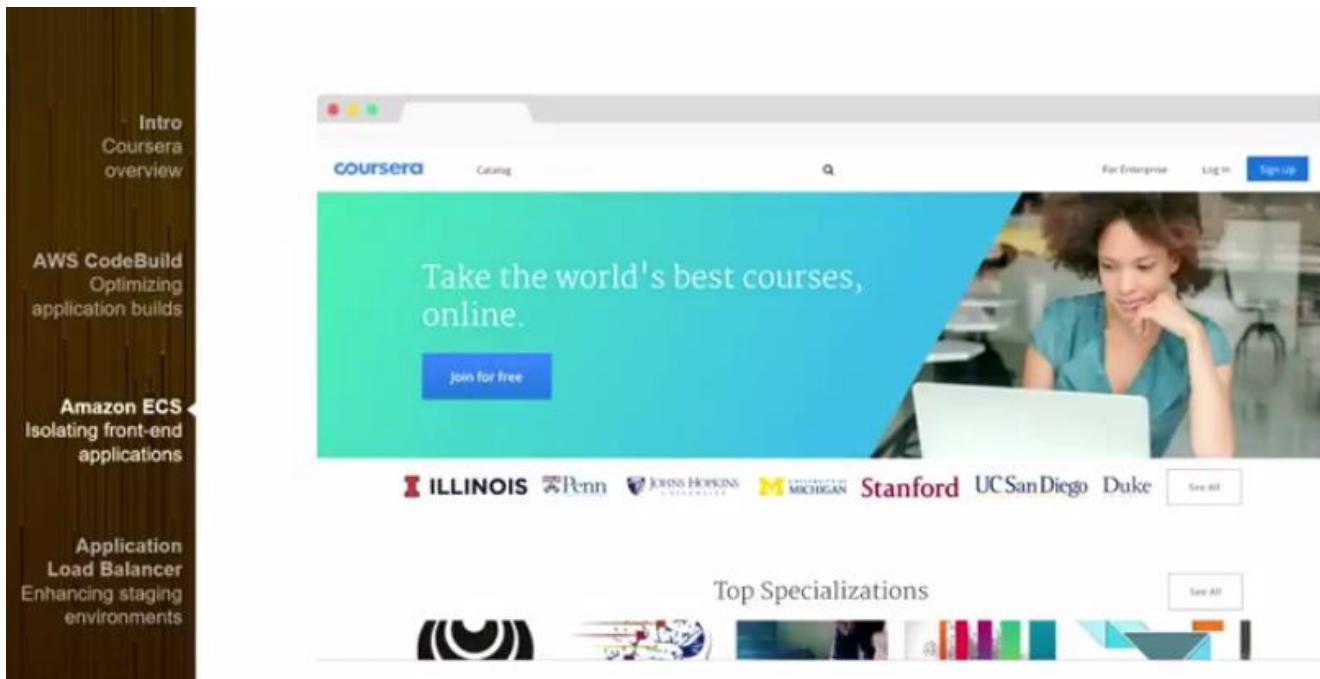
```
35. 09/27/2017 [38m Verifying integrity...[0m
36. 09/27/2017 gpg: keyring '/root/.gnupg/secring.gpg' created
37. 09/27/2017 gpg: key 86F5B310: public key "Yarn Packaging <yarn@dan.cx>" imported
38. 09/27/2017 gpg: Total number processed: 1
39. 09/27/2017 gpg: Imported: 1 (RSA: 1)
40. 09/27/2017 gpg: Signature made Mon Feb 27 15:37:57 2017 UTC using RSA key ID F02497F5
41. 09/27/2017 gpg: Good signature from "Yarn Packaging <yarn@dan.cx>"
42. 09/27/2017 gpg: WARNING: This key is not certified with a trusted signature!
43. 09/27/2017 gpg: There is no indication that the signature belongs to the
44. 09/27/2017 Primary key fingerprint: 72EC F46A 5604 AD39 C407 8087 1646 8018 80E5 D310
45. 09/27/2017 Subkey fingerprint: 6AB1 0C51 40B6 6599 AA17 F081 46C2 1380 F024 97F5
```

awslabs / aws-codebuild-jenkins-plugin


AWS CodeBuild
Optimizing
application builds

Amazon ECS
Isolating front-end

Fixing home-grown multi-tenancy with EC2 Container Service



Coursera's website is an SPA architecture built with JS and React



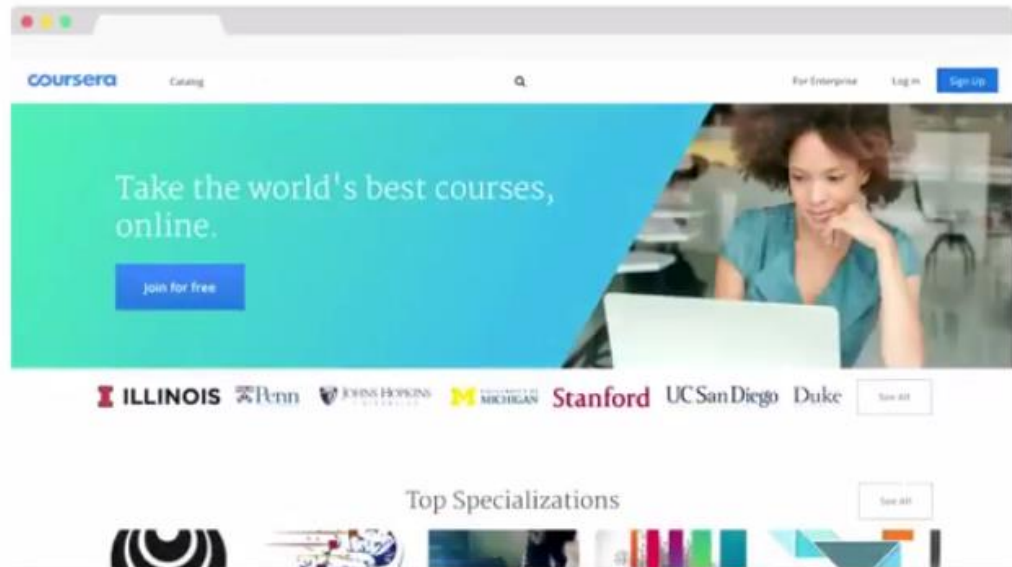
First you have JS downloaded, then it needs to make API calls for the content, then you get the full render

Intro
Coursera
overview

AWS CodeBuild
Optimizing
application builds

Amazon ECS
Isolating front-end
applications

Application
Load Balancer
Enhancing staging
environments



Intro
Coursera
overview

AWS CodeBuild
Optimizing
application builds

Amazon ECS
Isolating front-end
applications

Application
Load Balancer
Enhancing staging
environments

<pre> <html> <head> <script src="/builds/application.js" /> </head> <body> <h1>Loading...</h1> </body> </html> </pre>	<pre> <html> <head> <script src="/builds/application.js" /> </head> <body> <header>Sign up Log in</header> <main> <h1>Take the world's best courses</h1> ... </main> </body> </html> </pre>
---	---

Using server-side rendering in React allows us to execute NodeJS on the server and make all the needed API calls on the server but then send fully rendered HTML content to the client. The new server-side approach is on the right side.

- Intro
Coursera
overview
- AWS CodeBuild
Optimizing
application builds
- Amazon ECS**
Isolating front-end
applications
- Application
Load Balancer
Enhancing staging
environments




We now need to build a service to get this fully rendered pages

- Intro
Coursera
overview
- AWS CodeBuild
Optimizing
application builds
- Amazon ECS**
Isolating front-end
applications
- Application
Load Balancer
Enhancing staging
environments




The Render service was built as a Scala service, which can't really execute JS.

- Intro
Coursera
overview
- AWS CodeBuild
Optimizing
application builds
- Amazon ECS
Isolating front-end
applications
- Application
Load Balancer
Enhancing staging
environments



But NodeJS is good at executing JS, so we added some NodeJS to the Scala service. It is a NodeJS co-process to run alongside the Scala service with the NodeJS executing the JS and the Scala part doing everything else.

- Intro
Coursera
overview
- AWS CodeBuild
Optimizing
application builds
- Amazon ECS
Isolating front-end
applications
- Application
Load Balancer
Enhancing staging
environments




We do blue green deployment of the JS part before shifting traffic over to the new version

Intro
Coursera
overview

AWS CodeBuild
Optimizing
application builds

Amazon ECS
Isolating front-end
applications

Application
Load Balancer
Enhancing staging
environments




Intro
Coursera
overview

AWS CodeBuild
Optimizing
application builds


Amazon ECS
Isolating front-end
applications

Application
Load Balancer
Enhancing staging
environments



An application breaks occasionally

- Intro
Coursera
overview
- AWS CodeBuild
Optimizing
application builds
- Amazon ECS
Isolating front-end
applications
- Application
Load Balancer
Enhancing staging
environments




AWS
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



We want microservice isolation and not have faults bleed over to other microservice applications on the same EC2 host.

- Intro
Coursera
overview
- AWS CodeBuild
Optimizing
application builds
- Amazon ECS
Isolating front-end
applications
- Application
Load Balancer
Enhancing staging
environments



EC2 Container Service


We can now run multiple containers on a single host

Intro
Course
overview

AWS CodeBuild
Optimizing
application builds

Amazon ECS
Isolating front-end
applications

Application
Load Balancer
Enhancing staging
environments



Multi-tenancy *and* isolation


- Container management/orchestration system
- Runs Docker containers on EC2 hosts
- Easily handles:
 - Autoscaling
 - Dynamic resource reservations
 - Health checking
- Easy API interface for managing complex tasks



AWS CodeBuild
Optimizing
application builds

Amazon ECS
Isolating front-end
applications

Application
Load Balancer
Enhancing staging
environments



30% faster
*app response
times*



scales easily
*as traffic
increases*



66% cheaper
*from fewer
instances*

AWS CodeBuild
Optimizing
application builds

Amazon ECS
Isolating front-end
applications

Application
Load Balancer
Enhancing staging
environments



30% faster
*app response
times*



scales easily
*as traffic
increases*



66% cheaper
*from fewer
instances*

AWS CodeBuild
Optimizing
application builds

Amazon ECS
Isolating front-end
applications

Application
Load Balancer
Enhancing staging
environments



Oprah Winfrey
@Oprah

Follow

How to teach character! Fascinated by online
course organized by my friend

[@norman_atkins](#) [@RelayGSE](#).
coursera.org/course/teachin...

30% f.
*app response
times*

*as traffic
increases*

cheaper
*from fewer
instances*

AWS CodeBuild
Optimizing
application builds

Amazon ECS
Isolating front-end
applications

Application
Load Balancer
Enhancing staging
environments



30% faster
*app response
times*




scales easily
*as traffic
increases*



66% cheaper
*from fewer
instances*

Application isolation prevents cascading failures



one developer

migrated everything in

two months

Scale your containers *and* hosts

Choose one metric to scale your containers on:

pick either **CPU** or **memory**

Scale your containers *and* hosts

Combine the metrics for your ECS hosts:

scale based on *potential* number of
new containers that could be launched

[http://garbe.io/blog/2017/04/12/
a-better-solution-to-ecs-autoscaling/](http://garbe.io/blog/2017/04/12/a-better-solution-to-ecs-autoscaling/)



Automate your deployments

- Look into AWS CodeDeploy and AWS CloudFormation
- Use *services* for blue/green deployments

[https://blog.codeship.com/easy-blue-green-
deployments-on-amazon-ec2-container-service/](https://blog.codeship.com/easy-blue-green-deployments-on-amazon-ec2-container-service/)



Intro
Coursera
overview

AWS CodeBuild
Optimizing
application builds

Amazon ECS
Isolating front-end
applications

Application
Load Balancer
Enhancing staging
environments


Intro
Coursera
overview

AWS CodeBuild
Optimizing
application builds

Amazon ECS
Isolating front-end
applications

Application
Load Balancer
Enhancing staging
environments

Intro	Coursera overview
AWS CodeBuild	Optimizing application builds
Amazon ECS	Isolating front-end applications
Application Load Balancer	Enhancing staging environments





This is what we had earlier

Intro	Coursera overview
AWS CodeBuild	Optimizing application builds
Amazon ECS	Isolating front-end applications
Application Load Balancer	Enhancing staging environments



Now we have all the utilities like SumoLogic in a separate container that logs get forwarded to from our actual apps

- Intro
Coursera
overview
- AWS CodeBuild
Optimizing
application builds
- Amazon ECS
Isolating front-end
applications
- Application
Load Balancer
Enhancing staging



You might want to choose the correct EC2 instance types for your apps based on load tests against different instance types




We need to route requests to different URLs running on coursera.org. this requires an ALB from AWS, it is part of the ELB ecosystem

- Intro
Coursera
overview
- AWS CodeBuild
Optimizing
application builds
- Amazon ECS
Isolating front-end
applications
- Application
Load Balancer

What is Application Load Balancer?

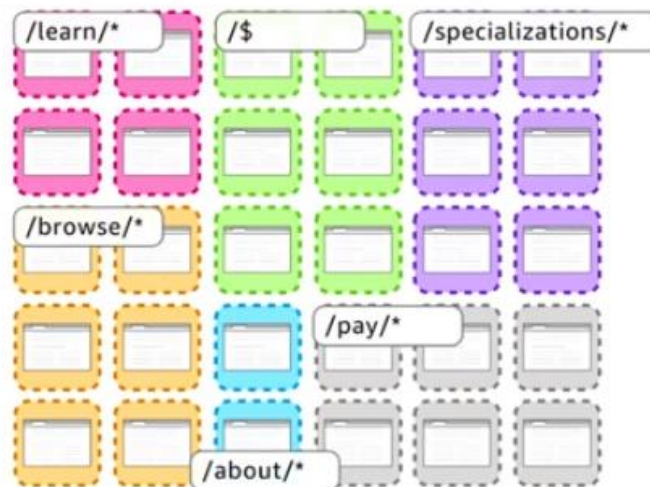
- distributes traffic across multiple targets in **target groups**
- support for container based targets in ECS
- content-based routing via **listener rules**
 - path-based routing
 - host-based routing
- set up **listener rules** and register targets via SDK/CLI

We have to set up listener rules to help check and direct traffic



Since we have many different apps running in ECS, how do we route requests to correct app?

Path-based Routing



Path-based routing is too limited

- authentication can affect on routing
- experiments can affect routing

Host-based routing *can be* flexible

- An edge tier allows you to generate a Host header

Edge Tier

Given a request, figure out what application needs to serve the request.

Edge Tier

Takes advantage of Host Based Content Routing:

- Given a request, rewrite the Host Header "`${appname}-${version}.coursera`"
- Forward the request to ALB

Edge




Application Load Balancer listener rules inspect the Host header to figure out to which target group it needs to forward the request

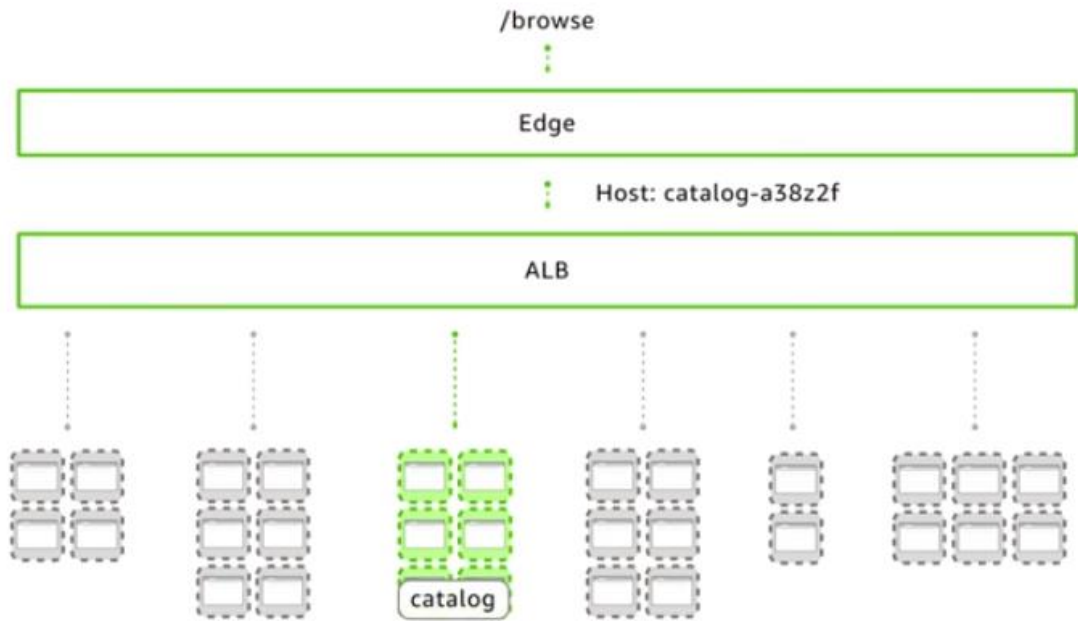
Edge



ALB



Intro Coursera overview
AWS CodeBuild Optimizing application builds
Amazon ECS Isolating front-end applications
Application Load Balancer Enhancing staging environments



When a request comes in for the catalog service, it hits the Edge tier and that generates a host header based on the intended path. Then the edge tier passes the request along to the ALB, the ALB can then read the host header data and forward the request to the appropriate app. This is how we implement routing for our production services.

Intro Coursera overview
AWS CodeBuild Optimizing application builds
Amazon ECS Isolating front-end applications
Application



We can also preview changes too

Intro
Coursera
overview

AWS CodeBuild
Optimizing
application builds

Amazon ECS
Isolating front-end
applications

Application
Load Balancer
Enhancing staging
environments

Production



Alice

Catalog@FeatA-1



Catalog@FeatA-2



Bob

Catalog@FeatB-1



Catalog@FeatB-2




If we change our build process to trigger on pull requests as well, we can then spin up containers running dev code also

Intro
Coursera
overview

AWS CodeBuild
Optimizing
application builds

Amazon ECS
Isolating front-end
applications

Application
Load Balancer
Enhancing staging
environments



We then have this type of scenario where we can have routing to the dev preview builds by setting up the correct host headers for the edge and the ALB in creating this temporary staging environments

Application Load Balancer Limitations

- Many hard limits exist, check documentation
 - Example: Max number of listener rules: 100
- Doesn't support traffic weighting
 - We implemented this in our edge tier

DevOps Lessons from Coursera:
Site Performance, Reliability, and Developer Productivity

Recap

Use AWS CodeBuild to scale your
build environment on-demand

Pay only for what you use,
and take advantage of elasticity

Amazon ECS keeps your applications independent without the overhead of Amazon EC2

Reduce costs, and scale faster by containerizing your applications

AWS re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



ALBs give you the flexibility to route traffic in creative ways

Target specific hosts or containers with dynamic rules for on-demand environments

AWS re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.





AWS re:Invent

Thanks for coming!

AWS
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

