

AWS re:INVENT

Digital Transformation

Adrian Cockcroft

AWS VP Cloud Architecture Strategy
@adrianco

November 28, 2017

AWS re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



Many industries are going through a digital transformation as their existing business models are being disrupted and new competitors emerge. The key driver is a need for faster time-to-value as a direct relationship with customers provides analytics that drive personalization and rapid product development. There's a cultural aspect to the change, as well as new organizational patterns that go along with a migration to cloud native services. Application architectures are evolving from monoliths to microservices and serverless deployments, and they becoming more distributed, highly available, and resilient. The highly automated practices that have built up around DevOps are moving to the mainstream, and some new techniques are emerging around security red teams and chaos engineering.

Digital Transformation



Disruption and opportunities



Microservices evolution



Cloud native

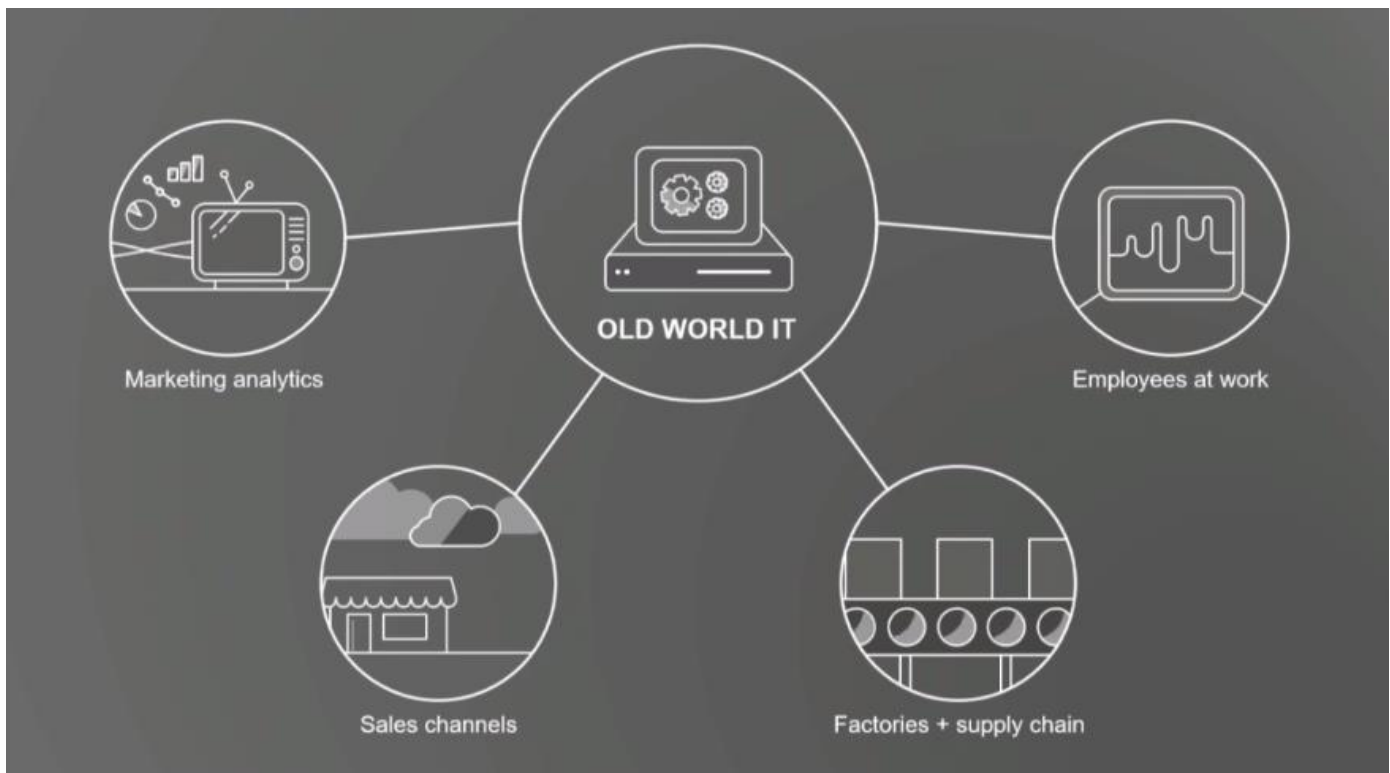


Dependency management



Chaos architecture

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



New Needs

Personalization

Customer tracking

New channels direct to customer

You need to build things that include the above, personalization because you can now know who is watching TV or seeing your ads or products in near real-time and you can now optimize their experiences.

Evolution of Business Logic



Monolith



Microservices



Functions

Splitting Monoliths Ten Years Ago




Monoliths usually pass around big fat XML and SOAP chunks between their sub-components and this made them slow

Splitting Monoliths Five Years Ago



Microservices Five Years Ago



This allows us to iterate and increase development time while working on the smaller services independently

Microservices to Functions

Standard building brick services provide standardized platform capabilities



We can now start replacing a lot of the common services with better managed services like SQS, DynamoDB, Kinesis, etc.

Microservices to Functions

Standard building brick services provide standardized platform capabilities



This are now our unique services that we now have to write.

Microservices to Functions



Microservices to Functions



Microservices to Functions



Microservices to Functions



Microservices to Ephemeral Functions



Microservices to Ephemeral Functions



These lambda functions are now ephemeral and can sleep while not being used

Microservices to Ephemeral Functions



When a request now comes in, it wakes up some lambda functions and services to do the particular work needed and then goes back to sleep

Microservices to Ephemeral Functions



Microservices to Ephemeral Functions



Microservices to Ephemeral Functions



This is the concept of the Serverless evolution

Evolution of Business Logic



The New De-Normal



We are trying to get to a de-normalized model

**Expensive,
Hard to Create
and Run**



**Database Schema
Entity Relationship**



They usually have a very complicated schema structure and becomes very hard to untangle the database

Kitchen Sink Analogy



Kitchen Sink Cleanup



You want to untangle the mess as below

Kitchen Sink Cleanup



Consistency Problem

How many complete sets are there?



But we are missing a spoon. It means we have to have a mechanism for keeping all our components in sync because microservices usually have their own individual databases and queues

Adding a New Use Case



We are going to have to create a new place to put a new dataset instead of lumping it together with other data

Cloud Makes It Easy to Add New Databases



Amazon
DMS



Amazon
DynamoDB




Amazon
RDS



You can now simply create your new database for a new service easily using a managed database service like DynamoDB

Untangle and Migrate Existing 'Kitchen Sink' Schemas



Untangle and Migrate Existing 'Kitchen Sink' Schemas



You now also have much simpler schema and systems, you don't need to worry about transactions and joins anymore because your data is now in a single database.

The New De-Normal



Cloud Native Architecture




Datacenter Native Architecture



DATACENTER

Datacenter Native Infrastructure Lives for years



DATACENTER

Cloud Migration

Pay as you go



DATACENTER



Applications and data

Pay up front and
depreciate over
3 years

Pay a month later
for the number of
seconds used



Cloud Native Principle

Pay for what you used last
month...

...not what you guess you will
need next year.

**File tickets and wait
for every step**




VS

**Self-service,
on-demand, no delays**



What you want is self-service, on-demand, API driven services

**File tickets and wait
for every step**



VS

**Self-service,
on-demand, no delays**



wait for every step

Deploy by filing a ticket and waiting weeks or months

on-demand, no delays

Deploy by making an API call self-service within minutes

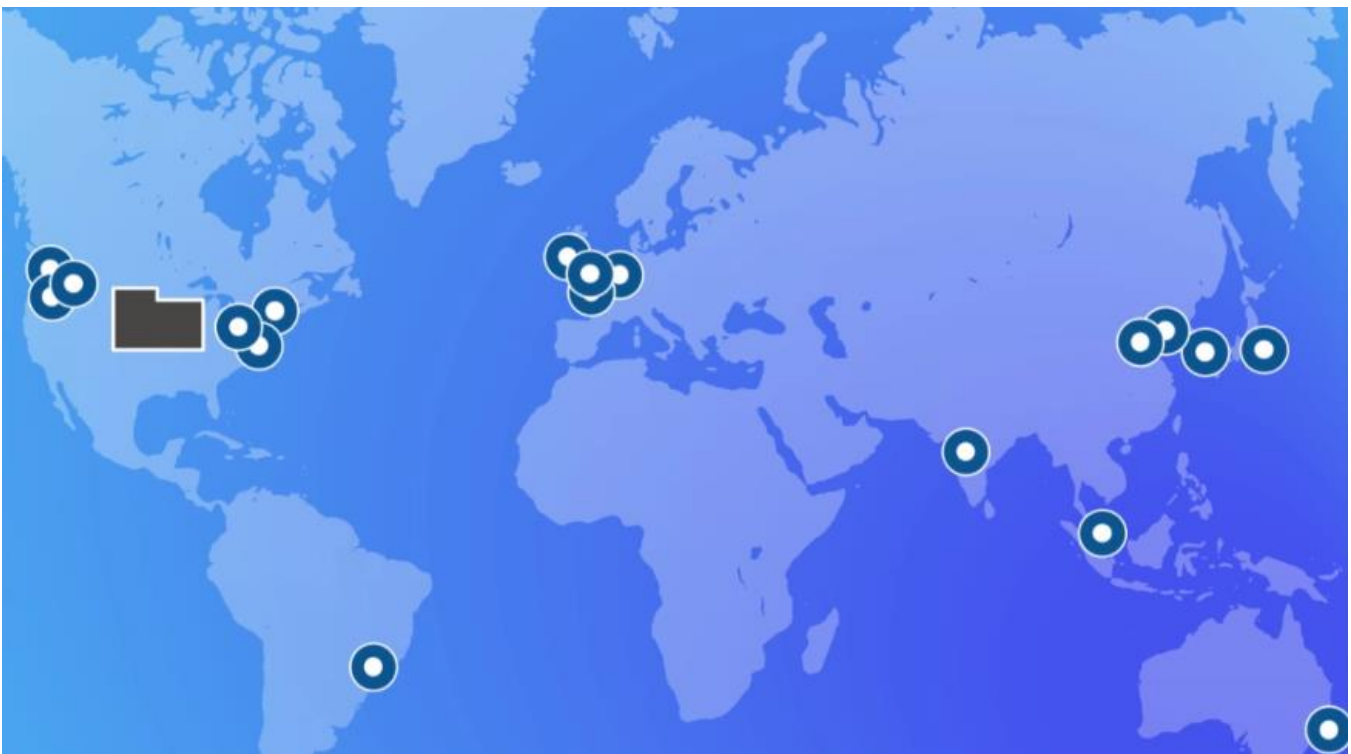
You need to make all your operational infrastructure be API call driven to be truly cloud-native



Cloud Native Principle

Self-service, API-driven, automated

Move from request tickets at every step to a tracking ticket that records what happened






Cloud Native Principle

Instant globally distributed deployments and data by default

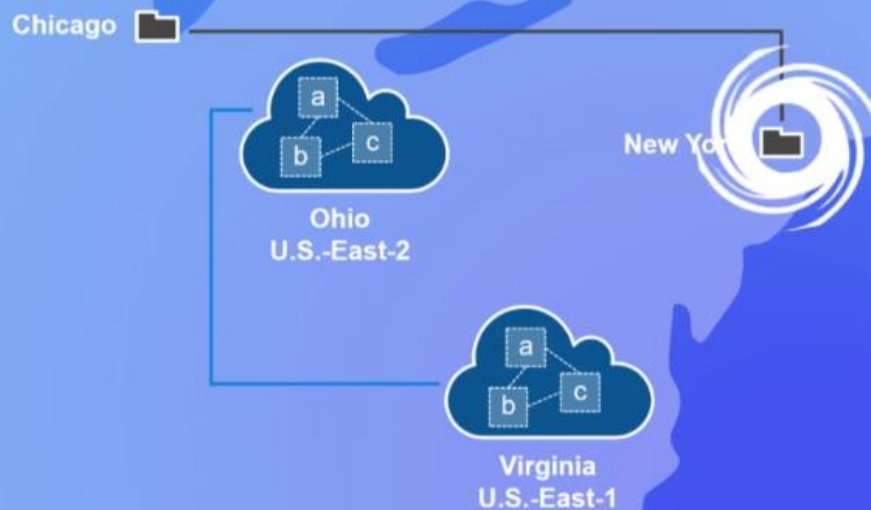
Regions and Zones

Typical cloud architecture
Zones a, b, and c are 10–100km apart




Regions and Zones

Hurricane Sandy



Regions and Zones

Hurricane Sandy



Regions and Zones

Datacenter native migration to cloud

Keep the same configuration and run MySQL on a cloud instance yourself



What you really want to do to have a distributed database is to use the Aurora service as below

Regions and Zones

Cloud native data migration

Amazon Aurora
Distribute over all three zones



Regions and Zones

Cloud native data migration

MySQL
Secondary

MySQL
Primary

Regions and Zones

Cloud native data migration

More resilient within each region

MySQL
Secondary

MySQL
Primary



Cloud Native Principle

Distribute over zones within a region by default

Elasticity



DATACENTER

Hard to get over 10 percent utilization—need extra capacity in case of peak



CLOUD

Target over 40 percent utilization—no capacity overload issues



Automatic scaling for predictable heavy workloads



Serverless for spiky workloads with idle periods



Cloud Native Principle

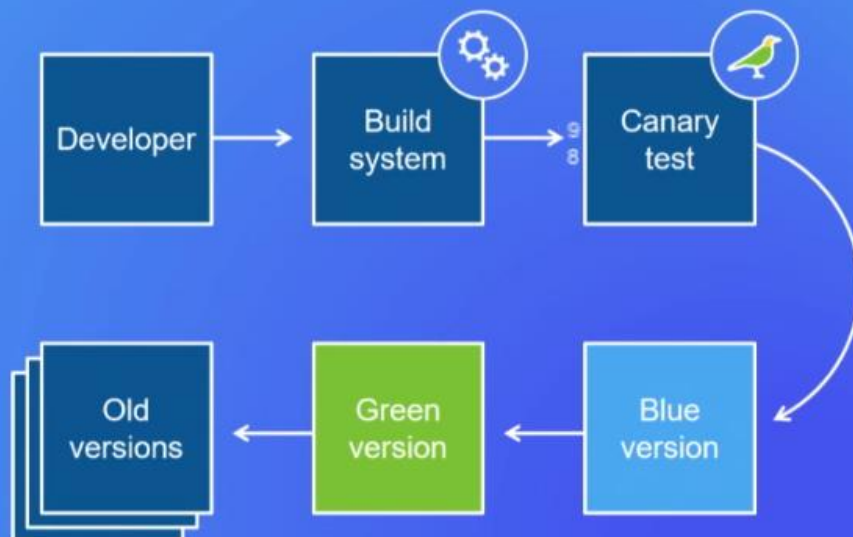
Turn it off when it's idle

Many times higher utilization

Huge cost savings

Avoids capacity overloads

Versioned Delivery Pipeline





Cloud Native Principle

Immutable code

Automated builds

Ephemeral instances, containers, and functions

Blue–Green deployments

Versioned services



Cloud Native Principles

Pay as you go, afterward

Self-service—no waiting

Globally distributed by default

Cross-zone/region availability models

High utilization—turn idle resources off

Immutable code deployments



Cloud Native Principles

Remain constant as practices evolve

Cloud Native Architecture



Principles and practices

Building on Cloud



Opportunities



Dependency
management



Chaos
architecture


Startup Spaces

Incremental

Gap fillers

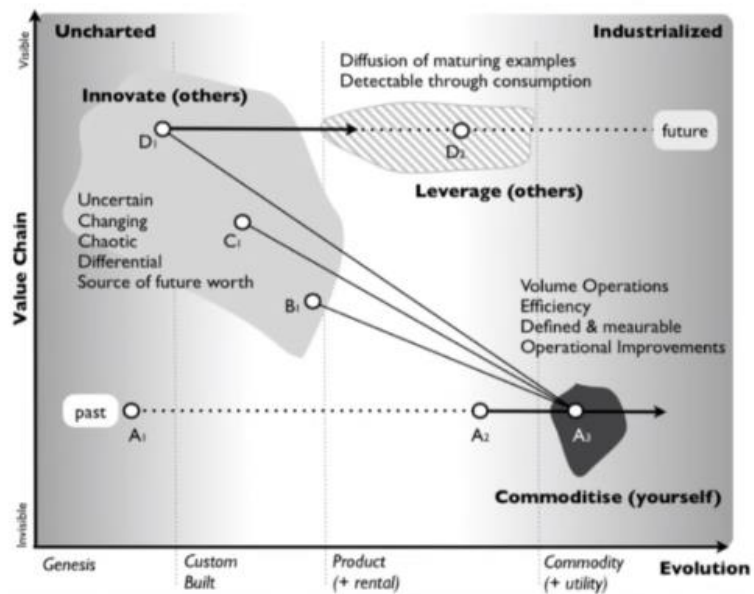
Leverage

Category creator



Leverage Wardley Maps

Move up the value chain



With Leverage Comes Dependencies...

Lock-in and the Lifecycle of Dependencies



Choosing, using, and losing

What is the return on investment (ROI) for each phase?



Choosing



Using



Losing

What is the ROI
for each phase?

How has ROI changed
with advances in
technology and
practices?



Choosing



Using



Losing



Choosing



Choosing



Investments

Negotiating, learning, experimenting

Hiring experts, building

Installing, customizing

Developing, training



Choosing

How much
time elapses?



“The best decision is the right
decision. The next-best decision
is the wrong decision. The worst
decision is no decision.”

— Scott McNealy



Choosing

Analysis paralysis

vs.

Snap judgement





Choosing



Making a commitment

Whenever development is frozen and the operations team takes over, the key is turned in the lock



Choosing—What Changed?

Old World

Monolith—all in one
Proof of concept install
Enterprise purchase cycle
Months
\$100K–Millions

New World

Microservice—fine grain
Web service/open source
Free tier/free trial
Minutes
\$0–\$1,000s



Using



Using



Investments

- Cost of setup
- Cost of operation
- Capacity planning
- Scenario planning
- Incident management
- Tuning performance and utilization



Using



Returns

- Service capabilities
- Availability, functionality
- Scalability, agility
- Efficiency



Using—What Changed?

Old World

Frozen installation
Ops specialist silo
Capacity upgrade costs
Low utilization
High cost of change

New World

Continuous delivery
Dev automation
Elastic cloud resources
High utilization
Low cost of change

The Using phase is about moving from a Project to a Product



Losing



Losing



Investments

Negotiating time
Contract penalties
Replacement costs
Decommissioning effort
Archiving, sustaining legacy



Losing



Returns

Reduced spending

More advanced technology

Better service, agility, scalability

Choose again, the cycle continues...



Losing—What Changed?

Old World

Monolithic—all or nothing

Frozen waterfall projects

Long-term contracts

Local dependencies

New World

Microservices—fine grain

Agile continuous delivery

Pay as you go

Remote web services

Old World

Monolithic on-prem waterfall lock-in

Years

Millions of dollars

Hundreds of dev years

Lock-in

Lawyers and contracts

New World

Agile cloud-native micro-dependencies

Weeks

Hundreds of dollars

A few dev weeks

Refactoring

Self-service

Bottom Line

ROI for choosing, using, losing has changed radically. Stop talking about lock-in, it's just refactoring dependencies

The cost of each dependency is far lower

Frequency of refactoring is far higher

Investment and return are much more incremental



Chaos Architecture




A Cloud Native Availability Model

Adrian Cockcroft

Chaos Architecture

Four layers
Two teams
An attitude



Infrastructure and Services

No single point of failure



No single
point of failure

No single
point of failure

Distributed
Automated

Replicated
Cloud



Switching and Interconnecting

Data replication
Traffic routing
Avoiding issues
Anti-entropy recovery



You now need to build an architecture that defines how does the data get to 2 places and how does the data stay in sync? You need a strategy for failovers and have a way of re-routing customers to the backup source when the main source fails and goes down reliably, how do you route back when the main source comes back on stream?



Who has a backup datacenter? What's the best description of it?

1. Availability theater—never tried to use it
2. Infrequent partial testing
3. Regular tests during maintenance
4. Frequent failovers during production to prove that no-one can tell it's happening



**Route updates and customer requests
to specific regions and services**



**Replicate data and re-route requests
during incidents**




**Switching mechanism must be far
more reliable than redundant
elements you are switching between**



Application Failures

Error returns
Slow response
Network partition




What happens when your app faces error returns from the APIs? Or slow responses? What do you want the reduced state or degraded state look like to your end user? Are you going to let them continue to do what they want to do if your service is down?



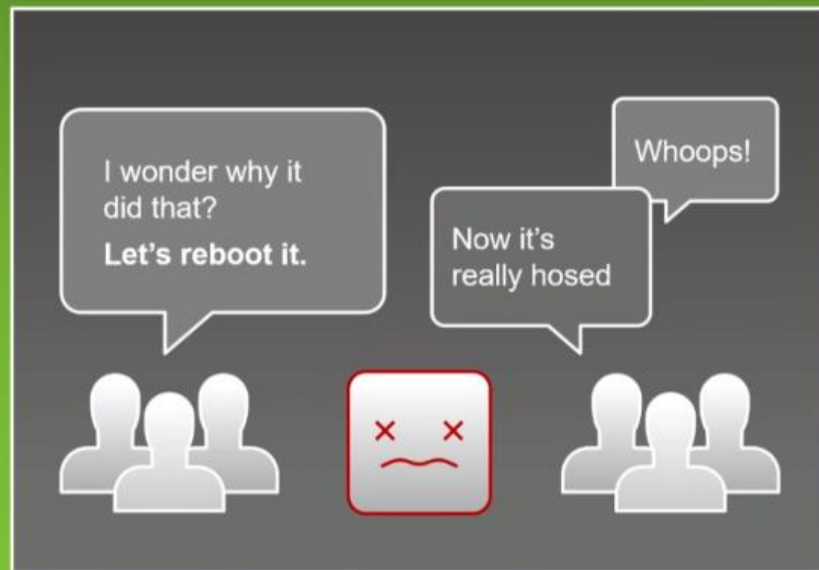
Microservices limit 'blast radius' for software incidents

Circuit breakers limit damage
Bulkheads prevent it from spreading
DITTO—Do Idempotent Things To Others
Avoid update and delete semantics




People

Unexpected application behavior often causes people to intervene and make the situation worse



People Training

A fire drill is a boring routine where we make everyone take the stairs and assemble in the parking lot



People Training

Fire drills save lives in the event of a real fire, because people are trained for how to react




Who runs the 'fire drill' for IT?

People

Application

Switching

Infrastructure



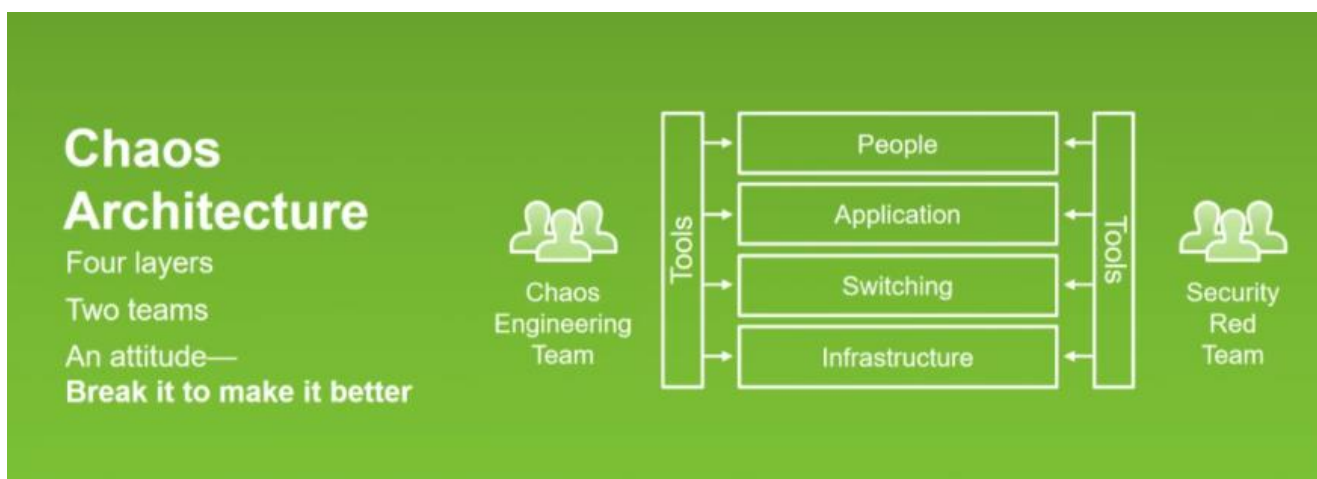
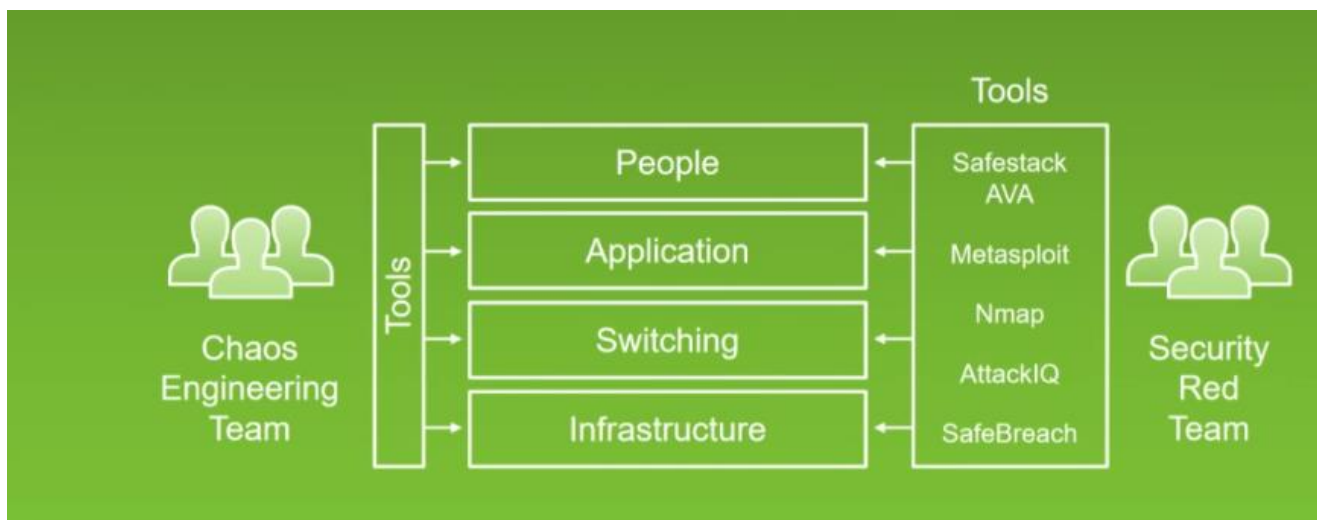
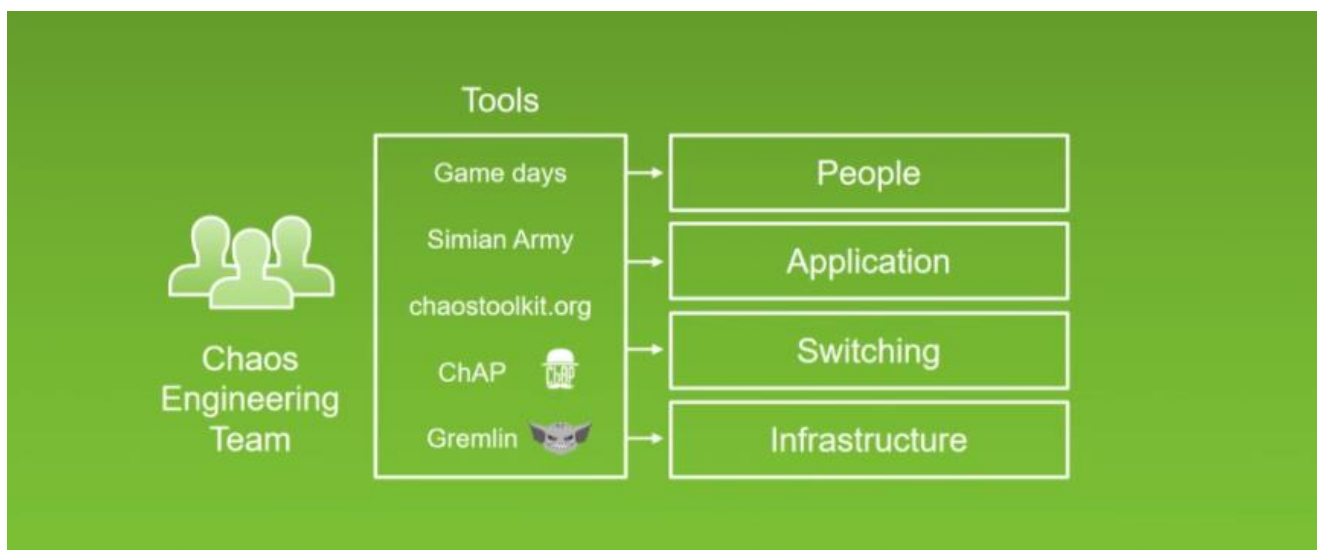
Chaos
Engineering
Team

People

Application

Switching

Infrastructure



Risk Tolerance

Who is at risk for what?

Is downtime a bigger risk?

Consistency
Secure
Downtime

← You Choose →

Availability
Higher risk
Permissive



Incident Lifecycle

Mitigate

Restore

Adapt

Antifragile



Break It to Make It Safer

The “new view” of safety based on:
Todd Conklin’s pre-accident podcast
John Allspaw’s stella.report
Sydney Dekker—Drift into Failure

Failures are a System Problem— Lack of Safety Margin

Not something with a root cause
of component or human error



**Blindfolded on a cliff edge,
what would you do?**



Hypothesis Testing

- We think we have safety margin in this dimension, let's carefully test to be sure
- In production
- Without causing an issue



How to Select a Test?

LDFI—Lineage-driven Fault Injection
Dependency graphs for important flows

Experienced staff

Robust applications

Dependable switching fabric

Redundant service foundation

Chaos Architecture



**A cloud native
availability model**

Digital Transformation



**Disruption and
opportunities**



**Microservices
evolution**



**Cloud
native**



**Dependency
management**



**Chaos
architecture**

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.