

DAT327

AWS re:INVENT

DynamoDB Adaptive Capacity

Rick Houlihan

Senior Practice Manager—NoSQL, AWS

November 30, 2017

AWS
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.




Database capacity planning is critical to running your business, but it's also hard. In this session we'll compare how scaling is usually performed for relational databases and NoSQL databases. We'll look behind the scenes at how DynamoDB shards your data across multiple partitions and servers. Finally, we'll talk about some of the recent enhancements to DynamoDB that make scaling even simpler, particularly a new feature called adaptive throughput that eliminates much of the throttling issues that you may have experienced.


Agenda

- Traditional approaches to database scaling
- How NoSQL databases scale compared to RDBMS
- Evolution of Amazon DynamoDB and adaptive capacity

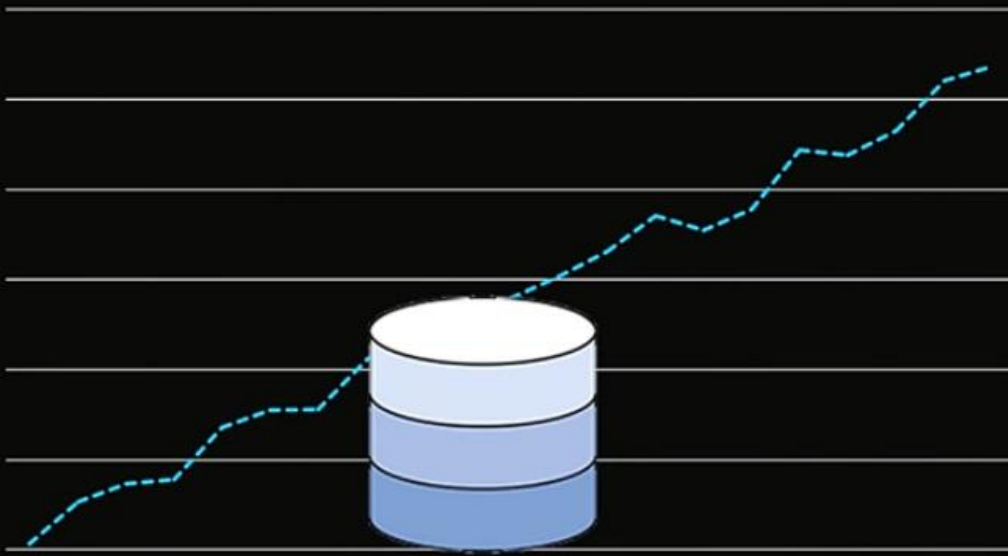
Capacity Planning Is Important (and Hard)



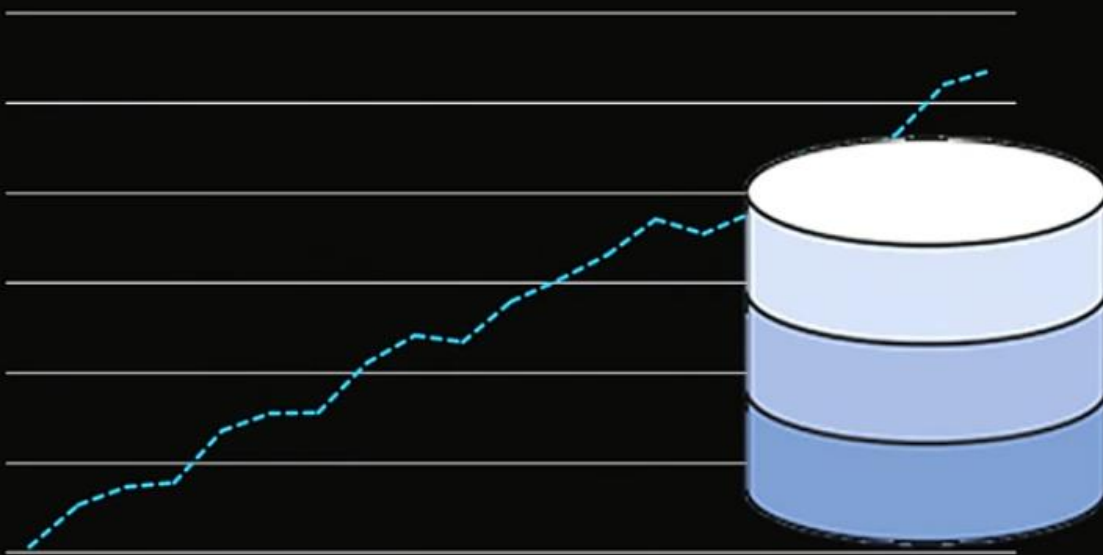
Scaling Relational DBs



Scaling Relational DBs




Scaling Relational DBs



There are several very large single database rack deployments that companies have implemented as data gets bigger


Sharded Relational DBs?



As the data gets even bigger, then you start to partition the database maybe by time or segment/categories to spread your queries as above

NoSQL Databases

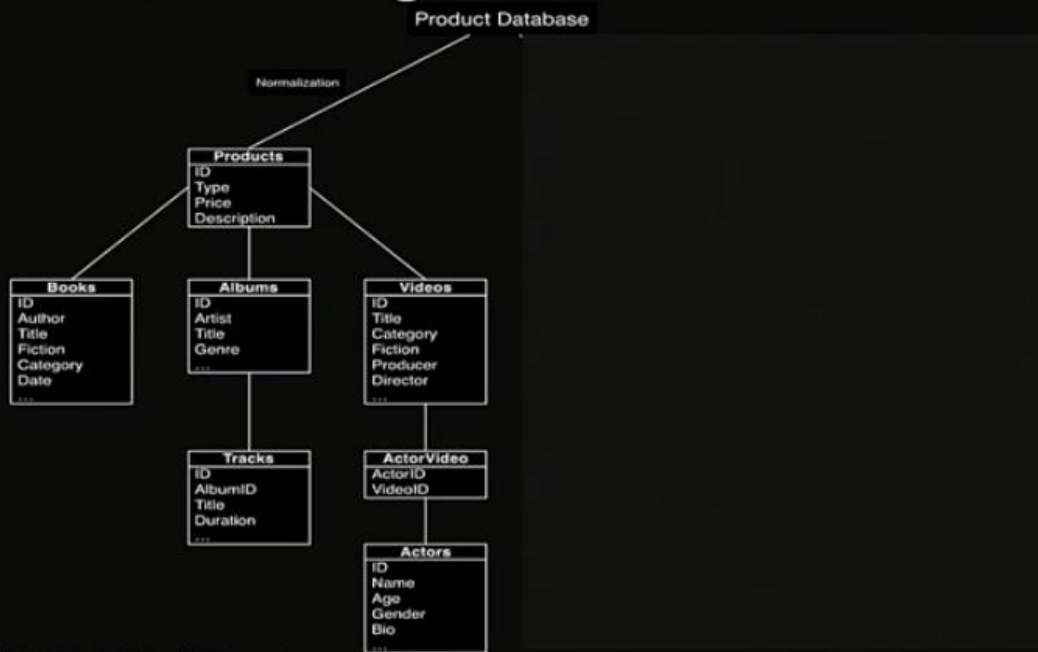
- Denormalize and shard to provide horizontal scale
- Near unbounded throughput and storage



The Iron Triangle of Data—All About CAP



SQL vs. NoSQL Design Pattern

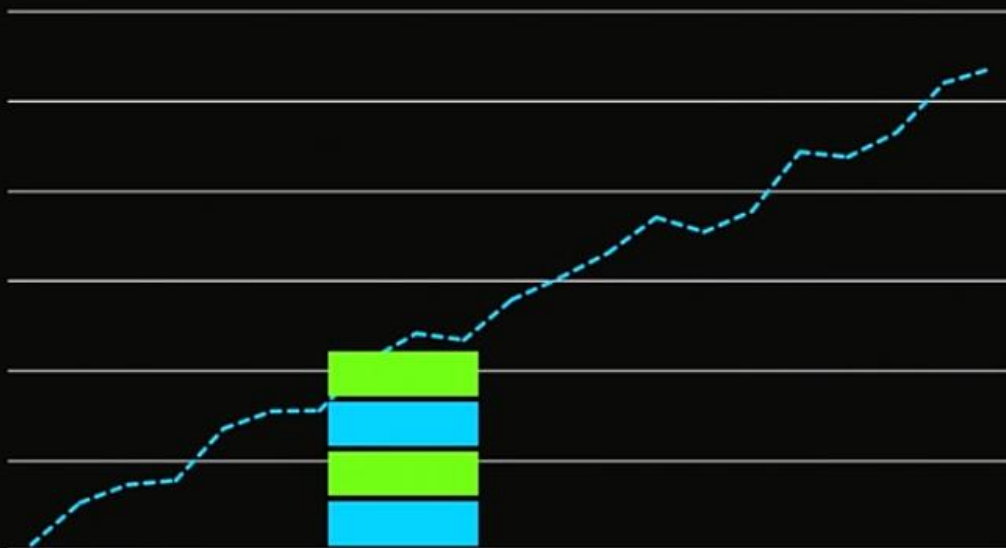


SQL vs. NoSQL Design Pattern



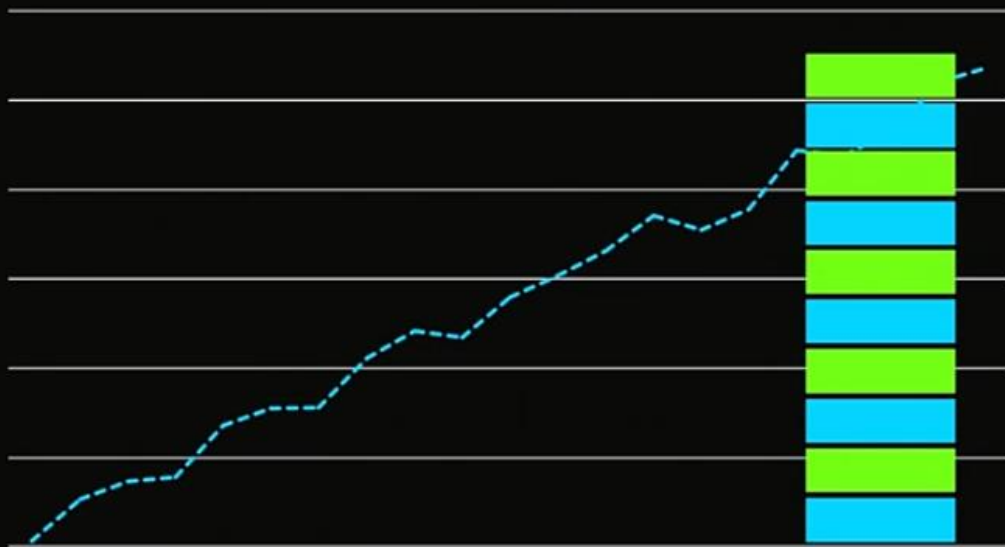
We can use a denormalized datastore by trying to flatten the datastore and push all our data into a set of documents items or collections.

Scaling NoSQL DBs



© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

Scaling NoSQL DBs



© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

There is a replication process or shard as you do this incremental process of replicating data from all the shards in a NoSQL database

Why NoSQL?

SQL

NoSQL

Optimized for storage	Optimized for compute
Normalized/relational	Denormalized/hierarchical
Ad hoc queries	Instantiated views
Scale vertically	Scale horizontally
Good for OLAP	Built for OLTP at scale

AWS re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.




Amazon DynamoDB



Fully managed NoSQL



Document or key-value



Scales to any workload



Fast and consistent



Access control



Event-driven programming

AWS
re:Invent

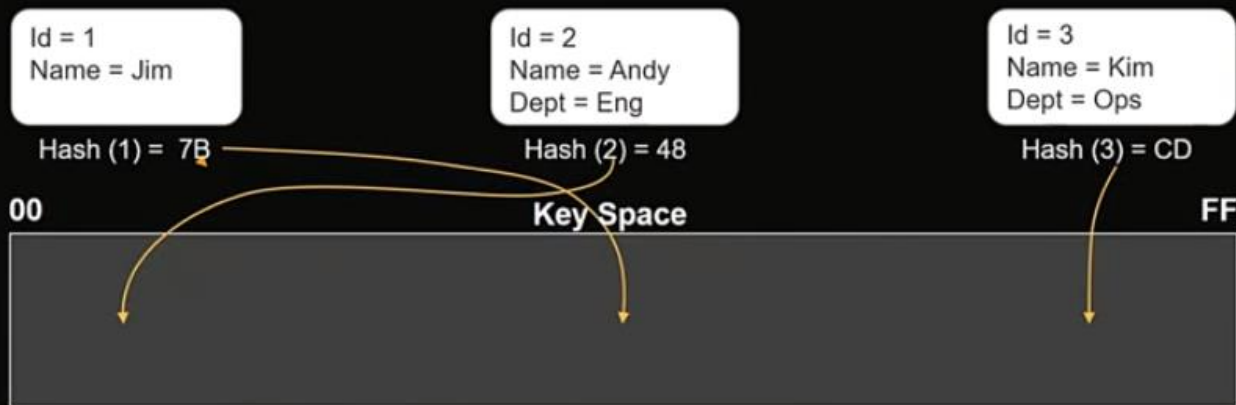
© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



Partition Keys in NoSQL

Partition Key uniquely identifies an item

Partition Key is used for building an unordered hash index




© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

Data is distributed in a NoSQL DB using some shard/partition key/attribute that is mandatory for all items, it uniquely identifies that particular item.

Partition Keys in NoSQL

Partition Key uniquely identifies an item

Partition Key is used for building an unordered hash index



© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

What we do is to hash the unique value and create an unordered item index and then lay out the values in an arbitrary key space as below

Partition Keys in NoSQL

Partition Key uniquely identifies an item

Partition Key is used for building an unordered hash index



© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

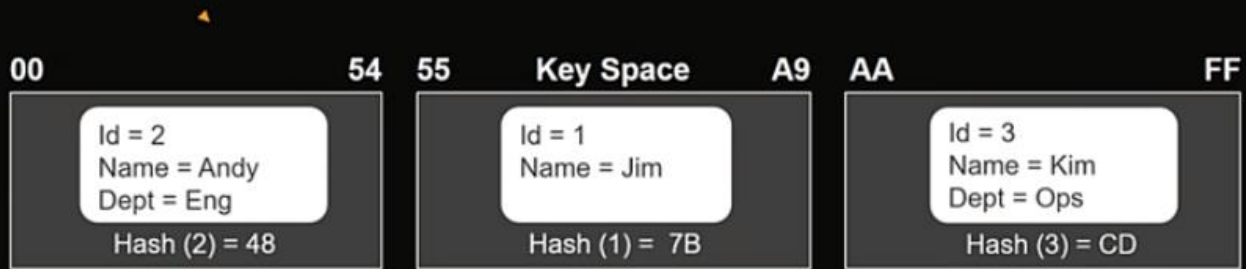
Once we've done this, as we add more items we are then going to start splitting up the arbitrary key space up into visible boxes/servers as below

Partition Keys in NoSQL

Partition Key uniquely identifies an item

Partition Key is used for building an unordered hash index

Allows table to be partitioned for scale



© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

This is how all NoSQL DBs scale

Partitions are three-way replicated



© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

Behind the scenes...





AWS
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



When you create a new DynamoDB table, AWS does not roll out new infrastructure for you. Instead they reclaim capacity for you within the Massively Multi-tenant infrastructure underlying the DynamoDB service.

Behind the scenes...



1K WCU or 3K RCU
Up to 10 GB

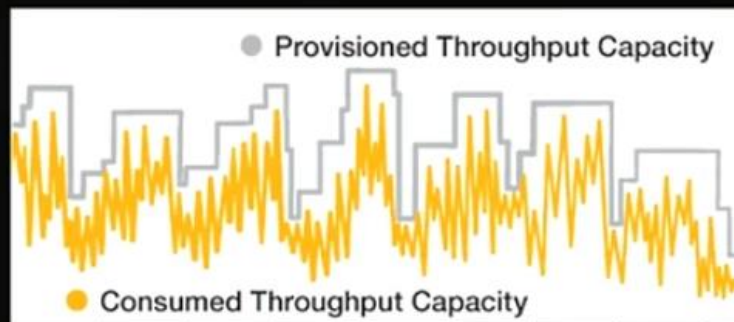
AWS
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



It simply takes your table and your data, chop it up and spread around the infrastructure and share it with other user's tables as above. WCU is write capacity unit, RCU is Read Capacity unit. They both define a partition in DynamoDB.

DynamoDB Scaling



In DynamoDB when the load rises you can simply dial up the provisioned capacity as needed and dial it down again when the load falls using DynamoDB's dynamic back plane.

Exact—forget whole server increments



Quick—additional capacity in minutes



Elastic—meet your needs, up and down



- Maximize availability
- Minimize cost
- No downtime

AWS
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



Results

Hundreds of thousands of customers

Millions of requests per second

Hundreds of billions of items

Petabytes of storage

Adaptive Capacity

Get it right, nobody notices.

Get it wrong, everybody will.

Dynamic Scaling—Throughput and Storage

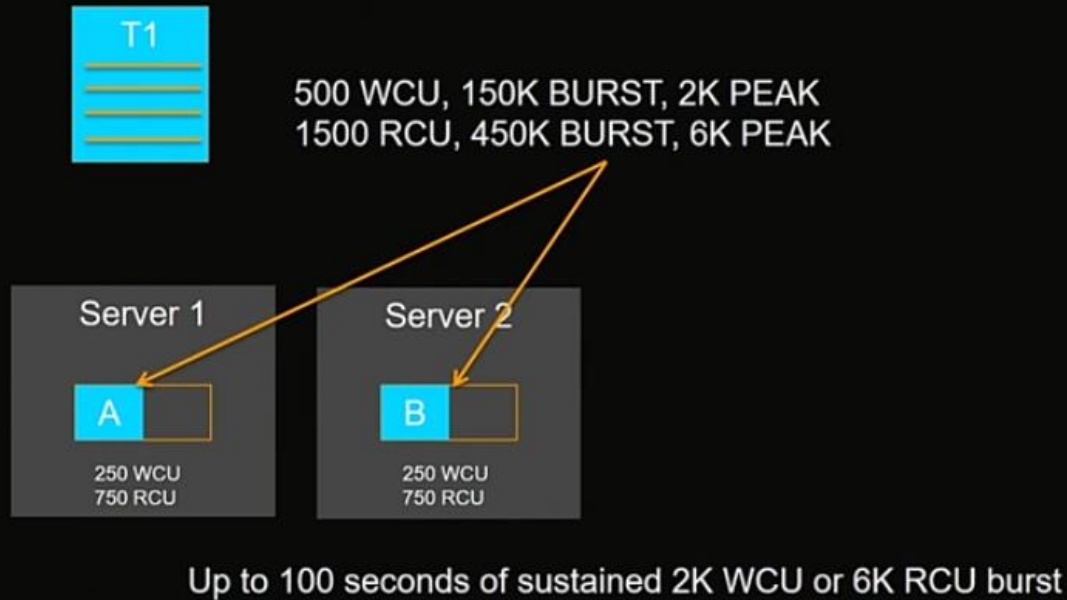


500 WCU, 150K BURST, 1K PEAK
1500 RCU, 450K BURST, 3K PEAK



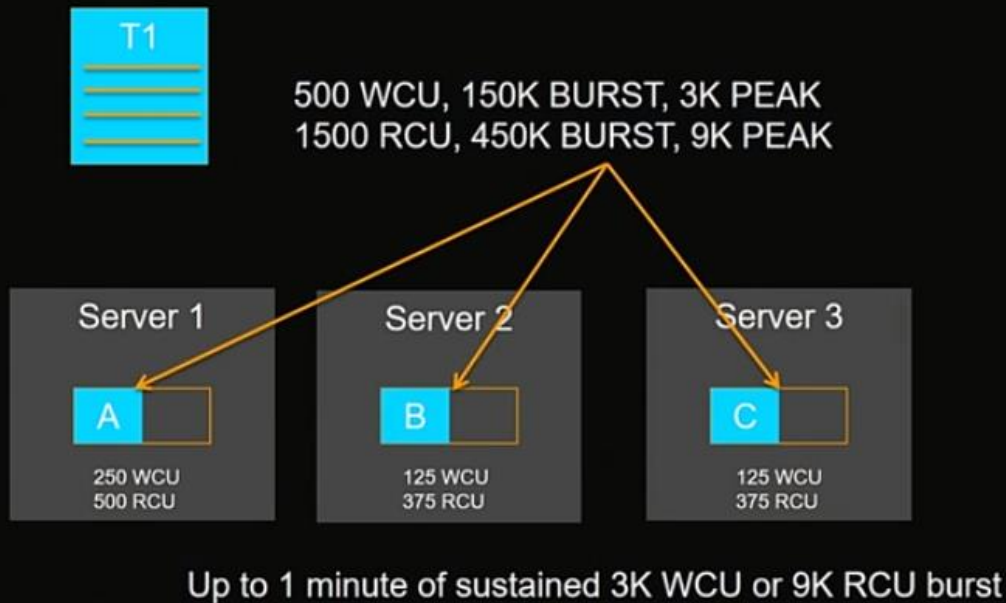
Up to 5 minutes of sustained 1K WCU or 3K RCU burst

Dynamic Scaling—Throughput and Storage



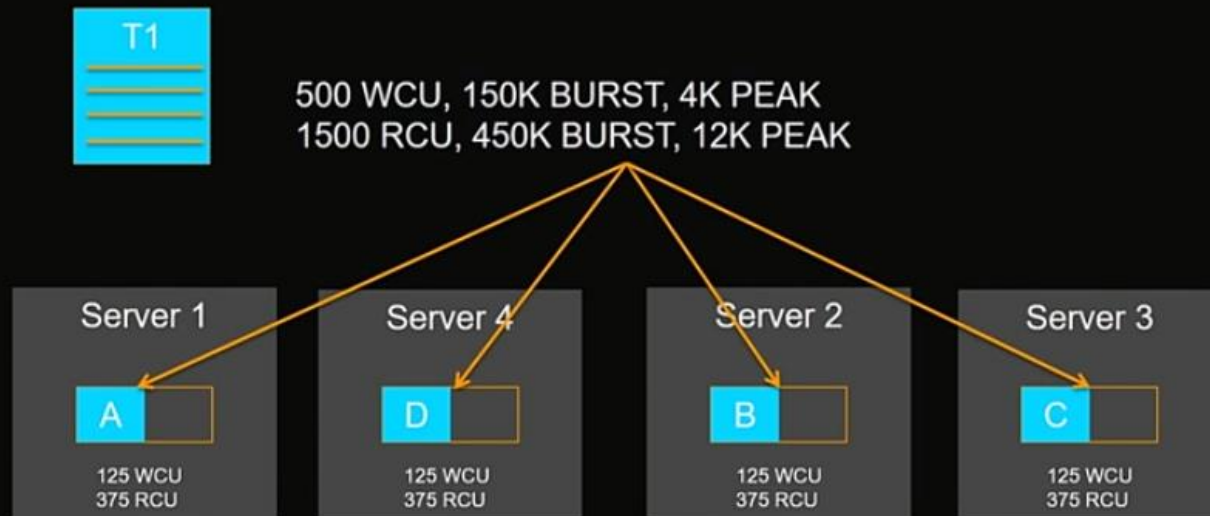
© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

Dynamic Scaling—Throughput and Storage



© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

Dynamic Scaling—Throughput and Storage



Up to 43 seconds of sustained 4K WCU or 12K RCU burst

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

Scenario: Census Application



Statistics Canada (Canada's national statistical agency) hires you to build an online census application

You choose DynamoDB with the following key schema:

Partition Key: **province**


Sort Key: **id**

AWS
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



What you didn't realize...




AWS re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

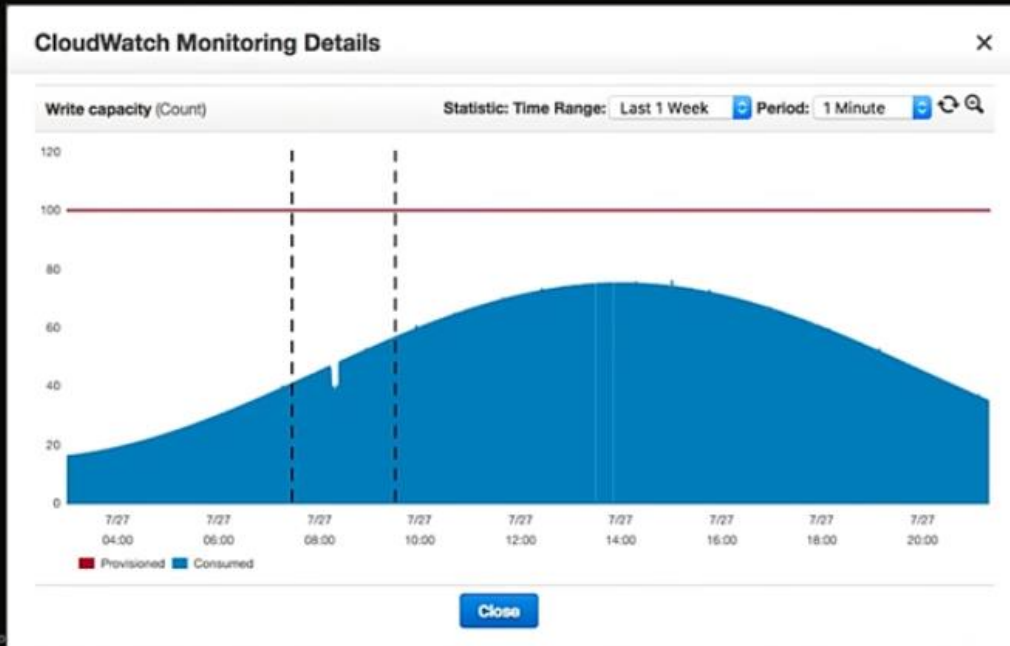


DynamoDB Throughput Allocation

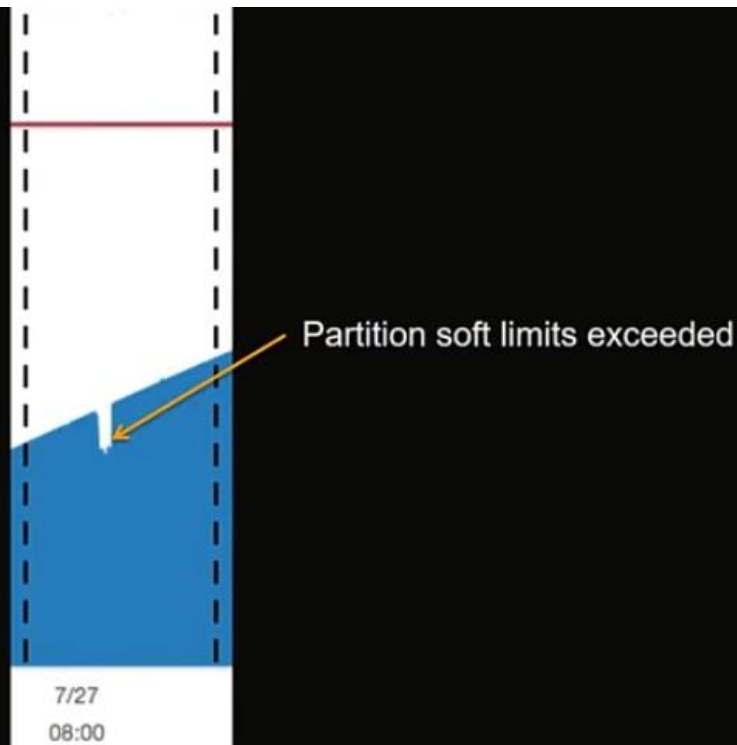


© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

So What Happens?




So What Happens?



At this point there might be a little bit of throttling on the table as we exceed the soft limit and also exceeded the burst bucket and we figure this table need sustained bursting, we then enable adaptive throughput mode and capacity and scaling and start borrowing throughput from the other table partitions WCUs

Adaptive Throughput Kicks In



You will only get adaptive throughput if you have capacity left on your table's other partitions WCU and have adaptive throughput enabled



Everybody wants a fully hands-off experience

Auto Scaling

Throughput automatically adapts to your actual traffic



Without Auto Scaling



With Auto Scaling

AWS re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



For auto-scaling, you just need to set a high-water mark and a low-water mark or thresholds for the system to increase or decrease capacity depending on the load



AWS re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



Time-to-Live (TTL)

Features

- **Automatic:** deletes items from a table based on expiration timestamp
- **Customizable:** user-defined TTL attribute in epoch time format
- **Audit log:** TTL activity recorded in Amazon DynamoDB Streams

Benefits

- **Reduce costs:** delete items when no longer needed
- **Performance:** optimize app performance by controlling table size growth
- **Extensible:** trigger custom workflows (e.g. auto-archive to Glacier) with DynamoDB Streams, AWS Lambda, Amazon Kinesis, etc.

You can use the TTL attribute on your processed data that can get stale and you want them removed from the table. When the time stamp expires, there is a sweeper that comes along to clear out all the stale event data items for you. This allows you to reduce your cost, TTL items are going to show up on the DynamoDB stream and can be processed by a lambda function.

TTL Results



Purged 85 terabytes of stale data and **reduced their costs by over \$200K per year**, while also simplifying their application logic

Amazon CloudWatch



- "Reduced our overall provisioned throughput by 75 percent,"
- "And with that reduction, data-retrieval **latencies were also reduced by up to 10 percent.**"
- "Expect to **save millions of dollars** annually."

Managing Databases

Servers

- Capacity planning
- Provisioning
- Monitoring
- OS patching
- Hardware upgrades
- Provision new regions

Databases

- Database upgrades
- Security patches
- Scaling
- Monitoring
- Performance tuning
- Replication across datacenters
- Re-replicate on server failure

AWS
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

With zero downtime???



Fully Managed, Adaptive, Database Service

Servers

- Capacity planning
- Provisioning
- Monitoring
- OS patching
- Hardware upgrades
- Provision new regions

Databases

- Database upgrades
- Security patches
- Scaling
- Monitoring
- Performance tuning
- Replication across datacenters
- Re-replicate on server failure

DYNAMO  DB

AWS
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

All with zero downtime!!!





AWS re:Invent

Thank you!

AWS
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

