

CLICK TO ADD TEXT

ENT332

# AWS re:INVENT

## Getting Started with Serverless Computing Using AWS Lambda

Chris Munns – Serverless Senior Developer Advocate – AWS

Nicky Joshi – Director Software Engineering – Capital One

November 30, 2017

AWS  
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



## About me

Chris Munns – [munns@amazon.com](mailto:munns@amazon.com), [@chrismunns](https://twitter.com/chrismunns)

- Senior Developer Advocate – Serverless
- New Yorker
- Previously:
  - AWS Business Development Manager – DevOps, July '15–Feb. '17
  - AWS Solutions Architect Nov. '11–Dec. '14
  - Formerly on operations teams @Etsy and @Meetup
  - Little time at a hedge fund, Xerox, and a few other startups
- Rochester Institute of Technology: Applied Networking and Systems Administration '05
- Internet infrastructure geek



# Why are we here today?

## Serverless means...



**No servers to provision or manage**



**Scales with usage**

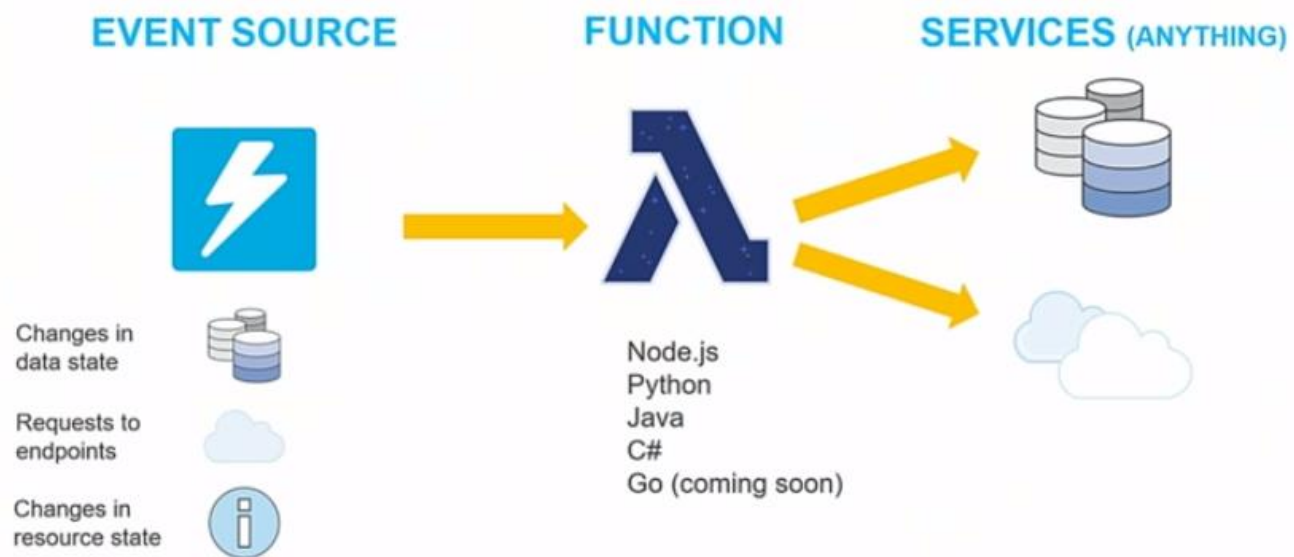


**Never pay for idle**



**Availability and fault tolerance built in**

# Serverless applications



© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

## CUSTOMERS LOVE SERVERLESS



**AWS re:Invent**

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.





# Common serverless use cases



## Web Applications

- Static websites
- Complex web apps
- Packages for Flask and Express



## Backends

- Apps & services
- Mobile
- IoT



## Data Processing

- Real time
- MapReduce
- Batch



## Chatbots

- Powering chatbot logic



## Amazon Alexa

- Powering voice-enabled apps
- Alexa Skills Kit



## IT Automation

- Policy engines
- Extending AWS services
- Infrastructure management

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



# Using AWS Lambda



## Bring your own code

- Node.js, Java, Python, C#
- Bring your own libraries (even native ones)



## Simple resource model

- Select power rating from 128 MB to 3 GB
- CPU and network allocated proportionately



## Flexible use

- Synchronous or asynchronous
- Integrated with other AWS services



## Flexible authorization

- Securely grant access to resources and VPCs
- Fine-grained control for invoking your functions

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

# Using AWS Lambda



## Authoring functions

- WYSIWYG editor or upload packaged .zip
- Third-party plugins (Eclipse, Visual Studio)



## Monitoring and logging

- Metrics for requests, errors, and throttles
- Built-in logs to Amazon CloudWatch Logs



## Programming model

- Use processes, threads, /tmp, sockets normally
- AWS SDK built in (Python and Node.js)




## Stateless

- Persist data using external storage
- No affinity or access to underlying infrastructure

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

# Fine-grained pricing



Buy compute time in 100ms increments

Low request charge

No hourly, daily, or monthly minimums

No per-device fees

## Free Tier


1M requests and 400,000 GB of compute.  
Every month, every customer.

**Never pay for idle**

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

# Lambda execution model

## Synchronous (push)



## Asynchronous (event)



## Stream-based



These are the 3 ways you can invoke your Lambda function.

# Event sources that trigger AWS Lambda



© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

## Lambda permissions model

**Fine-grained security controls for both execution and invocation:**

### Execution policies:

- Define what AWS resources/API calls can this function access through IAM
- Used in streaming invocations
- E.g., "Lambda function A can read from DynamoDB table users"

### Function policies:


- Used for sync and async invocations
- E.g., "Actions on bucket X can invoke Lambda function Z"
- Resource policies allow for cross-account access

```
1- {  
2-   "Version": "2012-10-17",  
3-   "Statement": [  
4-     {  
5-       "Effect": "Allow",  
6-       "Action": [  
7-         "logs:CreateLogGroup",  
8-         "logs:CreateLogStream",  
9-         "logs:PutLogEvents"  
10-      ],  
11-      "Resource": "*"   
12-    }  
13-  ]  
14- }
```




There are several ways to control use of your lambda function, the place where it sits like a VPC as well as its execution and operating policies.

# Amazon API Gateway



# Amazon API Gateway



Create a unified API front end for multiple micro-services



DDoS protection and throttling for your backend



Authenticate and authorize requests to a backend



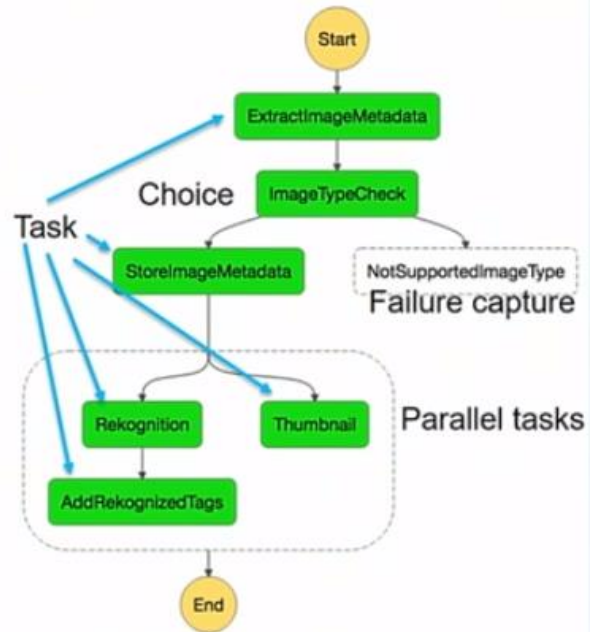
Throttle, meter, and monetize API usage by third-party developers



# AWS Step Functions

## "Serverless" workflow management with zero administration:

- Makes it easy to coordinate the components of distributed applications and microservices using visual workflows
- Automatically triggers and tracks each step, and retries when there are errors, so your application executes in order and as expected
- Logs the state of each step, so when things do go wrong, you can diagnose and debug problems quickly



## Amazon Lex



Service for building conversational interfaces into any application using voice and text

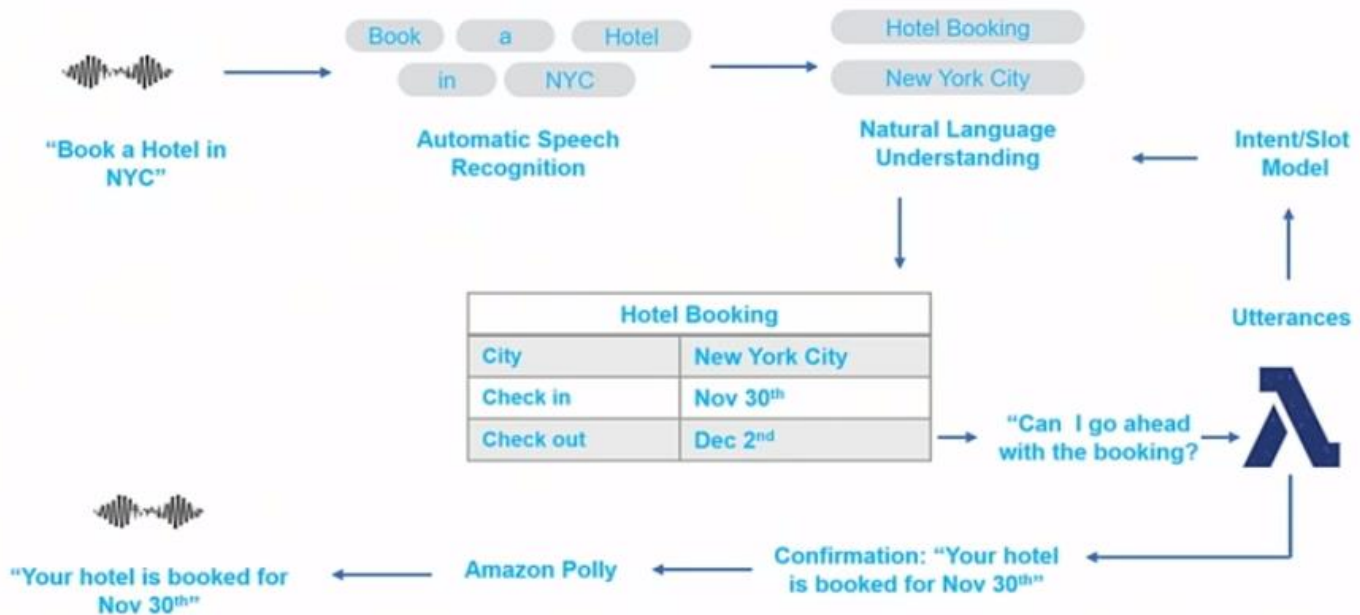
Automatic speech recognition (ASR) for converting speech to text

Natural language understanding (NLU) to recognize the intent of messages

Powered by the same deep learning technology as Alexa

Fully managed service


# "Book a Hotel"



## AWS Lambda + Amazon Kinesis

### Real-time data processing:

1. Real-time event data sent to **Amazon Kinesis**, allows multiple **AWS Lambda** functions to process the same events.
2. In **AWS Lambda**, Function 1 processes and aggregates data from incoming events, then stores result data in **Amazon DynamoDB**.
3. Lambda Function 1 also sends values to **Amazon CloudWatch** for simple monitoring of metrics.
4. In **AWS Lambda** function, Function 2 does data manipulation of incoming events and stores results in **Amazon S3**.



<https://s3.amazonaws.com/awslambda-reference-architectures/stream-processing/lambda-refarch-streamprocessing.pdf>


# Serverless distributed computing: PyWren

**PyWren Prototype** developed at University of California, Berkeley

Uses **Python** with **AWS Lambda** stateless functions for large-scale data analytics

Achieved @ **30-40 MB write and read** performance per-core to Amazon S3 object store

Scaled to **60-80 GB** across **2,800 simultaneous** functions



**Build PCI and HIPAA-compliant serverless applications!**



Serverless platform services that can be used in both:



AWS  
Lambda



Amazon  
S3



Amazon  
CloudFront



Amazon  
DynamoDB



Amazon  
Kinesis  
Streams



Amazon  
Cognito



Amazon API  
Gateway



Amazon  
SNS





## Frameworks

APEX



© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

Find a framework that fits the need you are looking to feel.



# Claudia.js



Node.js framework for deploying projects to AWS Lambda and Amazon API Gateway

- Has sub-projects for microservices, chatbots, and APIs
- Simplified deployment with a single command
- Use standard NPM packages, no need to learn Swagger
- Manage multiple versions

```
1 var ApiBuilder = require('claudia-api-builder')
2
3 var api = new ApiBuilder();
4
5 module.exports = api;
6
7 api.get('/hello', function () {
8   return 'hello world';
9 });
10
```

<https://claudiajs.com>

<https://github.com/claudiajs/claudia>

\$ claudia create --region us-east-1 --api-module app

# Chalice



Python serverless "microframework" for AWS Lambda and Amazon API Gateway

- A command line tool for creating, deploying, and managing your app
- A familiar and easy-to-use API for declaring views in python code
- Automatic Amazon IAM policy generation

```
1 from chalice import Chalice
2
3 app = Chalice(app_name="helloworld")
4
5 @app.route("/")
6 def index():
7     return {"hello": "world"}
8
9
```

<https://github.com/aws/chalice>

<https://chalice.readthedocs.io>

\$chalice deploy

# AWS Serverless Application Model (AWS SAM)



Template-driven resource management model optimized for serverless

New serverless resource types: Functions, APIs, and tables

Supports anything AWS CloudFormation supports

Open specification (Apache 2.0)

<https://github.com/awslabs/serverless-application-model>

## SAM template

```
AWSTemplateFormatVersion: '2010-09-09'
Transform: AWS::Serverless-2016-10-31
Resources:
  GetHtmlFunction:
    Type: AWS::Serverless::Function
    Properties:
      CodeUri: s3://sam-demo-bucket/todo_list.zip
      Handler: index.gethtml
      Runtime: nodejs4.3
      Policies: AmazonDynamoDBReadOnlyAccess
      Events:
        GetHtml:
          Type: Api
          Properties:
            Path: /{proxy+}
            Method: ANY

  ListTable:
    Type: AWS::Serverless::SimpleTable
```

Tells AWS CloudFormation this is a SAM template it needs to "transform."

Creates a Lambda function with the referenced managed IAM policy, runtime, code at the referenced .zip location, and handler as defined. Also creates an API Gateway and takes care of all necessary mapping/permissions.

Creates a DynamoDB table with five Read & Write units.

SAM is all about managing the AWS resources,

# SAM template



From: <https://github.com/aws-labs/aws-serverless-samfarm/blob/master/api/saml.yaml>

This piece of code will generate the 6 different AWS resources on the right for you automatically. It is better than using the CLI to create these same resources.

## Introducing SAM Local

CLI tool for local testing of serverless apps

Works with Lambda functions and "proxy-style" APIs

Response object and function logs available on your local machine

Uses open-source docker-lambda images to mimic Lambda's execution environment:

- Emulates timeout, memory limits, runtimes

<https://github.com/aws-labs/aws-sam-local>





The screenshot displays the AWS Cloud9 IDE interface. On the left, a file explorer shows a project named 'LambdaHome1' with a subdirectory 'sam-dance-serverless-stz' containing a 'README.md' file. The main editor window shows a file named 'index.js' with the following JavaScript code:

```
1 'use strict';
2 var moment = require('moment');
3
4 exports.handler = (event, context, callback) => {
5
6   var originURL = process.env.ORIGIN_URL || '';
7
8   emitLambdaAge();
9
10  // This variable can be updated and checked in to your repository
11  // to update the number of SAM squirrels on the screen.
12  var samCount = 10;
13
14  // Or you can update your Lambda function's environment variables
15  var samMultiplier = process.env.SAM_MULTIPLIER || 1;
16
17  var totalSAMs = samCount * samMultiplier;
18
19  console.log('The number of SAMs to show: ' + samCount);
20  console.log('Multiplier to apply to SAMs: ' + samMultiplier);
21  console.log('Total number of SAMs to show: ' + totalSAMs);
22
23  callback(null, {
24    "statusCode": 200,
25    "body": totalSAMs,
26    "headers": {
27      "Access-Control-Allow-Headers": "Content-Type,X-Amz-Date,Authorization,X-Api-Key,X-Amz-Security-Token",
28      "Access-Control-Allow-Methods": "GET,OPTIONS",
29      "Access-Control-Allow-Origin": originURL
30    }
31  });
32 }
33
34
35 function emitLambdaAge() {
36   var now = moment();
37   var lambdaAnnouncement = moment('2014-11-04');
38
39   var daysOld = now.diff(lambdaAnnouncement, 'days');
```

On the right, the 'Run' tab is active, showing the execution results for the function 'GetSAMPartyCount'. The test payload is set to '1'. The execution results show a successful response with a status code of 200 and a body of 10. The function logs show the following output:

```
2017-11-30 18:17:46.245 Lambda is 1122 days old!
2017-11-30 18:17:46.246 The number of SAMs to show: 10
2017-11-30 18:17:46.246 Multiplier to apply to SAMs: 1
2017-11-30 18:17:46.246 Total number of SAMs to show: 10
```

The request ID is 94551e04-8aac-1b8c-d21-2a822a274893.

# Capital One Auto – Journey to Serverless

**AWS re:Invent**

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



Let us see the way we migrated one of our marketing web applications to a Serverless architecture.




# Capital One



- Top 10 U.S. Bank
- Alexa Website Ranking says capitalone.com traffic is "81st in U.S. websites"
- Auto Financing-1M hits/month

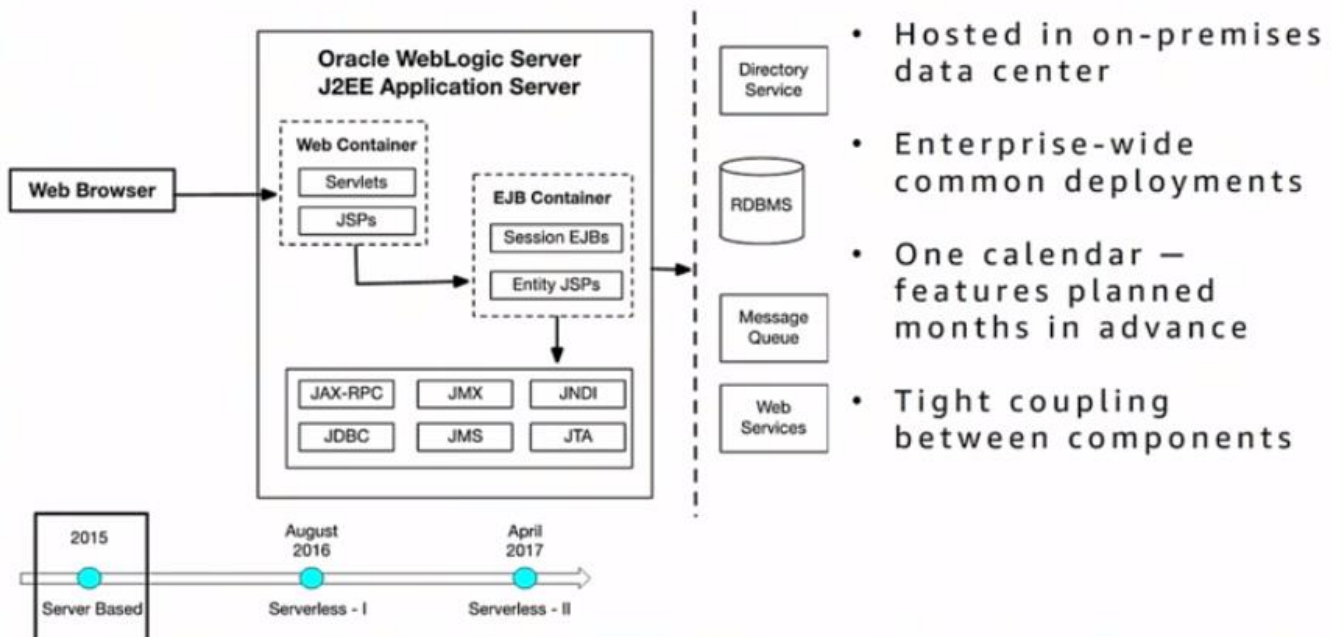
## Migration Journey Timeline



# Migration Requirements

- Full functionality and more
- Secure – yes, we actually are a Financial Institution
- Resilient (Active/Active)
- Responsive – front-facing Marketing Home Page, so initial load time is critical
- SEO friendly
- Continuous deployments – changes on demand
- Low maintenance – it just needs to run
- Use existing tools and processes in the enterprise – don't re-invent the wheel
  - Logging
  - Monitoring
  - Deployments

## The Painful Old Days

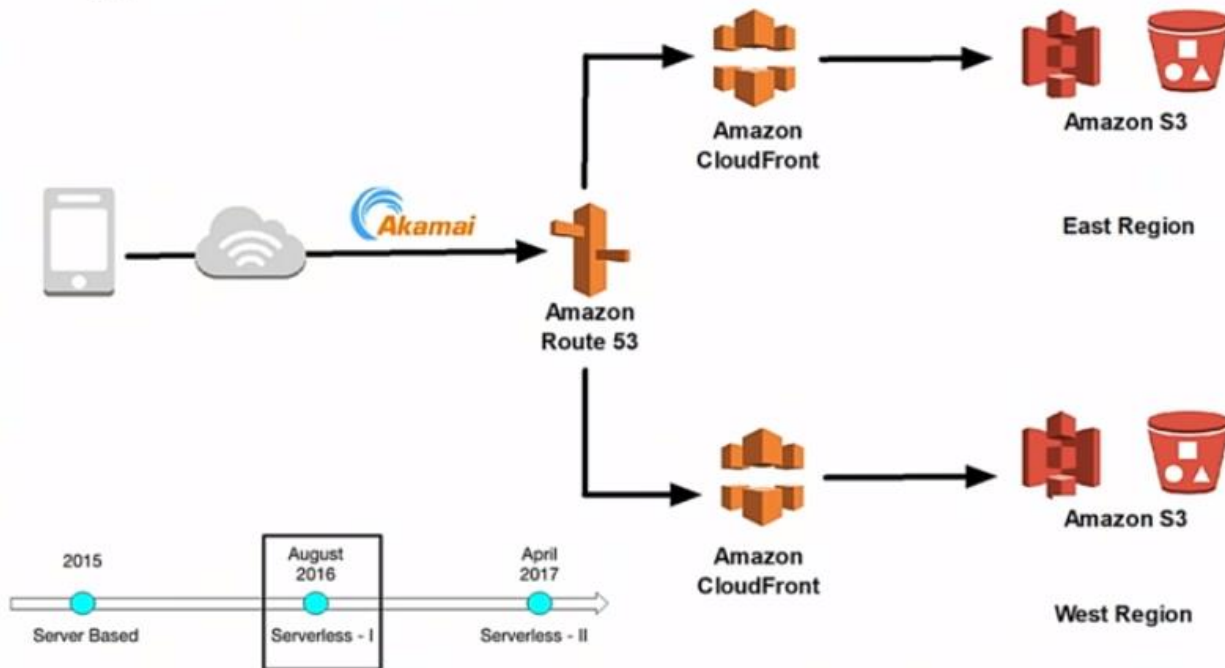


## Serverless Migration – I

Move to AWS

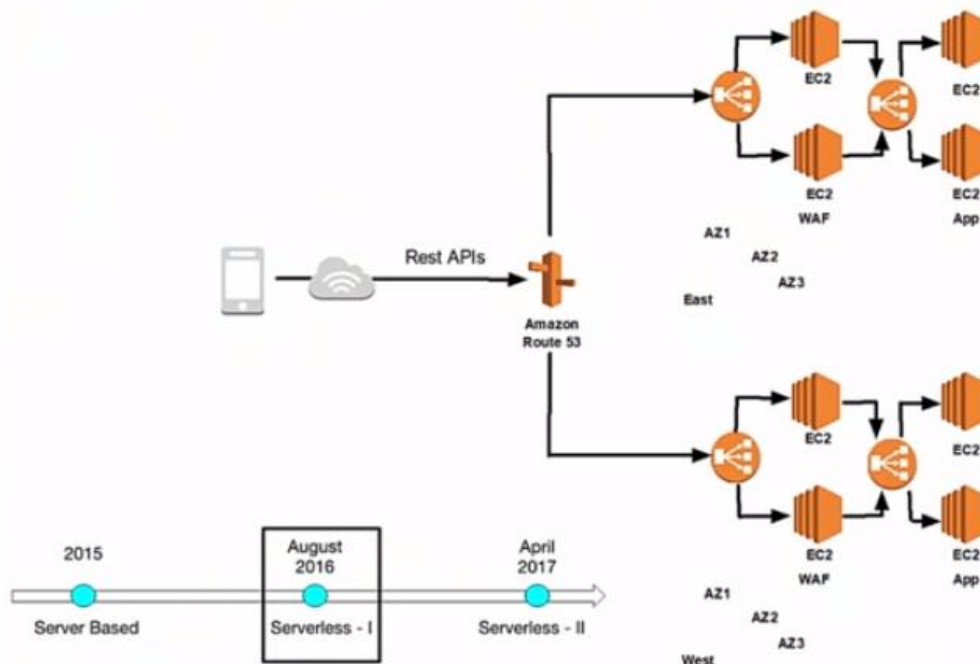
Serverless Content Strategy

## App architecture—content



We created an Angular SPA app that was bundled and stored in an S3 bucket with web hosting enabled. We then replicated the S3 bucket in both the East and West regions for an active-active deployment, then we fronted the web app using Cloudfront that has the certificates for SSL termination and also using Route53. We were also able to leverage non-AWS services like Akamai for content caching.


## App Architecture – APIs




This is the architecture that we used for our APIs. Due to internal company policies in 2016 we couldn't go all serverless for our backends and had to come up with the alternate EC2-based architecture above. The first set of EC2 instances are running Apache with security configurations and serve as our web application firewall while the 2 EC2 instances behind

the firewall are running Tomcat with our application logic running in them. Next, we then replicated this whole setup in another region for an active-active configuration.

## CI/CD for Content



- App templates in GitHub
- Content in CMS hosted service
- Pre-render angular content; SEO Friendly
- Use AWS CLI to push content to S3 bucket for web hosting

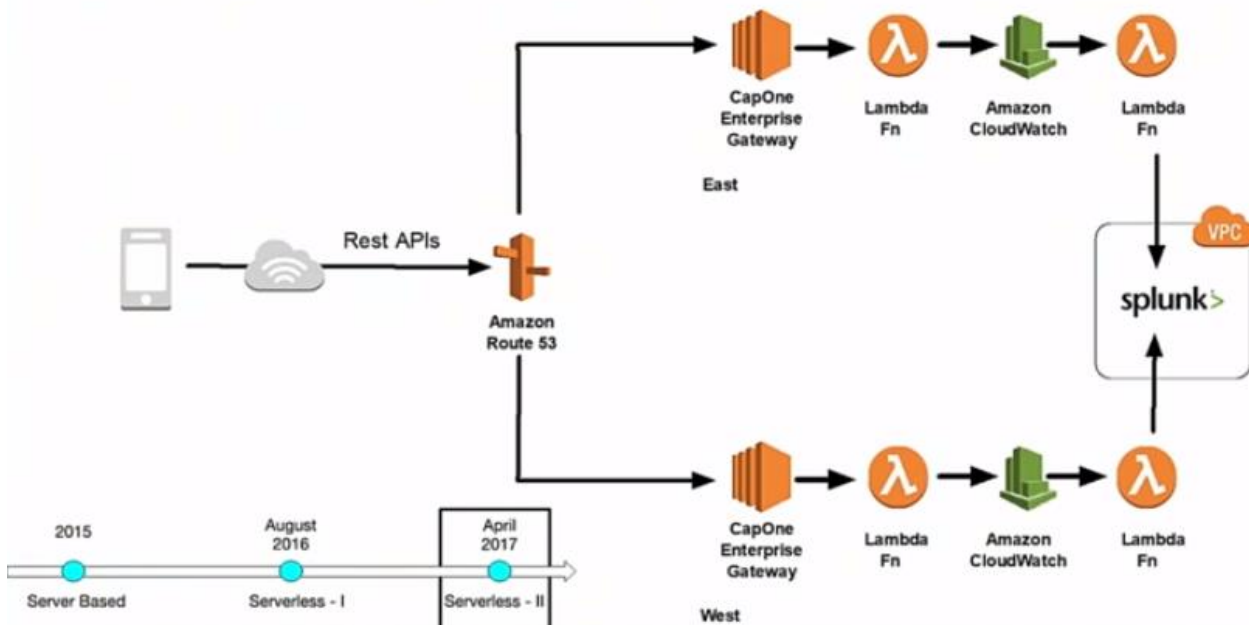


## Serverless Migration - II

Lambda-based APIs  
CI/CD Processes




# Move to Serverless APIs




We have now eliminated the use of EC2 instances in our web app hosting architecture

## CI/CD – Lambda Functions



- Lambda code in GitHub
- Jenkins triggered build
- Uses AWS CLI
  - Deploy in S3 bucket
  - Create/update Lambda function
- Other options:
  - SAM
  - Serverless framework



For our CI/CD process for our lambda functions, Jenkins is our enterprise build tool of choice. We store our lambda function code in GitHub that triggers a build job in Jenkins, the Jenkins job then takes that code, uses the AWS CLI to create a deployment bundle that it stores in S3. By leveraging functions like lambda create and lambda update, we are able to create and update our lambda functions accordingly.

## Benefits of Serverless

- Super simple architecture
- Runs itself, less monitoring
- Scales itself, no work on Auto Scaling needed
- No worries on AMI rehydration, which keeps compliance and operations happy
- **Cost savings:**
  - Conservative savings to tune of [\\$50K/year](#)
  - Amazon S3 web hosting eliminated any Amazon EC2 needs for website hosting
  - Able to eliminate 20 EC2 instances, CLBs, EBS volumes from API architecture
  - Countless hours saved on operations of application

## Lessons Learned

- Migration can be a journey, so plan accordingly
- Step-by-step progress is recommended
  - Start small
- Serverless architectures are flexible and fungible
  - No one size fits all
- Reuse tools and processes in the organization
- Get stakeholder buy-in early

# FIN, ACK

## Serverless:

- No servers to manage
- No cost for idle
- Automatic scaling
- High availability

## Use cases:


- Web applications
- Backends
- Data processing
- Chatbots
- Amazon Alexa
- IT Automation


## Integrated across AWS:

- Amazon API Gateway
- AWS Step Functions
- Amazon S3
- Amazon Kinesis
- Amazon DynamoDB
- Amazon SNS
- Amazon Cognito
- AWS CloudFormation
- AWS CodePipeline
- Amazon CloudWatch
- AWS IoT
- many more!

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

## aws.amazon.com/serverless

 Menu



[Contact Sales](#)

[Products](#) ▾

[Solutions](#)

[Pricing](#)

[More](#) ▾

English ▾

My Account ▾

[Sign in to the Console](#)

Serverless Computing

Overview

AWS Serverless Application Repository

Developer Tools

Resources

Partners

# Serverless Computing and Applications

Build and run applications without thinking about servers

Find serverless applications

Serverless computing allows you to build and run applications and services without thinking about servers. Serverless applications don't require you to provision, scale, and manage any servers. You can build them for [nearly any type of application](#) or backend service, and everything required to run and scale your application with high availability is handled for you.

Building serverless applications means that your developers can focus on their core product instead of worrying about managing and operating servers or runtimes, either in the cloud or on-premises. This reduced overhead lets developers reclaim time and energy that can be spent on developing great products which scale and that are reliable.

ENT332

AWS  
re:Invent

Thank you!

Remember to review this session!

AWS  
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

