

GPSTEC313

# AWS re:INVENT

GPS: Real-Time Data Processing with AWS  
Lambda Quickly, at Scale, and What Comes  
Next?

Juan Villa | Partner Solutions Architect

November 27, 2017

AWS  
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



Real-time data processing is a powerful technique that allows businesses to make agile automated decisions. This process is particularly powerful when applied to workloads like **security**, analyzing access logs, parsing audit logs, and monitoring API activity to detect behavior anomalies. Combined with automation, business can quickly take action to remediate security concerns, or even train a machine learning (ML) model. We explore different techniques for analyzing real-time streams on AWS using **Lambda**, **Amazon Kinesis**, Spark with **Amazon EMR**, and **Amazon DynamoDB**. We also cover best practices around short- and long-term storage and analysis of data and, briefly, the possibility of leveraging ML.

## About Me—Juan Villa



- Partner Solutions Architect
- Focus on IoT, serverless, and machine learning
- Love to build cool software and hardware projects
- Ex-Software Developer

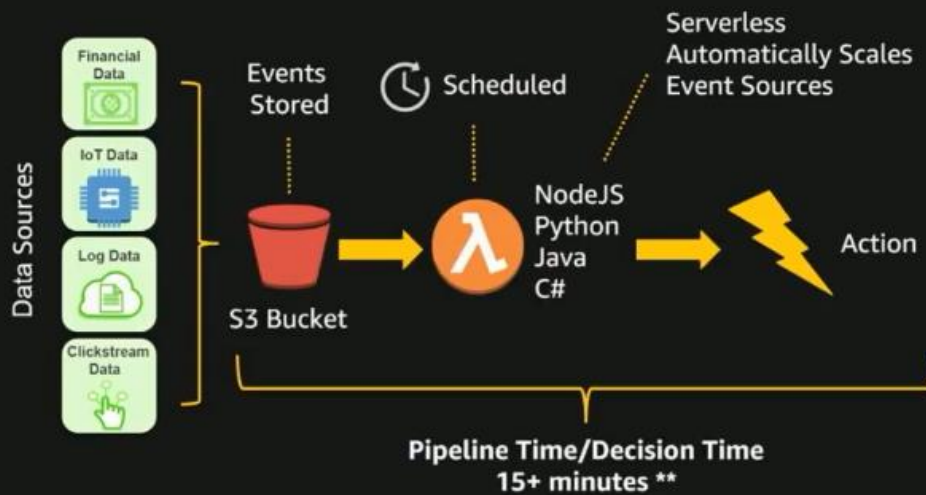
AWS  
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



# Real-Time vs Batch Processing?

## Batch Processing



AWS  
re:Invent

\*\* application specific, and can be less than 15 minutes, or more than 1 day

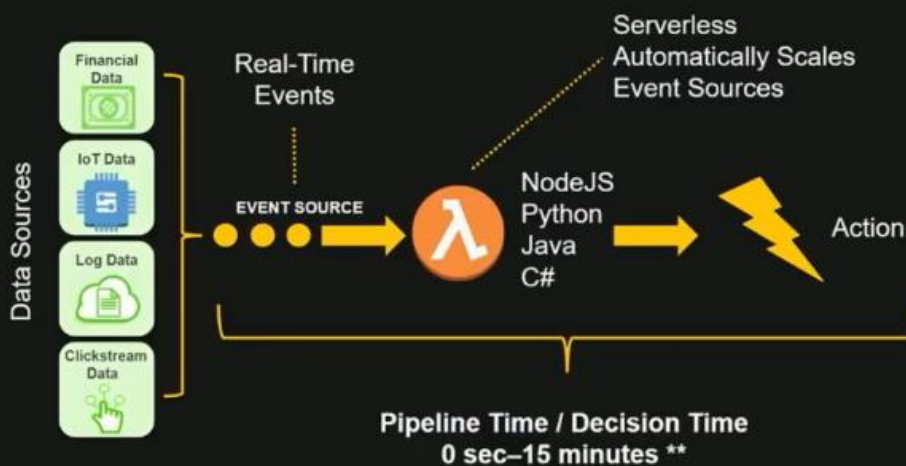


© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

**Batch processing** is about collecting events and processing them in some form of scheduled manner. **Clickstream data** events can be stored in an **S3 bucket** and then use lambda as the compute source to run specialized analytics on the events. All these data can be pulled and analyzed on some specified schedule.

# Real-Time vs Batch Processing?

## Real-Time (Event Driven)



AWS  
re:Invent

\*\* as fast as possible, depends on nature of event  
Near real time?



© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

For **Real-Time processing**, we are going to be processing the message in real-time before storing it in **S3** buckets.

# AWS Lambda

## Serverless Compute



**No Infrastructure to manage:** Compute without managing infrastructure like Amazon EC2 instances and Auto Scaling groups.

**Bring your own code:** Stateless, event-driven code with native support for Node.js, Java, Python, and C# languages.

**Triggered by events:** Direct Sync & Async API calls, AWS service integrations, and third-party triggers.

**Efficient performance at scale:** Easy to author, deploy, maintain, secure, and manage. Focus on business logic to build back-end services that perform at scale.

**Cost-effective:** Automatically matches capacity to request rate. Compute cost 100 ms increments.

## Data Sources and Applications



Load Balancer

..... Parse Load Balancer Access Logs looking for anomalous patterns



CloudTrail

..... Parse CloudTrail events to detect suspicious patterns such as logins from another country



Flow Logs

Parse VPC Flow Logs to detect anomalies and abuse



CloudWatch  
(logs, events, alarms)

Configure event triggering and alarms from sources such as CloudTrail or other event sources

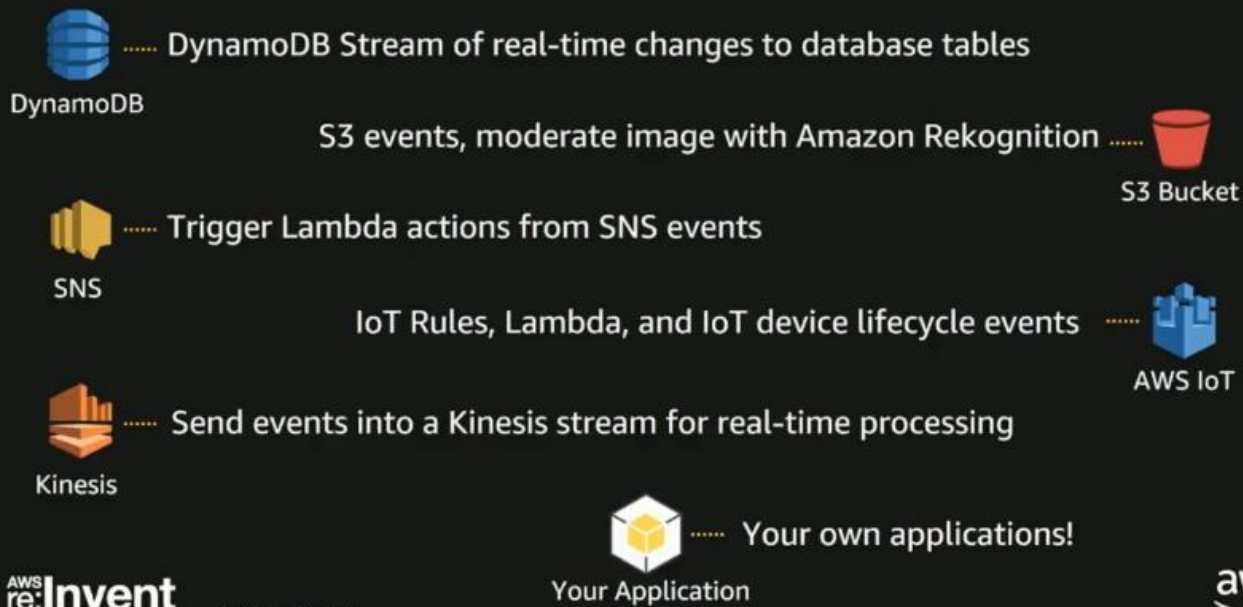
AWS  
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



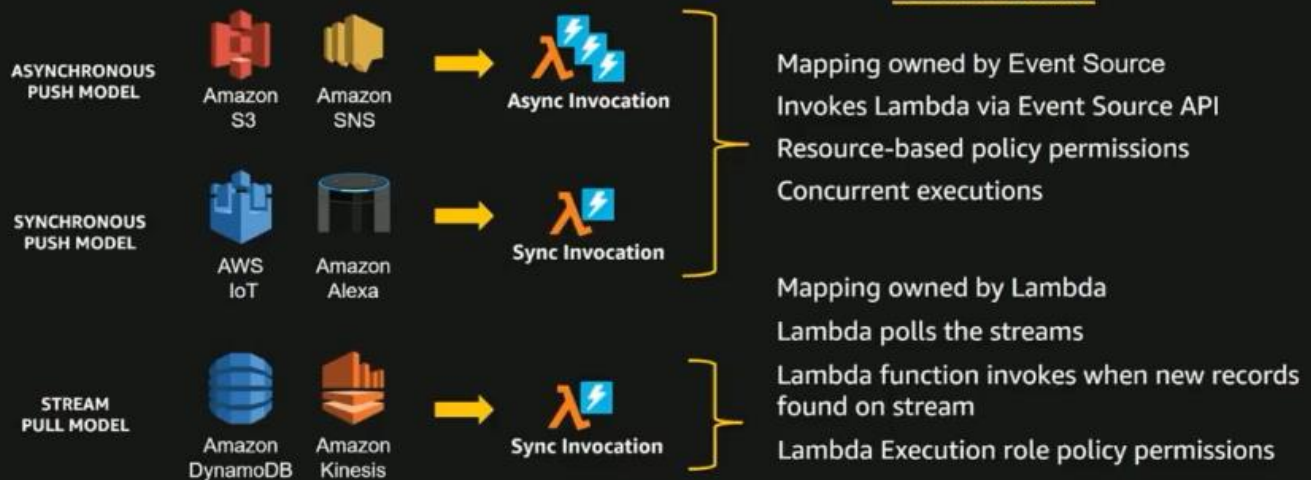


# Data Sources and Applications



## Lambda Event Driven Models

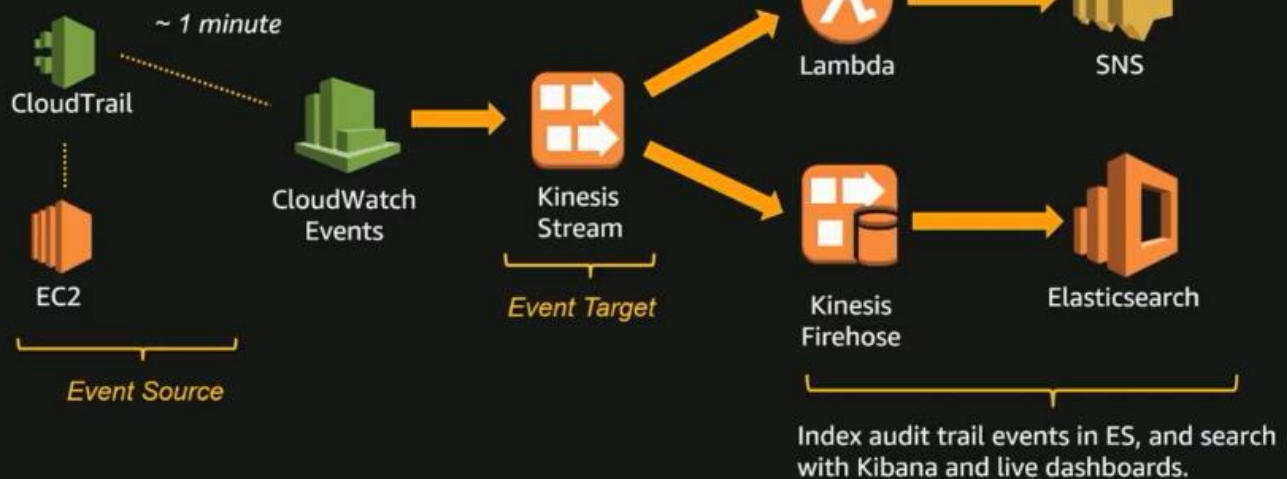
### HOW IT WORKS



These are the 2 types of invocations models for triggering a lambda function. Sync invocation is when you have to wait on the lambda to get the result back and use further down the pipeline.

# Pipeline Example #1

## Parsing Security Events



AWS re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



We are going to turn on and use **CloudTrail** monitor all your AWS **EC2** calls, this will let us know if something was turned off and by whom. We then set up a **CloudWatch event** on a specific event like **EC2 Prod instance termination** from the **CloudTrail logs** and put that event in a **Kinesis Stream**. We then have 2 consumers that are going to be consuming data coming from the Kinesis stream, a **lambda function** will put the message into an **SNS** topic, and a **Kinesis Firehose** that will put the event into **ElasticSearch** for searching and display using Kibana

# Configuring CloudWatch Event Rules

## Parsing Security Events

- Configure "Event Source"
  - EC2 CloudTrail "TerminateInstances" event
- Can have more than one event type in source filter
- Configure event "Targets"
  - Kinesis
  - SNS
- Can have more than one target

Configure Target

AWS re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



# Configuring CloudWatch Event Rules

## Parsing Security Events

### CloudWatch Event Source Pattern

```
{
  "source": [
    "aws.ec2"
  ],
  "detail-type": [
    "AWS API Call via CloudTrail"
  ],
  "detail": {
    "eventSource": [
      "ec2.amazonaws.com"
    ],
    "eventName": [
      "TerminateInstances"
    ]
  }
}
```

- Will get EC2 CloudTrail audit events on the account
  - CloudTrail audit events track **API events**, which includes the **principal**, **target resource**, and **action** that took place
  - Specifically filtering for **"TerminateInstances"** events
  - On average, events are emitted in ~1–2 minutes

AWS  
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



## What Does the Event Look Like?

### Parsing Security Events

*Summarized*

#### What do we need?

- IAM user that terminated the instance
- When the instance was terminated
- IP address of user that terminated the instance
- What instance was terminated

```
{
  "id": "aa7d2051-b082-6b8e-8d36-f7e72064ab5c",
  "detail-type": "AWS API Call via CloudTrail",
  "source": "aws.ec2",
  "account": "XXXXXXXXXXXX",
  "time": "2017-10-03T19:46:45Z",
  "region": "us-east-1",
  "detail": {
    "userIdentity": {
      "type": "IAMUser",
      "arn": "arn:aws:iam::332631645289:user/jcv123",
      "accountId": "XXXXXXXXXXXX",
      "userName": "jcv123"
    },
    "eventTime": "2017-10-03T19:46:45Z",
    "eventSource": "ec2.amazonaws.com",
    "eventName": "TerminateInstances",
    "awsRegion": "us-east-1",
    "sourceIPAddress": "324.21.198.71",
    "requestParameters": {
      "instancesSet": {
        "items": [ { "instanceId": "i-098aa789cd6b411c0" } ]
      }
    },
    "requestID": "3b6aeef7-c704-40f8-9886-a2091f04d290",
    "eventID": "e9a6003d-12e7-4efa-b71d-a13f4f171fd5",
    "eventType": "AwsApiCall"
  }
}
```

AWS  
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.





# Lambda Function

## Parsing Security Events

- Send SMS message to administrator when a user terminates an instance tagger as a production instance
- We could also send a message to a Slack Channel, or Pager Duty, or send an email
- Got notified in 30 seconds!

```
import base64
import json
import boto3

ec2_resource = boto3.resource('ec2')
sns_client = boto3.client('sns')

def handler(event, context):
    for record in event['Records']:
        # kinesis data is base64 encoded so decode here
        payload = json.loads(base64.b64decode(record['kinesis']['data']))
        print('Payload: {}'.format(payload))
        # process event
        user_name = payload['detail']['userIdentity']['userName']
        instance_id = payload['detail']['requestParameters']['instancesSet']['items'][0]['instanceId']
        ec2_instance = ec2_resource.Instance(instance_id)
        for tags in ec2_instance.tags:
            if tags['Key'] == 'Type' and tags['Value'] == 'Production':
                print('Production instance was terminated, notifying!')
                # a production instance was terminated, let's send a notification
                sns_client.publish(
                    PhoneNumber='+15552562562',
                    Message='User {} has terminated production instance {}'.format(user_name, instance_id))
```

Imports

Boto3 clients

Parse Kinesis Record

Extract Username and InstanceId

Check if instance is Production

Send SMS!

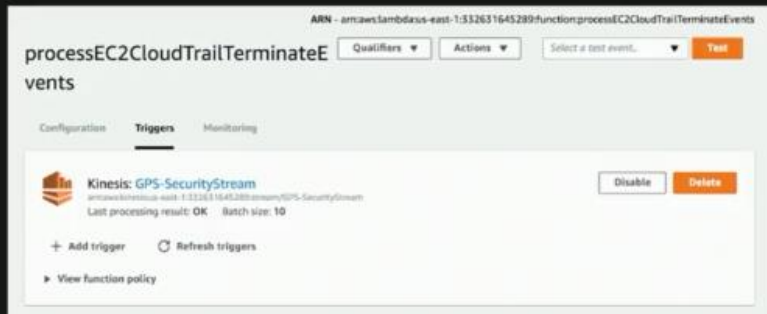
AWS  
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

# Executing the Lambda Function

## Parsing Security Events

- Configure a Kinesis trigger for our Lambda function
- Lambda service handles the Kinesis ingestion and execution of the Lambda function
- Also handles errors and retries



- Finally, we get our text message when an instance is terminated

Huzzah!

Today, 11:44 AM  
User jcv-iam has terminated production instance i-00196a3b35c509b33

AWS  
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



# Deeper Look into Amazon Kinesis




**Amazon Kinesis offering:** Managed services for streaming data ingestion and processing.

- **Amazon Kinesis Streams:** Build applications that process or analyze streaming data.
- **Amazon Kinesis Firehose:** Load massive volumes of streaming data into Amazon S3, Amazon Redshift, and Elasticsearch.
- **Amazon Kinesis Analytics:** Analyze data streams using SQL queries.

**Easy to use:** Focus on quickly launching data streaming applications instead of managing infrastructure.

**Real-time:** Collect real-time data streams and promptly respond to key business events and operational triggers. Real-time latencies.



AWS re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



Lambda functions are synchronously invoked, meaning that there is a guarantee of message order processing. This is why there is the concept of retries to enable the stoppage of the real-time processing of your pipeline messages before they fall off the Trim Horizon. You can fix the issue with retries before messages start getting dropped at the trim horizon, then your processing will continue as expected.

## Creating a Kinesis Stream

### Streams

- Made up of shards
- Each shard ingests/reads data up to 1 MB/sec
- Each shard emits/writes data up to 2 MB/sec
- Each shard supports 5 read transactions/sec

### Data

- All data is stored and is replayable for 24 hours (default)
- Retention window can be configured up to 7 days
- Partition key used to distribute PUTs across shards
- Even partition key distribution optimizes throughput

### Best practice

Determine an initial size/shards to plan for expected maximum demand

- ✓ Leverage "Help me decide how many shards I need" option in console
- ✓ Use formula for number of shards:

$$\max(\text{incoming\_write\_bandwidth\_in\_KB}/1000, \text{outgoing\_read\_bandwidth\_in\_KB} / 2000)$$

The screenshot shows the AWS Kinesis console configuration for a stream named 'demo-Kinesis'. The 'Stream Name' field is filled with 'demo-Kinesis'. Below it, there is a checkbox labeled 'Help me decide how many shards I need' which is checked. The 'Number of Shards' field is set to '1' and is circled in red. Below this, a table shows values calculated based on the number of shards entered above:

	Read:	Write:
Total Stream Capacity:	2 MB/s	1 MB/s
Max Transactions/second:	5	1000

AWS re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.





# Creating Lambda Function

## Memory

CPU allocation proportional to the memory configured  
Increasing memory makes your code execute faster (if CPU bound)

Increasing memory allows for larger record sizes processed

## Timeout

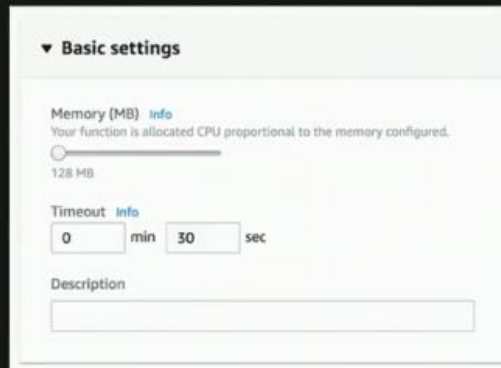
Increasing timeout allows for longer functions, but longer wait in case of errors

## Permission model

Execution role defined for Lambda must have permission to access the stream

## Retries

With Amazon Kinesis, Lambda retries until the data expires (24 hours)



▼ Basic settings

Memory (MB) [Info](#)  
Your function is allocated CPU proportional to the memory configured.  
128 MB

Timeout [Info](#)  
0 min 30 sec

Description

## Best practice

Write Lambda function code to be stateless  
Instantiate AWS clients and database clients **outside** the scope of the function handler to take advantage of connection re-use.



AWS  
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

# Configuring Event Trigger

Amazon Kinesis mapped as event source in Lambda

## Batch size

Max number of records that Lambda will send to one invocation

Not equivalent to effective batch size

Effective batch size is every 250 ms—calculated as:

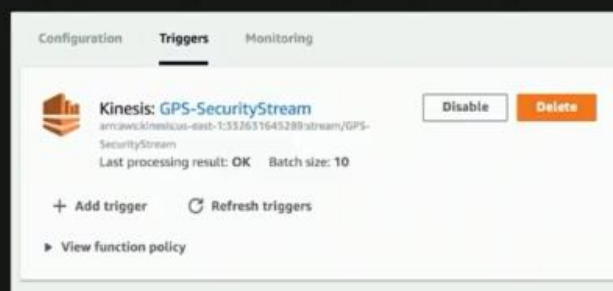
**MIN(records available, batch size, 6MB)**

Increasing batch size allows fewer Lambda function invocations with more data processed per function


## Best Practices

Set to "Trim Horizon" for reading from start of stream (all data)

Set to "Latest" for reading most recent data (LIFO) (latest data)



Configuration **Triggers** Monitoring

 **Kinesis: GPS-SecurityStream** [Disable](#) [Delete](#)  
arn:aws:kinesis:us-east-1:532651643289:stream/GPS-SecurityStream  
Last processing result: OK Batch size: 10

+ Add trigger Refresh triggers

► View function policy

AWS  
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

# How It Works

## Polling


Concurrent polling and processing **per shard**  
Lambda polls every 250 ms if no records found  
Will grab as much data as possible in one **GetRecords** call (Batch)

## Batching

Batches are passed for invocation to Lambda through function parameters  
Batch size may impact duration if the Lambda function takes longer to process more records  
Sub batch in memory for **invocation** payload

## Synchronous invocation

Batches invoked as **synchronous** RequestResponse type  
Lambda honors Amazon Kinesis **at least once** semantics  
Each shard **blocks** in **order** of synchronous invocation




**AWS re:Invent**

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



# Tuning Throughput



If put/ingestion rate is **greater** than the theoretical throughput, your processing is at risk of falling behind

## Maximum theoretical throughput

$\# \text{ shards} * 2 \text{ MB} / \text{Lambda function duration (s)}$

## Effective theoretical throughput


$\# \text{ shards} * \text{batch size (MB)} / \text{Lambda function duration (s)}$

**AWS re:Invent**

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



# Tuning Throughput with Retries



## Retries

Will retry on execution failures until the record is expired

Throttles and errors impact duration and directly impact throughput

## Best Practice

Retry with exponential back-off of up to 60 s

## Effective theoretical throughput with retries

$$(\text{\# shards} * \text{batch size (MB)}) / (\text{function duration (s)} * \text{retries until expiry})$$

**re:Invent**

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



# Monitoring Your Real-Time Pipeline

## Monitoring Amazon Kinesis Streams


- GetRecords: (effective throughput)
- PutRecord: bytes, latency, records, etc.
- GetRecords.IteratorAgeMilliseconds: how old your last processed records were

## Monitoring Lambda functions

- Invocation count: time function invoked
- Duration: execution/processing time
- Error count: number of errors
- Throttle count: number of time function throttled
- Iterator age: time elapsed from batch received and final record written to stream

## Debugging

- Review all metrics
- Make custom logs
- View RAM consumed
- Search for log events



**Log Streams**

2015/05/21/2184a6f4cb3e4c81a7265418e0c4078c

2015/05/21/3d07c7e60fc04acab257ee789bb667ef

2015/05/21/898f9b697db84b61a19879477d197679

**re:Invent**

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.





# Example Debugging with X-Ray



AWS X-Ray

## Retries with Lambda/Exponential Back-offs



- Total Duration = 6 seconds (five retries before success)
  - Lambda Invocation Duration = 2.2 seconds
  - Actual Throughput =  $1 * (10000/6) = 1666$  records/s
  - Ideal Throughput =  $1 * (10000/2.2) = 4545$  records/s
- 2.72x difference


AWS  
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



# Falling Behind

## Not Processing Fast Enough



- Producers are inserting 100k records per minute
- 2 shards
- Lambda functions take 1.7 seconds to process 1000 records
- $2 * (1000/1.7) = 1176$  records/s = ~70k/minute

100k produced per minute > ~70k processed per minute


AWS  
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



# Resharding to Increase Throughput

## Increase Number of Shards to Catch Up



- Increased number of shards to 4 (doubled)
- Lambda functions take 1.7 seconds to process 1000 records
- $4 * (1000/1.7) = 2352 \text{ records/s} = \sim 141\text{k} / \text{minute}$

100k produced per minute < ~141k processed per minute  
REAL TIME!

AWS  
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



## Netflix and AWS

“ Amazon Kinesis Streams processes multiple terabytes of log data each day, yet events show up in our analytics in seconds. We can discover and respond to issues in real time, ensuring high availability and a great customer experience. ”

—John Bennett  
Senior Software Engineer, Netflix

# NETFLIX

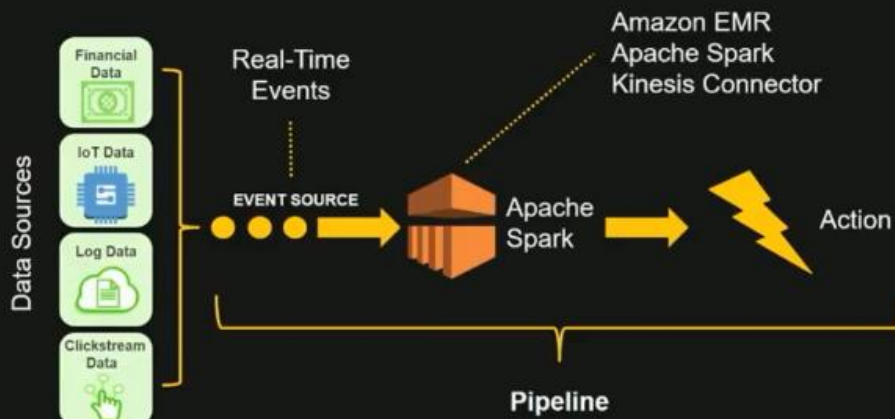
- Netflix needed a solution for ingesting, augmenting, and analyzing the multiple terabytes of data its network generates daily in the form of virtual private cloud (VPC) flow logs.
- Netflix uses AWS to analyze billions of messages across more than 100,000 application instances daily in real time, enabling it to optimize user experience, reduce costs, and improve application resilience.

AWS  
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



# What about EMR?



AWS  
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



## Why EMR + Apache Spark?

- Mature features and libraries
- Machine learning functionality with SparkML
- Interact with Data using SparkSQL
- Spark Streaming provides support for operators states such as sliding windows
- Fault tolerance/recovery out of the box
- Ability to combine batch and streaming
- Multiple languages (Scala, Java, Python, R, more)

### SparkSQL (Python)

```
context = HiveContext(sc)
results = context.sql(
    "SELECT * FROM people")
names = results.map(lambda p: p.name)
```

### SparkML (Python)

```
data = spark.read.format("libsvm")\
    .load("hdfs://...")

model = KMeans(k=10).fit(data)
```

### GraphX (Scala)

```
graph = Graph(vertices, edges)
messages = spark.textFile("hdfs://...")
graph2 = graph.joinVertices(messages) {
    (id, vertex, msg) => ...
}
```

AWS  
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.





# Zillow and AWS

“ We can compute Zestimates in seconds, as opposed to hours, by using Amazon Kinesis Streams and Spark on Amazon EMR. ”

—Jasjeet Thind  
Vice President of Data Science and Engineering



**Zillow®**

Zillow provides online home information to tens of millions of buyers and sellers every day.

- Needed to provide timely home valuations for all new homes
- Runs Zestimate, its machine learning-based home-valuation tool, on AWS
- Performs machine-learning jobs in hours instead of a day
- Gives customers more accurate data on more than 100 million homes
- Scales storage and compute capacity on demand

**AWS re:Invent**

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



## Next Steps

- ✓ Learn more about **AWS Serverless** at <https://aws.amazon.com/serverless>
- ✓ Explore the **AWS Lambda Reference Architecture** on GitHub:
  - Real-Time Streaming: <https://github.com/aws-labs/lambdarefarch-streamprocessing>
  - Distributed Computing Reference Architecture (serverless MapReduce) <https://github.com/aws-labs/lambdarefarch-mapreduce>

**AWS re:Invent**

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



## Next Steps

- ✓ Create an **Amazon Kinesis** stream. Visit the [Amazon Kinesis Console](#) and configure a stream to receive data Ex. data from Social media feeds.
- ✓ Create and test a **Lambda** function to process streams from Amazon Kinesis by visiting [Lambda console](#). First 1M requests each month are on us!
- ✓ Read the **Developer Guide** and try the Lambda and Amazon Kinesis Tutorial:
  - <http://docs.aws.amazon.com/lambda/latest/dg/with-kinesis.html>
- ✓ Send questions, comments, feedback to the [AWS Lambda Forums](#)

**AWS**  
**re:Invent**

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



**AWS**  
**re:Invent**

Thank you!

**AWS**  
**re:Invent**

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

