

With the increasing use of genomic sequencing for scientific discovery, the rate-limiting step for researchers is not in obtaining genetic code, but in having the capacity for storage and computing power to analyze it. In this session, you learn how Eagle Genomics built a cloud platform that uses an open-source workflow engine (eHive), Docker containers to process jobs, and a REST service to manage pipeline runs, all to help customers process genetic sequences up to 20 times faster without additional costs. You also learn how Eagle used Amazon EC2 Spot instances to provide low-cost compute power, and services such as Amazon S3 to power their cost-effective and scalable genetic processing platform.

The Big Data problem & Unilever case study

About Eagle Genomics

Founded in 2008
Based in Cambridge, UK, on the Wellcome Genome
Campus

Smart data management for Life Sciences—software and services

- Human and animal health
- Personal care and cosmeceuticals
- Food and nutriceuticals

Blue-chip customers across US and Europe

One of the first companies ever to run bioinformatics pipelines on AWS cloud, in 2008

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved

The challenge







Key requirements:

- Trustworthy and validated evidence
- Cater for future and unforeseen scientific questions
- Manage data as assets, quantify their quality and value
- No single data source can answer complex scientific questions

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved

Evolution of data processing





Present: Server/Cluster



Emerging: Secure Scalable Cloud











C 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved

Empowering data biologists into the future

Today: Simple to deploy and scalable on cloud







amazon

"Always On", access from anywhere

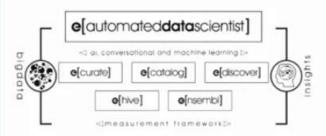
Effective collaboration across global teams

Empowering biologists to become data scientists

Streamlined workflows-"Click and Go"

Scalable storage and analysis power

Tomorrow: True data-driven innovation



Questions-driven, Al-enabled

Radically improved productivity of your scientist

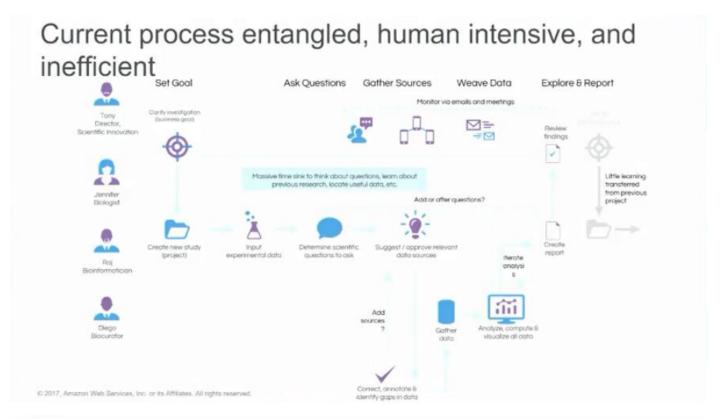
Transformational interaction between executive and scientist

Data governance by design

Data-driven innovation

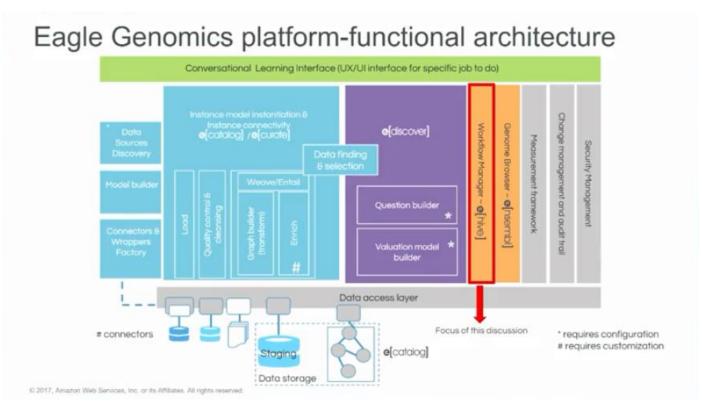
Presented by Unillover at July 2017 webinar

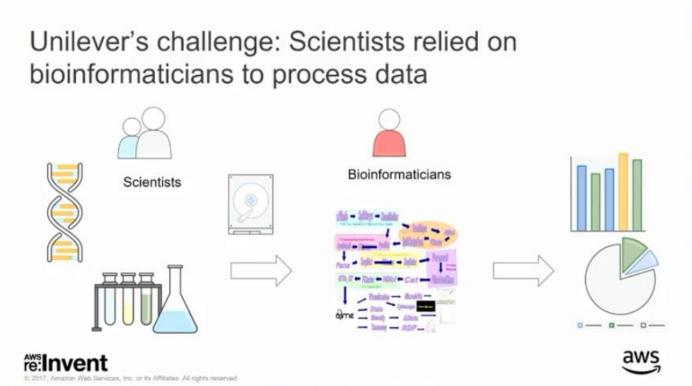
© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



The future—the automated data scientist—data science as a service Jennifer Refine / ask new questions Determine Conduct Ask Analyze all study report source report experiment questions studies data -1eL 62 mls AA Jennifer Biologist data Roj Tony Director, Bioinformatician

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved





We need to industrialized the middle part and reduce the number of bioinformaticians needed and make it self service for the scientists to use on their own with little help

Unilever's requirements

Background

New biology and informatics program promotes access to public data

Underlying architecture must keep pace with expanding scientific discoveries

Simple but robust solution combines Amazon Elastic Compute Cloud (Amazon EC2) and Amazon Simple Storage Service (Amazon S3) with the open-source workflow system, eHive

Results

Ten times as many scientists can process studies simultaneously, compared to noncloud architecture

Genetic sequence processing is twenty times faster, without increasing compute costs

Confidence that the AWS-based deployment helps Unilever's scientists create marketleading innovations



Inc. or its Affliates. All rights reserved



Successful outcome: Reference scientific project

First-of-a-kind project providing product competitiveness using microbiome data; Eagle's Microbiome Platform accelerated the project timeline across global teams

SCIENTIFIC REPORTS

OPEN A randomised clinical study to determine the effect of a toothpaste containing enzymes and proteins on plaque oral microbiome ecology

S. E. Adams', D. Smold', B. Wurphy', F. Carroll', A. S. T. Charl', R. E. Marrott', B. H. G. Brading'

"Unilever's digital data program now processes genetic sequences twenty times faster - without incurring higher compute costs. In addition, its robust architecture supports ten times as many scientists, all working simultaneously."

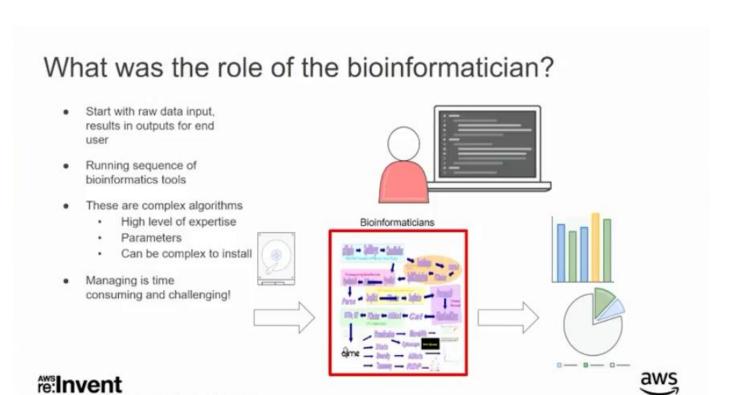
-Pete Keeley, eScience Technical Lead, R&D IT at Unilever

to 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

Background concepts

Workflows and workflow engines

Introducing the need for workflows and workflow engines in biology



These steps can be formalised as a workflow Define complex data-flow events Represent each step programmatically Enables reproducibility Process for a specific purpose Formalised workflow Workflow configuration file

The Workflow configuration file allows us to represent the formalized workflow in a programmatic and repeatable way using automation via workflow engines.

Workflow engines

- · Need a way of running the workflow configuration file
- Lots of workflow engines Azkaban CloudSlang GloudSlang GloudSlang Services

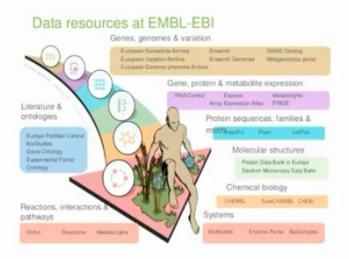
 RUNDECK WORKFLOWENGINE AIRFlow
- Key requirements:
 - Scalable
 - Reproducible
 - Traceable
 - Robust (e.g., handle failures)
 - Need to handle very complicated workflows



0.2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



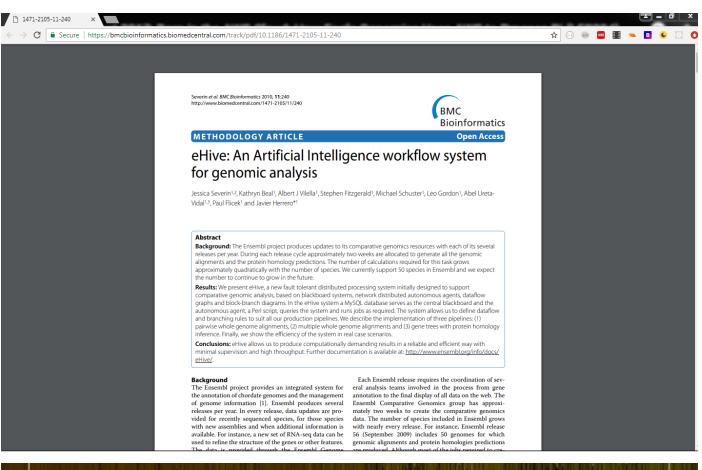
eHive has a track record in genomics workflows



- Since 2004 at the European Bioinformatics Institute (FBI)
- · Proven scalability on big data sets:
 - Approximately 450 CPU years of compute every year
 - Some pipelines reach > 6,000 of concurrent workers and manage millions of jobs
- Tried and tested with complex data sets and analysis







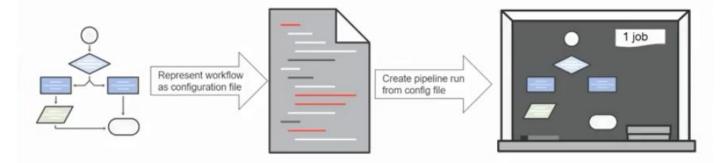
Background concepts

eHive key concepts

Key concepts about eHive that are essential for understanding dockerisation

eHive uses a blackboard system

- · Workflow analysis steps described by config file
- · During workflow run jobs are updated in a central blackboard database
- Represents flow structure and job list





aws

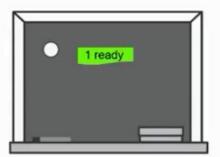
2017, Amazon Web Services, Inc. or its Attitates. All rights reserved.

Jobs are processed by autonomous workers



Analysis 'B'

s, Inc. or its Attitutes. All rights reserved





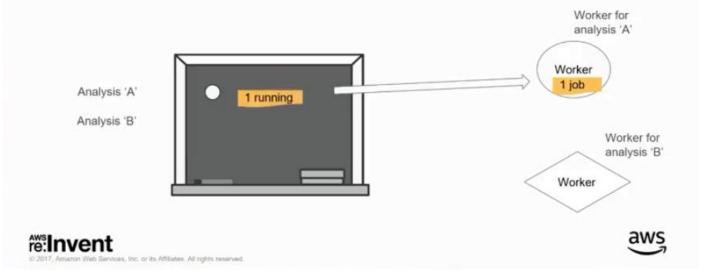




awş

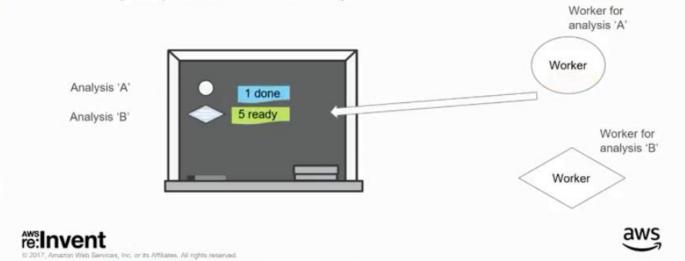
Jobs are processed by autonomous workers

- . The worker queries the blackboard for jobs of a type that it can run
- Workers process jobs that match the analysis type

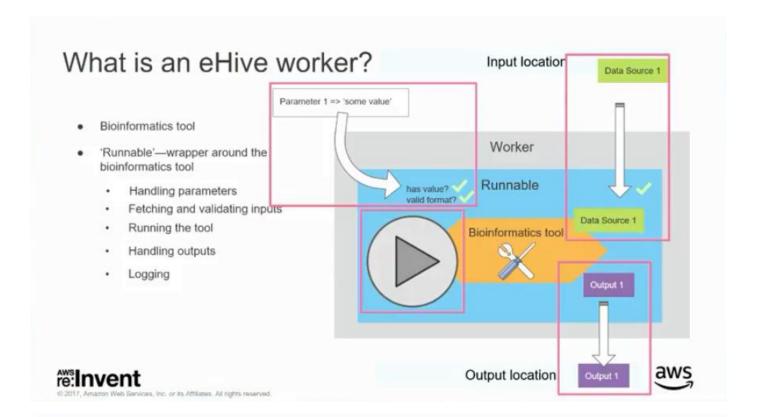


Jobs are processed by autonomous workers

- . The worker queries the blackboard for jobs of a type that it can run
- · Workers process jobs that match the analysis type
- When a job completes it can create 1 or more new jobs



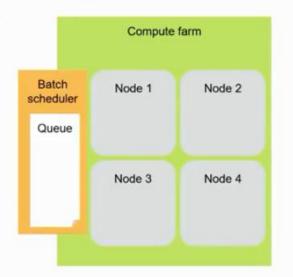
The job creation is a dynamic process in the Hive and this makes the process scalable



eHive runs workers using a batch scheduler

· eHive monitors jobs to see whether there are jobs to be done





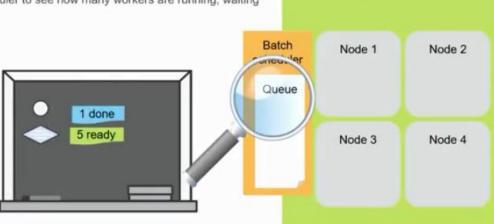




eHive runs workers using a batch scheduler

· eHive monitors jobs to see whether there are jobs to be done

· Queries job scheduler to see how many workers are running, waiting





C WITE Assessed Web Consisses for or its Afflictor All nobbs speeded



Compute farm

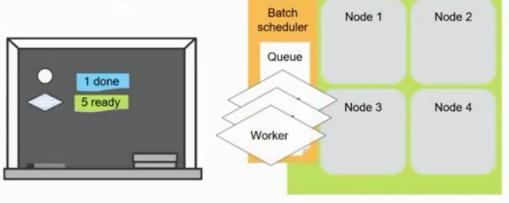
Compute farm

eHive runs workers using a batch scheduler

· eHive monitors jobs to see whether there are jobs to be done

Queries job scheduler to see how many workers are running, waiting

eHive places workers in the scheduler's queue



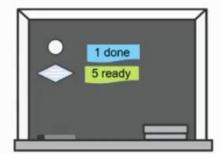


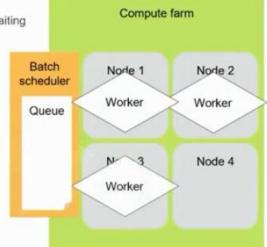


2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

eHive runs workers using a batch scheduler

- · eHive monitors jobs to see whether there are jobs to be done
- · Queries job scheduler to see how many workers are running, waiting
- eHive places workers in the scheduler's queue
- The scheduler finds available compute to run the worker





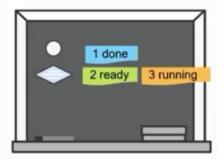


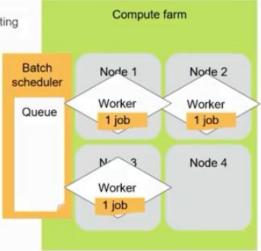
C 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



eHive runs workers using a batch scheduler

- · eHive monitors jobs to see whether there are jobs to be done
- Queries job scheduler to see how many workers are running, waiting
- · eHive places workers in the scheduler's queue
- The scheduler finds available compute to run the worker



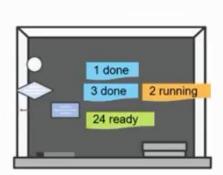


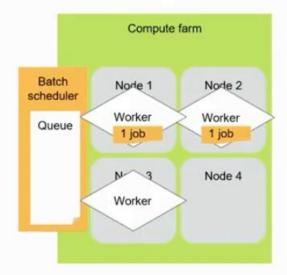




D 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

Any number and type of workers can run together



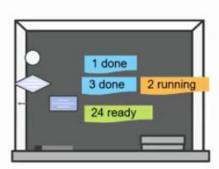


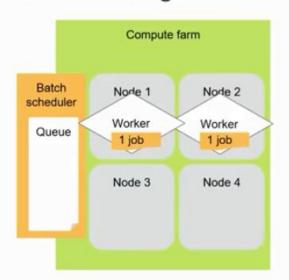


© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

Any number and type of workers can run together

Workers exit if they have no more jobs



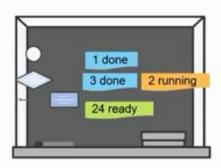


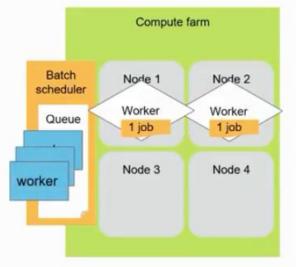




Any number and type of workers can run together

- · Workers exit if they have no more jobs
- More workers of the correct type are requested







© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

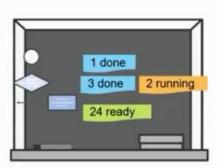


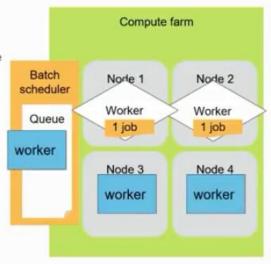
Any number and type of workers can run together

- Workers exit if they have no more jobs
- · More workers of the correct type are requested

c. or its Attiliates. All rights reserved

. Multiple workers from multiple analysis can run at the same time



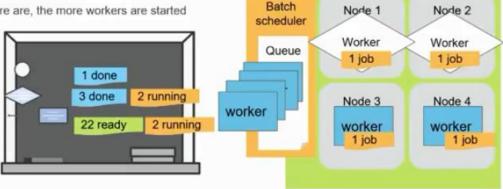






Any number and type of workers can run together

- Workers exit if they have no more jobs
- More workers of the correct type are requested
- Multiple workers from multiple analysis can run at the same time
- The more jobs there are, the more workers are started







Compute farm

Implementation approach From eHive to e[hive] Taking eHive to the AWS cloud...and back again

Making e[hive] a part of our platform



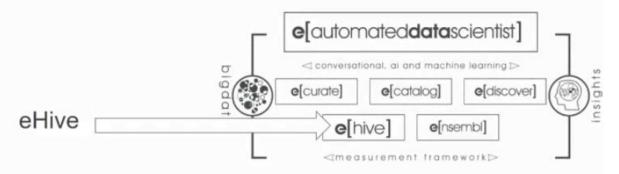
- Designed to work on site at EBI (LSF)
- Requires technical user

- Want to run in cloud environments
- Enable scientists to run pipeline without knowledge of e[hive]





Making e[hive] a part of our platform



- Designed to work on site at EBI (LSF)
- Requires technical user

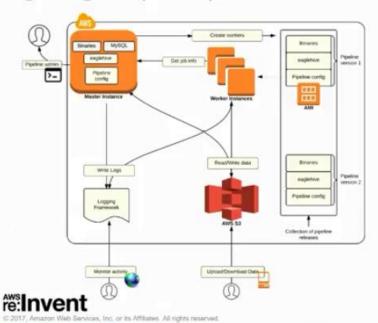
- · Want to run in cloud environments
- Enable scientists to run pipeline without knowledge of e[hive]
- Integrate with platform providing analytics capability



SHIVELL STATES



e[hive] v1 (2008)—Workers as AWS instances

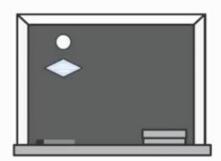


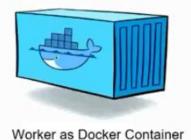
- · Each worker is an AWS instance
- Many production workflows implemented in AWS!
- Poor resource usage
 - Workflow steps specify instance type
 - Resource bottleneck
 - Same instance for multiple analysis
- · Whole workflow on AMI
 - Slows development
 - Difficult to test analysis
 - Poor reusability
- Automated dependency installation using Chef



e[hive] v2 (2015) - Workers as Docker containers?

- · Could e[hive] workers be made into a docker image?
- How can e[hive] manage scheduling workers across a cluster?





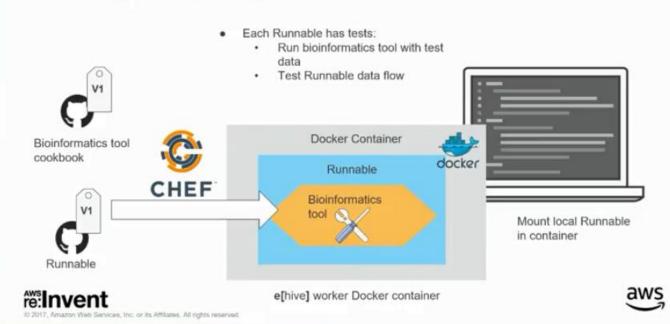


© 2017, Amazon Web Services, Inc. or its Affiliates, All rights reserved



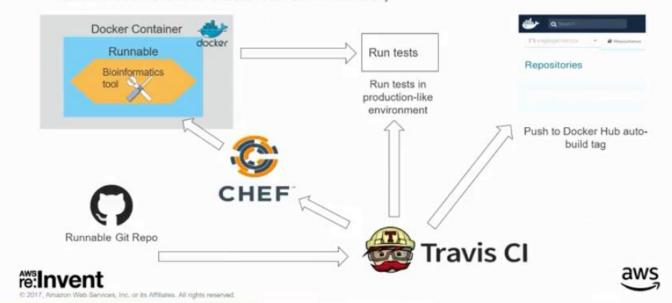


Building the e[hive] worker with Docker



e[hive] worker development cycle

· Runnable tests are run from CI tool on each commit and weekly

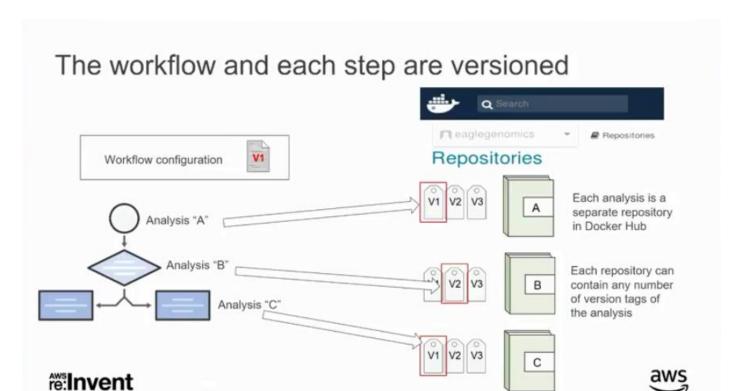


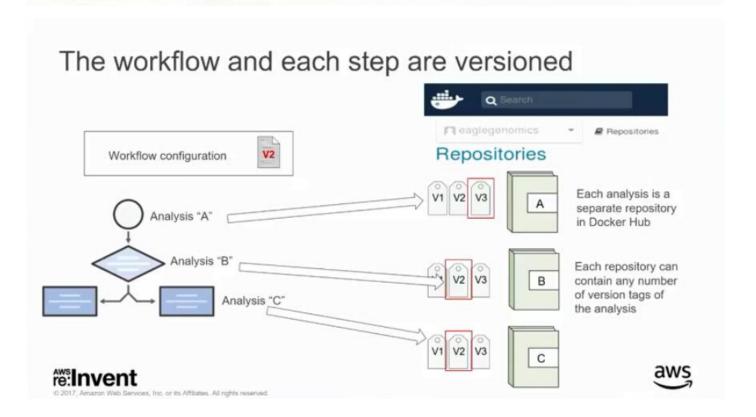
Some key questions about building Docker images

Question	Reason	Solution(s) Dockerfile, configuration management	
How was the image built? What is on the image?	We shouldn't trust something if we don't know how it is built		
Does the build still work now?	Or are we stuck with a 'golden image' we can't reproduce anymore?	Build regularly, e.g., via continuous integration tool	
Does the tool work? How do we test the image?	This is essential! Build may Continuous integration to tests install!		
Can we redeploy elsewhere? Or are we tied to Docker? Is this a problem?	In some cases we can't use Docker (e.g., on an HPC)	Need a suitable build method that takes this into account	









Versioning is critical

- Reproducibility in scientific workflows is essential
 - · Replicate other people's work
 - Necessary for quality control, compliance, etc.
 - Requires clear instructions
 - · Versioning at many levels is key for this





- Pipeline run report for end user
 - Workflow description
 - Each tool and version



2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



Workers log to a centralised logging server

- Workers log to Amazon Elasticsearch Service (Amazon ES) via FluentD—jobs execute on different hosts, but we see all the logs together
- Visualise and explore in Kibana
- Can quickly see any errors even with 100 1000s jobs

Job count/analysis step

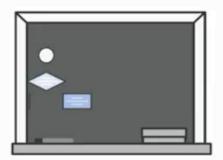


From eHive to e[hive]: Stage 2

Building the Docker Swarm Orchestrator

Creating an auto scaling container scheduler for e[hive]

Extending eHive's interfaces to schedulers



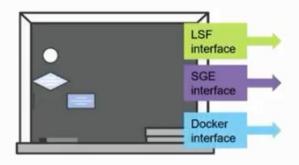


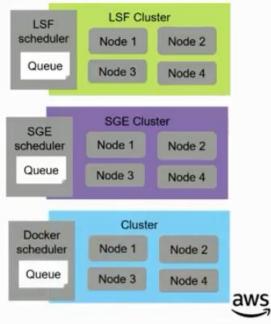


2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

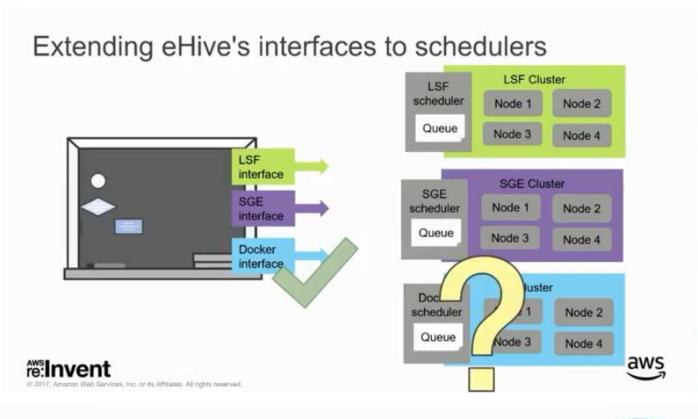


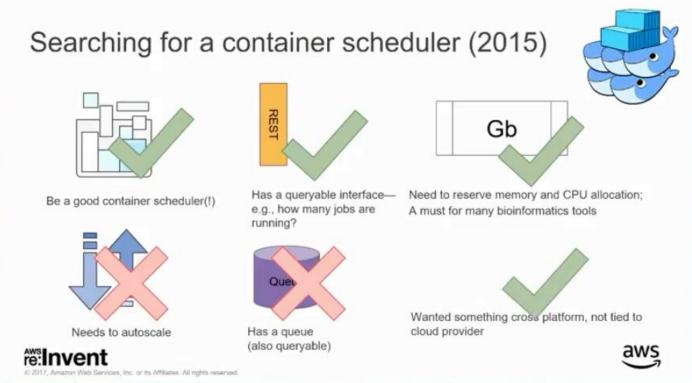
Extending eHive's interfaces to schedulers



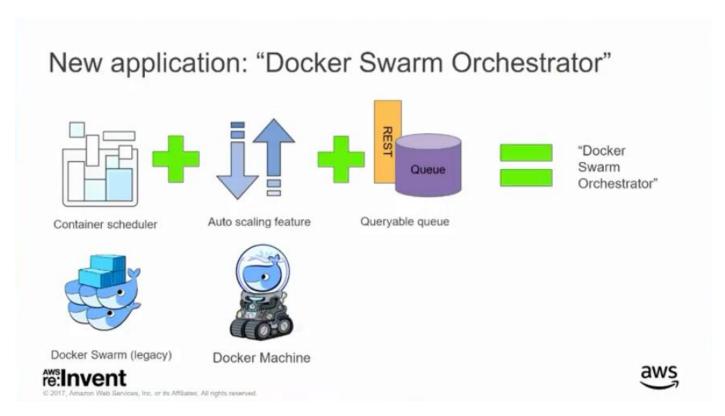


re:Invent

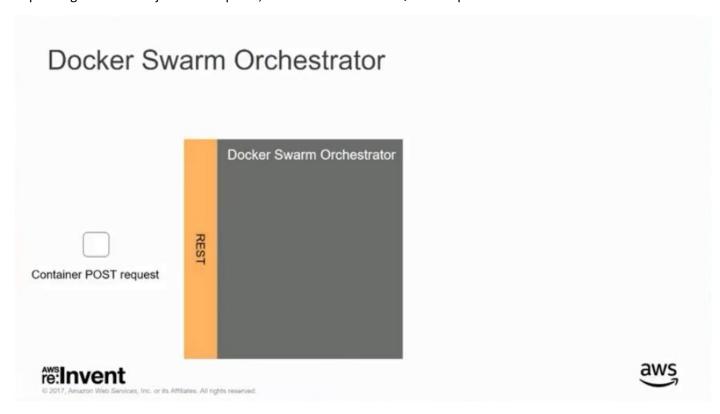




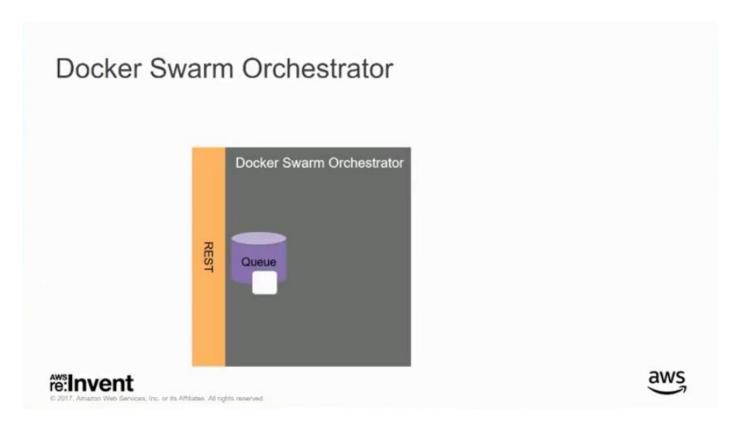
We need the above requirements fulfilled when eHive creates/schedules its jobs for the workers to pick up and run. This meant that we had to build our own tool that will have all features above.



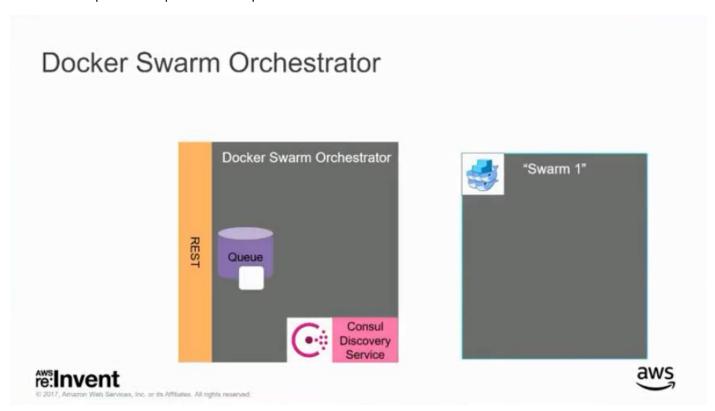
We used Docker Swarm as our orchestrator, Docker Machine API for autoscaling the swarm to create or delete instances depending on available jobs in the queue, and we used RabbitMQ for the queue.



When we make a container POST request to the REST interface of the Docker Swarm Orchestrator DSO,



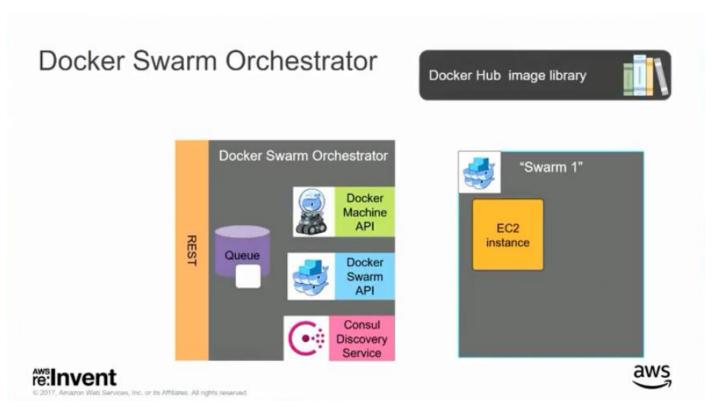
The DSO then puts the request into the queue



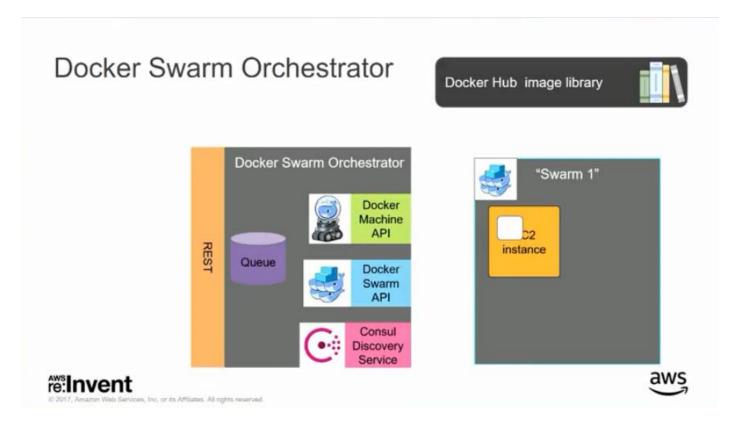
We then create a new Docker swarm that uses a discovery service like Consul to register this new swarm

Docker Swarm Orchestrator Docker Swarm Orchestrator Docker Machine API Consul Discovery Service Discovery Service 2017, Anazor Web Service, Inc. or #A Afflates. All rights seered.

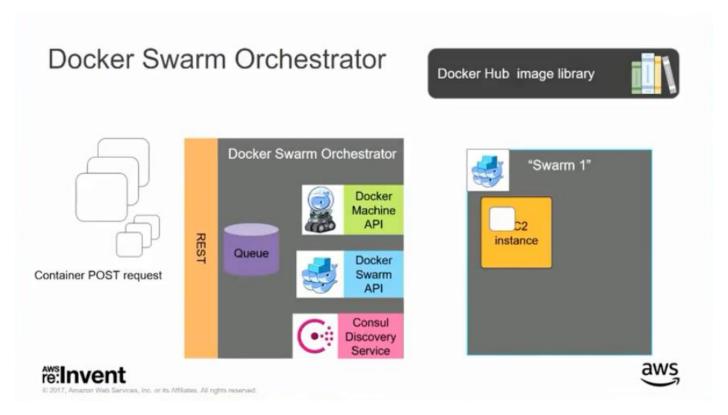
We then use the Docker Machine API to add instances into the swarm when we have jobs in the queue



The instance is now registered in Docker Swarm and Docker Swarm can then pull the needed image to run from Docker Hub



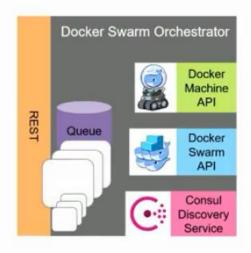
It then pulls the job details and starts running the job in the container

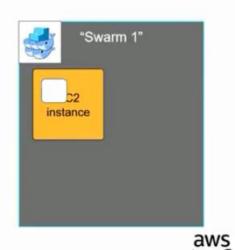


You can post any number of jobs into the queue with different memory requirements

Docker Swarm Orchestrator







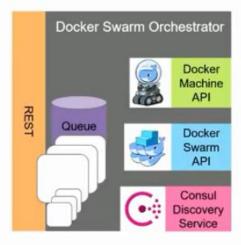
re:Invent

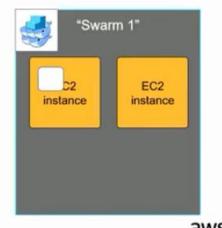
17017 Amoreon West Constraint Inc. on the Affiliation. All probles recognised

The DSO starts the containers using available resources as available









re:Invent

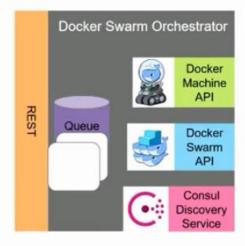
2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved

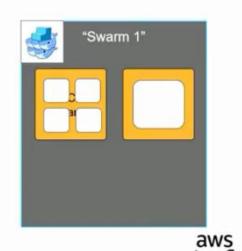
Docker Swarm Orchestrator



List GET request: /swarm/{swarmid}/container/count

- · 2 queued
- 5 running





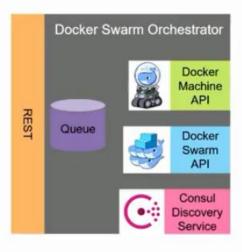


IS 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

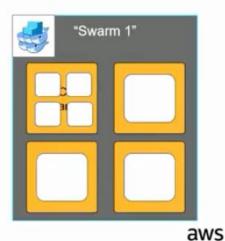
Docker Swarm Orchestrator

List GET request: /swarm/{swarmid}/container/count

7 running









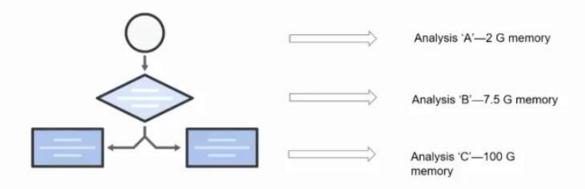
017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

The DSO then dynamically starts some more instances to run the jobs in the queue

Putting the e[hive] pieces together

How e[hive] runs a containerised workflow in the AWS cloud using the Docker Swarm Orchestrator

Workflow steps and resource requirements





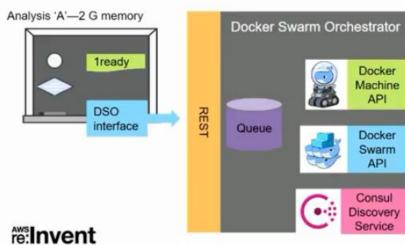




Assembling the e[hive] pieces

Docker Hub image library

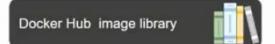


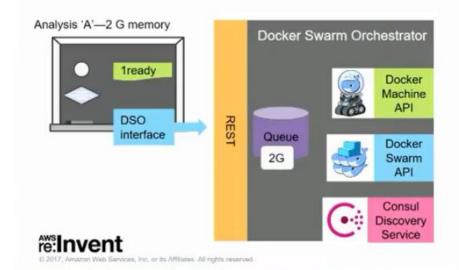




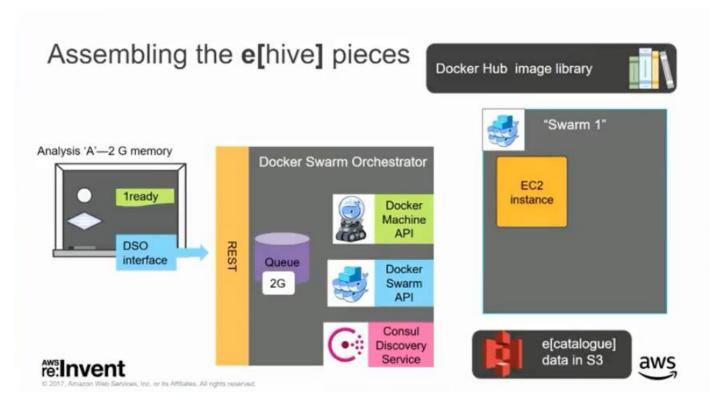


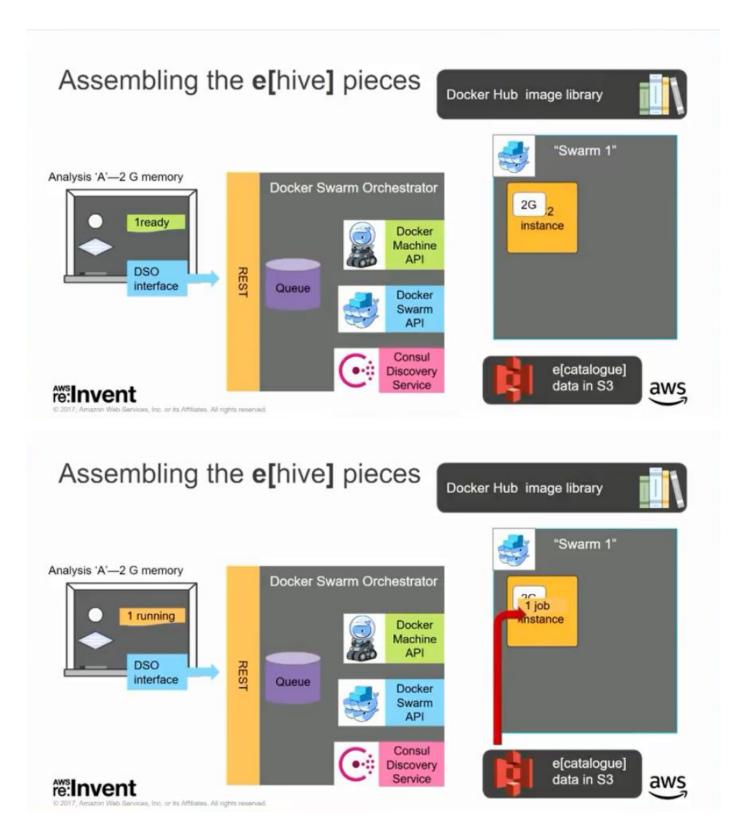
Assembling the e[hive] pieces



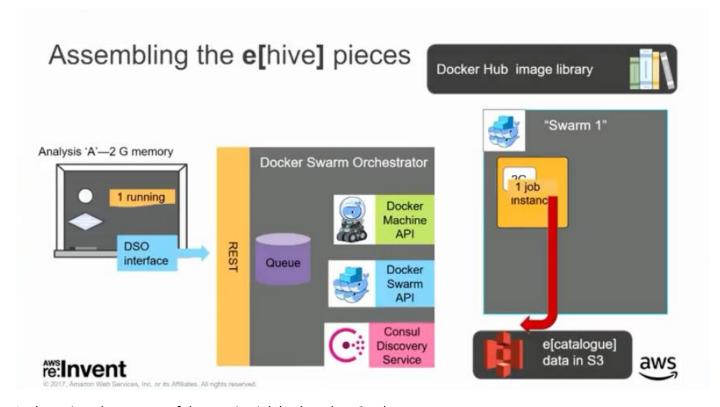




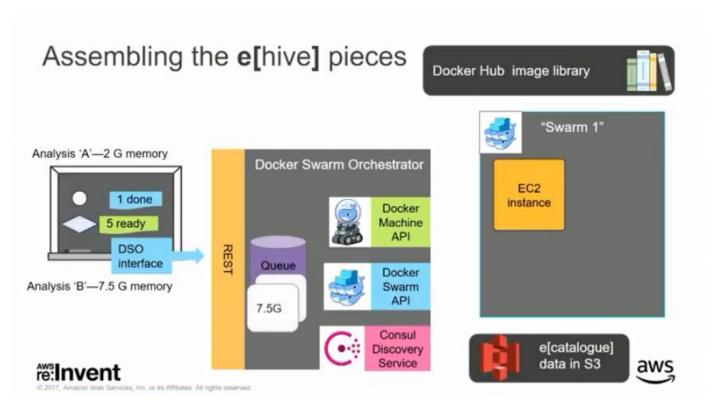


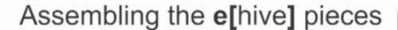


The running instance pulls the job details from the eCatalogue bucket in S3 and starts running the job

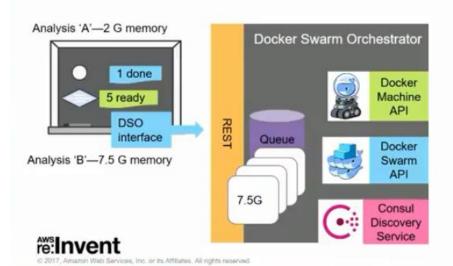


It also writes the outputs of the running job back to the eCatalogue.









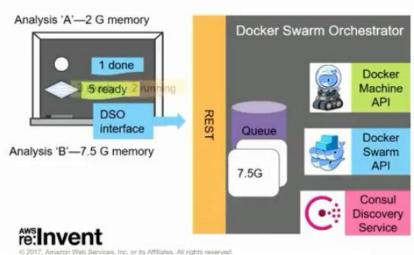






Assembling the e[hive] pieces

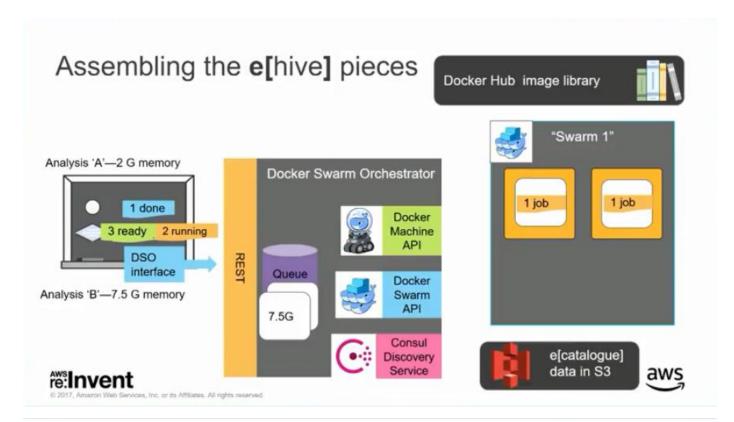


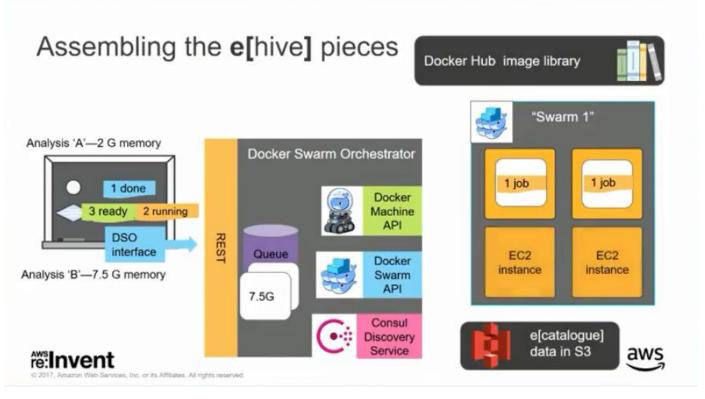


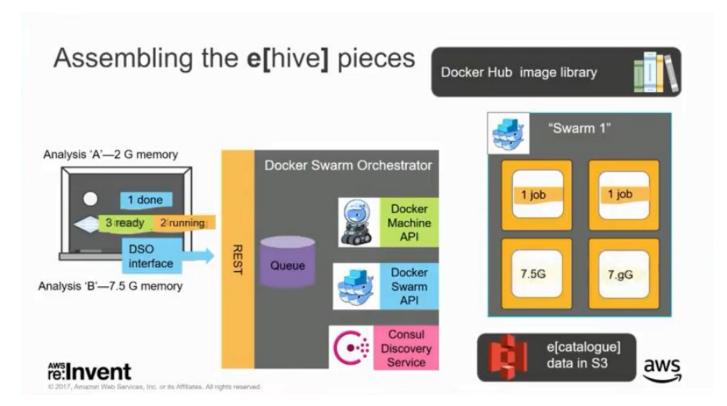


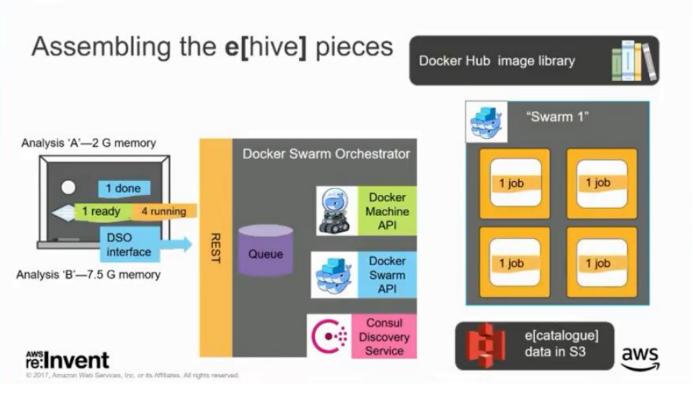




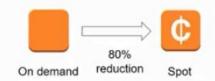


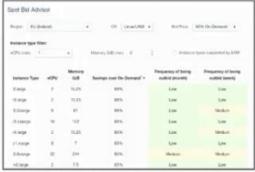




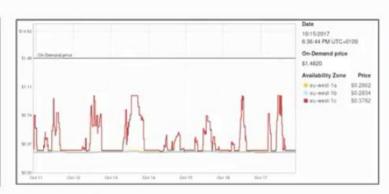


Use spot instances!





Instance types and reductions vary choose carefully



Availability Zone important



ID 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved



Spot instances and reliability for SLA



- Spot instances have risk
- Multi-layered fault tolerance mitigates against this
 - eHive and job/worker monitoring
 - · DSO and container queue
 - DSO and Docker Swarm







Putting e[hive] into action

Optimising resource usage with e[hive]

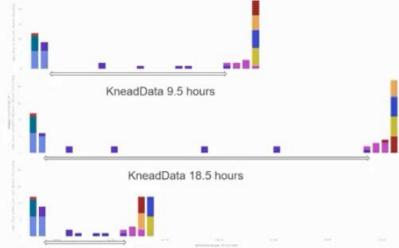
Using metrics to optimise analysis resource usage

Making use of the metrics: Optimising resource usage

7.5 G container, bowtie2 options: very sensitive

30G container, bowtie2 options: very sensitive

7.5 G container, bowtie2 options: very fast



re:Invent

ID 2017, Arsazon Web Services, Inc. or its Affiliates. All rights reserved



Time and cost of KneadData analysis

KneadData 5 hours

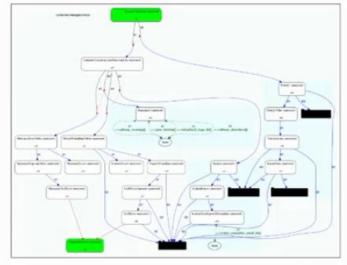
Options used	Average time per job (sample)	Cost per sample (on demand, r3.xlarge \$0.371 hourly)	Cost per 1000 sample (on demand, r3.xlarge \$0.371 hourly)	Cost per 1000 sample (Spot instances, r3.xlarge \$0.05 hourly)
7.5G container, bowtie2 options : very-sensitive	6.7 hr	2.4857	2485.7	335
30G container, bowtie2 options : very- sensitive	4.3 hr	1.5953	1595.3	215
7.5G container, bowtle2 options : very-fast	3.5 hr	1.2985	1298.5	175
KneadData with 30G memory container, bowtie2 options: 4 CPUs, very-sensitive	1.1 hr	0.4081	408.1	55
KneadData with 30G memory container, bowtie2 options: 4 CPUs, very-fast	50.2 min	0.308	308	41.5





Metagenomics workflow performance

- · Identify function and taxonomy for all sequences
- Input: ~1.7 billion read pairs (each read 125 base pairs)
- Took about 30 hours across ~50 instances (200 CPUs)
- Total cost < \$50

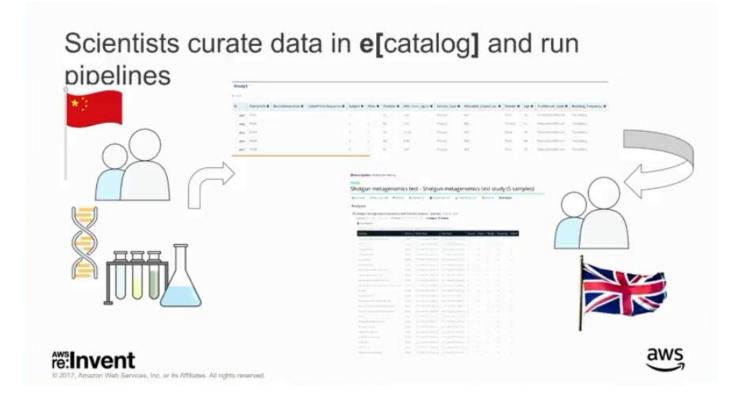




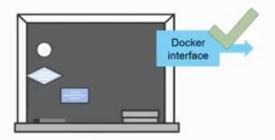
© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved

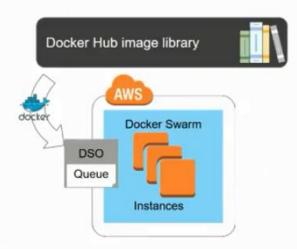


Conclusions and perspective



Bringing it back to the HPC

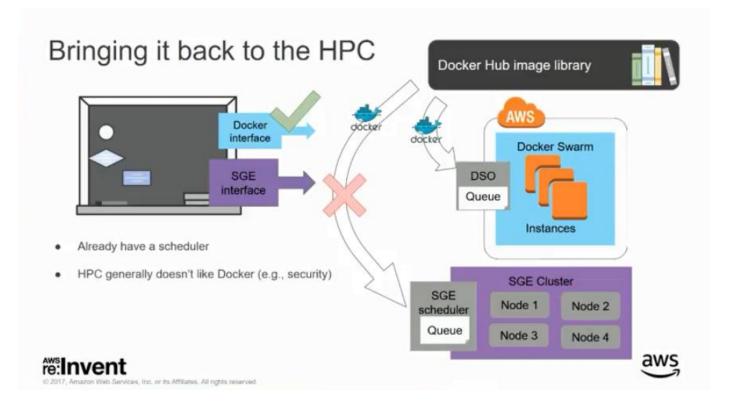


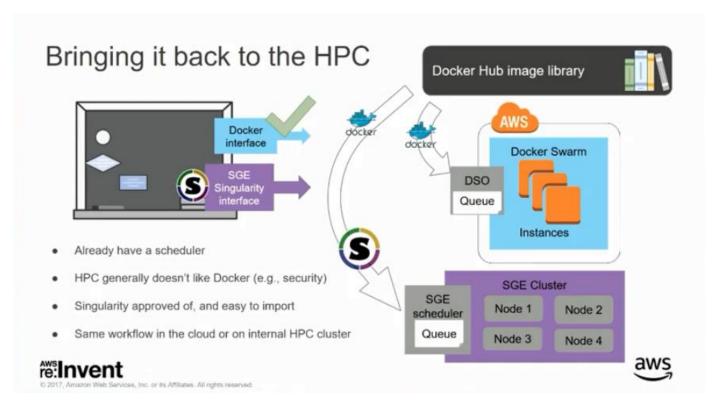




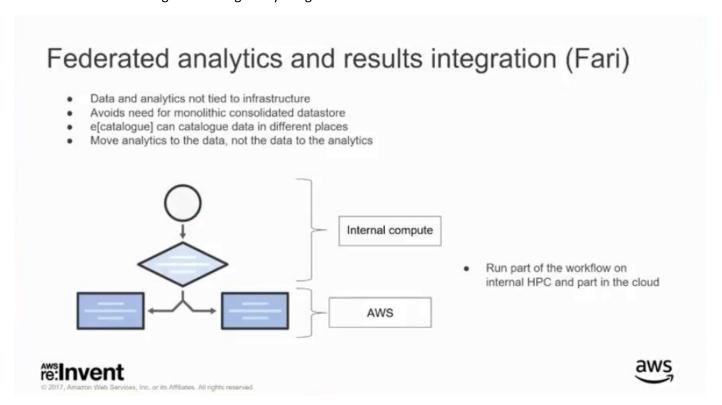
© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved



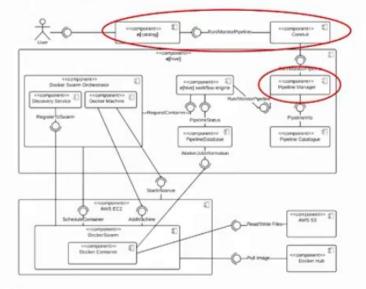




But we can now use another container type called Singularity that is well suited for within internal SGE clusters. It can translate our Docker image into a Singularity image and run that within the SGE cluster.



Putting it together: Enabling the end user



e[catalog] is separate application

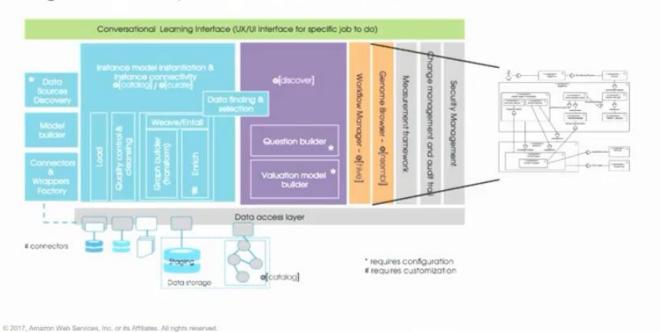
Implemented RESTful pipeline manager



ili 2017, Amazon Web Services, Inc. or its Affliates. All rights reserved



Eagle Genomics platform functional architecture



Unilever productivity & scale increased 20 times using this workflow



