aws SUMMIT TEL AVIV

Mastering Kubernetes on AWS

Arun Gupta, @arungupta
Principal Open Source Technologist, AWS

Amit Bezalel @amitbez
Chief System Architect, Microfocus
(github.com/amitbet)

Agenda

1. Kubernetes cluster setup
2. CI/CD with applications deployed on Kubernetes
3. Visibility

1 Kubernetes cluster set-up choices

Kubernetes cluster setup—choices

Install, operate, upgrade, delete a Kubernetes cluster

Development—Minikube

Community—Kops
- List: kubernetes-aws.io

Enterprise
- Elastic Container Service for Kubernetes (EKS)
- CoreOS Tectonic
- Red Hat OpenShift

Custom
- CloudFormation
- Terraform

AWS Partners: Docker, Heptio, Mesosphere

With EKS, AWS gives you a control plane and you bring your worker nodes and attach to the Master node you will get

Manage a Kubernetes cluster: kops

Community supported
- SIG AWS
- Kops office hours and Slack channel

Generate CloudFormation or Terraform scripts

github.com/kubernetes/kops

```
export AWS_AVAILABILITY_ZONES=${ZONES:-"us-east-1b,us-east-1c,us-east-1d"}
export KOPS_STATE_STORE="s3://kubernetes-aws-io"
kops create cluster cluster.k8s.local \
  --master-count 3 \
  --master-size m4.large \
  --node-count 5 \
  --node-size m4.large \
  --zones $AWS_AVAILABILITY_ZONES \
  --networking calico \
  --yes
```

Above is the command that you can use to create a K8s cluster using KOPS in 3 AZs, you specify the parameters above and the state of the cluster itself will be stored in a S3 bucket

Elastic Container Service for Kubernetes

Managed K8s control plane—highly available master and etcd

Bring your own worker nodes, like ECS

Core tenets

- Platform for enterprises to run production-grade workloads
- Provides a native and upstream Kubernetes experience
- Not forced to use additional AWS services, but offer seamless integration
- Actively contributes to the Kubernetes project

APIs

```
aws eks create-cluster \
  --cluster-name <> \
  --desired-master-version <> \
  --role-arn <>
```

You need to give the EKS K8s cluster an IAM role that it will use when creating those EC2 instances

EKS architecture

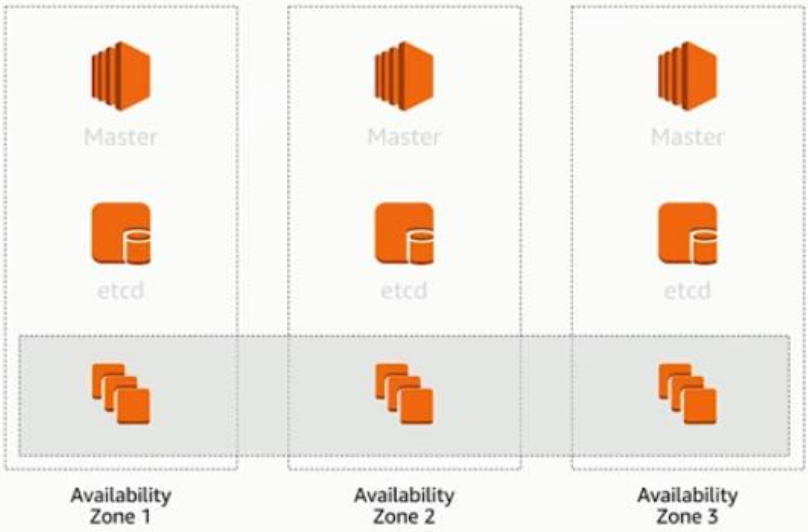

EKS abstracts away the control plane from you and instead gives you a control plane that includes an API server, Controller, Scheduler (all the key components of the K8s master), you juts need to bring your worker nodes

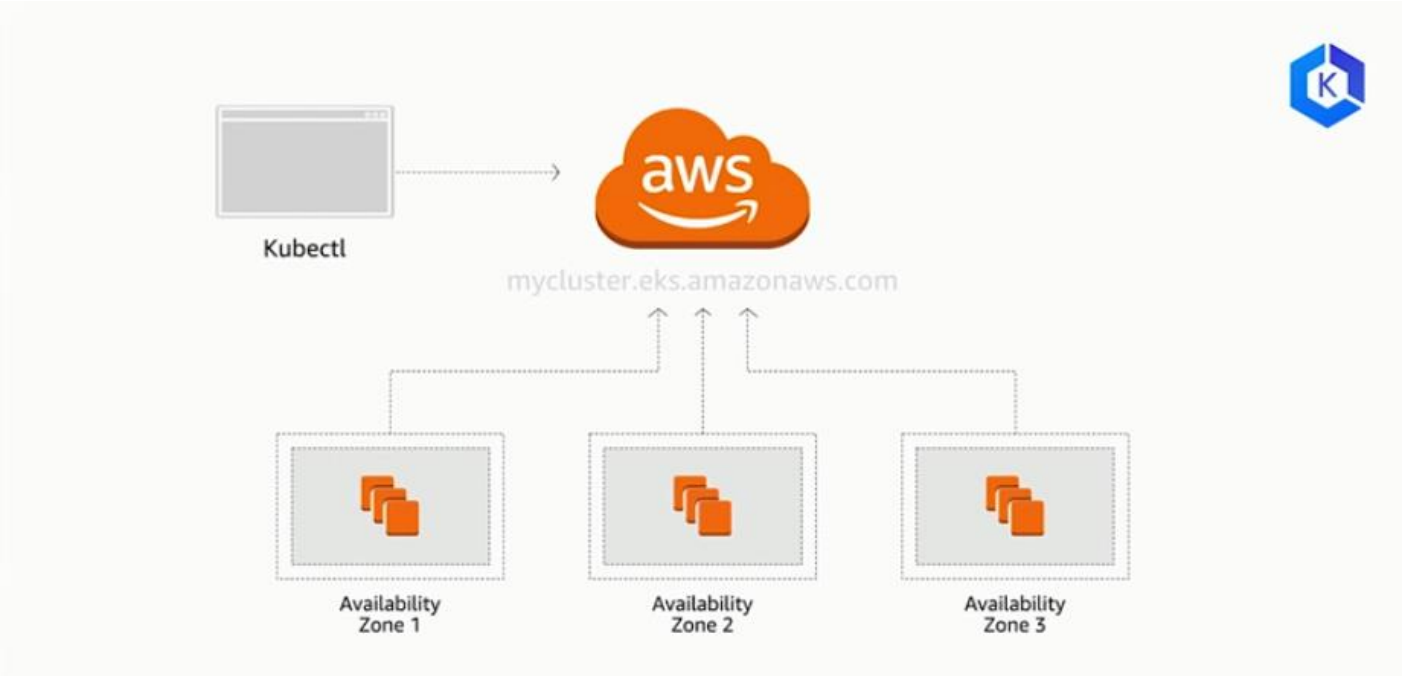

You then attach your worker nodes to the Master/control plane, then you can start using ***Kubectl*** API to interact and issue commands to the Master node for creating, managing and controlling your apps

Cluster setup at Microfocus

StormRunner Functional

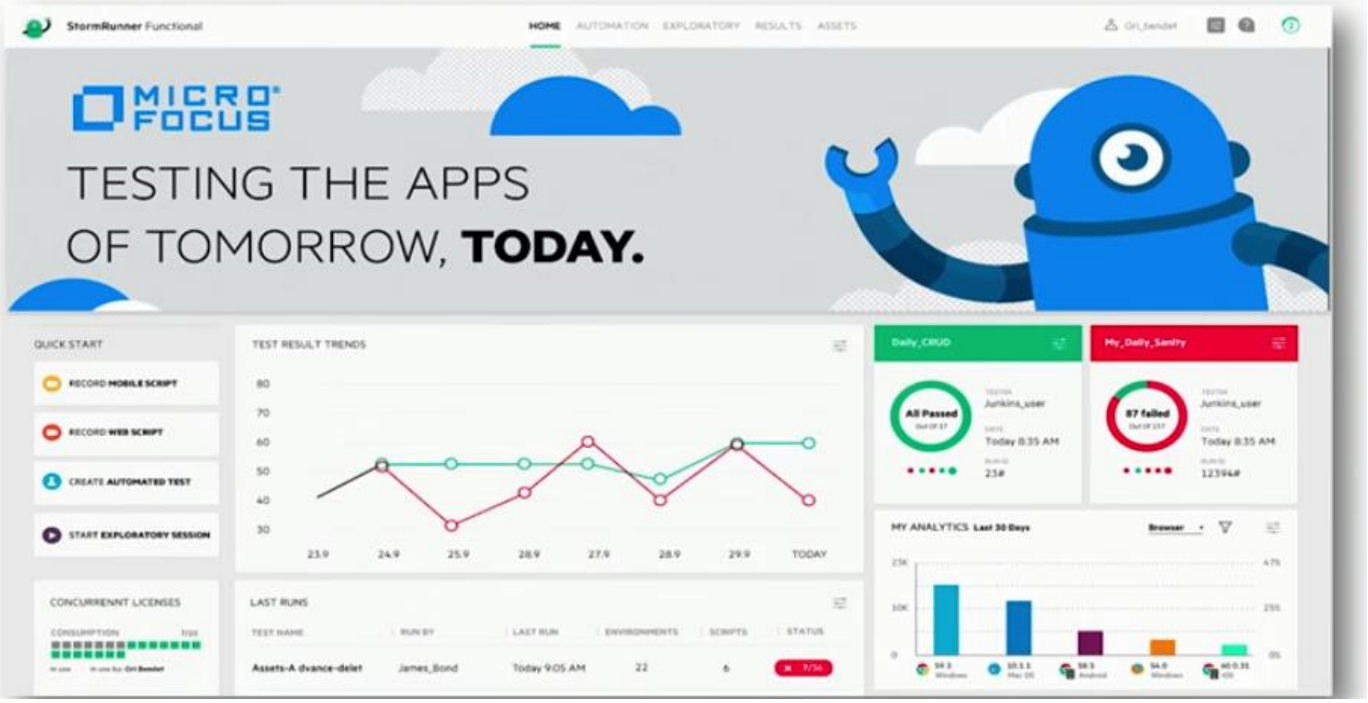

SRF High Level Solution

Analytics & Predictive
INTERNAL | PRODUCTION | INDUSTRY

Reporting (RCA)
STATUS | DEFECTS | TRENDS

Execution (Test & Scripts)
REMOTE | UPLOAD | RECORD

LAB Service (Cloud/On-prem & Hybrid)
WEB | MOBILE | THINGS*

Why Kubernetes?

VMs & AMIs	Containers	Kubernetes
Low density	Dense, Fast, Isolated, **but**:	
Spin up time	Connecting services ?	Service discovery
Large images	Several machines ?	Machine management
Unknown processes	Keep alive (recycle) ?	Replicate & Recycle
OS Updates	Scale out ?	Auto scales easily
	Logs / Monitoring ?	Integrated visibility

Kubernetes clusters

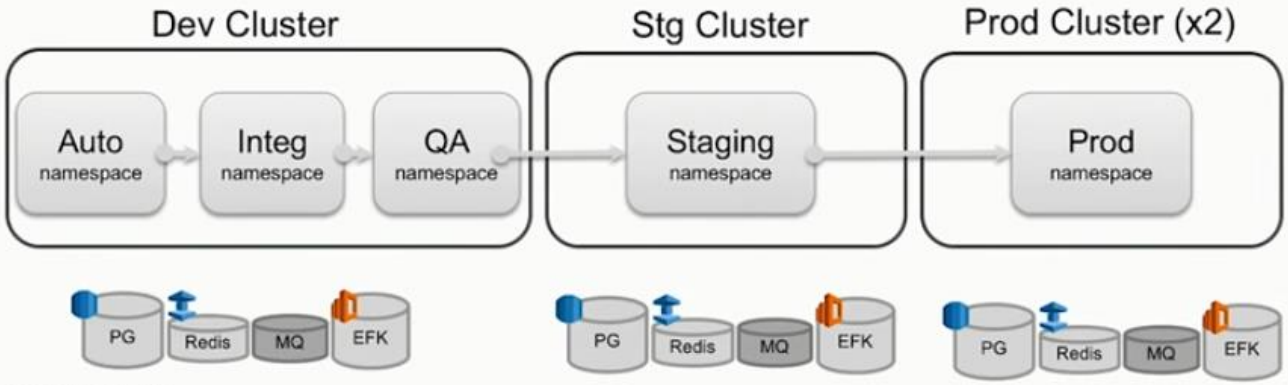

Namespaces:

- One machine cluster, multi envs
- Fast & Flexible
- Shared Infra
- Each has: DB schema / Redis prefix

Setup with kops

Kops to provision clusters

Kube-up is not flexible and is deprecated

Why kops

- **Customizable**:
  - Support for CNI Networks & k8s plugins
  - Accepts custom base images
  - Supports AWS, GCP
- **Popular & Maintained**:
  - Best of breed
  - Community supported
  - Supports newest k8s versions

SRF Application Services

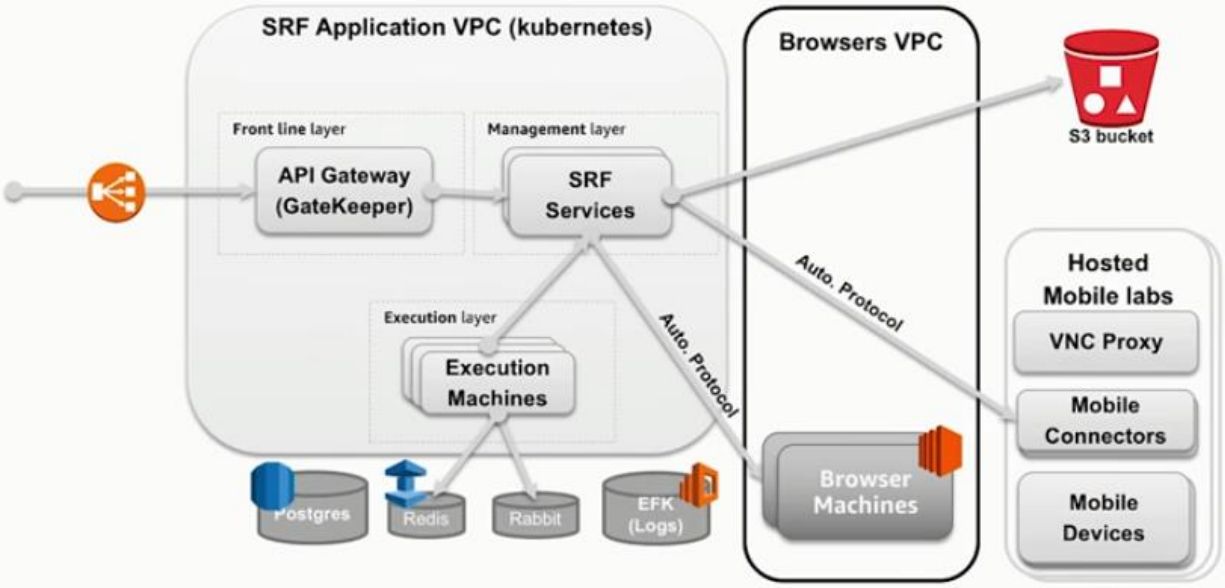

SRF Physical Cluster

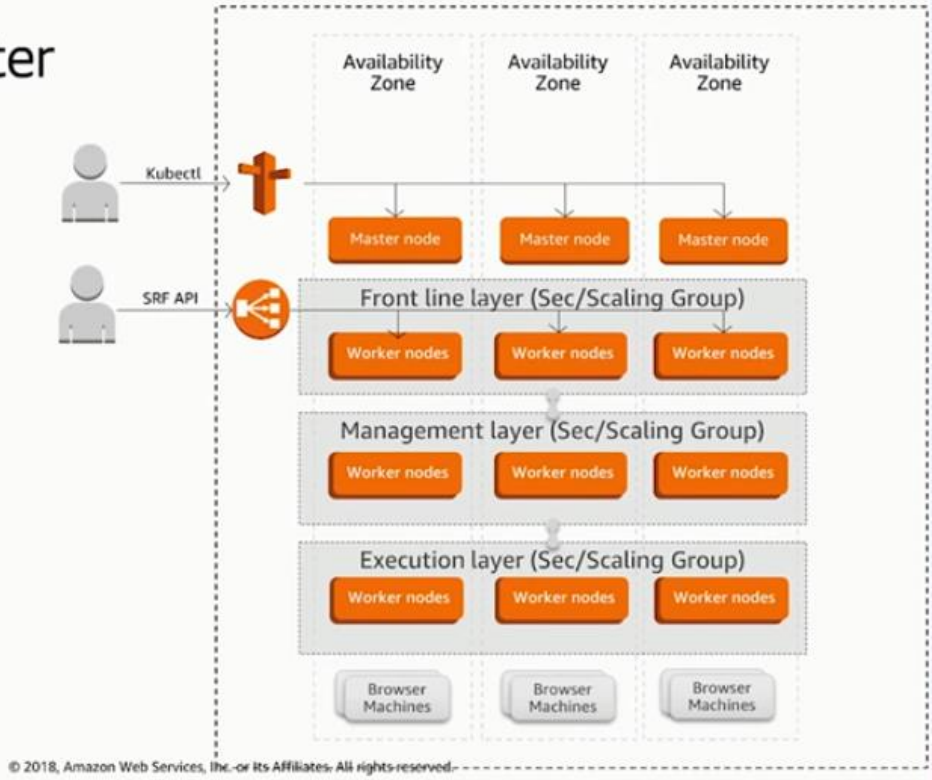

Base setup

Base image: Ubuntu 16.04 (hardened)

Networking:

- AWS Routing (kube-net)
- Researching calico for better security

Deployments:

- Development – Docker Compose (better debugging options)
- DevOps – Work with mini-kube for testing locally
- Cloud – Kubernetes, Pods spread across AZ

Special cases:

- Executor pools: manager pod as DaemonSet (easy visibility)
- Service config: ConfigMap + Hot configuration via queue

2 CI/CD of apps on Kubernetes

CI/CD of apps on Kubernetes—choices

Jenkins

AWS CodePipeline, AWS CodeCommit, AWS CodeBuild

AWS partners

- GitLab
- Shippable
- CircleCI
- Codeship

Jenkins – CI/CD with Kubernetes

```
node {

    stage 'Checkout'
    git 'https://github.com/omarlari/aws-container-sample-app.git'

    stage 'Build Dockerfile'
    docker.build('hello')

    stage 'Push to ECR'
    sh ("eval \$(docker run awscli aws ecr get-login --region ${REGION} --no-include-email | sed 's|https://||')")
    docker.withRegistry('https://${ECR_REPO}') {
        docker.image('hello').push('${BUILD_NUMBER}')
    }

    stage 'update application'

    kubernetes: { node {
    docker.image('kubectl').inside("--volume=/home/ec2-user/.kube:/config/.kube"){
        sh 'kubectl describe deployment ${APP}'
        sh 'kubectl set image deployment/${APP} hello=${ECR_REPO}/hello:${BUILD_NUMBER}'
        sh 'kubectl describe deployment ${APP}'
        }
    }}
}
```

Jenkins – CI/CD with Kubernetes

```
node {

    stage 'Checkout'
    git 'https://github.com/omarlari/aws-container-sample-app.git'

    stage 'Build Dockerfile'
    docker.build('hello')

    stage 'Push to ECR'
    sh ("eval \$(docker run awscli aws ecr get-login --region ${REGION} --no-include-email | sed '
    docker.withRegistry('https://${ECR_REPO}') {
        docker.image('hello').push('${BUILD_NUMBER}')
    }

    stage 'update application'

    kubernetes: { node {
    docker.image('kubectl').inside("--volume=/home/ec2-user/.kube:/config/.kube"){
        sh 'kubectl describe deployment ${APP}'
        sh 'kubectl set image deployment/${APP} hello=${ECR_REPO}/hello:${BUILD_NUMBER}'
        sh 'kubectl describe deployment ${APP}'
        }
    }}
}
```

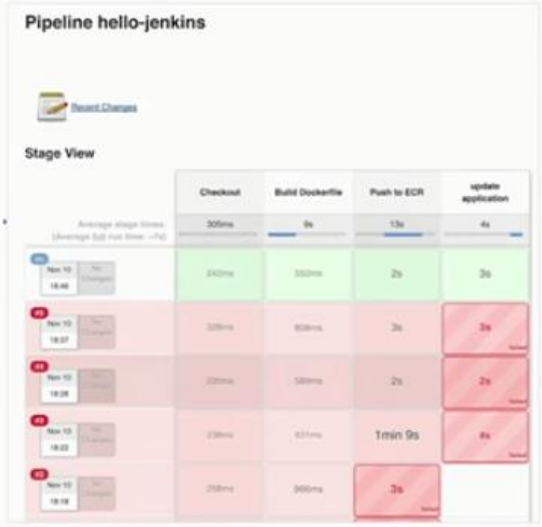

AWS CodePipeline – CI/CD with Kubernetes

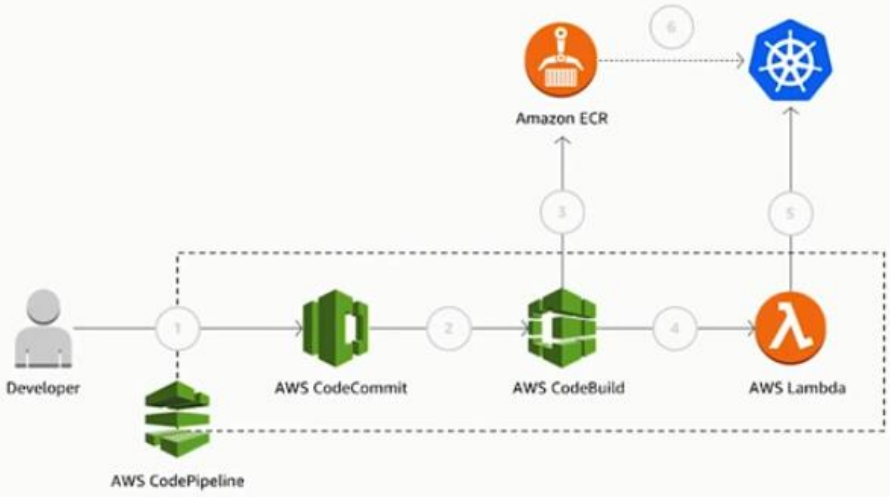

1. Developers continuously integrate changes into a main branch hosted within a repo
2. Triggers an execution of the pipeline when a new version is found, builds a new image with build ID
3. Pushes the newly built image tagged with build ID to ECR repo
4. Invokes a Lambda function to trigger application deployment
5. Leverages Kubernetes Python SDK to update a deployment
6. Fetches new container image and performs a rolling update of deployment

CI/CD at Microfocus on Kubernetes

SRF CI/CD with Kubernetes

Around 100 developers total

- 13 Dev Teams writing Java, NodeJS & Go
- Commits trigger auto deployment
- ~40 different git repos

SRF CI/CD with Kubernetes

A dedicated DevOps team

- Research into kubernetes & AWS features
- Upgrading to latest kubernetes version
- Implementing Jenkins pipelines
- Tweaking service auto-scaling & resource limits
- Improving Kubernetes security
- Deploying SRF to production

DevOps

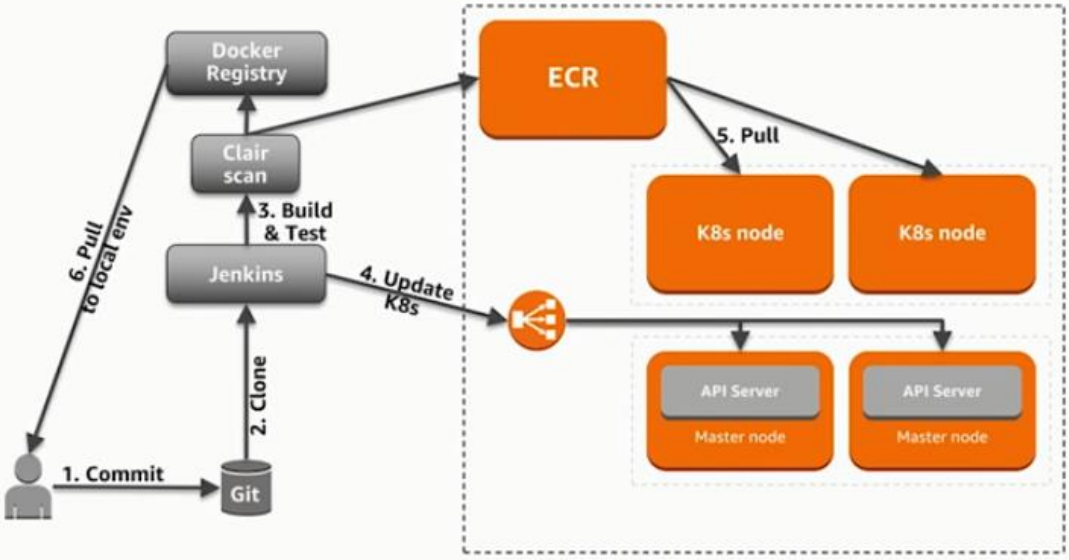

Pipelines I : Jenkins jobs

Parametrized Infra Jenkins Jobs

- Commit, Compile Docker Build, Deploy, Test, etc.
- Manager Jobs for each service
- A specialized build for:
  - Common modules
  - AMI Creation

Pipelines I : Issues with Jenkins jobs

- **Hard to trace:** Multiple cascading jobs with parameters
- **Dev needs DevOps:** for build changes
- **Slow pipeline creation**: hard to create new build processes (next versions / private branches)
- **No build versioning:** build structure is managed manually, backed up as a whole

Pipelines II : Build As Code

Dev + DevOps Cooperation: pipeline maintenance is a joint effort

- **Multi stage docker**
- **Jenkinsfile**
- **Gulpfile**
- **Kubernetes manifest** files
  - In each Git repo

Environment creation as code:

- **Helm for whole-system deployment script**
- **Terraform for environment creation (external infra & network)**
- **Packer for browser AMI construction (Windows / Linux / +Mac vSphere templates)**
  - **Kops for cluster creation**

Allowed us to take control of infrastructure from Jenkins

Benefits of "Build As Code"

- **Free Dev**: No need to wait for DevOps
- **Free Ops**: Less manual work on pipelines
- **Traceability**: one job per build
- **Carry-on Build**: Branching simply copies pipeline scripts
- **Revertability**: Build structure is versioned
- **Visibility**: Code visible to all (github search)

3

Visibility in cluster

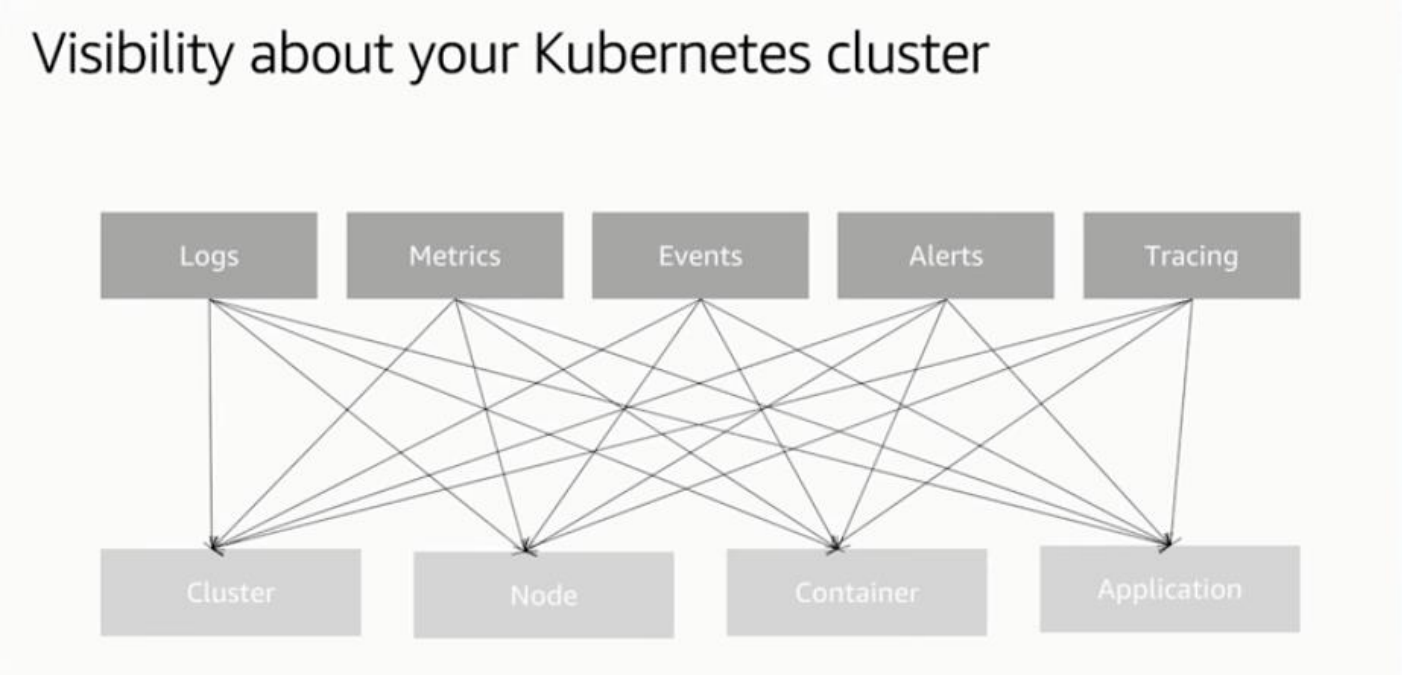

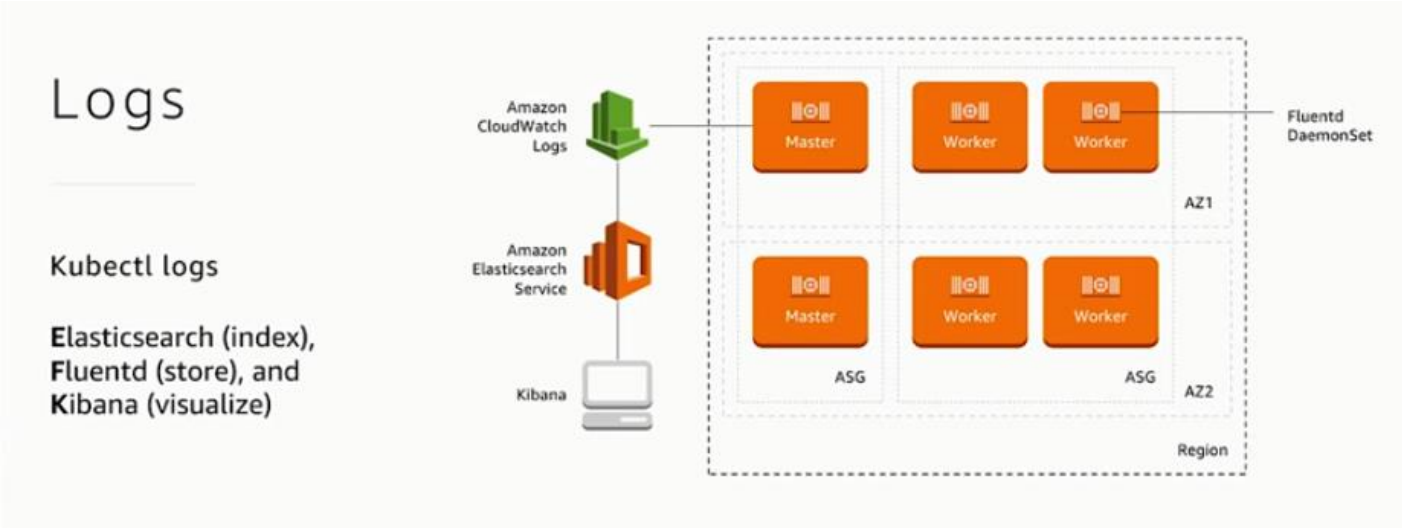

Fluentd DaemonSets scrapes the logs and sends them to CloudWatch Logs, you can then use a plugin from CloudWatch Logs to connect to Elasticsearch where you can set up your Kibana dashboards.

Metrics

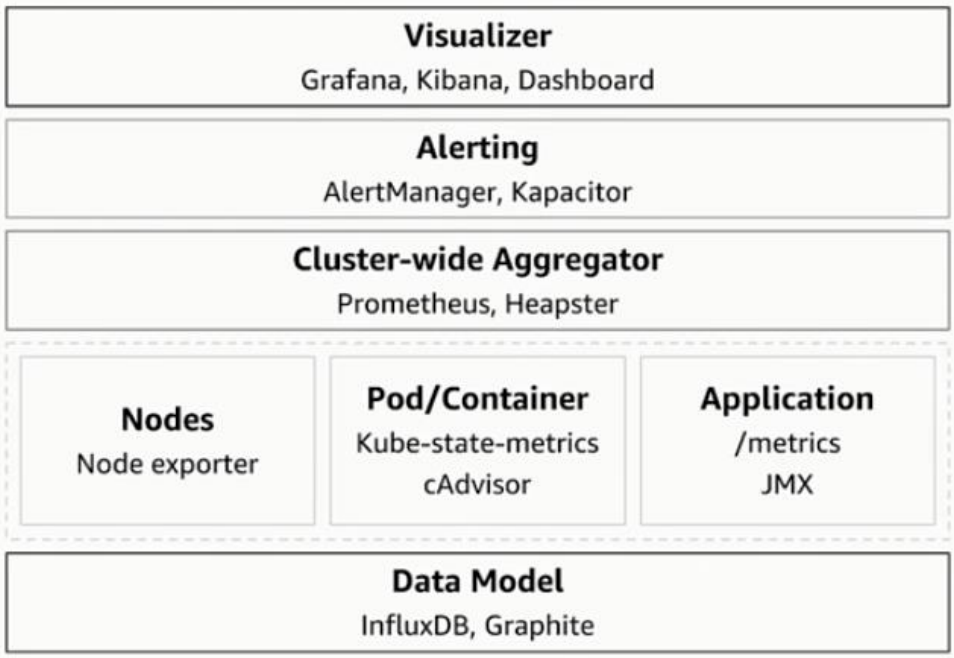

Application tracing with Kubernetes

Analyze and debug production, distributed applications, such as those built using a microservices architecture

Identify and troubleshoot the root cause of performance issues and errors

End-to-end view of requests as they travel through your application

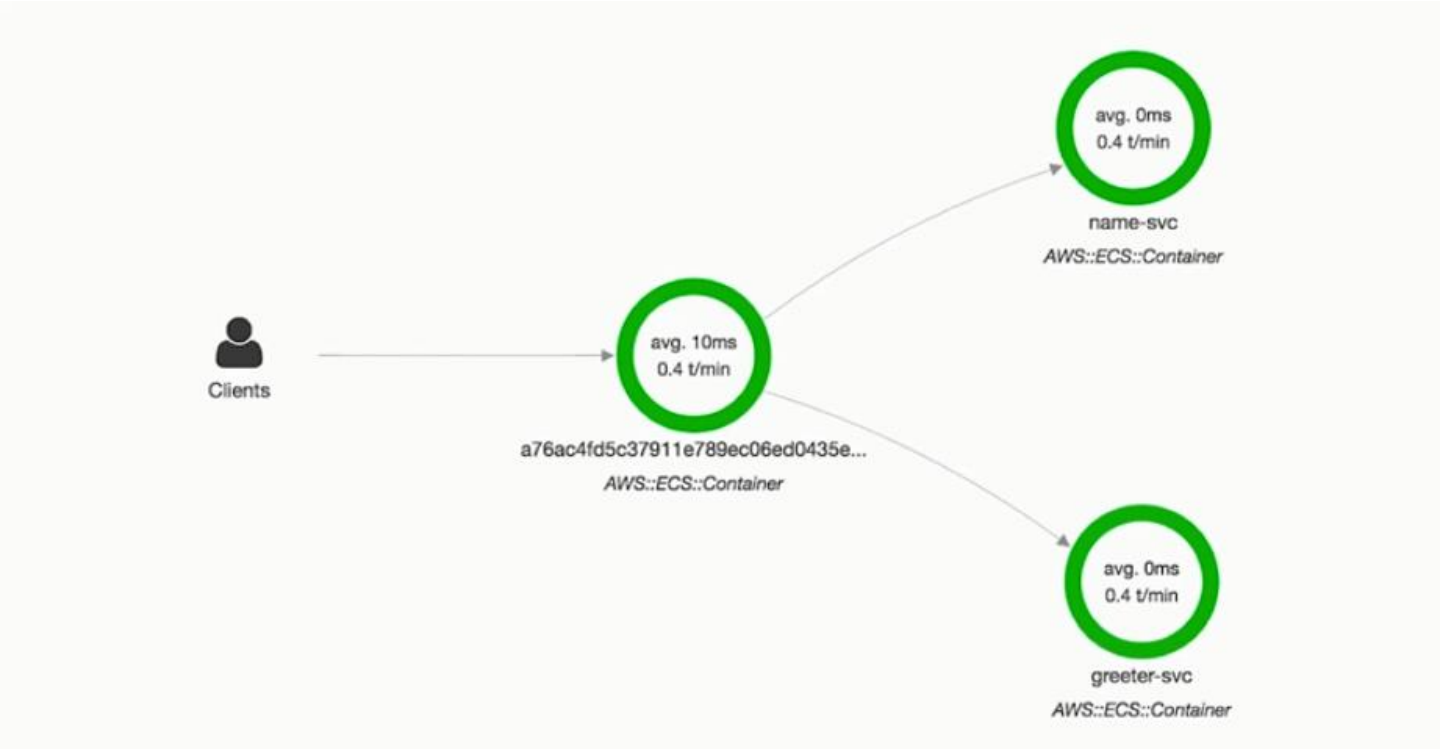

You can instrument with the XRay SDK that will use the XRay daemon and add to your nodes to collect traces

Visibility in cluster at Microfocus

Logging with Kubernetes

Centralized logging solution:

Elastic / FluentD / Kibana

Pods log to stdout

FluentD DaemonSet:

Sends logs to centralized ElasticSearch

Logtrail plugin in kibana:

External logs:

- Browser Machine Logs
- Script Execution Logs

Collected to FluentD Directory (applicative)

Collector uses FluentD on its machine

Monitoring Setup

Monitoring/alerting tool: **Prometheus**

Node exporter as DaemonSet for system metrics

External Prom dockers – outside of k8s cluster

cAdvisor collects pods metrics

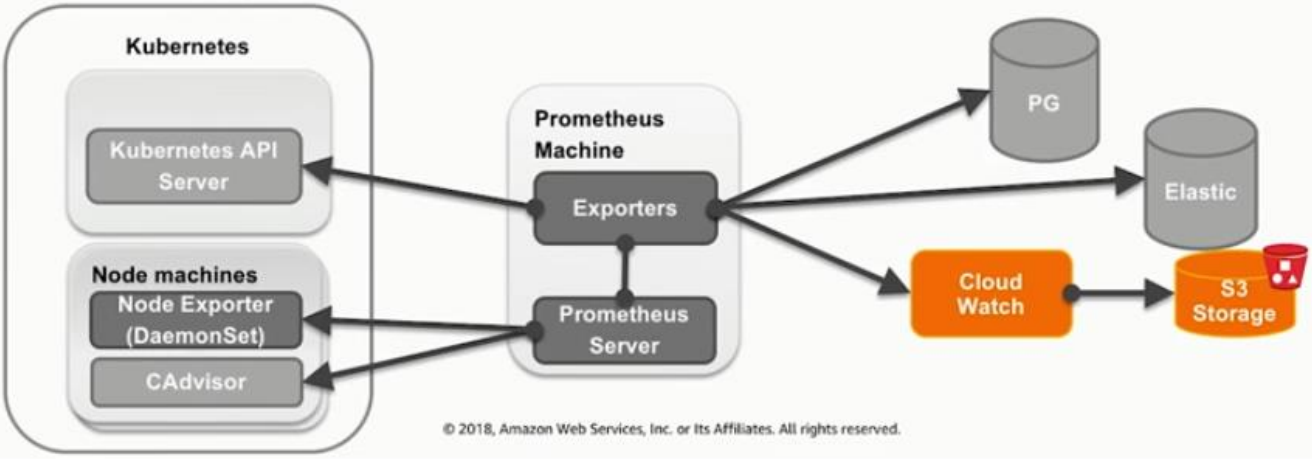

Application Monitoring & Alerts

Sytem Metrics tracked:

- **Liveliness**: Per System & Per service
- **System**: CPU / Memory / IO latency
- **API**: Latency, response rate, errors
- **Custom:** Browser pools & Execution slots

Alerts:

Alerts on system limits

Autoscaling Thresholds at 75% of limits

Infra Monitoring

Infra Metrics tracked:

- **Prometheus** (self monitoring)
- **Redis** (Elastic Cache)
- **Postgres** DB
- **ElasticSearch**
- **RabbitMQ**

Grafana – Cluster dashboard

Grafana – API Dashboard

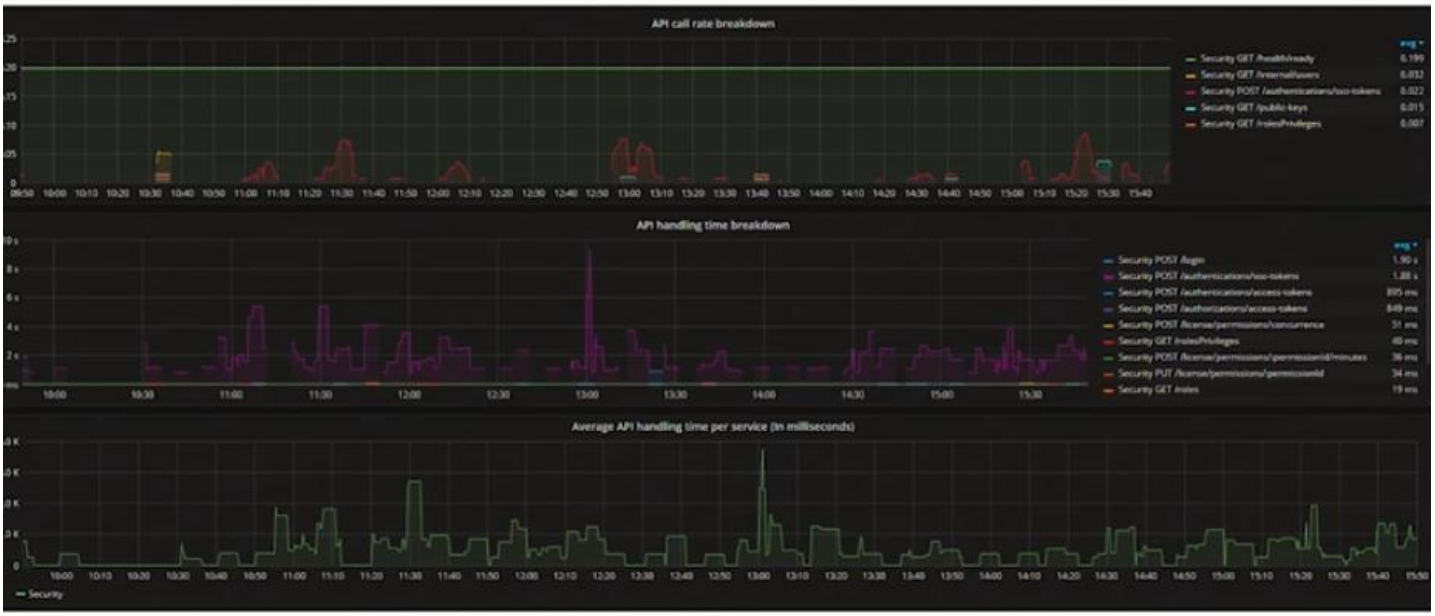

References

github.com/aws-samples/aws-workshop-for-kubernetes