

STG315

AWS re:INVENT

Case Study: Learn How Thomson Reuters Uses Amazon EFS to Deliver Billions of Pieces of Content to Hundreds of Millions of Visitors Every Year

Edward Zhitnitsky - Production Infrastructure Manager, Thomson Reuters

Andy Crudge - Amazon EFS Product Manager

November 30, 2017

AWS re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



In the fast-paced world of news media, having a high performance, reliable storage service is a critical component for delivering content at scale. Learn how Thomson Reuters leverages Amazon EFS to host their rich media content, saving them time and money while delivering unmatched reliability. Come hear about their journey, their approach, their architecture, and the considerations they took when building out their environment. Best practices and tips for success will be shared throughout.

Agenda:

Amazon Storage and EFS Overview

Case Study: Thomson Reuters

Reference and Other Sessions

Q&A

What do you think about when choosing a storage solution?



Storage type




Features
and performance



Economics

Three types of storage



File



Block



Object

Three types of storage



File

Data stored as files in a
directory hierarchy

Shared over a network

Three types of storage



Block

Data stored as blocks on a disk or disks

Locally attached

Three types of storage



Object

Data is stored as an object that's identified by a key in a flat space

Simple API to get and put data based on key

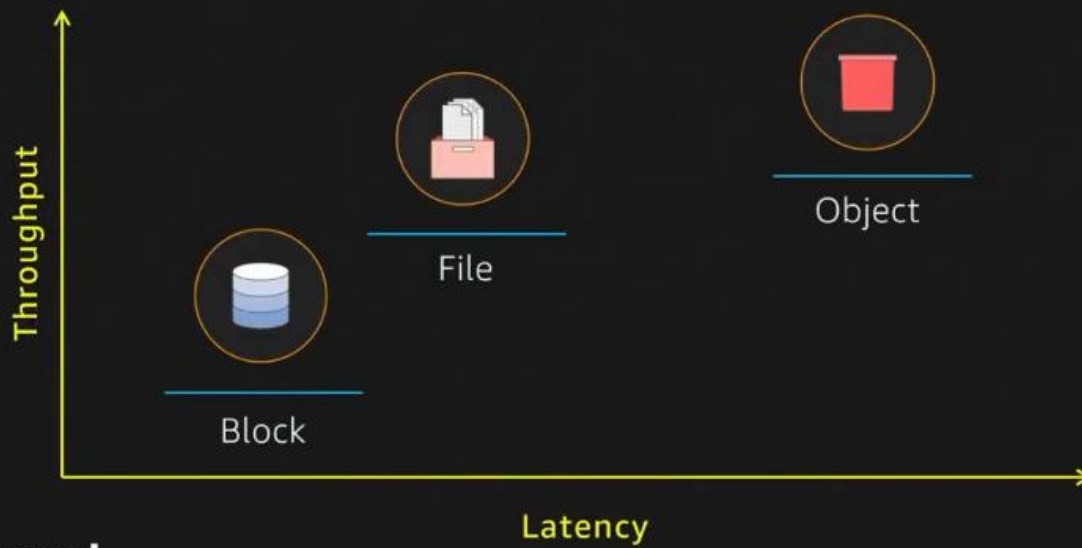
Why is file storage so popular?

Works natively with operating systems

Provides shared access while providing consistency guarantees and locking functionality

Provides hierarchical namespace

How does performance compare?




AWS re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



Before EFS... DIY file storage



AWS re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.





Amazon EFS

A fully managed file service

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

Key features of Amazon EFS



Simple



Elastic



Scalable



Highly available and durable

AWS
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

aws

Designed for a wide spectrum of needs




Analytics
M&E workflows



Web serving
Content management
Database backups
Enterprise apps
& messaging



Dev tooling
Home directories
Container storage



AWS
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



Amazon EFS customers

Database backups



J.D. POWER Sysco

Media and entertainment workflows

bambuser

zynga

Enterprise applications



Cardinal Health

Atlassian

infor

Developer tools



coursera



Home directories

Container storage

NETFLIX

here

Web serving + content management

Big data and analytics



THOMSON REUTERS




workfront



PlanGrid



finra



toast



ZS



MONSANTO




SAP



Google Cloud



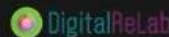
SPOKEO




BeSoftware



Seeking Alpha



DigitalReLab



custora



kissmetrics



Cognome

AWS
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



Security model and security features



Control network traffic
using VPC security groups and network ACLs



Control file and directory access
using POSIX permissions



Control administrative access (API access)
using AWS Identity and Access Management (IAM)
(action-level and resource-level permissions)




Encrypt data at rest
using keys managed in AWS KMS

AWS re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



Access EFS from various environments



AWS re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



EFS economics



No minimum commitments
or up-front fees



No need to provision
storage in advance



No other fees, charges,
or billing dimensions


\$0.30/GB-Month (US Regions)
\$0.33/GB-Month (EU Ireland)
\$0.36/GB-Month (EU Frankfurt)
\$0.36/GB-Month (AP Sydney)

AWS
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



Before EFS... Costs of DIY file storage




AWS
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



Before EFS... Costs of DIY file storage



AWS
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



TCO example: For storing 500 GB,
Amazon EFS is 70% less than DIY

EFS cost: $(500\text{GB} * \$0.30/\text{GB-month}) = \underline{\$150 \text{ per month}}$

For DIY, you might provision 600 GB of EBS (i.e., ~85% utilization):

Storage (2x 600 GB EBS gp2 volumes):	\$120 per month
Compute (2x m4.xlarge instances):	\$290 per month
Inter-AZ data transfer costs (est.):	\$130 per month
Total	<u>\$540 per month</u>

AWS
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



Where is EFS available today?



- US West (Oregon)
- US East (N. Virginia)
- US East (Ohio)
- EU (Ireland)
- EU (Frankfurt)
- Asia Pacific (Sydney)


More coming soon!

AWS
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



Complete set of data building blocks



AWS
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



Introducing EFS File Sync

Sync data from existing file systems into Amazon EFS file systems



Simple

Set up and manage easily
from the AWS Console



Fast

Up to 5x faster than
standard Linux copy tools



Secure

Encrypted parallel data
transfer to AWS

Use EFS File Sync to copy...

file systems from on-premises to EFS

DIY in-cloud file systems to EFS

EFS file systems between AWS Regions

AWS re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



AWS re:Invent

Thomson Reuters Journey to EFS

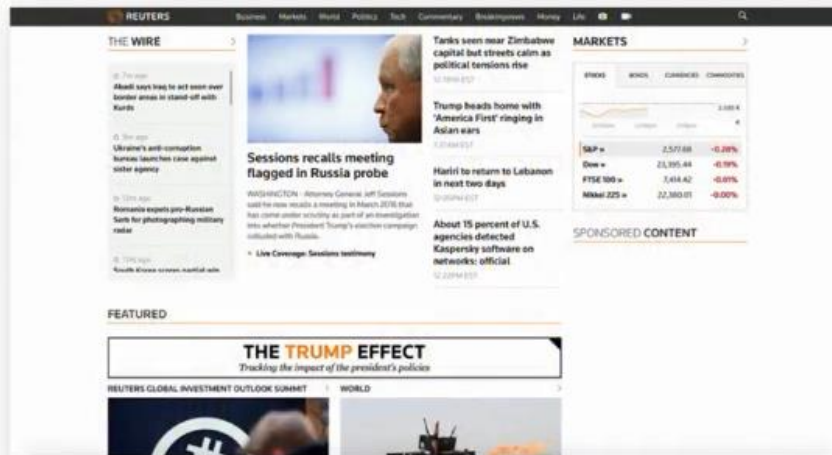
Reuters.com Re-Hosting – Edward Zhitnitsky

AWS re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



Reuters.com News, Financial Info, Video



AWS re:Invent

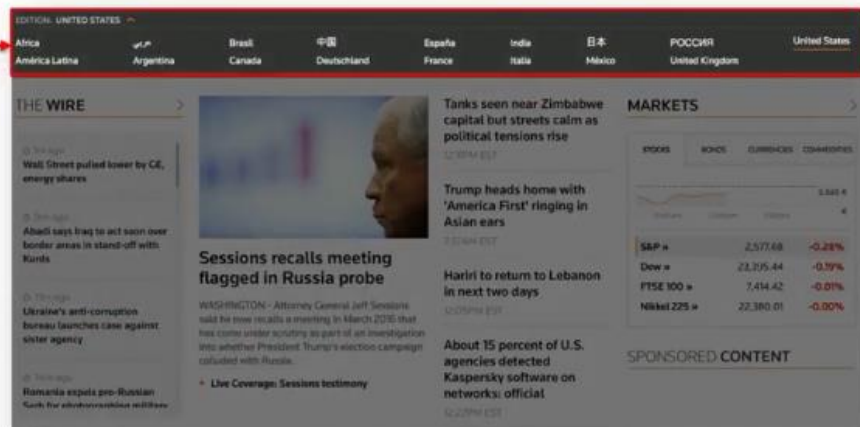
© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



Reuters.com – Usage Metrics



- 221,710,480 total views/month
- 48,546,796 unique visitors/month
- 17 desktop editions
- 4 mobile editions
- News
 - Articles
 - Video
 - Pictures
- Financial data
 - Stocks
 - Company Look-Ups




AWS re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



Previous Implementations

Previous Implementations



- Classic 3-tier architectures
- 120+ servers per environment
- GTMs
- Perimeter & Internal Firewalls
- 10 VLANs
- Load Balancers
- Edge Caching
- Master/replica news DB
- NAS/NFS for static content
- CMS
- Fixed assets require over provisioning



© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



Previous Implementations

- Aging architecture – 7+ years old
- Bandwidth commitments/constraints
- Performance
- Supportability
- Cost



© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

The Amazon Journey

The Amazon Journey – 3 possible solutions

Existing Provider Proposal

- Higher costs, particularly to design and commission service
- Unproven tools
- Lock-in (sizing was locked into approximating existing footprint)

Internal Data Centers

- Data centers location
- Shared resource and support limitations
- Difficulty quantifying true costs
- Lead times to scale, augment, or limits on system access

Amazon

- Learning curve
- Application migration
- Most cost effective
- Capacity on demand
- Full system control
- Clarity in costs

AWS re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

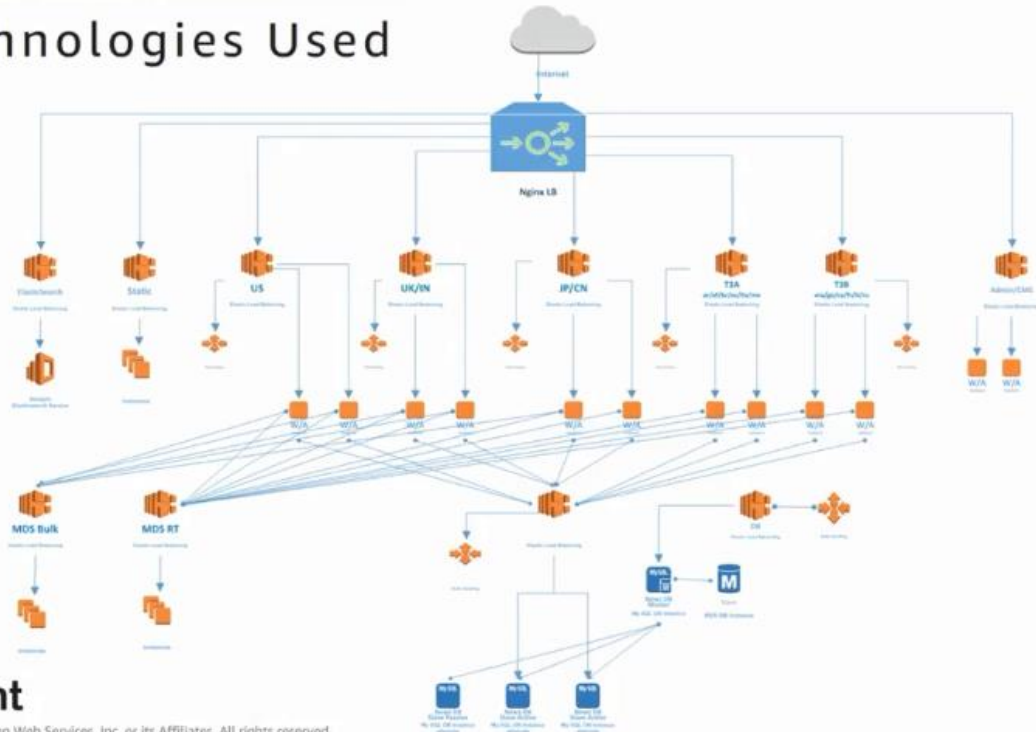


Decision: We're moving to AWS!!!



Technologies Used

Technologies Used




AWS re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



Technologies Used



AWS re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



Technologies Used

Migration to Amazon EFS

- Great performance
- Redundancy
- Metadata protection
- No capacity limitations
- Easily mounted
- Cost savings



Technologies Used

Challenge: How to migrate existing content to EFS?

AWS Snowball

- Prompt delivery
- Data security
- Metadata protection



The Bottom Line

The Bottom Line



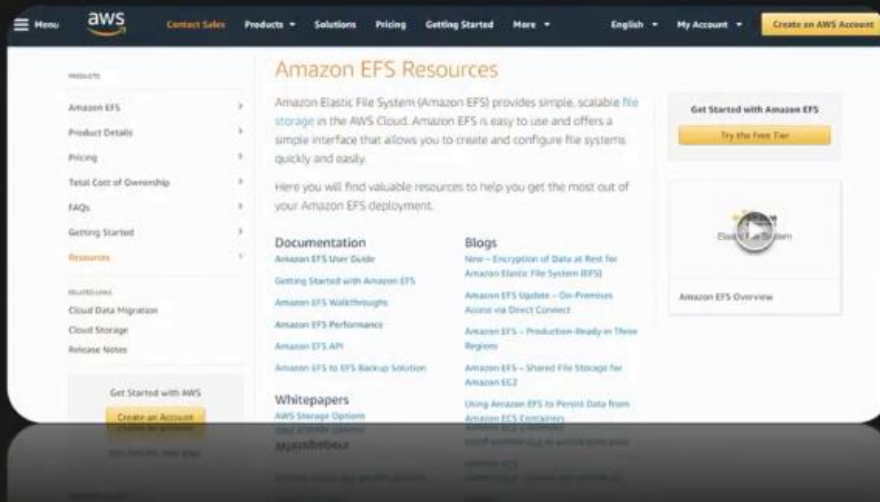
- Embraced new technology
- Using capacity on demand
- No more NAS Volume expansion
- Current hardware
- Bandwidth capacity eliminated
- Additional security added
 - DDoS protection
 - DNS scaling protection
- 67% cost savings




© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



How to learn more: aws.amazon.com/efs



Feature blogs
Whitepapers
Ref architectures
TCO calculator
10-minute tutorials
Documentation



© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



re:Invent EFS Sessions

Monday, 11:30AM	STG314: Case Study: Learn How HERE Uses JFrog Artifactory with Amazon EFS
Tuesday, 12:15PM	STG401: This Is My Architecture - Lightning Round: Most Innovative Storage Solutions
Tuesday, 2:00PM	STG327: Chalk Talk: Learn How to Maximize Throughput to EFS by Using Parallel Data Access Methods
Tuesday, 2:30PM	STG324: Learn to Build a Cloud-Scale WordPress Site That Can Keep Up with Unpredictable Changes
Tuesday, 4:45PM	STG307: Deep Dive on Amazon Elastic File System (Amazon EFS)
Wednesday, 12:15PM	STG325: Case Study: Learn How SiriusXM & Digital ReLab Leveraged Amazon EFS for their Media Workflows
Wednesday, 4:00PM	STG310: Files in AWS: Overcoming Storage Challenges for Common File Use Cases - Special Guest Celgene
Thursday, 12:15PM	STG307-R: Deep Dive on Amazon Elastic File System (Amazon EFS) [REPEAT]
Thursday, 3:15PM	STG315: Case Study: Learn How Thomson Reuters Uses Amazon EFS to Deliver Billions of Pieces of Content
Thursday, 3:15PM	STG321: Amazon EFS: Leverage the Power of a Distributed Shared File System in the Cloud [WORKSHOP]
Thursday, 5:30PM	STG326: How to Build a Scalable & Highly Available Messaging Service with IBM & TIBCO Using Amazon EFS

New Storage Training



For Enterprise Storage Engineers

- Learn how to architect and manage highly available solutions on AWS storage services
- Advance toward AWS certifications
- Help your organization migrate to the cloud faster

Online at www.aws.training

- Access 100+ new digital training courses including advanced training on storage
- Deep Dives on S3, EFS, and EBS
- Migrating and Tiering Storage to AWS (Hybrid Solutions)

At re:Invent

- Visit Hands-on Labs at the Venetian
- Attend a proctored "Introduction to EFS" Spotlight Lab on Thursday at 3pm at the Venetian
- Meet Storage experts at the Ask the Experts in Hands-on Labs room at the Venetian

AWS
re:Invent
THANK YOU!