

# AWS re:INVENT

## DAT311: Technology trends in Data Processing

**Anurag Gupta**, Vice President, Amazon Web Services  
[awgupta@amazon.com](mailto:awgupta@amazon.com)

**AWS re:Invent**

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



In this talk, Anurag Gupta, VP for AWS Analytic and Transactional Database Services, will talk about some of the key trends we see in data processing and how they shape the services we offer at AWS. Specific trends will include the rise of machine generated logs as the dominant source of data, the move towards Serverless, API-centric computing, and the growing need for local access to data from users around the world.

## Agenda

- Managing explosion of data
- Serverless, API-centric computing
- Global users, local access experience

## Managing Data Explosion with Data Lakes

## Traditionally, analytics used to look like this



Relational data

TBs-PBs scale


Schema defined prior to data load

Operational reporting and ad hoc


Large initial capex + \$10k-\$50k / TB

## Explosion of machine-generated data

Machine-generated data is growing **10x faster** than business data



Network connected smart devices drive variety and volume of data



Micro-services architecture increases need for real-time monitoring and analytics



Transition from IT to DevOps increases rate of change

# Data lakes extend the traditional approach



Relational and non-relational data

TBs-EBs scale

Schema defined during analysis

Diverse analytical engines to gain insights

Designed for low-cost storage and analytics

## Data lakes on AWS



Most ways to bring data in

Unmatched durability and availability at EB scale

Best security, compliance, and audit capabilities

Run any analytics on the same data without movement

Scale storage and compute independently

Store at \$0.027 / GB-month; Query for \$0.05/GB scanned

## Layers of a data lake



## AWS Glue – Serverless Data catalog & ETL service



Automatically discovers data and stores schema

Data searchable, and available for ETL

Generates customizable code

Schedules and runs your ETL jobs

Serverless

# Crawlers: Automatic schema inference

enumerate  
S3 objects

identify file type  
and parse files




**custom classifiers**

Apache log parser


**built-in classifiers**

JSON parser  
CSV parser  
Parquet parser  
...

semi-structured  
per-file schema



semi-structured  
unified schema



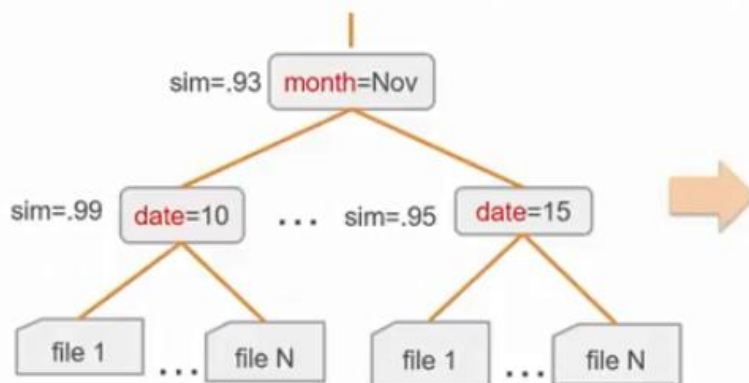
**AWS re:Invent**

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



# Crawlers: Automatic partitions detection

**S3 bucket hierarchy**



**Table definition**

Column	Type
month	str
date	str
col 1	int
	float
⋮	⋮

Estimate schema similarity among files at each level to  
handle semi-structured logs, schema evolution...

**AWS re:Invent**

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.





# Amazon Redshift – Data Warehousing

Fast, powerful, simple, and fully managed data warehouse at 1/10 the cost

Massively parallel, scale from gigabytes to petabytes

## Fast at scale



Columnar storage technology to improve I/O efficiency and scale query performance

## Open file formats



Analyze optimized data formats on the latest SSD, and all open data formats in Amazon S3

## Secure



Audit everything; encrypt data end-to-end; extensive certification and compliance

## Inexpensive



As low as \$1,000 per terabyte per year, 1/10th the cost of traditional data warehouse solutions; start at \$0.25 per hour

# Amazon EMR – Big data processing

Analytics and ML at scale

19 open-source projects: Apache Hadoop, Spark, HBase, Presto, and more

Enterprise-grade security

## Latest versions



Updated with the latest open source frameworks within 30 days of release

## Low cost



Flexible billing with per-second billing, EC2 spot, reserved instances, and auto-scaling to reduce costs 50-80%

## Use S3 storage



Process data directly in the S3 data lake securely with high performance using the EMRFS connector

## Easy



Launch fully managed Hadoop & Spark in minutes; no cluster setup, node provisioning, cluster tuning

# Amazon Elasticsearch Service

Easy to deploy, secure, operate, and scale Elasticsearch

Customers use Elasticsearch for log analytics, full-text search, and application monitoring

## Easy to Use



Fully-managed.  
Deploy production-ready  
clusters in minutes.

## Open



Direct access to  
Elasticsearch open-source  
APIs. Supports Logstash  
and Kibana.

## Secure



Secure access with VPC  
to keep all traffic within  
AWS network.

## Available



Zone awareness replicates  
data between two AZs;  
automatically monitors and  
replaces failed nodes.

# Amazon Kinesis – Real time

Easily to collect, process, and analyze video and data streams in real time

New



## Kinesis Video Streams

Capture, process,  
and store video  
streams for analytics



## Kinesis Data Streams

Build custom  
applications that  
analyze data streams



## Kinesis Data Firehose

Load data streams  
into AWS data stores



## Kinesis Data Analytics

Analyze data streams  
with SQL

# Highly connected data, best represented in a graph

## Relational model

Foreign keys used to represent relationships  
Queries can involve nesting & complex joins  
Performance can degrade as datasets grow




## Graph model

Relationships are first-order citizens  
Easy to write queries that navigate the graph  
Results returned quickly, even on large datasets



# Building apps with highly connected data


## Relational databases



Unnatural for querying graph



Inefficient graph processing



Rigid schema, inflexible for changing graphs

## Existing graph databases



Difficult to scale



Difficult to maintain high availability



Too expensive



Limited support for open standards

## Amazon Neptune

NEW!

Fully managed graph database

### Open



Supports Apache TinkerPop & W3C RDF graph models

### Fast



Query billions of relationships with millisecond latency

### Reliable



Six replicas of your data across three AZs, with full backup and restore

### Easy



Gremlin  
SPARQL

Build powerful queries easily with Gremlin and SPARQL

## Serverless, API-Centric Computing



# Serverless Analytics


Deliver cost-effective analytic solutions faster

  
Serverless. Zero  
Infrastructure. Zero  
Administration.

  
Never pay for  
idle resources

  
Automatically  
scales resources  
with usage

  
Availability and  
fault tolerance  
built in



**AWS re:Invent**

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



## Amazon Athena—interactive analysis

Interactive query service to analyze data in Amazon S3 using standard SQL


No infrastructure to set up or manage and no data to load

Ability to run SQL queries on data archived in Amazon Glacier (coming soon)


### Query Instantly

  
Zero setup cost; just  
point to S3 and start  
querying

### Pay per query

  
Pay only for queries run;  
save 30-90% on per-query  
costs through  
compression

### Open


  
ANSI SQL interface,  
JDBC/ODBC drivers, multiple  
formats, compression types,  
and complex joins and data  
types

### Easy

  
Serverless: zero  
infrastructure, zero  
administration  
Integrated with QuickSight

# Redshift Spectrum

Extend the data warehouse to your S3 data lake



Exabyte Redshift SQL queries against S3

Join data across Redshift and S3

Scale compute and storage separately

Stable query performance and unlimited concurrency

CSV, ORC, Grok, Avro, and Parquet data formats

Pay only for the amount of data scanned

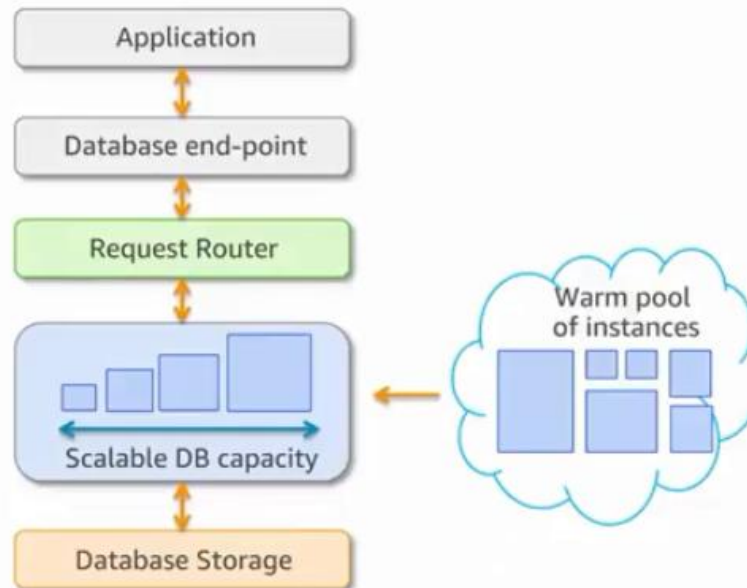
**AWS re:Invent**

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



# Aurora Serverless

- Starts up on demand, shuts down when not in use
- Automatically scales with no instances to manage
- No impact to applications during scaling events
- Pay per second for the database capacity you use



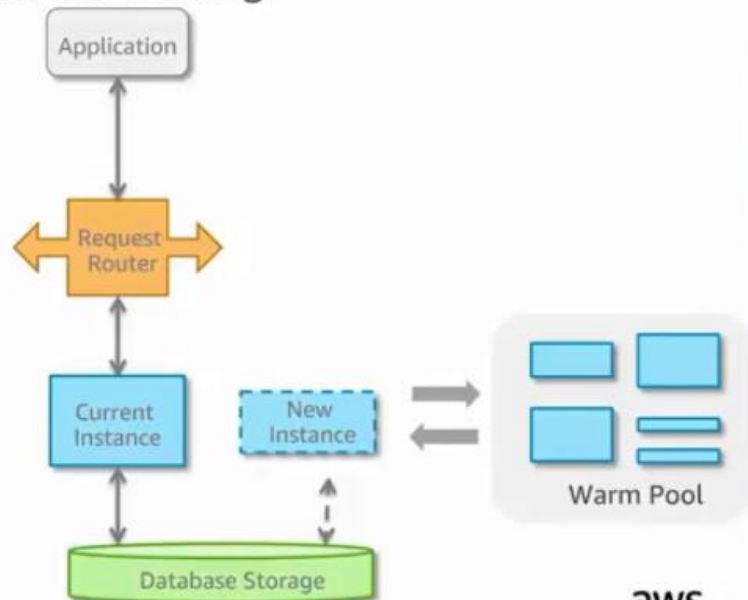
**AWS re:Invent**

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



# Instance provisioning and scaling

- First request triggers provisioning of a database instance. It typically takes about 5-10 secs.
- Instances scale-up and scale-down automatically in response to changes in workloads. Instance scaling takes about 1-3 secs.
- Instances are hibernated after a user-defined period of inactivity
- Scaling operations are transparent to the application – user sessions are not terminated
- Database storage is persisted until explicitly deleted by user



AWS  
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

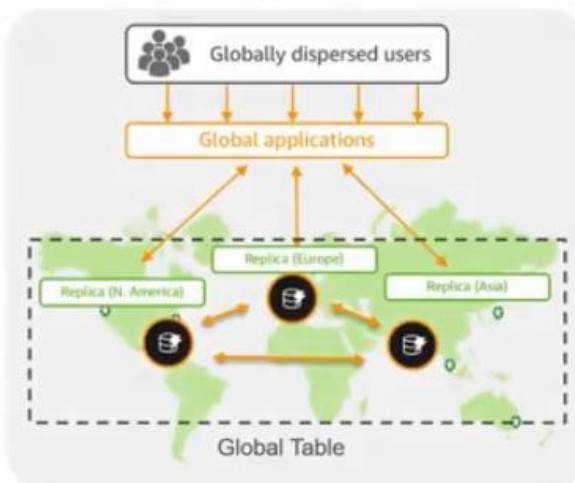
aws

## Global Users, Local Processing

### DynamoDB Global Tables (GA)

NEW!

First fully managed, multi-master, multi-region database




Build high performance, globally distributed applications

Low latency reads & writes to locally available tables

Disaster proof with multi-region redundancy

Easy to set up, and no application rewrites required

# Distributed Lock Manager




## Pros

- All data available to all nodes
- Easy-to-build applications
- Similar cache coherency as in multi-processors

**aws re:Invent**

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.




## Cons

- Heavyweight cache coherency traffic, on per-lock basis
- Networking can be expensive
- Negative scaling when hot blocks

**aws**

# Consensus with two phase or Paxos commit




## Pros

- Query broken up and sent to data node
- Less coherence traffic – just for commits
- Can scale to many nodes

**aws re:Invent**

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



## Cons

- Heavyweight commit and membership change protocols
- Range partitioning can result in hot partitions, not just hot blocks. Re-partitioning expensive.
- Cross partition operations expensive. Better at small requests

**aws**



# Conflict resolution using distributed ledgers

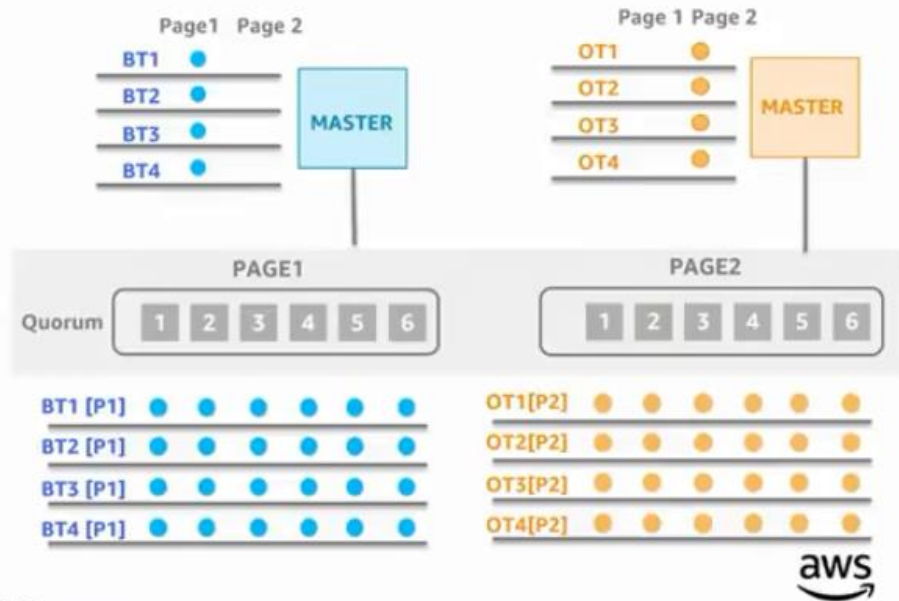
There are many "oases" of consistency in Aurora

The database nodes know transaction orders from that node

The storage nodes know transactions orders applied at that node

Only have conflicts when data changed at both multiple database nodes AND multiple storage nodes

Much less coordination required



**AWS re:Invent**

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



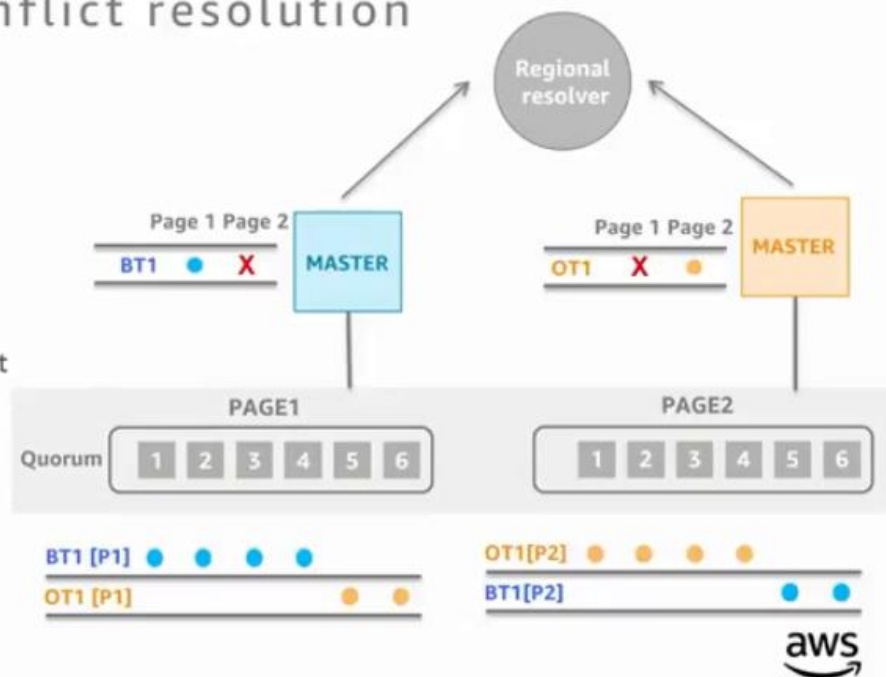
## Hierarchical conflict resolution

Both masters are writing to two pages P1 and P2

**BLUE** master wins the quorum at page P1; **ORANGE** master wins quorum at P2

Both masters recognize the conflict and have two choices: (1) roll back the transactions or (2) escalate to regional resolver

Regional arbitrator decides who wins the tie breaker.



**AWS re:Invent**


© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.





# Crash recovery in multi-master

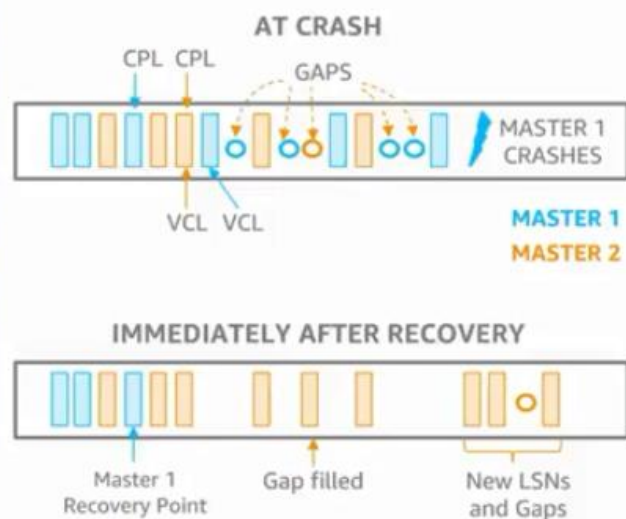
## SINGLE MASTER



**aws re:Invent**

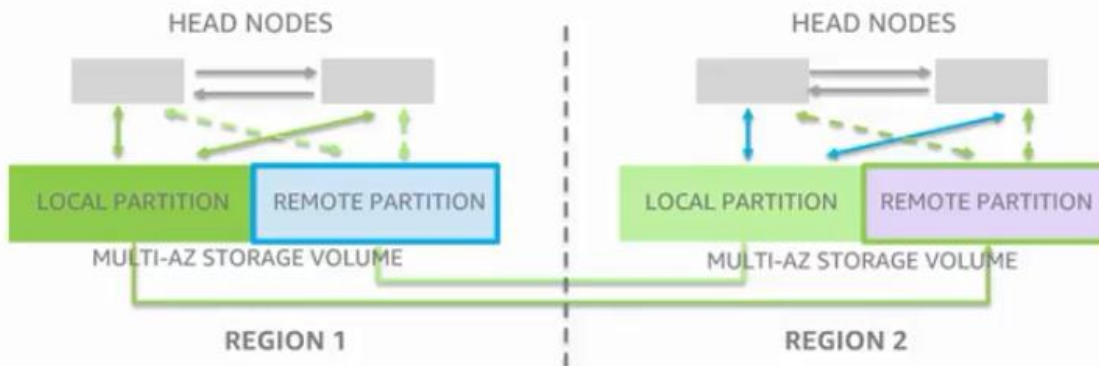
© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

## MULTI-MASTER



**aws**

# Multi-region Multi-master



Write accepted locally

Optimistic concurrency control – no distributed lock manager, no chatty lock management protocol

Conflicts handled hierarchically – at head nodes, at storage nodes, at AZ and region level arbitrators

Near-linear performance scaling when there is no or low levels of conflicts

**aws re:Invent**

**aws**

AWS  
re:Invent

Thank you!

AWS  
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

