

SRV409

A Serverless Journey: Under the Hood of AWS Lambda

Marc Brooker
Senior Principal Engineer
Serverless

Holly Mesrobian
Director of Engineering
AWS Lambda

aws
re:Invent

© 2018, Amazon Web Services, Inc. or its affiliates. All rights reserved.



Serverless computing allows you to build and run applications and services without thinking about servers. Serverless applications don't require you to provision, scale, and manage any servers. However, under the hood, there is a sophisticated architecture that takes care of all the undifferentiated heavy lifting for the developer. Join Holly Mesrobian, Director of Engineering, and Marc Brooker, Senior Principal of Engineering, to learn how AWS architected one of the fastest-growing AWS services. In this session, we show you how Lambda takes care of everything required to run and scale your code with high availability.

SERVERLESS AT SCALE IS THE NEW NORM



THOMSON REUTERS

processes **4,000 requests**
per second



processes **half a trillion**
validations of stock
trades daily



executes **16 million**
requests a month



ingests, analyzes and
stores **17+ petabytes of**
data per season



API traffic to register and license
more than **47 million driver**
records in Great Britain,



processes **tens of**
billions of data
points monthly

aws
re:Invent

© 2018, Amazon Web Services, Inc. or its affiliates. All rights reserved.



Running Highly Available Large Scale Systems Is a Lot of Work




Load Balancing

aws
re:Invent

© 2016 Amazon.com, Inc. or its affiliates. All rights reserved.



You need to have load balancing in every layer of your architecture for scalability and availability. You also need to configure the needed routing rules.



Scaling Up and Down

aws
re:Invent

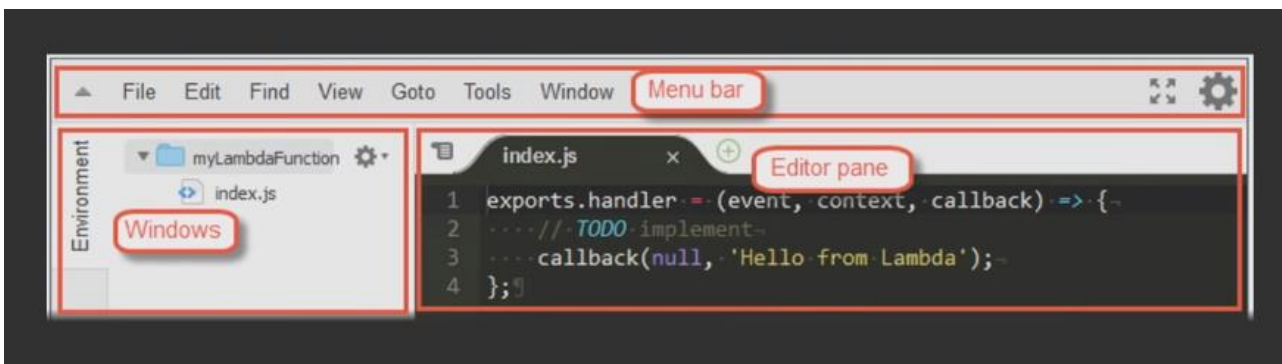
© 2016 Amazon.com, Inc. or its affiliates. All rights reserved.





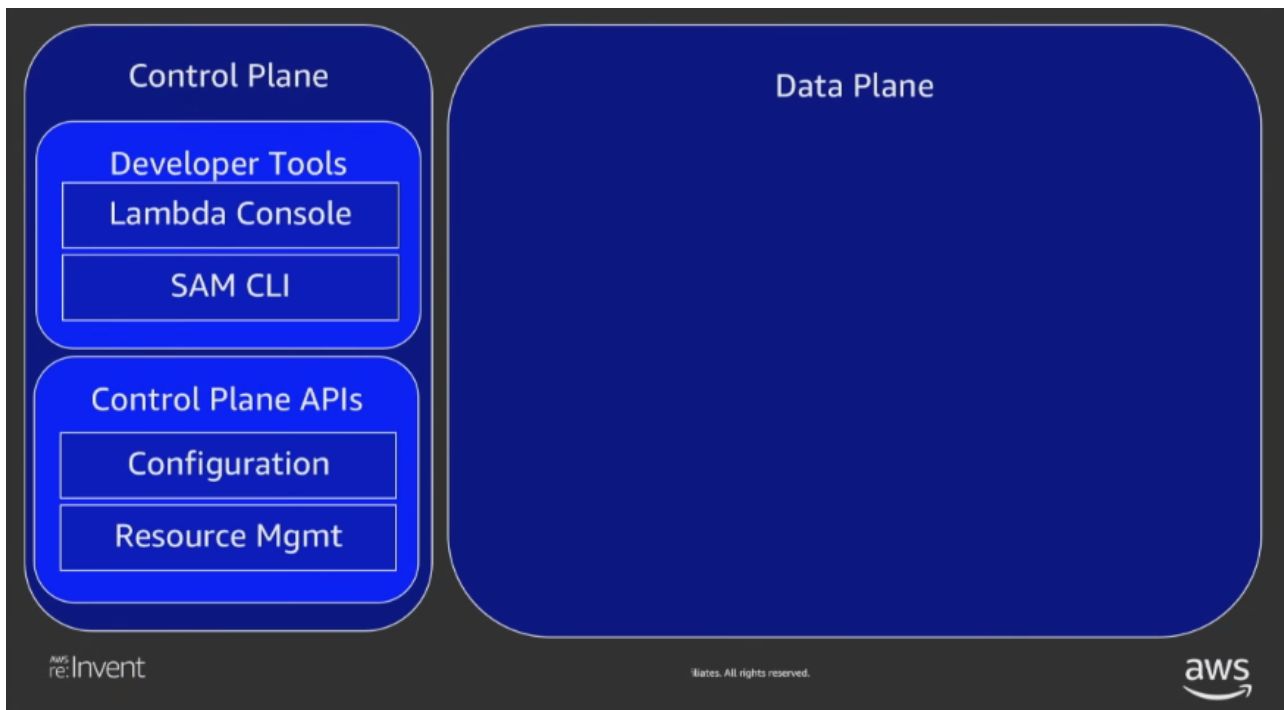
Handling Failures

You need to consider when hosts fail by configuring health-checks

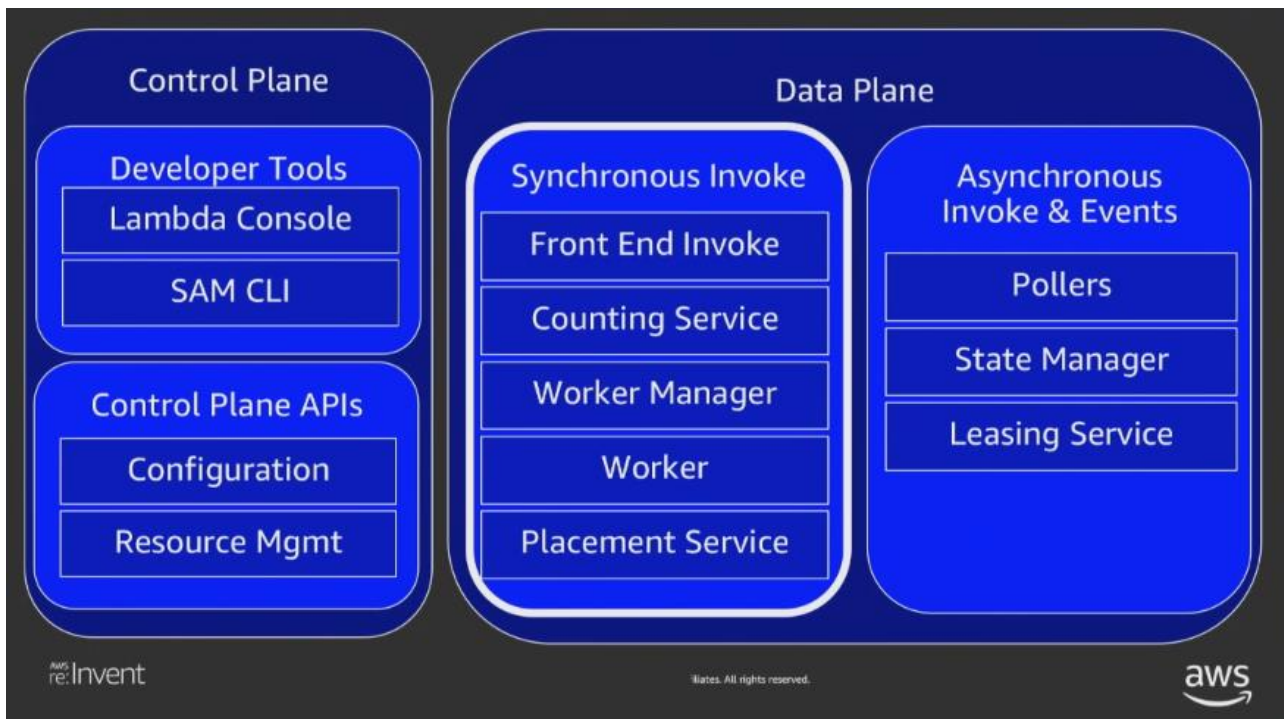


Lambda takes care of admin duties for you

Lambda Handles
Load Balancing
Auto Scaling
Handling Failures
Security Isolation
Managing Utilization
(and many other things) for you



The Lambda architecture is split into a control plane (CP) and a data plane (DP), the CP is where engineers and developers interact with the Lambda service using tools that end up calling the control plane APIs.



The data plane does the async invokes and inter-operates with other services that help with the event processing.

Front End Invoke

Orchestrate both synchronous and asynchronous Invokes

The first thing this service does is to authenticate the caller of the invoke function, it then retrieves and loads function metadata and code, confirms the concurrency with the counting service, then maps the function to the worker manager that controls the actual worker.

Counting Service

Provides a region wide view of customer concurrency to help enforce set limits

This service uses a quorum-based protocol and designed for high throughput, resilient, and HA.

Worker Manager

Tracks container idle and busy state and schedules incoming invoke requests to available containers

It spins up or down sandboxes to prevent waste and interacts with the placement service.

Worker

Provisions a secure environment for customer code execution

This creates and manages a collection of sandboxes needed for function executions, downloads customer code and executes them, manages communicating with other AWS services like CloudWatch.

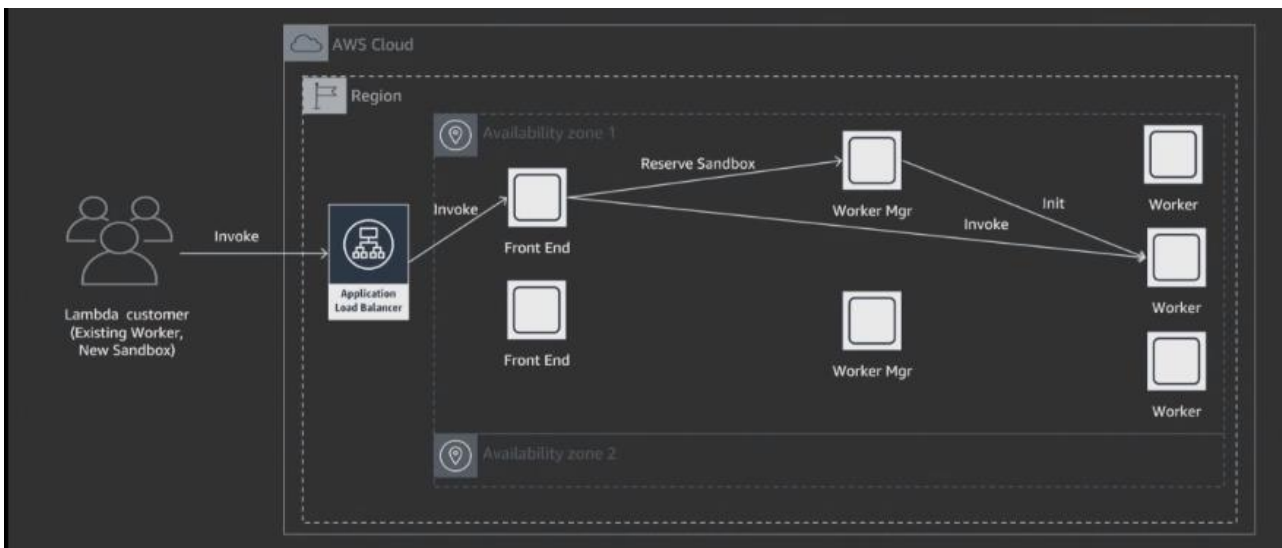
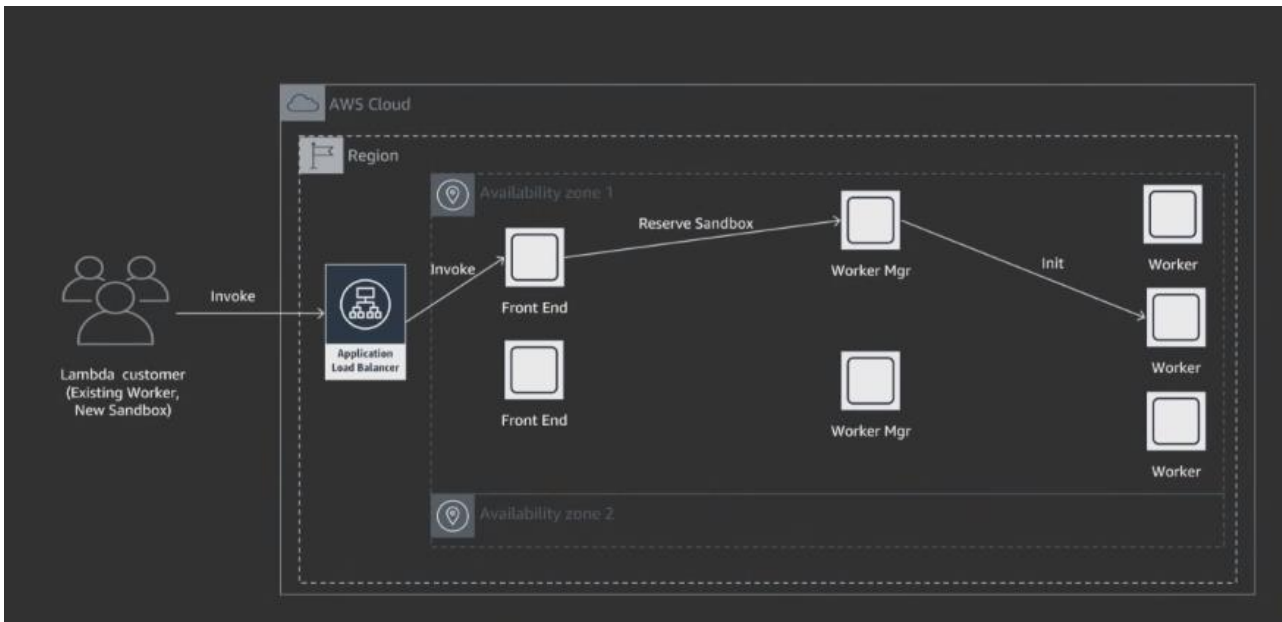
Placement Service

Places sandboxes on workers to maximize packing density without impacting customer experience or cold-path latency

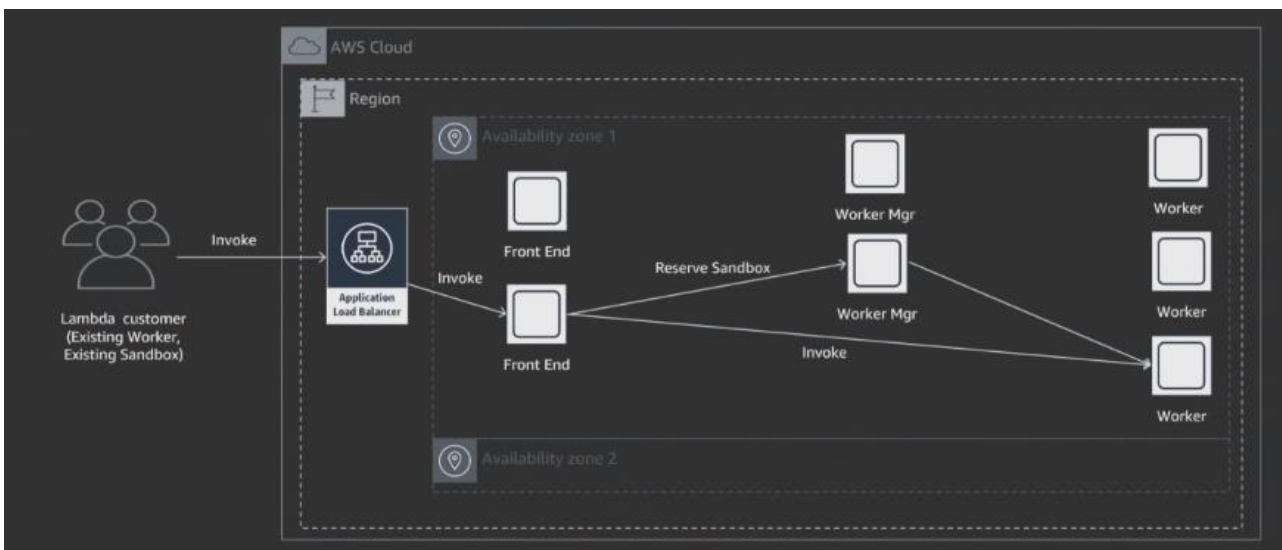
It also monitors worker health and removes bad workers.

Load Balancing

Routing function traffic across hosts distributed across Availability Zone



We now have a warm sandbox that we can call invoke on



Assuming we already have a warm sandbox when a call comes in, it gets faster to calling the init of your code

SERVERLESS CUSTOMERS



"What took us just a few days to build using a serverless solution based on AWS Lambda would have taken us six months to build from scratch. Our CTO and the rest of the project stakeholders were really happy with how much money and time we saved."

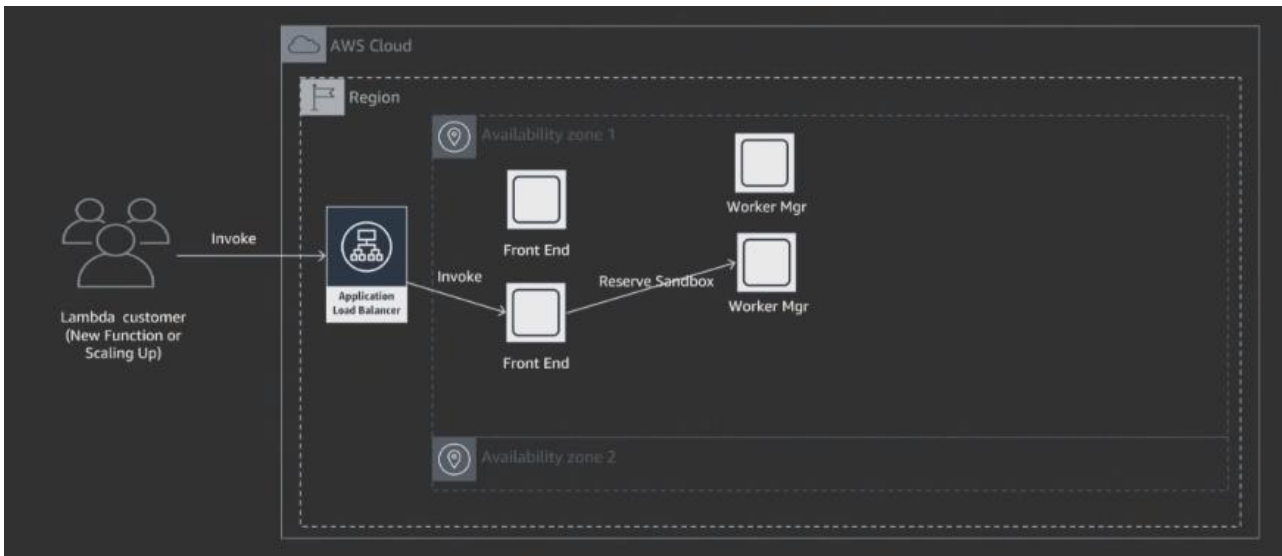
Nitin Mahajan

Executive Director for Service Engineering

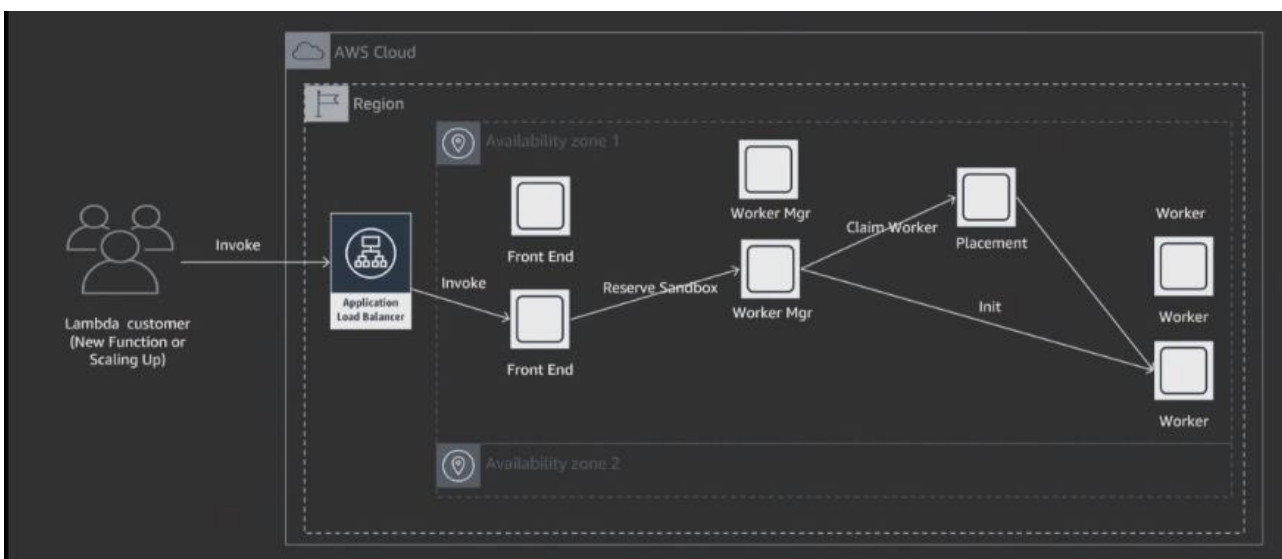
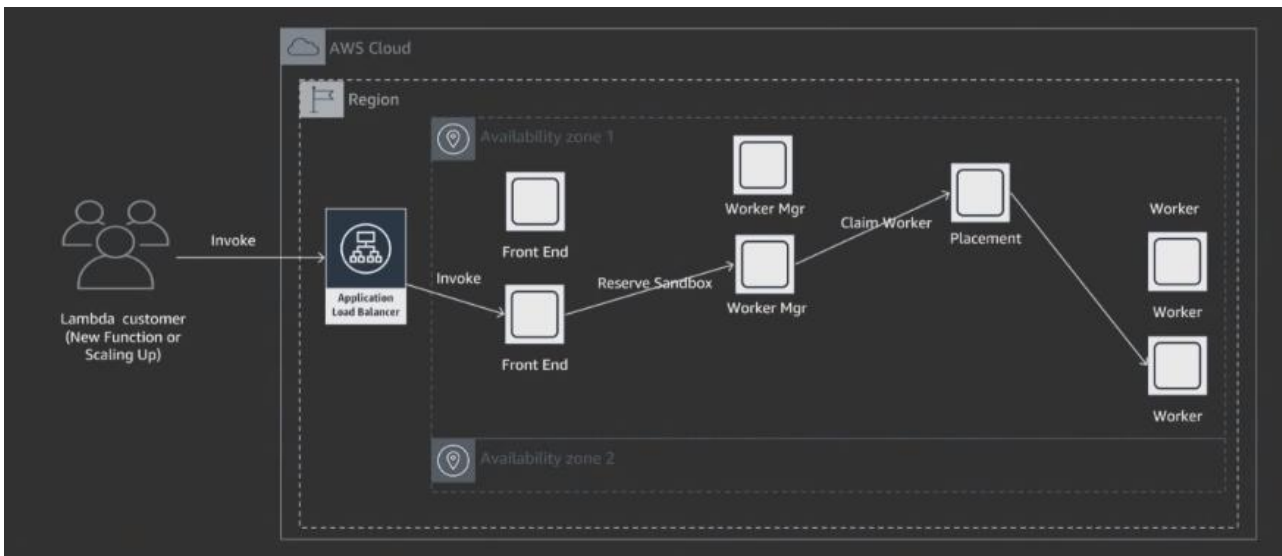
Auto Scaling

Provision function capacity when needed and releasing when not needed

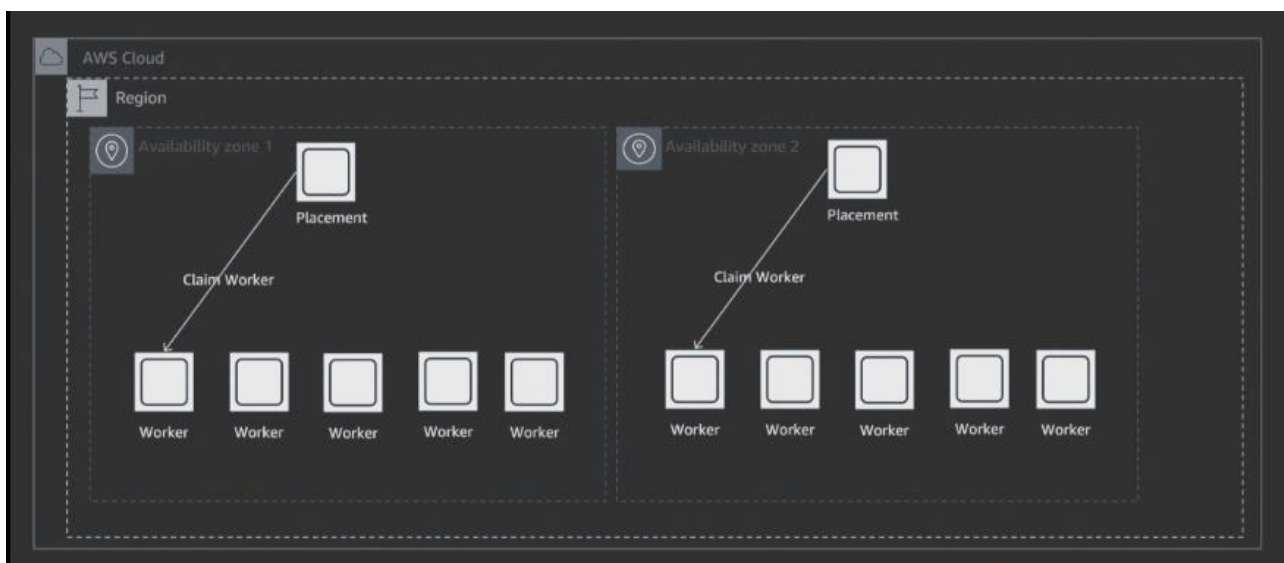
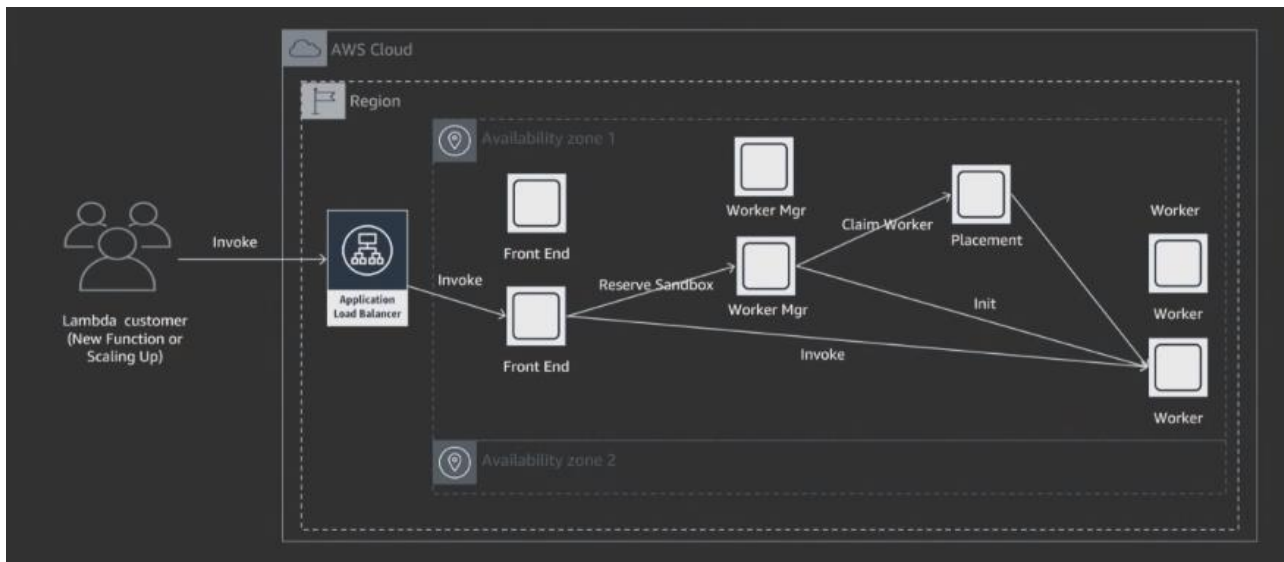
What happens when we need to scale up new workers quickly?



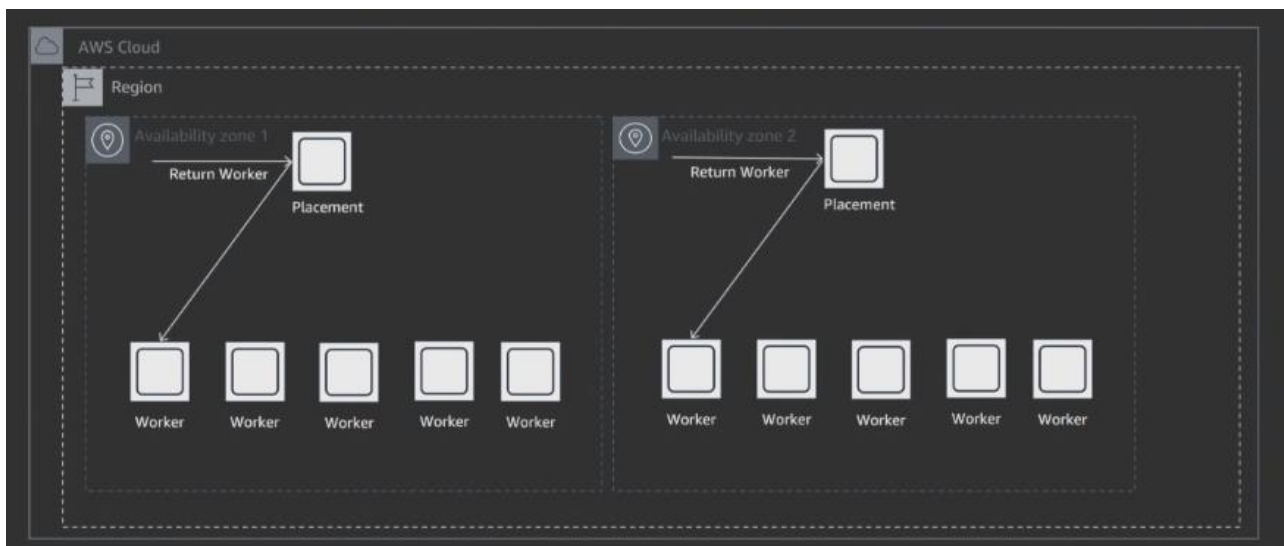
But this time the WM says it does not have a worker



We now have a warm sandbox that the front end can call invoke on



The Placement service ensures sufficient worker capacity to continue to fulfill Worker Manager requests for Worker hosts for a lease of between 6 – 10 hours.



The WM must return the Worker with an expiring lease.

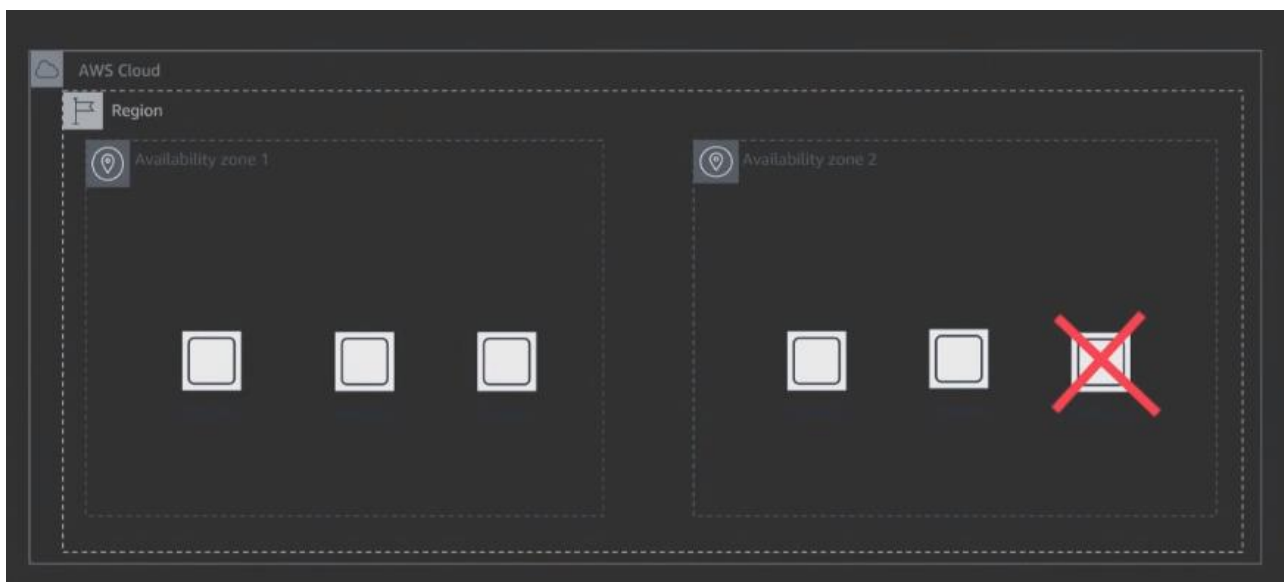


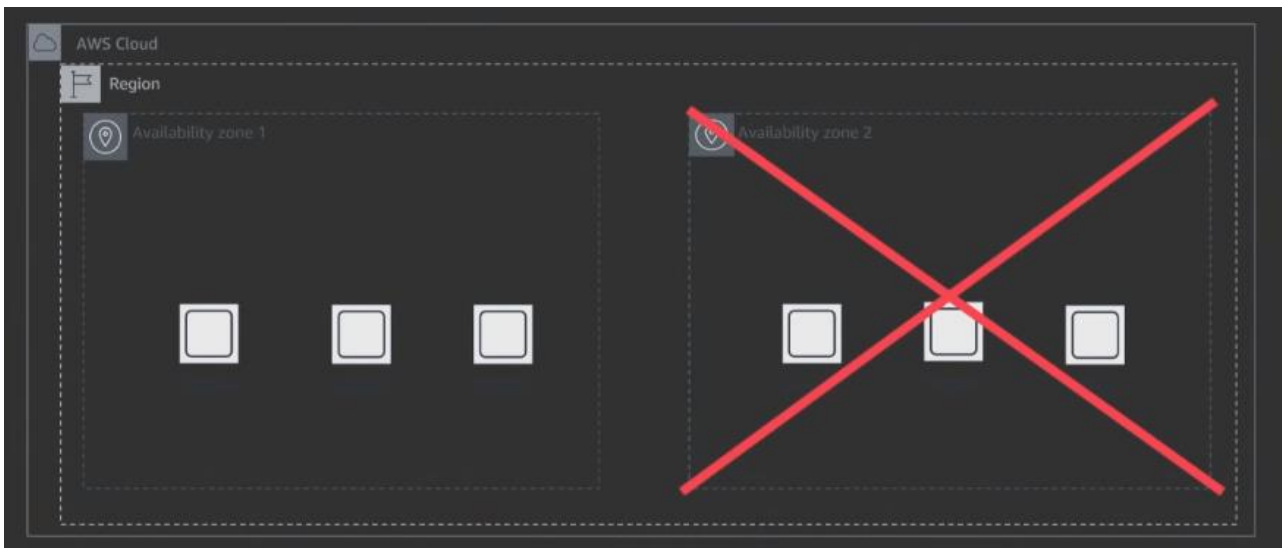
Workers are returned when their leases expire

Handling Failures

Handling Host and Availability Zone failure

With Lambda:
Always have a healthy host

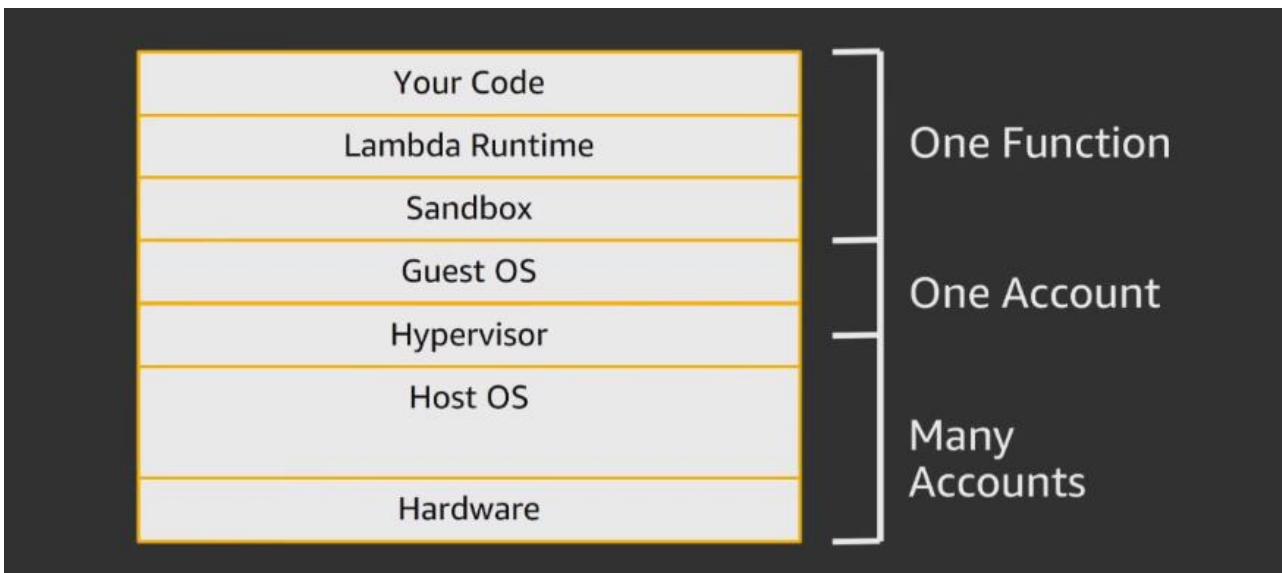




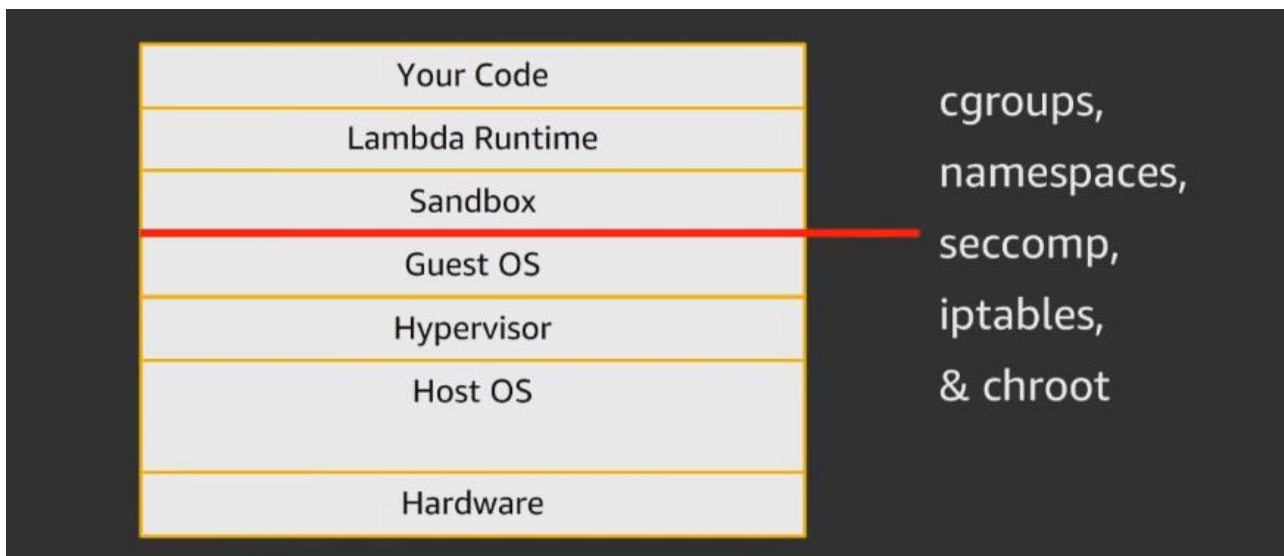
Isolation

Keeping Workloads Safe and Separate

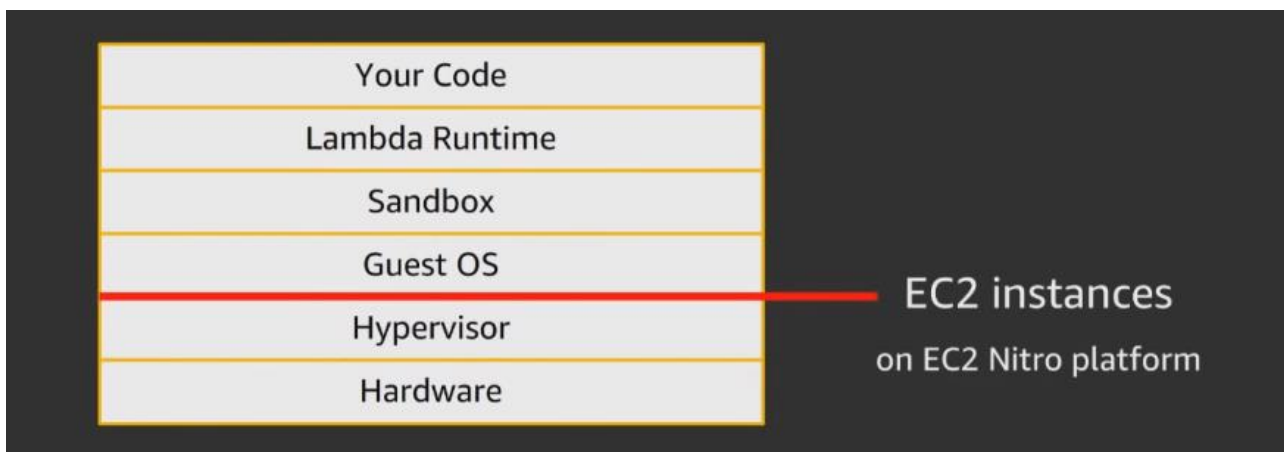
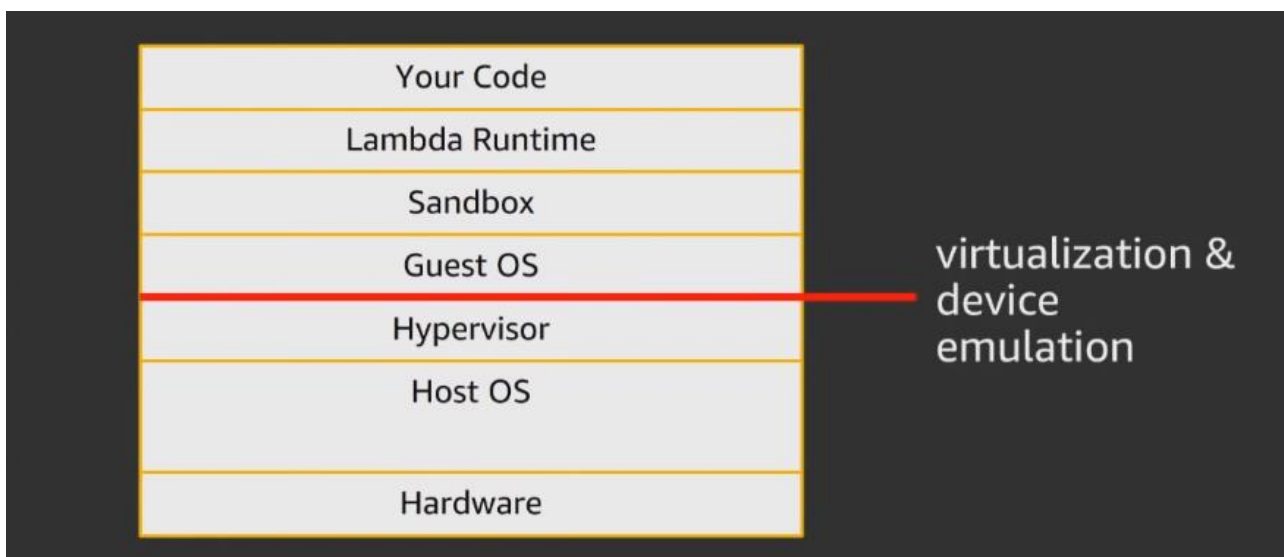
We run multiple Lambda functions on a single host at times, we keep them isolated and running with a consistent performance.

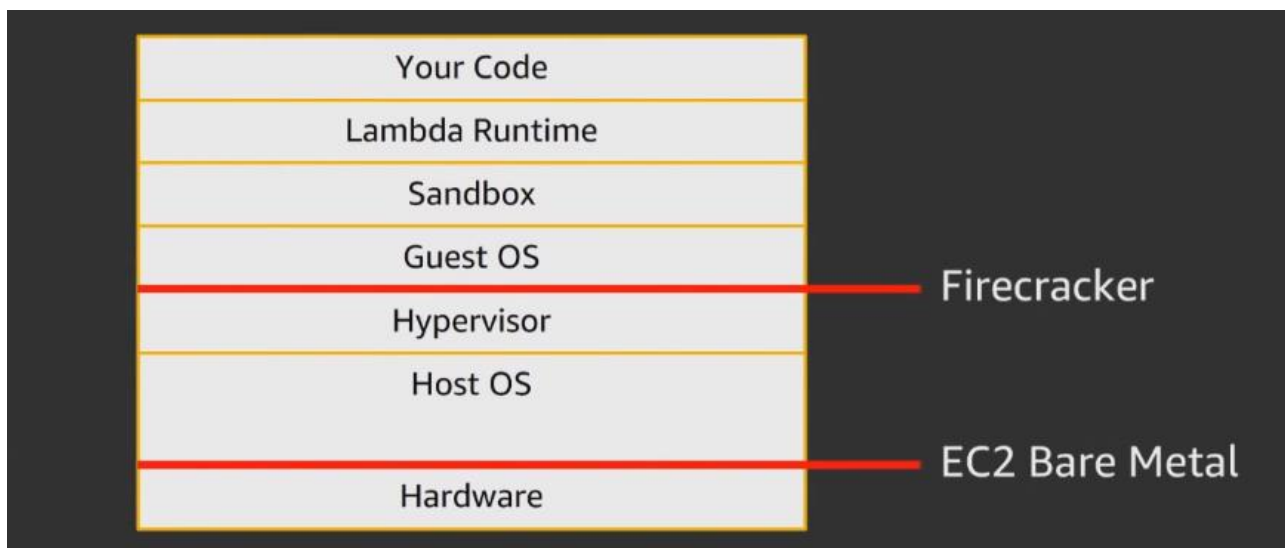


This is the stack running on a Worker node or host. The guest OS is Amazon Linux that is isolated from each other using a Hypervisor. Multiple same functions will run on the same sandbox but not overlap.

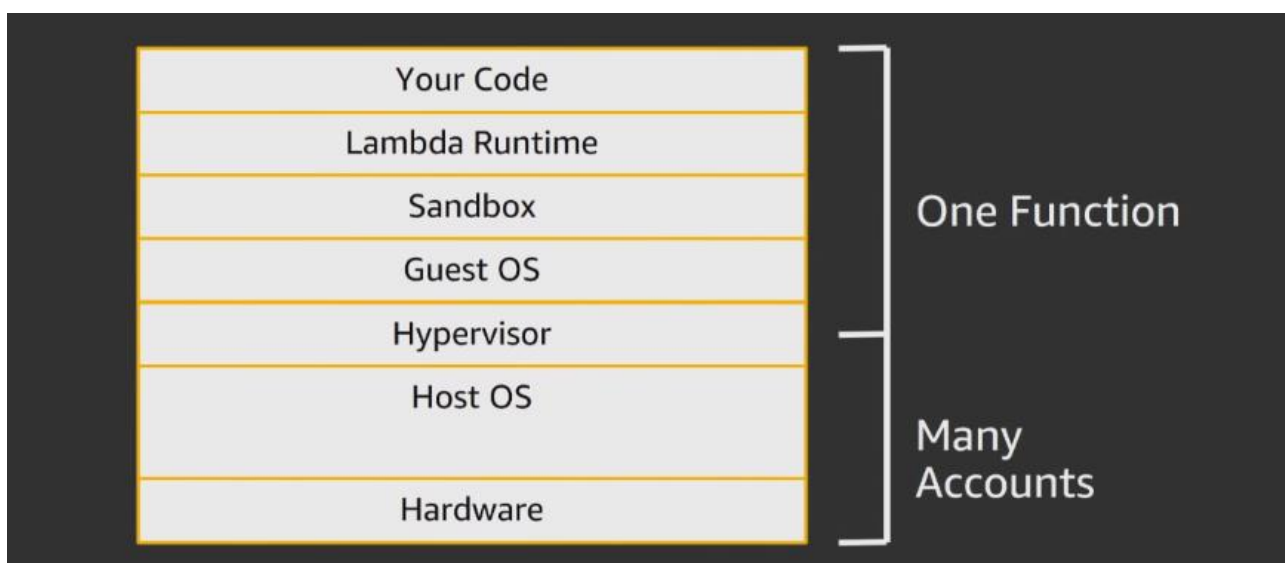


Containers are groupings of functionalities within the Linux toolbox, these include things like **cgroups**, namespaces, **seccomp**, iptables, chroot. The **cgroups** are sticky and a process added cannot get out of the cgroup. The process that your lambda function runs on is PID in its process namespace. The **seccomp** is like a firewall for your process to determine allowed **syscalls** that can be called, it helps restrict your lambda function to certain tasks.

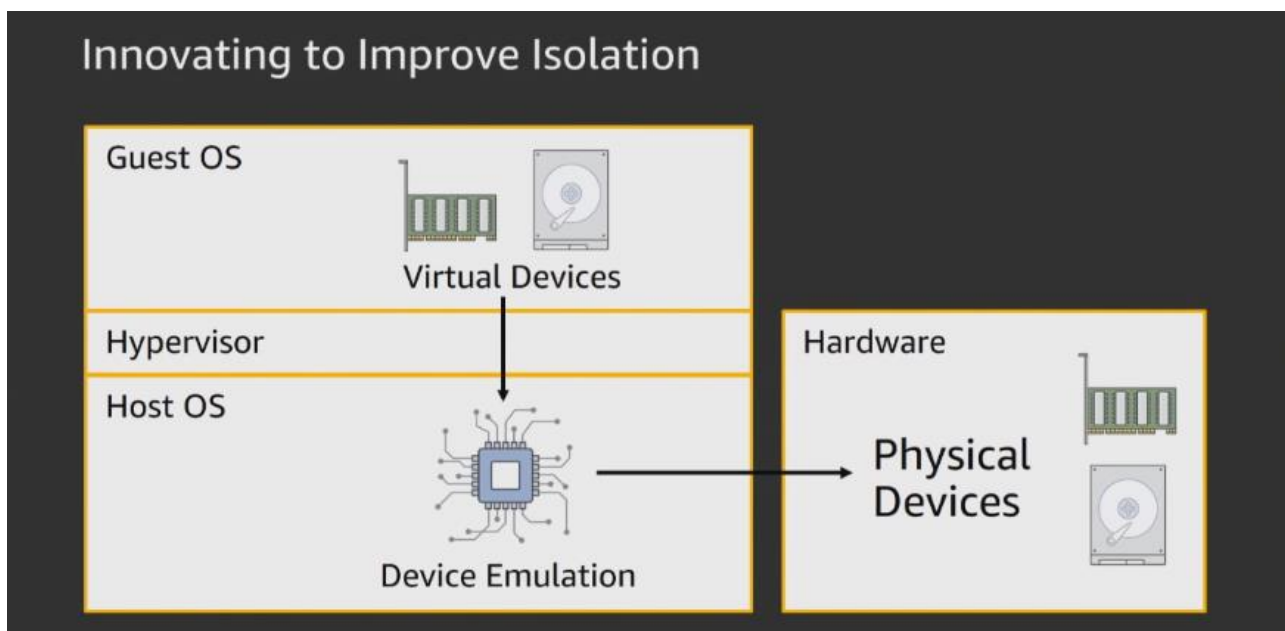




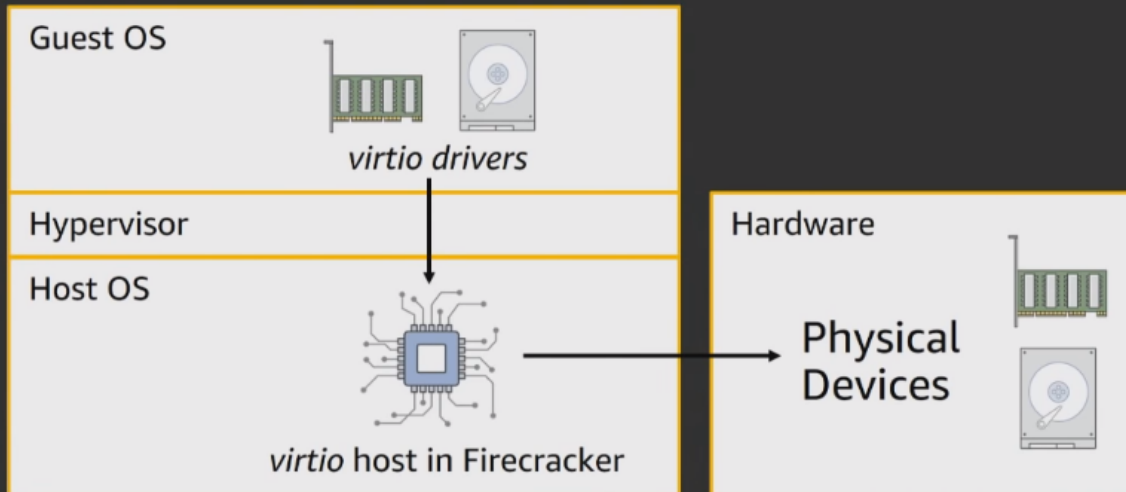
We use Firecracker to launch 100,000s of micro VMs on EC2 bare metal instances.



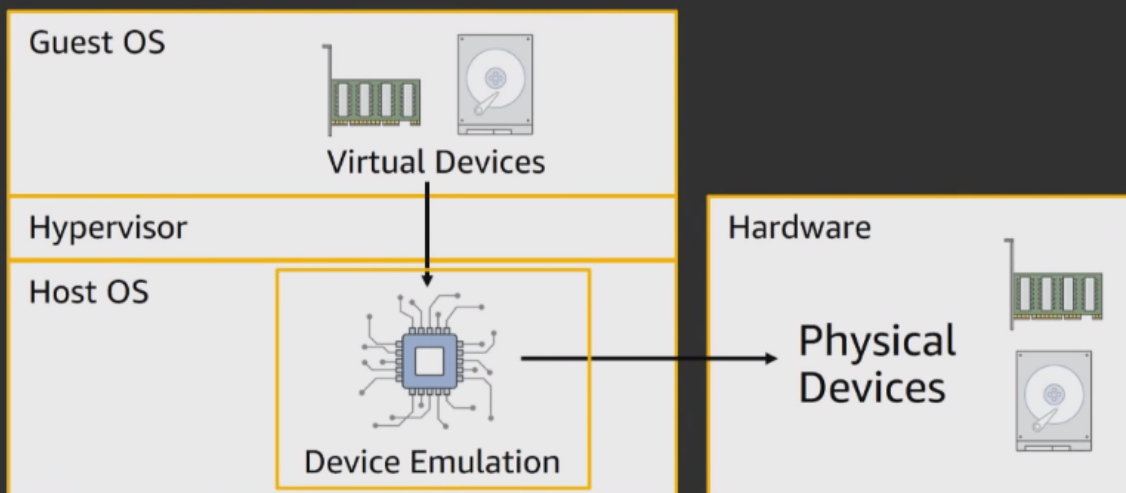
This provides strong isolation even between multiple Lambda functions when we are running in the Firecracker mode.



Innovating to Improve Isolation



Innovating to Improve Isolation



Utilization

Keeping Servers Busy


% of Resources Doing Useful Work

(vs. idle or waste)

With Lambda:
Pay only for useful work.

Inside Lambda:
Optimize To Keep Servers
Busy

Available Sandboxes For a Function



Bad: Balance The Load

60%

60%

60%

60%

60%

60%

60%

Good: Concentrate The Load

99%

99%

99%

99%

0%

0%

0%

Cache Locality

Ability to Autoscale

Bad: Pack Server With One Workload

Server

Workload

Workload

Workload

Workload

Workload

Workload


This will have very correlated load that results in same spikes or memory demands from the workloads of the same type running on the server

Better: Pack With Many Loads




Take advantage of
*Statistical
Multiplexing*


Throwing a d20




Throwing a d10, twice



Throwing a d10, ten times




Best: Placement Optimization




Pick workloads that pack together well.

Minimize contention.

Improving VPC start-up and scaling: now



Improving VPC start-up and scaling: 2019



This allows us to use 1 ENIs to use with many lambda functions

Improving VPC start-up and scaling: benefits



Faster
Scaling



Lower
Latency



Easier To
Use

Firecracker Hypervisor vs. Others

- ↓ Startup time
- ↓ Memory overhead
- = Performance
- ↑ Flexibility

Firecracker Unlocks
Higher Utilization and Scale

In Conclusion



Please complete the session
survey in the mobile app.