

ANT398

Amazon Managed Streaming for Kafka: A Fully Managed, Highly Available, and Secure Service for Apache Kafka

Damian Wylie
Principal Product Manager
Amazon Data Streaming

aws
re:Invent

© 2018, Amazon Web Services, Inc. or its affiliates. All rights reserved.



Discover the power of running Apache Kafka on a fully managed AWS service. In this session, we describe how Amazon Managed Streaming for Kafka (Amazon MSK) runs Apache Kafka clusters for you, demo Amazon MSK and a migration, show you how to get started, and walk through other important details about the new service.

Agenda

- Real-time data
- Apache Kafka
- Introducing Amazon Managed Streaming for Kafka (Amazon MSK)
- How to get started
- Comparing Amazon MSK with Amazon Kinesis Data Streams
- Demo
- Q&A

Data is produced continuously



Mobile Apps




Web Clickstream

```
[Wed Oct 11 14:32:52  
2018] [error] [client  
127.0.0.1] client  
denied by server  
configuration:  
/export/home/live/ap/h  
tdocs/test
```

Application Logs



Metering Records



IoT Sensors

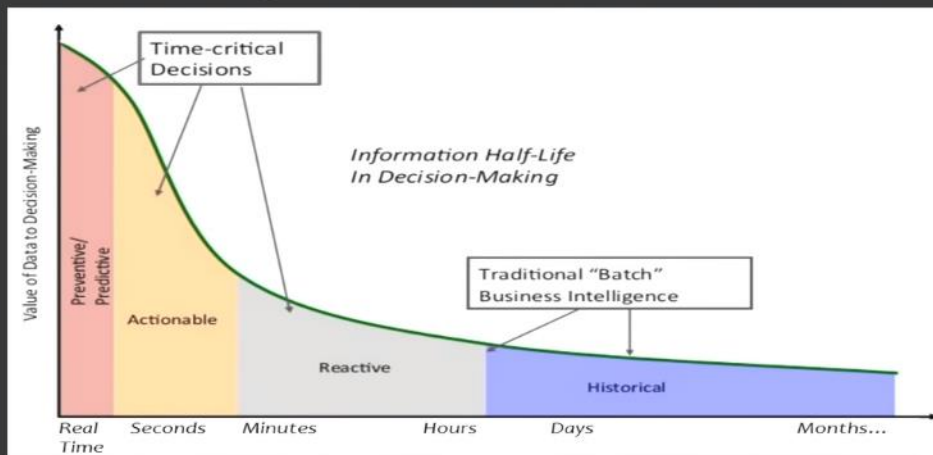


Smart Buildings

Data can be transformed continuously



The diminishing value of data over time



aws
reInvent

© 2018, Amazon Web Services, Inc. or its affiliates. All rights reserved.

aws



Apache Kafka



Apache Kafka use cases

Real-time web and log analytics

Messaging

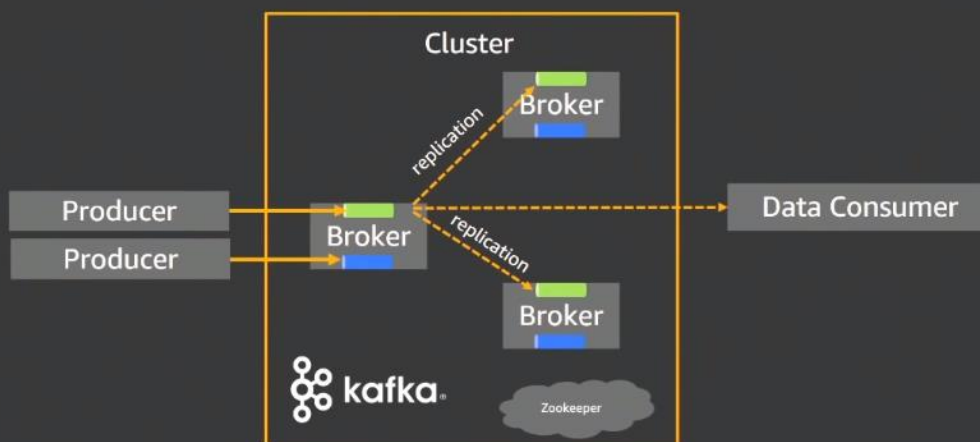
Transaction and event sourcing

Decoupled microservices

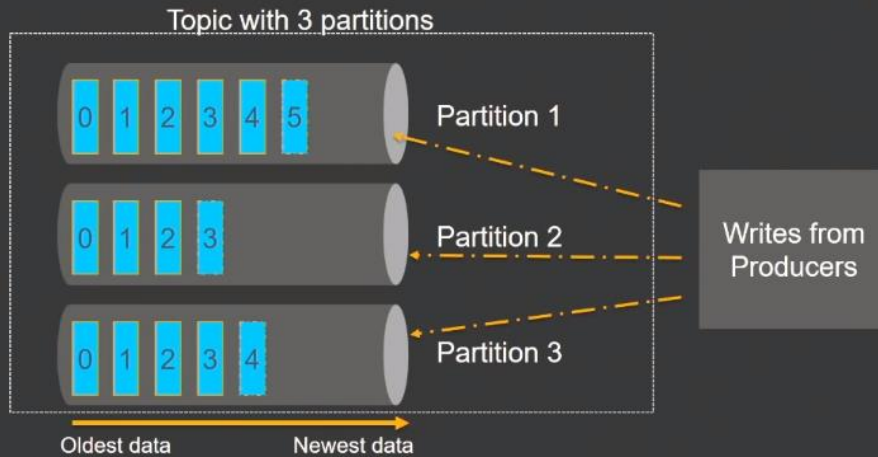
Streaming ETL



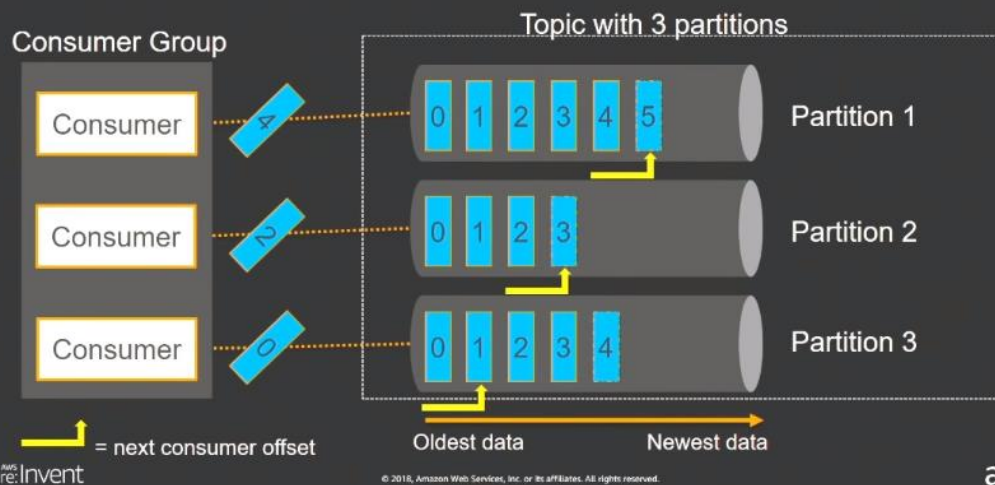
Apache Kafka Anatomy 101



Apache Kafka Anatomy – Writes to partitions



Apache Kafka Anatomy – Reads from partitions



Challenges operating Apache Kafka

Difficult to setup



Tricky to scale



Hard to achieve high availability



AWS integrations = development



No console, no visible metrics



$$f(kafka_{usage}) = \sum_{n=1}^{\infty} (SRE)$$



Introducing

Amazon Managed Streaming for Kafka (Amazon MSK)

A fully managed, highly available, and secure service for Apache Kafka



Now available in public preview in the US East (N. Virginia) Region

aws
re:Invent

© 2018, Amazon Web Services, Inc. or its affiliates. All rights reserved.

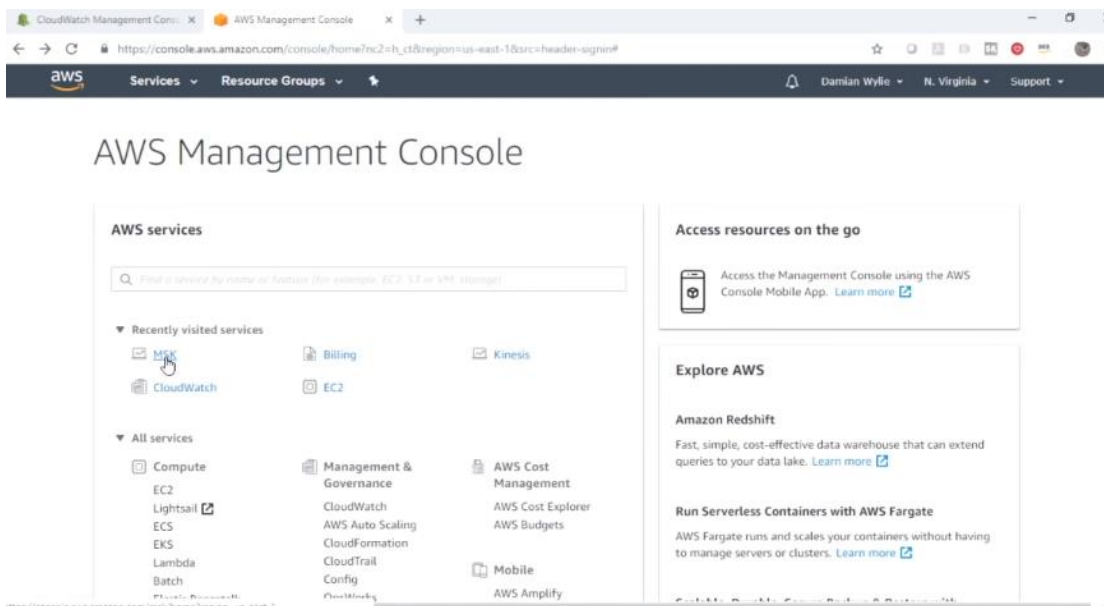


Getting started with Amazon MSK is easy



- Fully compatible with Apache Kafka v1.1.1
- AWS Management Console and AWS API for provisioning
- Clusters are setup automatically
- Provision Apache Kafka brokers and storage
- Create and tear down clusters on-demand

Create cluster,
producer/consumer
demo



CloudWatch Management Console | Amazon Managed Streaming for Kafka

Services | Resource Groups

Amazon Managed Streaming for Kafka

Clusters

MSK > Clusters

Clusters (1)

Use the CLI to retrieve cluster and broker connection details. [Learn more](#)

Search clusters

Name	Status	Apache Kafka version	Brokers per Availability Zone	Availability Zones	Creation time
DamiansDemoCluster	Active	1.1.1	1	3	Tue, 27 Nov 2018 05:52:57 GMT

Feedback | English (US) | © 2008 - 2018, Amazon Web Services, Inc. or its affiliates. All rights reserved. | Privacy Policy | Terms of Use

CloudWatch Management Console | Amazon Managed Streaming for Kafka

Services | Resource Groups

MSK > Clusters > Create cluster

Create Kafka cluster

General

Cluster name
You can't change it after you create the cluster.
DamianReinventDemo
The name must be unique and can have a maximum of 64 characters.

VPC
Defines the virtual networking environment for this cluster. You can't change this setting after you create the cluster.
Choose VPC

Apache Kafka version
Software version for deployed Kafka brokers.
1.1.1

Availability Zones
Specify where Kafka brokers are deployed in your VPC.

Feedback | English (US) | © 2008 - 2018, Amazon Web Services, Inc. or its affiliates. All rights reserved. | Privacy Policy | Terms of Use

CloudWatch Management Console | Amazon Managed Streaming for Kafka

Services | Resource Groups

MSK > Clusters > Create cluster

9 total brokers will be deployed in your cluster, distributed evenly across your 3 availability zones.

Advanced details
To customize advanced settings, use the CLI to create a cluster. [Learn more](#)

Broker instance type
kafka.m5.large

EBS storage volume
1000 GiB

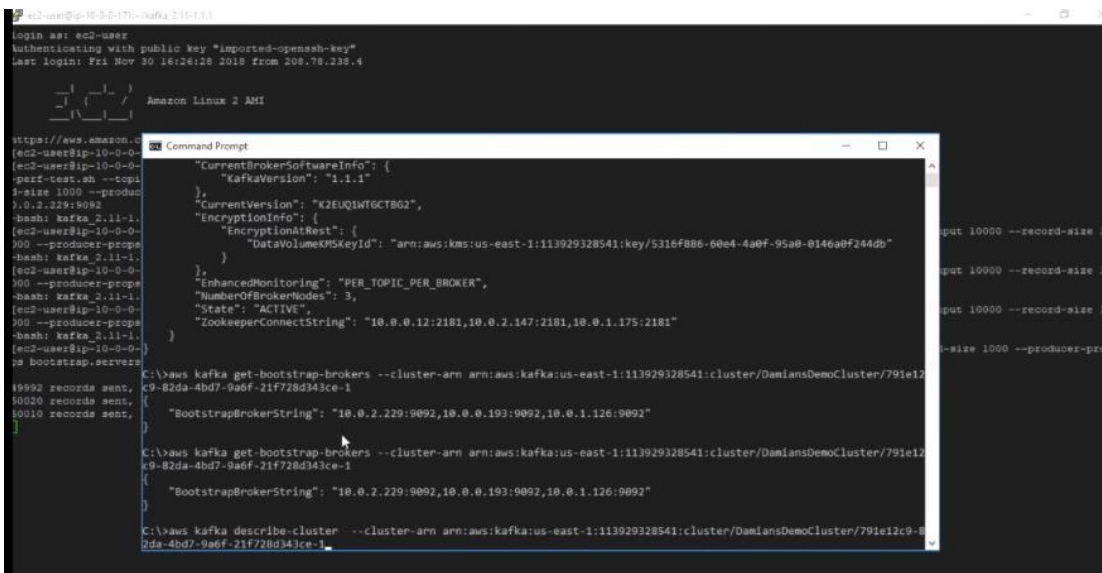
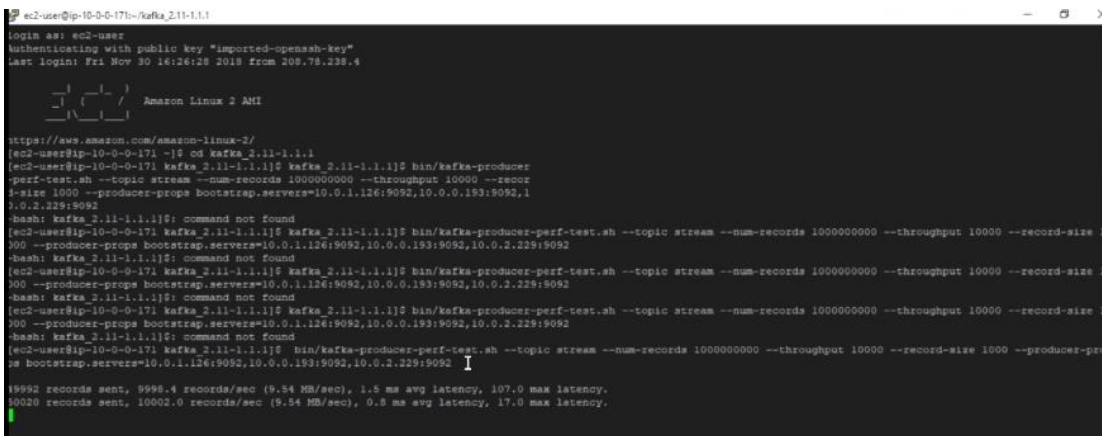
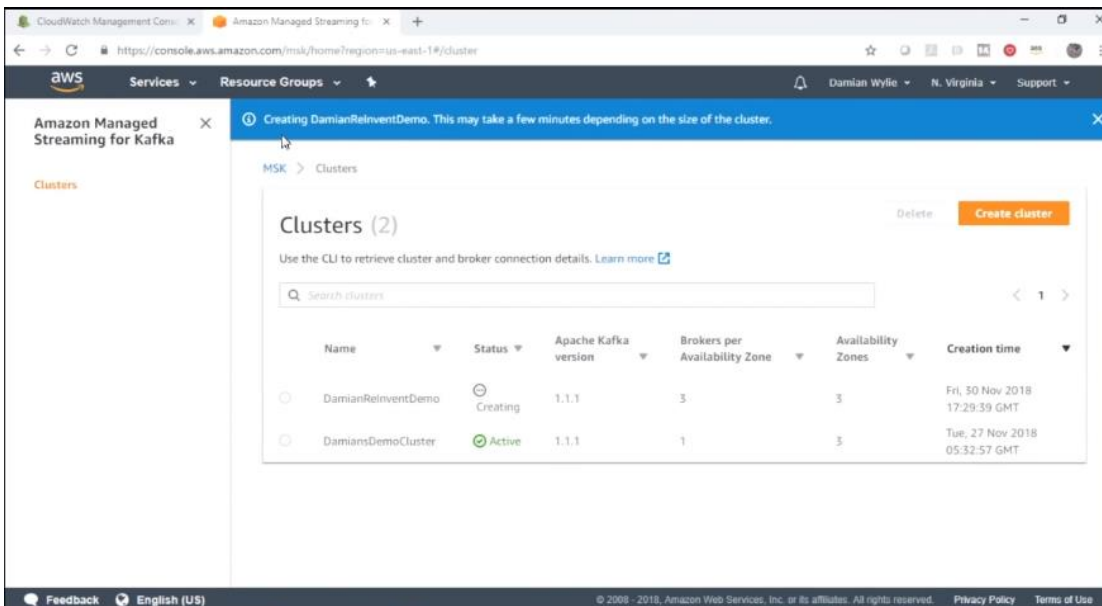
Customer master key
Default

Apache Kafka configuration
Default

Monitoring
Default

Cancel | Create cluster

Feedback | English (US) | © 2008 - 2018, Amazon Web Services, Inc. or its affiliates. All rights reserved. | Privacy Policy | Terms of Use



```
ec2-user@ip-10-0-0-171:~/kafka-2.11-1.1.1
login as: ec2-user
Authenticating with public key "imported-openssh-key"
Last login: Fri Nov 30 16:26:28 2018 from 208.78.238.4

┌───┐ ┌───┐
└───┘ └───┘ Amazon Linux 2 AMI

https://aws.amazon.com/amazon-linux-2/
[ec2-user@ip-10-0-0-171 ~]$ cd kafka_2.11-1.1.1
[ec2-user@ip-10-0-0-171 kafka_2.11-1.1.1]$ bin/kafka-producer-perf-test.sh --topic stream --num-records 1000000000 --throughput 10000 --record-size 1000 --producer-props bootstrap.servers=10.0.1.126:9092,10.0.0.193:9092,10.0.2.229:9092
-bash: kafka_2.11-1.1.1: command not found
[ec2-user@ip-10-0-0-171 kafka_2.11-1.1.1]$ bin/kafka-producer-perf-test.sh --topic stream --num-records 1000000000 --throughput 10000 --record-size 1000 --producer-props bootstrap.servers=10.0.1.126:9092,10.0.0.193:9092,10.0.2.229:9092
-bash: kafka_2.11-1.1.1: command not found
[ec2-user@ip-10-0-0-171 kafka_2.11-1.1.1]$ bin/kafka-producer-perf-test.sh --topic stream --num-records 1000000000 --throughput 10000 --record-size 1000 --producer-props bootstrap.servers=10.0.1.126:9092,10.0.0.193:9092,10.0.2.229:9092
-bash: kafka_2.11-1.1.1: command not found
[ec2-user@ip-10-0-0-171 kafka_2.11-1.1.1]$ bin/kafka-producer-perf-test.sh --topic stream --num-records 1000000000 --throughput 10000 --record-size 1000 --producer-props bootstrap.servers=10.0.1.126:9092,10.0.0.193:9092,10.0.2.229:9092
19992 records sent, 9998.4 records/sec (9.94 MB/sec), 1.5 ms avg latency, 107.0 max latency.
10020 records sent, 10002.0 records/sec (9.94 MB/sec), 0.8 ms avg latency, 17.0 max latency.
10010 records sent, 10002.0 records/sec (9.94 MB/sec), 0.8 ms avg latency, 17.0 max latency.
10010 records sent, 10002.0 records/sec (9.94 MB/sec), 0.8 ms avg latency, 17.0 max latency.
10010 records sent, 10000.0 records/sec (9.94 MB/sec), 0.8 ms avg latency, 20.0 max latency.
]
```

```
ec2-user@ip-10-0-0-171:~/kafka-2.11-1.1.1
login as: ec2-user
Authenticating with public key "imported-openssh-key"
Last login: Fri Nov 30 16:26:28 2018 from 208.78.238.4

┌───┐ ┌───┐
└───┘ └───┘ Amazon Linux 2 AMI

https://aws.amazon.com/amazon-linux-2/
[ec2-user@ip-10-0-0-171 ~]$ cd kafka_2.11-1.1.1
[ec2-user@ip-10-0-0-171 kafka_2.11-1.1.1]$ bin/kafka-producer-perf-test.sh --topic stream --num-records 1000000000 --throughput 10000 --record-size 1000 --producer-props bootstrap.servers=10.0.1.126:9092,10.0.0.193:9092,10.0.2.229:9092
-bash: kafka_2.11-1.1.1: command not found
[ec2-user@ip-10-0-0-171 kafka_2.11-1.1.1]$ bin/kafka-producer-perf-test.sh --topic stream --num-records 1000000000 --throughput 10000 --record-size 1000 --producer-props bootstrap.servers=10.0.1.126:9092,10.0.0.193:9092,10.0.2.229:9092
-bash: kafka_2.11-1.1.1: command not found
[ec2-user@ip-10-0-0-171 kafka_2.11-1.1.1]$ bin/kafka-producer-perf-test.sh --topic stream --num-records 1000000000 --throughput 10000 --record-size 1000 --producer-props bootstrap.servers=10.0.1.126:9092,10.0.0.193:9092,10.0.2.229:9092
-bash: kafka_2.11-1.1.1: command not found
[ec2-user@ip-10-0-0-171 kafka_2.11-1.1.1]$ bin/kafka-producer-perf-test.sh --topic stream --num-records 1000000000 --throughput 10000 --record-size 1000 --producer-props bootstrap.servers=10.0.1.126:9092,10.0.0.193:9092,10.0.2.229:9092
19992 records sent, 9998.4 records/sec (9.94 MB/sec), 1.5 ms avg latency, 107.0 max latency.
10020 records sent, 10002.0 records/sec (9.94 MB/sec), 0.8 ms avg latency, 17.0 max latency.
10010 records sent, 10002.0 records/sec (9.94 MB/sec), 0.8 ms avg latency, 17.0 max latency.
10010 records sent, 10002.0 records/sec (9.94 MB/sec), 0.8 ms avg latency, 17.0 max latency.
10010 records sent, 10000.0 records/sec (9.94 MB/sec), 0.8 ms avg latency, 20.0 max latency.
]
```

```
Command Prompt: aws kafka describe-cluster --cluster-arn arn:aws:kafka:us-east-1:113929328541:cluster/DamiansDemoCluster/791e12c9-82da-4bd7-9a6f-21f728d343ce-1
{
  "kafkaVersion": "1.1.1",
  "currentVersion": "K2EUQ1MTGCTBG2",
  "encryptionInfo": {
    "EncryptionAtRest": {
      "DataVolumeKMSKeyId": "arn:aws:kms:us-east-1:113929328541:key/5316f8b6-60e4-4a0f-95a0-0146a0f244db"
    }
  },
  "enhancedMonitoring": "PER_TOPIC_PER_BROKER",
  "numberOfBrokerNodes": 3,
  "state": "ACTIVE",
  "zookeeperConnectString": "10.0.0.12:2181,10.0.2.147:2181,10.0.1.175:2181"
}

C:\>aws kafka get-bootstrap-brokers --cluster-arn arn:aws:kafka:us-east-1:113929328541:cluster/DamiansDemoCluster/791e12c9-82da-4bd7-9a6f-21f728d343ce-1
[
  {
    "bootstrapBrokerString": "10.0.2.229:9092,10.0.0.193:9092,10.0.1.126:9092"
  }
]

C:\>aws kafka get-bootstrap-brokers --cluster-arn arn:aws:kafka:us-east-1:113929328541:cluster/DamiansDemoCluster/791e12c9-82da-4bd7-9a6f-21f728d343ce-1
[
  {
    "bootstrapBrokerString": "10.0.2.229:9092,10.0.0.193:9092,10.0.1.126:9092"
  }
]

C:\>aws kafka describe-cluster --cluster-arn arn:aws:kafka:us-east-1:113929328541:cluster/DamiansDemoCluster/791e12c9-82da-4bd7-9a6f-21f728d343ce-1
```

```
ec2-user@ip-10-0-0-171:~/kafka-2.11-1.1.1
login as: ec2-user
Authenticating with public key "imported-openssh-key"
Last login: Fri Nov 30 16:26:28 2018 from 208.78.238.4

┌───┐ ┌───┐
└───┘ └───┘ Amazon Linux 2 AMI

https://aws.amazon.com/amazon-linux-2/
[ec2-user@ip-10-0-0-171 ~]$ cd kafka_2.11-1.1.1
[ec2-user@ip-10-0-0-171 kafka_2.11-1.1.1]$ bin/kafka-producer-perf-test.sh --topic stream --num-records 1000000000 --throughput 10000 --record-size 1000 --producer-props bootstrap.servers=10.0.1.126:9092,10.0.0.193:9092,10.0.2.229:9092
-bash: kafka_2.11-1.1.1: command not found
[ec2-user@ip-10-0-0-171 kafka_2.11-1.1.1]$ bin/kafka-producer-perf-test.sh --topic stream --num-records 1000000000 --throughput 10000 --record-size 1000 --producer-props bootstrap.servers=10.0.1.126:9092,10.0.0.193:9092,10.0.2.229:9092
-bash: kafka_2.11-1.1.1: command not found
[ec2-user@ip-10-0-0-171 kafka_2.11-1.1.1]$ bin/kafka-producer-perf-test.sh --topic stream --num-records 1000000000 --throughput 10000 --record-size 1000 --producer-props bootstrap.servers=10.0.1.126:9092,10.0.0.193:9092,10.0.2.229:9092
-bash: kafka_2.11-1.1.1: command not found
[ec2-user@ip-10-0-0-171 kafka_2.11-1.1.1]$ bin/kafka-producer-perf-test.sh --topic stream --num-records 1000000000 --throughput 10000 --record-size 1000 --producer-props bootstrap.servers=10.0.1.126:9092,10.0.0.193:9092,10.0.2.229:9092
19992 records sent, 9998.4 records/sec (9.94 MB/sec), 1.5 ms avg latency, 107.0 max latency.
10020 records sent, 10002.0 records/sec (9.94 MB/sec), 0.8 ms avg latency, 17.0 max latency.
10010 records sent, 10002.0 records/sec (9.94 MB/sec), 0.8 ms avg latency, 17.0 max latency.
10010 records sent, 10002.0 records/sec (9.94 MB/sec), 0.8 ms avg latency, 17.0 max latency.
10010 records sent, 10000.0 records/sec (9.94 MB/sec), 0.8 ms avg latency, 20.0 max latency.
]
```

```
Select Command Prompt
{
  "securityGroups": {
    "sg-090b6221d3bd2e7c1"
  },
  "storageInfo": {
    "efsStorageInfo": {
      "volumeSize": 1000
    }
  },
  "clusterArn": "arn:aws:kafka:us-east-1:113929328541:cluster/DamiansDemoCluster/791e12c9-82da-4bd7-9a6f-21f728d343ce-1",
  "clusterName": "DamiansDemoCluster",
  "creationTime": "2018-11-27T05:32:57.792Z",
  "currentBrokerSoftwareInfo": {
    "kafkaVersion": "1.1.1"
  },
  "currentVersion": "K2EUQ1MTGCTBG2",
  "encryptionInfo": {
    "EncryptionAtRest": {
      "DataVolumeKMSKeyId": "arn:aws:kms:us-east-1:113929328541:key/5316f8b6-60e4-4a0f-95a0-0146a0f244db"
    }
  },
  "enhancedMonitoring": "PER_TOPIC_PER_BROKER",
  "numberOfBrokerNodes": 3,
  "state": "ACTIVE",
  "zookeeperConnectString": "10.0.0.12:2181,10.0.2.147:2181,10.0.1.175:2181"
}
```


Automation drives higher availability



@ Preview

- Cluster lifecycle is fully automated
 - Brokers and Apache Zookeeper nodes auto-heal
 - IPs remain intact
 - Patches are applied automatically

@ GA

- Service level agreement (SLA)
- Apache Kafka version upgrades

Where's Apache Zookeeper?



- Apache Zookeeper is under the hood, highly available, and included with each cluster at no additional cost



Scalability and configurability



@ GA

- Scale a cluster
 - Horizontally (add more of the same)
 - Vertically (add larger brokers) scale a cluster
- Supports Apache Kafka partition reassignment tooling
- Define custom cluster configurations
- Auto scale storage

Deeply integrated with AWS services

@ Preview

- Amazon Virtual Private Cloud (Amazon VPC) for network isolation
- AWS Key Management Service (AWS KMS) for at-rest encryption
- AWS Identity and Access Management (IAM) for control-plane API control
- Amazon CloudWatch for Apache Kafka broker, topic, and ZK metrics
- Amazon Elastic Compute Cloud (Amazon EC2) M5 instances as brokers
- Amazon EBS GP2 broker storage
- Offered in the US-East (N. Virginia) AWS Region

@ GA

- Tagging
- AWS CloudTrail
- AWS CloudFormation
- Offered worldwide



What Amazon MSK does for you

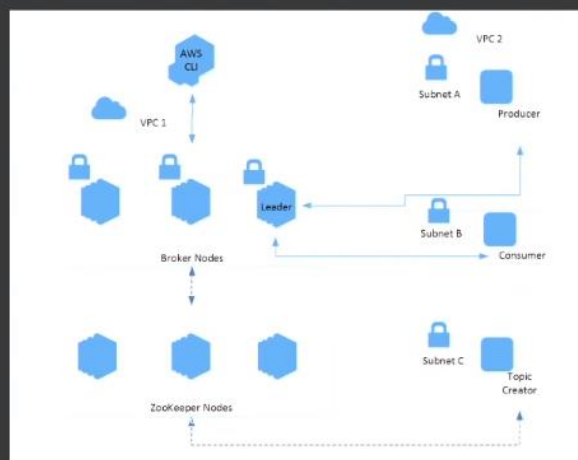
- Makes Apache Kafka more accessible to your organization

$$f(kafka_{usage}) = \sum_{n=1}^{\infty} Streaming\ Apps$$

- Drives best practices through design, defaults, and automation
- Allows developers to focus more on app development, less on infrastructure management
- Amazon MSK is committed to improving open-source Apache Kafka



How it works



Amazon MSK defaults

Config	Default Setting
offsets.topic.replication.factor	3
transaction.state.log.replication.factor	3
transaction.state.log.min.isr	2
auto.create.topics.enable	False
default.replication.factor	3
min.insync.replicas	2
unclean.leader.election.enable	True
auto.leader.rebalance.enable	True
authorizer.class.name	kafka.security.auth.SimpleAclAuthorizer
group.initial.rebalance.delay.ms	3000
log.retention.hours	168

How pricing works




- On-demand, hourly pricing prorated to the second
- Broker and storage pricing
 - Broker pricing starts with kafka.m5.large @ \$0.21/hr
 - Storage pricing is \$0.10 per GB-month

Comparing Amazon Kinesis Data Streams to MSK



Amazon Kinesis Data Streams

Stream with 3 shards




re:invent



Amazon MSK

Topic with 3 partitions



aws

© 2018, Amazon Web Services, Inc. or its affiliates. All rights reserved.

Comparing Amazon Kinesis Data Streams to MSK



Amazon Kinesis Data Streams

- AWS API experience
- Throughput provisioning model
- Seamless scaling
- Typically lower costs
- Deep AWS integrations



Amazon MSK

- Open-source compatibility
- Strong third-party tooling
- Cluster provisioning model
- Apache Kafka scaling isn't seamless to clients
- Raw performance

aws
re:Invent

© 2018, Amazon Web Services, Inc. or its affiliates. All rights reserved.



Why add Amazon MSK?

Thank you!

Damian Wylie
wylled@amazon.com
Follow @DamianWylie on Twitter
for live updates

aws
re:Invent

© 2018, Amazon Web Services, Inc. or its affiliates. All rights reserved.

