DAT320

Becoming a Nimble Giant: How Amazon DynamoDB Serves Nike at Scale

Zack Owens
Principal Architect
Nike

Adam Farrell
Lead Software Engineer
Nike

aws re:Invent

aws

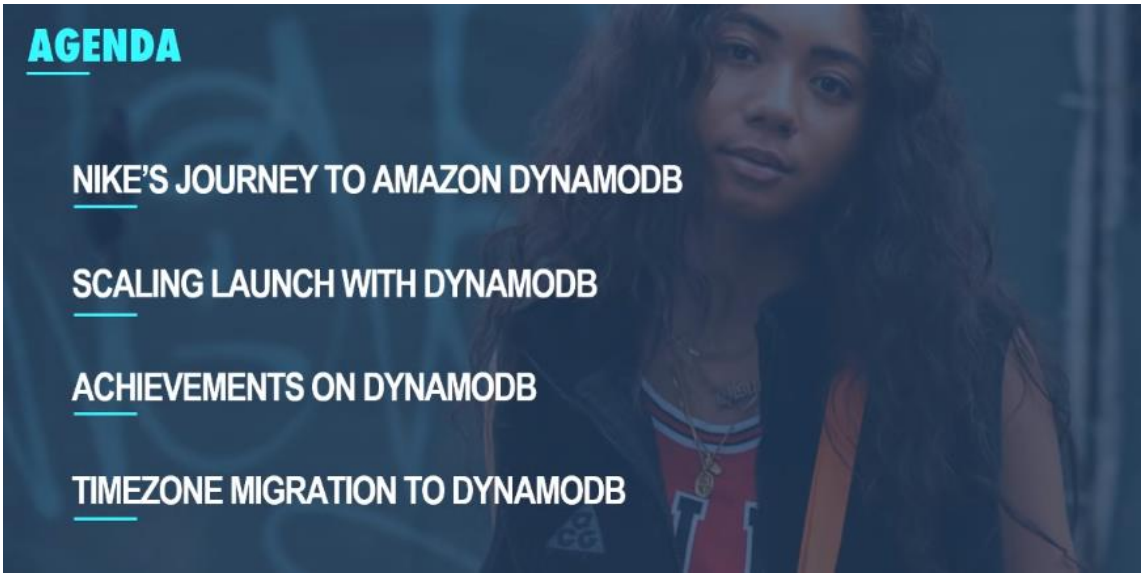

ARCHITECTURE

MICROSERVICE ARCHITECTURE

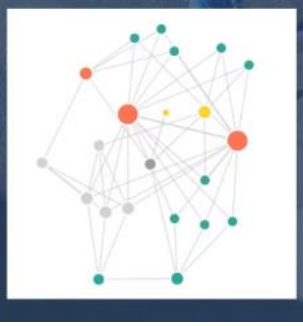

SERVED FROM
AMAZON WEB SERVICES (AWS)

NIKE'S JOURNEY TO AWS

PRE-2013	2013	2015
DATA CENTER MONOLITHS	HYBRID CLOUD	CLOUD NATIVE
PROBLEMS SCALING	MANUAL DEPLOYMENTS	AUTOMATED CI/CD
LARGE DEPLOYS		

CORE PRINCIPLES

Nike Digital Engineering | JDE

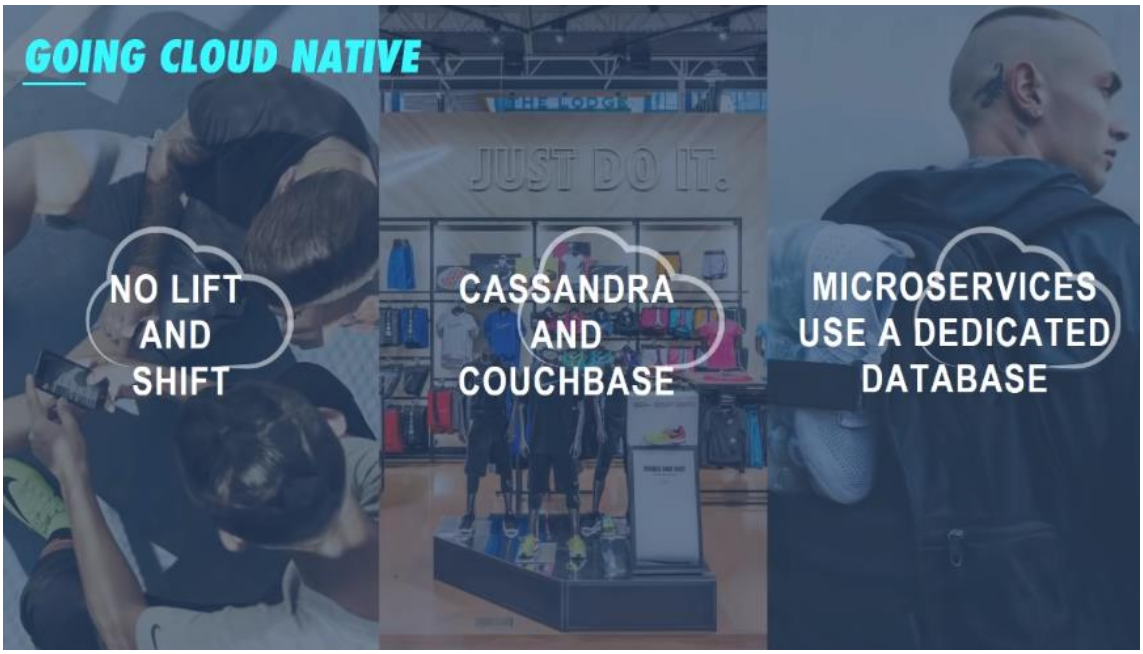
GOING CLOUD NATIVE
JUST DO IT.
NO LIFT AND SHIFT
CASSANDRA AND COUCHBASE
MICROSERVICES USE A DEDICATED DATABASE

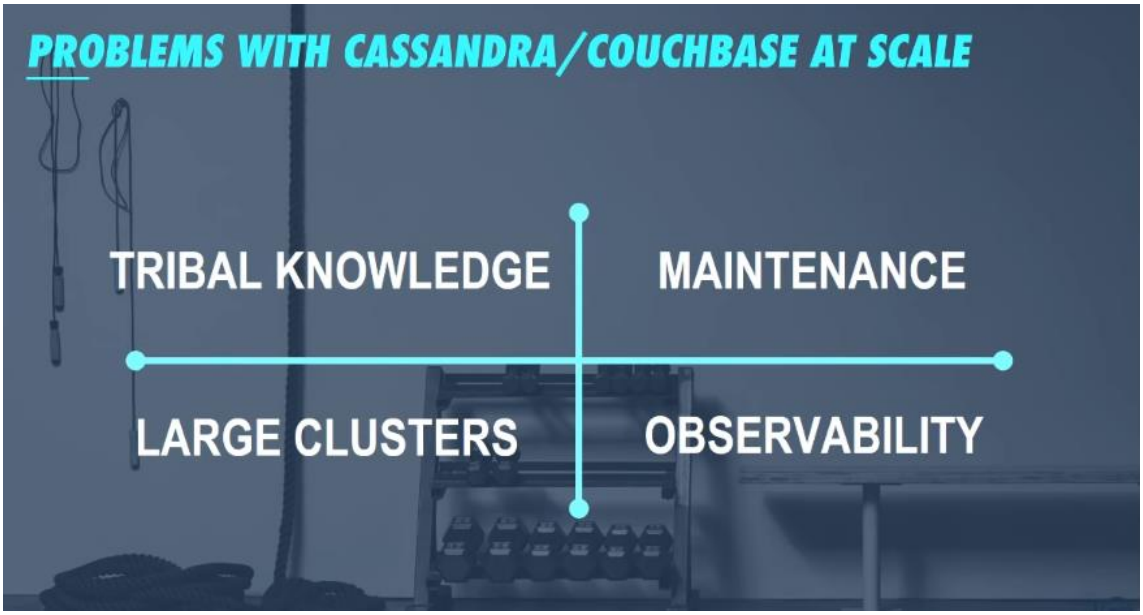
PROBLEMS WITH CASSANDRA/COUCHBASE AT SCALE
TRIBAL KNOWLEDGE
MAINTENANCE
LARGE CLUSTERS
OBSERVABILITY


SCALING LAUNCH
WITH DYNAMODB

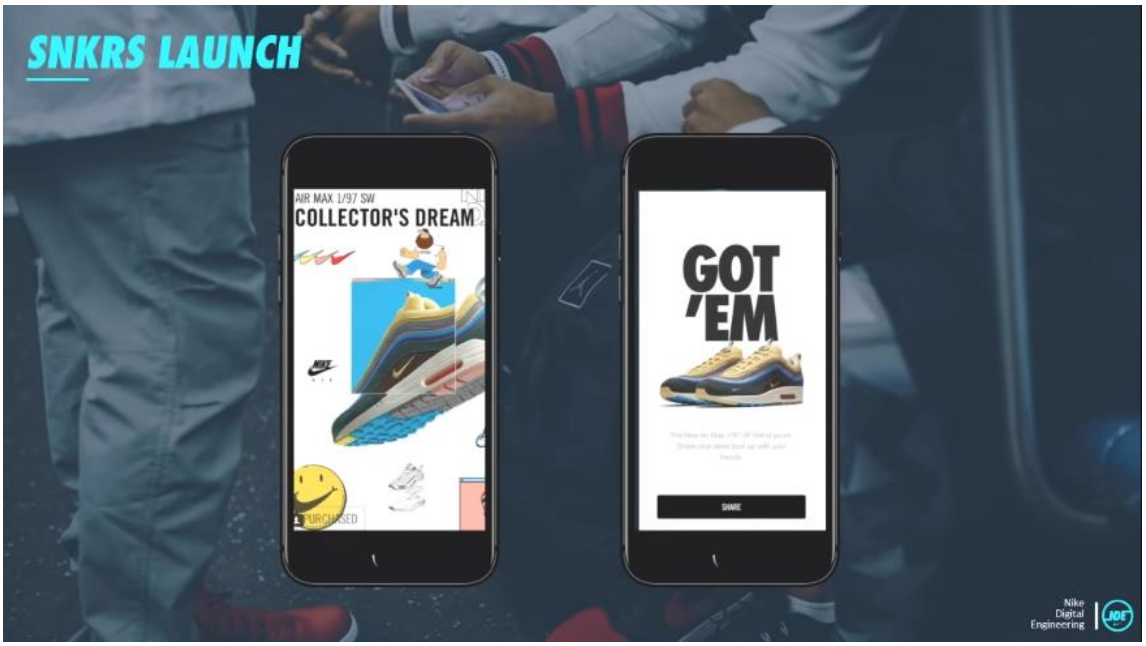
SNKRS LAUNCH
AIR MAX 1/97 SW
COLLECTOR'S DREAM
GOT 'EM
SHARE
Nike Digital Engineering
JDE


LAUNCH IN THE CLOUD
HIGH-DEMAND PRODUCT
HIGH-THROUGHPUT TRAFFIC

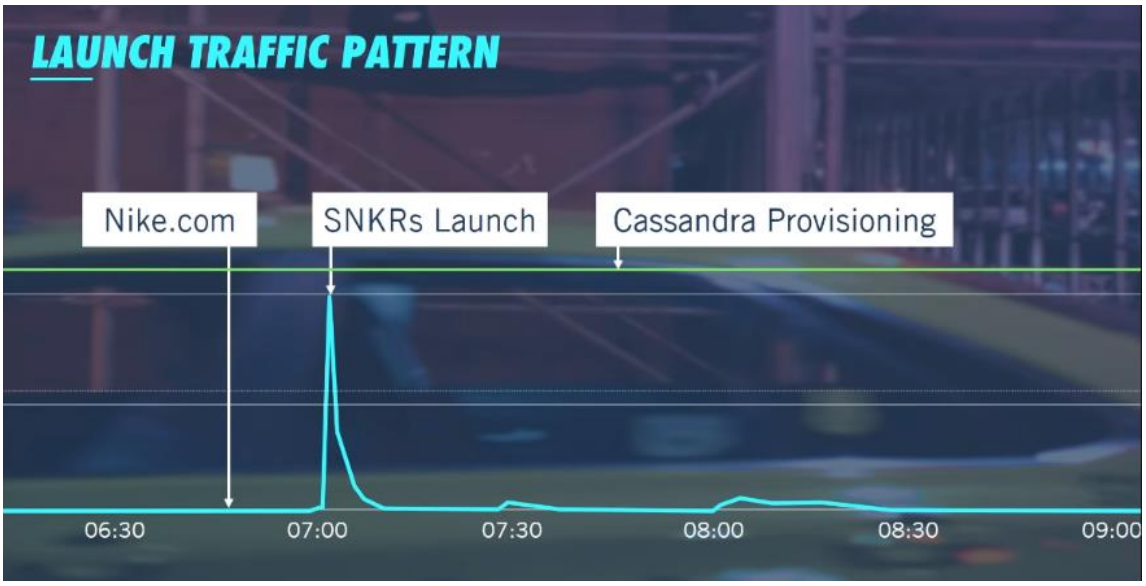
LAUNCH TRAFFIC PATTERN
Nike.com
SNKRs Launch
Cassandra Provisioning
06:30
07:00
07:30
08:00
08:30
09:00

The DynamoDB feature set is extensive and solves a lot of our problems. **GSI's allows you to query data using a hash and range key that is different from your table**, this allows multiple different, efficient query patterns to be possible. **GSI's can also be applied after the table has been created and has data inside it**, you can **add new indexes with a single API call**. This allows us to update our data models as needed.

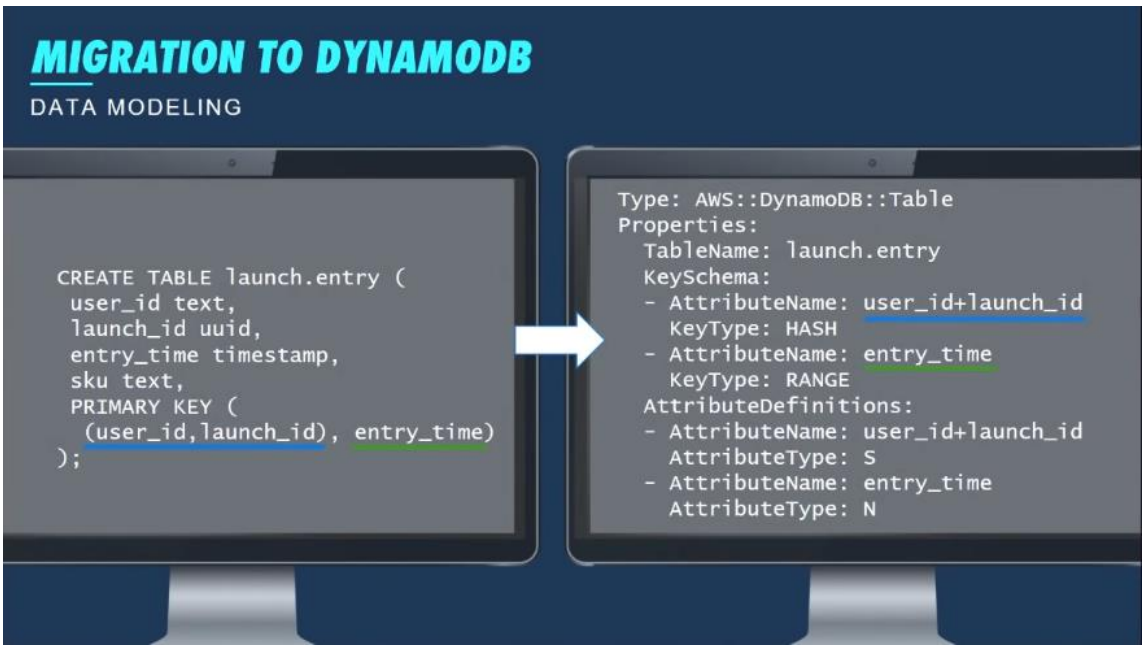

In DynamoDB on the right, we can use 2 keys in our query by using string concatenation to create a unique key like **user_id+launch_id**.

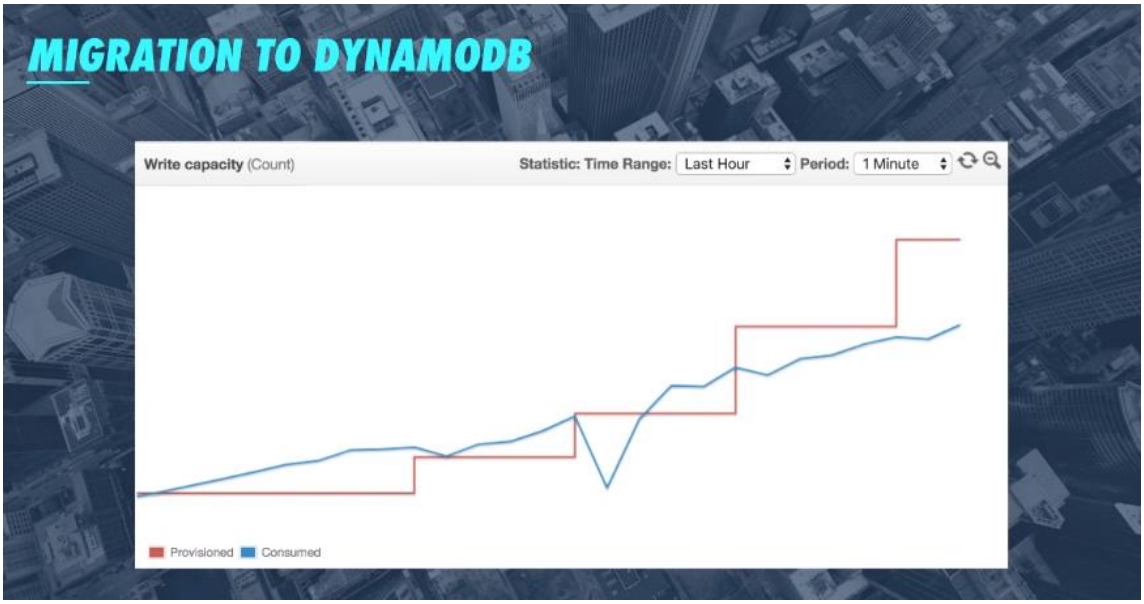

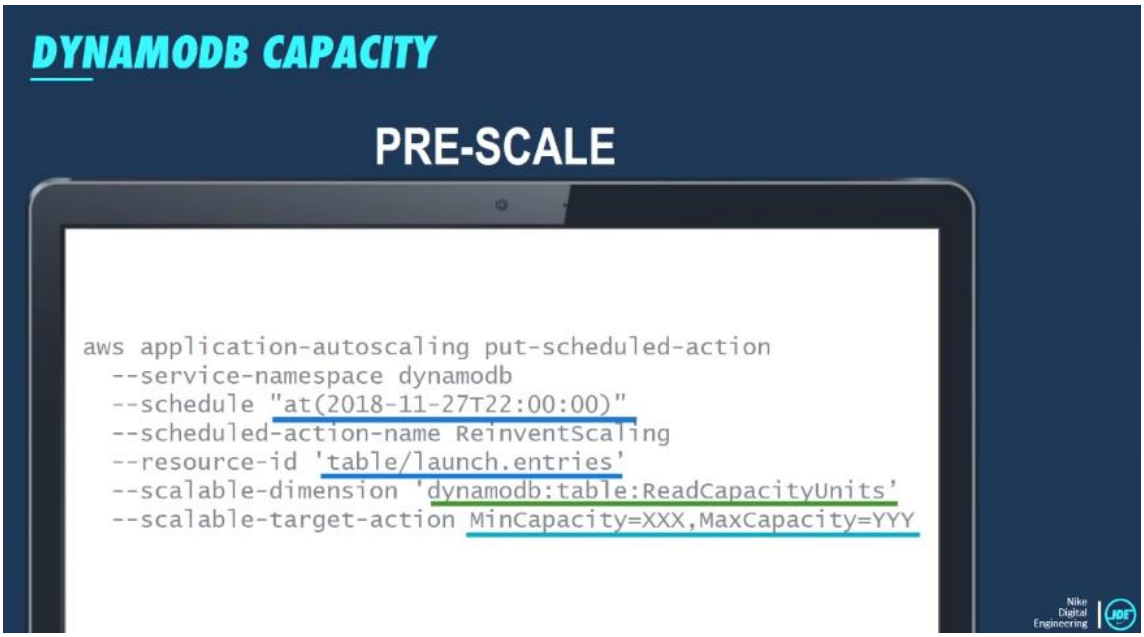

We also use pre-scaling because auto-scaling is not fast enough to respond to our user traffic spike

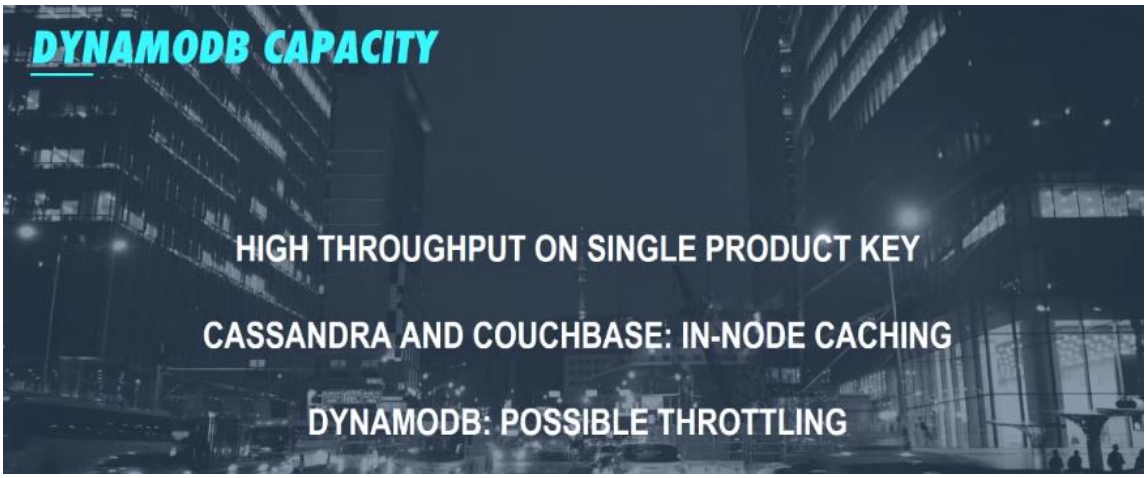

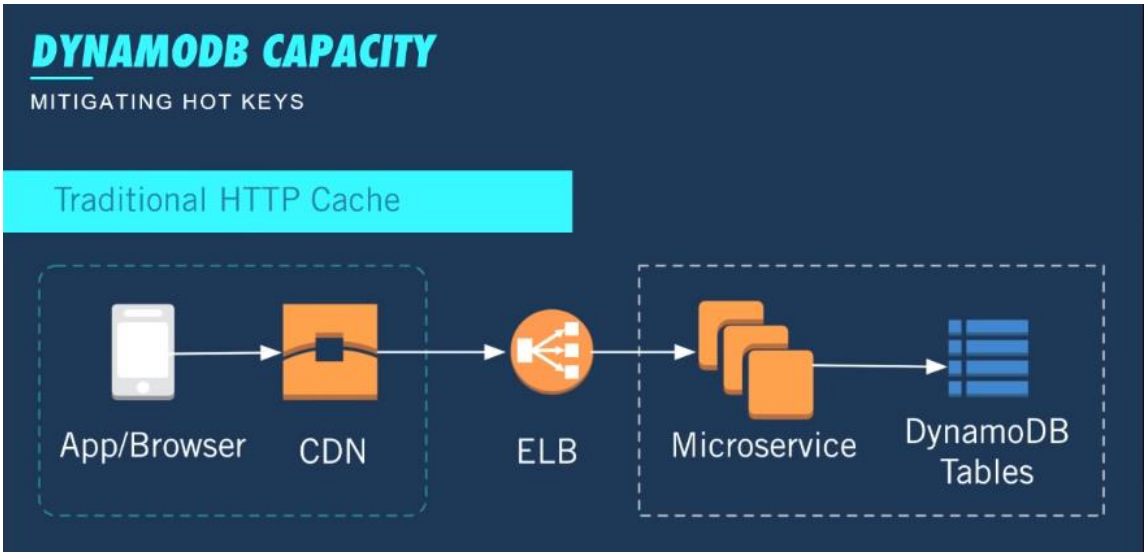

The first place we cache is for full HTTP responses at the CDN layer on different data centers around the world

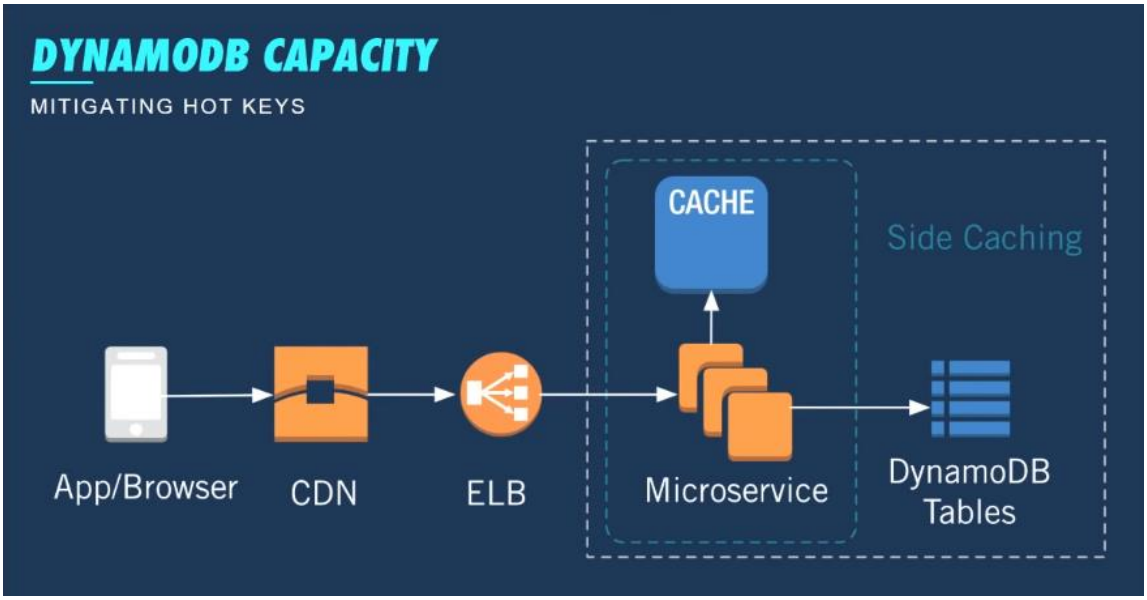

We then use a side-cache using memcache to keep some data in-memory for each microservice

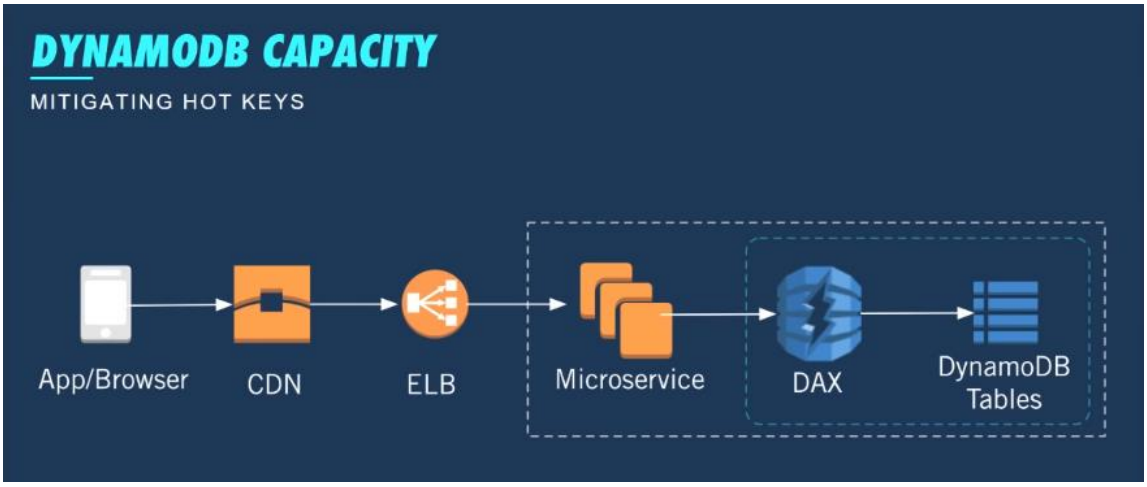

DAX sits in front of your DynamoDB and can do the caching for you effectively

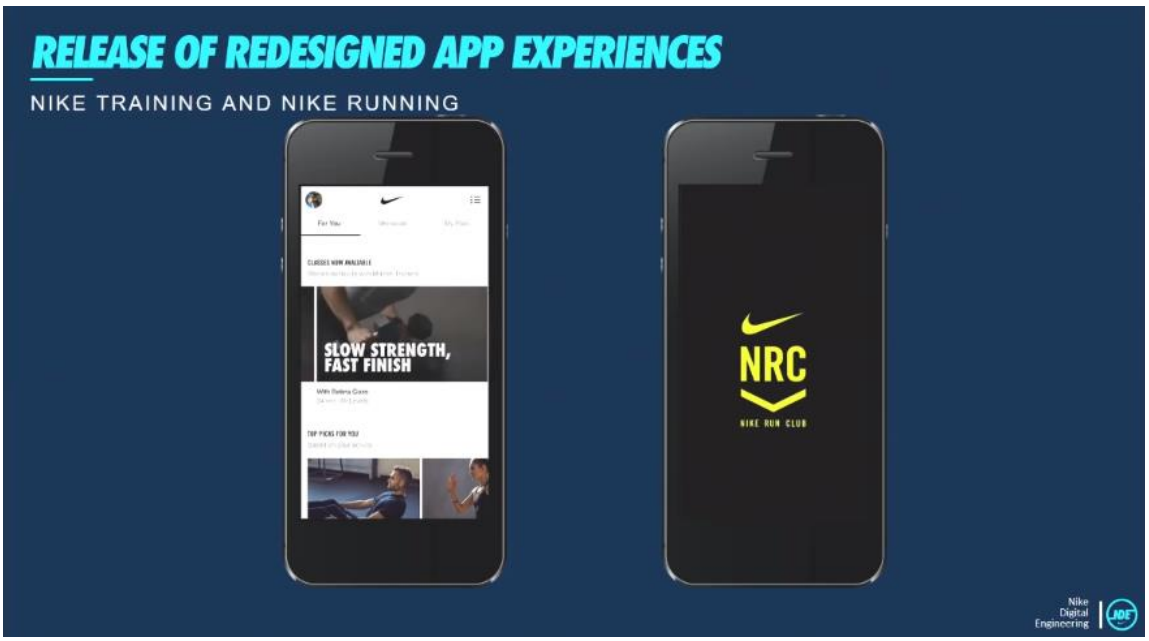

RELEASE OF REDESIGNED APP EXPERIENCES

NIKE TRAINING AND NIKE RUNNING

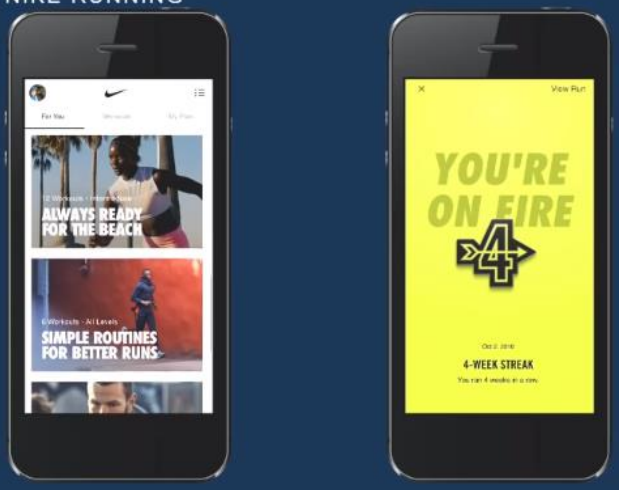

Nike Digital Engineering

ACHIEVEMENTS WAS NOT IN THE CLOUD IN FIRST RELEASE

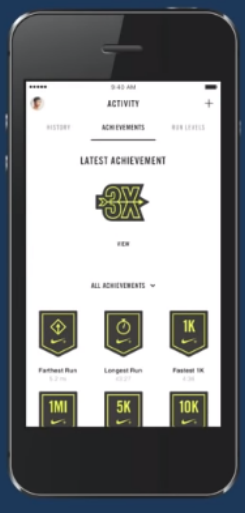

ACHIEVEMENTS RUNS WITH DYNAMODB

- SCALABLE
- COST EFFICIENT
- DURABLE
- OBSERVABLE
- ELASTIC
- CONDITIONAL UPDATES

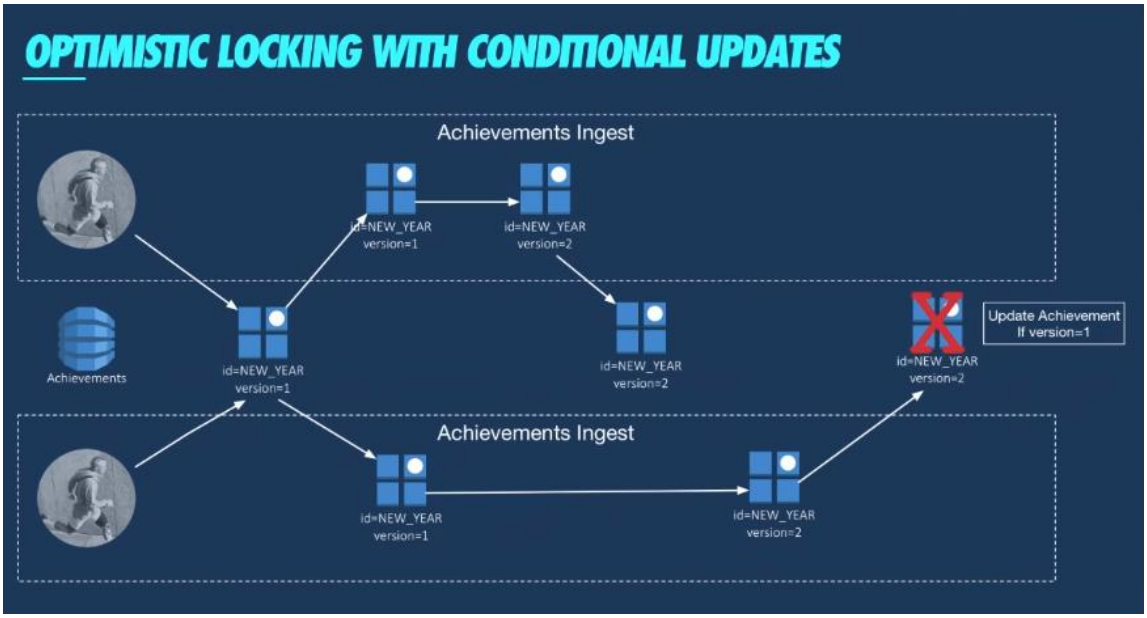
OPTIMISTIC LOCKING WITH CONDITIONAL UPDATES
Achievements Ingest
id=NEW_YEAR version=1
id=NEW_YEAR version=2
Achievements
id=NEW_YEAR version=1
id=NEW_YEAR version=2
id=NEW_YEAR version=2
Update Achievement If version=1
Achievements Ingest
id=NEW_YEAR version=1
id=NEW_YEAR version=2

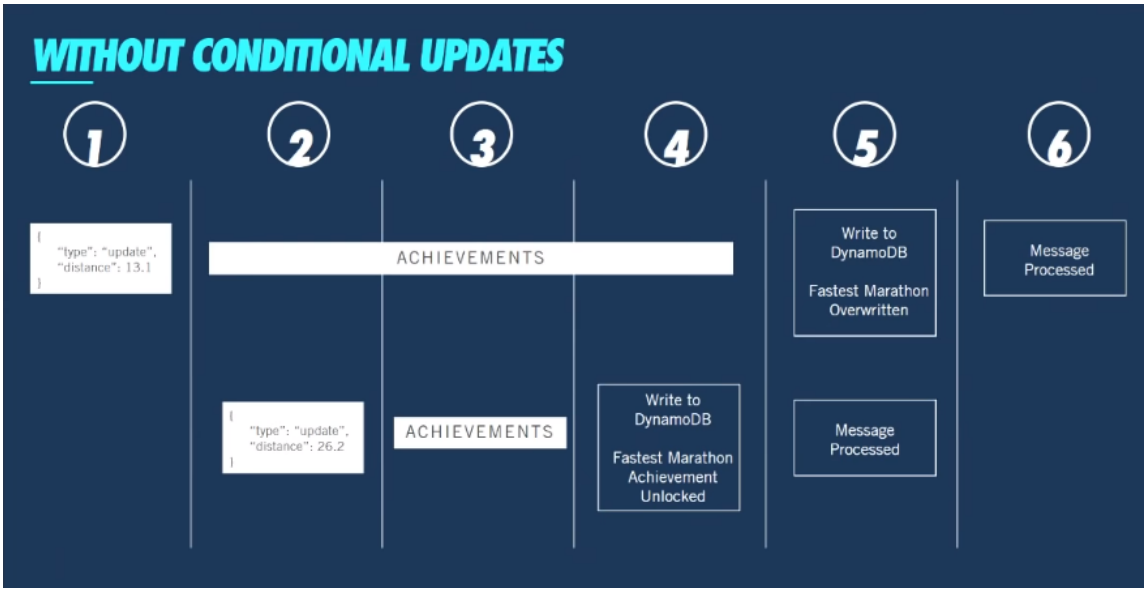
WITHOUT CONDITIONAL UPDATES
1
2
3
4
5
6
{ "type": "update", "distance": 13.1 }
ACHIEVEMENTS
Write to DynamoDB
Fastest Marathon Overwritten
Message Processed
{ "type": "update", "distance": 26.2 }
ACHIEVEMENTS
Write to DynamoDB
Fastest Marathon Achievement Unlocked
Message Processed


WITH CONDITIONAL UPDATES
1
2
3
4
5
6
{ "type": "update", "distance": 13.1 }
ACHIEVEMENTS
Conditional Update Failed
ACHIEVEMENTS
Message Processed
{ "type": "update", "distance": 26.2 }
ACHIEVEMENTS
Write to DynamoDB
Fastest Marathon Achievement Unlocked
Activity has been updated to 26.2 miles, no update occurs
VERSION 1
VERSION 2

TIMEZONE SERVICE

BUILT WITH CASSANDRA

Responsible for keeping track of an athlete's time zones as they move through the world.

All our activity data is stored in UTC time. We use time zones to adjust the user's activity data when needed.

The left contains a Cassandra schema and the right is the DynamoDB table schema where we just need to specify the key columns only.

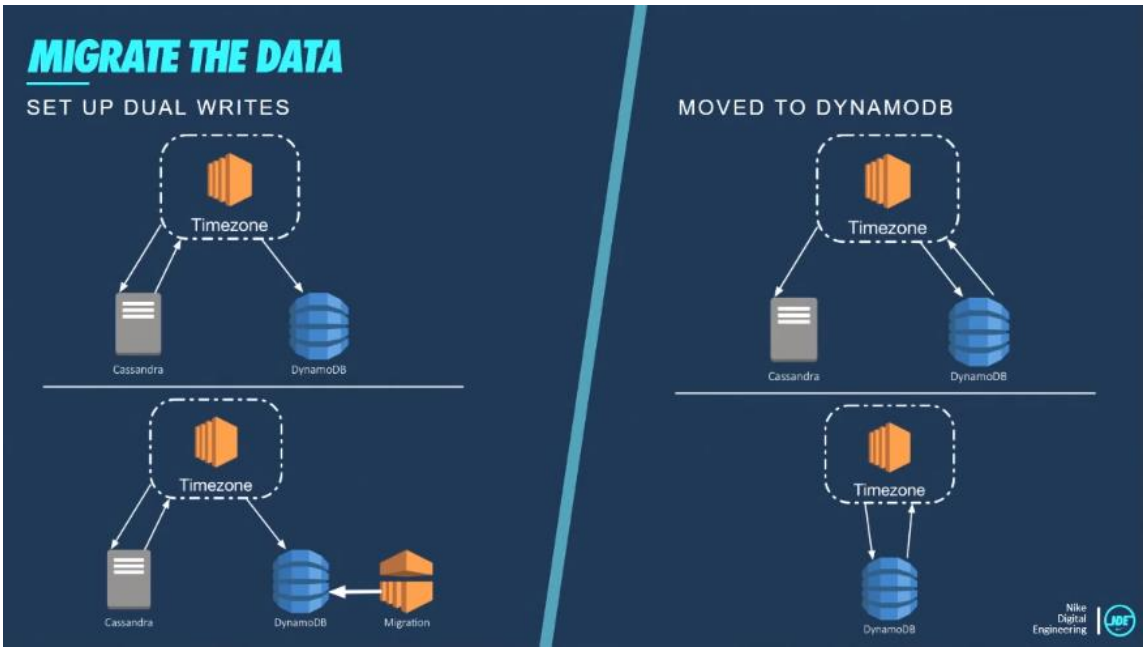
MIGRATE THE DATA
SET UP DUAL WRITES
Timezone
Cassandra
DynamoDB
Timezone
Cassandra
DynamoDB
Migration
MOVED TO DYNAMODB
Timezone
Cassandra
DynamoDB
Timezone
DynamoDB
Nike Digital Engineering


98% COST SAVINGS


DYNAMODB
IS THE NEW DEFAULT
SCALE
CONTINUOUS
INNOVATION


FIND US ON MEDIUM
MEDIUM.COM/NIKEENGINEERING
Nike
Digital
Engineering