

WPS305

How Fannie Mae Processes over a Quarter Million Loans per Day with Amazon S3

Harsha Nippani
Solutions Architect
Amazon

Oliver Mathias
Senior Architect
Fannie Mae

aws
re:Invent

© 2018, Amazon Web Services, Inc. or its affiliates. All rights reserved.



In this session, Fannie Mae discusses how they completely **re-architected a mission-critical application using AWS native services** that **process hundreds of thousands of mortgage loans every day in a highly scalable and reliable manner**. The **transaction-heavy workload uses over 20+ million Amazon S3 transactions a day, each within 150-millisecond response times**, thus providing increased uptime and faster response.



Harsha Nippani

- Solutions Architect, AWS
- Help customers deploy successfully on AWS
- 18+ years in IT systems engineering and Ops



Oliver Mathias

- Senior Architect, Fannie Mae
- Enthusiastic about advanced technology
- 15+ years solving complex technical problems

What You'll Get Out of this Session

- Amazon Simple Storage Session (Amazon S3) best practices to build low latency apps
- Front row seat to our journey at Fannie Mae
- Practical takeaways you can try at home



Amazon S3 – Object Store at Scale

The AWS Storage Portfolio



Amazon EBS
(persistent)



Amazon EC2
Instance Store
(ephemeral)

Block



Amazon EFS

File



Amazon S3



Amazon Glacier

Object

Data Transfer



AWS
Snow Family



AWS Storage
Gateway



EFS
File Sync



3rd Party
Connectors



AWS Direct
Connect



S3 Transfer
Acceleration



Amazon
Kinesis

aws
re:Invent

© 2018, Amazon Web Services, Inc. or its affiliates. All rights reserved.



Amazon S3 by the Numbers

One of first
three AWS
services (2006)



44 Availability Zones
(16 more coming in 2018)



16 Regions
(5 more coming
in 2018)



99.999999999
% Durability



Millions of
requests per
second



Trillions of
objects

Benefits of Amazon S3 & Glacier



Durable, Available, & Scalable



Security & Compliance



Low-Cost




Flexible Management



Ecosystem

Amazon S3 Architecture



Amazon S3 Security, Encryption, & Compliance

One of the broadest set of tools in the industry



Security

- AWS Identity Access and Management (IAM) and bucket policies
- Access control lists
- Audit logging with AWS CloudTrail & alerts with Amazon CloudWatch
- Secure AWS CloudFormation templates
- Amazon Macie
- Amazon S3 Console permission checks *New*



Encryption

- Encryption in transit with TLS
- SSE-S3 – Amazon S3 manages data & keys
- SSE-C – Customer managed keys
- SSE-KMS – Master keys in AWS Key Management Service (AWS KMS)
- CSE – 100% Customer managed *New*
- Default bucket encryption *New*
- Encryption status in inventory



Compliance

- PCI-DSS
- HIPAA/HITECH
- FedRAMP
- FISMA
- EU Data Protection Directive

AWS Storage Customers



Enterprise Partner Integration (Storage)



High Request Rates

- S3 automatically scales to high request/sec (RPS) rates
 - 3,500 RPS for PUT/POST/DELETE (Per Bucket Prefix)
 - 5,500 RPS for GETS (Per Bucket Prefix)
 - Add prefixes to increase read and write performance exponentially
- SSE-KMS encryption
 - 5,500 – 10,000 RPS for KMS Encrypt/Decrypt actions

Using a Three or Four Character Hash (Entropy)

Due to recent Amazon S3 performance enhancements, most customers no longer need to worry about introducing entropy in key names

examplebucket/232a-2017-26-05-15-00-00/cust1234234/photo1.jpg
examplebucket/7b54-2017-26-05-15-00-00/cust3857422/photo2.jpg
examplebucket/921c-2017-26-05-15-00-00/cust1248473/photo2.jpg



A bit more LIST friendly:

examplebucket/animations/232a-2017-26-05-15-00-00/cust1234234/animation1.obj
examplebucket/videos/ba65-2017-26-05-15-00-00/cust8474937/video2.mpg
examplebucket/photos/8761-2017-26-05-15-00-00/cust1248473/photo3.jpg



Random hash should come before patterns such as dates and sequential IDs
Always first ensure that your application can accommodate

Parallelizing GETs

For large objects, use range-based GETs - align your get ranges with your parts

```
GET /example-object HTTP/1.1
Host: example-bucket.s3.amazonaws.com
x-amz-date: Fri, 28 Jan 2016 21:32:02 GMT
Range: bytes=0-9
Authorization: AWS
AKIAIOSFODNN7EXAMPLE:Yxg83MZaEgh3OZ3l0rLo5RTX11o=
```



For content distribution and edge caching, enable Amazon **CloudFront**

- Caches objects from S3 at the edge
- Low latency data transfer to end user
- Multiple endpoints globally

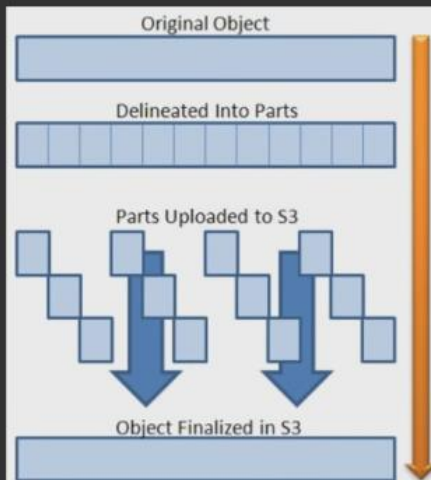
aws
re:invent

© 2016, Amazon Web Services, Inc. or its affiliates. All rights reserved.



If you are repeatedly asking for the same objects, you might want to consider using edge caching using CDNs like CloudFront for low latency.

Parallelizing PUTs



- Multi part upload for large objects
- Increase aggregate throughput by parallelizing PUTs on high-bandwidth networks
 - Move the bottleneck to the network, where it belongs
- Increase resiliency to network errors; fewer large restarts on error-prone networks

Key Amazon S3 best practices – Performance

- ✓ Faster upload over long distances with S3 Transfer Acceleration
- ✓ Faster upload for large objects with S3 multipart upload
- ✓ Optimize GET performance with Range GET and Amazon CloudFront
- ✓ SQL Query on S3 with Amazon Athena
- ✓ Distribute key name for high TPS workload
- ✓ TCP Window Scaling for long, fat networks
- ✓ TCP SACK for fast, lossy connections like mobile

How do you successfully modernize your important customer's highly visible, high throughput customer-facing application?

And live to talk about it?

Loan Processing at Scale

Amazon S3

Store any file, any size, any volume

For any amount of time

With stellar durability

Fast, reliable access to data

And predictable scalability

All for a cost that does not break the bank?



"Usual" List

Static web site hosting
File sync and storage solutions
Image, video, & multimedia sharing
Data store for big data workloads
Data lake & AI storage
Social web and SaaS use cases



"Unconventional" Thinking

High frequency requests
Blocking transactions
Quick response times
Atomicity and consistency
Guaranteed delivery



Amazon S3 as blazing fast storage system for high throughput transactional apps?

Age-old solutioning

Traditional solution

- Build out DC
- Upfront high CapEx
- High Complexity
- Peg to peak capacity



Costly, unmaintainable, and inflexible.
Upgrade treadmill

Cloud 1.0 strategy

Cloud 1.0 solution

- Lift and shift
- Custom cloud storage
- Poor resiliency
- Peg to peak storage capacity

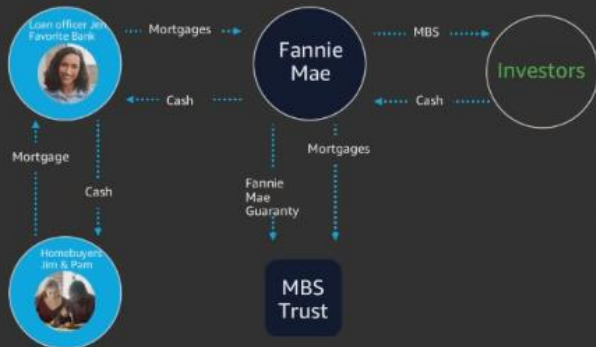


Slightly better, cloud-enabled solution. Checks off "Migrate to the cloud"

Cloud 2.0 Thinking : Get a little innovative, let's see how Fannie Mae did it!

Fannie Mae's Use Case

Fannie Mae is at the Heart of the Housing Industry



aws
re:Invent



Fannie Mae, America's Most Valued Housing Partner

1 in 3

Homes in the country are
financed by Fannie Mae

\$570 B

Mortgage Financing in 2017

1.2 M

New Homes

1 M

Refinancing

770 K

Multifamily Units

76%

Of Americans seeking a first
lien mortgage chose
30-Year fixed rate mortgage
In 2017

aws
re:Invent

© 2018, Amazon Web Services, Inc. or its affiliates. All rights reserved.

aws



Really Simple View of Loan Processing



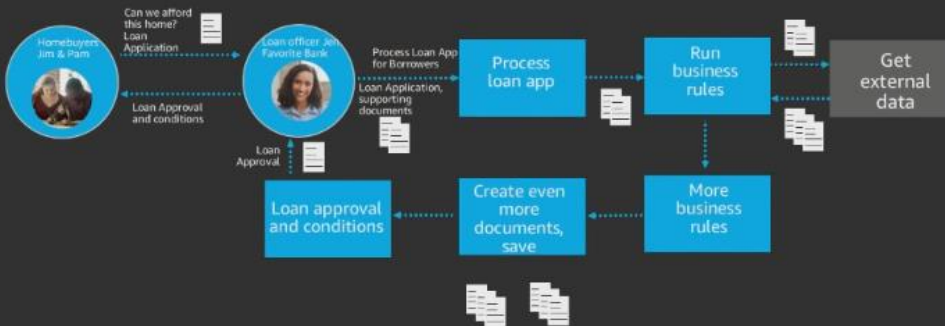
A prospective loan applicant works with a loan officer, submits a loan application along with a set of loan documents.

Really Simple View of Loan Processing



The loan application and the documents are then processed by Fannie Mae systems using a set of business rules to understand the applicant

Really Simple View of Loan Processing



More business rules are run in addition to collected external data, then a result is given about the loan application.

Three
very important goals

Near zero downtime

Near zero data loss

Millisecond responses

Key Architecture Challenges

Maximizing performance, availability and durability all at the same time

Incredibly low latency

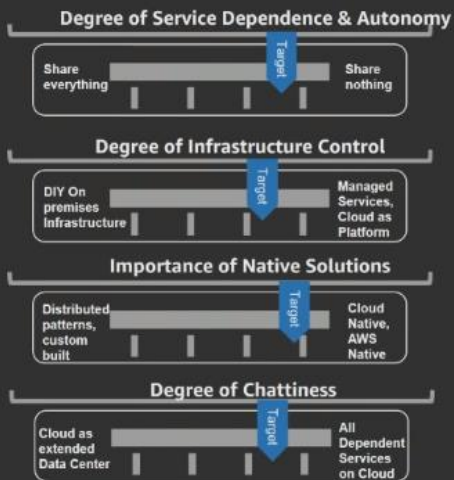
Billions of objects/year

Work with current platform

Reduce Operating Expenses(OpEx)



Create Guardrails Early

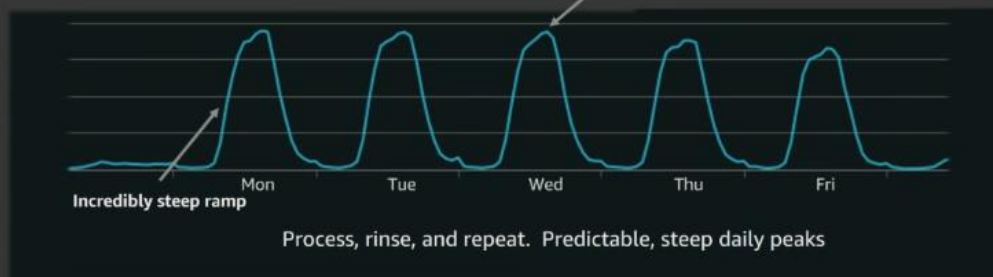


The Solution


Just Another Week

Object PUTs/GETs over one week

800/1600+ TPS at peak volume



Simplified Architecture View



Initial Observations – GET Latency



Key architecture challenges

Best practices, key first steps

Second phase, shaving off milliseconds



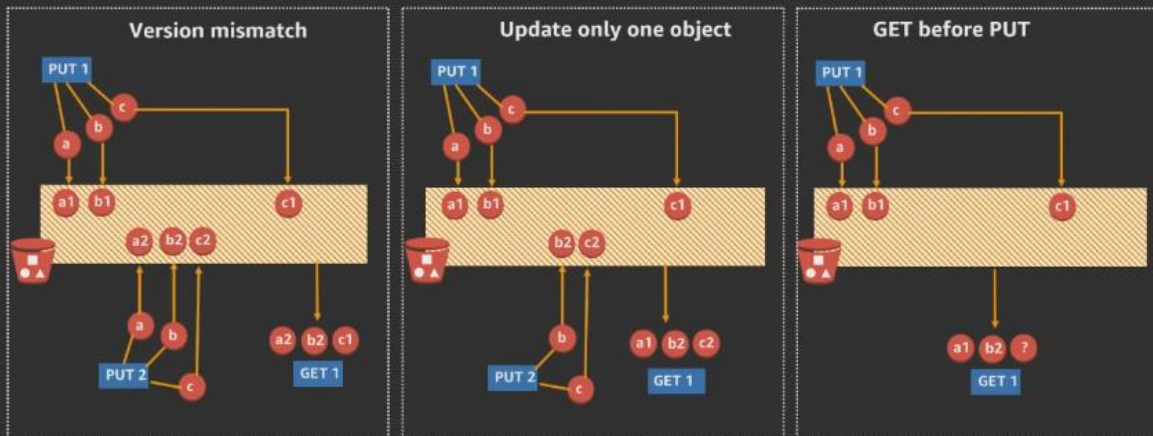
Standard Best Practices

- + Entropy (randomness)
- + Parallelizing PUTs & GETs
- + For troubleshooting use Request ID, S3 Extended Request ID & Host ID
- + Externalizing markers
- + RDS as an indexing mechanism for markers

Pay attention to client execution timeout

```
try {
    ClientConfiguration configuration = new ClientConfiguration();
    configuration.setMaxConnections(...);
    configuration.setRetryPolicy(...);
    configuration.setClientExecutionTimeout(...);
    AmazonS3 s3Client = AmazonS3ClientBuilder.defaultClient();
    S3Object object = s3Client.getObject(new GetObjectRequest(...));
} catch (ClientExecutionTimeoutException ce) {
    log.error(...);
    throw ce;
}
// the caller of this method, may retry to get the object successfully
```

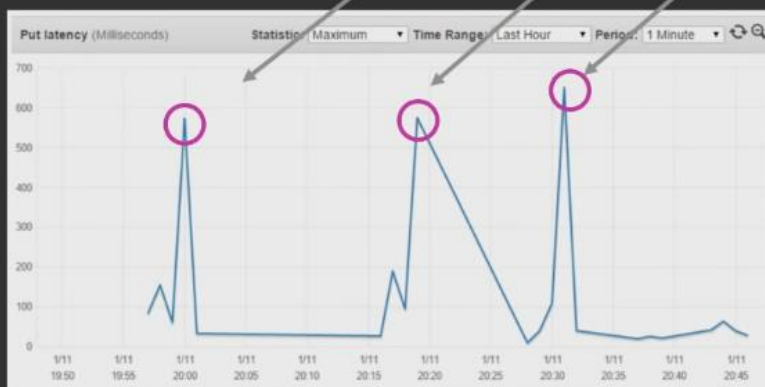
Working Around Amazon S3 Eventual Consistency



Our application was strongly consistent but S3 is eventually consistent, so we had to augment with an RDS database with Postgres

PUT Latency

Consistently reproducible response time spikes every 15 min





Unpredictable, could add several seconds to customer response times

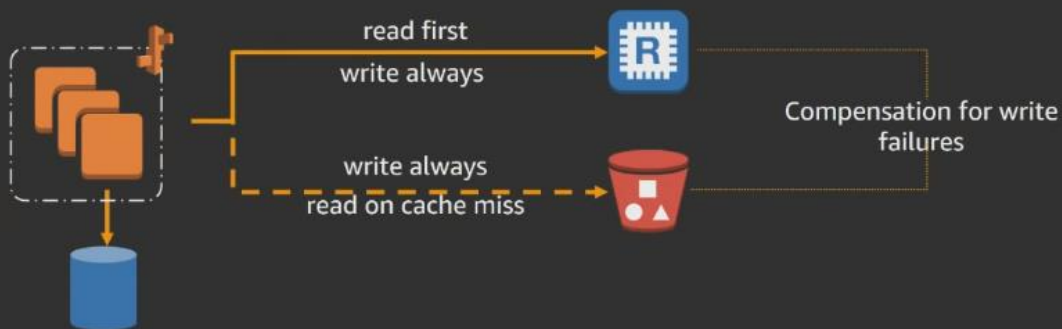
Best practices, key first steps

Second phase, shaving off milliseconds

Third phase, pushing the limits, literally. Improving resiliency

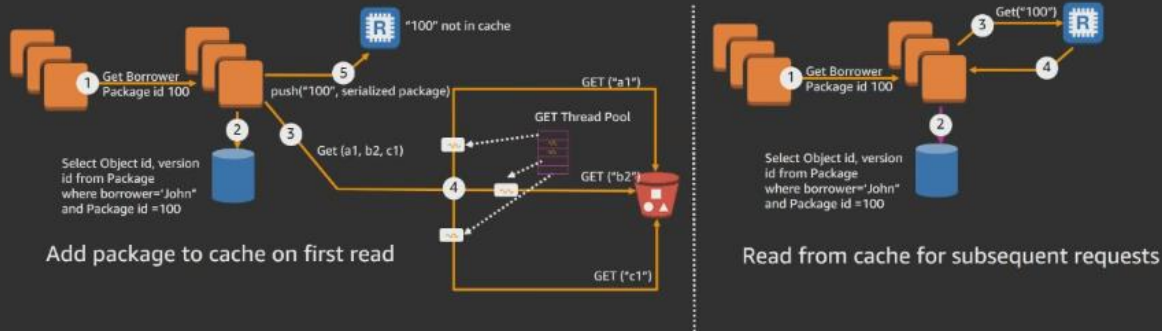


Caching In



Uncoupling Amazon S3 stores to improve PUT latency
Needs compensation to ensure continued high durability

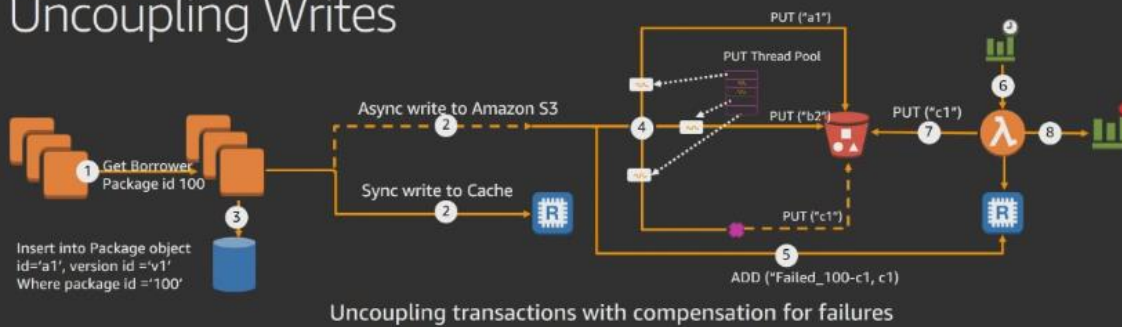
Cached Reads



Saves 100-150ms per transaction


Over 200,000 seconds of processing time/latency saved every day

Uncoupling Writes



Cache writes are blazing fast, saves 250-300ms per transaction
Over 400,000 seconds of processing time/latency saved every day

Encrypting Data at Rest



- ✓ SSE-KMS provides many operational advantages, so you might want to stick to it
- ✓ AWS KMS observed availability is pretty high, 99.999%
- ✓ For high variable workloads from 1000s to millions, SSE-S3 might be a better fit

Switching to CSE can save 15-25ms per transaction
Over 100,000 seconds of processing time/latency saved every day




Goal : Delivering Incredibly low latency - 99% within 200-300ms

What's our percentile?

☐ 95% ☒ 97% ☐ 99%

Optimized Response Times

S3 GETs



S3 PUTs

0.17	(60%)
0.22	(70%)
0.27	(80%)
0.29	(85%)
0.32	(90%)
0.36	(95%)
0.38	(97%)
0.17	Average



Improve Store response time by over 100% with uncoupled cache



< 300ms GET latency for 97% of transactions



< 300ms PUT latency for 97% of transactions (*not quite, but not relevant due to uncoupling*)


Second phase, shaving off milliseconds

Third phase, pushing the limits, literally. Improving resiliency


More key challenges



Slightly Less Simple View of the Architecture

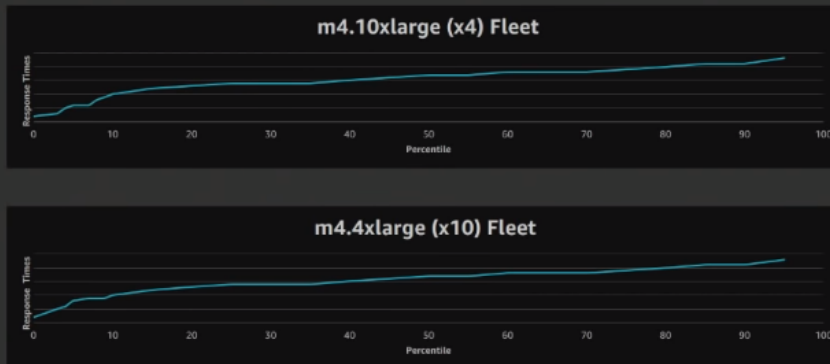


Retries for Occasional GET Latency Spikes



Distributed nature of Amazon S3 = some GETs take over several seconds
Retry resolves time outs, may take 2-3 attempts

Rightsizing Instances



Smaller instances have similar, may be even better performance


Goal : Delivering incredibly low latency - 99% within 200-300ms

What's our percentile?


☐ 97% ☐ 98% ☒ 99%

Results

S3 GETs



S3 PUTs with Caching



99% of stores and reads within 275 milliseconds

More Key Challenges

High(er) availability

Region failure immunity

Minute scale recovery point objective (RPO) & recovery time objective (RTO)



Achieving Resiliency

Assume everything in the app could fail

Dependencies could fail

Network & connectivity could fail


Availability Zones and Regions could fail

Assume everything will fail all the time


App & Dependent Service Resiliency, Briefly




Resiliency in a Region



Multi-region Options



RPO vs. Cost & Complexity



Typical Multi-region Options

	All native	Custom DB	Custom Amazon S3	Custom DB & Custom Amazon S3
Compute	Native	Native	Native	Native
	Active / Active	Active / Active	Active / Active	Active / Active
	Async	Async	Async	Sync
Data - RDS	Native	Custom	Native	Custom
	Active / Standby	Active / Standby	Active / Standby	Active / Active
				Sync
Objects - Amazon S3	Native	Native	Custom	Custom
	Active / Standby	Active / Standby	Active / Active	Active / Active
RPO	High	Medium	Low	Very Low
Cost	VL	H	M	VH
Complexity	L	M	L	H
Response Time Impact	None	Network & lag dependent	None	Slowest due to sync n/w & lag

Our Simple High Availability Plan



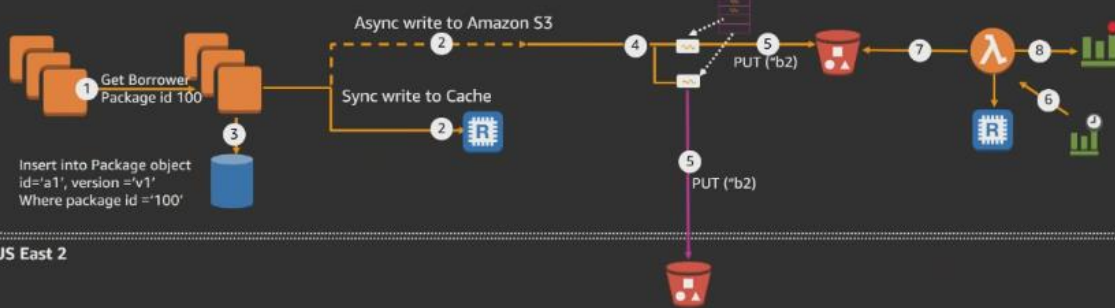
Uncoupled Multi-Region Writes



US East 2

Write to both regions, build redundant compensation for failures

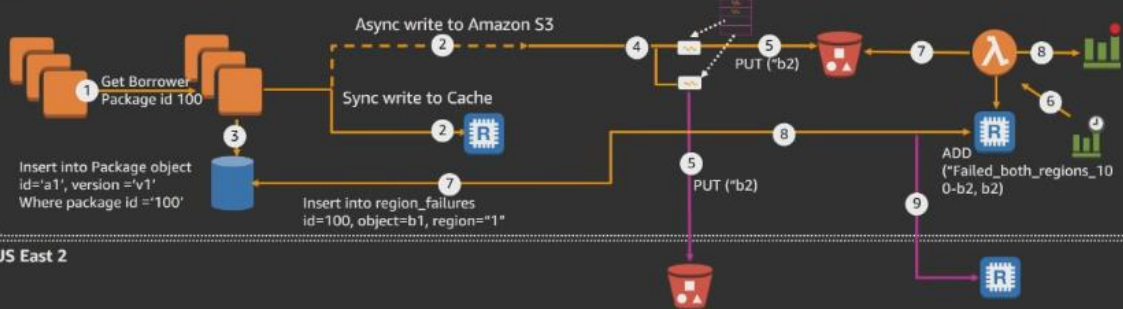
Uncoupled Multi-Region Writes



US East 2

Write to both regions, build redundant compensation for failures


Uncoupled Multi-Region Writes



US East 2

Write to both regions, build redundant compensation for failures

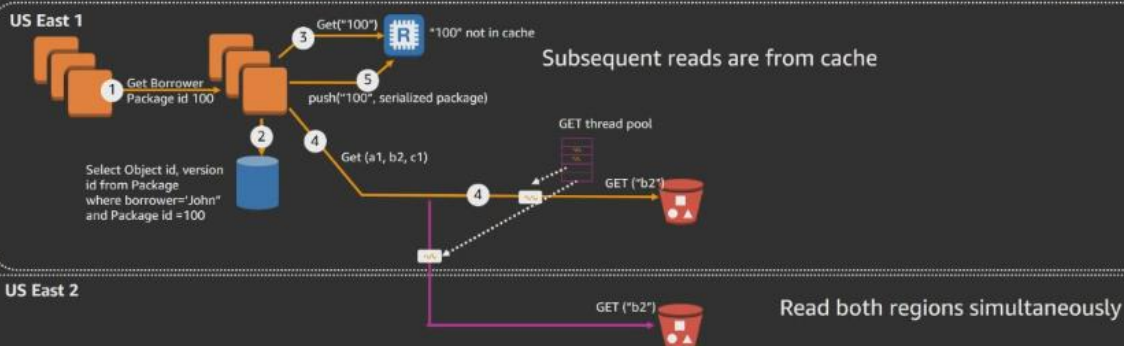
Multi-region Reads



US East 2

Eliminates sporadic S3 latency spikes, region failure immunity

Multi-region Reads



Eliminates sporadic S3 latency spikes, region failure immunity

Key Takeaways

- + 1000's of TPS is possible
- + Parallelize GETs & PUTS, add retries
- + Last 1% will take significant effort
- + SSE KMS good for almost every use case, < 20ms overhead
- + Remember: Bigger instances are not always better !
- + Objects and metadata can't stand each other
- + Cache in for predictable low latency
- + 4-9s+, write and read from both regions

Be a little unconventional with Amazon S3, serious benefits are in store

99.9% of transactions < 300 milliseconds

Better availability

Better TCO

Thank you!

Oliver Mathias

Oliver_mathias@fanniemae.com

Harsha Nippani

nippanih@amazon.com

AWS
re:Invent

© 2018, Amazon Web Services, Inc. or its affiliates. All rights reserved.

aws