

ARC414-S

Accelerated analytics: Building the next-gen data platform for Hertz

Lucy Meyo

VP Architecture and
Integration
Hertz

Rohit Sinha

Senior Manager
Analytics and Cognitive
Deloitte Consulting LLP

Gowtham Ramu

Specialist Leader
Cloud Engineering
Deloitte Consulting LLP

aws
re:Invent

© 2019, Amazon Web Services, Inc. or its affiliates. All rights reserved.



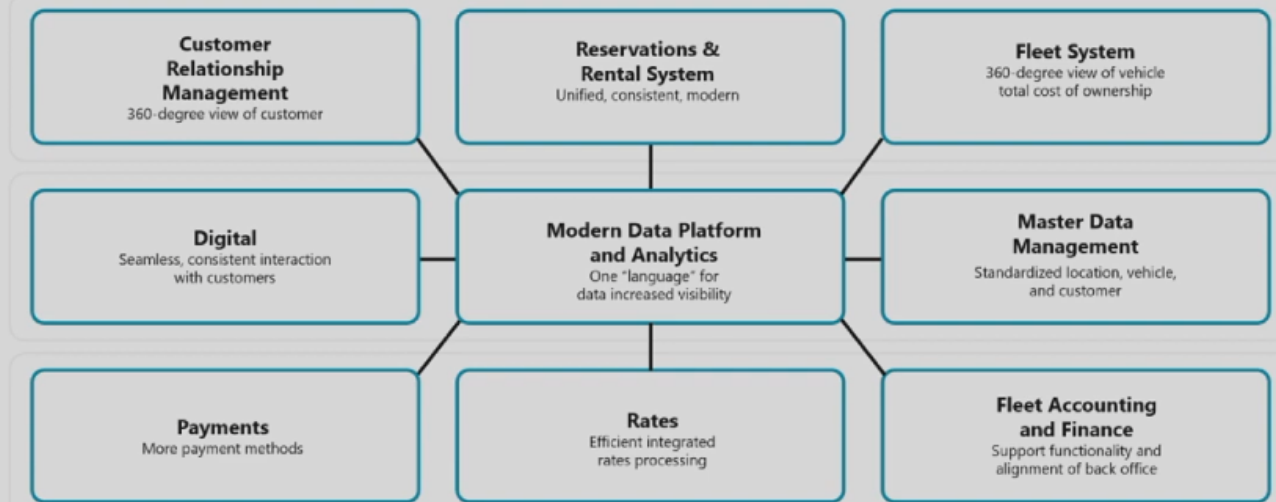
Hertz is undertaking a massive digital transformation to evolve its technology landscape. This move provides an opportunity to extract valuable insights from a large amount of data produced by the new systems. In this session, learn how Deloitte collaborated with Hertz to build a next-generation data platform, which includes an integration hub, a unified reporting layer, and an ecosystem of tools used to build advanced analytics models. Discover how the solution leveraged all native AWS services, such as Amazon S3, Amazon Kinesis, Amazon EMR, and AWS Lambda, to enable cross-functional insights and accelerate the cloud journey in under 12 months. This presentation is brought to you by Deloitte, an APN Partner.

Agenda

- 1 The opportunity
- 2 Accelerated journey to the cloud
- 3 Role of next-generation data platform
- 4 Cloud-native architecture
- 5 Key takeaways

Digital Transformation

The goal of the Digital Transformation program at Hertz is to deploy an integrated global solution that provides a better view of our customer and greater visibility into our assets



We have over 1.4 million cars, 19 million global customers...we need to provide a seamless new digital journey using our data powered modern data platform and analytics.

The opportunity

A *modern data and analytics platform* was critical to the success of the overall *digital transformation at Hertz* to extract value and insights from large amount of data being produced from new systems and create a common platform for cross-functional insights and answering new business questions using the advanced analytics and data science models.



Sharply reduce TCO of data management and analytics with the move to AWS



Leverage scalable elastic compute of cloud infrastructure for modern analytics and data science use cases



Centralize data assets in AWS data lake



Improve regulatory and security controls with the process of centralizing data assets in AWS data lake



Use cloud infrastructure scalability and elasticity as well as spot instances to optimize cost of compute-intensive operations and answer new business questions



With pattern-based ingestion and component-based architecture, enable quicker time to implement new use cases



Enable digitization through new business models around data and analytics

Hertz journey of how Deloitte helped build the MVP for the data platform in less than 10 months

Data platform journey so far

Data and reporting requirements

Defined data and reporting requirements across all functions at Hertz and programs within digital transformation to support unified reporting, self-service BI, and analytics



- Conducted workshops with 30+ functional groups/sub-groups to finalize Day1/Day2 reporting and analytics requirements
- Rationalized ~1,800+ reports down to ~280 prioritized Day1 reports/dashboards

Data source discovery

Leveraged data and reporting requirements to do source system discovery; mapped new source system integrations as well as legacy systems and catalog integration gaps



- Identified 30+ data sources outside of new systems to identify integration gaps based on Day 1/Day 2 critical report requirements and associated data needs; catalogued total of 200+ integrations coming in and out of data platform from different systems

Architecture alignment and infrastructure standup

Achieve alignment across a scalable, forward-looking data platform architecture, and stand up infrastructure in AWS to support integration, reporting, analytics, and data science use cases



- Finalized data platform architecture blueprint and technical bill of materials
- Successfully established the data platform architecture on AWS to support in-time integration with Hertz digital transformation program

Infrastructure automation

Created automated infrastructure provisioning scripts to standardize infrastructure components across the data and analytics stack




- Built all environments (Dev, QA, UAT, and prod) using CFTs and Ansible scripts for infrastructure automation

Copyright © 2019 Deloitte Development LLC. All rights reserved. 5

Solution overview and role of data platform at Hertz

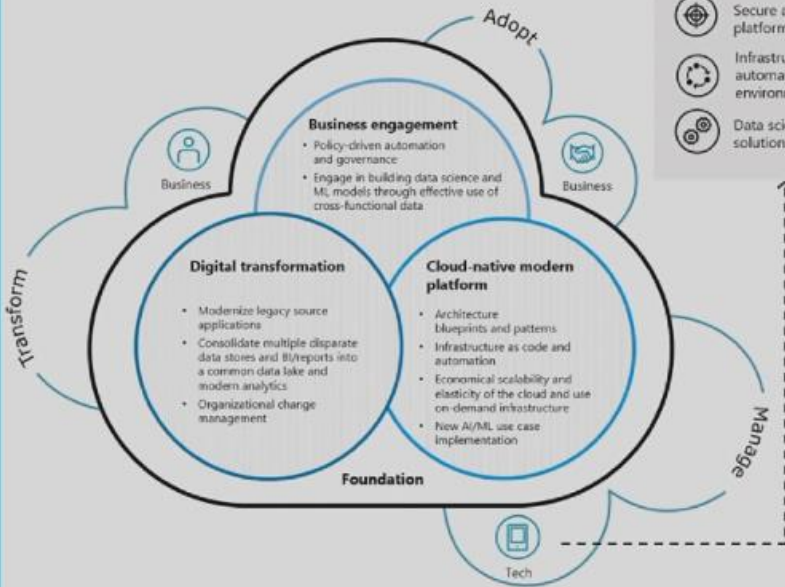
Key purpose and underlying architecture



Solution overview

Cloud-native data platform

Enabling digital transformation journey and providing consolidated data across multiple systems for accelerated insights generation

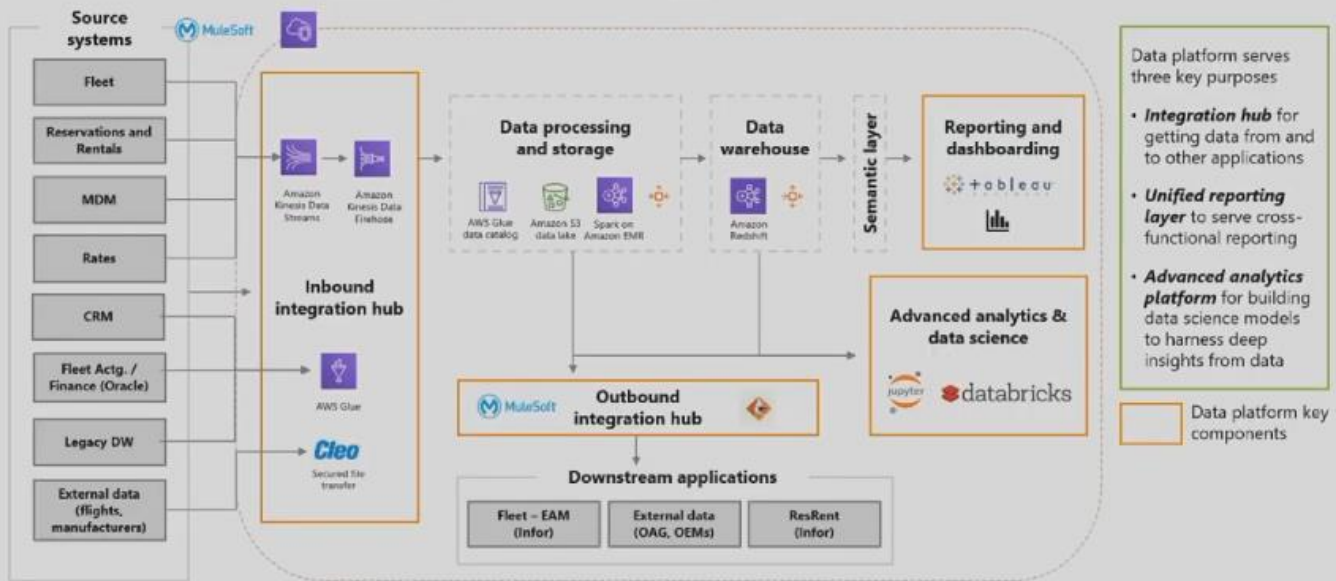


Cloud-native solution architectures

- Batch and streaming ingestion patterns
- Secure and scalable platform
- Infrastructure automation for all environments
- Data science/ML solutions

The role of the next-generation data platform at Hertz

High-level architecture illustration of the data platform at Hertz and the three key purposes of the modern platform



Data platform serves three key purposes

- Integration hub** for getting data from and to other applications
- Unified reporting layer** to serve cross-functional reporting
- Advanced analytics platform** for building data science models to harness deep insights from data

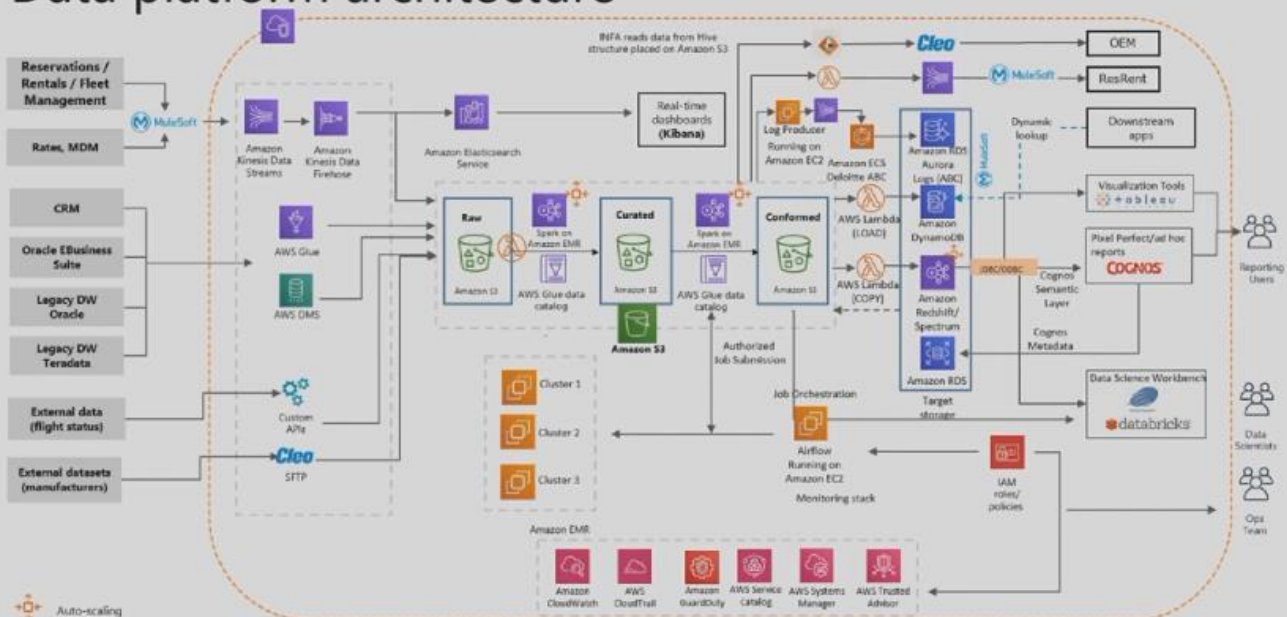
Orange boxes represent Data platform key components

Cloud-native architecture

Built with all native AWS services to stand up a modern data platform




Data platform architecture



Repeatable patterns in architecture

The patterns below illustrate high-level reusable components within the data platform at Hertz that was deployed using native AWS services

- Real-time stream ingestion in Amazon S3 with Amazon Kinesis Data Streams, Amazon Kinesis Data Firehose, and Amazon S3 for all event-driven ingestion in the data platform
- Real time to batch use cases where data-generating system can only produce events and consuming systems can only take batches (orchestrated using Kinesis Data Streams, Kinesis Data Firehose, and AWS Glue or Informatica reading from Hive)
- Real-time API call with external system (Flight Aware) to tie reservation data for optimizing fleet management
- Low-latency fast lookups by transactional applications to look up in real time performed using lookup tables hosted in Amazon DynamoDB with defined read and write capacities

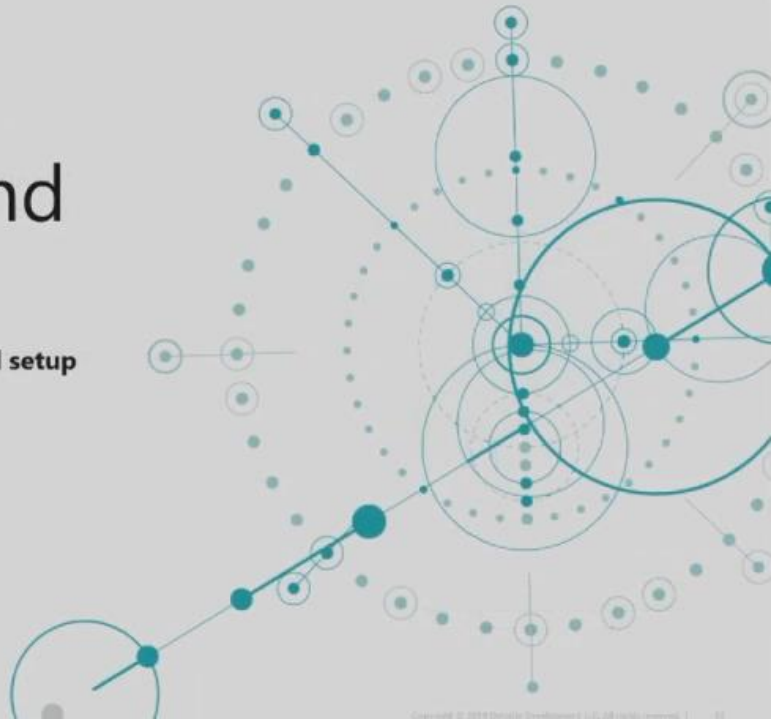


Impact

Patterns simplify the ingestion and consumption out of the data platform and avoid point-to-point integration with different components and designs, making it modular code

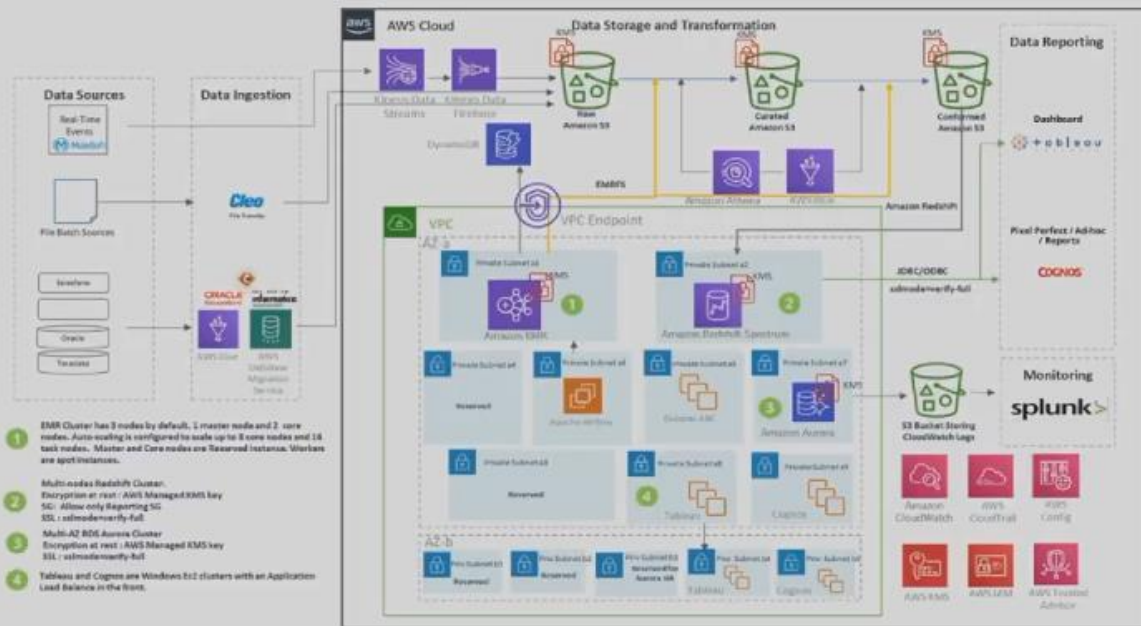
Infrastructure provisioning and automation

Under the hood infrastructure design and setup

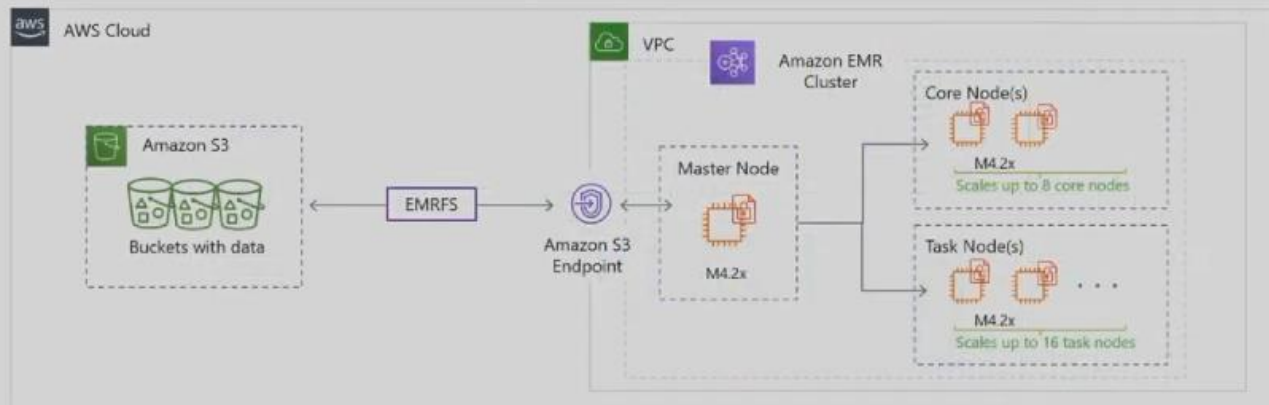


Copyright © 2019 Amazon Development LLC. All rights reserved. | 11

Infrastructure blueprint

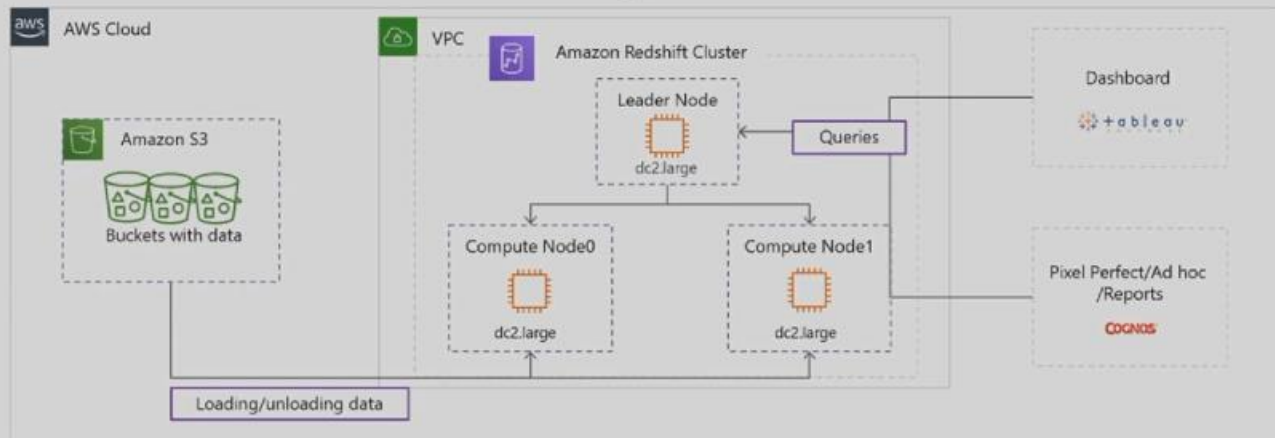


Amazon EMR stack diagram



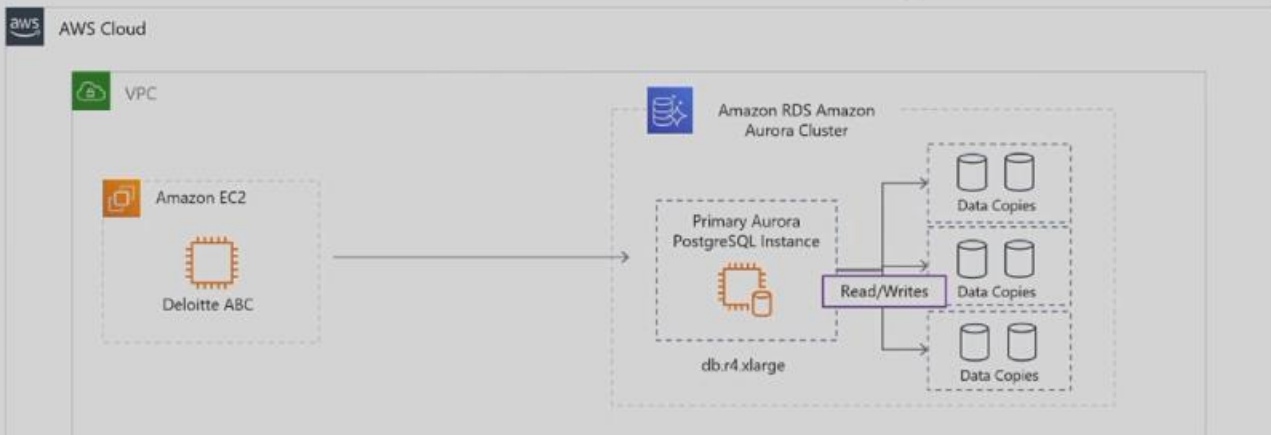
- Auto-scaling is configured to scale up to 8 core nodes and 16 task nodes
 - Scale-Out** "Add 1 Instance if **ContainerPendingRatio** is greater than 0.7 for 1 five-minute period with a cooldown of 300 seconds"
 - Scale-In** "Terminate 1 Instance if **ContainerPendingRatio** is less than 0.2 for 3 five-minute periods with a cooldown of 300 seconds"
 - Scale-Out** "Add 1 Instance if **HDFSUtilization** is greater than 90 for 1 five-minute period with a cooldown of 300 seconds"
 - Scale-In** "Terminate 1 Instance if **HDFSUtilization** is less than 70 for 3 five-minute periods with a cooldown of 300 seconds"
- The setup includes YARN Resource Manager, Spark History Server, and JupyterHub notebooks

Amazon Redshift stack diagram



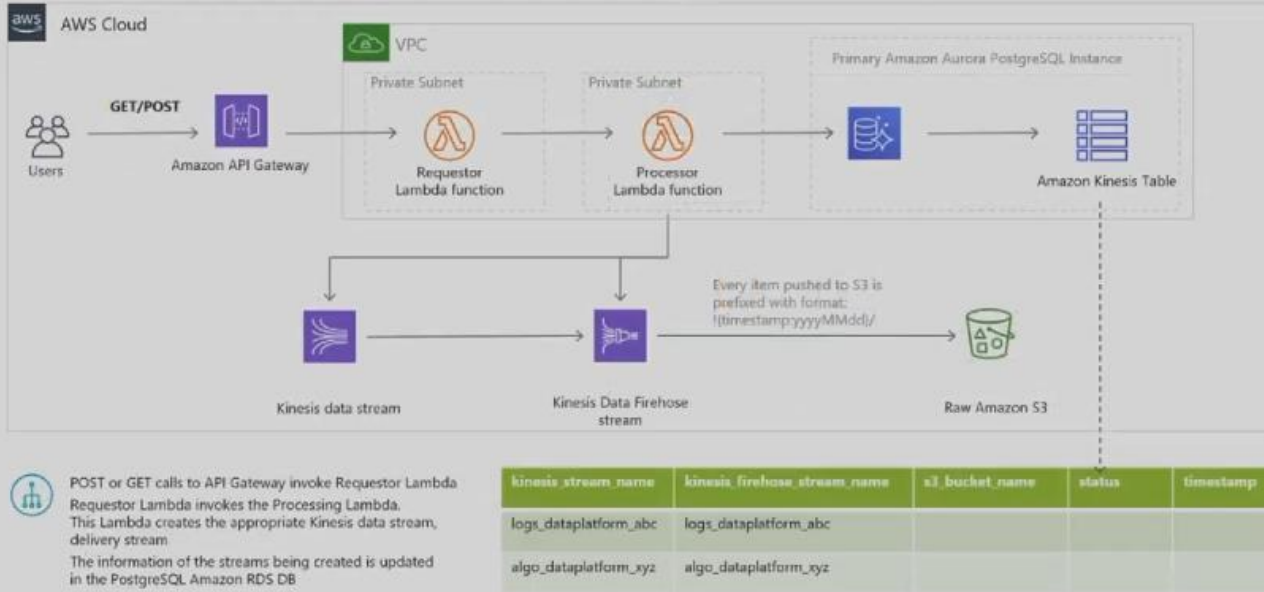
- spectrumNodeType: dc2.large
- spectrumClusterType: multi-node
- Publicly accessible: No

Amazon RDS Amazon Aurora stack diagram



- Deloitte's Audit-Balance-Control (ABC) Framework uses Amazon Aurora Postgres to write errors and logs entries for every data movement with the platform in an asynchronous manner
- DBInstanceClass: db.r4.xlarge
- DbEngineVersion: 9.6.9
- MultiAZ: True
- RDSDBParameterGroup: Customer CFT

Amazon Kinesis API stack diagram



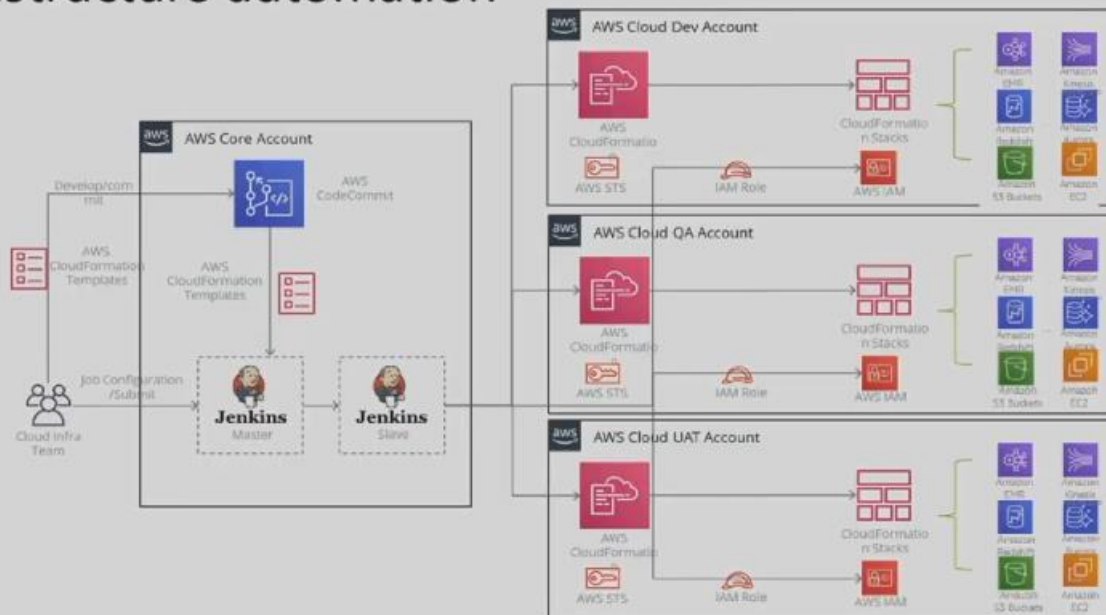
Copyright © 2019 Global to Development LLC. All rights reserved. | 34

AWS Glue stack diagram



- S3 bucket naming structure: environment-platform-appID-type-clientname-ID-region
- S3 encryption: AES-256 (SSE-S3)
- Amazon EMR cluster encryption: KMS key
- AWS Glue database naming structure: environment_platform_appID_SeqNum
- Amazon Athena will use the AWS Glue catalog for querying

Infrastructure automation



Demo

```
AWS::CloudFormation::Template
Description: AWS CloudFormation Template for S3 bucket creation.

Parameters:
  S3BucketName1:
    Type: String
    Description: Provide a unique bucket name. Bucket names must follow naming standard.
    AllowedPattern: ".*"
    MinLength: 6
    MaxLength: 63

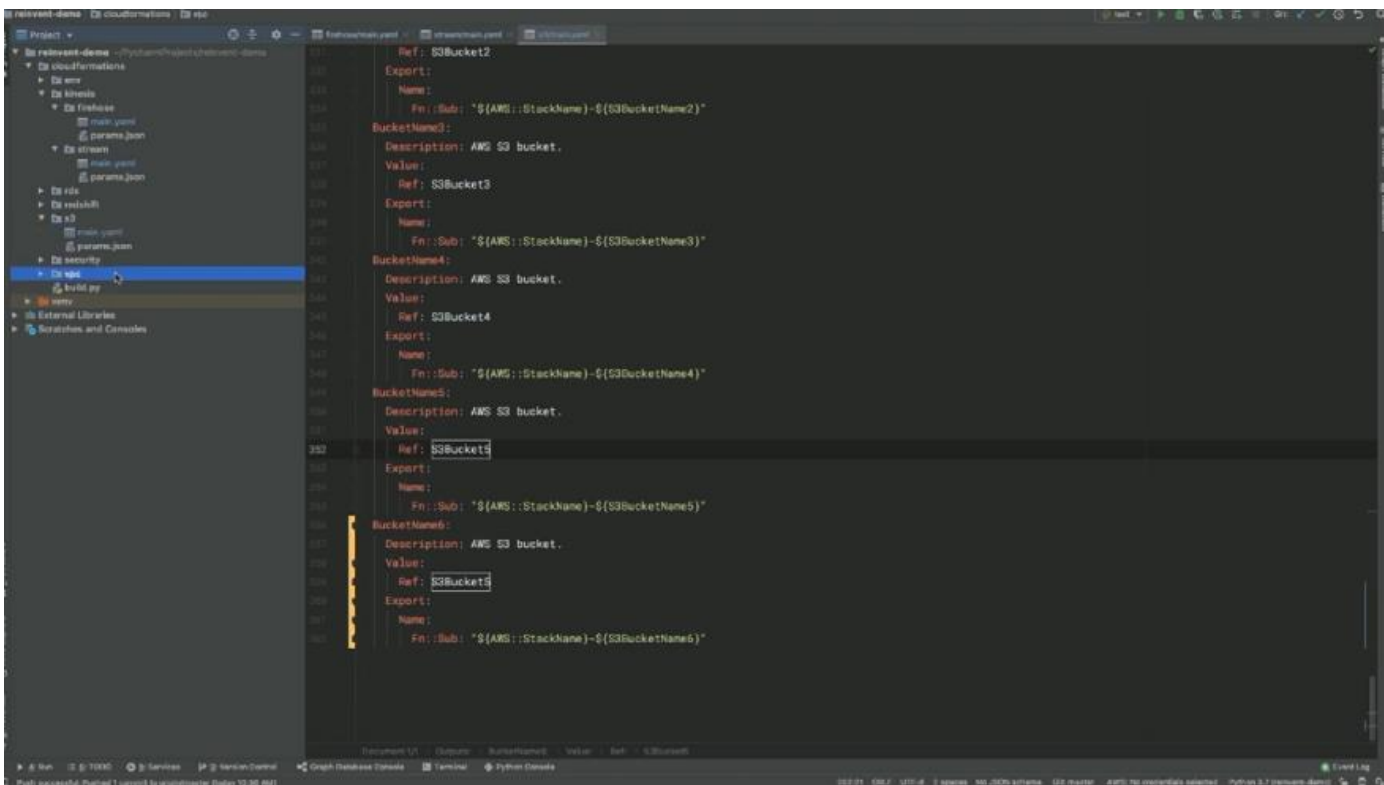
  S3BucketName2:
    Type: String
    Description: Provide a unique bucket name. Bucket names must follow naming standard.
    AllowedPattern: ".*"
    MinLength: 6
    MaxLength: 63

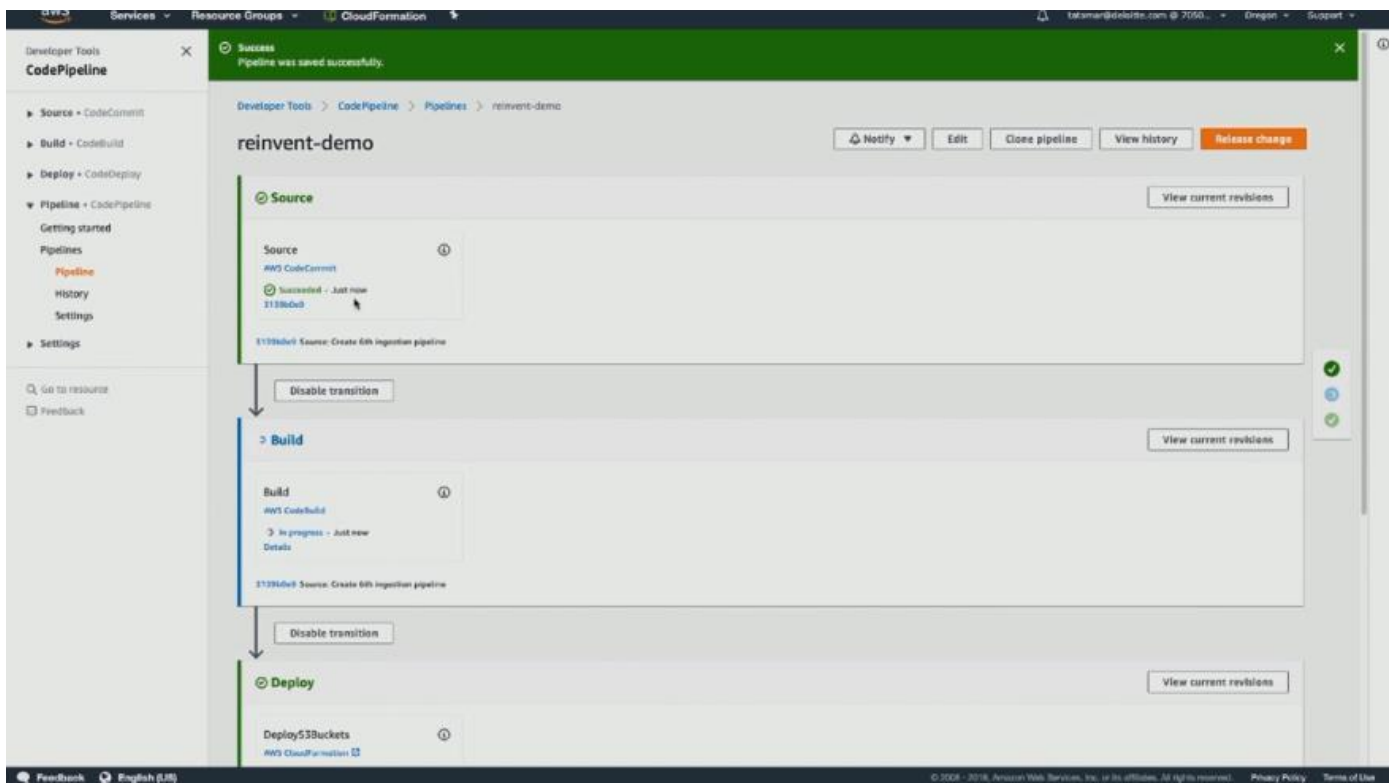
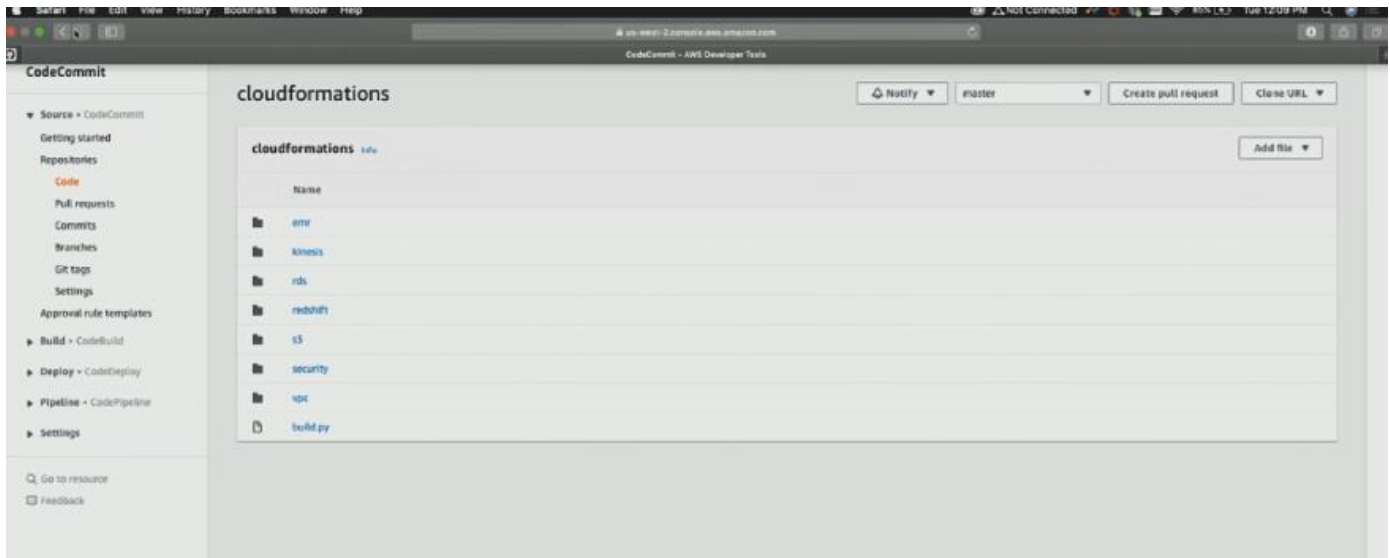
  S3BucketName3:
    Type: String
    Description: Provide a unique bucket name. Bucket names must follow naming standard.
    AllowedPattern: ".*"
    MinLength: 6
    MaxLength: 63
```

```

280     value: !Ref Lifecycle
281
282 S3Bucket6:
283   Type: AWS::S3::Bucket
284   Properties:
285     BucketName:
286       Ref: S3BucketName6
287     VersioningConfiguration:
288       Fn::If:
289         - HasVersioning
290         - Status: Enabled
291         - Ref: AWS::NoValue
292     BucketEncryption:
293       ServerSideEncryptionConfiguration:
294         - ServerSideEncryptionByDefault:
295             SSEAlgorithm: AES256
296     PublicAccessBlockConfiguration:
297       BlockPublicAcls: true
298       BlockPublicPolicy: true
299       IgnorePublicAcls: true
300       RestrictPublicBuckets: true
301     LifecycleConfiguration:
302       Rules:
303         - Id: GlacierRule
304           Prefix: !Ref LifeCyclePrefix
305

```





CloudFormation > Stacks > reinvent-demo-kinesis-firehoses

Stacks (7)

- reinvent-demo-kinesis-firehoses (UPDATE_COMPLETE)
- reinvent-demo-s3 (UPDATE_COMPLETE)
- reinvent-demo-err (CREATE_COMPLETE)
- reinvent-demo-kinesis-streams (UPDATE_COMPLETE)
- reinvent-demo-ids (CREATE_COMPLETE)
- reinvent-demo-redshift (CREATE_COMPLETE)
- reinvent-demo-vcpc (CREATE_COMPLETE)

reinvent-demo-kinesis-firehoses

Delete Update Stack actions Create stack

Stack info Events Resources Outputs Parameters Template Change sets

Overview

Stack ID	arn:aws:cloudformation:us-west-2:705068161531:stack/reinvent-demo-kinesis-firehoses/8d3f20d-158f-11ea-8a17-0a10741d0c78	Description	AWS CloudFormation Template to Kinesis Firehose delivery Stream
Status	UPDATE_COMPLETE	Status reason	-
Root stack	-	Parent stack	-
Created time	2019-12-02 21:43:42 UTC-0800	Deleted time	-
Updated time	2019-12-03 12:11:29 UTC-0800		
Drift status	NOT_CHECKED	Last drift check time	-
Termination protection	Disabled	IAM role	arn:aws:iam::705068161531:role/reinvent-demo-codedeploy

Tags (0)

Stack-level tags will apply to all supported resources in your stack. You can add up to 200 unique tags for each stack.

Search tags

Key	Value
No tags	

There are no tags defined for this stack.

Key learnings



Leverage AWS native services to enable a more seamless deployment (AWS pipeline and deploy)



Amazon RDS and Amazon Redshift don't enforce encryption in transit by default, and effort for encryption needs to be factored in



Streamline instance type to leverage reserved instance (RI) discount for better cost optimization



Don't hesitate to deep dive and ask questions or seek help; e.g., Customize AWS Secrets Manager password spec

Key takeaways

Benefits realized



Improvements realized



Stood up a modern data and analytics platform on AWS to create a consolidated data lake for value and insights generation from a large amount of data being produced from Hertz digital transformation and new systems



Enabled enhanced reporting and visualization of cross-functional data to support multiple business functions (sales, marketing, ops, customer care, etc.)



Created a platform for cross-application integration and cross-functional insights, and delivered on Hertz's analytical needs



Onboarded selected high-value data science use cases (e.g., fleet optimization, capacity substitution model, net price, etc.) on the data platform that needed scalable compute

Thank you!