

SVS308-R

Moving to event-driven architectures

Tim Bray

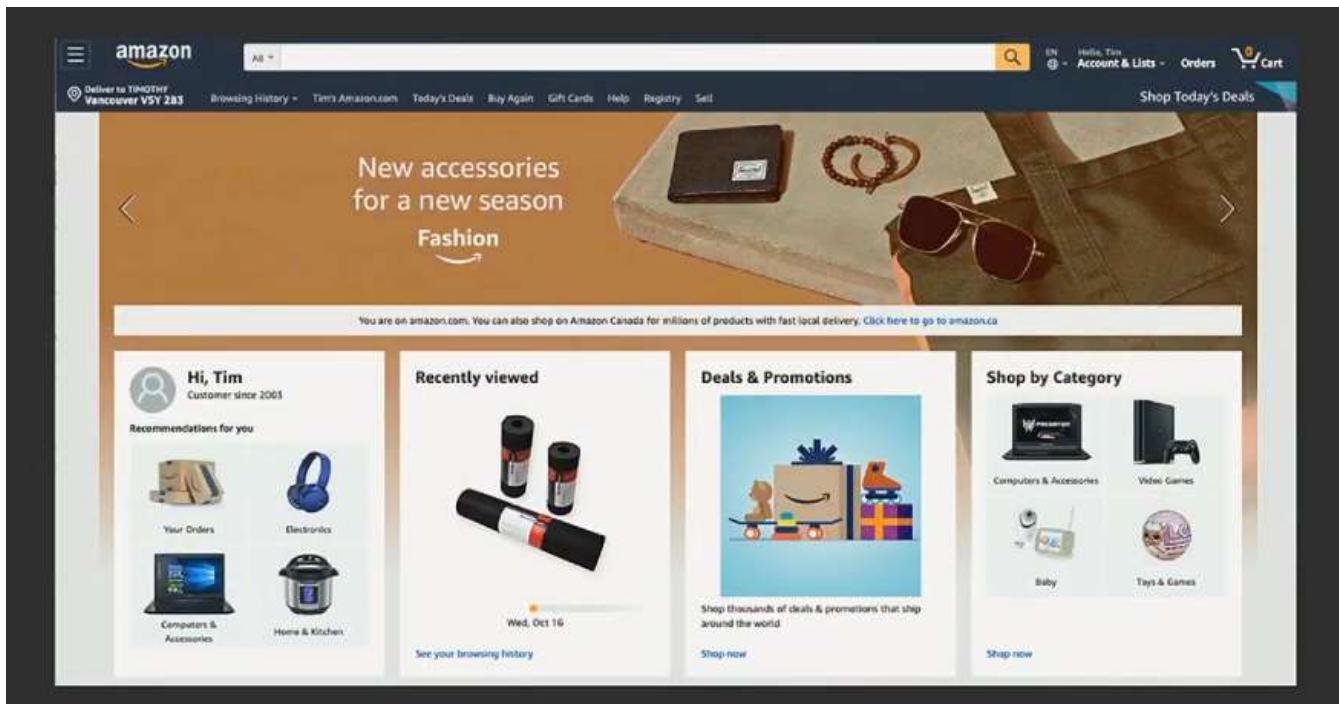
VP/Distinguished Engineer
Amazon Web Services

aws
re:Invent

© 2019, Amazon Web Services, Inc. or its affiliates. All rights reserved.



Event-driven architectures are getting a lot of attention. We have recently invested in tools and infrastructure to make event-driven architectures easier to implement and operate. In this session, we discuss what events are, why the community is increasingly interested in event-centric applications, and what's new in the domain and with AWS in particular. In addition, we discuss the challenges that still face us and our customers. By the end of this session, you understand the key principles and benefits of being event-driven.



This website is built as an event-driven architecture. A lot of things happen when the user clicks the buy button to generate the 'purchase' event.



A lot of async things need to get done after the customer clicks the 'buy' button, a 'purchase' event is generated and dropped in the queue with details saying that the customer bought something. Another system then picks up the event, opens the payload and then does its own stuff like creating an 'order' or 'packing' or 'shipping' event.

Amazon.com needs:

1. Scalable ingestion
2. Reliable storage
3. Fan-out

Amazon.com doesn't need:

1. Strong ordering
2. Deduplication
3. Push delivery
4. Ultralow latency

Amazon.com does not need strong ordering guarantees, it knows how to sort and filter out duplicates, does not do push into other backend systems and instead lets them ingest from queues at their own rates. These are the types of questions you need to ask when building an event-driven architecture.

Is event-driven
for you?

Is event-driven for you?

1. Are you passing around self-contained transactions?



"Account 0973482 bought an Instapot ID 238479r8732 for \$131.32, Visa confirmation 341513, in Prime, sending to 510 W Georgia St, Vancouver."



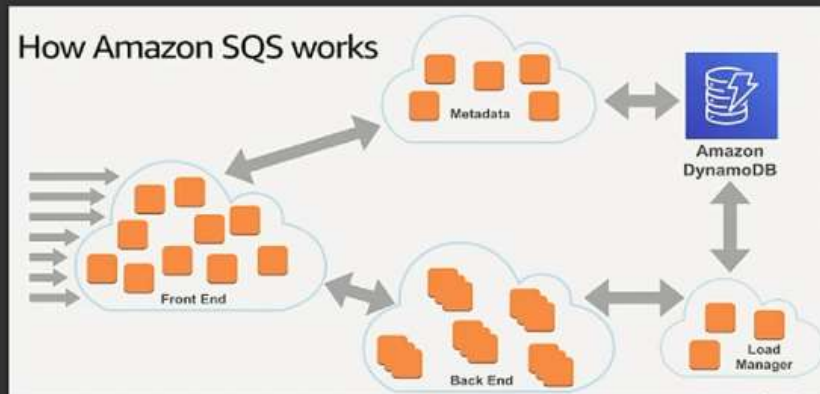
Is event-driven for you?

1. Are you passing around self-contained transactions?
2. Are useful events available for free?

A screenshot of the AWS Events console interface. The window is titled "Events" and has a close button (X) in the top right corner. Below the title bar, there are tabs for "Add notification", "Delete", and "Edit". The main content area is divided into sections for configuring a new event. The "Name" section has a text input field with the placeholder "e.g. MyFirstEventPut". The "Events" section contains a list of event types with checkboxes: PUT (checked), POST, COPY, Multipart upload completed, All object create events (checked), Object in RRS lost, Permanently deleted, Delete marker created, All object delete events, Restore initiated, and Restore completed. The "Prefix" section has a text input field with the placeholder "e.g. images/". The "Suffix" section has a text input field with the placeholder "e.g. .jpg". The "Send to" section has a dropdown menu with "Lambda Function" selected, and a list of other options: "SNS Topic", "SQS Queue", and "Lambda Function". At the bottom, there is a checkbox for "Active notifications" and "Cancel" and "Save" buttons.

Is event-driven for you?

1. Are you passing around self-contained transactions?
2. Are useful events available for free?
3. Do you need to strongly decouple your microservices?



Is event-driven for you?

1. Are you passing around self-contained transactions?
2. Are useful events available for free?
3. Do you need to decouple your microservices?
4. Do you need publish-and-subscribe?



Photo: Terje Skjerdal

What is an "event"?



A small AWS event

```
{
  "version": "0",
  "id": "f7a39f75-eff9-a823-5534-1075b196edd3",
  "detail-type": "EC2 Instance State-change Notification",
  "source": "aws.ec2",
  "account": "*****",
  "time": "2018-08-21T20:55:26Z",
  "region": "us-east-1",
  "resources": [ ],
  "detail": {
    "instance-id": "i-00b414b880501ae45",
    "state": "running"
  }
}
```

This is an event that tells you about a state change in EC2 from our EventBridge service. The **detail** field tells you what happened. The top-level part that don't change is called the event envelope

A larger AWS event

```
{
  "version": "0",
  "id": "5af0d99b-0841-2766-e5d5-06a865895fdf",
  "detail-type": "Support Ticket: Status Changed",
  "source": "aws.partner/zendesk.com/9242270/default",
  "time": "2019-05-25T01:23:45Z",
  "region": "us-east-1",
  "resources": [ ],
  "detail": {
    "ticket_event": {
      "type": "Status Changed",
      "previous": "open",
      "current": "solved",
      "ticket": {
        "id": 35436,
        "created_at": "2019-05-20T22:55:29Z",
        "updated_at": "2019-05-25T01:23:45Z",
        . . .
      }
    }
  }
}
```

This is another event generated from Zendesk to AWS for injecting events over a special API.

Example CloudEvent

```
{
  "specversion" : "1.0",
  "type" : "com.example.someevent",
  "source" : "/mycontext",
  "id" : "C234-1234-1234",
  "time" : "2018-04-05T17:31:00Z",
  "comexampleextension1" : "value",
  "comexampleothervalue" : 5,
  "datacontenttype" : "application/json",
  "data" : {
    "appinfoA" : "abc",
    "appinfoB" : 123,
    "appinfoC" : true
  }
}
```



CloudEvents

CloudEvents

<http://cloudevents.io>

This needs to have some data in it as below

AWS event → CloudEvent

```
func toCloudEvent(ae *awsevents.AWSEvent) (*cloudevents.Event, error) {
    ce := cloudevents.NewEvent();

    ce.SetType(ae.GetDetailType())
    awsSource := ae.GetSource()
    if (strings.HasPrefix(awsSource, "aws.")) {
        awsSource = awsSource[4:]
    }
    ce.SetSource("com.amazonaws/" + string(ae.GetRegion()) + "/" + awsSource)
    ce.SetID(ae.GetId())
    ce.SetTime(ae.GetTime());
    ce.SetExtension("awsregion", ae.GetRegion())
    ce.SetExtension("awsaccount", ae.GetAccount())
    ce.SetDataContentType("application/json")
    err := ce.SetData(ae.GetDetail())
    if err != nil {
        return nil, err
    }
    return &ce, nil
}
```

AWS event → CloudEvent

```
{
  "awsaccount": "123456789012",
  "awsregion": "us-east-2",
  "contenttype": "application/json",
  "data": {
    "instance-id": "i-00b414b880501ae45",
    "state": "running"
  },
  "id": "B977EE77-75E0-7A64-DD73-73A159BB4FF2",
  "source": "com.amazonaws/us-east-2/ec2",
  "specversion": "0.2",
  "time": "2019-10-28T19:56:14.195503Z",
  "type": "EC2 Instance State-change Notification"
}
```

This is the equivalent EC2 state change event wrapped up as a CloudEvent

What is an "event"?



An event can be a command, a fact that something happened, a record of what somebody/something did, etc. What we generally care about is what we can do regarding processing, reformatting etc of events.

Event processing facets

Facets allow us to structure our events

Event facet: Strict ordering



May vary



FIFO

Strict FIFO is only meaningful if you have a single sender and banking transactions will generally be FIFO but more expensive and less scalable

Event facet: Deduplication



Eventing
service



Does an event created get available and consumed exactly once or are you implementing at-least-once? Even when things go down. You can make sure APIs idempotent like the SQS createQueue API, or you can make your database backends transactional that can detect the state of a record and ignore other prior states.

Event facet: Point-to-point vs. publish/subscribe




Event facet: Push vs. pull



You can get the events by polling the source periodically or pushing on to topics or lambda endpoints. Polling allows you to consume data when you have the capacity to do that.

Event facet: Latency profile



Event facet: Serverless vs. broker/cluster

Amazon Simple Queue Service

Fully managed message queues for microservices, distributed systems, and serverless applications

Get started for free

Amazon Kinesis

Easily collect, process, and analyze video and data streams in real time

Get started with Amazon Kinesis

Amazon Simple Notification Service

Fully managed push/pull messaging for microservices, distributed systems, and serverless applications

Get started with Amazon SNS

RabbitMQ

ActiveMQ

APACHE
kafka®

There are 2 kinds of **eventing and messaging infrastructure**, Serverless and Broker/Cluster, which are actual servers that needs maintaining and will require scaling. Brokers mostly don't use **HTTP** (that is slower since it sets up a connection and tears it down) and instead use **TCP** or some framing protocols like **AMQP** that guarantees lower latency. There is now **HTTP long polling** and **HTTP2** for multiple transactions. **Quic or HTTP3** is UDP that is coming with very different latency profile and very fast.

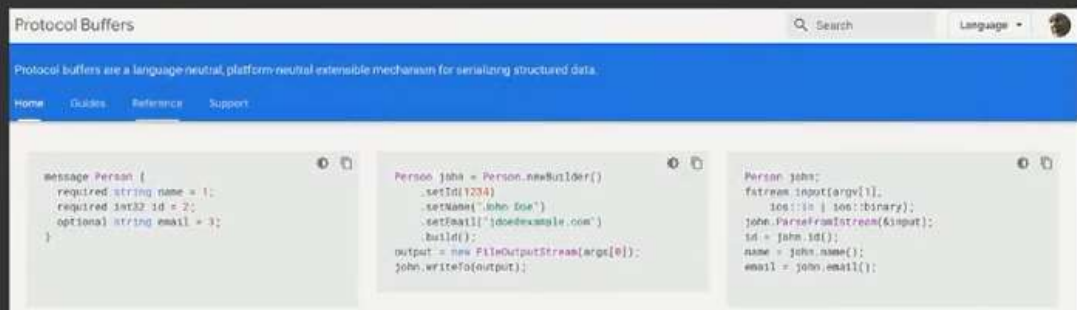
Event facet: Filtering vs. firehose

```
{
  "vendor": [ "vitamix", "instapot" ],
  "product_detail": {
    "price_usd": [
      { "numeric": ["<=", 150] }
    ]
  }
}
```

Most Pub/Sub and eventing APIs now come with filtering capabilities to reduce the streams data coming through to the event subscriptions

Event facet: Data blobs vs. structured objects

```
{
  "name": "John Doe",
  "id": 1234,
  "email": "johndoe@example.com"
}
```



(Also Avro, Thrift, Parquet, Cap'n Proto, Ion, etc.)

You can wrap/unwrap your data in JSON or blobs, binary formats like Avro, Protocol Buffers are beginning to be used even though they are smaller to transport but can't be unpacked easily.

Event facet: Records vs. documents

Activity Code	Activity Name
23	Patching
24	Overlay
25	Crack Sealing

Key = 24

Activity Code	Date	Route No.
24	01/12/01	I-95
24	02/08/01	I-66

Date	Activity Code	Route No.
01/12/01	24	I-95
01/15/01	23	I-495
02/08/01	24	I-66

```
{  
  "version": "0",  
  "id": "5af0d99b-0841-2766-e5d5-06a865895fdf",  
  "detail-type": "Support Ticket: Status Changed",  
  "source": "aws.partner/zendesk.com/9242270/default",  
  "time": "2019-05-25T01:23:45Z",  
  "region": "us-east-1",  
  "resources": [ ],  
  "detail": {  
    "ticket_event": {  
      "type": "Status Changed",  
      "previous": "open",  
      "current": "solved",  
      "ticket": {  
        "id": 35436,  
        "created_at": "2019-05-20T22:55:29Z",  
        "updated_at": "2019-05-25T01:23:45Z",  
        . . .  
      }  
    }  
  }  
}
```

A lot of the event content is business data that can be complex and spread over relational data. These are different from having JSON-like data that are deeply nested data.

Event facet: Uniform vs. heterogeneous events



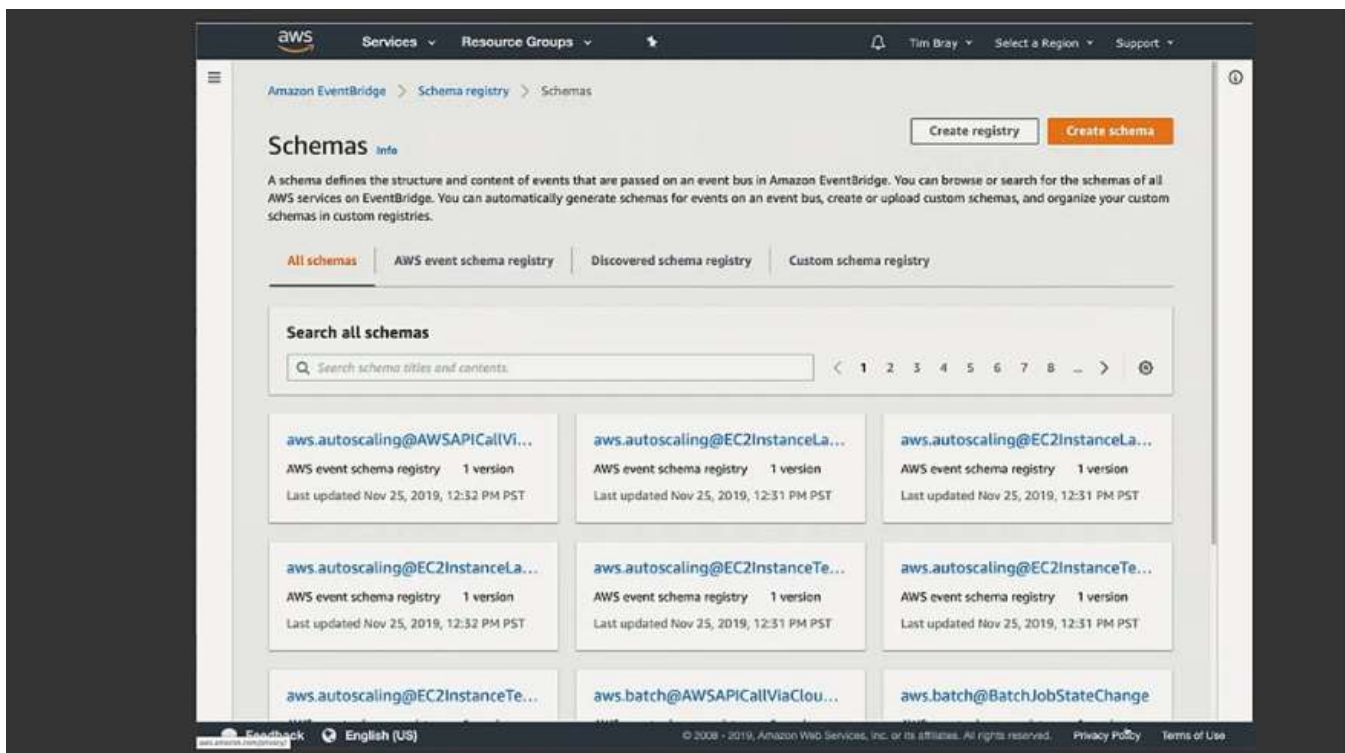
Are your events more the same or very different events and formats. This will affect the design of your event bus.


```

boolean isPassNode(final JsonNode node) {
    if (node.isObject()) {
        final JsonNode child = node.get(Constants.TYPE_FIELD);
        if (child != null) {
            if (child.isTextual()) {
                return Constants.PASS_TYPE.equals(child.asText());
            }
        }
    }
    return false;
}

```

Complexity means you might have to do things like above, finding what type your event payload is made of and then checking for fields in it. This is an anti-pattern.



We can use a schema registry to store schemas you use and make things easier. You can also use Athena to automatically generate your schemas for you automatically and upload them into the service registry.

aws

Services

Resource Groups

Tim Bray

Select a Region

Support

Amazon EventBridge > Schema registry > Schemas

Schemas

Info

Create registry

Create schema

A schema defines the structure and content of events that are passed on an event bus in Amazon EventBridge. You can browse or search for the schemas of all AWS services on EventBridge. You can automatically generate schemas for events on an event bus, create or upload custom schemas, and organize your custom schemas in custom registries.

All schemas

AWS event schema registry

Discovered schema registry

Custom schema registry

Search all schemas

ec2 instance state change

< 1 2 3 4 5 6 7 8 > ⚙

aws.ec2@EC2InstanceStateChan...

aws.opsworks@OpsWorksInstan...

aws.ssm@EC2StateManagerInst...

aws.emr@EMRInstanceFleetStat...

aws.dlm@DLMPolicyStateChange

aws.emr@EMRInstanceGroupSta...

aws.codedeploy@CodeDeployIn...

aws.emr@EMRClusterStateChange

aws.codepipeline@CodePipeline...

Feedback

English (US)

© 2008 - 2019, Amazon Web Services, Inc. or its affiliates. All rights reserved. Privacy Policy Terms of Use

aws

Services

Resource Groups

Tim Bray

Select a Region

Support

Amazon EventBridge > Schema registry > Schemas > aws.ec2

aws.ec2@EC2InstanceStateChangeNotification

Schema details

Schema name

Last modified

Schema ARN

aws.ec2@EC2InstanceStateChangeNotification

Nov 25, 2019, 12:32 PM PST

aws.events

Description

Number of versions

Event type

=

1

OpenApi3

Version 1 Created on Nov 25, 2019, 12:32 PM PST

Action

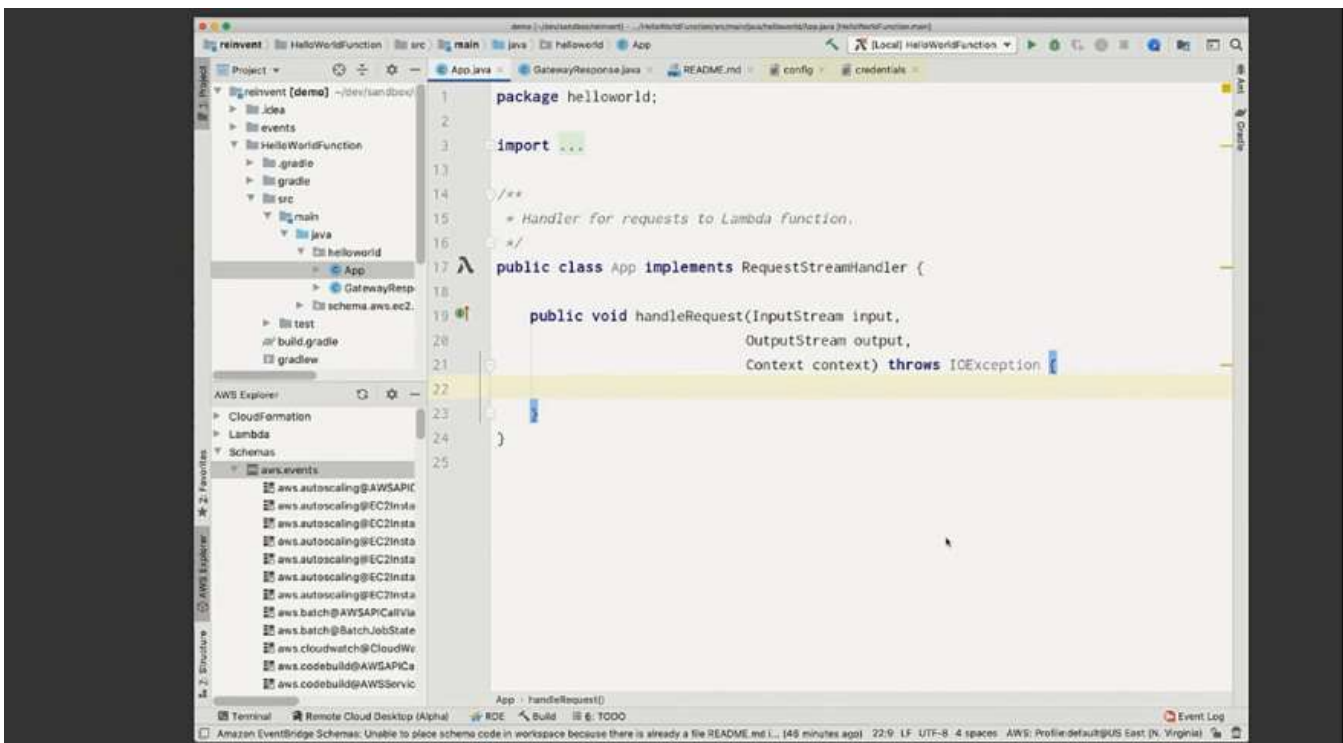
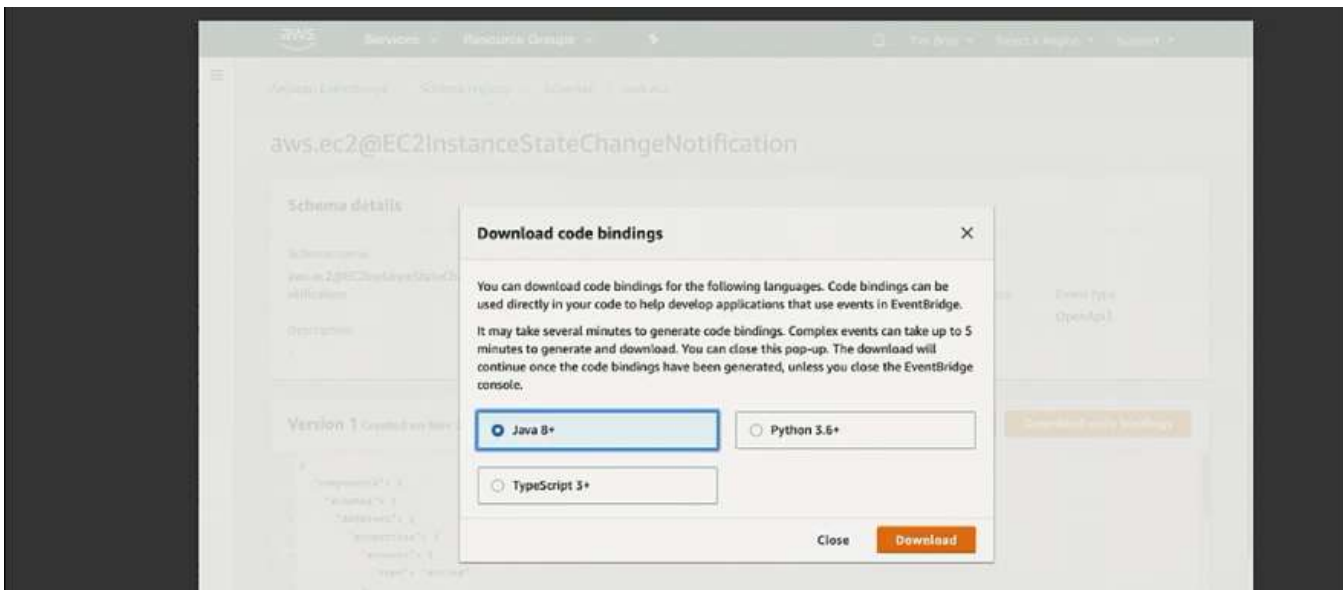
Download code bindings

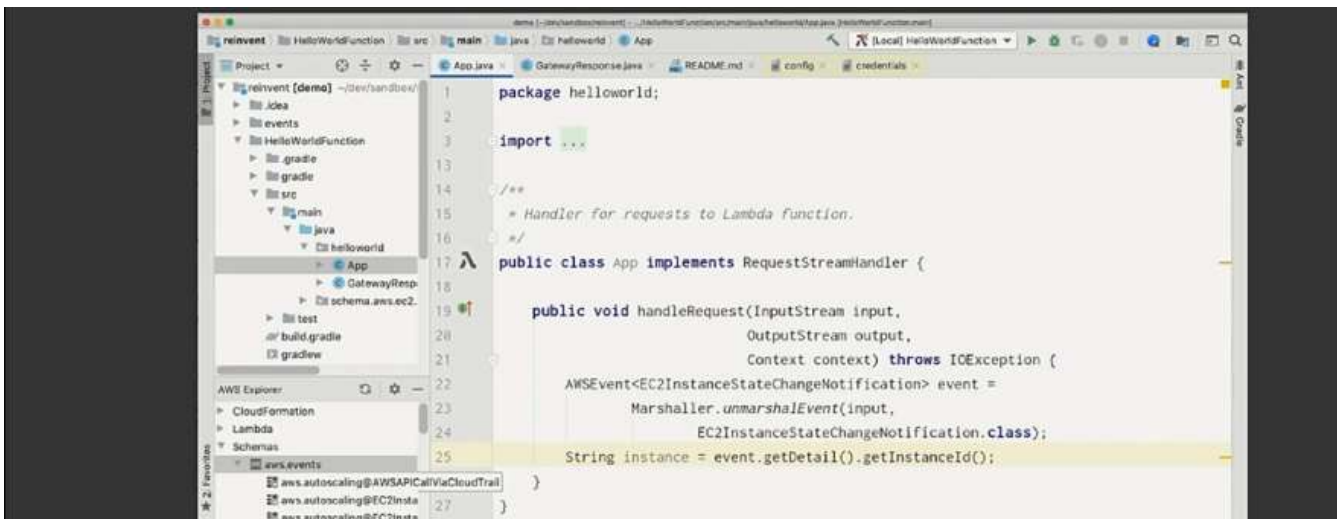
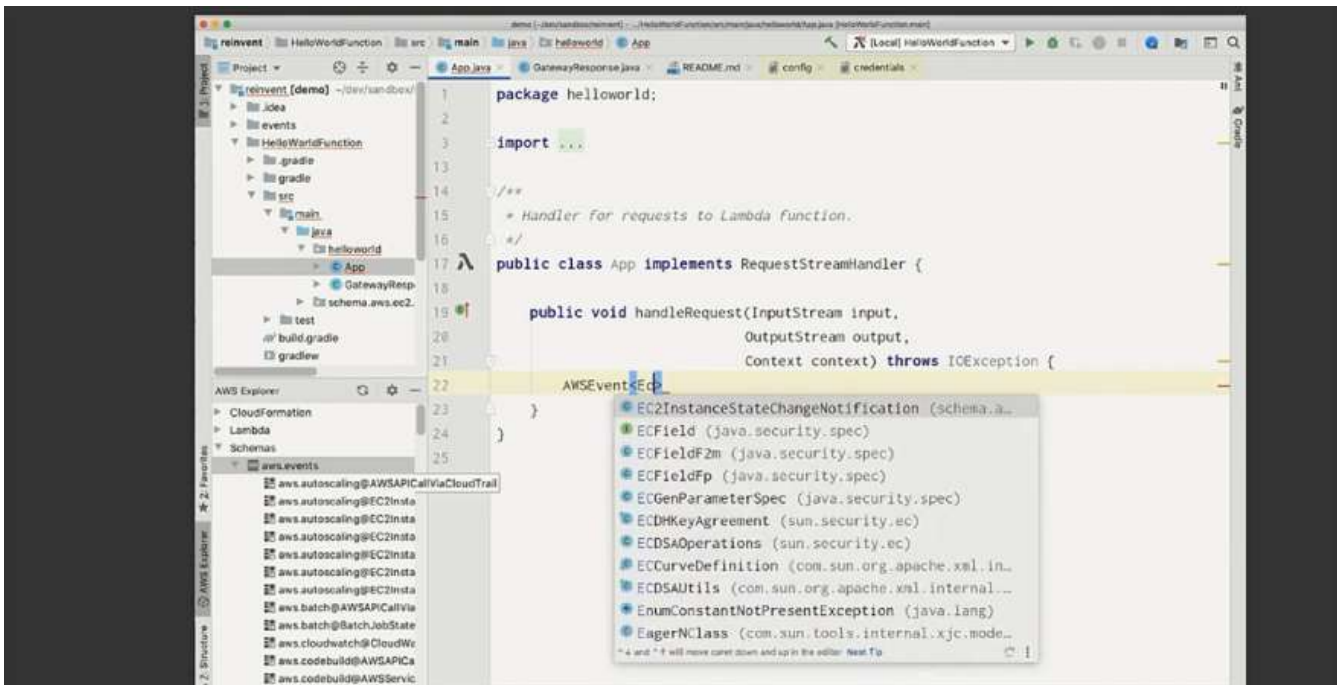
```
1 {
2   "components": {
3     "schemas": {
4       "AWSEvent": {
5         "properties": {
6           "account": {
7             "type": "string"
8           },
9           "detail": {
10            "ref": "#/components/schemas/EC2InstanceStateChangeNotification"
11          },
12          "detail-type": {
13            "type": "string"
14          },
15          "id": {
```

Feedback

English (US)

© 2008 - 2019, Amazon Web Services, Inc. or its affiliates. All rights reserved. Privacy Policy Terms of Use





Event facets

	ActiveMQ	Artemis	Event Bridge	Kafka	Amazon Kinesis	RabbitMQ	Amazon SNS	Amazon SQS	Amazon SQS FIFO
Ordering	Yes	Yes	No	Yes	Yes	Yes	No	No	Yes
Dedupe	*	Yes	No	Yes	No	No	No	No	Yes
P2p vs. Pub/sub	P/S	P/S	P/S	P/S	P/S	P/S	P/S	P2p	P2p
Push vs. Pull	Both	Both	Push	Pull	Pull	Both	Push	Pull	Pull
Serverless	No	No	Yes	No	Yes	No	Yes	Yes	Yes
Filtering	Yes	Yes	Yes	*	No	No	Yes	No	No

*: "It's complicated"

Events vs. ...

APIs

API thinking is associated with synchronous thinking, but we can use asynchronous features to it like making async call or using futures or doing request/response while doing things. You can also now fire off an API call and get the response back from a different queue.

Events vs. ...

**Service
mesh**



AWS App Mesh
Application-level networking for all your services

Get started with AWS App Mesh

You can also use a Service Mesh to help your work with APIs. These are smart proxies that allows you to skip calling APIs with DNS or IP addresses and instead call the generic service name like 'Customer Info Service'.

Eventing things that still need work:

1. Loops!
2. Too much pipefitting.
3. Easier and safer access control.
4. Blobby events, discovery, and autocomplete needed.


Learn serverless with AWS Training and Certification

Resources created by the experts at AWS to help you learn modern application development



Free, on-demand courses on serverless, including

- Introduction to Serverless Development
- Getting into the Serverless Mindset
- AWS Lambda Foundations
- Amazon API Gateway for Serverless Applications
- Amazon DynamoDB for Serverless Architectures



Additional digital and classroom trainings cover modern application development and computing

Visit the Learning Library at <https://aws.training>

© 2019, Amazon Web Services, Inc. or its affiliates. All rights reserved.



Thank you!

Tim Bray

timbray@amazon.com

[@timbray](https://twitter.com/timbray)

<https://tbray.org>

aws
re:Invent

© 2019, Amazon Web Services, Inc. or its affiliates. All rights reserved.

