

CON330

# Running Kubernetes clusters at scale: Bird

## Connor Poole

Principal Engineer  
Bird

## John Heroy

Engineering Manager  
Bird

## Omar Lari

Business Development Manager  
Amazon Web Services

aws  
re:Invent

© 2019, Amazon Web Services, Inc. or its affiliates. All rights reserved.



Launched in fall 2017, Bird is a micromobility company that enables access to shared e-scooters and lightweight electric vehicles in 100+ locations worldwide. Join us to hear how ***building a modern stack on top of Amazon EKS*** has enabled Bird to quickly ramp up its development in order to provide business value in a stable and secure manner. Further, learn how Bird's backend utilizes native AWS services like Amazon S3 and Amazon SQS, open-source streaming systems like Kafka and Flink, and a modern microservices architecture to turn terabytes of geospatial data into the mobility revolution of the future.

## Agenda



EKS and how it helps customers

Application evolution and EKS Adoption at Bird

Building a low latency always on IoT layer

Managing a platform on top of EKS

Running data streaming workloads in k8s

# AWS best practices for building modern applications

- Create a culture of innovation by **organizing into small DevOps teams**
- Continually evaluate your security posture by **automating security**
- Componentize applications using **microservices**
- Update applications & infrastructure quickly by **automating CI/CD**
- Standardize and automate operations by **modeling infrastructure as code**
- Simplify infrastructure management with **serverless technologies**
- Improve application performance by **increasing observability**

## What is Kubernetes?



Open source container management platform



Helps you run containers at scale



Gives you primitives for building modern applications

## How are customers using Amazon EKS?



Microservices



Platform as a service



Enterprise App Migration



Machine Learning



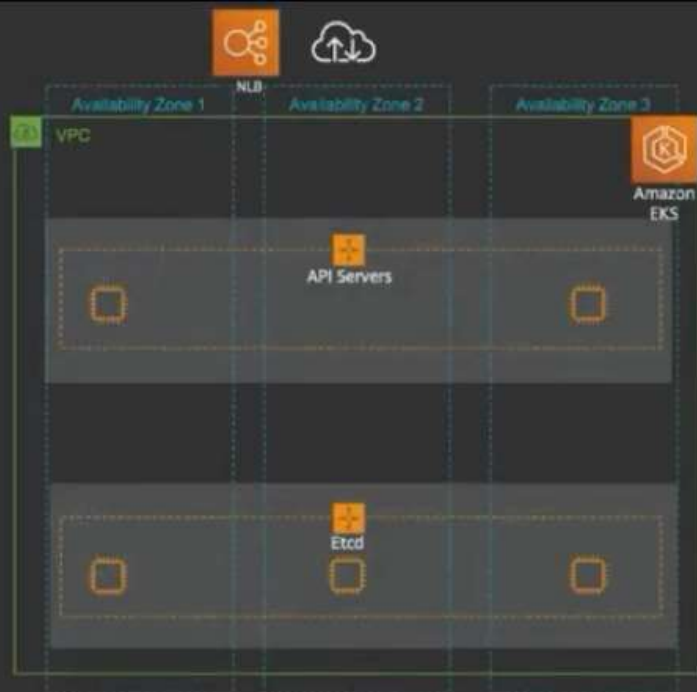
## Amazon Elastic Kubernetes Service

### Kubernetes control plane

Highly available and single  
tenant infrastructure

All “native AWS” components

Fronted by a Network Load  
Balancer



This is what you get when you run a ‘\$ eks create cluster’ command in your account, you get a HA etcd cluster to store cluster data and a HA k8s control plane (the API Servers, Controller Managers, etc.) that gets exposed to you via a Network LB to connect you client tooling like kubectl or your other apps. These are all encrypted.

### Our tenets

1. Amazon EKS is a platform to run **production-grade workloads**.
2. Amazon EKS provides a **native and upstream Kubernetes** experience.
3. If you want to use additional AWS services, **integrations** are as **seamless** as possible.
4. The Amazon EKS team in AWS **actively contributes** to the **upstream Kubernetes** project.

# Amazon EKS, a year in review

**June – December 2018:**

- Amazon EKS achieves K8s conformance, HIPAA-eligibility, Generally available
- Amazon EKS AMI build scripts and AWS CloudFormation templates available in GitHub.
- Support for GPU-enabled EC2 instances, support for HPA with custom metrics.
- Amazon EKS launches in Dublin, Ireland
- Amazon EKS simplifies cluster setup with update-kubeconfig CLI command
- Amazon EKS adds support for Dynamic Admission Controllers (Istio), ALB Support with the AWS ALB ingress controller
- Amazon EKS launches in Ohio, Frankfurt, Singapore, Sydney, and Tokyo
- Amazon EKS adds Managed Cluster Updates and Support for Kubernetes Version 1.11, CSI Driver for Amazon EBS

## 2019:

Amazon EKS launches in Seoul, Mumbai, London, and Paris  
Amazon EKS achieves ISO and PCI compliance, announces 99.9% SLA, cluster creation limit raised to 50  
API Server Endpoint Access Control, AWS App Mesh controller  
Windows support (preview), Kubernetes version 1.12,  
CSI Drivers for Amazon EFS, Amazon FSx for Lustre, Control Plane Logs, A1 (ARM) instance support (preview)  
Deep Learning Benchmark Utility, Public IP Address Support,  
Simplified cluster authentication, SOC compliance, Kubernetes 1.13, PodSecurityPolicies,  
Container Insights, CNI 1.5.0, Amazon ECR, AWS PrivateLink Support  
Managed Nodes

## Open-source roadmap

[illegible]



# Amazon EKS services roadmap: Highlights

## Shipped

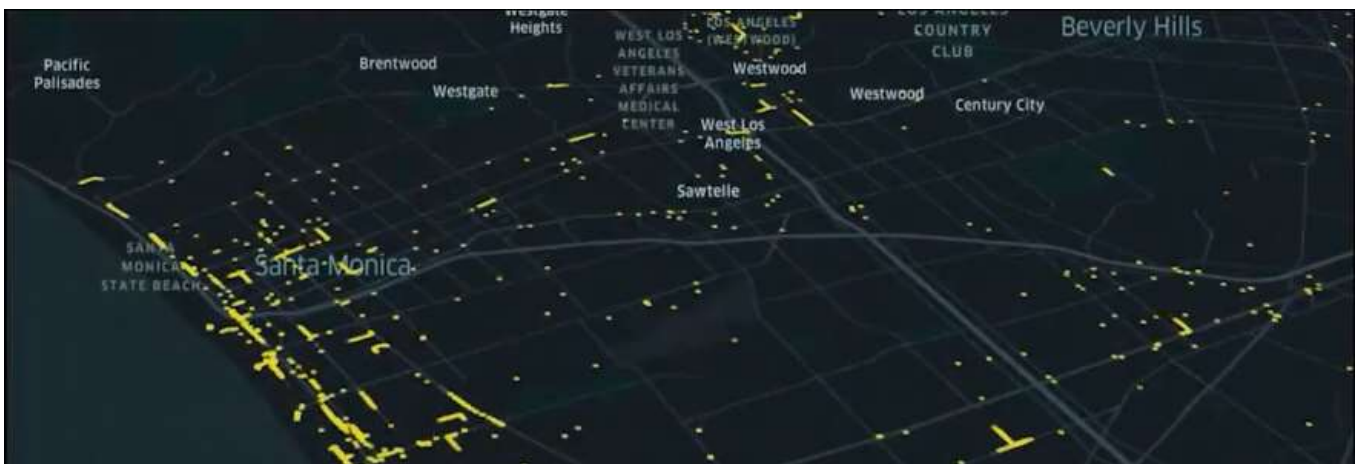
- IAM Roles for Service Accounts (pods)
- Managed Node Groups
- K8s version 1.13, 1.14
- ECR Private Link
- New Amazon EKS Regions: Paris, London, Mumbai, Hong Kong, Sao Paulo, Montreal

## Coming soon

- Fargate for EKS
- Service Linked Role
- KMS Encryption Provider
- New Amazon EKS Regions: Beijing, Ningxia

## Working on it

- Managed add-ons
- DNS resolution of Amazon EKS private endpoints
- New Amazon EKS Regions
- Next-generation CNI plugin



# Building a platform on a deadline



## Reduce, reuse, recycle



Infrastructure:  
Terraform



CI/CD:  
Jenkins

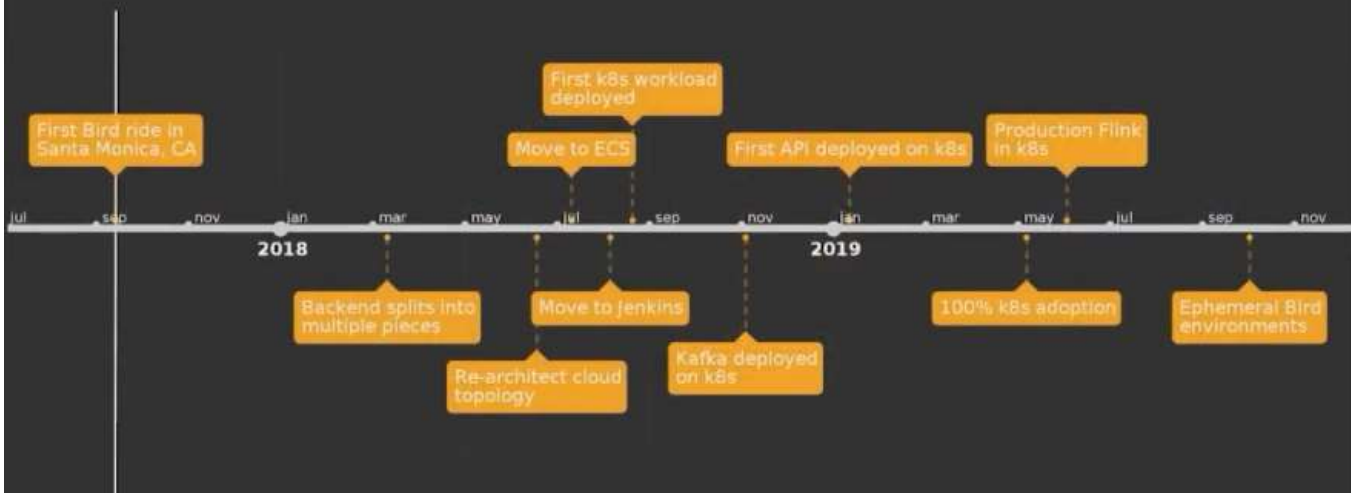


Orchestration:  
Amazon EKS



Managed  
Services:  
AWS

## Timeline



## September 2017

- First Bird ride in Santa Monica
  - Modified retail scooters
  - Phone apps and a single backend



## March 2018

- Second deployable emerges
  - Fleet size has grown considerably
    - Thousands of Birds
    - Multiple markets in the US
  - Began decoupling batch processes into "workers"


## June 2018

- Start fresh in AWS
  - Re:Invent the wheel



# Cloud topology


- Hub-and-spoke topology
- Accounts are mostly 1:1 with spokes



# VPC design

Standard three-level VPC design

Public, private, Fort Knox





# Roll with the punches, and embrace detours

- Outages frequently
- Deploying jars from Amazon S3 wasn't working for us
- Multiple environments + CI environments needed more encapsulation
- Amazon EKS was still in beta at this time, so we migrated fully to Amazon Elastic Container Service (Amazon ECS) in under 30 days



## Why Amazon ECS?

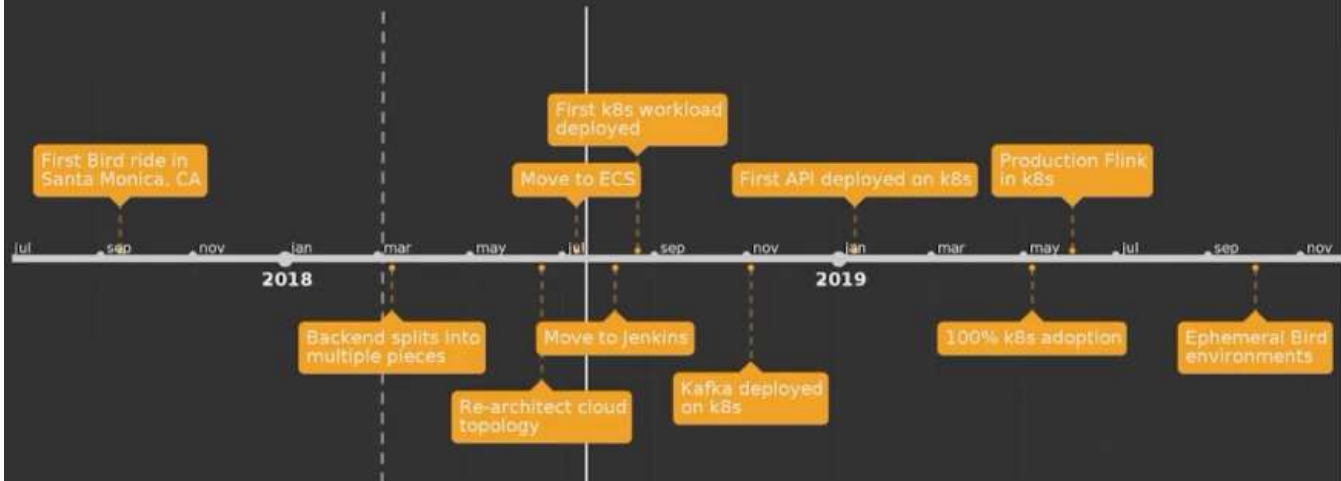
We needed to get onto containers ASAP

Amazon ECS is simple to set up and had native AWS Identity and Access Management (IAM) integration

Was a smaller jump from Amazon EC2 to Amazon EKS



## Timeline recap



## Why Amazon EKS?

Supports advanced stateful workloads

Amazon EKS removes all the headaches of k8s control plane management

Native VPC integration



## Managing Amazon EKS

- Prune your clusters to a known state on startup
- Use large workers for bin packing
- Rely on Cluster Autoscaler + Horizontal Pod Autoscaler
- Calico for network policy

# Who orchestrates the orchestrator?

```
module "eks" {  
  vpc_id          = module.vpc.id  
  security_group_ids = [ ... ]  
  
  efs_security_group = module.vpc.aws_security_group-allow_all_within_vpc.id  
  efs_subnets       = module.vpc.private_subnets.*.id  
  
  subnet_ids = concat(module.vpc.public_subnets.*.id,module.vpc.private_subnets.*.id)  
  
  name_prefix      = var.eks_cluster_prefix  
  
  private_ingress_ssl_cert      = aws_acm_certificate.cert.arn  
  private_ingress_security_groups = module.vpc.aws_security_group-vpc_local.id  
  public_ingress_ssl_cert       = aws_acm_certificate.cert.arn  
  public_ingress_security_groups = module.vpc.aws_security_group-external_access.id  
  
  vault_auth_enabled = 1  
}
```

We use Terraform to manage our K8s clusters across our account, this is what a module looks like in 20 lines

## Partitioning a cluster

- Cluster sprawl is real
- Namespaces help clamp the chaos
- Choose your divisions wisely; too many namespaces is equally bad



Use namespaces to make pods and items naming and location easy

## Choose a k8s deployment strategy

- Minimize cruft by handling all object creation in a unified manner
- We chose Helm for its flexibility and community support
- Deployment history through configmaps provides a clean API to code against

Use I deployment strategy, we use Helm because it provides a clean API that we can code against and work with

"You'd have to be crazy to run stateful mission-critical systems on Kubernetes."

**Kubernetes Community – 2019**

Except you are doing it with K8s on EKS

## Data pipeline begins

- 216 nodes per cluster
- 500 GB per disk
- 10 Kafka clusters per EKS cluster
- Be prepared for disk space resizing at 11 p.m. on your laptop, pulled over to the side of the road

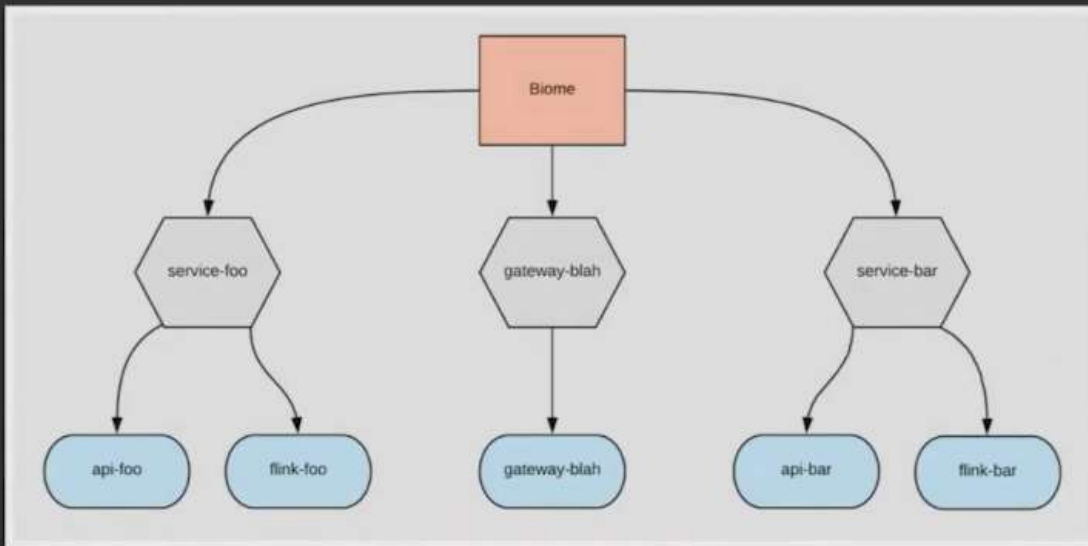


This is the first workflow we deployed on EKS, we run multiple Kafka clusters on EKS




# Monolith decomposition begins

- Organize your apps, data, and configs in a standard structure




Colocate your data sources



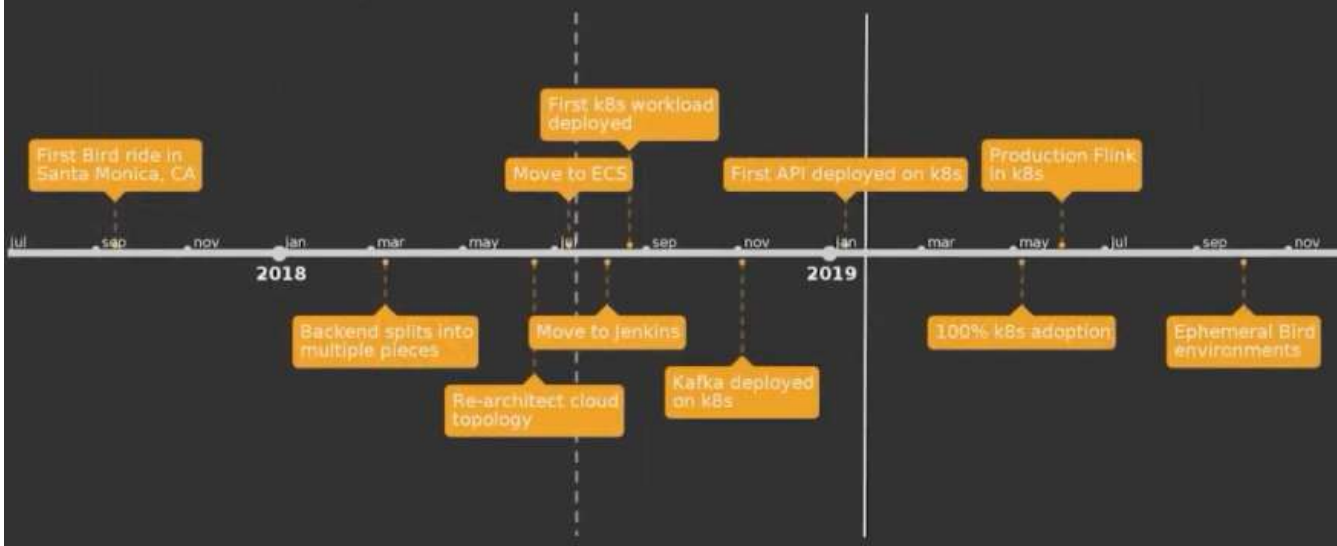
Use namespace properly to grant restricted access to teams

## To mesh or not to mesh

- Do you really need a full mesh?
- Ingress-nginx was good enough for us
- It is simple to set up, scale, and run locally




## Timeline recap



## Let's talk about vehicles

- Bird has built a large network of always-connected IoT devices
- The backend should always be able to connect to our vehicles




## Controlling TCP

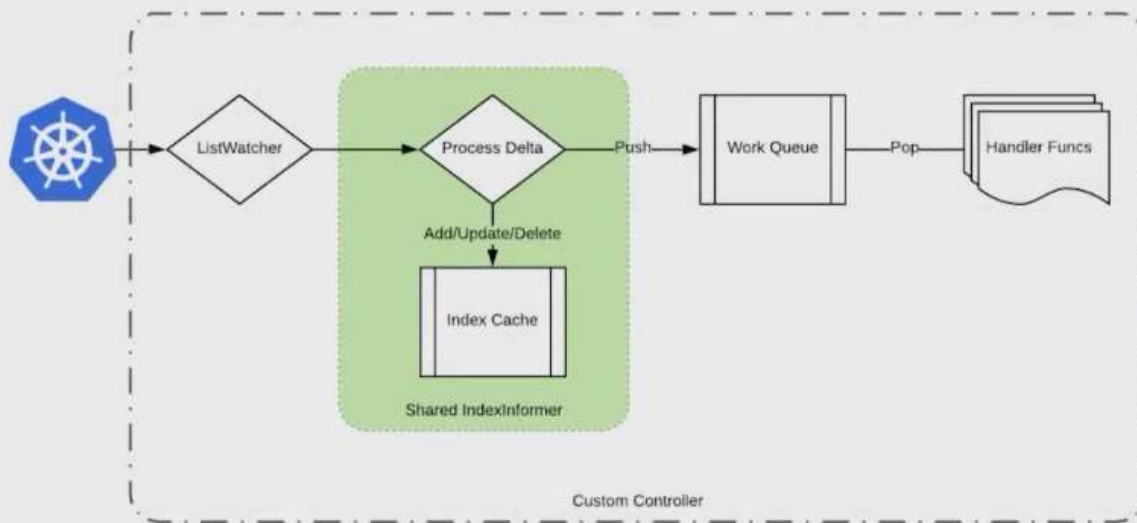
- We needed to ensure infinite socket life for our IoT gateway
- NLB integration with Kubernetes wasn't enough

# NLB attacher

- Direct from NLB to pods
- Lower latency, fewer moving pieces
- Most important – infinite socket life



## Building a k8s controller



We extended our K8s with custom controllers



# Core controller

```
// Controller - the primary struct responsible for all cluster actions
type Controller struct {
    clientset      kubernetes.Interface
    queue          workqueue.RateLimitingInterface
    informer       cache.SharedIndexInformer
    eventHandler   handlers.Handler
    config          config.Config
    serverStartTime time.Time
}

const maxRetries int = 5
const enableLabelValue string = "nlb-attacher.bird.co/enabled"
const targetGroupAnnotationKey string = "nlb-attacher.bird.co/target-groups"
```

# Handler interface

```
// Handler is implemented by any handler.
// The Handle method is used to process event
type Handler interface {
    Init(tgAnnotation string, annotationEnabledValue string) error
    PodCreated(created *v1.Pod)
    PodDeleted(deleted *v1.Pod)
    PodUpdated(oldPod, newPod *v1.Pod)
    TestHandler()
}
```

# Adding pods

```
// PodCreated - Handle the creation event of a pod and ensure it's attached to all of the specified target groups
func (handler *Handler) PodCreated(created *v1.Pod) {
    if created.DeletionTimestamp != nil {
        log.Infof("Discovered pod %s with deletionTimestamp already set. skipping...", created.Name)
        return
    }

    if created.Status.PodIP == "" {
        log.Infof("Recieved event for pod %s without an ip address yet. skipping...", created.Name)
        return
    }

    log.Debugf("created object: %v", created.Name)
    handler.addToTargetGroups(created)
}
```

# From Bird with love

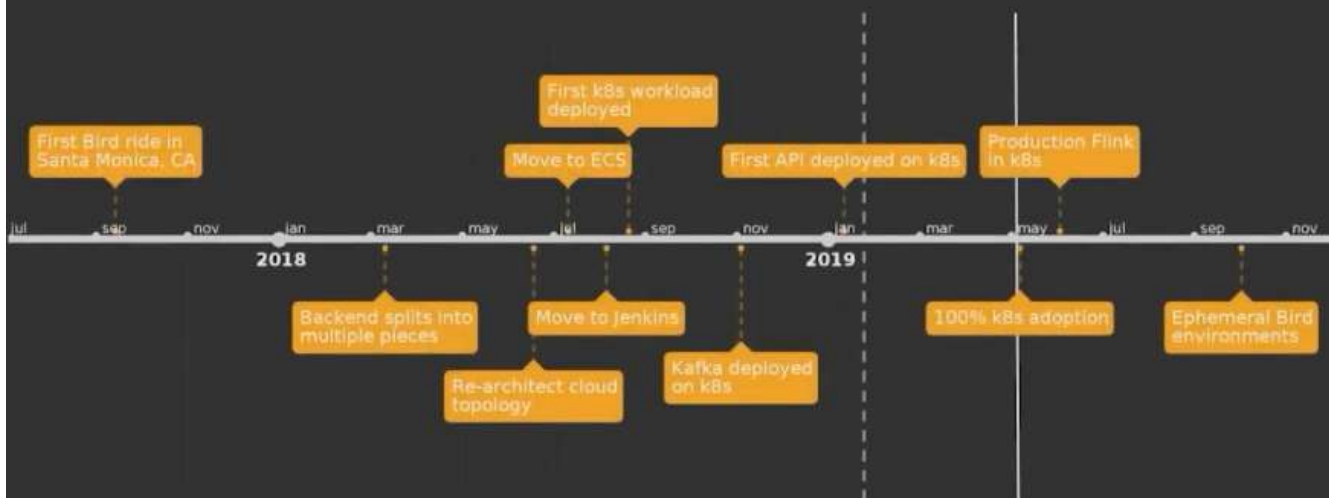
Open-sourcing this tool today:

<https://github.com/birdrides/nlb-attacher>

Future roadmap:

- Pod UDP support
- Target group creation

## Timeline recap



## Clamping state drift

- You're never going to fully solve this
- Carrot, not the stick
- Centralize all deployment pieces in as few graphs as possible
- Consider using CRDs to represent as many pieces of critical infrastructure as possible

## But be realistic

- Terraform is still our favorite tool today
  - A single graph can combine Helm, Kubernetes, Vault, CDN providers, etc.
- The interface between Helm and Terraform is poor
  - A brittle, weakly typed, whitespace-sensitive contract will cause headaches
  - Perform your logic inside common Helm templating libraries or inside Terraform modules that wrap Helm
- Look to the future:
  - <https://github.com/aws/aws-service-operator-k8s>
  - <https://github.com/rancher/terraform-controller>

## A platform is born

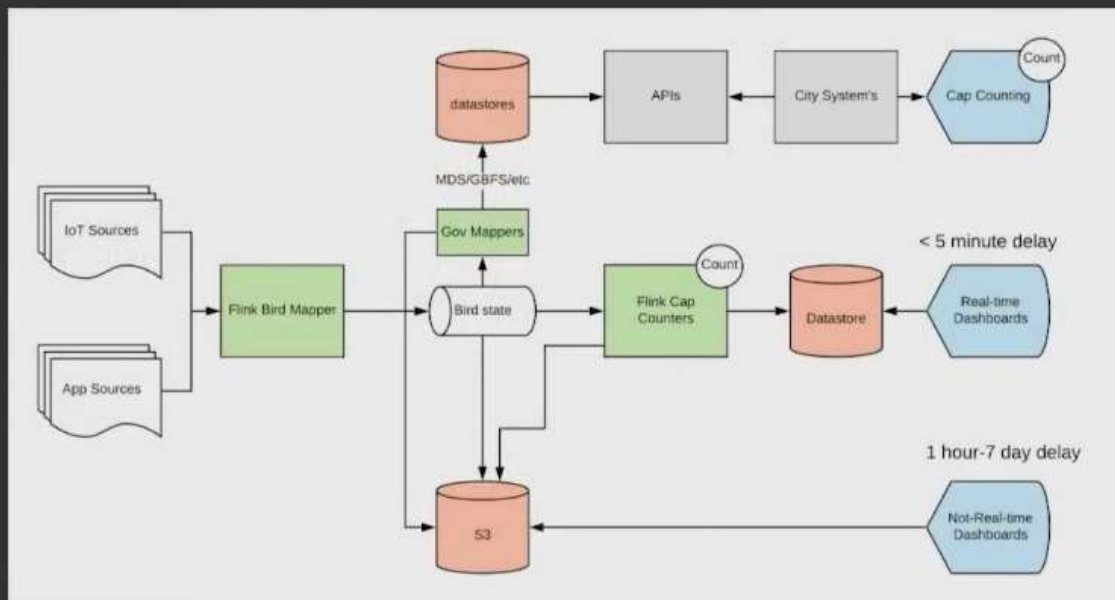
- Understand your users
  - "DevOps" is a cultural goal in moderation
- Kubernetes evolves quickly
  - Abstraction is your friend – keep your abstraction
    - Keep your abstraction thin – shim cloud native
- Tighten the feedback loop
  - Instant gratification is key
    - Users must see results/feedback during their testing loops

## Interacting with our platform

```
{
  "name": "service-habitat",
  "dependencies": [
    "sqs",
    "postgres"
  ],
  "deployables": [
    {
      "name": "api-habitat",
      "dependencies": [
        "service-habitat"
      ],
      "language": {
        "name": "kotlin",
        "properties": {
          "framework": "vertx",
          "main_class": "co.bird.habitat.api.HabitatVerticle"
        }
      },
      "type": {
        "name": "http-api"
      }
    }
  ]
}
```

This is a Bird service that we abstracted

## Flink





## But why Flink?



We need data  
in real time



State transitions  
are infrequent  
but latency-sensitive



Native k8s  
support



Too many vehicles  
for batch




Best in memory state  
management, snapshot,  
and debugging tooling

## Why not Flink?

- Flink is still new, bugs exist, and we regularly find them
- Kafka + Flink have silent partition issues
- You're not already on k8s


IoT signals are messy

Where do you think this Bird is?



IoT signals are messy

The Bird is in a superposition of states!



Accurate signals are critical

Location:


- Within polygons
- Distance from other POI

Stream processing  
is flexible

Quickly grow our ecosystem

- Events
- Publishers
- Consumers

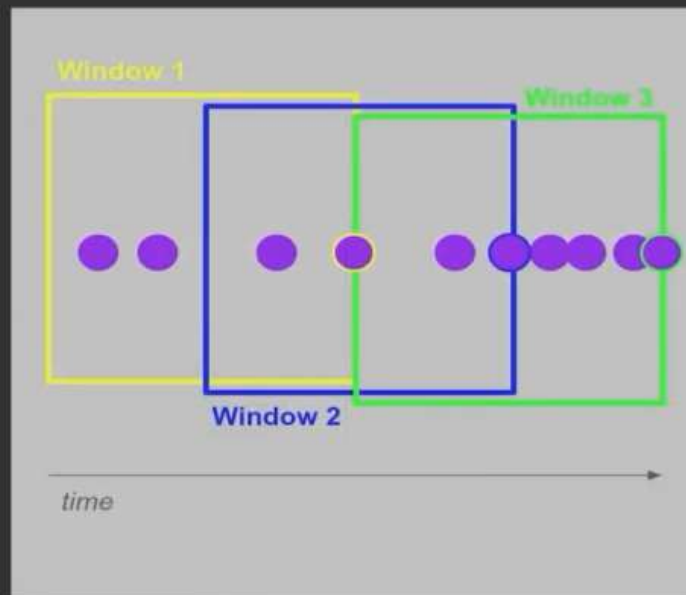
Derived and raw events



Kafka is our global event bus


## Let's process locations

- Sliding windows (Flink)
- Compute average of points




We declare a Flink Job

Turn this



Into this!



```
Activities  Terminal
WEDNESDAY 12/09/19  2% 14%
Terminal
connor:~/git/github.com/birdrides/api/api/service-derived-location
$ ls
manifest.json manifest.json_bak
connor:~/git/github.com/birdrides/api/api/service-derived-location
$ vim manifest.json
```

```
Activities  Terminal
WEDNESDAY 12/09/19  2% 14%
Terminal
3  "dependencies": [
4
5  ],
6  "deployables": [
7    {
8      "name": "location-averager",
9      "dependencies": [],
10     "language": {
11       "name": "kotlin"
12     },
13     "type": "f",
14     "name": "flink"
15   }
16 ]
17 }
18 }
```

```
Activities  Terminal
WEDNESDAY 12/09/19  2% 14%
Terminal
connor:~/git/github.com/birdrides/api/api/service-derived-location
$ ls
manifest.json manifest.json_bak
connor:~/git/github.com/birdrides/api/api/service-derived-location
$ vim manifest.json
connor:~/git/github.com/birdrides/api/api/service-derived-location
$ clear^C
connor:~/git/github.com/birdrides/api/api/service-derived-location
$ ../../cli/bin/bird tooling service update --service-name=service-derived-location
```

```
Activities  Terminal
WEDNESDAY 12/09/19  2% 14%
Terminal
⚠ CREATING file @ /home/connor/git/github.com/birdrides/api/api/service-derived-location/location-averager/src/main/resources/logback.xml
⚠ CREATING directory /home/connor/git/github.com/birdrides/api/api/service-derived-location
⚠ CREATING file @ /home/connor/git/github.com/birdrides/api/api/service-derived-location/helm/service-derived-location/Chart.yaml
⚠ CREATING file @ /home/connor/git/github.com/birdrides/api/api/service-derived-location/helm/service-derived-location/values.yaml
⚠ CREATING file @ /home/connor/git/github.com/birdrides/api/api/service-derived-location/Jenkinsfile
⚠ CREATING file @ /home/connor/git/github.com/birdrides/api/api/service-derived-location/terraform/main.tf
⚠ CREATING file @ /home/connor/git/github.com/birdrides/api/api/service-derived-location/terraform/locals.tf
⚠ CREATING file @ /home/connor/git/github.com/birdrides/api/api/service-derived-location/terraform/remote.tf
⚠ CREATING file @ /home/connor/git/github.com/birdrides/api/api/service-derived-location/terraform/vars.tf
⚠ CREATING file @ /home/connor/git/github.com/birdrides/api/api/service-derived-location/terraform/namespace.tf
⚠ CREATING file @ /home/connor/git/github.com/birdrides/api/api/service-derived-location/terraform/extra_vars.tf
⚠ CREATING file @ /home/connor/git/github.com/birdrides/api/api/service-derived-location/terraform/global/ecr.tf
⚠ CREATING file @ /home/connor/git/github.com/birdrides/api/api/service-derived-location/terraform/global/local.tf
⚠ CREATING file @ /home/connor/git/github.com/birdrides/api/api/service-derived-location/terraform/global/main.tf
⚠ CREATING file @ /home/connor/git/github.com/birdrides/api/api/service-derived-location/terraform/global/vars.tf
⚠ CREATING file @ /home/connor/git/github.com/birdrides/api/api/service-derived-location/terraform/global/vault.tf
connor:~/git/github.com/birdrides/api/api/service-derived-location
$ ls
helm Jenkinsfile location-averager manifest.json manifest.json_bak terraform
connor:~/git/github.com/birdrides/api/api/service-derived-location
$
```

We now have Helm in this directory with the Helm chart



```

$ tree
.
├── helm
│   └── service-derived-events
│       ├── Chart.yaml
│       └── values.yaml
├── Jenkinsfile
├── location-averager
│   ├── build.gradle.kts
│   ├── Dockerfile
│   ├── extras.groovy
│   └── helm
│       └── location-averager
│           ├── Chart.yaml
│           ├── templates
│           │   ├── configmap.yaml
│           │   ├── deployment.yaml
│           │   ├── helpers.tpl
│           │   ├── ingress.yaml
│           │   ├── NOTES.txt
│           │   ├── pod_disruption_budget.yaml
│           │   ├── rbac.yaml
│           │   └── service.yaml
│           └── values.yaml
├── Jenkinsfile
├── src
│   └── main
│       ├── kotlin
│       └── ...

```

```
└─ values.yaml
Jenkinsfile
location-averager
└─ build.gradle.kts
Dockerfile
extras.groovy
helm
└─ location-averager
    └─ Chart.yaml
        └─ templates
            └─ configmap.yaml
            └─ deployment.yaml
            └─ helpers.tpl
            └─ ingress.yaml
            └─ NOTES.txt
            └─ pod_disruption_budget.yml
            └─ rbac.yaml
            └─ service.yaml
            └─ values.yaml
Jenkinsfile
src
└─ main
    └─ kotlin
        └─ co
            └─ bird
                └─ flink
                    └─ demo
                        └─ DemoSplitFirehoseJob.kt
                        └─ DemoTemplateJob.kt
    └─ resources
        └─ logback.xml
```

```
└─ resources
    └─ DemoTemplateJob.kt
    └─ logback.xml
terraform
└─ extra_helm_yaml.tf
└─ extra_vars.tf
└─ helm.tf
└─ iam_json.tf
└─ locals.tf
└─ main.tf
└─ remote.tf
└─ vars.tf
version.json
manifest.json
manifest.json_bak
terraform
└─ extra_vars.tf
    └─ global
        └─ ecr.tf
        └─ locals.tf
        └─ main.tf
        └─ vars.tf
        └─ vault.tf
└─ locals.tf
└─ main.tf
└─ namespace.tf
└─ remote.tf
└─ vars.tf

.7 directories, 42 files
```

```

├── extra_vars.tf
├── helm.tf
├── iam_json.tf
├── locals.tf
├── main.tf
├── remote.tf
├── vars.tf
├── version.json
├── manifest.json
├── manifest.json_bak
└── terraform

├── extra_vars.tf
├── global
├── ecr.tf
├── locals.tf
├── main.tf
├── vars.tf
├── vault.tf
├── locals.tf
├── main.tf
├── namespace.tf
├── remote.tf
├── vars.tf

17 directories, 42 files
connor:~/git/github.com/birdrides/api/api/service-derived-location
$ ls
helm Jenkinsfile location-averager manifest.json manifest.json_bak terraform
connor:~/git/github.com/birdrides/api/api/service-derived-location
$
```

```

1 package co.bird.flink.tracks
2
3 import ...
4
51 private const val DEFAULT_JOB_NAME = "Bird Location Averager Job"
52
53 /**
54  * Averages locations from Bird tracks over sliding (overlapping) time windows.
55  */
56
57 class LocationAveragerJob {
58
59     fun build(trackStream: SingleOutputStreamOperator<BirdTrack>): SingleOutputStreamOperator<AveragedBirdLocation> {
60         return trackStream
61             .filter { it?.location != null && it.location!!.isValidLatLong() }
62             .returns(BirdTrack::class.java)
63             .keyBy({ it.birdId.toString() }, Types.STRING)
64             .window(SlidingProcessingTimeWindows.of(Time.seconds( seconds: 60L), Time.seconds( seconds: 30L)))
65             .process()
66     }
67
68     companion object {
69         @JvmStatic
70         fun main(args: Array<String>) {
71             LocationAveragerJob().build()
72         }
73     }
74 }
```

This is our Flink Job code

```

37 class LocationAveragerJob {
38
39     fun build(trackStream: SingleOutputStreamOperator<BirdTrack>): SingleOutputStreamOperator<AveragedBirdLocation> {
40         return trackStream
41             .filter { it?.location != null && it.location!!.isValidLatLong() }
42             .returns(BirdTrack::class.java)
43             .keyBy({ it.birdId.toString() }, Types.STRING)
44             .window(SlidingProcessingTimeWindows.of(Time.seconds( seconds: 60L), Time.seconds( seconds: 30L)))
45             .process(AverageLocationFunction())
46     }
47 }
```

```

$ kubectl get pods -n service-derived-locations
NAME                                READY   STATUS    RESTARTS   AGE
link-location-averager-jobmanager   1/1     Running   0           12h
link-location-averager-taskmanager-79746fdbdc-kfw7  1/1     Running   0           12h
link-location-averager-taskmanager-79746fdbdc-tk222  1/1     Running   0           12h
$ kubectl get pods -n service-derived-locations
NAME                                READY   STATUS    RESTARTS   AGE
link-location-averager-jobmanager   1/1     Running   0           14h
link-location-averager-taskmanager-79746fdbdc-kfw7  1/1     Running   0           14h
link-location-averager-taskmanager-79746fdbdc-tk222  1/1     Running   0           14h
$

```

```

[Error Code: 404 Not Found; Request ID: DF83B4A3B880992C5; S3 Extended Request ID: F9dLHtPrpSLNN0VwuA1DWE0K65kREiUk0SG0LEP3aFGlNr0gmUu6Phy4uo5n0KKn0tyHokw); S3 Extended Request ID: F9dLHtPrpSLNN0VwuA1DWE0K65kREiUk0SG0LEP3aFGlNr0gmUu6Phy4uo5n0KKn0tyHokw); RequestType: GetObject(MetadataRequest); AWSRequestID=[DF83B4A3B880992C5], HttpClientPoolPendingCount=0, RequestCapacityConsumed=0, HttpClientPoolAvailableCount=0, RequestCount=1, Exception=1, HttpClientPoolLeasedCount=0, ClientExecuteTime=[26.871], HttpClientSendRequestTime=[0.049], HttpClientRequestTime=[26.416], RequestSigningTime=[0.148], CredentialsRequestTime=[0.003, 0.0], HttpClientReceiveResponseTime=[16.138]].
C
error:-
$ kubectl logs -f -n service-derived-locations flink-location-averager-taskmanager-79746fdbdc-kfwm7

```

```
Terminal
[{"bird_id": "46842bc1-5a0f-4f29-b616-5c5a5840d4b9", "created_at": "1575493656000", "location": {"string": "POINT(-118.47282333333332 34.03032666666666)"}, "event_id": "6a33bc7e-2e81-44a0-85e4-95931190e027", "tracks_count": 2, "period_seconds": 600, "event_time": "1575494480000"}, {"bird_id": "46842bc1-5a0f-4f29-b616-5c5a5840d4b9", "created_at": "1575494011000", "location": {"string": "POINT(-118.47282333333332 34.030315)"}, "event_id": "416065ba-035c-41ca-98ad-aa4469a700f9", "tracks_count": 2, "period_seconds": 600, "event_time": "1575494780000"}, {"bird_id": "46842bc1-5a0f-4f29-b616-5c5a5840d4b9", "created_at": "1575494363000", "location": {"string": "POINT(-118.47282333333332 34.030303333333336)"}, "event_id": "46842bc1-5a0f-4f29-b616-5c5a5840d4b9", "tracks_count": 1, "period_seconds": 600, "event_time": "1575495000000"}, {"bird_id": "46842bc1-5a0f-4f29-b616-5c5a5840d4b9", "created_at": "1575494668000", "location": {"string": "POINT(-118.47282333333332 34.03030833333333)"}, "event_id": "088e4080-b44c-4630-838b-88ed5374309d", "tracks_count": 1, "period_seconds": 600, "event_time": "1575495600000"}, {"bird_id": "46842bc1-5a0f-4f29-b616-5c5a5840d4b9", "created_at": "1575495226000", "location": {"string": "POINT(-118.47282333333332 34.030303333333336)"}, "event_id": "98188706-b0c3-4fa3-9505-1750e1c180ff", "tracks_count": 2, "period_seconds": 600, "event_time": "1575495600000"}, {"bird_id": "46842bc1-5a0f-4f29-b616-5c5a5840d4b9", "created_at": "1575495608000", "location": {"string": "POINT(-118.47282333333332 34.030308000000004)"}, "event_id": "46842bc1-5a0f-4f29-b616-5c5a5840d4b9", "tracks_count": 1, "period_seconds": 600, "event_time": "1575495600000"}]
```

We are reading a stream of events coming out of Kafka on the left, then we strip out the tracks for an individual Bird

