

Events to the rescue: solving distributed data problems in a microservice architecture

Chris Richardson

Microservice architecture consultant and trainer

Founder of Eventuate.io

Founder of the original CloudFoundry.com

Author of POJOs in Action and Microservices Patterns

• @crichardson

chris@chrisrichardson.net

<http://adopt.microservices.io>



To deliver a large complex application rapidly, frequently and reliably, you often must use the microservice architecture. The microservice architecture is an architectural style that structures the application as a collection of loosely coupled services. One challenge with using microservices is that in order to be loosely coupled each service has its own private database. As a result, implementing transactions and queries that span services is no longer straightforward.

How to use events to implement transactions and queries in a microservice architecture

In this EDA Summit presentation, you will learn how event-driven microservices address this challenge. Chris Richardson, creator of microservices.io, describes how to use sagas, which is an asynchronous messaging-based pattern, to implement transactions that span services. You will learn how to implement queries that span services using the CQRS pattern, which maintain easily queryable replicas using events.

- Why loosely coupled microservices?
- Distributed data challenges in a microservice architecture
- Transactions in a microservice architecture
- Querying in a microservice architecture

Microservice architecture = architectural style

`createCustomer(creditLimit)`

`createOrder(customerId, orderTotal)`

`findOrdersForCustomer(customerId)`

`findRecentCustomers()`

REST API

Small team

Order team

Online store

API
Gateway

Customer
Service

Order
Service

<15 minute lead time

Customer team

Loosely coupled

Independently deployable

Why microservices: success triangle

Process: Lean + DevOps/Continuous Delivery & Deployment



ACCELERATE



Businesses must be
nimble, agile and
innovate faster

IT must deliver software
rapidly, frequently, reliably and sustainably:

- Deployment frequency: \geq once per developer/day
- < 15 minute lead time
- ...

Enables:
Testability
Deployability

Organization:
Network of small,
loosely coupled, teams

Enables:
Loose
coupling

Architecture:
Microservices
(sometimes)@crichardson

Loose coupling is essential

Services collaborate, e.g. Order Service must reserve customer credit

⇒

Some coupling - degree of connectedness - is inevitable

BUT

You must design services to be loosely coupled



Runtime coupling

=

Both services must be available

⇒

Reduced availability



Design-time coupling

=

lockstep changes

⇒

Reduced productivity

Infrastructure coupling

=

Competing for same resources

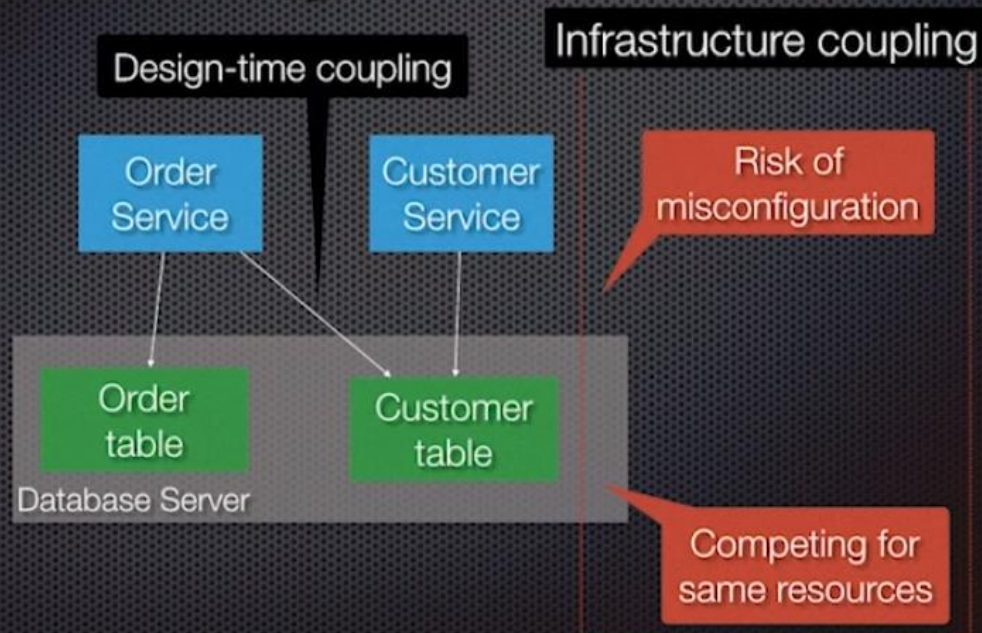
Misconfiguration bugs

⇒

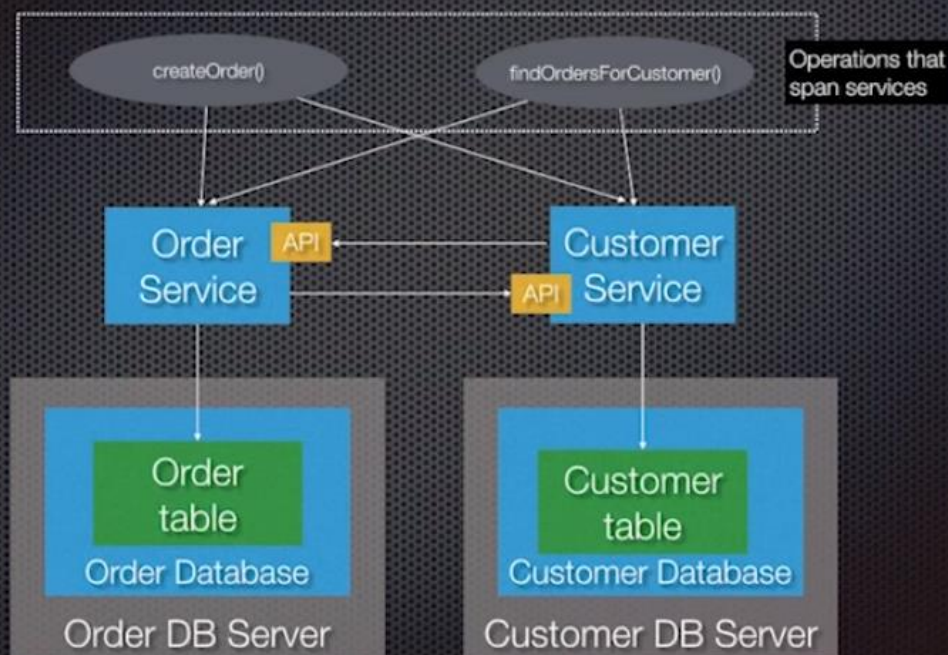
Reduced availability

- Distributed data challenges in a microservice architecture

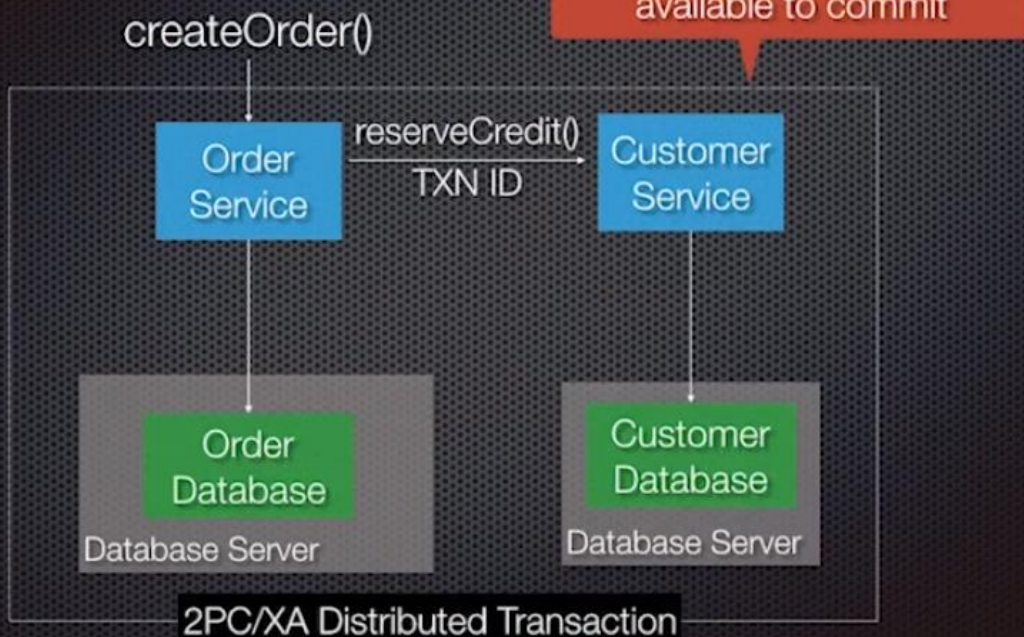
Sharing database tables = tight coupling



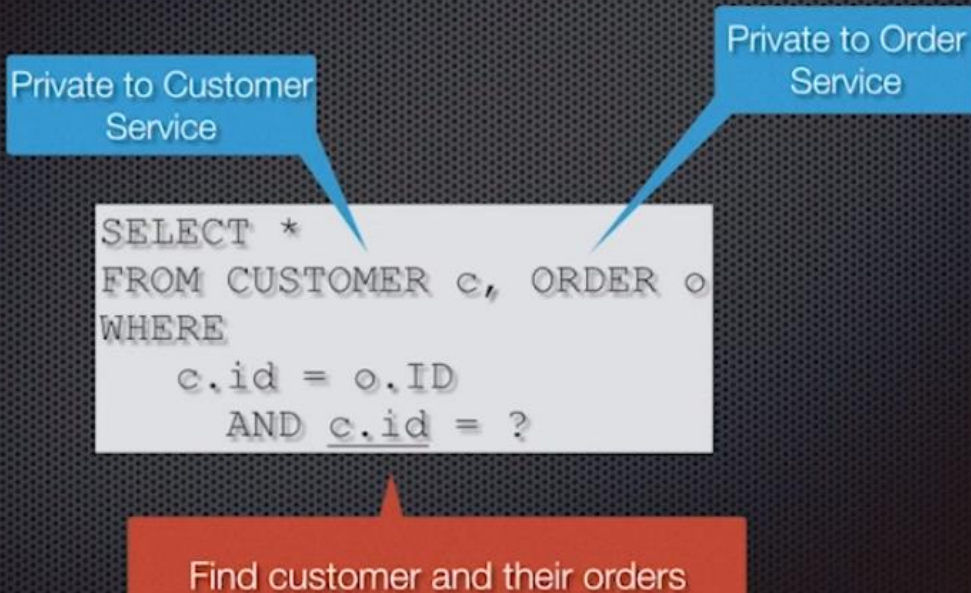
Use a Database-per service



Traditional distributed transactions = runtime coupling



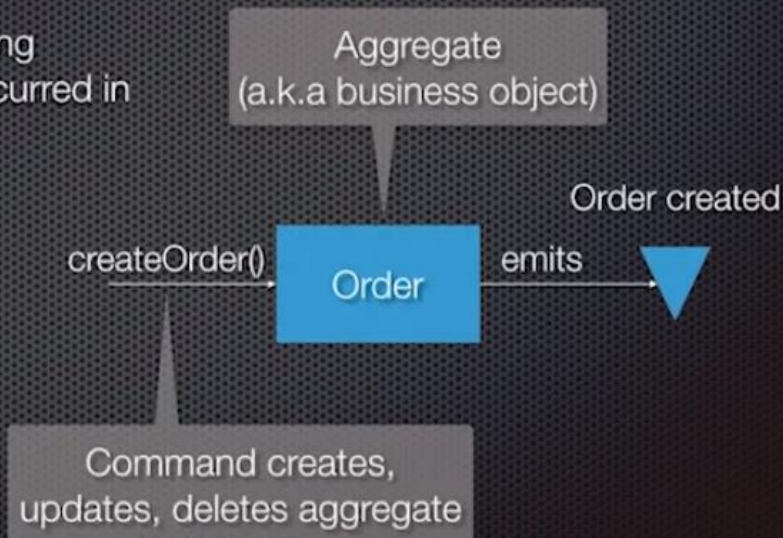
Queries that do cross-database joins violate encapsulation



Events to the rescue

An event is something notable that has occurred in a domain, e.g.


- Order Created
- Order Confirmed
- Order Cancelled
- ..



An event is a type of message

Channel types:

- Point-to-point - deliver to one recipient
- Publish-subscribe - deliver to all recipients

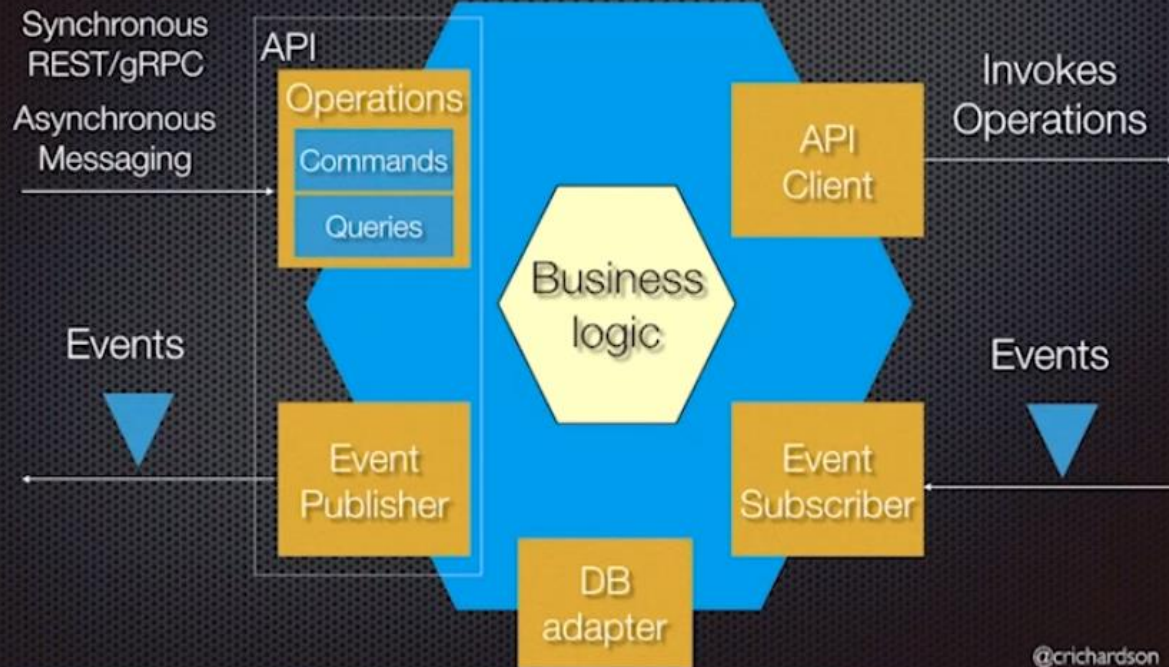


Message types:

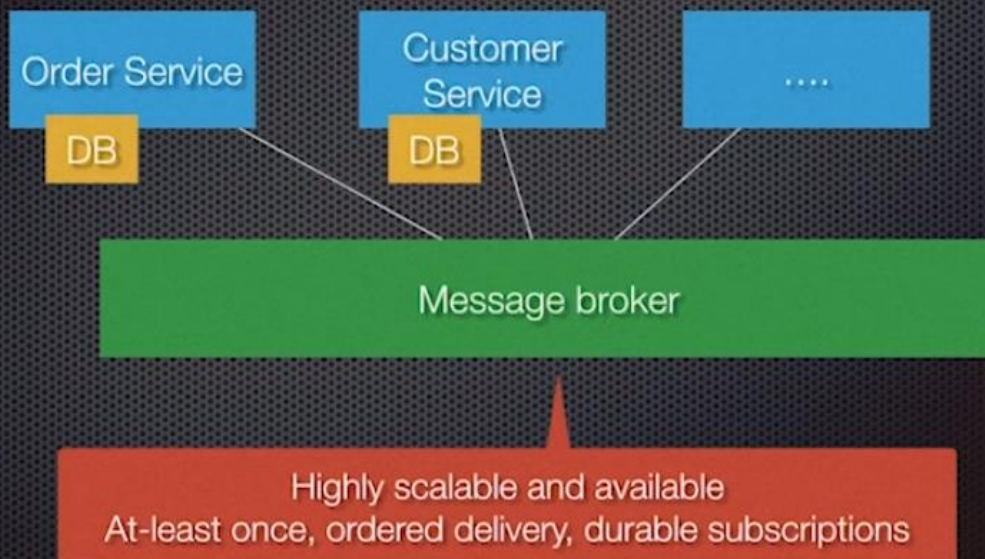
- Message - generic
- **Event - something that has happened**
- Command - a request to perform an operation
- Reply - a reply to a command
- ...




Services exchange events



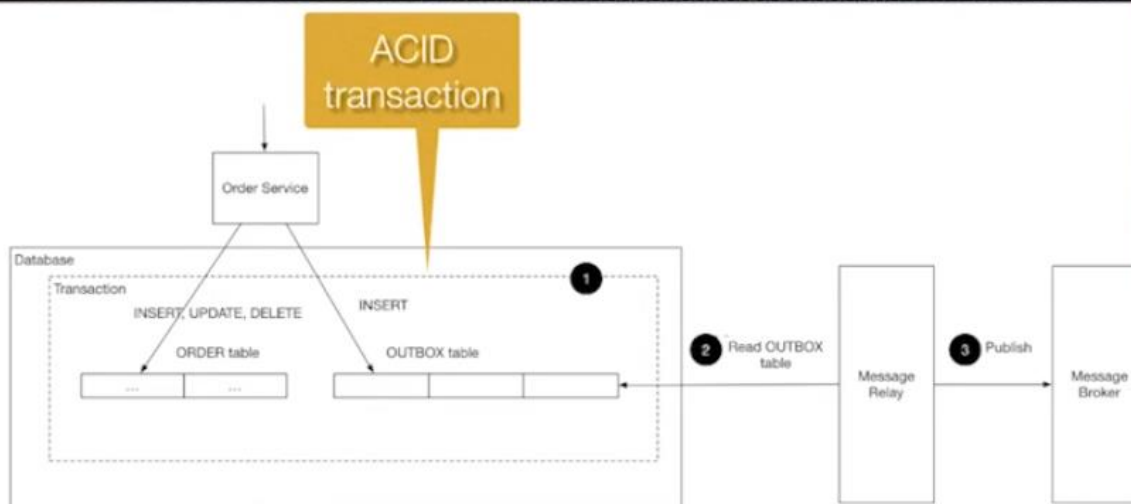
Use asynchronous, broker-based messaging



Reliable messaging requires atomicity



Atomically updating state and publishing events



<https://microservices.io/patterns/data/transactional-outbox.html>

- Transactions in a microservice architecture

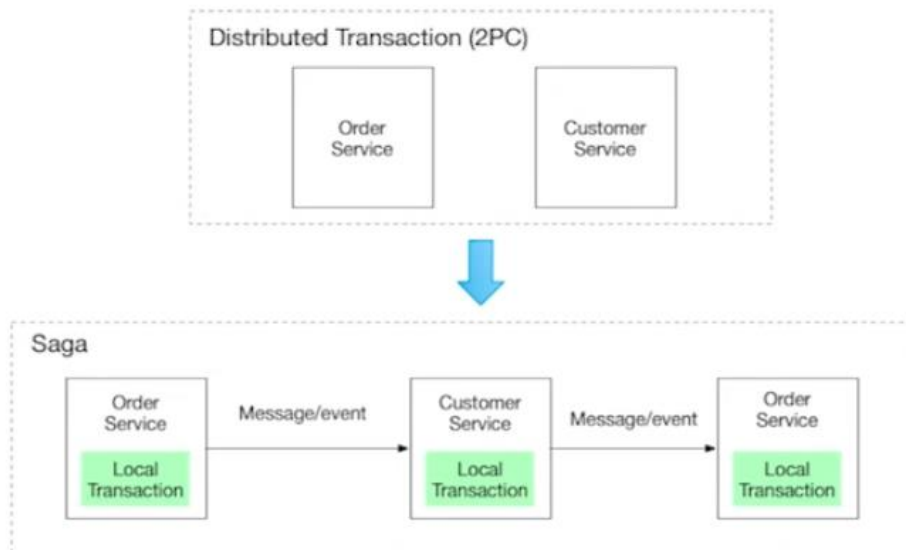
From a 1987 paper

SAGAS

Hector Garcia-Molina
Kenneth Salem

Department of Computer Science
Princeton University
Princeton, N J 08544

Sagas: event-based transactions



<https://microservices.io/patterns/data/saga.html>

Create Order Saga

createOrder()

Initiates saga

Order Service

@crichardson

Create Order Saga

createOrder()

Initiates saga

Order Service

Local transaction

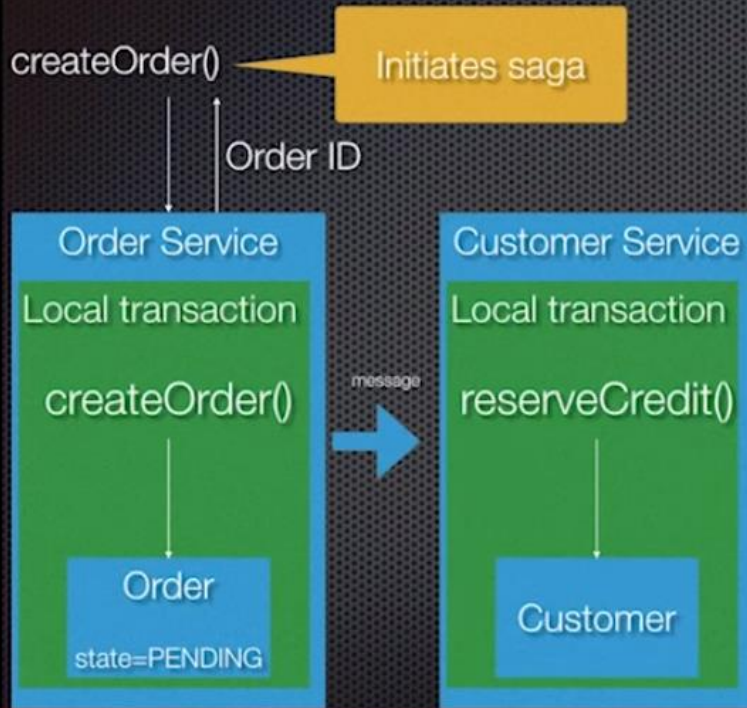
createOrder()

Order

state=PENDING

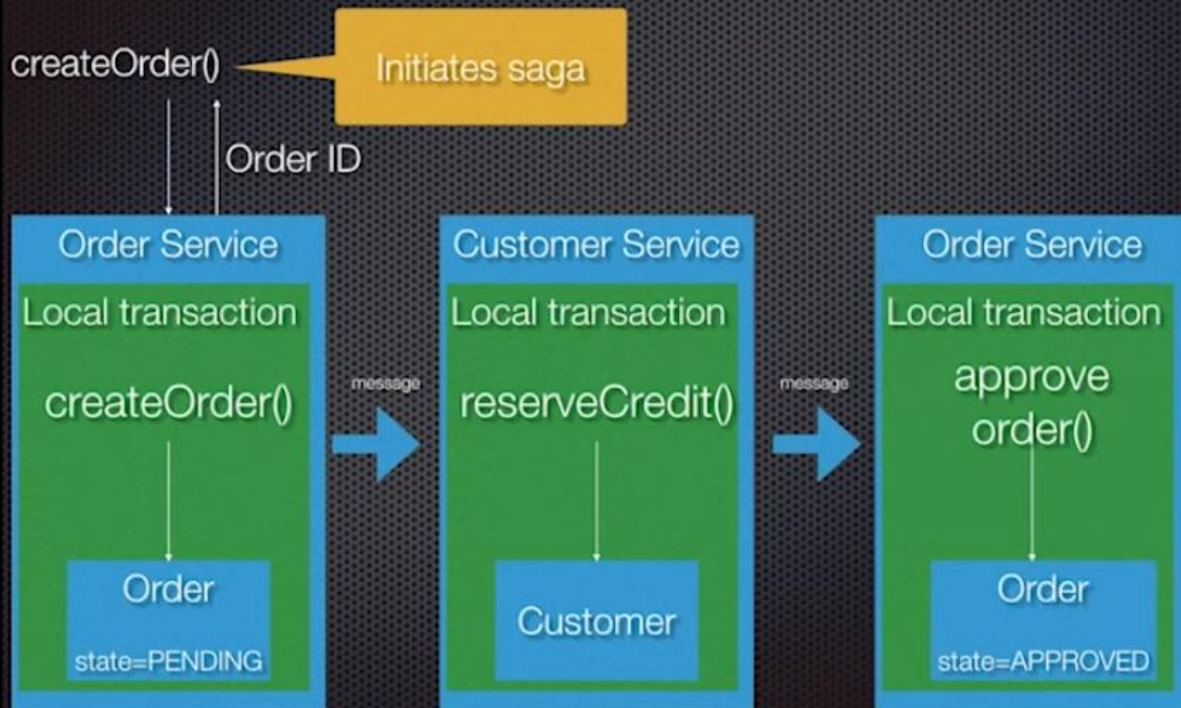
@crichardson

Create Order Saga

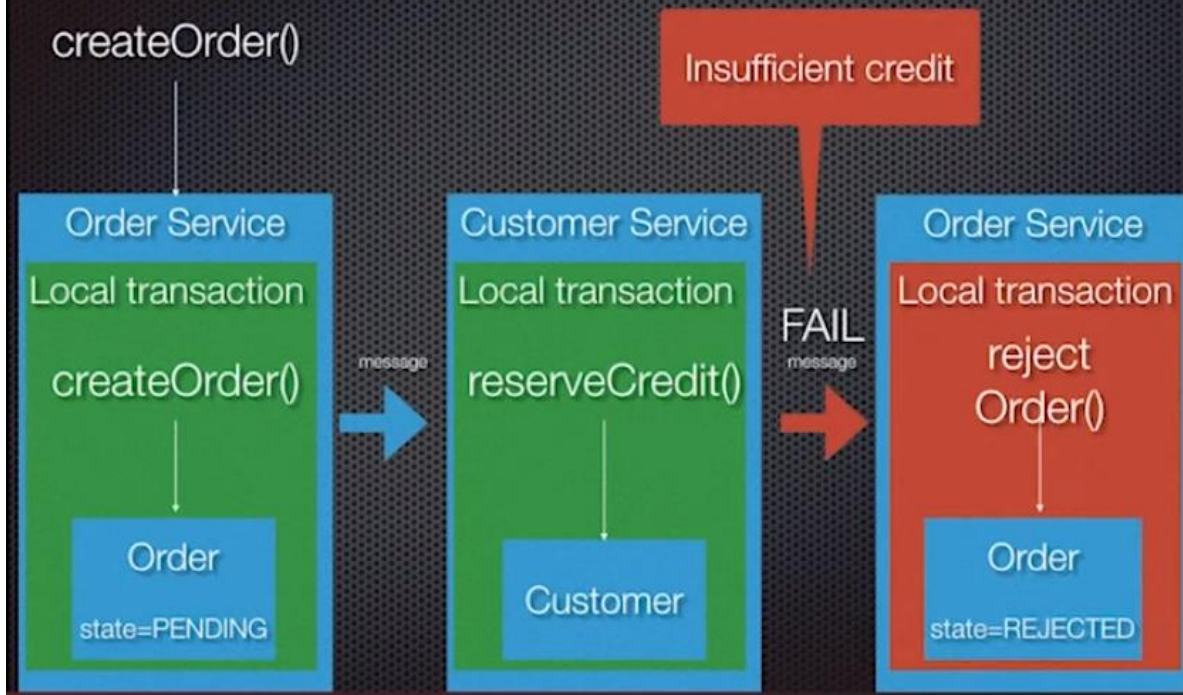


@crichardson

Create Order Saga




Create Order Saga - “rollback” using a **compensating** transaction



How to sequence the steps of a saga?

- After the completion of transaction T_i/C_i “something” must decide what step to execute next
- Success: $T_{(i+1)}$
- Failure: $C_{(i-1)}$



Choreography: **distributed** decision making

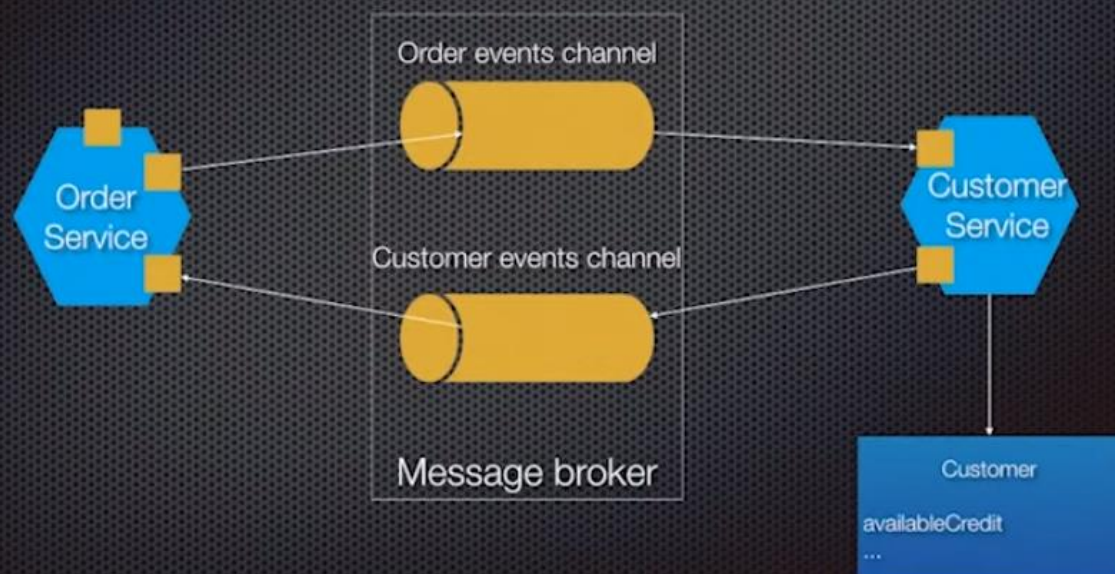
VS.

Orchestration: **centralized** decision making


Choreography = event-driven sagas

- Coordination logic = code that publishes events + event handlers
- Whenever a saga participant updates a business object, it publishes (domain) events announcing what it **has done**
- Saga participants have event handlers that update business objects => ...

Choreography-based Create Order Saga



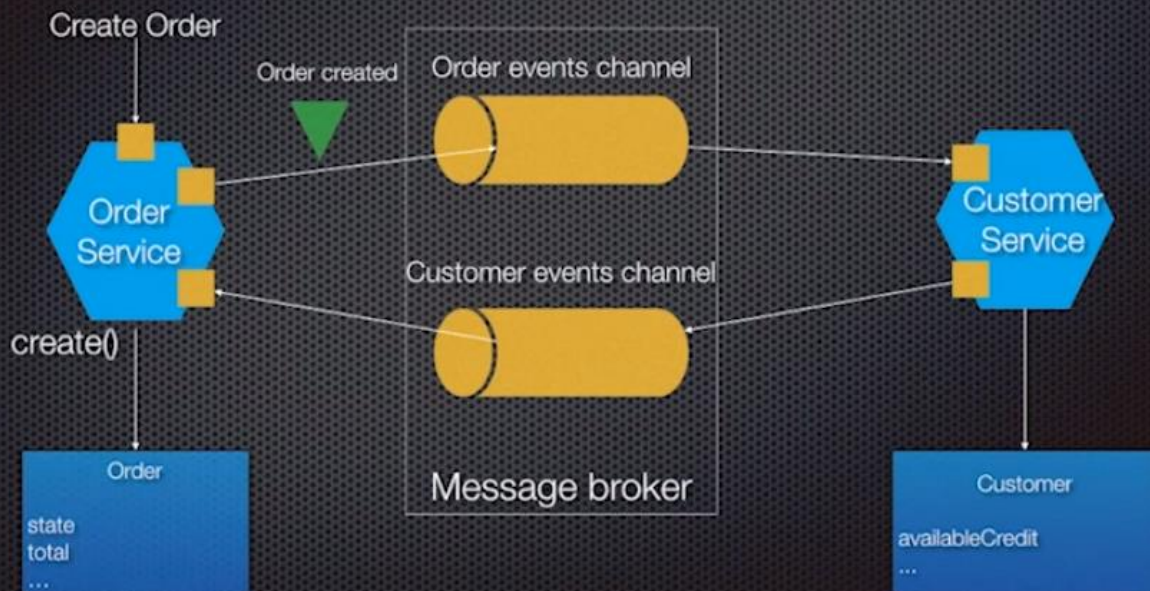
Choreography-based Create Order Saga



<https://github.com/eventuate-tram/eventuate-tram-examples-customers-and-orders>

@crichardson

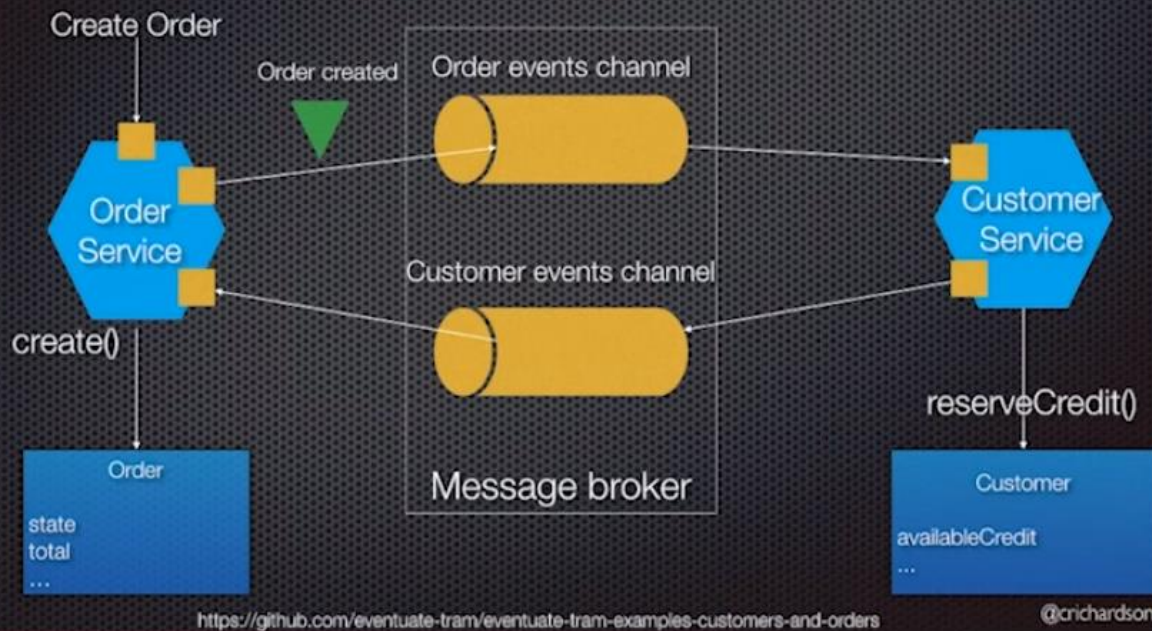
Choreography-based Create Order Saga



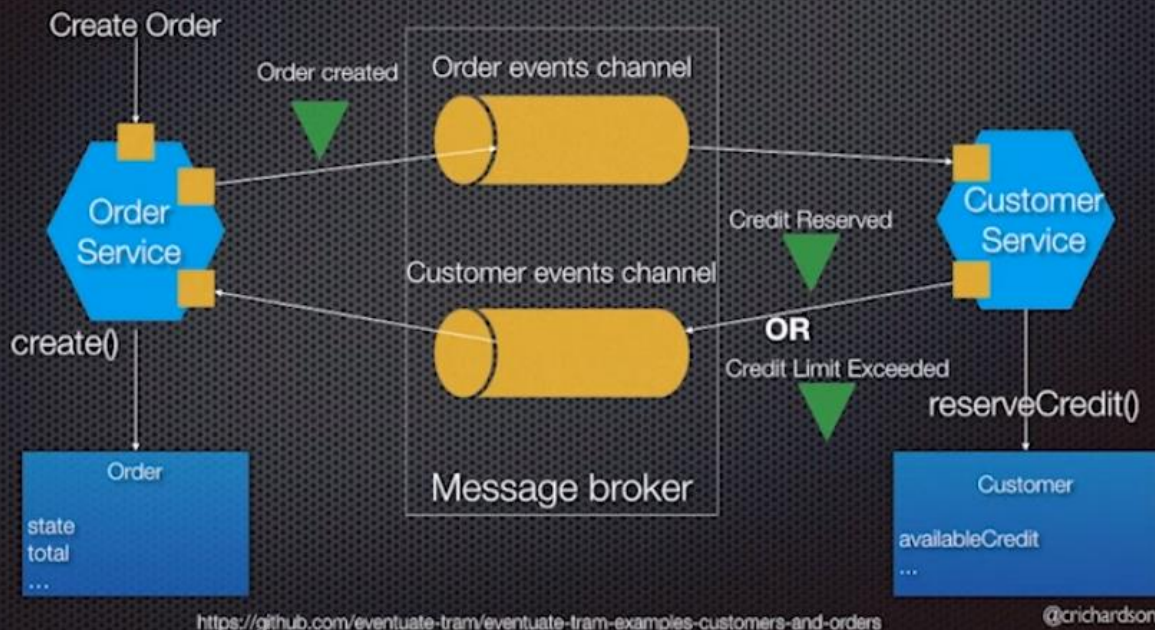
<https://github.com/eventuate-tram/eventuate-tram-examples-customers-and-orders>

@crichardson

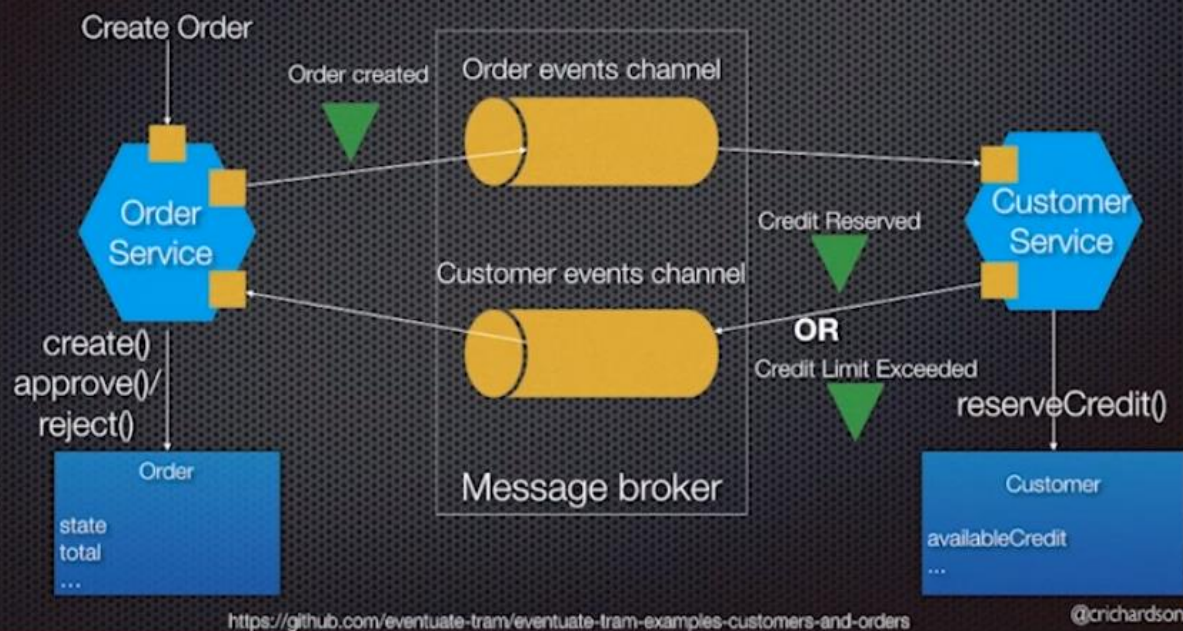
Choreography-based Create Order Saga



Choreography-based Create Order Saga

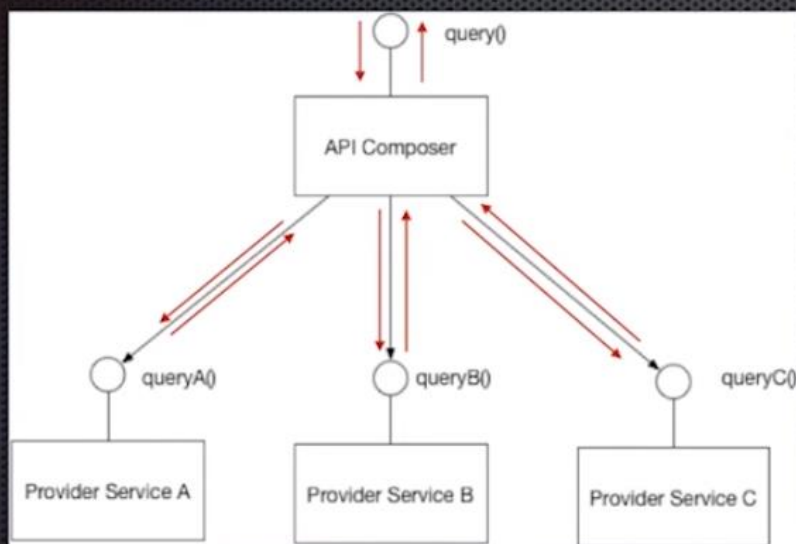


Choreography-based Create Order Saga



- Querying in a microservice architecture

Querying using the API Composition pattern

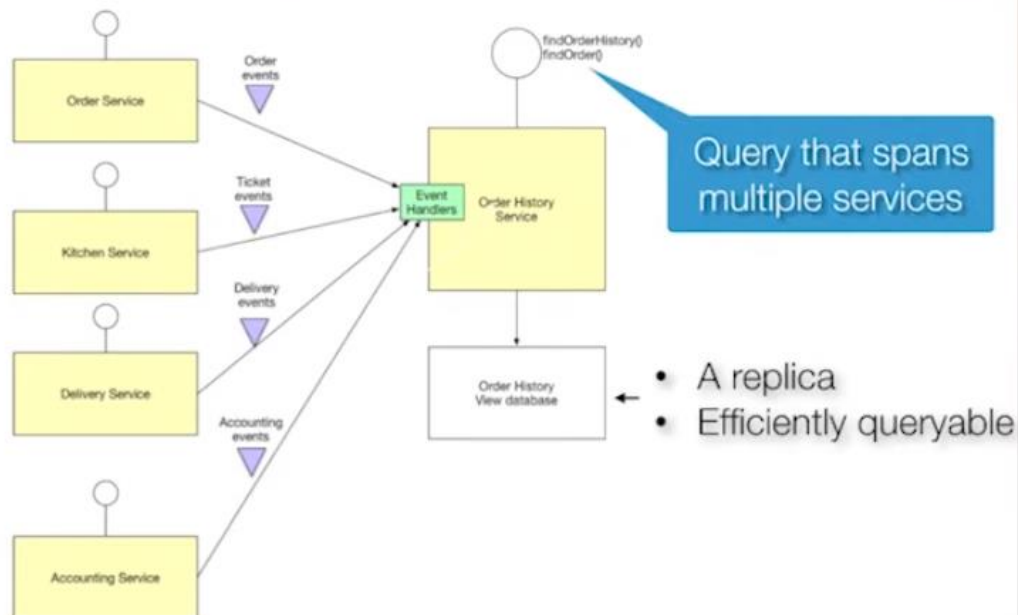


Simple but involves runtime coupling and is sometimes too inefficient

<https://microservices.io/patterns/data/api-composition.html>

Using circuit breakers and callback strategies in your microservices can help limit the defects of this approach.

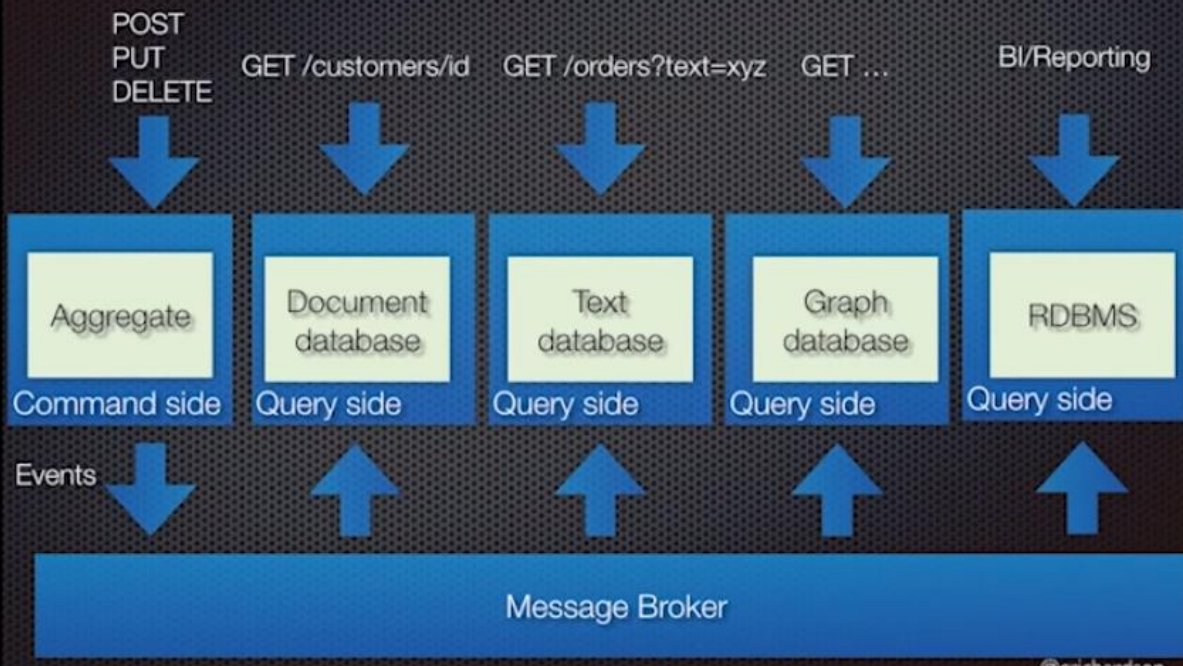
CQRS: event-based queries



<https://microservices.io/patterns/data/cqrs.html>

For CQRS, you will need separate data models for the querying and updates

Pick a database that efficiently supports the queries



Summary

- Rapid, frequent, reliable and sustainable software delivery requires services to be loosely coupled
- Sharing databases = tight design-time coupling \Rightarrow Database per service
- Traditional distributed transactions = tight runtime coupling
- Use the Saga pattern to implement commands that span services
- Use the CQRS pattern to implement queries that span services